

DISJUNCTIVISM: PERCEPTION, ACTION, KNOWLEDGE

Disjunctivism has attracted considerable philosophical attention in recent years: it has been the source of a lively and extended debate spanning the philosophy of perception, epistemology, and the philosophy of action. Adrian Haddock and Fiona Macpherson present seventeen specially written essays, which examine the different forms of disjunctivism and explore the connections between them. This volume will be an essential resource for anyone working in the central areas of philosophy, and the starting point for future research in this fascinating field.

Adrian Haddock is Lecturer in Philosophy at the University of Stirling.

Fiona Macpherson is Senior Lecturer in Philosophy at the University of Glasgow.

Disjunctivism: Perception, Action, Knowledge

Edited by

ADRIAN HADDOCK AND FIONA MACPHERSON

OXFORD
UNIVERSITY PRESS

OXFORD
UNIVERSITY PRESS

Great Clarendon Street, Oxford OX2 6DP

Oxford University Press is a department of the University of Oxford.
It furthers the University's objective of excellence in research, scholarship,
and education by publishing worldwide in

Oxford New York

Auckland Cape Town Dar es Salaam Hong Kong Karachi
Kuala Lumpur Madrid Melbourne Mexico City Nairobi
New Delhi Shanghai Taipei Toronto

With offices in

Argentina Austria Brazil Chile Czech Republic France Greece
Guatemala Hungary Italy Japan Poland Portugal Singapore
South Korea Switzerland Thailand Turkey Ukraine Vietnam

Oxford is a registered trade mark of Oxford University Press
in the UK and in certain other countries

Published in the United States
by Oxford University Press Inc., New York

© The several contributors 2008

The moral rights of the author have been asserted
Database right Oxford University Press (maker)

First published 2008
First published in paperback 2011

All rights reserved. No part of this publication may be reproduced,
stored in a retrieval system, or transmitted, in any form or by any means,
without the prior permission in writing of Oxford University Press,
or as expressly permitted by law, or under terms agreed with the appropriate
reprographics rights organization. Enquiries concerning reproduction
outside the scope of the above should be sent to the Rights Department,
Oxford University Press, at the address above

You must not circulate this book in any other binding or cover
and you must impose this same condition on any acquirer

British Library Cataloguing in Publication Data

Data available

Library of Congress Cataloging in Publication Data

Data available

Typeset by Laserwords Private Limited, Chennai, India

Printed in the United Kingdom by
Lightning Source UK Ltd., Milton Keynes

ISBN 978-0-19-923154-6 (Hbk.)

ISBN 978-0-19-969308-5 (Pbk.)

Preface and Acknowledgements

Many of the contributions to this volume began their life as papers presented to a conference held at the University of Glasgow on 4th and 5th June 2005, entitled 'Disjunctivism: Perception, Action, Knowledge'. This conference brought together theorists of perception, action, and knowledge from the UK, US, and Canada, to discuss the impact of disjunctivism on their respective areas of interest. We would like to thank the Royal Institute of Philosophy, the Mind Association, the Analysis Trust, Blackwell Publishing, and *The Philosophical Quarterly* for supporting this conference. We would also like to thank the conference speakers and commentators: Alex Byrne, Jonathan Dancy, Jennifer Hornsby, Heather Logue, Mike Martin, Mohan Matthen, Alan Millar, Ram Neta, Matthew Nudds, Richard Price, Duncan Pritchard, Sonia Sedivy, and Finn Spicer.

The conference was held under the auspices of The Centre for the Study of Perceptual Experience, which is based in the Department of Philosophy at the University of Glasgow. The primary aim of the Centre is to conduct and facilitate analytical philosophical research into the nature of perceptual experience. A secondary aim is to facilitate communication and collaboration between researchers in philosophy and other disciplines whose research remit includes perceptual experience. At the time of writing, in its two-and-a-half-year history, the Centre has organized four international conferences: 'The Individuation of the Senses', 'Disjunctivism: Perception, Action, Knowledge', 'Graduate Interdisciplinary Conference on Perception', and 'The Admissible Contents of Experience', in addition to numerous smaller workshops and events.

We would like to thank all of the contributors to this volume for sending us their essays in good time, and always responding to our numerous requests. We are also, of course, very grateful to them for producing such an excellent, thought-provoking, series of essays.

We would also like to thank Professor Luis M. Valdés-Vilanova for granting us permission to reproduce John McDowell's 2006 article 'The Disjunctive Conception of Experience as Material for a Transcendental Argument' from the journal *Teorema*, volume 25.

Finally, we would like to thank the following people for their help and advice: Michael Brady, Jane Calvert, Antony Duff, Jennifer Hornsby, John McDowell, Alan Millar, Dimitris Platchias, Duncan Pritchard, Matthew Soteriou, Peter Sullivan, and, above all, Susanna Siegel. We both owe much thanks to Peter Momtchiloff, and his staff at Oxford University Press, for their efficiency and encouragement in producing the book, and bringing it to press.

AH

FM

Glasgow, March 2007

Contents

| | |
|--|----|
| Introduction: Varieties of Disjunctivism <i>Adrian Haddock and Fiona Macpherson</i> | 1 |
| Analytical Table of Contents | 25 |

PART I: PERCEPTION

| | |
|--|-----|
| 1. Hinton and the Origins of Disjunctivism <i>Paul Snowdon</i> | 35 |
| 2. Either/Or <i>Alex Byrne and Heather Logue</i> | 57 |
| 3. Against Disjunctivism <i>E. J. Lowe</i> | 95 |
| 4. Disjunctivism About Visual Experience <i>Scott Sturgeon</i> | 112 |
| 5. Disjunctivism, Indistinguishability, and the Nature of Hallucination <i>William Fish</i> | 144 |
| 6. How to Account for Illusion <i>Bill Brewer</i> | 168 |
| 7. Disjunctivism and Discriminability <i>A. D. Smith</i> | 181 |
| 8. The Epistemic Conception of Hallucination <i>Susanna Siegel</i> | 205 |

PART II: ACTION

| | |
|---|-----|
| 9. Disjunctive Theories of Perception and Action <i>David-Hillel Ruben</i> | 227 |
| 10. A Disjunctive Conception of Acting for Reasons <i>Jennifer Hornsby</i> | 244 |
| 11. On How to Act—Disjunctively <i>Jonathan Dancy</i> | 262 |

PART III: KNOWLEDGE

| | |
|---|-----|
| 12. McDowellian Neo-Mooreanism <i>Duncan Pritchard</i> | 283 |
| 13. In Defence of Disjunctivism <i>Ram Neta</i> | 311 |
| 14. Perceptual-Recognitional Abilities and Perceptual Knowledge <i>Alan Millar</i> | 330 |
| 15. Starting Afresh Disjunctively: Perceptual Engagement with the World <i>Sonia Sedivy</i> | 348 |
| 16. The Disjunctive Conception of Experience as Material for a Transcendental Argument <i>John McDowell</i> | 376 |
| 17. Comment on John McDowell's 'The Disjunctive Conception of Experience as Material for a Transcendental Argument' <i>Crispin Wright</i> | 390 |
| <i>Index</i> | 405 |

Contributors

Bill Brewer is Professor of Philosophy at the University of Warwick. He is author of *Perception and Reason* (Oxford University Press, 1999), and various papers in philosophy of mind, metaphysics, and epistemology. He is working on a second book for Oxford University Press, *Perception and its Objects*, which is about the natures of our perceptual experience and the physical objects which we perceive.

Alex Byrne is Professor of Philosophy at MIT. He has written on a variety of topics in philosophy of mind, philosophy of language, and metaphysics and epistemology. He is currently writing a book on self-knowledge.

Jonathan Dancy taught at Keele University for twenty-five years. He now teaches at the University of Reading, UK, and at the University of Texas at Austin. His main interests lie in moral theory, the general theory of reasons, and the philosophy of action.

William Fish is Lecturer in Philosophy at Massey University, New Zealand. He has published on disjunctivism and other topics in the philosophy of mind. He is currently in the process of completing a book providing a more comprehensive defence of a disjunctive naive realism.

Adrian Haddock is Lecturer in Philosophy at the University of Stirling. His research is primarily in the philosophy of mind and action.

Jennifer Hornsby is Professor of Philosophy, Birkbeck College, University of London, and a Co-director of the Rational Agency project in the Centre for the Study of Mind in Nature, Oslo. Her research is mainly in philosophy of mind and action.

Heather Logue is a graduate student in the Department of Philosophy at MIT. She is currently working on her dissertation, which explores various issues in the philosophy of perception.

E. J. Lowe is Professor of Philosophy at Durham University. He specializes in metaphysics, the philosophy of mind and action, philosophical logic, the philosophy of language, and the philosophy of John Locke. His recent books include *Locke* (Routledge, 2005) and *The Four-Category Ontology* (Oxford University Press, 2006).

John McDowell is a University Professor of Philosophy at the University of Pittsburgh. Before moving to Pittsburgh in 1986, he taught at University College, Oxford. He has held visiting appointments at Harvard University, the University of Michigan, UCLA, and Princeton University. He was the John Locke Lecturer at Oxford University in 1991. His major interests are Greek philosophy, philosophy of language, philosophy of mind, metaphysics and epistemology, and ethics. He is

a fellow of the British Academy and a fellow of the American Academy of Arts and Sciences.

Fiona Macpherson is Lecturer in Philosophy at the University of Glasgow and Director of the Centre for the Study of Perceptual Experience. She was recently Research Fellow at the Centre for Consciousness at the Australian National University. She has research interests in philosophy of mind, psychology, and perception.

Alan Millar, FRSE, studied at the Universities of Edinburgh and Cambridge. He has taught philosophy for many years, and is now a Professor of Philosophy at the University of Stirling. He is the author of *Reasons and Experience* (Oxford University Press, 1991), *Understanding People: Normativity and Rationalizing Explanation* (Oxford University Press, 2004), and articles mainly in the philosophy of mind and the theory of knowledge.

Ram Neta received his PhD in philosophy from the University of Pittsburgh, and then taught philosophy at the University of Utah. He is currently Assistant Professor of Philosophy at the University of North Carolina at Chapel Hill, and his research is primarily in epistemology.

Duncan Pritchard is a Professor of Philosophy at the University of Edinburgh. His work is mainly in the area of epistemology, in which he has published widely. His books include *Epistemic Luck* (Oxford University Press, 2005) and *What is this Thing Called Knowledge?* (Routledge, 2006).

David-Hillel Ruben is Professor of Philosophy at Birkbeck, University of London, and Director of NYU's London campus. His research is in the philosophy of social science and metaphysics. He has published extensively in both areas. His recent books include *Explaining Explanation* (Routledge, 1990) and *Action and its Explanation* (Oxford University Press 2003).

Sonia Sedivy received her PhD in philosophy from the University of Pittsburgh. Since then she has taught philosophy at the University of Toronto. Her research is primarily in the philosophies of mind, perception, and art.

Susanna Siegel is Professor of Philosophy at Harvard University. She has published papers on disjunctivism, the contents of experience, the phenomenology of agency, and the semantics of demonstrative expressions.

A. D. Smith is Professor of Philosophy at the University of Warwick. He works and publishes in a number of areas of philosophy, including the philosophy of mind and action, metaphysics, philosophical logic and the philosophy of language, the history of philosophy, and phenomenology. His recent books include *The Problem of Perception* (Harvard University Press, 2002) and *Husserl and the Cartesian Meditations* (Routledge, 2003).

Paul Snowdon has been Grote Professor of Mind and Logic in the Department of Philosophy at University College London since 2001. Before that he was, for thirty

years, a Fellow and Lecturer in Philosophy at Exeter College, Oxford. His publications include articles about the philosophy of mind, perception, and personal identity.

Scott Sturgeon is a University Lecturer in Philosophy at Oxford University and a Tutorial Fellow of Wadham College. He has helped edit the *British Journal for the Philosophy of Science*, *Mind*, and *Philosophical Studies*. He has authored *Matters of Mind* (Routledge 2001) and he is currently writing a book on epistemic norms—usefully titled *Epistemic Norms*—for Oxford University Press.

Crispin Wright, FBA, FRSE, is Wardlaw Professor at the University of St Andrews, where he directs the philosophical research centre, *Arché*, and Global Distinguished Professor of Philosophy at New York University.

Introduction: Varieties of Disjunctivism

Adrian Haddock and Fiona Macpherson

Inspired by the writings of J. M. Hinton (1967a, 1967b, 1973), but ushered into the mainstream by Paul Snowdon (1980–1, 1990–1), John McDowell (1982, 1986), and M. G. F. Martin (2002, 2004, 2006), disjunctivism is currently discussed, advocated, and opposed in the philosophy of perception, the theory of knowledge, the theory of practical reason, and the philosophy of action. But what is disjunctivism?

A good way to answer this question is to consider the conceptions of experience advanced by Hinton, Snowdon, Martin, and McDowell. Snowdon's contribution to this volume offers an excellent introduction to Hinton's work. So, in this introduction, we will concentrate on the more well-known, and more influential, views of Snowdon, McDowell, and Martin. As we will see, these views have a number of features in common. But, as we will also see, these commonalities must not be allowed to obscure the equally important differences. In fact, the views of Snowdon, McDowell, and Martin serve to exemplify three distinct varieties of disjunctivism.

It is not unusual, and not always unjustified, to speak of a position called 'disjunctivism', and to refer to each of the philosophers mentioned above as its advocates. But doing so carries the danger of eliding important differences and engaging in unjust criticism. To mention only three possibilities: disjunctivism about the nature of experience may be taken to task for failing to establish its epistemological advantages over alternative positions, when in fact it was never intended to have any such advantages; disjunctivism about the epistemic warrant that experience can provide will be attacked on the grounds that it is compatible with a causal theory of perception, when in fact it was never intended to oppose such a theory; and disjunctivism about experience's phenomenal character may be criticized for failing to undermine Cartesian scepticism, when in fact it had no such aim.

It is because of the seriousness of these misunderstandings, and the apparent ease of falling into them, that this introduction takes the form that it does. Rather than offering a summary of each of the essays that appear in the volume, it presents in detail some key essays by Martin, McDowell, Snowdon, and others. In so doing it elucidates, compares, and provides a much-needed taxonomy of the ideas that are most often discussed under the disjunctivist heading, and makes clear which of the

Many thanks are due to Jennifer Hornsby, John McDowell, Alan Millar, and Susanna Siegel. We would like to thank Susanna in particular, for her numerous helpful suggestions and comments.

essays that appear in the volume discuss which of these ideas. We will discuss not only established disjunctive accounts in the theories of perception and knowledge, but also emerging disjunctive accounts in the theory of action. Summaries of the volume's essays, grouped thematically, and together providing an overview of the entire volume, can be found in the 'Analytical Table of Contents' that succeeds this introduction.

1 PERCEPTION: SNOWDON AND EXPERIENTIAL DISJUNCTIVISM

Snowdon is interested in the nature of visual perception, and the analysis of its concept, and wishes to argue against the causal theory of visual perception—advocated by H. P. Grice (1961), P. F. Strawson (1979), and many others—and *inter alia* makes use of (what he calls) a disjunctive theory in order to do so.

The causal theory comes in two versions. In the first, it concerns the nature of perception, and consists in two claims—(1) “the causal thesis” and (2) “the effect thesis”:

- (1) The causal thesis: If a subject *S* sees a public object *o* then *o* causally affects *S* (in the appropriate way).
- (2) The effect thesis: If a subject *S* sees a public object *o* then *o* produces in *S* an L-state—namely, “a state reportable in a sentence beginning ‘It looks to *S* as if . . .’ where these words are interpreted phenomenologically (rather than ascribing, say, a tentative judgement by *S*)” (Snowdon 1980–1: 176).¹

The second version also concerns the concept of visual perception, and appends the following claim:

- (3) The causal thesis and the effect thesis are requirements of our ordinary concept of vision.

Grice and Strawson's main argument for both versions of the causal theory begins by pointing to cases in which the following conditions are satisfied: “(i) *S* is in an L-state appropriate to seeing *o*; (ii) *o* is in [*S*'s] environment; (iii) the L-state is not causally dependent on *o*; and (iv) *o* is not seen” (Snowdon 1980–1: 181). For example, it looks to *S* as if there is a pig in front of him; there is a pig in front of him; the pig is not causally responsible for its looking to *S* as if there is a pig in front of him; and *S* does not see the pig. The argument concludes that the causal theory is true.

¹ A further requirement, which Snowdon does not mention, is that the *o* and the L-state ‘suitably match’. What this requirement amounts to is a matter of some controversy, and will depend upon one's theory of perceptual experience. For example, if one believes that experiences are representational states one obvious way to spell out the relevant notion of matching is to say that it occurs when the L-state represents the *o* or a similar object. There is more on this conception of experience as a representational state in section 3 below.

However, Snowdon claims that the possibility of such cases does not entail the impossibility of a case in which (i) to (iii) are satisfied but *o* is seen (by S). It may be that it looks to S as if there is a pig in front of him; there is a pig in front of him; the pig is not causally responsible for its looking to S as if there is a pig in front of him; but S does see the pig—seeing objects does not depend on there being causality between L-states and the object seen. Part of the point of Snowdon’s disjunctive theory is to establish the possibility of this case.

According to the disjunctive theory, the explanation of S’s failure to see *o* is not that the L-state is not caused by the object seen but (simply) that S is in the *wrong kind* of L-state; and, in the same way, the explanation of S’s success in seeing *o* would be (simply) that S is in the right kind of L-state. The theory offers the following account of what it is to be in an L-state: it is *either* to be in an L-state that is “intrinsically independent of surrounding objects” *or* to be in an L-state that intrinsically “involves the surrounding objects” (Snowdon 1980–1: 186). If S is in an L-state of the latter kind then they do see objects (and the objects they see are the very objects the state ‘involves’), and if they are in an L-state of the former kind then they do not. Snowdon (1980–1: 185) expresses this theory by means of the following formula:

[D] It looks to S as if there is an F; (there is something which looks to S to be an F) \vee (it is to S as if there is something which looks to S to be an F).²

Here the sentence in front of the semicolon reports the L-state to be explained, and the sentences that form the disjunction report the more fundamental states in which they consist.

Snowdon (1980–1) does not attempt to establish the explanatory primacy of the disjunctive theory relative to the causal theory. Instead he makes two conditional claims. First, if the disjunctive theory is true, then both versions of the causal theory are false. Second, if the disjunctive theory is false, and the reasons why it is false appeal to data of a certain sort—namely, data that are not “relatively immediately acknowledgeable by any person, whatever their education, who can count as having the concept in question” (Snowdon 1980–1: 185)—then the second version of the causal theory is also false. The central burden of Snowdon’s (1980–1) essay is to show that the first claim is true. He does not attempt to establish its antecedent.³

The crucial difference between the causal theory and the disjunctive theory is their contrasting accounts of the *intrinsic nature* of L-states. According to the disjunctive theory, it is not possible to characterize the intrinsic nature of certain L-states without mentioning objects seen, because it is part of the intrinsic nature of these L-states to be states of seeing objects. By contrast, the causal theory assumes that it is always possible to characterize the intrinsic nature of L-states without mentioning the object seen, because it is no part of the intrinsic nature of any L-state to be a state of seeing an

² We might wonder why Snowdon says [D] and not [D*]: It looks to S as if there is an F; (S sees an F) \vee (it is to S as if there is something which looks to S to be an F). Snowdon’s objection to [D*] is that it is possible for S to see an F without it looking to S as if there is one. In his (1980–1) he took this to be an objection to Hinton; but some recanting occurs in his essay for this volume.

³ Attempts to do so can be found in both Snowdon (1990–1) and Martin (2002).

object. It is this assumption which ensures that, if the disjunctive theory is true, then the causal theory is false.

We will call disjunctivism, of the variety Snowdon advocates, ‘experiential disjunctivism’. Issues pertaining to experiential disjunctivism are taken up in the essays in this volume by E. J. Lowe, Alan Millar, and Snowdon himself. Lowe defends the causal theory of perception against experiential disjunctivism; Millar addresses some arguments for experiential disjunctivism; and Snowdon offers an extensive discussion, and assessment, of Hinton’s writings—widely regarded as disjunctivism’s founding texts.

In the next section, we will examine disjunctivism of a rather different variety.

2 KNOWLEDGE: McDOWELL AND EPISTEMOLOGICAL DISJUNCTIVISM

McDowell presents his disjunctive conception of appearances in two essays—‘Criteria, Defeasibility, and Knowledge’ (1982), and ‘Singular Thought and the Extent of Inner Space’ (1986). In both, his concerns are epistemological and transcendental: he is concerned with the nature of perceptual knowledge, and with the very possibility of states that are directed towards objects.⁴ In the first of these essays, McDowell advocates a position that, following Alex Byrne and Heather Logue (this volume), we will call ‘epistemological disjunctivism’. In the second essay, his position is somewhat different, and sometimes seems closer to experiential disjunctivism. The exegetical issues are complex here, and for this reason we will discuss both essays in detail.

2.1 ‘Criteria, Defeasibility, and Knowledge’

In this essay, McDowell’s ambitions are threefold: to cast doubt on the adequacy of a certain response to the traditional problem of other minds; to question an interpretation of Wittgenstein that yields this response; and to deploy a disjunctive conception of appearances in order to combat an argument that stands in the way of his favoured epistemological outlook.⁵ It is the first and third of these ambitions that concern us here.

The response to the problem of other minds that McDowell targets takes the following points for granted: first, if *S* knows that *p*, then *S* has a reason to believe that *p*; and, second, this reason is *defeasible*—it is consistent with *S*’s having the reason that *p* is not the case. McDowell describes the second of these points, in its application to the possibility of experientially acquired knowledge of other minds, as follows:

⁴ The latter is what McDowell (1998) calls a ‘transcendental’ concern. This is not the only thing philosophers have meant by this word.

⁵ This outlook is also relevant to other aspects of McDowell’s philosophy. Its relevance to his take on Michael Dummett’s anti-realism in general and the ‘manifestation requirement’ in particular is apparent in McDowell (1981) and (1984), and to scepticism about perceptual knowledge of the external world in McDowell (1995) and his essay for this volume.

[Even] on the occasions that seem most favourable for a claim [that] someone else is in some 'inner' state, the reach of one's experience falls short of that circumstance itself—not just in the sense that the person's 'inner' state is not itself embraced within the scope of one's consciousness, but in the sense that what is available to one's experience is something compatible with the person's not being in the 'inner' state at all. (McDowell 1982: 371)⁶

This passage is supposed to capture the idea that when we acquire knowledge of other minds through experience, we do not literally *perceive* that the other is in an inner state (such as that of being in pain). This seems to be what McDowell is getting at when he says that, according to the targeted view, the other's inner state is not "embraced within the scope of one's consciousness".

McDowell's doubts attach to this idea of an experientially acquired defeasible reason. His target is what he (1995: 402) calls a "hybrid conception" of experientially acquired knowledge, according to which such knowledge comes in two parts: possession of an experientially acquired defeasible reason for believing *p*, and *p*'s being the case.

Consider a pair of cases, in both of which someone competent in the use of some claim [possesses an experientially acquired defeasible reason for the claim], but in only one of which the claim is true. [The] story is that the scope of experience is the same in each case; the fact itself is outside the reach of experience. And experience is the only mode of cognition—the only mode of acquisition of epistemic standing—that is operative . . . How can a difference in respect of something conceived as cognitively inaccessible to both subjects, as far as the relevant mode of cognition goes, make it the case that one of them knows how things are in the inaccessible region while another does not—rather than leaving them both, strictly speaking, ignorant on the matter? (McDowell 1982: 373–4)

These remarks point towards the alternative outlook that McDowell wants to recommend: in order to have experientially acquired knowledge (of other minds or anything else), the putatively known facts *must* be "cognitively accessible" to the subject via the "mode of cognition" by which she is supposed to acquire the relevant knowledge; in other words, and taking the visual case as our example, in order to have perceptual knowledge that *p*, *S* must *see that p*.

The idea of seeing that *p* is central to McDowell's epistemological outlook; it is the idea of an *experiential*, *factive*, and *epistemic* state. The state is experiential in that if *S* sees that *p* then it looks to *S* as if *p*—where the consequent is interpreted phenomenologically, rather than ascribing a tentative judgement by *S*. It is factive in that if *S* sees that *p* then *p* is so. And it is epistemic in the following two respects: that *S* sees that *p* provides an indefeasible reason for *S* to believe that *p*; and, if *S* does see that *p*, then *S* is in a position to know that *p*—or, as McDowell (1982: 390) sometimes puts it, "knowledge of the fact" that *p* is "made available" to him.

Being in a position to know something is not the same as knowing it: if "someone has been misled into thinking his senses are out of order . . . we might then hesitate

⁶ We have elided some of McDowell's words here. He speaks, not simply of a claim that someone else is in some inner state, but of a claim that one sees that someone else is in such a state. It might be thought that, even if it is granted that seeing that someone else is in such a state can give one a reason for the former, there is a question as to whether it can also give one a reason for the latter. In this exegesis, we will concentrate on the former, and discuss the latter *en passant*.

to say he possesses the knowledge that his senses (in fact functioning perfectly) make available to him" (McDowell 1982: 390 fn. 37). So, it seems that even if *S* sees that *p*, *S* will not know that *p* if he does not possess (at least) the appropriate beliefs.⁷ And yet, once *S* is in a position to know, he needs "no extra help from the world to count as knowing" (McDowell 1995: 406). In order to ensure that seeing that *p* can put one in a position to know, many of the conditions that "post-Gettier" epistemology considers central to knowing must be written in to the conditions for seeing that *p*. This ensures that cases we might be tempted—pre-theoretically—to count as ones in which we see that *p* will not count as such, as the following quotation illustrates.

If one's senses are actually out of order, though their operations are sometimes unaffected . . . an experience subjectively indistinguishable from that of being confronted with a tomato, even if it results from confrontation with a tomato, need not count as experiencing the presence of a tomato [that is, need not count as seeing that there is a tomato before one]. Another case in which it may not count as that is one in which there are lots of tomato façades about, indistinguishable from tomatoes when viewed from the front. (McDowell 1982: 390, fn. 37)

The justification provided by such seeing is supposed to be non-inferential as well as indefeasible. There *is* an excellent inference from "the fact that someone sees that things are thus and so to the fact that things are thus and so" (McDowell 1982: 416), and it is in part because there is that states of seeing that *p* have their epistemic importance. However, when the reason in virtue of which someone knows that *p* is provided by the fact that she sees that *p*, her knowledge is not the result of an inference that *she* makes, from this reason, to *p*. Exactly how we are to understand the idea of non-inferential justification is another matter; but one consequence is clear: that one sees that *p* is not "something . . . of which one can assure oneself independently of the claim" that *p* is so. This "flouts an idea we are prone to find natural, that a basis for a judgement must be something on which we have a firmer cognitive purchase than we do on the judgement itself; but although the idea can seem natural, it is" —McDowell thinks— "an illusion to suppose it is compulsory" (McDowell 1982: 385).

McDowell wants to ensure that the possibility of this epistemological outlook is not obscured by (a version of) the argument from illusion. The argument runs as follows. On any occasion on which we attempt to acquire knowledge by looking, deception is possible, in the following sense: for any fact *p* that concerns public objects, it can look to *S* as if *p* even though *p* is not the case. It follows that "any capacity to tell by looking how things are in the world independent of [*S*] is at best fallible" (McDowell 1982: 386). Non-deceptive cases are also possible, of course; it can look to *S* as if *p* when *p* is the case. But, according to this argument, the reason that *S* acquires through experience must be the same in the deceptive and non-deceptive cases alike.⁸

⁷ Having appropriate beliefs falls under the general head of what McDowell (1993: 429) calls "doxastic responsibility". There is more on this in McDowell (1993).

⁸ McDowell's exact words are: "since there can be deceptive cases experientially indistinguishable from non-deceptive cases, one's experiential intake—what one embraces within the scope of one's consciousness—must be the same in both kinds of case" (McDowell 1982: 386). And, according

It follows that S's experientially acquired reason can be no better than: that it looks to S as if p . It can be no better than it is in the deceptive case, and as a result can be no more than defeasible, and so cannot be that S sees that p . When McDowell speaks of the *highest common factor* conception of experience, he simply means this idea that the reason S's experience makes available to S can never be any better than the reason it makes available to S in deceptive cases.⁹

There are various ways of responding to this argument. We might attempt to deny its premise, by claiming that there is a class of facts concerning public objects about which we cannot be mistaken: for such p , if it looks to S as if p then p is the case. Another strategy would be to try to provide a safe haven for cases of seeing, by insisting that no non-factive position is present in such cases: if S sees that p then it does *not* look to S as if p . McDowell employs neither strategy. His disjunctive conception is designed to deny the move from fallibility (of perceptual knowledge) to defeasibility (of experientially acquired reasons); in other words, to deny the move from the fact that it can look to S as if p , when not p , to the fact that none of the experientially acquired reasons in virtue of which S knows that p can be any better than what is shared between deceptive and non-deceptive cases—namely, that it looks to S as if p . So, he suggests the following: a state in which it looks to S as if p is *either* a state of S's seeing that p , and thereby being put in a position to know that p , and so acquiring an indefeasible reason for believing that p , *or* a state in which it merely looks to S as if p , in which S acquires no such reason, and so is not put in any such position.¹⁰ *This* is the disjunctive conception of appearances of (McDowell 1982).¹¹ It denies the highest common factor conception, precisely because it refuses to understand the reason in (at least some) non-deceptive cases as defeasible.

When he introduces his disjunctive conception, McDowell invites us also to look at the discussion of a disjunctive account of 'looks statements' in Snowdon (1980–1). We are now in a position to see what these two accounts have in common, aside from the fact that they are expressible in disjunctive form.

to McDowell, the argument from illusion claims that what one can embrace in either case can be no better than that it looks to one as if p . So, because one's experientially acquired reason is what is constituted by what is so embraced, the argument ensures that in neither case can this reason be any better than it looks to one as if p .

⁹ This brings out how Wright's (2002) claim that S's justification consists in the whole disjunction (that is, in the fact that either S is seeing that p or it merely looks to S as if p) is itself a version of the highest common factor conception.

¹⁰ In both his (1994), and his essay for this volume, McDowell suggests we understand the second disjunct as reporting a case in which it looks to S as if p but p is not so. This may be harmless; but, if cases in which S sees that p , and cases in which it merely looks to S as if p , are supposed to exhaust the options, it will serve to occlude cases in which S does not see that p , but still it looks to S as if p , and p is so. It is clear that McDowell (1982) recognizes cases of this kind (in his example of the tomato and the tomato façades discussed above). And McDowell (1993) flags two other considerations each of which make for a case of this kind: (i) the fact that it looks to S as if p results from the fact that p , but not "in the way that is characteristic of seeing"; and (ii) the subject of the experience is 'doxastically irresponsible' (for which see footnote 7).

¹¹ There are various ways in which McDowell chooses to express this conception. We employ the present formulation so as to bring out its epistemological import.

Both accounts refuse to credit a certain kind of philosophical significance to the fact that a pair of states can each be truly described as a state of its looking to S as if things are some way (where the way is the same in both cases, and the looking is phenomenological). Snowdon refuses to let this fact ensure that the states have the same intrinsic nature, because he makes room for the possibility that, when S's being in the L-state is S's seeing *o*, the L-state is constituted, in part, by the object seen. McDowell refuses to let the fact ensure that the states have the same epistemological status, because he makes room for the possibility that at least one of the states puts the subject in a position to know a fact about public objects (by being a state of seeing that *p*).

There is no reason to saddle Snowdon with McDowell's epistemological commitments. Is there reason to think that McDowell is committed to Snowdon's experiential disjunctivist conception of experience, but with seeing that *p* in place of seeing *o*? That is to say, must McDowell insist that those states of its looking to S as if *p* that are states of S's seeing that *p* are constituted by the fact seen?

On the one hand, we might think so. If (as McDowell claims) some states of its looking to S as if *p* are states of S's seeing that *p*, we might suppose him to be making an identity claim with the following implication: just as it is part of the intrinsic nature of states of seeing that *p* to involve the fact that *p*, it must also be part of the intrinsic nature of those states of its seeming as if *p* that are identical to states of seeing that *p* to involve the fact seen. This would bring McDowell's view much closer to Snowdon's experiential disjunctivism.

On the other hand, we might think not. One might hold that a state of its looking to S as if *p* is a state of S's seeing that *p* just in case it is (appropriately) caused, but not constituted, by the fact that *p*. And there is some evidence that McDowell accepts this. He speaks of certain experiences as having the 'fact itself' as their 'object'—that is to say, as being cases of seeing that *p*—“when we have them as the upshot (in a suitable way) of the fact” (McDowell 1982: 388–9). And he also writes of “how things look to one” being the “result” of “how things are . . . in the way that is characteristic of seeing” (McDowell 1993: 430, fn. 25). This should not surprise us, for it is not clear that McDowell's purposes in this essay provide any reason to object to the idea that cases of its looking to S as if *p* are “intrinsically independent” of surrounding facts. And there is no problem with the idea that a looking is a seeing, not because of its intrinsic nature (which it shares with mere lookings), but because of its extrinsic nature—its relatedness to the facts. Consider the following analogy. We say of red marks on the skin that some are sunburns and others not; the red marks count as sunburns in virtue of their having an appropriate causal history; but their having such a history is compatible with the fact that their intrinsic nature is shared with red marks that do not count as sunburns. So, on the basis of what McDowell (1982) says, it seems there are grounds for thinking that there is considerable distance between his disjunctive conception of appearances and Snowdon's experiential disjunctivism.¹²

McDowell (1982) also introduces a transcendental concern that does not show up anywhere in Snowdon's thinking. He appears to consider (with a view to rejecting)

¹² It is worth noting that, in other papers, McDowell *does* appear to commit himself to the Snowdon-style view, as we will see in section 2.2 below.

the following thought, which applies to all experiences: although we do not ever see that p , if we have an experience in which it looks to us as if p , and this experience is suitably caused by the fact that p , then we are in a position to know that p , even though the fact that p is not something we have actually perceived. This thought depends on the assumption that we can have experiences with representational content; the assumption, that is to say, that there can be cases in which it looks to us as if p , where p is capable of truth and falsity (or some appropriate analogue). McDowell appears to argue that there can be experiences with such content “only because when [our experiences] are the upshot (in a suitable way) of the fact” that p (McDowell 1982: 389), our experiences *are* states of seeing that p . In other words, he seems to think it a condition of the possibility of states of its looking to S as if p —that is, of experiences with representational content—that these states can, in the right circumstances, be ones of S ’s seeing that p . If McDowell is right about this, then there is a problem at the heart of the present thought: in claiming that states of its looking to S as if p cannot be states of S ’s seeing that p , it denies a condition of the possibility of experiences’ having content, and, as a result, its proponents are no longer entitled to assume that there can so much as be such states. It seems right to call this argument transcendental, in McDowell’s sense, because it is concerned with the very possibility of representational content. (It is interesting that it also seems to presuppose causation between experiences and the world.)¹³ McDowell returns to this transcendental concern in his essay for this volume.¹⁴

McDowell’s epistemological outlook raises a host of questions. States of seeing that p are supposed to yield indefeasible reason for believing that p ; but does S need to believe that p in order to know that p ? Is his a view on which knowledge is true justified belief? Or is it one on which factive states, such as seeings that p , are themselves states of knowing given appropriate beliefs (which may or may not include the belief that p)? Do we require an indefeasible reason to believe *every* fact we know?¹⁵ And what exactly *is* the reason in visually based cases? Is it the fact that S sees that p ? Or is

¹³ The text upon which this reading is based is as follows: “It seems unproblematic that if his experience is in a suitable way the upshot of the fact that it is raining, then the fact itself can make it the case that he knows that it is raining. But that seems unproblematic only because the content of the appearance is the content of the knowledge. And it is arguable that we find that match in content intelligible only because we do not conceive the objects of such experience as in general falling short of the meteorological facts [that is, on our reading, only because we do conceive some experiences as states of seeing that p]. That is: such experiences can present us with the appearance that it is raining only because when we have them as the upshot (in a suitable way) of the fact that it is raining, the fact itself is their object; so that its obtaining is not, after all, blankly external [that is, on our reading, only because when our experience is the suitable upshot of the fact it is a state of seeing that p]” (McDowell 1982: 389).

¹⁴ Is the transcendental argument he develops in this volume the transcendental argument of McDowell (1982)? According to the 1982 argument, it is intelligible that experiences have content only if it is also intelligible that, if the experiences are caused (in a suitable way) by environmental facts, then the experiences are cases of seeing that p . This might look like a different argument from the one he presents in his essay for this volume. But the 1982 argument surely implies at least the following: it is intelligible that experiences have content only if it is also intelligible that there can be experiences that are cases of seeing that p . And that is McDowell’s ‘new’ transcendental argument.

¹⁵ McDowell (1982) seems to endorse the possibility of scientific knowledge that requires mediation by theory. Must this knowledge be based on indefeasible reasons as well?

it the fact that p , which S only has as her reason when she takes it in by seeing? Different answers to these last two questions would give different ways for McDowell's epistemology to make sense of the requirement that S 's reason must be *accessible* to S , when S knows. If S 's reason is p , then we might think accessibility amounts simply to S 's seeing it. If S 's reason is that she sees that p , it seems accessibility must amount to something like second-order knowledge to the effect that she sees. McDowell's outlook makes clear sense of the first sort of accessibility, but a question remains as to how it is to make sense of the second.

There is a passage in McDowell's essay for this volume in which he speaks of the fact that S sees that p as S 's indefeasible reason, and that suggests that he needs to insist on the possibility of the relevant second-order knowledge, if he insists on an accessibility requirement. But the context of this passage suggests that he might intend this indefeasible reason to be, not (or perhaps not simply) the reason in virtue of which S knows that p , but (also) the reason in virtue of which S is entitled to *claim* that she sees, and so knows, that p . He says the following:

What does entitle one to claim that one is perceiving that things are thus and so, when one is so entitled? The fact that one is perceiving that things are thus and so. That is a kind of fact whose obtaining our self-consciously possessed perceptual capacities enable us to recognize on suitable occasions, just as they enable us to recognize such facts as that there are red cubes in front of us, and all the more complex types of environmental facts that our powers to perceive things put at our disposal. (this volume: section 5)

Later in the same section of his essay, McDowell says something about how we are to understand this second-order knowledge:

If the animal in front of me is a zebra, and conditions are suitable for exercising my ability to recognize zebras when I see them (for instance, the animal is in full view), then that ability, fallible though it is, enables me to see that it is a zebra, *and to know that I do*. (our italics)

So, in McDowell's account, the fact that I see that p , and the fact that I know that I see that p , are equally the upshots of the operation of my abilities to recognize F s when I see them.¹⁶ It is worth noting that this point connects up with two features of McDowell's earlier work: his (1982: 390) claim that we count as seeing that p "only in the exercise of an ability to tell that [p]", and his (1986) claim that, once the inner and the outer realms are pictured as "interpenetrating", so-called introspective knowledge to the effect that one sees that p must be understood as a "by-product" of the exercise of our *perceptual* capacities.

Epistemological disjunctivism is discussed further in the essays by Alex Byrne and Heather Logue, John McDowell, Alan Millar, Ram Neta, Duncan Pritchard, and Crispin Wright. These essays fall into two overlapping groups. The first group consists of the essays by McDowell, Pritchard, Neta, and Wright, and considers epistemological disjunctivism's prospects as a way of undermining skepticism about knowledge of the external world. The second consists of the essays by Byrne and

¹⁶ Alan Millar's essay for this volume offers an extensive treatment of the idea of a recognitional ability.

Logue, Millar, Neta, and Pritchard, and considers its prospects as an account of perceptual knowledge.

2.2 ‘Singular Thought and the Extent of Inner Space’

In this paper, McDowell’s ambition is to present a picture of the mind that has the potential to liberate us from (epistemological and transcendental) Cartesian problems. In order to do this, he introduces both a conception of object-dependent Fregean thoughts and a version of his disjunctive conception of appearances.

A thought is object-dependent only if its availability for thought, or expression, depends upon the existence of the object it concerns. As McDowell (1986: 228) puts it, following Russell: genuinely referring expressions “combine with predicates to express [thoughts] that would not be available to be expressed at all if the objects referred to did not exist”. A thought is Fregean in the present sense if and only if it conforms to the Intuitive Criterion of Difference, according to which “thoughts differ if a single subject can simultaneously take rationally conflicting attitudes towards them (say, any two of acceptance, rejection and neutrality) without thereby standing convicted of irrationality” (McDowell 1994: 180).

We get the idea of an object-dependent Fregean thought by putting these two ideas together.

In its directly psychological application, the resulting idea ensures that the “configurations a mind can get itself into [are] partly determined by what objects there are in the world” (McDowell 1986: 230). And, according to McDowell, the resulting picture of the mind stands in opposition to Descartes’s fundamental contribution to philosophy: a picture of the mind (or “subjectivity”, or the “inner”) “as a region of reality that is transparent—accessible through and through—to [our] capacity for [infallible] knowledge” (McDowell 1986: 240). On this Cartesian picture, *every* fact about the mind’s layout is supposed to be knowable in this (infallible) way.

McDowell also opposes this picture with a version of his disjunctive conception of appearances, which he describes as follows: ‘of facts to the effect that things seem thus and so to one . . . some are cases of things being thus and so within the reach of one’s subjective access to the external world [in the visual case, cases of seeing that *p*], whereas others are mere appearances’ [that is, cases of its merely seeming as if *p*] (McDowell 1986: 241). Through its combination of the factivity of cases of seeing that *p* (for example), and the insistence that such cases obtain in the inner realm, this conception invites us to “picture the inner and outer realms as interpenetrating, not separated from one another by the characteristically Cartesian divide” (McDowell 1986: 241).

As before (in section 2.1), there is an interesting question as to whether McDowell’s disjunctivist conception is distinct from Snowdon’s disjunctive theory.

On the one hand, perhaps McDowell agrees with the causal theorist that we can explain what it is for *S* to see that *p* in terms of there being a causal relationship between a purely outer fact and a purely inner state (of its seeming to *S* as if *p*), for it is not clear why doing so should rule out an understanding of seeing as a state that occupies a region of the inner realm that interpenetrates with the outer. (In doing so, we are insisting that some inner states do not occupy the interpenetrating region; there remain

purely inner states (of its seeming to S as if p), whose intrinsic nature does not consist in anything from the outer realm.) If this is what McDowell is saying, then when he says that some cases of seeming are cases of seeing, we need to understand him—once again—as saying that cases of seeming count as such in virtue of their aetiology.

On the other hand, when McDowell returns to summarize the conception, he says the following:

Short of the fully Cartesian picture, the infallibly knowable fact—its seeming to one that things are thus and so—can be taken disjunctively, as *constituted* either by the fact that things are manifestly thus and so [for example, that one sees that p] or by the fact that that merely seems to be the case (our emphasis). (McDowell 1986: 242)

In the visual case, this passage is naturally interpreted as saying that what it is for it to look to S as if p is either for S to see that p (and so to be in a state that “intrinsically involves” a fact), or for it merely to look to S as if p (and so for S to be in a state that is ‘intrinsically independent’ of a fact). It turns McDowell’s disjunctive conception into a variant of Snowdon’s disjunctive theory (but with ‘seeing that p ’ in place of ‘seeing o ’). And it also offers a more complete rejection of the Cartesian divide between the inner and the outer, by making it impossible to characterize the intrinsic nature of certain paradigmatically inner states—namely, certain states of seeming—without mentioning facts from the outer realm.

McDowell (1994) clearly favours this more complete rejection in another essay:

Compare the psychological feature that is unsurprisingly shared between someone who sees that such-and-such is the case and someone to whom it merely looks as if such-and-such is the case It is not compulsory to conceive seeing that such-and-such is the case as constituted by the common feature together with favourable facts about embedding in the environment. We can understand things the other way round: the common feature—its being to all intents and purposes as if one sees that such and such-is the case—intelligibly supervenes on each of the divergent “wide” states. And it is better to understand things this way round.

Here by “‘wide’ states” McDowell seems to have in mind states of S’s seeing that p , and states of its merely seemingly to S as if p , which—precisely because they are wide—are such that they can be had only if p is so (in the former case), and only if p is not so (in the latter).¹⁷

McDowell seems to go on the following journey: in McDowell (1983) he does not commit himself to this rejection of the causal theory’s picture of experiences as intrinsically independent of the world; but he moves towards it in McDowell (1986), and finally makes it explicit in McDowell (1992). However, he has not yet made clear if, and if so why, we need to reject this picture in order to secure the epistemological goals of McDowell (1982). And this is a significant omission, in light of the fact that it is far from obvious that we need to reject this picture of experiences in order to deal with these worries in the way he recommends.

At a very general level, the point of both this version of the disjunctive conception of appearances, and the conception of object-dependent Fregean thought, is to provide a non-Cartesian picture of the mind. But, at a less general level, we face the question

¹⁷ See footnote 10 for some of McDowell’s own reasons for being suspicious of this idea.

of whether the Cartesian picture under attack is simply the idea that mental items are not object-involving, or the idea that mental items are object-involving but only as a matter of their extrinsic, rather than their intrinsic, nature. If the first, McDowell does not need to reject the causal theory. If the second, he does.

The essays by Bill Brewer and Sonia Sedivy in this volume have more to say about McDowell's approach to perception, and its implications for our understanding of the mind. Brewer suggests that McDowell's disjunctivism (as he understands it) is best combined with a conception of experience as wholly lacking in representational content, whereas Sedivy, by contrast, argues that it should be seen as essential to a form of direct realism which understands experience as having a representational content that is conceptual through-and-through.

3 PERCEPTION AGAIN: MARTIN AND PHENOMENAL DISJUNCTIVISM

Up to this point, we have compared and contrasted experiential disjunctivism and epistemological disjunctivism. We will now turn our attention to what we will call 'phenomenal disjunctivism'. This form of disjunctivism is widely associated with the work of M. G. F. Martin. And as we will see, phenomenal disjunctivism is somewhat different from the disjunctive positions articulated by Snowdon and McDowell.

Martin is concerned predominantly, but not exclusively, with the phenomenology of experience. He is interested in defending (what he calls) *naïve realism*, the view that "mind-independent objects are present to the mind when one perceives, [and] that when one has such experience, its objects must actually exist" (Martin 2002: 393). Martin claims that naïve realism is the best theory of how our experiences strike us (Martin 2004: 42); namely, as phenomenally "transparent":

At heart, the concern is that introspection of one's perceptual experience reveals only the mind-independent objects, qualities and relations that one learns about through perception. The claim is that one's experience is, so to speak, diaphanous or transparent to the objects of perception at least as revealed to introspection. (Martin 2002: 378)

A good way to understand naïve realism is to set it against the background of some opposing views.

The *sense-datum* theory of perception and the *intentional* theory of perception are moved in the same way as each other by the consideration that it seems possible for there to be cases of hallucination in which the ways things seem to the subject are exactly as they would be in a case of veridical perception. Both theories assume that this consideration shows that the hallucinatory experience and the veridical perceptual experience have the same phenomenology.¹⁸ In what seems like a natural move,

¹⁸ By 'veridical perceptual experience' we mean to refer to the experiences involved in veridical perception, as opposed to those involved in hallucination. This usage contrasts with another in which 'veridical perceptual experience' refers to experiences that accurately reflect how things are and that may, or may not, be involved in cases of hallucination. In this latter usage, 'perceptual' is

they hold that the experiences in both cases are therefore identical *qua* experiences, and thus that we should give the very same account of the nature of each. In which case, the mind-independent objects that we perceive cannot explain their natures precisely because there are no such objects answering to how the world seems in the hallucinatory case.

The sense-datum theory posits mind-dependent objects—*sense-data*—as the objects of both the hallucinatory and veridical perceptual experiences. It claims that in having these experiences sense-data are perceived—and, indeed, that such objects cannot be misperceived. The alleged common nature of hallucinatory and veridical perceptual experiences is thereby explained. This theory does some justice to the phenomenology of hallucination in which it seems to us as if we are perceiving objects, for it says that we are actually perceiving objects. The caveat is that these objects are mind-dependent. It is possible to think that the theory gets the phenomenology of hallucination wrong to some degree if one thinks that the phenomenology of hallucination involves the seeming presentation of mind-independent objects in public space. In the case of the phenomenology of veridical perceptual experiences, realist versions of the sense-datum theory hold that there are mind-independent objects, and that when we perceive such objects we indirectly perceive them in virtue of directly perceiving sense-data. To the extent that one thinks that experience is transparent, the sense-datum theory will have to claim that the phenomenology of experience is misleading. Of course it is possible for a sense-datum theorist to try to deny the transparency claim. This might be done by claiming that the apparent direct contact that we have with mind-independent objects is not part of experiences' phenomenology, but is instead part of some judgement or similar cognitive reaction that is made in response to our experience. Such judgement might be held to occur automatically, and rather quickly, and not on the basis of conscious inference, thus leading to the false, but understandable, claims about the phenomenology of experience. It is also worth noting that many philosophers have found the sense-datum theory unattractive in light of its seemingly metaphysically extravagant and non-naturalistic commitments.

'The intentional theory' labels many different views of experience. One such view holds that the common phenomenal character and nature of hallucinatory and veridical perceptual experiences are to be explained by holding that both have the same representational content—in other words, both represent the world to be a certain way, and the world may, or may not, be that way.

Propositional attitudes are the paradigm of states with content. It is instructive to compare experience with one type of propositional attitude—namely, belief. A subject who believes that *p* must be committed to the truth of its content. It is often thought that a subject of experience need not have quite the same commitment to the truth of the content of the experience. This is on account of cases where a subject believes themselves to be hallucinating or undergoing a perceptual illusion (whether or not they are) and at the same time believes that the world is not as it seems; or, where the subject suspends belief as to how the world is, perhaps on account of

taken to mark out such experiences from experiences of other types such as emotional experiences, or experiences involved in mere sensations.

believing that conditions for viewing are not good; or, where the testimony of others leads the subject to believe that things are not how they seem. Instead, the commitment of the subject of an experience might be, to use the terminology of Millar (1991), that *in the absence of countervailing considerations* they will be committed to the truth of the content; or, in the terminology of Armstrong (1961), that they *have a prima facie inclination to believe* the truth of the content. There are a number of other ways in which the content of experience is often thought to differ from that of belief, thereby further distinguishing the two states.

The core idea of the intentional theory is that a subject can bear the same relationship to the same content in introspectively indiscriminable perceptual and hallucinatory cases. And the subject's doing so explains the alleged common phenomenology and nature of hallucinatory and veridical perceptual experience. What is more, the phenomenology of experience is completely explained by its content.

The intentional theory accounts to some degree for the phenomenology of experience, and offers a particularly good account of veridical perceptual experience's apparent transparency because it claims that in such cases we do directly see mind-independent objects. There are no intermediaries of the kind that the sense-datum theory posits. Of course this theory must avoid saying that our being aware of the content of experience is our being aware of properties of the experience in virtue of which it represents, because according to the transparency claim, when experiencing we seem to be aware, not of properties of our experience, but of mind-independent objects in public space. Whether intentionalism can avoid saying this is a matter of some dispute.

Those who think that intentionalism cannot do so may hold something of a hybrid view that attempts to combine aspects of sense-datum theory with aspects of intentionalism of the kind we have been discussing. Such a view is also, unhelpfully, liable to be labelled an intentionalist view. We will call such a view a *weak intentionalist* view and distinguish it from the previous view, which we will hereafter call a *strong intentionalist* view. Weak intentionalist views hold that experience has intentional contents but that such contents do not alone explain its phenomenal character, although they may do so in part. Experience also has non-intentional properties that help to explain its phenomenal character. These properties are most often held to be intrinsic properties of experience. Different versions of the view will hold that such properties may or may not be the properties in virtue of which an experience represents. If such a view held that it is in virtue of being aware of such properties that the subject of the experience was aware of the phenomenology of the experience then the view would share some of the commitments of the sense-datum theory as it would have either to deny transparency or to hold that the phenomenology of experience is not a good guide to its nature. Such a view avoids the particular perceived metaphysical extravagances of sense-datum theory, but some will argue that it nonetheless is landed with others that are no less troubling. Weak intentionalism can garner the explanatory resources of strong intentionalism concerning those phenomenological aspects of experience, if any, that it takes to be explained by content in the manner of strong intentionalism.

Strong intentionalism must of course think that the phenomenology of experience is in error in some respects. In hallucination we appear to be directly aware of

mind-independent objects; but strong intentionalism must hold that we are not aware of any such objects. And yet, it offers an explanation of why we seem to be so aware: we are in a state that represents the world to be a certain way—it is just that on this occasion the state misrepresents the way things are.

We can now see that the sense-datum and intentional theories employ a ‘highest common factor’ conception of experience, according to which what explains visual phenomenology must be the same in the two cases. The sense-datum theory posits shared perceived objects as the highest common factor, the strong intentional theory posits shared contents, and the weak intentional theory posits shared contents and shared intrinsic properties of experience. (Of course, it is essential not to confuse *this* highest common factor conception with that advocated by McDowell, which is exclusively concerned with the justification that experience can provide.)

Naïve realism offers an alternative to sense-datum theory and to intentionalism in its strong and weak varieties. According to naïve realism, in the case of veridical visual experience:

we should explain the phenomenal transparency in terms of the [mind-independent] objects of perception, and not in terms of the experience’s representational content: the objects have actually to be there for one to have the experience, and indeed one may claim that they are constituents of the experiential situation. (Martin 2002: 393)

However, accounting for hallucinatory experiences and their nature provides a notorious stumbling block for naïve realism. Martin’s position is distinctive for the way that it tries to avoid it.

Martin does not deny the possibility of introspectively indiscriminable cases of perception and hallucination. His crucial move is to insist that the experiences involved in both cases belong to different “fundamental kinds” (Martin 2002: 404; 2004: 43, 54, 60), because this allows him to claim that what explains the phenomenal character of the experience in the perceptual case is different from what explains it in the hallucinatory. In the perceptual case, it is explained by the perceived mind-independent stuff. In the hallucinatory case, there is no perceived mind-independent stuff, and the phenomenal character is explained simply by the fact that the state is introspectively indiscriminable from a situation of perceiving such stuff. Martin says:

A common explanation is not offered of the three cases [hallucination, illusion, and veridical perception]—we explain the veridical perception by reference to the relational properties it alone possesses, and we explain the other two by reference to their indiscriminability from this. So, the particular situation of veridical perception is fundamental to the explanation of the character of all cases of perceptual experience. (2002: 402)

In his later work, Martin goes further and commits himself to the idea that the phenomenal character of experiences in the perceptual case need not be the same as the phenomenal character of the experience in the hallucinatory case (Martin 2006: 366–72).

This response to the argument from hallucination denies the highest common factor conception, which claims that what explains the phenomenal character of the experiences is the same in all cases. And it does so by denying that the experiences

have the same constituents and are of the same fundamental kinds. Martin offers numerous arguments in defence of this view, which can be found in his (2002, 2004, and 2006).

So, what, if anything, does this phenomenal disjunctivism have in common with epistemological disjunctivism and experiential disjunctivism, and how does it differ?

Phenomenal disjunctivism seems to entail experiential disjunctivism. If the experiences (or L-states) involved in cases of perception and those involved in cases of hallucination belong to different fundamental kinds, then what it is to have an experience in a perceptual case must differ from what it is to have an experience in a hallucinatory case.

According to the disjunctivist, the phenomenological character of all perceptual experience requires us to view the transparency and immediacy of perceptual experience as involving actual relations between the subject and the objects of perception and their features. In just the case of veridical perception, the experience is a matter of certain objects being presented as just so. (Martin 2002: 402)

Thus, we could not have that experience, and according to Martin (2006) not even that phenomenology, in the absence of the mind-independent stuff, and one way to capture this point would be to say that the experience “intrinsically involves” the stuff. So, at least part of what it is to have an experience in the perceptual case is to be in a state that “intrinsically involves” such stuff. By contrast, at least part of what it is to have an experience (with the relevant phenomenology) in a hallucinatory case is to be in a state that does not involve the perception of mind-independent stuff.

But experiential disjunctivism arguably does not entail phenomenal disjunctivism. It seems to be consistent with the fact that an experience (or L-state) “intrinsically involves” public objects that these objects do not explain the phenomenology of the experience. Nor need an experiential disjunctivist think that the phenomenal character of the hallucination and of the veridical perceptual experience are different. Further, unlike a phenomenal disjunctivist who promotes naïve realism, an experiential disjunctivist can endorse intentionalism, and even accept some versions of sense-datum theory—namely, versions that took appeals to content and sense-data merely as a way to explain and conceive of the phenomenal character of experience, and not as a way to conceive of the entire nature of the experience. An experiential disjunctivist could hold that the phenomenology of veridical visual experiences is determined by a content or mind-dependent object that it shares with the hallucinatory case but claim that, nonetheless, the veridical experience *in toto* is partly constituted by mind-independent objects. On such a theory, the fact that an experience intrinsically involves an object is compatible with its making no difference to the experience’s phenomenal character. The phenomenal character of an intrinsically object-involving perceptual experience and the character of an object-free hallucination would have the same explanation, and could actually be the same, if the experiences shared the appropriate content or sense-data.

Phenomenal disjunctivism seems to differ substantially from epistemological disjunctivism. For example, the former need make no claim about the epistemic status of the different experiences, and the latter need make no claim about their

phenomenology. It seems clear that McDowell, our prime example of an epistemological disjunctivist, thinks experiences in perceptual and hallucinatory cases can share representational content that is not object-involving. In his (1998) he speaks of them both as cases in which we seem to see, or ostensibly see, that there is a red cube over there, for example. Nonetheless, it does not follow that such an epistemological disjunctivist is committed to thinking that they share phenomenal character, for there is no reason for them to endorse (or deny) any kind of strong intentional theory. Their concerns are simply different.

The final point to note is that phenomenal disjunctivism conforms to the *fundamental disjunctivist commitment*. It refuses to credit a certain kind of philosophical significance to the fact that a pair of states can each be truly described as a case of its looking to S as if things are some way (where the way is the same in both cases, and the looking is phenomenological). That is to say, it refuses to let this fact ensure that the states' phenomenal characters have the same explanations (or even that the states have the same phenomenal characters) because it makes room for the possibility that what it is for at least one of the states to have its phenomenology is for S to see mind-independent stuff in their environment.

The phenomenal disjunctivist position raises a host of fascinating questions. What notion of indiscriminability ought to be employed in spelling out a plausible version of the position? Can one claim both that hallucinatory experiences have phenomenal character and that a hallucinatory experience is *simply* a situation that is introspectively indiscriminable from a situation of perception? Is it possible to claim that some or all hallucinations *lack* a phenomenology? The essays in this volume by Scott Sturgeon, William Fish, A. D. Smith, and Susanna Siegel discuss these issues and either elaborate, or criticize, phenomenal disjunctivism.

An interesting question for phenomenal disjunctivism (indeed for any disjunctive view) is how to handle cases of illusion. Should they be treated more like hallucination or more like veridical perception? The essays by Bill Brewer and Alex Byrne and Heather Logue explore the answers to these questions. In addition, Byrne and Logue go on to address the issue of whether phenomenal disjunctivism ought to be the default position and, indeed, whether it is true at all.

4 ACTION

Disjunctivism in the philosophy of action is in its infancy. In this section, we trace its roots in McDowell's work, and distinguish its various varieties.

McDowell was the first person to offer a disjunctivist thesis in the philosophy of action. He suggests that S can acquire knowledge of another's mind by seeing that the other's behaviour is "giving *expression* to [the other's] being in [a certain] 'inner' state" (McDowell 1982: 387), where 'expression' is used in such a way that it is not possible for someone's behaviour to give expression to her being in an inner state unless she is in the inner state in question. McDowell refers to this suggestion as an aspect of his disjunctive conception of appearances, and draws attention to how it "flouts" the idea of an "interface" between the subject and "external reality" (McDowell 1982: 392–3).

In the context of visually acquired knowledge about the external world in general, this “interface” idea is the idea that our visually acquired justification is restricted to what is yielded by facts to the effect that it looks to us as if p . In the context of visually acquired knowledge of other minds in particular, the idea is that such justification is restricted to our perception of “psychologically neutral” facts about the “facing surfaces of other human bodies” (McDowell 1982: 392–3); that is to say, of facts about other human bodies that fall short of facts about inner states. The disjunctive conception of appearances introduced in section 2.1 flouts the first of these ideas, and the current suggestion flouts the second.

Although the disjunctive conception of appearances does not deny that visually acquired knowledge involves its looking to S as if p , it nevertheless insists that some of the cases in which it looks to S as if p are cases in which S acquires an indefeasible justification to believe p , precisely because they are cases in which S sees that p . The current suggestion makes an analogous epistemological point. It does not deny that visually acquired knowledge of other minds requires seeing how things are with the other’s body, but it nevertheless insists that some of the cases in which S sees this are cases in which S sees that the other’s behaviour expresses the fact that she is in a certain inner state. To illustrate: I might come to know that your hand is moving backwards and forwards by seeing that it is. In some cases such bodily movement may express your inner state of intending to greet me. In other cases it may not, and may instead express the fact that you intend to bid farewell, or that you desire to get rid of unwelcome flies, or that you are in another inner state. In yet other cases it may express no fact about your inner states at all. The view under consideration claims that, in some of the cases in which I see a bodily movement, I see that your behaviour expresses the fact that you intend to greet me. We might capture this idea by saying that the bodily movements available to S ’s perception are *either* movements that are expressive of the other’s being in a certain inner state, and which thereby help to put S in a position to know that the other is in that state when they are present in S ’s perception; *or*, mere movements, which are not so expressive, and whose perceptual presence to S cannot put S in a position to have perceptual knowledge of such facts. ‘Epistemological disjunctivism about bodily movement’ seems an apt name for this suggestion.

We might think of epistemological disjunctivism about bodily movement as, *inter alia*, a conception of the bodily behaviour involved in intentional bodily action. And, just as McDowell’s epistemological disjunctivism about experience sometimes seems to be consistent with a causal theory of visual perception, and so with the denial of experiential disjunctivism, the epistemological disjunctivist conception of bodily movements under discussion seems to be consistent with a causal theory of bodily action. According to such a theory, part of what it is for S to act intentionally is for there to be a causal relation between items of S ’s bodily behaviour that are intrinsically independent of S ’s mental states, and S ’s mental states. Is there any reason to think that the truth of this theory will prevent S ’s perceptions from being perceptions to the effect that others’ behaviour expresses facts about their mental states?

However, just as there can be a disjunctive theory of appearances that denies that all L-states—that is to say, all experiences—are intrinsically independent of public

objects, and thereby endorses experiential disjunctivism, so there could also be a disjunctive theory of bodily movement that does an analogous thing. This theory would tell us what it is to be in (what we might call) an M-state—a state that consists in one's body moving in a certain way. It is to be in a state that *either* intrinsically involves mental states (such as intentions), *or* is intrinsically independent of these states. Various philosophers have endorsed a theory of this kind, and for various reasons. To take two examples: in his essay for this volume, David-Hillel Ruben does so in order to avoid the difficulties that he claims to discern in causal and volitional theories of action; and, in her (1997), Jennifer Hornsby does so in order to show that certain bodily movements do not belong in a physicalist ontology, because they are essentially mind-involving.

So, as with disjunctivism about experience, disjunctivism about bodily movement comes in epistemological and metaphysical varieties.

A rather different way of employing disjunctivist thinking in the philosophy of action is to defend, and render plausible, a certain account of reasons for acting. This account consists of two components: (I) it must be possible that the reasons for which a subject acts are, in certain cases, the good reasons that there are for the subject to act; and (II) good reasons are facts.

Forms of disjunctivism that address this issue do not come in epistemological varieties, because they are concerned *exclusively* with what it is for us to act for a good reason, and *not* with what it is for us to possess knowledge (of our reasons, or anything else). They are developed in response to the following putative problem for the above account of reasons for acting. If the fact that *p* is both a good reason for the subject to act, and the reason for which the subject acts, it looks as if we can explain why the subject acts as they do with an explanation of the following shape: he acts because *p*. However, it also looks as if the agent can act for some kind of reason even though it is not a fact that *p*: in such a case, he acts because he believes that *p*. For example, when someone goes to the doctors because he believes he has the flu, he can go to the doctors for a reason even though, in fact, he does not have the flu, but merely a hang-over. And now a tempting highest common factor assumption looms into view: the reason for which the agent acts, and the rational explanation of his action, must be the same in cases where *p* is so as it is in cases where *p* is not so.

Two distinct disjunctive conceptions are employed to target this assumption. The first concerns the idea of a reason for which agents act, and says the following: the reason for which *S* acts is *either* that *p*, *or* that *S* believes that *p*. (Jonathan Dancy toys with, and rejects, this conception in Dancy (2000), and in his essay for this volume.) The second concerns the idea of a rational explanation of action: rational explanations to the effect that *S* acts because she believes that *p* entail explanations to the effect *either* that *S* acts because she knows that *p*, *or* that *S* acts because she merely believes that *p*—where mere belief is simply belief that falls short of knowledge. (Jennifer Hornsby endorses this conception in her contribution to this volume.) It may be possible to accept the first conception, and reject the second, and vice versa. For example, Dancy seems to think the first does not need the second because, although a sentence of the form '*S* acts because she believes that *p*' is an acceptable way of expressing a rational explanation, it is really a form of loose talk; strictly speaking, the way to

express such an explanation is with a sentence of the form ‘S acts for the reason that p ’. In similar vein, Hornsby seems to think the second conception does not require the first, because the first obscures the fact that the reasons for which agents act are *propositions*—true ones if the reasons are good, and false ones otherwise.

Exactly how disjunctivist thinking can help to defend this account of reasons for acting, and whether it is needed to do so, are points of contention that Dancy and Hornsby discuss in their essays in this volume.

5 CONCLUSION

According to Byrne and Logue (this volume), disjunctivism comes in epistemological and metaphysical varieties. We should now be able to see the truth in this claim, even as we can also see that it would be a mistake to think that McDowell—the chief proponent of epistemological disjunctivism—has never committed himself to at least one of the two sub-varieties of metaphysical disjunctivism: namely, experiential disjunctivism (as advocated, for example, by Snowdon) and phenomenal disjunctivism (as advocated, for example, by Martin). But can we say anything about what unites these disjunctivisms—what it is in virtue of which they deserve to be thought of as varieties of the same kind?

Philosophy has tended to credit a certain kind of significance to the fact that a pair of states can both be truly described as states of its looking to S as if things are some way (where the way is the same in both cases, and the looking is phenomenological). One such significance is that S’s experientially acquired reasons for belief must be the same in both cases. Another is that the intrinsic nature of such states must be the same in both cases, and cannot involve anything ‘outer’. And another is that, in both cases, experiential phenomenology, and what explains it, must be the same. It is a mark of disjunctivism about experience to question these assumptions.

Philosophy has also tended to credit a certain kind of significance to the fact that a pair of states can both be truly described as ones in which someone’s body moves in some way. One such significance is that S’s perception of the other’s bodily movement must yield the same knowledge in both cases. Another is that the intrinsic nature of each bodily movement must be the same, and cannot involve anything ‘inner’. It is a mark of disjunctivism about bodily movement to question these assumptions.

Finally, philosophy has tended to credit a certain kind of significance to the fact that a pair of cases can both be truly described as ones in which someone acts for a reason. One such significance is that in both cases the reason must be the same. Another is that rational explanations of the relevant actions must, in both cases, be exhausted by: the agent acts because they believe that p . It is a mark of disjunctivism about reasons for acting to question these assumptions.

But apart from the fact they are questioned by disjunctivisms, is there anything else that these assumptions have in common?

The Cartesian picture of the mind understands the inner world as constitutively independent of anything outer: it is possible that the outer world differs radically,

and the inner world remains exactly as it is. There are various ways of understanding the notion of an inner world. We might understand it as the idea of (amongst other things) the subject's experiences (what Snowdon calls their L-states). We might understand it as (simply) the idea of a subject's experientially acquired reasons—those reasons that a subject enjoys when it looks to him as if things are a certain way. And we might also understand it as the idea of the phenomenology of experience, perhaps for the reason McDowell (1986) suggests, namely, that it is the idea of what is within the scope of the subject's capacities for infallible knowledge.

In this light, we can see that Snowdon's metaphysical disjunctivism opposes the Cartesian picture once the idea of the inner world is understood to include a subject's experiences, for it understands some of these as constituted by outer objects and facts. McDowell's epistemological disjunctivism also opposes it, once the idea of the inner world is understood as the idea of the subject's experientially acquired reasons, for this variety of disjunctivism insists that, in some cases, these reasons can be to the effect that she sees that *p*. And Martin's metaphysical disjunctivism also opposes it in the experiential case, once the idea of the inner world is understood as the domain of experiential phenomenology, for it insists that this can sometimes be explained (and even constituted) by outer objects and facts.

Something similar can be said for disjunctivism about bodily movement. We might understand the idea of the outer as consisting at least of the subject's bodily movements. We might understand it as (simply) the idea of those of a subject's reasons that concern the behaviour of others—reasons that specifically concern how things are with the bodies of others. Metaphysical disjunctivism about the surface sees some bodily movements as partly constituted by inner states. And its epistemological variant sees some of the reasons that concern behaviour as also concerning the state of the other's inner world (as when subjects see that the other's behaviour is expressing the fact that they are in pain). We might think of disjunctivism about bodily movement as an inverted, but structurally analogous, version of disjunctivism about experience, because it asserts, not the suffusion of the inner by the outer, but that of the outer by the inner. Put together, both kinds of disjunctivism yield a conception of the inner and the outer as each suffusing the other.

Disjunctivism about reasons for acting has a similar significance. Reasons for which agents act might be thought to belong to the subject's inner world. By insisting that some such reasons are facts (such as the fact that the subject has the flu), it serves to place aspects of the outer within the scope of the inner. By the same token, the *explanans* of rational explanations might be thought to show the agent to be moved only by states of the inner world—namely, only by facts that the agent believes that *p*. Disjunctivism about reasons for acting ensures that some explanations that appeal to such states *ipso facto* guarantee that the agent is moved by aspects of the outer: they act because *p*. Again, we have a picture in which the inner and the outer suffuse one another.

We can, therefore, say the following: the mark of disjunctivism, in all of its varieties, is a conception of the inner and the outer as suffused. Each brand of disjunctivism—be it about experience, bodily movement, or reasons for acting—exemplifies this central mark in distinct, and sometimes very different, ways.

REFERENCES

- Armstrong, D. M. (1961) *Perception and the Physical World*, London: Routledge and Kegan Paul.
- Dancy, J. (2000) *Practical Reality*, Oxford: Clarendon Press.
- Grice, H. P. (1961) 'The Causal Theory of Perception', *Proceedings of the Aristotelian Society* Supplementary Volume 35: 121–52.
- Hinton, J. M. (1967a) 'Experiences', *Philosophical Quarterly*, 17: 1–13.
- (1967b) 'Visual Experiences', *Mind*, 76: 212–27.
- (1973) *Experiences*, Oxford: Clarendon Press.
- Hornsby, J. (1997) *Simple Mindedness: A Defence of Naïve Naturalism in the Philosophy of Mind*, Cambridge, MA: Harvard University Press.
- McDowell, J. (1981) 'Anti-Realism and the Epistemology of Understanding', in H. Parret and J. Bouveresse (eds.) *Meaning and Understanding*, Berlin: Walter de Gruyter. [Reprinted in his *Meaning, Knowledge, and Reality*, Cambridge, MA: Harvard University Press, 1998.]
- (1982) 'Criteria, Defeasibility, and Knowledge', *Proceedings of the British Academy*, 68: 455–79. [Reprinted in his *Meaning, Knowledge, and Reality*, Cambridge, MA: Harvard University Press, 1998.]
- (1984) 'Wittgenstein on Following a Rule', *Synthese*, 58: 325–63. [Reprinted in his *Mind, Value, and Reality*, Cambridge, MA: Harvard University Press, 1998.]
- (1986) 'Singular Thought and the Extent of Inner Space', in P. Petit and J. McDowell (eds.) *Subject, Thought, and Context*, Oxford: Clarendon Press. [Reprinted in his *Meaning, Knowledge, and Reality*, Cambridge, MA: Harvard University Press, 1998.]
- (1992) 'Putnam on Mind and Meaning', *Philosophical Topics*, 20: 35–48. [Reprinted in his *Meaning, Knowledge, and Reality*, Cambridge, MA: Harvard University Press, 1998.]
- (1993) 'Knowledge by Hearsay', in B. K. Matilal and A. Chakrabati (eds.) *Knowing from Words: Western and Indian Philosophical Analyses of Understanding and Testimony*, Dordrecht: Kluwer. [Reprinted in his *Meaning, Knowledge, and Reality*, Cambridge, MA: Harvard University Press, 1998.]
- (1994) *Mind and World*, Cambridge, MA: Harvard University Press.
- (1995) 'Knowledge and the Internal', *Philosophy and Phenomenological Research* 55: 877–93. [Reprinted in his *Meaning, Knowledge, and Reality*, Cambridge, MA: Harvard University Press, 1998.]
- (1998) 'Having the World in View: Sellars, Kant, and Intentionality', *Journal of Philosophy*, 95: 431–91.
- (this volume) 'The Disjunctive Conception of Experience as Material for a Transcendental Argument'.
- Martin, M. G. F. (2002) 'The Transparency of Experience', *Mind and Language*, 17: 376–425.
- (2004) 'The Limits of Self-Awareness', *Philosophical Studies*, 120: 37–89.

- Martin, M. G. F. (2006) 'On Being Alienated', in T. S. Gendler and J. Hawthorne (eds.) *Perceptual Experience*, Oxford: Oxford University Press.
- Millar, A. (1991) *Reasons and Experience*, Oxford: Clarendon Press.
- Snowdon, P. (1980–1) 'Perception, Vision, and Causation', *Proceedings of the Aristotelian Society*, 81: 175–92.
- (1990–1) 'The Objects of Perceptual Experience', *Proceedings of the Aristotelian Society Supplementary Volume*, 64: 121–50.
- Strawson, P. F. (1979) 'Perception and its Objects' in G. MacDonald (ed.) *Perception and Identity: Essays Presented to A. J. Ayer with his Replies*, London: Macmillan.
- Wright, C. (2002) 'Anti-sceptics Simple and Subtle: G. E. Moore and John McDowell', *Philosophy and Phenomenological Research*, 65: 330–48.

Analytical Table of Contents

PART I: PERCEPTION

1 Hinton and the Origins of Disjunctivism

Paul Snowdon

The aim of the essay is to investigate the central argument of Hinton's book *Experiences*. Hinton is credited with being the originator of disjunctivism, but his own writings are neglected. It is suggested that the idea of disjunctivism pre-dates Hinton's writings, but his achievement was being the first to focus on perception—illusion disjunctions, and to employ that kind of statement in an argument against a certain view of experience. It is argued that his account of them is not open to some objections that have been brought in the past, but it remains, in some respects, questionable. His main employment of them in an argument against the common element hypothesis is also questionable, but it is argued that he is right to object to that hypothesis on one central conception of it. It is hard to determine what his attitude is to that hypothesis on another conception, and that uncertainty represents the main mystery about his view.

2 Either/Or

Alex Byrne and Heather Logue

This essay surveys the varieties of disjunctivism about perceptual experience. Disjunctivism comes in two main flavours, metaphysical and epistemological. Metaphysical disjunctivism is the view usually associated with the disjunctivist label. After some initial discussion of (metaphysical) disjunctivism, epistemological disjunctivism is explained. The rest of the essay is solely concerned with explaining and assessing metaphysical disjunctivism, a theory of the nature of perceptual experience. The main (and provisional) conclusion is that although there is considerable insight in the vicinity, metaphysical disjunctivism is false.

3 Against Disjunctivism

E. J. Lowe

A version of the causal theory of perception that is incompatible with disjunctivism is formulated, argued for, and defended against criticisms typically levelled at such a theory by disjunctivists, such as that it promotes scepticism and that it is unfaithful to the phenomenology of perception. It is then argued that, far from disjunctivism being ontologically less extravagant than that causal theory of perception,

the reverse is true, so that, all things considered, the causal theory of perception is the superior theory and disjunctivism in the philosophy of perception should be rejected.

4 Disjunctivism about Visual Experience

Scott Sturgeon

This essay explores an approach to visual experience known as ‘disjunctivism’. Section 1 lays out a conception of visual experience focused on two charged aspects of the phenomenon—its conscious portrayal of the world, and its perceptual contact with the world. Section 2 lays out the backbone of disjunctivism and relates it to that conception of visual experience. Section 3 lays out a range of disjunctive positions and shows how they share the backbone. Section 4 canvasses major motivations for disjunctivism and shows how they confer a dialectical advantage on one version of the approach. Sections 5–6 sketch that version, explore its details, and show how they can be used to rebut criticism found in the literature. Section 7 presents new worries for the view; and section 8 explores a pure kind of disjunctivism—one which avoids these worries—sketching when it counts as well motivated.

5 Disjunctivism, Indistinguishability, and the Nature of Hallucination

William Fish

Critics of disjunctivism have argued that the disjunctivist needs to provide a plausible explanation of just how two distinct mental states might be indistinguishable for their subject without simply taking the fact of indistinguishability for granted. In the present essay, the everyday notion of indistinguishability is examined alongside relevant empirical findings, and an account of hallucinatory indistinguishability is then offered which suggests that the indistinguishability of hallucination from veridical perception is grounded in the similarity of their effects, rather than in any phenomenal similarities between the two states. This analysis is then elaborated to show how it can account for the first-person aspects of hallucinations, to show how it is possible for conceptually unsophisticated creatures such as animals and infants to hallucinate, and to show how a subject may suffer from hallucinations whilst being aware that their experiences are non-veridical.

6 How to Account for Illusion

Bill Brewer

Early modern empiricists think of perceptual experience as the presentation of an object to a subject. Phenomena of illusion suggest that such objects must be mind-dependent things. Alternatively, perceptual experience may be characterized instead in terms of its representational content (this is the Content View). In that case, illusion is simply false perceptual content. It is argued here that the early modern

empiricists had a key insight: the idea that the core subjective character of perceptual experience is to be given by citing the object presented is more fundamental than any appeal to perceptual content, and can account for illusion, and indeed hallucination, without resorting to the problematic postulation of any mind-independent such objects distinct from the mind-independent physical objects we all know and love (this is the Object View). It is also suggested that the Object View provides a more promising context for the basic commitments of disjunctivism than the current orthodoxy of the Content View.

7 Disjunctivism and Discriminability

A. D. Smith

Disjunctivists typically claim that something is a hallucination if it is an experience that is not a perception, but that the subject cannot tell, just in virtue of having the experience and ‘introspecting’ it, whether it is not a perception. Against this it is argued that there are states that are not hallucinations, as this term is usually understood, but that meet the condition of being thus subjectively indiscriminable from perception. Such states are not hallucinations, since they are not even sensory in character. A modification to the condition is finally introduced that improves matters, but does not entirely avoid counter-examples. In the course of the essay a position that is termed ‘extreme disjunctivism’ is discussed. This position is, it is argued, particularly unfitted to accommodate the counter-examples.

8 The Epistemic Conception of Hallucination

Susanna Siegel

Since disjunctivists about perception deny that hallucinations and veridical perceptions have a common fundamental nature, they need some other way to account for the fact that these kinds of experiences can ‘seem the same’ from the inside. A natural response is to give a purely epistemic account of hallucination, according to which there is nothing more to hallucinations than their indiscriminability from veridical perceptions. It is argued that the epistemic conception of hallucination falters in its treatment of cognitively unsophisticated creatures, and that it cannot respect all the facts about what we can know on the basis of introspection.

PART II: ACTION

9 Disjunctive Theories of Perception and Action

David-Hillel Ruben

In order to state a disjunctive theory of action, one needs two terms of art: ‘intrinsic event’ and ‘mere event’. One also needs two senses of ‘event’: a wide sense and

a narrow sense. In the wide sense, all actions are events by classificatory fiat. In the narrow sense, mere events and intrinsic events are events, but basic actions are events only if they are identical to their intrinsic events. If basic actions were events in the narrow sense, it would not be true by classificatory fiat but by the expenditure of honest philosophical labour. The argument of the essay is that if a basic physical action occurs, no event in the narrow sense occurs, either one with which the action is identical, or one which the action causes, or one which is a part of the action. A disjunctive theory of action claims that, if an item *i* is an event in the broad sense, either *i* is a mere event in the narrow sense or *i* is an intrinsic event in the narrow sense or *i* is an action.

10 A Disjunctive Conception of Acting for Reasons

Jennifer Hornsby

A disjunctivist conception of acting for reasons is introduced by way of showing that an account of acting for reasons should give a place to *knowledge*. This disjunctive conception is claimed to have a role analogous to that of the disjunctive conception that John McDowell recommends in thinking about perception; and it is shown that the two conceptions have work to do in combination when they are treated as counterparts. It is also claimed that the disjunctive conception of acting for reasons safeguards the connection between what moves us to act (sometimes called ‘motivating reasons’) and what favours our acting (sometimes called ‘normative reasons’).

11 On How to Act—Disjunctively

Jonathan Dancy

This essay reconsiders Dancy’s (2000) rejection of a disjunctive account of acting for a reason. It starts with a brief account of the marks of disjunctivism in general, to be used as a template, with special attention to issues raised in Dancy (1995) where it was suggested that the second disjunct of perceptual disjunctivism might be expressed substantially rather than merely as a state indistinguishable from the first disjunct. The essay then considers the motivations behind a disjunctive account of acting for a reason, and also various other ways in which such accounts might be phrased, before concluding that none of them is successful. They are all either empirically false or philosophically challengeable. The suggestions made by Hornsby in her paper for the conference in Glasgow, of which this volume is in part a record, which amount to a disjunctive conception of ‘A φ -ed because she believed that p’, are then considered and it is argued that this account cannot be sustained. In a final section, Martin’s reasons for insisting that the second disjunct in disjunctive accounts of experience must be conceived as whatever is indistinguishable from a perception, and not in more substantial terms, is considered one more time, and rejected.

PART III: KNOWLEDGE

12 McDowellian Neo-Mooreanism*Duncan Pritchard*

It is claimed that McDowell's treatment of scepticism offers a potential way of resurrecting the much-derided 'Moorean' response to scepticism in a fashion that avoids the problems facing classical internalist and externalist construals of neo-Mooreanism. In the course of evaluating the prospects for a McDowellian neo-Mooreanism, further support for the view is offered.

13 In Defence of Disjunctivism*Ram Neta*

McDowell has offered a particular epistemological argument in favour of one version of disjunctivism about perception. This argument has been prominently criticized by Crispin Wright, and the conclusion of the argument has been prominently criticized by Mark Johnston. The present essay rebuts both of those criticisms.

14 Perceptual-Recognitional Abilities and Perceptual Knowledge*Alan Millar*

A conception of recognitional abilities and perceptual-discriminative abilities is deployed to make sense of how perceptual experiences make it possible to have cognitive contact with objects and facts. It is argued that accepting the emerging view does not commit one to the claim that perceptual experiences are essentially relational, as they are conceived to be in disjunctivist theories. The discussion explores some implications for the theory of knowledge in general. In particular, it considers how we can shed light on the nature of knowledge if we do not aim to provide a conceptual analysis of knowledge in terms of true belief plus something else. Consideration is also given to how best to make sense of the practical value that knowledge has.

15 Starting Afresh Disjunctively: Perceptual Engagement with the World*Sonia Sedivy*

It is claimed that a non-objectifying explanatory approach to perception shows how the disjunctive insight works together with conceptualist direct realism to give the best explanation of perception. This emphasizes the role of the disjunctive insight as an integral component within a conceptualist approach that emphasizes non-objectifying explanation. Taking an explanatory approach highlights the

sui generis nature of perceptual content: determinate and object and property involving—or genuinely singular—and conceptual. It shows that perception is a mode of engagement—because perceptual content is determinate and such content involves its objects and properties—and that disengaged states such as dreams or hallucinations are peripheral and limiting conditions of the central condition. Developing this explanation addresses three principal concerns regarding disjunctive, direct realism: (i) A new argument from the fact that perceptual content cannot be retained in its original determinate detail shows that the fabric of perceptual consciousness is relational. (ii) To explain how it is that the mind-independent individuals we perceive have the properties that we perceive—even though the scientific picture might be taken as suggesting otherwise—broadly interpretivist and Wittgensteinian resources are deployed and argued to be tenable alongside scientific explanations of sub-personal perceptual processes. (iii) The explanatory approach taken in this essay also addresses the fundamental worry that disjunctive approaches simply side-step the fundamental sceptical challenge that perceptual experiences might be indistinguishable from veridical experiences.

16 The Disjunctive Conception of Experience as Material for a Transcendental Argument

John McDowell

A diagnostic approach to scepticism about perceptually acquired knowledge of the external world is considered, and it is suggested that this approach can be pursued through a transcendental argument that belongs neither to the ambitious type (associated with P. F. Strawson) nor to the modest type (associated with Barry Stroud). The diagnosis is that this scepticism expresses an inability to make sense of the idea of direct perceptual access to objective facts about the environment. And the argument aims to show that this idea must be intelligible, because it is a necessary condition for it to be intelligible that experience has a characteristic that, for the purposes of the argument, is not in doubt. The disjunctive conception of experience insists upon this idea. Crispin Wright has accused this conception of being “dialectically quite ineffectual” as a response to scepticism; this charge is considered, and rejected.

17 Comment on John McDowell’s ‘The Disjunctive Conception of Experience as Material for a Transcendental Argument’

Crispin Wright

In his contribution to the present volume, John McDowell proposes a distinctive kind of ‘transcendental’ argument for the disjunctive conception of experience, and renews his claim that the latter can be deployed to defuse certain kinds of sceptical doubt, responding *obiter* to the misgivings advanced in Wright (2002) about its credentials for the latter task. This essay queries the putative ‘transcendental’ authority

of disjunctivism, and reinforces the misgivings. It is argued that the root of sceptical doubt has less to do with a 'highest common factor' conception of the commonality between perceptions and illusions than with the possibility of phenomenological matching; and that scepticism can take a direct realist conception of sense experience in its stride once proper heed is given to the gap between direct awareness of a situation and the possession of warrant to believe that it obtains.

PART I
PERCEPTION

1

Hinton and the Origins of Disjunctivism

Paul Snowdon

The use of the term ‘disjunctivism’ has evolved so that people now talk about disjunctivist approaches in, for example, the theory of action, the theory of reasons, the theory of memory, and the theory of knowledge, as well as in the theory of perception. The idea is that in relation to each area (or domain) views are possible which share a single structure. ‘Disjunctivism’ then acts as a name for this recurring structure. There is something careless, or at least misleading, in such talk. It is not that, say, perception itself or knowledge itself is treated disjunctively, but rather that the more general categories which tend to be employed in analysing such phenomena are treated disjunctively. In the mentioned cases they are the notions of experience (or appearance) and of belief. The proposed accounts treat *them* disjunctively. What remains to be properly explained is, of course, what treating such general features disjunctively amounts to. In this essay, however, I shall concentrate on disjunctivism in the philosophy of perception. It has, after all, to be in that domain that the structure (or structures) meriting the name ‘disjunctivism’ is (or are) initially picked out.

Within that domain, though, I want to concentrate on the work of Michael Hinton, who, according to most people who write about disjunctivism, should be counted as the inventor of the idea.¹ Despite being credited as the originator of the idea, and in that role frequently cited, Hinton himself is not widely read or discussed (nor has he ever been). I think that this neglect is both unjustified and regrettable

An early version of this essay was presented at a conference on disjunctivism, organized by Tim Crane and Marcus Willaschek, held in Frankfurt in 2004. I am very grateful to Tim Crane for his comments on the paper presented there, and also to the other participants in the conference, especially, Katalin Farkas, M. G. F. Martin, Howard Robinson, A. D. Smith, and Charles Travis. Some recent comments by two anonymous reviewers and by Adrian Haddock have been helpful. I must also thank Fiona Macpherson and Adrian Haddock for their invitation to contribute to the present book, and their patience in waiting for my contribution. The final revisions were made during my stay in the philosophy department in RSSS at ANU. I am very grateful for the Visiting Fellowship which enabled me to work in the wonderful environment they provide, as I am to UCL for allowing me to be absent.

¹ It needs pointing out that although Hinton is credited with inventing the position called ‘disjunctivism’ he did not invent the name.

(though explicable), and I want to do something to remedy it. It is not, in fact, an automatically merited consequence of being the inventor of a view that subsequent discussion should focus on, or indeed pay particular attention to, one's views. But I do think that Hinton's writings merit consideration because such attention illuminates a large amount in current debates about disjunctivism. I have to confess though that Hinton's work is rich and difficult and in this essay I can attend to only a part of it. I should also confess that some elements in the interpretation of Hinton's work, to which I find myself, after repeated readings of it, inclined, are not quite what I expected when I set out.

In this discussion of his views I want to pay particular attention to Hinton's book *Experiences*, in which he gives the fullest exposition of his thinking, in all its complexity. Hinton's book was published in 1973, and some of the leading ideas in it has been adumbrated by him in earlier articles (primarily Hinton 1967a and Hinton 1967b), but it is my impression that there was no manifest engagement either in the literature or in informal discussions at the time with Hinton's writings until the early 1980s, when two of us, both at that time within the broad philosophical community of Oxford which also included Hinton himself, invoked his ideas to support projects about which it is by no means obvious what Hinton's attitude was or would have been.² I am referring to my own article 'Perception, Vision and Causation', published in 1981, and to John McDowell's paper 'Criteria, Defeasibility and Knowledge', published in 1982. There is no doubt that McDowell's radically 'epistemologized' version of disjunctivism presented both there and in his subsequent highly influential publications can be primarily credited with inspiring the current interest in the idea whose career Hinton had supposedly started earlier.³

1 DISJUNCTIVISM AS A COMMITMENT OF NAÏVE REALISM

It is reasonable, I think, to feel somewhat puzzled by the idea that Hinton invented disjunctivism in the late 1960s. The reason there is a puzzle is that disjunctivism is,

² I want, in this footnote, to provide a few details of Hinton's life. My information comes from the *Worcester College Record* of 2000, which contains two addresses delivered at his funeral, one by his colleague Sabina Lovibond (to whom I am very grateful for providing me with a copy) and one written by his wife Jenny. Hinton was born in 1923, and educated at the famously hardy public school Gordonstoun, where he was, unsurprisingly, unhappy. A contemporary of his there was the Duke of Edinburgh, who on meeting Hinton later in life is reported to have said; "Ah, Hinton, you're supposed to be the clever one." Hinton's character was, in part, shaped by his horrific experiences when, as a pacifist, he was sent, in 1945, to clean up Belsen. He had to drag bodies to pits left by the SS. As his wife says, "This made a deep scar on Michael." It gave him a political seriousness, which expressed itself in a commitment to socialism and CND. After a philosophical education at Cambridge, being a pupil of Wisdom, he became a Fellow at Worcester College Oxford. Hinton published one book and about twenty articles, which it is my impression are primarily about perception, the mind/body problem, and philosophical method. He retired early owing to bad health and died in 2000. He is reputed to have said of himself: "I may seem composed, even dignified—but underneath, I'm Donald Duck."

³ In describing McDowell's disjunctivism as highly 'epistemologized' I am following a reading of his work which I sketched in Snowdon (2004).

in a nutshell, the idea that *apparent* perceptual experiences should be thought of as belonging to two alternate, disjoint classes. They belong, that is, to one or the other case, where one case is the kind of experience that constitutes genuine perception of an item or feature in the percipient's environment, and the other is an experience which seems to be that way but which is actually not that way, and is, rather, an hallucination. We might, therefore, using a somewhat resonant term, describe disjunctivism as experiential *dualism*. Now, the conclusion of most reflection about perception in our philosophical tradition has been that all such experiences have a single nature, usually taken to be the apprehending, or directly perceiving, of some inner item. Consequently it can be said that the standard philosophical approach has been in favour of experiential *monism*. There has, however, always been a tradition of thought in which naïve realism about perception is defended. It seems, though, that this naïve realist tradition must accept the claim that there are, besides the perceptual experiences which can be described in accordance with the naïve realist conception (whatever it is), other sorts of experiences, say hallucinations, which seem to the subject of them to be of the genuinely perceptual sort, but which are certainly *not* of the same sort. This seems to mean that they must be experiential *dualists*, and so disjunctivists, of some sort.⁴

It might seem that this general argument would be strengthened by finding explicit affirmation of such a dualism about experience in the tradition defending naïve realism. In fact this is not easy to find (though I have not gone out of my way to search for it). The reason, I conjecture, is that defenders of naïve realism put their intellectual efforts into denying that the analysis of experience that monists were arguing for was correct for the perceptual case. Having successfully, as they thought, held the line over that, there was no need to say much about the non-perceptual cases. They could, as it were, be analysed as one wished. Such defences of naïve realism were, however, committed to denying the thought that a similarity, to the experiencing subject, between one experience and another, necessarily means that they are the same sort of experience.⁵ That denial is, of course, the (or, certainly, a) crucial claim that a dualist or disjunctivist must make. Here, for example, on one reading at least, is Austin's famous exposition of precisely this point. "If I am told that a lemon is generically different from a piece of soap, do I expect that no piece of soap could look just like a lemon? Why should I?" (Austin 1964: 50). Again, on a certain reading, Austin is making the same point later in these words:

But if we are prepared to admit that there may be, even that there are, *some* cases in which 'delusive and veridical perceptions' really are indistinguishable, does this admission require us to drag in, or even to let in, sense-data? No. For even if we were to make the prior admission (which we have so far found no reason to make) that in 'abnormal' cases we perceive sense-data, we should not be obliged to extend this admission to the 'normal' cases too. For why on earth should it *not* be the case that, in some few instances, perceiving one sort of thing is exactly like perceiving another? (Austin 1964: 52)

⁴ A conversation with Scott Sturgeon a few years ago brought home to me, or at least made me think, that there is a case for being puzzled over this matter, and suggested some of the responses I cite. I have been considering the issue from time to time ever since.

⁵ This denial is an essential element in response to the so-called argument from hallucination.

Austin here looks to be precisely making room for what I have been calling experiential dualism.⁶

Here, next, is Pitcher making a dualist defence in a rather different way. The monist argument is driven on his understanding by something that he calls principle P, which says, roughly, that if two occurrences are physically the same then the experiences they are (or amount to) must be the same sort of experience. Pitcher is quite happy to discuss the issue on the assumption that physicalism is correct. Against principle P Pitcher (1971: 56) simply says:

Nothing in the argument prevents the direct realist from denying that the sense-datum description applies in the normal case: nothing prevents him from asserting that, on the contrary, in the normal case it is the incompatible additional description 'the subject's direct awareness of a (real) red triangle' that properly applies.

There are, it has to be admitted, very important differences here between Austin's attitude and Pitcher's and consequently, no doubt, differences in how to assess their respective proposals, but they exhibit, or seem to exhibit, the same idea of denying the soundness of an inference from a similarity of some sort between two experiences (say that of seeming the same to the subject, or being physically the same) to the conclusion that they are the same sort of experience.

The conclusion to which these very general considerations seem to lead is that disjunctivism must be counted an old idea, which could have been extracted from defenders of naïve realism fairly easily. It seems to represent something that they thought, or at least something that what they thought implies.

Is there really, in this way, a *prima facie* puzzle in the role ascribed to Hinton? It would be right to concede straight off that the puzzle has been generated on the basis of fairly meagre resources. In particular it has invoked a rather unspecific, but still perhaps restricted, conception of disjunctivism. However, although it is relatively impoverished, it should be agreed that it is close to the way disjunctivism is often (if not always) introduced. The puzzle is not, then, dissolved by this observation. Now, my aim, given the rather restricted evidence I am presenting (and, indeed, have available) is, at most, to make a suspicion look reasonable. Is there more that can be said against the suspicion?

One objection to the reasoning is to claim that not all defenders of naïve realism have in fact been disjunctivists. There were, and are, it might be said, defenders of naïve realism who should be counted as *monists* about experience. Thus, some theorists hold that all experiences (of the kind we are concerned with) involve items of some sort or other *appearing* to a subject, the only difference is that in some cases these items are external objects, and then we have regular perceptual cases, whereas

⁶ Austin's attitude to the philosophy of perception, that the correct response to it is one of "unpicking, one by one, a mass of seductive (mainly verbal) fallacies . . . an operation which leaves us, in a sense, just where we began" (Austin 1964: 5–6), makes it seriously misleading to classify him as naïve realist. His official attitude is that he is not proposing a theory. Austin is not a pre-Hintonian disjunctivist, but, one might say, one of the fallacies he exposes (or supposes he exposes) is that which holds monism in place.

in others the items are internal objects, call them images, and then we have hallucinations and other non-perceptual experiences. Such theorists are, therefore, monists about experience but also defenders of naïve realism.

Two points can be raised about this response. The first question is whether the uniformity amongst the characterization of all experience which such theories propose is actually inconsistent with describing the view as a version of experiential dualism (and hence disjunctivism). It is not at all obvious that the claim that we can characterize the two sorts of case as being the same at *some level* makes it incorrect to call the view disjunctivism. Although, in one sense, in the two cases the same thing is going on, namely items appearing, in another sense it is *not* the same that is going on, because the appearing items have a quite different nature and stand in quite different relations to the subject. I do not see, therefore, that such an approach should not be counted as disjunctivism. The second point to make is that the existence of the puzzle does not really depend on the claim that every version of naïve realism must count as committed to disjunctivism, but rather that certainly some do and they pre-date the late 1960s. This means that it is not, in a sense, crucial how we adjudicate the first point.

Although I am, somewhat tentatively, refusing to count the view just described as a genuinely monistic account, there is certainly one way which has been taken by its proponents as being a defence of naïve realism that is purely monistic and not disjunctivist at all, and that is idealism of the sort propounded by Berkeley and his followers. Berkeley precisely offers an idealist account of spatial objects in order to combine *monism* about experience with *naïve realism*. That approach does not, indeed, fit my general claim, which should be restricted to naïve realism which incorporates *genuine realism about the spatial world*.

So far, then, it seems that disjunctivism must precede the late 1960s. This conclusion is not, I think, modified, but an important aspect of the issue is brought out, when we stress a distinction not properly highlighted so far. My discussion of this distinction will be both intuitive and un-nuanced, but I hope that it has the ring of truth. When I have talked about experiential dualism I have explained that as the claim that there are two types of experience (which are considered simply in themselves indistinguishable to the subject). What, however, counts as a type of experience? There are basically two different ways to explain types here. The first way corresponds to what Pitcher says in the earlier quotation. Thus he distinguishes between one type of experience and another in terms of a difference in their causes. His thought is that an experience belongs to the type known as direct perceptions of the external world in virtue of it, the experience, being caused in a suitable way. Other experiences count as belonging to another type in virtue of having other causes. We might call this an external principle for typing. In contrast we can understand the idea of what we might call 'internal' typing. Thus, considered in itself, maybe one type of experience has a fundamentally different nature or make-up to another type of experience. Perhaps, one type consists of a relation to an external object whereas another type consists of something different, and not such a relation. This is, and will remain, an intuitive contrast, but it seems to me to make sense. Now, it further seems to me that when people talk of disjunctivism postulating two types of experiences they have in mind a

difference *internally considered*. If this is right it follows that a defence of direct realism along the lines of Pitcher does not amount to an endorsement of experiential dualism *in the appropriate sense*. With this more restricted understanding of experiential types in place it follows that the existence of a tradition of naïve realists who endorse different types of experience in itself need not amount to an anticipation of the view. This observation does not dissolve the problem because it seems to me quite fair to suggest that it was not common for philosophers before the 1960s who thought of themselves as naïve realists to suppose that the relation between subject and object which amounted to the subject's directly perceiving the object was a matter of a distinctive causal source for an experience. Such philosophers seem committed to a dualism of experience in the way that disjunctivism itself endorses.

The interim conclusion is that the puzzle we started with remains. But that is very much an interim conclusion and more that is relevant to it will, I hope, emerge from a detailed engagement with what Hinton himself says, to which it is high time we turn.

2 HINTON'S DISJUNCTIONS

Why, then, are Hinton's writings treated as the source of an idea which taken one way there is a case for thinking probably has to have been around much longer? To approach an answer to this I want to look at Part II of his book *Experiences*. In Part I of the book Hinton elucidates what he thinks of as the standard notion (or notions) of an experience. You will, for example, have the experience of falling down some stairs if that sort of event should befall you while awake and conscious. Although his observations are interesting and acute I shall not engage with them because Hinton's purpose is to move on from this notion to a more philosophical, or as he says 'a very special', notion of experience, and that special notion becomes the focus of intense critical scrutiny in Part II.

Hinton begins Part II with the following resonant remark.

Even if few things are certain, it is certain that there are what I shall call perception–illusion disjunctions: sentences or statements like 'Macbeth perceives a dagger or is having that illusion', which you can compose by adding words like '... or x is having that illusion' to a sentence which says that a particular person, x, perceives a thing of some particular kind. (Hinton 1973: 37)

The presence and prominence of this resounding sentence provides the, or a part of the, answer to the question I have just asked. What is special about Hinton and disjunctivism? The answer is that Hinton was (as far as I know) the first person to demarcate perception–illusion disjunctions as a sort of sentence or statement, and to investigate the logical, and to some extent the epistemological, properties of such sentences. He was the first person, that is, to look at those particular sentential beasts under a philosophical microscope. The study of perception–illusion disjunctions as a sort of statement starts with Hinton. Saying this is not to say why (or even that) perception–illusion disjunctions are important objects of study in the philosophy of perception, nor even why Hinton thinks they are. In fact, and, as one might say, in

a somewhat characteristic way, Hinton, when he introduces perception–illusion disjunctions, does not attempt to say what *role* his consideration of such sentences will have for the overall case he is developing. (From now on, unless quoting Hinton himself, I shall abbreviate ‘perception–illusion disjunction’ to ‘P–I D’.)

Hinton’s opening remark which has just been quoted merits three comments. (1) When Hinton says that “it is certain” that there “are” P–I Ds he cannot mean that it is certain that there are such disjunctions in the sense that they are formulated and considered by people. That is hardly certain. He must mean, rather, that it is certain there are such sentences in the language. That is indeed not a matter for doubt. (2) It would be a complete misunderstanding of Hinton to think that the reference here to certainty reflects the thought that P–I Ds should be counted as specially certain claims playing a crucial role in human cognition. In the first place, what Hinton claims is certain is that there *are* P–I Ds, not that what P–I Ds claim is certain. Second, Hinton is himself manifestly not engaging with epistemology. As Hinton (1973: 45) says; “[T]o start a new ‘theory of knowledge’, according to which everything is based on perception–illusion disjunctions, is far from being my wish”. (3) The specification of P–I Ds in the quoted passage is not particularly precise. Does Hinton mean to exclude such a sentence as ‘there is someone such that either he is seeing an O or is having that illusion’? The problem is that this is in no clear sense a sentence which says something ‘about a particular person, x’. Again, is ‘Blair is seeing my favourite book or is having that illusion’ a P–I D? The problem here is whether the kind ‘my favourite book’, not being visually determinate in itself, is of the right sort. Also, is ‘Blair saw an orange or was having that illusion’ a P–I D? The problem here is the tense of the main verb. Hinton’s examples are all in the present tense. This lack of precision is no cause for great concern, but it should, perhaps, be registered.

In considering Hinton’s treatment of P–I Ds in *Experiences* there are two features that he explicitly builds into the notion that deserve stressing. One is to do with the perception side, the other to do with the illusion side. The first is that, as he puts it, the word ‘see’ when it is used on the perception side has the sense of ‘plainly sees’ (Hinton 1973: 42). What does that mean? Hinton does not mean, of course, with his talk of plainly seeing that it is plain (or as one might say, obvious) to us or to the percipient *that* he or she sees an O. He means rather that the use of ‘see’ has three features. ‘S plainly sees (an) O’ entails ‘there is an O that S sees’, but it also has to be true that, as Hinton puts it, “the unambiguous testimony of his visual sense is of O”, or as he more tentatively puts it, ‘it entails seeming to see [an] O’ (Hinton 1973: 32).⁷ Hinton adds that plainly seeing an O does not require that the percipient takes or believes what is seen to be an O. Now, there are certainly questions to be asked as to what precisely Hinton means by his explanation (and I shall return to that), but these questions do not interfere with the significance I wish to claim for how he characterizes P–I Ds. The second important aspect that Hinton builds in is that the illusion disjunct is to be taken as claiming that the illusion is of the same sort of experience and apparent object as the perception disjunct reports.

⁷ These quotations are slightly inaccurate, in that I have substituted ‘O’ as the term for the object in place of ‘X’ which is what Hinton uses. This is, of course, insignificant.

These two aspects about which Hinton is explicit in his book mean that two criticisms I made of him in my 1980–1 article are mistaken. There I wrote as if Hinton's disjunctions were: either S sees an O or S is having the illusion of an O. As a candidate close to the claim that it (visually) appears as if there is an O this disjunction fails because an O can be seen without looking like an O, and, further, an illusion of an O need not be a *visual* illusion of looking to be an O. In the light of these criticisms the disjunction that I suggested as avoiding the problems was: either there is something which looks (to S) to be an O or it is to S as if there is something which looks to S to be an O (Snowdon 1980–1: 202). Now, this criticism depended on two things. The first is the assumption that Hinton wanted it to be a property of the disjunctions he was stipulating and examining that they were more or less equivalent to a looks-judgement such as 'It looks to S as if there is an O'. The second assumption is that the disjunctions I was assessing were the ones he had in mind. The problem with the criticism concerns not the first assumption, which is, surely, a fair one, but the second one. It is clear that Hinton's understanding of the content of each disjunct in his P–I Ds makes the double objection unsound. If the seeing of an O has to be a what Hinton calls a 'plain' seeing then it is ruled out that the O might not appear *as an O*, and if the illusion disjunct has to claim that there is an illusion as of the same sort of experience then it could not be a different sort of illusion.⁸

Hinton (1973: 61) does later point out that his understanding of P–I Ds permits the description or specification of what is seen in terms of how the thing looks and of the corresponding illusion as of something which looks that way. Thus, Hinton allows as versions of his P–I Ds such a sentence as 'S sees something which looks red or is having that illusion'. This means that Hinton allows disjunctions in which appearance or looks claims are embedded in the left-hand disjunct, as the ones I suggested also involve, but with the difference that the perception disjunction contains a perceptual verb and that the other disjunct contains the notion of illusion.

Now, although Hinton's understanding of P–I Ds means that the two reasons I had for thinking that he was focusing on the wrong disjunctions are faulty, it does not at all follow that the ones he focuses on are unproblematic for his purposes. There are two initial observations. (1) Given that Hinton asks us to read the 'see O' in the left disjunct as 'plainly sees O' and given that that is explained as entailing that the object seems to be O then it cannot, it seems, be Hinton's intention to offer the disjunctions as analyses of 'looks' or 'seems' statements. Given his own explanation of their content that would be circular. (2) It seems to me that Hinton's understanding of P–I Ds has the consequence that in plenty of cases of genuine perception no disjunction of his sort will be true. The problem is that the left-hand side can be true only if it is *plain* what is seen and the right-hand side can be true only if the subject is under

⁸ My unfairness to Hinton arose because my formulation of the disjunctions he was interested in came from his earlier (1967a) article 'Visual Experiences'. There he concentrates on the example 'Either I see a flash of light or I have an illusion of a flash of light'. My criticisms do apply to that if it is suggested as (more or less) equivalent to 'It looks to me as if there is a flash of light'. Since it was this earlier formulation which attracted my attention and excited me originally and about which I had tended to think I quite failed to notice the modification presented in *Experiences* (Hinton 1973).

the illusion of a *plain* seeing. Surely, though, there are cases of perception where neither disjunct applies, namely where the perception is in some way not plain nor does it seem plain. Imagine the experience of a short-sighted person without glasses when an object zooms across his field of vision. Nothing was plainly seen, or seemed to be.

Are these serious problems for Hinton? Problem (1) is a genuine problem only if Hinton is proposing to analyse 'seems' or 'appearance' sentences in terms of P–I Ds in a sense of 'analysis' according to which analysis can only go in one direction. In that case the fact that he clarifies the meaning of 'plain seeing' in terms of 'seems' would rule out clarifying the meaning of 'seems' in terms of his P–I Ds. In fact, it is quite unclear that in *Experiences* Hinton is proposing any such analysis. Hinton is extremely cautious and seems to prefer advancing analytic claims in terms of a relation which he expresses as "rough synonymy or meaning-similarity" (Hinton 1973: 128). He argues at length that the claim "'A' is roughly synonymous to 'B'", said in illumination of 'A', is consistent with its being correct to say "'B' is roughly synonymous with 'A'", said in illumination of 'B' (Hinton 1973: 126–44). It can, therefore, be said that Hinton's own conception of any analytic role for P–I Ds is one that he tries to isolate from circularity restrictions. I do not, therefore, want to pursue this problem any further.

Problem (2) raises two issues. The first is whether it is correct to claim that Hinton's P–I Ds, because of the involvement of 'plain' seeing (or perception), does not cover all cases. The second is whether, if they do not, that matters. Now, since it does not seem that all perceptual experiences must be 'plain perceptions' and it does not seem that all illusions must be illusions of 'plain perception', Hinton's P–I Ds do not exhaust the range of cases. This means that if our aim as theorists is to say something about the whole range of perceptual and apparently perceptual experiences Hinton's disjunctions are not adequate for that purpose. On the face of it this is a serious restriction. It is not clear, though, that Hinton himself had such a goal, nor is it clear precisely what role P–I Ds have for him. So it is very hard to estimate at this stage whether the restriction is serious for Hinton. To pursue this, but also to provide more details of Hinton's arguments, I want next to describe its overall structure, as I see it, and then to scrutinize it.

3 THE OVERALL STRUCTURE OF HINTON'S DISCUSSION

Having introduced and initially pinned down, to his own satisfaction, the class of disjunctive sentences on which he wishes to focus, Hinton's overall discussion has the following three-stage structure. (i) In section IIa he introduces, and to some extent, categorizes, P–I Ds, and then advances various claims about such sentences, pointing out things that he thinks can be said about them and some things that cannot be said. For example, Hinton notes (or claims) that a P–I D of the form 'S sees an O or has that illusion' does not entail that S believes there is an O. It should be noted, though, that in this section the role of P–I Ds is simply that of providing the subject matter for certain positive or negative classificatory claims. About this section two main questions can be raised. The first is whether Hinton defines a clear enough notion of

P–I Ds. The second is whether the classifications he proposes are correct. (ii) Next, in section IIb, which is on my reading, the argumentative heart of Hinton's book, he identifies a view about experience, the 'special' view to which he has alluded earlier. This is the idea of a type of experience which is "supposed to occur when you perceive a flash of light and when you have the illusion of seeing one" (Hinton 1973: 62). Hinton links this idea to talk of the experience which can be shared as an 'inner experience', and he also calls it, very helpfully, the idea of a 'common element' in perception and illusion. Now, Hinton seems to argue that this is at least *not a view that we should accept*. This way of putting Hinton's main thesis is carefully chosen. It is intended to leave it open whether Hinton wishes to claim that idea is actually false. Hinton thinks that the common element hypothesis can be held in different ways. He employs the notion of P–I Ds in an argument against one version, and also objects to the other that it has no support. This extremely complex section raises two questions. The first is interpretative: what exactly is Hinton saying and how is he defending it? The second is whether his arguments are convincing and solid. It can be pointed out that Hinton nowhere explains how the assessment of the common element hypothesis links with other issues. He does not allege that its adoption would make a philosophical account of perception or of knowledge more difficult. He provides us with no reason to be suspicious of the common element view other than the reasons that its direct investigation generates. (iii) In IIc Hinton returns to P–I Ds and does two things: he first says things about perception ascriptions and illusion ascriptions separately, and second, as I have already mentioned, tries to counter some assumptions about conceptual analysis which might lend support to the common element thesis. I shall not discuss this section here, except to note that when he returns to the topics of perception and of illusion Hinton adopts a resolutely conceptual focus. For example, in discussing perception his dominant goal is to persuade us that there is a reading of 'S sees that P' which does not entail 'S believes P'. This, it seems to me, defies the natural expectation one has when reaching this stage of the argument. The expectation is surely that Hinton will say more about the nature of the different sorts of experiences that the rejection of the common element view involves. Hinton does not travel in that direction.

I have briefly indicated the three stages or parts of Hinton's argument. I want to say something about the first stage before looking at the crucial second stage.

4 SOME COMMENTS ON STAGE (I)

There are five comments I wish to make about the first stage of Hinton's discussion. Some of my comments express agreement with Hinton, whereas others are, in effect, suggestions that Hinton was some way from accurately describing P–I Ds (in his sense).

(a) In debates since Hinton wrote it has sometimes been observed that when we are thinking about apparent (but not necessarily actual) perceptual experiences we have three categories we can apply. The first is that of accurate or veridical perception,

where for example, I see a cup and it looks or seems to be what it is; second, there are illusions where, for example, I see a cup which looks different from how it is; the third is that of hallucination, where nothing is actually perceived but it seems as if something is. If there are these three cases and we have an interest in formulating disjunctions involving two disjuncts which we want to cover the whole range of cases it seems that we need to align two of the cases together. But which? Hinton in effect brackets together the second two. It might therefore seem strange and a problem for his procedure that he uses the term 'illusion' to cover both of them since I have just employed it as the way to pick out the *middle category*. However, Hinton's approach here is, I believe, quite defensible. Hinton uses the example of an electric current being passed through someone's head, which produces a phosphene or a light appearance. Now, this is not an example of an illusion where that term is understood in the sense picked out above. But we are, as English speakers, I think, quite happy to describe the result as *the illusion of seeing a light flash*. Hinton then seems correct to employ the term 'illusion' to cover both cases. I do not mean that Hinton's choice of disjuncts is appropriate, but that he can express *his* second disjunct in the way he does.⁹

(b) Hinton attaches importance to a distinction between two sorts of P-I Ds, which he calls respectively 'pointed' and 'blunt'. A disjunction seem to count as pointed if the first disjunct contains after the noun the words 'of a certain sort'. For example, 'I see a light flash of a certain sort or ...', whereas 'I see a flash of light or ...', and 'I see a flash of bluish jagged light or ...' are blunt. I conjecture that Hinton chose the name 'pointed' here because adding 'of a certain sort' can be described as pointing at but not explicitly saying what the further features are. (I do not know why he chose the term 'blunt' to act as a contrast.) Why, though, does Hinton attach importance to this contrast? The best answer that I can provide to that question is that he believed that blunt perception-illusion disjunctions were genuine disjunctions, and so they can be represented formally as ' $P \vee Q$ ', whereas he seems to have thought that pointed ones were not (Hinton 1973: 40, 125). It seems to me that pointed perception-illusion disjunctions are logically complex because the logical treatment of '... of a certain sort', is complex. But if one can deal with the sentence 'S plainly saw an O of a certain sort' on its own, which it seems we can, then it seems that forming a P-I D from it would not be any the less a disjunction than the blunt case. Indeed there is a reason to suspect that *neither* is logically a pure disjunction, in that the second disjunct is always incomplete. The second disjunct involves reference back to the first when it talks of 'having the illusion of doing so'. Further, even if pointed disjunctions have an extra logical complexity it remains quite unclear to me why Hinton felt it necessary to grapple with it. Why should we be so interested in it?

⁹ A. D. Smith raised at a conference in Frankfurt the problem that Hinton's P-I Ds are not appropriately expressed because the illusion disjunct would apply to someone who, as we say, is simply under the illusion that he or she saw an O. An illusion in this context is simply a false belief, and so, the P-I Ds would cease to be essentially experience ascriptions. The answer, I think, is that in interpreting Hinton P-I Ds we need to register the distinction between having an illusion of seeing an O and being under the illusion of seeing an O. The former is an experience ascription and not a belief ascription.

(c) Of P–I Ds Hinton wishes to claim that they are *not* what he calls narrowly doxastic, that is, they do not entail that the subject to whom they apply believes that he sees an O (or that there is an O). He does claim that they are what he calls broadly doxastic, and that means that they entail that there is a *temptation* to believe that an O is seen (or that there is an O) (Hinton 1973: 41, 104–12, 113).

Now Hinton is right in the negative claim here: the truth of a P–I D does not entail the presence of a belief to the effect that an O is seen, nor that there is an O. But does it entail the presence of a temptation? There are three problems. (i) It is hard to know what exactly a belief temptation is, so assessing the claim that a P–I D entails the presence of such a thing is difficult to assess or validate. However, as far as interpreting it is concerned, I regard a temptation to believe something as a positive thing, which may be resistible. A temptation to believe is not legitimately ascribable in a certain context simply if it is true that if things had been different there *would have been* a belief. This is simply the application of a general point about temptation. John might have a temptation at *t* to eat a cake as he passes it, but he would not count as having such a temptation at *t* simply because it is true that if he had taken a different route at *t* he would have seen (and desired) the cake. (ii) The temptation is supposed to be to the perceptual belief (that is, that *S* sees an O). But it seems that perceivers who do form their beliefs, and temptations to believe, in response to such experiences, need not themselves be thinking in terms of the notion of seeing. There is no reason to suppose that they must have the concept of the perceptual state itself. So I do not think it is true that there has to be a temptation to move into *that* cognitive state. (iii) It really does not seem odd to me to believe that it is possible to be the subject of a P–I D without in any way being tempted to believe either that one is seeing an O or that there is an O. This is because there seem to be states in which one might be when the P–I D applies which mean that no temptation at all arises. Thus, one might be hypnotized, very confused, pathologically cautious, or mad. It seems to me better to say simply that it will be *likely* that the subject of a P–I D will believe that there is an O or, almost as likely, that an O is seen. I am therefore tempted to say that, strictly speaking, P–I Ds are *not* even broadly doxastic.

(d) Hinton (1973: 45) says of P–I Ds (or at least of most P–I Ds that he demarcates) that they are “things that seem oddly prefixed by ‘I am inclined to believe that . . .’”. In fact, although Hinton is not explicit about this, he is operating at this point with a restricted form of P–I D, namely those which are *first personal*. There is nothing odd in saying that one is inclined to believe that either Fred is seeing an O or having that illusion. There is, though, some inclination to agree with Hinton about the existence of this oddity in the first personal case.

There is though a difficulty which Hinton does not consider. The problem with what he says arises because I might not be absolutely sure what an O looks like. For example, I might be looking at a flag which is plainly in view and the colour and shape of which is quite clear. Is it the Swiss flag or the Swedish flag? Suppose that I am not sure, but incline to thinking it the Swiss flag. Then I can perfectly intelligibly say that I am inclined to believe that either I am seeing a Swiss flag or having the illusion of doing so. Now, the reply that Hinton would have to make here is that my

uncertainty is inconsistent with the ‘plainness’ of the seeing built into P–I Ds. But this is not clearly correct. Consider the following. I see the flag, at the time not being sure which country’s flag it is. I then learn which it is. Cannot I say subsequently that I was there and then plainly seeing the Swiss flag, even though at the time I did not realize that? Indeed, someone could quite correctly say about my very first sight of the Swiss flag (when I was not sure what flag it was, although I suspected that it was the Swiss flag): ‘You cannot see a flag plainer than that.’ This is the uncertainty about what Hinton means by plainness at which I earlier gestured. There are, it seems, two possible notions that he might have intended. The first notion of plain seeing is this: it is to see O in such a way that *one could not get a better look at O*. Thus, to see the flag plainly in this sense one must be able to see what colours are there and what their layout is in a way that cannot be improved. If I call this ‘plainly₁ seeing’ then I think it is clear that one can plainly₁ see an O without being sure that it is an O, since one is unsure what an O looks like. This in fact seems to me to be the normal meaning of ‘plainly seeing’. The second notion of ‘plainly seeing’, call it ‘plainly₂ seeing’, would be defined as the conjunction of the first notion plus the requirement that the subject takes it, as we might say, that if he is plainly₁ seeing anything then it is an O he is seeing. If plainly seeing is plainly₂ seeing then the claim by Hinton that I am analysing in this section would be correct. Hinton in fact is not clear enough about the ‘plainness’ requirement for this issue to be decided. I think, therefore, it remains doubtful that Hinton’s claim is correct.

(e) Hinton divides his consideration of what can be said about P–I Ds into two sorts, called respectively, “some descriptions that apply . . .” (Hinton 1973: 38) and “some descriptions that do not apply . . .” (Hinton 1973: 46). This is an artificial distinction in that if ‘des’ stands for a description that does *not* apply then ‘not-des’ stands for a description that *does* apply. However, the description that it is most important as far as Hinton is concerned to realize does not apply is that of “answering the question as to what is happening to the subject” (Hinton 1973: 46). Hinton is extremely cautious and careful in his discussion but I think that it is fairly obvious that he is correct in his claim. The basic reason for thinking that he is right is that by “answering the question as to what is happening to the subject” Hinton means saying in some specific or precise or exact way what is going on in the subject given that the subject is either perceiving or having an illusion. Now, the disjunction obviously does not do *that*. Typically, Hinton provides no indication why this fact is worth noting. As we shall see, there are questions about that.

Of the comments I have just advanced two represent agreements with Hinton, two minor disagreements, and only the problem about plainness is at all serious.

5 HINTON’S TARGET AND HIS OBJECTION TO R-REPORTS

In part IIa, Hinton has identified P–I Ds and ascribed certain properties to them. He offers no explanation at all as to why such types of propositions are interesting. However, in part IIb Hinton identifies a *philosophical thesis* which it is, or appears to be, his

main aim to oppose. This thesis says that there is a type of experience which is present both in perception and in what Hinton is calling illusions. The name he uses to pick out what the theory believes in is 'The Common Element in Perception and Illusion'. So we might say that he is opposed to what I shall call the *Common Element Hypothesis* (abbreviated hereafter to CEH). This gives us a fix on disjunctivism, at least in so far as it is intended to be a position deriving from Hinton. It does so, of course, only if we can work out what CEH is supposed to be. Viewing it this way encourages us to think of Hinton's disjunctivism as an essentially negative thesis; it can be read precisely as the denial of CEH. This would mean that it has the status which for a time it was popular to ascribe to cognitivism in ethics. That was said to be best understood as anti-non-cognitivism. Its content was read as a denial of the positive thesis of non-cognitivism. We might, then, similarly read disjunctivism as *anti-CEH-ism*. I myself am not recommending this way of understanding the debate, but it seems to be the way Hinton viewed it.

Assuming for the moment that this is more or less accurate as an account of Hinton's understanding of the debate, two fundamental questions arise: (a) What exactly is CEH? (We cannot understand the course of Hinton's argument unless we can say.) (b) Does Hinton persuasively show that the CEH is false (or is, perhaps, *not* to be accepted)?

At this point, however, we have located *one* role within the major debate for P–I Ds, something about which, as I have indicated, Hinton is not very explicit. A P–I D states the two alternatives which, *according to CEH*, share a common element. So CEH itself is a thesis about the states of affairs which obtain when a P–I D is true. This of course merely gives P–I Ds a role in fixing the *range of a thesis that disjunctivism denies*. Further, it can also be said that there are equivalent ways of fixing the range of cases that CEH covers without employing P–I Ds.¹⁰

Hinton (1973: 62) is aware, though, that supporters of CEH might, for all that their affirmation of CEH says, hold that the so-called common element present whenever a P–I D is true could itself well be present also on other occasions, occasions, that is, when a P–I D is not true. One possibility that Hinton might have in mind is a theorist who thinks that there is a common experiential element but who thinks that what is called 'perception' and what is called 'illusion', that is to say, what is imported by those descriptions, necessarily involves some other ingredients, say concept application, which need not occur simply because the common experiential element is present. Hinton is right to allow this. He is also right to allow that the supporter of CEH need not say that it *is* the common element which *is* the perception when it is a case of perception or *is* the illusion when it is a case of illusion (Hinton 1973: 62–3). Maybe the common experiential element is not what counts as the perception or illusion but is rather an element involved in there being a perception or an illusion. We should, I think, admire Hinton's subtlety in avoiding unnecessary restrictive assumptions in spelling out CEH. (A theme in Hinton's writings generally, of which these remarks are one relatively minor case, is the care he took over deciding

¹⁰ You could say, for example, that the CEH states that there is a certain type of experience in common to cases of perception *and* to cases of illusion.

when identity claims are appropriate ways to formulate philosophical theses.) But what, then according to Hinton is CEH?

Hinton's initial approach to the formulation of CEH is to introduce the idea of the existence of certain sorts of *reports*. As he puts it, "the question will be whether there is such a thing as a report of such an experience" (Hinton 1973: 63). By "such an experience" Hinton means such as the CEH supposes there is. Hinton's idea, therefore, is, or seems to be, to explicate CEH in terms of the existence of certain types of *reports*. He initially specifies what he calls *R-statements* (Hinton 1973: 63–6) and then he introduces a weaker notion of *E-reports* (Hinton 1973: 74). As far as I can tell, his assumption is that the existence of R-reports would be a sufficient condition for the truth of CEH, but their existence is not necessary for its truth. In contrast the existence of E-reports is both necessary and sufficient for CEH (Hinton 1973: 100). Hinton's strategy is then to deny or query the *existence* of such reports. On encountering this argumentative strategy it would be natural to respond by wondering why, if the existence of E-reports is what CEH is equivalent to, we should consider R-reports. I hope that some sort of answer to this will emerge in the course of this discussion.

I have introduced Hinton's R-reports and E-reports terminology, but I need to indicate what his conception of these two sorts of reports is.¹¹ I shall begin with the R-report category. The conditions for a statement's qualifying as an R-report are, more or less, the following (Hinton 1973: 63–4). (I shall use the numbering that Hinton employs for the separate conditions.) R1: the statements must be true when a P–I D is true. R2: the statement must answer the question 'what happened to the subject?' It must, as he sometimes puts it, give the 'what-it-is of the event' in question. R3: it must be a statement which it is incongruous for a subject to preface by the words 'I am inclined to believe that . . .'. R4: it must report an experience. R5: the subject must be able, without relying 'on undischarged physical assumptions', to determine that it meets R2. Hinton (1973: 65) adds: it must *not* be a hypothesis or conjecture that it was true when affirmed by the subject.

The idea of an R-report seems intended to be something that is, or is close to being, a genuine experience report, which is to say something that a normal subject can determine as true simply in virtue of having an experience (an example of such a report being 'I am in pain'). It has been quite common for philosophers to stress the contrast between reports and non-reports as types of judgements. Some of Hinton's conditions seem to be attempts to spell this out. However, as well as being a report, the sentence must specify or capture the nature, in some sense, of the happening which is the same whichever disjunct in the true P–ID obtains. Hinton resorts to talk of the 'what-it-is' of the occurrence, which is reminiscent of Wiggins' talk of 'answering the "what is it?" question' as a way to fix on sortal predication. In line with this, one might say that Hinton's idea of an R-report is that of a *report of an experience sort*.

Later Hinton introduces the weaker notion of an E-report. The conditions for being such a statement are—R1: that it be true whenever a P–I D is true; and, basically, R2: that it specifies the nature of an experience common to both cases.

¹¹ I am grateful to Tim Crane for suggesting that this is the place to explain somewhat these technical terms.

Although Hinton still talks of reports the idea has been completely dropped that an E-report is in any normal sense a report of an experience. (I have not completely adhered to Hinton's language here, but I hope that it does not affect things.)

Hinton, then, transforms CEH into a thesis claiming that certain reports exist, and construes his own arguments against CEH as arguments against the existence of reports of such a character. Even before scrutinizing his arguments there are reasons for being at the very least puzzled by Hinton's way of developing the debate.

One worry is that there is a viewpoint about experience which Hinton's strategy overlooks. It is not an unknown thought that experiences are ineffable, inexpressible, and indescribable. According to such a view, although there are experiences there are no experience reports (even if there seem to be). It appears to me coherent to hold both CEH and the inexpressibility thesis. Such a theorist claims there is a common sort of experience across the range of cases even though the sort is not describable or reportable. Of course, anyone expressing this thesis must allow that some things can be said about experiences, but it need not be *very much*. This possibility, if it is a possibility, seems to show that Hinton's formulation in terms of the existence of reports is not strictly equivalent to CEH.

Hinton (1973: 100) does pick up on this problem but responds briefly by expressing scepticism about the inexpressibility claim. We can add, I think, that the philosophers Hinton was targeting, in so far as we can conjecture who they are, are not themselves in any way inclined to endorse the inexpressibility thesis. So, although Hinton's procedure is not strictly necessitated by the hypothesis he is criticizing, it would be excessive to attach importance to this fact.

There is, however, another problem to voice. The worry might be put this way. CEH is fundamentally an existential hypothesis, and it is not normally the case that existential hypotheses about real things can be sensibly tied to the existence of reports or statements with a certain character in our language. Thus, if I advance the claim that there are adders in both Oxfordshire and Yorkshire no one would debate that in terms of the existence of reports in English with a certain character. Why then does Hinton take the issue along these apparently linguistic lines?

In response to the question I can think of two things that might be said. First, it might be pointed out that CEH is an existential hypothesis about *experiences* (and not about types of snakes). Now, the notion of experience is an extremely general one in its application. We ascribe experiences to normal adult humans, but also to human infants, other animals, and also to abnormal humans. But it is quite plausible and standard to tie its application in the normal *adult* case to the idea of a *report*, in a strong sense of 'report' that I have already alluded to and shall simply take for granted. To ascribe an experience is to ascribe something which a normal adult must be capable of reporting. Suitably restricted, we might say that experiences are, essentially, *reportables*. If Hinton subscribed to this slogan then he might have thought that a thesis about a common element which is an *experience* should be formulated in terms of *reports* of a certain character.¹²

¹² One place in Hinton's discussion where the link between experience and reports seems to surface is when he says, in elucidating a possible remark that the effect of a certain cause is an

If Hinton's strategy is based on this reasoning then I think that it is based on a fallacy. When there is talk of a report, as there is here, the idea is of a report that the subject him- or herself can make. But the sense in which experiences are reportables is that the occurrence of an experience can be reported *in certain respects* by the subject, not that every truth about an experience need be so reportable. Thus if someone proposes that each experience is a physical happening then although it is a claim about experiences it does not imply the existence of reports (in the strong sense) of experiences *corresponding to this feature*. Rather the form of the thesis is this: concerning experiences, which are therefore in some respects reportables, they are all physical. Similarly there seems no reason given by the fact that experiences are reportables and the claim is that there is a common sort of experience across a certain range that there *must* be distinctive reports corresponding to that fact about experience. So there is here really no ground for Hinton's dialectical strategy.

I have suggested that one *possible* ground for the temptation that Hinton seems to have felt to formulate the debate in terms of the existence of reports was the link between experiences and reports. But I want to suggest a second and rather simpler reason that might have influenced Hinton. This reason starts from the question: who are the supporters of CEH that Hinton supposed himself engaging with? Hinton (1973: 66) says something that is relevant to this in the following important passage:

it was widely held that the existence of such things as sense data, . . . could be established by familiar arguments. These would indirectly establish that there were statements about such entities The present philosophical and psychological situation is one of widespread scepticism about such arguments and such entities Yet unless I am quite mistaken it is now often believed, and more often taken for granted than argued, that there are statements which meet all these requirements without being about such entities.

This passage reveals Hinton's conviction that many philosophers just do accept that there are R-reports, though not on the basis of arguments similar to those for the existence of sense-data, which they also reject. So, naturally, he first considers this conviction, hence his focus on R-reports.

How, then, does Hinton cast doubt on the reality of R-reports? Hinton's discussion is extremely rich and cautious but, it seems to me that his main argument, or thought, can be represented as involving three premises. (1) P-I Ds are *not* R-reports, because they do not satisfy condition R2. Being disjunctive they merely say that one or another thing is going on. This premise is in effect one of the claims that he so carefully argues for, without explaining quite why, when he is trying to describe P-I Ds. So, intermediate claim (2): if there is any reason to believe of a certain sentence that it is an R-report then it must be possible to show that it is not equivalent to a P-I D. (3) It is not possible to show, for any candidate, that comes close to satisfying the condition for being an R-report, that they are not equivalent to a P-I D. The conclusion is that we have no reason to believe that there are any R-reports.

experience, that this "means, that the immediately following event which the subject can explicitly and truly report must be the same" (Hinton 1973: 76).

Now, before engaging with this line of thought, it can be said that if I have extracted the correct account of Hinton's central argument (or one of his central arguments) against the existence of R-reports, then it involves a second role in Hinton's discussion for P-I Ds. That is that P-I Ds are themselves *not* R-reports, and so unless a suggested report can be shown *not* to be more or less equivalent to a P-I D then there is no right to view it as an R-report.

There is much that could be said about this schematic argument, and much more that could be said about the details of Hinton's presentation of what I am suggesting is this argument. I shall, though, restrict myself to two remarks. (i) If we allow it is correct in its first two premises, it is bound to strike us that there is something rather general and hard to justify about premise (3). How can we know that for any candidates it is not possible to show that it is not equivalent to a P-I D? (ii) More seriously, there is, I suggest, a fallacy in Hinton's general argument. We should, surely, agree with premise (1). P-I Ds are not R-reports. But does (2) follow from this? A problem that arises in deciding this question is what the term 'equivalent' as employed in premise (2) means. I think that the most obvious reading of 'equivalent' which is suitable in this context is: 'P' is equivalent to 'Q' if and only if it is a priori true that necessarily (P if and only if Q). Now, the problem is that there is, or seems to be, no obvious reason why the sentence 'P' cannot satisfy the conditions for being an R-report even though it is equivalent to something that is not an R-report (e.g. a P-I D). The crucial reason why a P-I D is not an R-report is that it is disjunctive; it says that either one sort of thing or another sort of thing is happening. It cannot, in consequence, count as saying *what* sort of thing is occurring. However, being disjunctive does not on its own mean that it cannot be equivalent to a non-disjunctive claim, which, as a consequence of being non-disjunctive may well count as saying *what* is occurring. Here is a mathematical case which may not be completely persuasive but which will indicate the kind of possibility that there is as yet no reason to rule out. If I say that 'X wrote down a number which is either 2 or an uneven prime number' it can be said that my remark does not actually say what type of number X wrote. I give two alternatives. However, if I say that 'X wrote down a prime number' I do say what type of number he wrote. Despite that difference the two remarks are provably equivalent. If this is correct then premise (2) is incorrect, and Hinton's whole argumentative strategy is mistaken.

Although, I think, there is a problem with the logic (as I understand it) of Hinton's argument, and with it the role he assigns in this stage of his discussion to P-I Ds, there is, I want to suggest, a very important insight that Hinton is providing us with. As we have seen, Hinton divided his discussion of the CEH into two parts, one arguing against the existence of R-reports and the other casting doubt on the existence of E-reports. As I have suggested it is not entirely simple to understand Hinton's tendency to express the debate about CEH in terms of the existence of reports. But what does the R-report/E-report distinction really amount to?

The distinction seems to be based on the contrast between basing a belief in the truth of CEH on a substantial philosophical or scientific argument and taking CEH to be obviously true. Given that distinction, Hinton's view seems to be that if one grounds CEH on a theoretical argument one will hold merely that corresponding

to its truth there are E-reports, whereas if one thinks it is, so to speak, simply obvious then one will hold that some ordinary report that we can already make will have the status of an R-report. This last assumption of Hinton's linking the idea that the truth of CEH is obvious to the commitment to there being R-reports may not be correct. However, where Hinton is right, I believe, is in thinking that many philosophers analyse the structure of our thought about perception and experience simply *assuming* that CEH is correct. I would suggest that this is the position of people who endorse the standard argument for the causal analysis of the concept of perception, such as Grice, Strawson, Pears, and Peacocke. Now, Hinton challenges such people to show that there are R-reports which can be demonstrated not to be equivalent to P-I Ds. If the criticism I made above is correct then this is not the appropriate challenge. They do, however, face the closely related challenge of justifying their conviction that CEH is obvious. And as far as I can see there are two possible reactions they might have. The first is simply to say that somehow or other CEH is obviously correct and it is legitimate to assume its truth when analysing our notion of perception or perception itself. Against this Hinton has brought the very real difficulty that it is *not* obvious that the disjunctive description is not the most accurate one available. Why is anyone entitled simply to *assume* that is not so? Alternatively, the proponent of the view that it is obvious can try to argue that there are sentences that everyone uses and understands, which are true in both sorts of case that the disjunctions cover, and which entail CEH. To this claim, Hinton's challenge should have been: show that the sentences in question *do* entail (are committed to) the truth of CEH. In Hinton's remarks designed to cast doubt on the thought that certain sentences are not equivalent to P-I Ds there are many that are relevant to the querying of the claim that these sentences do entail or are committed to CEH. The most obvious question is: what is it about such sentences, candidate examples being, of course, 'It looks to me as if P', or 'I am having an experience as of O', or 'It sensibly seems to me that P', that makes it *obvious* that their use commits us to CEH?

I suggest that engagement with Hinton, despite some unclarities and mistakes in the direction of his argument, makes it clear to us that CEH cannot be taken as either obvious in itself or obvious as a commitment of our perceptual reporting language.

6 HINTON'S TREATMENT OF E-REPORTS AND SOME MORE QUESTIONS

If the previous argument is correct then it means that the CEH must be regarded as a claim requiring a theoretical justification for its acceptance. When Hinton moves the argument away from R-reports to what he calls E-reports he, in effect, simply makes the issue whether CEH is true. However he remains wedded to expressing it in terms of the existence of reports (this time E-reports) even though that is unnecessary, as far as I can see.

Hinton's (1973: 74–100) procedure is to formulate the arguments that occur to him in favour of CEH and to try to undermine them. Clearly, the adequacy of this procedure turns on two things: does Hinton present the best arguments? Does he

correctly assess the arguments that he does give? Now, before making a few remarks about Hinton's discussion I want to point out that at no stage does Hinton approach the issue of CEH by giving arguments *against* it, reasons for thinking that it is false, or reason why its inclusion in our thinking might generate some problems. Nor does Hinton ever explain why it is a legitimate procedure in considering the question of the acceptability of CEH to restrict the range of relevant considerations in the way he does.

If these remarks are fair then I think that we have to agree that there is a serious lacuna in Hinton's approach. The picture that I am proposing of Hinton's argument is that he is opposed, or seems opposed, to what I am calling CEH. Now part of his opposition is to those who simply take CEH to be obvious, and, perhaps, in a sense, so obvious that it is part of our basic understanding of the world within which, as one might say, our concepts are embedded. That is, roughly, the attitude he is engaging with when discussing R-reports. I have suggested that Hinton's opposition to this view amounts to a real insight (whatever the force of some of the points that he makes). However, if we are to follow him in opposing CEH *in general* we surely need some indication as to why it is a bad idea, some indication of the problems that await us in trying to incorporate it into our thinking. As far as I can see, Hinton provides nothing of this sort at all.

I want to highlight two moments in Hinton's engagement with what he is calling the existence of E-reports. Hinton, it seems to me, thinks that CEH and hence the claim that there are E-reports will be supported in two ways. The first is an argument based on considerations about causation (Hinton 1973: 75–82). The second is the vaguer thought that CEH is somehow vindicated by science and its study of perception.

The causal argument, in the form that Hinton seems worried by it, is roughly this: since the immediate cause of the occurrence of a perception is the same as the immediate cause of the occurrence of the illusion, and since the principle 'same cause, same effect' is true, then it follows that the same type of experience is caused, and so CEH is true. This is undoubtedly a plausible argument. It is, though, highly abstract and based on certain assumptions, so there is room to try to devise ways to resist it. Hinton, it has to be said, in engaging with this type of argument is engaging with a central line of thought relevant to the assessment of CEH. His discussion is, however, fairly hard to follow in that he tends to try to locate assumptions about *language* that the proponent of the argument is making, which, according to Hinton, are not assumptions we are under any obligation to accept. The question about this approach is why Hinton thinks that assumptions about language *are* being made. For example, Hinton (1973: 77) ascribes to the proponent of the argument the assumption that 'we should not know what we are talking about—indeed, should not be talking about anything, real' if we merely say that what is caused in the two cases covered by a P–I D is either a perception or an illusion. But, why should someone who thinks that the causal argument is sound make any such comment? What, then, besides this does Hinton offer as a basis to resist the argument? The core of his response seems to be to suggest that it rests on what he calls a principle "of the narrow identifiability of effects" (Hinton 1973: 80). Hinton claims that this principle need not be

accepted. Now, this is interesting but, I want to suggest, rather imponderable and difficult to sort out. My question is: what exactly is the principle about identifiability on which Hinton thinks the argument rests? One possible interpretation would treat 'identifiability' as requiring the possibility of someone's actually identifying what the effect is in terms which are independent of its cause. If this is what the principle means then a proponent of the argument could simply deny that it rests on any such assumption. The argument, it will be said, does not assume that we can identify effects; it assumes, rather, that the same cause will produce the same effect. If Hinton says that the narrow identifiability principle is the same cause/same effect principle then the supporter of the argument will simply demand that it be explained to him what is wrong with it. It seems to me that nothing very much has been clarified in this exchange.

In response to the more general claim about science Hinton makes the perfectly reasonable remark that he is not himself persuaded that the scientific examples he has looked at do involve CEH. However, having said that, he then merely remarks that "the progress of science, rather than the animadversions of philosophers, will decide whether this idea [i.e. CEH] prospers" (Hinton 1973: 86). This remark, if I have understood it, has the effect of plunging the reader of his text into considerable confusion. It leaves Hinton having no serious objection to CEH, nor any reason for thinking it will not get accepted. It seems to mean that Hinton is saying that he has not yet been persuaded that CEH is true. This is, perhaps, the most surprising aspect of Hinton's book.

7 CONCLUSION

In the central sections of *Experiences*, the ones I have been analysing, Hinton has, I believe, three massive achievements to his credit. The first is that he isolates a thesis about perceptual experience, which, abbreviating his name, I have called CEH, and about which he is right to suggest that it is simply assumed by many philosophers but without proper justification. There may be a problem with one major argument he uses but both his general thesis, and many of the remarks he made in the course of defending it, seem to me to have the ring of truth. Second, he isolates and analyses a type of disjunctive sentence, P–I Ds, which is of colossal interest in the philosophy of perception. I remain somewhat unclear as to the precise role Hinton himself ascribes to such sentences, but considering what role they, or other similar sentences, might have, has proved very fruitful. Third, Hinton identified and explored many of the considerations which are relevant to arguments in favour of CEH, even though he left much to say. Not surprisingly, his treatment of P–I Ds in *Experiences*, though avoiding some well known problems, faces difficulties given his employment of the notion of 'plainness'. There are also, I have suggested, problems with some of his leading arguments. The main lacuna in his treatment overall is that he gives us no real insight into the significance of CEH and, perhaps because of his extreme caution, ends up simply arguing that *so far* CEH has not been supported. This is a far weaker conclusion than he is normally understood to be advancing.

My conclusions, though, relate primarily to the central section of Hinton's book. I have not engaged much with the other sections, and have also had to ignore much in the central part. I hope these facts, together with the considerable interest, as it seems to me, of the parts I have looked at, will go some way to making Hinton read as well as cited.

REFERENCES

- Austin, J. L. (1964) *Sense and Sensibilia*, Oxford: Oxford University Press.
 Dancy, J. (ed.) (1988) *Perceptual Knowledge*, Oxford: Oxford University Press.
 Hinton, J. M. (1967a) 'Experiences', *Philosophical Quarterly*, 17: 1–13.
 — (1967b) 'Visual Experiences', *Mind*, 76: 217–27.
 — (1973) *Experiences*, Oxford: Clarendon Press.
 McDowell, J. (1982) 'Criteria, Defeasibility and Knowledge', *Proceedings of the British Academy*, 68: 455–79.
 Pitcher, G. (1971) *A Theory of Perception*, Princeton: Princeton University Press.
 Snowdon, P. F. (1980–1) 'Perception, Vision and Causation', *Proceedings of the Aristotelian Society*, 80: 175–92. [Page references are to the reprint in (Dancy 1988).]
 — (2004) 'The Formulation of Disjunctivism: A Response to Fish', *Proceedings of the Aristotelian Society*, 105: 129–42.

2

Either/Or

Alex Byrne and Heather Logue

Perhaps it has sometimes occurred to you, dear reader, to doubt the correctness of the familiar philosophical proposition that the outward is the inward . . .

(Kierkegaard, *Either/Or*)

This essay surveys the varieties of disjunctivism about perceptual experience. Disjunctivism comes in two main flavours, metaphysical and epistemological. Metaphysical disjunctivism is the view usually associated with the disjunctivist label, and whenever ‘disjunctivism’ occurs here unprefixd, it refers to this view. After some initial discussion of (metaphysical) disjunctivism (sections 1–3), we explain epistemological disjunctivism in section 4. The rest of the essay is solely concerned with explaining and assessing metaphysical disjunctivism, a theory of the nature of perceptual experience. Our main (and provisional) conclusion is that although there is considerable insight in the vicinity, metaphysical disjunctivism is false.

1 DISJUNCTIVISM: THE BASIC IDEA

Like most disjunctivists, we will concentrate exclusively on visual experience. Borrowing some terminology from Williamson (2000), let a *case* be a centred possible world—a world with a marked subject and time. A *good* case is one in which the subject sees an object, and sees that it has such-and-such properties (at the time of the

Thanks to audiences at MIT and the Universities of Glasgow and Connecticut for discussion; and to Tom Dougherty, Adrian Haddock, Benj Hellie, Rae Langton, Fiona Macpherson, Adam Pautz, Jonathan Schaffer, Susanna Siegel, Scott Sturgeon, Mike Thau, Charles Travis, and an anonymous referee for extremely helpful comments.

case). For example, there is a good case in which the subject is J. L. Austin who (at the time of the case) sees a single tomato against a plain background, and sees that it is red and spherical. In a good case of this sort, Austin's visual experience is *veridical* with respect to the tomato's colour and shape and (we may suppose) is not in any respect *non-veridical*; similarly for good cases in general.

Starting with a particular good case—say, involving Austin and the tomato—we may define a case α to be *subjectively indistinguishable from the good case* ('subjectively indistinguishable' for short) if and only if, in α , the subject is not in a position to know by 'introspection' alone that he is not in the good case. More exactly, there is no proposition p incompatible with the subject's being in the good case such that he is in a position to know p by introspection alone. ('Introspection' is simply a placeholder for the normal first-person method of finding out about one's experiences and perceptual states, whatever that is, exactly.¹) The good case itself is trivially subjectively indistinguishable, because the subject cannot be in position to know what's false—in particular, that he's not in the good case. Since normal human perceivers are not in a position to know everything, innumerable other good cases are also subjectively indistinguishable; potential omniscience would not block indistinguishability anyway, because no subject is in a position to know everything by introspection alone. Another subjectively indistinguishable case is an *illusory* case—Austin sees the tomato, but it is not the way it looks. For example, the tomato isn't red, despite looking so. Yet another is an *hallucinatory* case—Austin 'seems to see' a red tomato, but in fact sees nothing, or at least no material object.² Illusory and hallucinatory cases are *bad* cases. Although bad cases involve illusions and hallucinations of any kind, when we speak of 'bad cases' in what follows we typically mean the specific ones just mentioned, corresponding to the good case involving Austin and the tomato.

It is important to note that we are not *defining* this restricted kind of bad case as one that is subjectively indistinguishable and non-good. All such bad cases have this feature, but the converse is obviously problematic. For example, a case in which the subject is comatose is subjectively indistinguishable and non-good, yet is not one in which Austin is hallucinating or illuding a red tomato. On the other hand, we are not

¹ Presumably discovering that one sees a tomato by 'introspection' somehow involves exercising one's *visual* capacities—one discovers one sees a tomato by attending to the tomato before one's eyes (Evans 1982: 226–8; Dretske 2003). Just how this works is controversial, but we do not need to take a stand on the issue here.

² In some hallucinatory cases, the very same tomato present in the good case, with the very same colour and shape, is before the subject—a phenomenon known as *veridical hallucination*. Likewise, there are subjectively indistinguishable cases of *veridical illusion*: the subject sees the (red) tomato, which *looks* red, but only because of (say) some fluky occurrence in the subject's brain—given the lighting conditions, the (red) tomato would not look red to a normal subject. The (abnormal) subject does not see *that* the tomato is red. Note that veridical illusions are not 'illusory cases' as defined in the text; for simplicity we will set them aside. See, in particular, Johnston (2006); for discussion in connection with the disjunctivism defended by M. G. F. Martin, see Hawthorne and Kovakovich (2006: 162–3).

Obviously many real-life examples of perceptual experience are 'mixed'—they involve some combination of veridical perception, illusion, and hallucination. Again for simplicity these will be set aside.

precluding the possibility that a more refined account of subjective indistinguishability could supply a definition. Even though we have not given necessary and sufficient conditions for a case to be of the restricted bad kind, we assume that we have said enough to make the notion tolerably clear.³

On one view, the good case and the bad cases have a common mental core—in all such cases, the subject is having an experience of a certain kind, or is in a certain (experiential) mental state (at the time of the case). For instance, the good case and the bad cases might each involve experiences that represent that there is a red spherical thing before the subject, or experiences of seeing a red ‘circular’ sense-datum (the ‘red-prime’ terminology is from Peacocke 1983). As these examples illustrate, the common core is supposed to be quite specific—it is absent in dissimilar cases of perceptual experience. We can finesse various distracting details by saying that the alleged common core is at least absent in any good case in which Austin sees a green ovoid tomato. Although the nature of the common core is disputable, that there is one can seem nothing short of obvious.

The basic idea of disjunctivism, as it is predominantly characterized by its proponents, is that this ‘obvious’ view is false. At least some bad cases are mentally radically unlike the good case. All disjunctivists agree that hallucinatory cases are of this sort; they are divided on whether illusory cases should also be included. Since hallucinatory cases (at least) share no mental core with the good case, disjunctivists hold that the most perspicuous characterization of the class comprising the good and the bad cases is disjunctive—like the class of ravens-or-writing-desks.

Although traces of disjunctivism may be found in earlier authors, J. M. Hinton was the first to make the position explicit.⁴ In his article ‘Visual Experiences’ he contrasts seeing a flash of light with having “an illusion of a flash of light”, noting that we can truly say and think the following:

(A) I see a flash of light: actual light, a photic flash . . .

(B) I have an illusion of a flash of light . . .

(A \vee B) Either I see a flash of light, or I have an illusion of a flash of light. (Hinton 1967a: 217)

(Clearly Hinton is assuming that bad B-cases are subjectively indistinguishable from good A-cases.) He contrasts (A \vee B), which does not say “definitively what is happening”, with:

(Q) I psi—possible wordings: I see a flash, I have a visual experience of a flash—which does give a definite answer to the question ‘What is happening?’ (1967a: 220)

(Q) is supposed to describe something that happens when (and only when) (A) or (B) is true—the occurrence of a “perception/illusion-neutral entity” (in the

³ The extreme caution in this paragraph anticipates Martin’s view that a more sophisticated and elaborate account of subjective indistinguishability can be used to define bad cases. See section 6.2 and footnote 31 below.

⁴ Hinton does not use ‘disjunctivism’, though—that coinage appears to be due to Robinson (1985); see Martin (2006: 356, fn. 2).

terminology of Locke 1975: 467). That is, I psi iff I see a flash of light or I have an illusion of one. Unlike (A \vee B), (Q) says “definitively what is happening”—namely, that I am the subject of a perception/illusion-neutral event of psi-ing.⁵ The difference between (A \vee B) and (Q) is that while (Q) is true only if there is a mental event common to seeing a flash and having an illusion of one, the truth of (A \vee B) is compatible with perceptions and illusions having no kind of mental event in common.⁶

Although Hinton’s reasoning is hard to follow, he does think that (A \vee B) is the best that can be done: “I do not see”, he says, “how it can be shown that there is such a thing as my psi-ing” (1967a: 220). As he puts it in *Experiences*, there is no “kind of experience common *and peculiar*” to the good case and the bad cases—they have no “common element” (1973: 62). Essentially the same view is (apparently) expressed by later writers in slightly different terms:

Good and bad cases have no “highest common factor”. (McDowell 1982/1998: 386)

There is no “single sort of state of affairs” obtaining in good and bad cases. (Snowdon 1980–1: 186)

There is no “distinctive mental event or state common to these various disjoint situations”. (Martin 2004: 37)

(Various retractions and qualifications will be made below, in section 4, section 7.1, and section 7.2.)

We said above that there is some disagreement among disjunctivists about the status of illusory cases. Some, like Snowdon, group them with the good case, on the left-hand side of the disjunction; others group them with hallucinatory cases, on the right-hand side, as illustrated in Table 2.1.

Table 2.1 Types of disjunctivism

| | |
|----------------------------------|---|
| R. H. S. disjunct: hallucination | R. H. S. disjunct: hallucination and illusion |
| Snowdon | McDowell, Hinton, Martin |

⁵ Since (Q) (assuming there is such a proposition) and (A \vee B) are plainly supposed to be *necessarily* equivalent, Hinton’s claim that (Q) (if it exists) says “definitively what is happening”, while (A \vee B) doesn’t, is confusing. Hinton’s point is better put as follows: (Q) says that a particular kind of mental event is happening, while (A \vee B) doesn’t. Contrast ‘A horse is being born’ with ‘Either a steed is being born, or a mare is being born’; the former, but not the latter, says that the birth of a particular kind of creature is happening.

⁶ Criticizing Hinton, Snowdon (1980–1: 184) points out that ‘I see an F’ is an unhappy way of characterizing the good case, because it carries no implications whatsoever about how things look—one may see what is in fact a big red tomato, even though it looks exactly like a small purple grape. It is harder to tell a similar story in which one sees a flash of light which does not look like a flash—this is one respect in which Hinton’s main example is ill-chosen. (To add to the confusion, Hinton considers—as (Q) illustrates—the possibility that *both* sides of the disjunction could be characterized by “I see a flash”; see Hinton 1967a: 221–3.) In fact, Snowdon’s point is made earlier in Hinton’s book *Experiences*: Hinton accommodates it by stipulating that ‘see’ in perception-illusion disjunctions means what he calls “plainly see” (1973: 42, 112).

A cursory reading of Snowdon and McDowell will confirm their position in this table. The placement of Hinton and Martin is not straightforward, so we will briefly note some evidence.

One might think that Hinton requires no comment at all—haven't we just seen that he puts illusions ("I have an illusion of a flash of light") on the right-hand side? But Hinton's main example of a flash of light might as well have been purposely designed to blur the crucial distinction between illusions and hallucinations. For, of course, a typical 'illusion of a flash of light' is quite unlike a typical illusion of a bent stick—when one has an illusion of a flash of light by having a phosphene experience, say, one is not seeing *anything*. That is, an 'illusion' of a flash of light is often an *hallucination* of a flash of light.⁷

However, in his paper 'Experiences', Hinton says he is "inclined to believe" that "I see blue" . . . is indistinguishable from the mere disjunction, 'Either I actually see an optical object that is blue in colour, or I am in some situation or other that is to me like that one' " (1967b: 12). Here only the veridical case—a blue object looks blue to the subject—is on the left-hand side. Hence, for Hinton, illusions are grouped with hallucinations.⁸

Martin explicitly contrasts the good case with hallucinatory cases, and mentions illusions only in passing. But his emphasis is very much on the *veridical* nature of the good case, which suggests that it does not belong with illusory cases. And in 'On Being Alienated', he says that "perceptions fail to be the same kind of mental episode as illusions or hallucinations" (Martin 2006: 360; see also 361, 362), thus classifying illusions with hallucinations, not with (veridical) perceptions.⁹

2 THE ARGUMENT FROM ILLUSION/HALLUCINATION

So far, we have somewhat imprecisely outlined disjunctivism about perceptual experience, and indicated one choice point for disjunctivists. Nothing has yet been said

⁷ Indeed, Hinton's illustration of an "illusion of a flash of light" is a phosphene experience (1967c; 1973: 40). (See also his 1973: 61: ". . . if there is such an optical object").

⁸ See also Hinton (1973: 116), where Hinton says that illusions (in "the ordinary general-cultural sense", which is apparently his official understanding of the word) comprise both illusions and hallucinations (in our 'standard philosophical' sense).

There is, incidentally, one apparently recalcitrant passage in *Experiences*: "the perception-proposition in a perception-illusion disjunction can very well be a proposition about how something looks: one kind of perception-illusion disjunction is exemplified by: 'Either I visually perceive an optical object which looks (a great deal, a little, hardly at all) like a two-dimensional coloured shape, or I am having the illusion of doing so'" (1973: 61).

⁹ This quotation is perhaps misleading, because all Martin may have in mind is that perceptions fail to be the same "fundamental" kind of mental episode as illusions or hallucinations (for an explanation of Martin's "fundamental kind" terminology, see section 7.1); if so, then the quotation does not imply that there is no common element to the good case and the illusory cases (see again section 7.1). The probative point, though, is that Martin's *arguments* (see section 6 and section 7), although focused on hallucinations, may also be adapted for illusions—a task that we have left for the reader. See also Martin (2002a: 395, fn. 24).

about why the view might be thought attractive. The central motivations for disjunctivism found in the literature will be discussed much later, in section 6 and section 7. In this section and the next, we will consider an alternative route to disjunctivism, as a way of escaping from the argument from illusion (or the argument from hallucination).¹⁰ Although this route is not appealing, it will enable us to clear up something of an interpretive muddle concerning McDowell's 'disjunctivism'. We will argue that McDowell's view is very different from the kind of Hintonesque disjunctivism discussed so far.

Return to the particular good case in which Austin sees a red tomato. In a subjectively indistinguishable illusory case, Austin sees a non-red material thing (a green tomato, suppose); in a subjectively indistinguishable hallucinatory case he does not see anything material. The argument from illusion (hallucination) argues from a claim about illusory (hallucinatory) cases to the paradoxical claim that the good case is not a genuine possibility—that is, the good case does not exist. One version of the argument from illusion is the following.

1. The good case exists. (Assumption for *reductio*.)

Given (1), various subjectively indistinguishable illusory cases defined in terms of the good case also exist. Consider a particular illusory case in which Austin sees a green tomato:

2. In the illusory case, there is a red object *o* that Austin sees. (Premise)

Since the tomato in the illusory case is green, not red, and there are no other red material things Austin might see:

3. *o* is an 'immaterial' thing.
4. If Austin sees a red object in the illusory case, then he sees a red object of the same kind (material/immaterial) in the good case. (Premise)

Hence, from (2), (3), (4):

5. In the good case, Austin sees a red immaterial object.
6. If, in the good case, Austin sees a red immaterial object, he does not also see the tomato. (Premise)

Hence, from (5), (6), and the stipulation that Austin sees the tomato in the good case:

7. In the good case, Austin sees and does not see the tomato.

Since the good case is a *possibility*, (7) is a contradiction. Our starting assumption (1) is supposedly the culprit—so the good case does not exist.

The argument from hallucination is structurally similar. In the more traditional version of the argument(s), the bitter contradiction of (7) is sweetened by the insertion of 'directly' in front of 'see' throughout the premises. Thus the traditional conclusion is simply that Austin *sees* but does not *directly* see the tomato in the good case.

¹⁰ For a recent discussion of both, see Smith (2002).

This is not supposed to threaten the existence of the good case, because allegedly one may see tomatoes ‘indirectly’. Unfortunately no one was ever able to explain satisfactorily what ‘indirectly’ was supposed to mean.¹¹ So if the traditional argument works at all, it can apparently be strengthened to the argument above, where the conclusion is obviously unacceptable.

The argument has three premises, and denying the third—(6)—does not seem promising. The culprit must therefore be either (2) or (4). Adapting A. D. Smith’s terminology slightly, (2) is the “sense datum” premise and (4) is the “generalizing” premise (2002: 25–6). On the usual diagnosis, the argument is unsound because the sense-datum premise is false. One may see something that *looks* red without seeing anything that *is* red. Similarly, one may hallucinate a red tomato, even though one sees nothing at all.

Although the sense-datum premise is widely thought to be false, in both its illusory and hallucinatory versions, there is at least a case to be made for the hallucinatory version.¹² So another option is to accept the sense-datum premise but deny the generalizing premise. Without too much wriggling, this option can be extracted from Austin’s *Sense and Sensibilia*.

3 ‘AUSTINIAN’ DISJUNCTIVISM

Austin’s flirtation with accepting the sense-datum premise (in its hallucinatory version), while denying the generalizing premise, occurs in this passage:

[In the] mirage case . . . we are supposing the man to be genuinely deluded, he is *not* “seeing a material thing”. We don’t actually have to say, however, even here that he is “experiencing sense-data”; for though, as Ayer says above, “it is convenient to give a name” to what he is experiencing, the fact is that it already has a name—a *mirage*. (Austin 1962: 32, footnote omitted)

Here Austin could be read as conceding (perhaps only for the sake of the argument) that one is aware of an ‘immaterial object’ when one experiences a mirage, while implicitly denying that this has any untoward implications for the good case—one encounters immaterial objects like mirages very rarely; most of the time one sees material objects like oases.

Is this a kind of Hintonesque disjunctivism, where there is no mental “common element”? Not necessarily: the view is primarily a disjunctivism about the *objects* of experience, rather than a disjunctivism about the experiences themselves. (The distinction is due to Thau 2004: 195.) To adapt an example from Austin (1962: 50), a lemon and a lemon-like bar of soap are very different kinds of thing; that is consistent with the experience of the soap and the experience of the lemon being mentally exactly the same.

¹¹ The classic discussion of early attempts to explain it is Dretske (1969: 62–75); a notable later attempt is Jackson (1977: ch. 1), which relies on contentious claims about the analysis of sentences like ‘S sees o’; see also Armstrong (1976).

¹² See Smith (2002: 195).

Still, Hintonesque disjunctivism is not far away. Traditionally, one bears an especially intimate and somewhat mysterious relation of awareness—“acquaintance”—to sense-data (Russell 1912). So the Austinian disjunctivist might add that one is acquainted with immaterial mirages, but one merely sees material oases. On this Austinian/Hintonesque view, presumably the two sorts of experiences lack a common mental element.¹³

Despite the flirtation, Austin is not an Austinian disjunctivist.¹⁴ However, a similar position seems to be suggested in McDowell’s ‘Criteria, Defeasibility and Knowledge’, which is mostly taken up with a discussion of Wittgenstein’s views on ‘criteria’ and the problem of other minds. That problem starts with the assumption that our “best warrant for a psychological judgment about another person is defeasible evidence constituted by his ‘behavior’ and ‘bodily’ circumstances” (McDowell 1982/1998: 383). Given the assumption, the sceptic about other minds questions why this evidence, which is consistent with the absence of any psychology, is *ever* good enough for knowledge. And it is this assumption that Wittgenstein is said to deny: he “rejects the sceptic’s conception of what is given” (McDowell 1982/1998: 385). Sometimes one simply perceives that another person is sad, for instance, without this knowledge resting on evidence only contingently connected with the fact itself.

However, “[t]he possibility of such a position is liable to be obscured from us by a certain tempting line of argument”, namely the “argument from illusion” (McDowell 1982/1998: 385). That argument starts from the fact that an illusory bad case is subjectively indistinguishable from the good case, and derives the lemma that “one’s experiential intake—what one embraces within the scope of one’s consciousness—must be the same” (1982/1998: 386) in both cases. With this lemma in hand, the argument concludes that one’s “experiential intake” is *never* the “fact made manifest” (1982/1998: 390):

In a deceptive [illusory] case, one’s experiential intake must *ex hypothesi* fall short of the fact itself [that the tomato is green, for instance], in the sense of being consistent with there being no such fact. So that must be true, according to the argument, in a non-deceptive case too. (1982/1998: 386)

But accepting this argument is not—in a McDowellian word—‘compulsory’. Instead, we can:

say that an appearance that such-and-such is the case can be *either* a mere appearance *or* the fact that such-and-such is the case making itself perceptually manifest to someone. As before, the *object of experience* in the deceptive cases is a *mere appearance*. But we are not to accept

¹³ The Austinian/Hintonesque disjunctivist’s distinction between *seeing* and (*visual*) *acquaintance* may be thought of as equivalent to the traditional sense-datum theorist’s distinction between *indirectly seeing* and *directly seeing*.

¹⁴ See Thau (2004: 200–1). Alston (1999) comes close enough to the position, holding that in hallucinations “what appears to the subject is a particularly vivid mental image” (191); see also Locke (1967: 111–2), and Robinson (1994: 153). The “Selective Theory” discussed by Price (1950: 39–51) is a Berkeley-*sans*-God version of Austinian disjunctivism (compare Martin 1997: 95; Johnston 2004: 145).

that in the non-deceptive cases too the object of experience is a mere appearance, and hence something that falls short of the fact itself. On the contrary, *the appearance that is presented to one* in those cases is a matter of *the fact itself* being disclosed to the experiencer. So appearances are no longer conceived as *in general intervening* between the experiencing subject and the world . . . (1982/1998: 386–7, footnote omitted, boldface emphasis ours)

This passage can easily be read as saying that some strange entity, a “mere appearance”, “intervenes” between the subject and the world in an “experientially indistinguishable” (1982/1998: 386) illusory case, but is absent in the good case, in which something quite different, the “fact made manifest” (1982/1998: 390), is the object of experience. McDowell says little more about “mere appearances”, but given that this passage is part of his discussion of the traditional argument from illusion, it would not be unreasonable to take them to be sense-data. And here is Thau doing just that:

The argument from illusion has two stages. First, we’re asked to conclude that since the object of a non-veridical (i.e., inaccurate) perception isn’t a material thing, non-veridical (i.e., inaccurate) perceptions have some kind of odd nonmaterial object—John McDowell calls them “mere appearances”; but, of course, they’ve also been called “sense-data” . . . (Thau 2004: 194)

According to Thau, then, McDowell is an Austinian disjunctivist. Although Thau himself does not attribute Hintonesque disjunctivism to McDowell, that would be a natural extension of the position, as we have seen. And commentators frequently take McDowell and Hinton to be batting on the same team.¹⁵ These interpretations may be understandable but—as we will now argue—they are incorrect.

4 EPISTEMOLOGICAL DISJUNCTIVISM

In ‘Criteria, Defeasibility, and Knowledge’, the possibility of a Wittgensteinian position on other minds is said to be obscured by the “argument from illusion”. Obviously McDowell has a version of the traditional argument in mind, which is about the perception of objects like tomatoes, not another’s sadness. Equally obviously, McDowell thinks that the traditional argument (as he construes it) obscures the possibility of a Wittgensteinian position *when it is adapted to the example of other minds*. In the *adaptation* of the argument the good case is one in which the subject is observing another person, and sees that she is sad, and an (illusory) bad case is one in which the Other is *not* sad, but is behaviourally the same as in the good case—hence *appearing* sad. (So the subject in a bad case is not in a position to know that he is not in the good case.) Now, what could McDowell mean by *agreeing* that “the *object of experience* in the deceptive cases is a *mere appearance*”? Since this claim is supposed to apply, not just to the traditional bad case (in which a non-red

¹⁵ See, for example, Byrne (2001: 202, fn. 7); Child (1994: 143–4); Crane (2005: section 3.4); Langsam (1997: 57, fn. 10); Martin (2006: 356–7, fn. 2); Robinson (1994: 247, fn. 6); Smith (2002: 197); Sturgeon (2000: 41).

tomato looks red), but also to the ‘other minds’ bad case, he cannot mean that the subject is aware of a *sense-datum* in these bad cases. A sense-datum in the traditional bad case would bear the apparent properties of the tomato; likewise, a sense-datum in the ‘other minds’ bad case would bear the apparent properties of the Other—and sad sense-data are not, we may safely say, in McDowell’s ontology. So McDowellian “mere appearances” are not sense-data.

In fact, McDowell means nothing alarming by talk of “appearances”: it is just a way of talking about how things look (sound, taste, etc.). And a “mere appearance” is “an appearance that such-and-such is the case, falling short of the fact [in the sense of being consistent with there being no such fact]” (1982/1998: 386). What is potentially confusing, of course, is McDowell’s claim that sometimes a “mere appearance” is the “object of experience”. But here he is charitably construed as saying that sometimes “what one embraces in the scope of one’s consciousness” is not the fact that such-and-such, but merely the fact that it appears to one that such-and-such.

It needs to be emphasized that in ‘Criteria, Defeasibility and Knowledge’ McDowell is chiefly concerned with the *epistemology* of perception, not its metaphysical structure. He is not particularly interested in the question of whether there is a Hintonesque “common element” to the good and bad cases.¹⁶ This concern with epistemology explains why McDowell focuses on locutions of the form ‘S sees that p’ and ‘It looks to S as if p’, unlike Hinton, who mostly concentrates on ‘S sees an F’.¹⁷ And McDowell’s main point is that in the good case, the (perceptual) *evidence* (or, as he sometimes says, the “epistemic warrant”) one has for believing that the world is thus-and-so is considerably *better* than the (perceptual) evidence one has in the bad cases.

The view McDowell opposes can be motivated by the fact that the bad cases are subjectively indistinguishable from the good case—in a bad case, the subject is not in a position to know that he is not in the good case. If the subject’s evidence in a bad case is weaker than his evidence in the good case, then in the bad case he is not in a position to know what his evidence is—else he would be in a position to know that he is not in the good case. But surely one always *is* in a position to know what one’s evidence is. So the subject’s evidence in the good case must be the same in any bad case. In other words, “the true starting point in the space of reasons must be something common to the favourable and the potentially misleading cases (like having it look to one as if things are thus and so)” (McDowell 1995/1998: 397). (The classic critical exposition and discussion of this line of thought is Williamson 2000: ch. 8.)

Thus McDowell’s “highest common factor”, “what is available to experience in the deceptive and the non-deceptive cases alike” (1982/1998: 386), is *perceptual evidence*. And one problem with the highest common factor view is, of course, that it tends to lead to scepticism. The subject’s (perceptual) evidence in the bad cases is very impoverished, being solely a matter of how things appear. That evidence plainly does not

¹⁶ As Snowdon recognizes (2005: 139–40).

¹⁷ This difference between McDowell and Hinton is noted in Snowdon (1990: 131–2), Martin (2002a: 395, fn. 24), and Martin (2004: 44).

entail that (for example) there is a red spherical thing before the subject—but might it nevertheless *justify* the belief that there is a red spherical thing before the subject? “Anyone who knows the dreary history of epistemology knows that this hope is rather faint” (McDowell 1995/1998: 396). Hence, even if “the world is doing me a favor” (1995/1998: 396) by placing a ripe tomato on the table in good light with nothing funny going on, I do not know there is a red spherical thing before me.¹⁸

On the view McDowell recommends, the subject’s evidence in the good case is much better than in the bad cases: “When someone has a fact made manifest to him, the obtaining of the fact contributes to his *epistemic standing* on the question” (McDowell 1982/1998: 390–1, emphasis ours). Let E be the strongest perceptual evidence that the subject has in the bad cases. Say that E is *good enough* evidence if and only if, in the good case—an ordinary perceptual situation where the world is not playing any tricks—the subject knows propositions about the external world (that there is a red spherical thing before her, say) on the basis of E. The good case and the bad cases have a *common justifying element* if and only if E, the evidence the subject has in the bad cases, is good enough.

Now, since McDowell agrees with the sceptic that E is not good enough, he denies that there is a common justifying element. *Unlike* the sceptic, McDowell also denies that E is the strongest evidence in the good case; that is, he denies that there is a highest common factor. According to McDowell, either one knows that there is a red spherical thing before one, or it merely seems that one has evidence good enough for knowing there is a red spherical thing before one (the bad cases). We can call this view *epistemological disjunctivism*.¹⁹

Epistemological disjunctivism is not a rival to metaphysical disjunctivism; in fact, as we will explain immediately below (section 5), the latter leads naturally if not inexorably to the former. However, epistemological disjunctivism is quite compatible with the denial of metaphysical disjunctivism. Indeed, in ‘Criteria, Defeasibility, and Knowledge’, McDowell seems to presuppose that metaphysical disjunctivism is false—more exactly, that there is a Hintonesque common element at least in the *illusory* cases. It “look[s] to one exactly as if things [are] a certain way” (1982/1998: 385) in the good case and the illusory cases, and there is the strong suggestion that this is a perfectly proper mental respect of similarity.²⁰ The presence of a Hintonesque common element is confirmed in McDowell’s contribution to this volume: the central thesis of “the disjunctive conception” is that “the two sides of the disjunction differ in epistemic significance . . . This difference in epistemic significance is of course

¹⁸ For more on the sceptic’s ‘argument from underdetermination’, see Byrne (2004).

¹⁹ Williamson (2000) rejects the highest common factor view (and notes that McDowell agrees (Williamson 2000: 169, fn. 2)), and is no sceptic, but does not officially take a stand on whether one’s evidence is good enough in the bad cases. Although one may guess he thinks that it probably isn’t. In any event, he is not clearly an epistemological disjunctivist.

²⁰ It should not be assumed that McDowell holds that there is a common element in the *hallucinatory* cases. He might think that demonstratives should figure in instances of ‘It looks to one exactly as if things are a certain way’ (e.g. ‘It looks to one exactly as if *this* is red and spherical’), and if so then since the demonstrative does not refer in hallucinatory cases, the existence of a common element is problematic (compare McDowell 1986/1998). For more discussion see section 5, section 9, and footnote 56.

consistent with all sorts of commonalities between the disjuncts. For instance, on both sides of the disjunction it appears to one that, say, there is a red cube in front of one” (this volume: fn. 15). On *this* point—about the Hintonesque common element—Putnam has McDowell exactly right:

McDowell does indeed insist on the existence of *this* kind of “common factor.” Part of the content of a nonveridical experience can indeed be the same as part of the content of a veridical experience. Both experiences can “tell one” (incorrectly, in the nonveridical case) that there is a yellow door in front of one, for example. (Putnam 1999: 154)²¹

McDowell, though, deserves a slap on the wrist for encouraging the widespread tendency to misinterpret his views. The relevant section of ‘Criteria, Defeasibility, and Knowledge’, he unwisely conjectures in a footnote, “grew out of an unconscious recollection of Hinton’s articles ‘Experiences’ and ‘Visual Experiences’” (1982/1998: 387, fn. 33).

5 METAPHYSICAL DISJUNCTIVISM AND SOME FOILS

We now begin our assessment of metaphysical (Hintonesque) disjunctivism—the view that there is no kind of (reasonably specific) mental state or event common to the good case and the bad cases. Admittedly, this is only as clear as the notion of a mental state or event, but we will assume that it is clear enough. Leading candidates for mental statehood include *x believes that the tomato is red*, *x sees the tomato*, and *x sees that the tomato is red*. (One might demur at the last two examples on the ground that a subject’s being in the state entails, respectively, that the tomato exists and that it’s red; we will touch on this later in this section.) Examples of states specified in mental vocabulary that are *not* mental states include *x either believes that the tomato is red or sees the tomato*, *x does not see that the tomato is red* and—an example from Williamson (2000: 27–8)—*x truly believes that the tomato is red*.

As we saw in the previous section, epistemological disjunctivism can readily be combined with the rejection of metaphysical disjunctivism. But what about the converse? That is, could metaphysical disjunctivism be combined with the rejection of epistemological disjunctivism? Not happily. For illustration, suppose that Austinian-cum-Hintonesque metaphysical disjunctivism is true: in a certain hallucinatory bad case, Austin is acquainted with sense-data; in the good case, he simply sees the tomato.²² Consider Austin’s evidence for the proposition that there is a red spherical thing before him: in the bad case, his perceptual evidence is nothing stronger than

²¹ Although Putnam nails his colours to the disjunctivist mast (1999: 153), his use of this terminology is idiosyncratic. (Perhaps this is explained by Putnam’s unholy trinity of ‘disjunctivists’—James, Austin, and McDowell.) A Putnamian disjunctivist is simply someone who takes a standard ‘intentionalist’ line on illusions. There is no red (or red’) sense-datum or experience in an illusory case; instead, the illusory experience “*portrays the environment as containing*” a red thing, and so has redness “intentionally”, not “adjectivally” (Putnam 1999: 154). Far from being (metaphysical) disjunctivism, this is the view’s chief opposition (see the following section).

²² Austinian disjunctivism is not necessarily Hintonesque; in other words, it is not necessarily a version of metaphysical disjunctivism (see section 3).

facts about sense-data; in the good case, his perceptual evidence must be quite different, since there are no sense-data in the good case. What is that perceptual evidence? Perhaps it is, simply, that there is a red spherical thing before Austin, or something else that entails this proposition (e.g. that *this* is red' spherical, and before Austin). Alternatively, perhaps it is nothing stronger than the fact that *this* (the tomato) looks to be red, spherical, and before Austin. Either way, these are not items of evidence available in the hallucinatory bad case, where the tomato is absent. In order to *reject* epistemological disjunctivism, two claims must be established. First, that there is a certain overlap of evidence E in both cases. Note that the common evidence E can only be as strong as the *disjunction* of the evidence peculiar to both cases: for instance, either there is something red and spherical before Austin, or such-and-such facts about sense-data obtain.

(Further note that the disjuncts should not concern the particular tomato—that *this* looks red and spherical—or *particular* sense-data—that *this* is red' and circular'. The first proposition is not entertainable by Austin in the bad case, in which the tomato is entirely absent; the second is not entertainable in the good case, in which sense-data are entirely absent. Hence a disjunction with either of these disjuncts cannot be E, because it is not available as evidence in both cases.)

Second, and more heroically, it must be established that E is *good enough* evidence—in the good case, Austin knows that there is something red and spherical before him *on the basis of E*, and not *solely* on the basis of evidence peculiar to the good case. If uphill struggles are sensibly avoided, the road from metaphysical disjunctivism will lead to epistemological disjunctivism.

Here are two distinctions between varieties of metaphysical disjunctivism, both of which we touched on earlier, and which we will appeal to in what follows. One is a distinction between a disjunctivism that classifies the illusory cases with the good case—*Veridical Illusory v Hallucinatory (VI v H)* disjunctivism—and one that classifies them with the hallucinatory case—*Veridical v Illusory Hallucinatory (V v I/H)* disjunctivism. The other distinction concerns what can be said about the right-hand side of the disjunction. According to *positive* (metaphysical) disjunctivism: “there may be available a more direct characterization of the second disjunct, and in a totally explicit version of the theory it would indeed be characterized in that better way” (Dancy 1995: 436). That is, there is a specific kind of mental state or event that characterizes the right-hand side of the disjunction (and, hence, according to the disjunctivist, not the left-hand side). Austinian-cum-Hintonesque disjunctivism is an example, where the hallucinatory case supposedly involves acquaintance with sense-data. According to *negative* (metaphysical) disjunctivism, there is (at least sometimes) no such further characterization of the right-hand side. Martin is the main proponent of negative disjunctivism: certain hallucinations, he thinks, have “no positive mental characteristics other than their epistemological properties of not being knowably different from some veridical perception” (2004: 82).²³

²³ It is pretty clear that Hinton's view is similar: “A perfect . . . illusion of a given perception or other reality should be thought of as follows. It is an illusion of that reality, such that if you

Disjunctivism can be brought into sharper focus by considering some alternatives. The polar opposite position is what we may call the *Cartesian view*, which takes indiscriminable experiences to be mentally exactly alike, and conversely:

CV Experience e is mentally exactly alike experience e^* iff e is indiscriminable from e^* .²⁴

Let us assume that ‘ a is indiscriminable from b ’ is to be elucidated as in Williamson (1990: ch. 1), in terms of the ‘inability to activate the knowledge that a is distinct from b ’, and that experiences are particular (‘token’) events. The Cartesian view as expressed in CV needs qualification. Even if experiences e and e^* are mentally alike, a subject may not be in a position to tell that e and e^* are distinct because e and e^* are the experiences of other people. Alternatively, even if she undergoes both e and e^* , she may not be in a position to tell that they are distinct because she has forgotten undergoing e . On the other hand, she may be able to tell that e and e^* are distinct because she can tell that they occur at different times, or because an informant says so—it would not follow that e and e^* are *mentally* different. Suitably qualified, the Cartesian view is something along these lines:

CV+ Let S be an alert conceptually competent normal human subject with an excellent memory, and let e and e^* be two of her experiences. Then: e is mentally exactly alike e^* iff S is never in a position to know by introspection alone that e is mentally different from e^* .

Let us adapt our running choice of the good case by supposing that the tomato seen by Austin is later annihilated, keeping the proximal causes of Austin’s brain states constant. In the terminology of Johnston (2004), Austin undergoes a “subjectively seamless transition” from seeing the tomato to having a tomato-hallucination. Austin is never in a position to know by introspection that his two experiences differ mentally. According to disjunctivism, they are mentally completely different; according to the Cartesian view, they are mentally exactly alike. The Cartesian view thus rejects the claim that x *sees the tomato* is a mental state.

Although the Cartesian view has some intuitive appeal, for familiar reasons it is incorrect. Consider an alert (etc.) subject S who is looking at a patch that is darkening very slowly. Initially the patch is white, and after five minutes has changed to black. An empirical fact: S will never be in a position to know that her present experience differs mentally from her experience a short while ago. Yet this cannot imply that the two experiences are mentally exactly alike, otherwise her initial experience (as of a white patch) would be mentally exactly alike her final experience (as of a black

are involved in the illusion then you cannot tell, simply by being involved in it, that it is not that reality” (1973: 145; see also 76).

²⁴ Compare Millar (1996: 77): “the very idea of an experience is such that if A has an experience E and B has an experience E' , and being in E is experientially indistinguishable from being in E' , then E and E' are the same experience”. The epistemological expression ‘indistinguishable’ strongly suggests something like CV, but we are not sure that this is Millar’s intent.

patch), which is absurd. The problem is, of course, that the Cartesian view has the obtaining of an intransitive relation defined over S's experiences (i.e., *indiscriminability in mental respects by S*) being a necessary and sufficient condition for the obtaining of a transitive one (i.e., *sameness in mental respects*).²⁵

Why would one think that some bad cases are mentally *exactly* alike the good case? If there is some justification for this other than the Cartesian view, it is not clear what it is. So, given that *x sees the tomato* and *x sees that the tomato is red and spherical* are, on the face of it, mental states in good standing, opponents of disjunctivism should adopt the *moderate view*: the good case and the bad cases (both illusory and hallucinatory) are different in significant mental respects, despite having a common mental element.

What is that common element? If we set sense-data aside, then there are two options: first, qualia, taken to be introspectible non-intentional properties of experiences; second, the representational content of experience. Since the idea that experience has representational content is a good deal less controversial than the idea that experiences have (non-intentional) qualia, we will exclusively consider the second option.

There are two ways of spelling out a content-based version of the moderate view. One is *abstract intentionalism*: the representational content of visual experience is not object-dependent (at least with respect to objects in the perceiver's environment). (That is, the content is "abstract" in the sense of Tye 2000: 62.) For the sake of a simple example, we may pretend that, according to abstract intentionalism, the content of Austin's experience in the good case is that $\exists x(x$ is red, spherical, and before Austin). In the bad cases, the content is exactly the same. This is not to say that the good case and the bad cases are mentally exactly alike—for that, sameness of representational content is necessary but not sufficient. The abstract intentionalist, we may suppose, holds that Austin is in the mental state of seeing that the tomato is red in the good case—Austin is not in this state in any hallucinatory case, and is not in this state in any illusory case where the tomato is a different colour.

The other moderate view is *particular intentionalism*: the representational content of visual experience is object-dependent. For the sake of a simple example, we may pretend that, according to particular intentionalism, the content of Austin's experience in the good case includes the proposition that Tom is red, spherical, and before Austin, where Tom is the seen tomato. In the *illusory* cases, the content also includes this proposition. However, this proposition can hardly be included in the content in every hallucinatory case—in many of these, Tom does not even exist. It might thus appear that the particular intentionalist is forced to adopt a version of VI v H

²⁵ Graff argues that, "contrary to widespread philosophical opinion, phenomenal indiscriminability is transitive" (2001: 905). It might superficially appear that Graff's argument is at odds with the claims of the above paragraph. This would be a misreading, however. Graff's conclusion is simply that "*looking the same as* (in a given respect)" (2001: 905) is transitive. As Graff points out, this is not (or not obviously) an epistemic relation; in particular, it is not the relation of indiscriminability as characterized by Williamson (see Graff 2001: 910 and fn. 6). For helpful further discussion of indiscriminability, see Hellie (2005).

metaphysical disjunctivism: there is a common element across the good case and the illusory cases, but not across hallucinatory cases in general. However, as we will explain in section 9, particular intentionalists need not be friends of disjunctivism.

There are three contenders in the ring, then: disjunctivism, abstract intentionalism, and particular intentionalism. We now turn to the claim that disjunctivism should start the contest as the odds-on favourite.

6 METAPHYSICAL DISJUNCTIVISM AS THE DEFAULT VIEW

In the following section, we assess Martin's elaborate argument for disjunctivism; in this section we examine a weaker strategy, designed to secure disjunctivism as the starting point of any philosophical exploration of perception, if not its final destination. Since we will be conspicuously ignoring Snowdon's important early papers, a word of explanation is in order.

Snowdon is sometimes classed with Hinton and McDowell as a disjunctivist (for example see Martin 2004: 43). But he has not endorsed it in print, at least. The main burden of 'Perception, Vision, and Causation' is to argue against the causal theory of perception, considered as "in some sense, a conceptual truth" (Snowdon 1980–1: 176). In that paper, Snowdon argues as follows: if disjunctivism is true, then the causal theory is false; disjunctivism is not a *conceptual* falsehood; hence the causal theory is not a conceptual truth.²⁶ As he points out at the end of his paper, "an assertion of the disjunctive theory is not needed for a rejection of the present causalist view" (1980–1: 191), and indeed he makes no such assertion. Snowdon reaffirms his lack of public commitment to the disjunctivist cause in this recent retrospective comment:

There is the issue of whether disjunctivism, so explained, is true or false. But there is also the issue whether efforts at conceptual analysis of perceptual concepts are entitled to assume the falsity of disjunctivism. My own earlier papers in effect were arguing that they were not so entitled. (2005: 137, fn. 15)

6.1 Hinton's challenge

Hinton takes disjunctivism to embody the default view of perceptual experience. In 'Phenomenological Specimenism' he writes:

Although I contend that no neutral experience report satisfies Condition PS [that is, refers to "psi-ing"—a kind of mental event that occurs in both the good and bad cases], I do not undertake to prove that. I am not adopting the method of pointing to something that I can demonstrate to be a mistake, and suggesting or showing that it is common to a certain class of doctrines. Instead I am just taking a stand until overwhelmed by contrary reasoning. (Hinton 1980: 38)

²⁶ For an argument that disjunctivism is compatible with the causal theory, see Child (1994: ch. 5).

A clue to Hinton's motivation may be found in passages like the following, where Hinton describes the alternatives to disjunctivism:

The view that in all cases of seeing (and visual hallucination) we actually see a mental picture or image is now widely rejected, and I for one would also want to reject the marginally different view that a sort of as-it-were-picture-seeing occurs as a common constituent of illusion and true perception. (Hinton 1967b: 10)

Hinton seems to be assuming that if disjunctivism is false, then either the sense-datum theory or a marginally different as-it-were-sense-datum theory is true.²⁷ Of course, this assumption would be rejected today (see section 5). Arguably, whatever the default view is, it should *rule out* the sense-datum theory—in which case, given Hinton's assumption, the prize goes to disjunctivism. But this is only a reconstruction of Hinton's motivation, because he gives no explicit argument.

6.2 Martin's reinforcement of Hinton's challenge

In 'The Limits of Self-Awareness' Martin undertakes to provide what Hinton did not, an argument that disjunctivism is innocent until proven guilty:

When Michael Hinton first introduced the idea, he suggested that the burden of proof or disproof lay with his opponent, that what was needed was to show that our talk of how things look or appear to one to be introduced more than what he later came to call perception-illusion disjunctions . . . The aim of this paper is [in part] . . . to explain the way in which Hinton was correct in his challenge. Properly understood, the disjunctive approach to perception is the appropriate starting point for any discussion of the nature of perceptual experience. (2004: 37–8)

According to Martin, the disjunctivist's opponent, but not the disjunctivist, is committed to a "substantive epistemic principle" (2004: 50), and hence a non-disjunctivist position "carries more theoretical burdens" (2004: 51). It will be convenient to divide Martin's overall argument into three parts. The first part defends a claim about a "conception of perceptual experience" (2004: 48) that, Martin thinks, should be accepted on all sides. The second part characterizes the disjunctivist's opponent. And the third part purports to show that the first part's "conception of perceptual experience", together with the characterization of the second part, implies that the disjunctivist's opponent must embrace the "substantive epistemic principle".

According to the first part's conception of perceptual experience, "some event is an experience of a street scene just in case it couldn't be told apart through introspection from a veridical perception of the street as the street" (2004: 48). Or, as Martin puts it in 'On Being Alienated':

The notion of a visual experience of a white picket fence is that of a situation being indiscriminable through reflection from a veridical visual perception of a white picket fence as what it is. (Martin 2006: 363)

²⁷ A slightly different—and perhaps better—interpretation: the appeal of the "common element" is a hangover from the sense-datum theory, and should not have survived the theory's demise (compare Hinton 1973: 66).

A situation is “indiscriminable through reflection [introspection] from a veridical visual perception of a white picket fence” just in case the subject in the situation is not in a position to know via introspection that she is not having a veridical visual perception of a white picket fence (Martin 2006: 364–5).²⁸ It will be helpful to represent this using the apparatus of cases. A case, recall, is a centred possible world—a world with a marked subject and time. Let ‘ $E_p(\alpha)$ ’ abbreviate ‘in case α , the subject has a visual experience as if p ’, ‘ K_α ’ abbreviate ‘in α , the subject is in a position to know via introspection that’, and ‘ V_p ’ abbreviate ‘the subject (veridically) visually perceives that p ’.²⁹ Then the claim of the first part may be expressed as follows:

Exp $E_p(\alpha)$ iff $\sim K_\alpha \sim V_p$

(Compare (Martin 2006: 11); in Exp and the theses below, universal quantification over cases is tacit if it is not explicitly supplied.³⁰)

There are some pressing and difficult questions about the exact intended interpretation of Exp and whether, so interpreted, it is true—let alone whether it is common ground between the disjunctivist and her opponent. But here we can simply work with an intuitive understanding of Exp, leaving it largely unchallenged.³¹

²⁸ Compare the explanation of a ‘subjectively indistinguishable case’ in section 1. This notion of indiscriminability is like Williamson’s in its use of ‘not knowably not’; it is unlike Williamson’s in that the relevant items of knowledge are not negations of identities (‘this apple \neq that orange’) but rather negations of predications (‘this apple is not an orange/lacks the property *being an orange*’) (Martin 2006: 363–4, fn. 15). It would therefore be a mistake to take Martin’s notion of indiscriminability to be a relation between particular experiences; see Sturgeon (2006). (Having said that, sometimes Martin does talk of two perceptual episodes being indiscriminable, for example his 2004: 76; compare Siegel 2004: 109, fn. 5.)

A terminological note: Martin uses ‘perception’ and ‘perceives’ so they apply only to the good case, not the hallucinatory cases; we have followed him in this.

²⁹ ‘At the time of α ’ has been left tacit; ‘the subject’ is stipulated to refer directly to the subject of the case when it appears in the scope of ‘knows’.

³⁰ Thus (Exp) is a universally quantified schema, with ‘ p ’ being the sole schematic letter.

³¹ Most of the literature on Martin’s version of disjunctivism has focused on Exp, which as Martin recognizes, has some problems. The most immediate are these:

Right-to-left:

(RLi) Due to “excess of alcohol or lack of interest in the matter” (2004: 76; cf. 2006: 380), the subject may not be in a position to know that she is not veridically perceiving a tomato, which we may suppose she isn’t. Yet Exp implies that she is having an experience of a tomato.

(RLii) Due to conceptual incapacities, a dog (2004: 76; 2006: 379) is not in a position to know that he is not veridically perceiving a tomato, which we may suppose he isn’t. Yet Exp implies that he is.

(RLiii) Similarly, a comatose subject, or an “insentient stone” (2006: 383), is not in a position to know that she/it is not veridically perceiving a tomato. Yet Exp implies that she/it is.

(RLiv) ‘Zombies’, notoriously, are not in a position to know that they are entirely experience-free. Exp implies wrongly that the idea of a zombie is incoherent (compare 2006: 375–6).

(RLv) Arguably, when one has an experience solely of a determinate shade (red₁₉, say), one is not in a position to know that one is not experiencing (a fortiori not veridically perceiving) a slightly different shade (red₁₈ or red₂₀, say) (Williamson 2000: ch. 4). In which case, according to Exp one is *not* having an experience solely of red₁₉.

Now to the second part. Martin “further elaborates” the position of the disjunctivist’s opponent as follows:

A perceptual experience is a kind of event which has certain distinctive features $E_1 \dots E_n$ [e.g. “being the presentation of such and such mind-dependent qualities, as the sense datum theory supposes. Or we might instead take them to be representational properties, as an intentional theory would press”] . . . [a] the possession of these features [is] necessary and sufficient for an event to be an experience . . . [b] in addition, an event’s possession of them is introspectible by the subject of the experience. (Martin 2004: 47)

It is clear from the context that the “distinctive features $E_1 \dots E_n$ ” are supposed to be necessary and sufficient, not for being a perceptual experience in general, but for being a relatively specific kind of perceptual experience—“as of a street scene” (49), to use Martin’s main example.³² If we take the opponent to hold the schematic thesis that features $E_1 \dots E_n$ are necessary and sufficient for having an experience as if p , and let ‘ $F_p(\alpha)$ ’ abbreviate ‘in case α , the subject is undergoing an event with features $E_1 \dots E_n$ ’, then the opponent holds:

(a) $E_p(\alpha)$ iff $F_p(\alpha)$

The condition that the features are “introspectible by the subject of the experience” seems to amount to the claim that (necessarily) when they are present, the subject is in a position to know that they are (cf. Martin 2006: 390). So, letting ‘ F_p ’ (without the attached ‘ (α) ’) abbreviate ‘the subject is undergoing an event with features $E_1 \dots E_n$ ’, the opponent also holds:

(b) if $F_p(\alpha)$ then $K_\alpha F_p$

As with the first part, various questions of interpretation arise, but let us set them to one side and see how Martin puts the two parts together to force the opponent to accept the “substantive epistemic principle”.

Left-to-right:

(LR) “Consider an hallucination of an Escher-like scene with an impossible staircase” (2004: 80); one presumably could know by introspection that one was not *veridically* perceiving such a staircase, in which case Exp implies that one is not hallucinating an Escher-like scene after all; see also Siegel (2004: 94).

In ‘The Limits of Self-Awareness’, Martin responds to (RLi) and (RLii) by appealing to an “impersonal” conception of indiscriminability (2004: 75). In ‘On Being Alienated’, he gives the same reply to (RLi) (2006: 381); the reply to (RLii) (which, we take it, contains the seeds of his reply to (RLiii)) still appeals to impersonal indiscriminability, although the dialectic is considerably more serpentine (2006: 383–96). For criticism see Hawthorne and Kovakovich (2006) and Siegel (2004); for a more sympathetic view see Sturgeon (2006). Martin responds to (RLiv) as one might expect, by denying the zombie intuition (2006: 378–9). For his response to RLv and related problems see his (2004: 76–9) and (1997: 98–9). Here there are complications because in ‘The Limits of Self-Awareness’ Martin explains the difficulty in terms of the “intransitivity of indiscriminability”, which does not apply to indiscriminability as Martin explicitly characterizes it in ‘On Being Alienated’ (see footnote 29 above, and the corresponding main text). For criticism see Hawthorne and Kovakovich (2006). Martin’s response to (LR) is to switch to talk of the indiscriminability of “aspects of the experience” from aspects of a veridical perception (2004: 81). (Compare Siegel 2004: 110, fn. 6.)

³² *Pace* Siegel (2004: 102); of course, having *some* “distinctive features” or other of this kind is supposed to be necessary and sufficient for being a perceptual experience in general.

The argument of the third part begins by noting that Exp and (a) imply that “there cannot be any situation which is indistinguishable for its subject from actually perceiving a street scene and which lacks the relevant properties [the “distinctive features”]” (2004: 50). That is, Exp and (a) imply that there is no case α such that $\sim K_\alpha \sim V_p$ and $\sim F_p(\alpha)$, or, equivalently, for all cases α :

- (1) if $\sim F_p(\alpha)$ then $K_\alpha \sim V_p$

Martin’s argument then proceeds as follows:

[i] In turn, one must assume that a subject couldn’t but be in a position to discriminate a situation which lacked $E_1 \dots E_n$ from one which possessed [them]. [ii] Here I just assume that for one situation to be indiscriminable from another requires only that it not be possible to know that it is distinct in kind. [iii] Therefore to deny it is possible that a situation which is distinct in kind from an event possessing $E_1 \dots E_n$ is not possibly knowable as distinct in kind, is to claim that for any situation distinct in kind from an event possessing $E_1 \dots E_n$ it is possible to know that it is distinct. (Martin 2004: 50, endnote omitted)

- (i) can be recast as follows:

- (2) necessarily, a situation lacking $E_1 \dots E_n$ is discriminable from a situation with $E_1 \dots E_n$

(ii) gives the operative account of indiscriminability, which we explained above when introducing Exp.³³ Given that, (2) amounts to:

- (3) if $\sim F_p(\alpha)$ then $K_\alpha \sim F_p$

And (iii), despite the misleading ‘therefore’, simply notes a logical equivalence:

- (4) $\sim \exists \alpha (\sim F_p(\alpha) \text{ and } \sim K_\alpha \sim F_p) \text{ iff } \forall \alpha (\text{if } \sim F_p(\alpha) \text{ then } K_\alpha \sim F_p)$

Martin’s argument ends here, with the following paragraph stating the conclusion:

Adopting this position is to attribute a privileged epistemic position to the subject of experience. For, according to it, a responsible subject who wishes to determine how things are with him or herself through reflection must not only correctly identify phenomenal properties of a specific sort [the “distinctive features”] when they are present, but also they cannot be misled into judging them present when they are not. It is not merely that the properties which determine an event as an experience are held to be self-evident on this view—that the presence of such properties indicates to the subject that they are present when they are present. It must also be the case that the absence of such properties when they are absent is equally detectable by the subject, so that there is always some way that a subject could tell that he or she was not so experiencing when not doing so. It is to attribute to responsible subjects potential infallibility about the course of their experiences. (Martin 2004: 50–1, emphasis ours)

³³ Restricted to perceptual examples, (ii) amounts to this:

A situation (case) α lacking F_p is *indiscriminable* from any case with F_p iff $\sim K_\alpha \sim F_p$

Equivalently:

A situation (case) α lacking F_p is *discriminable* from any case with F_p iff $K_\alpha \sim F_p$

We can now identify the substantive epistemic principle to which the disjunctivist's opponent is committed. It is "potential infallibility about the course of [experience]", namely:

SEP if $F_p(\alpha)$ then $K_\alpha F_p$ & if $\sim F_p(\alpha)$ then $K_\alpha \sim F_p$

And SEP is derived from one claim that the opponent is said to hold at the outset, namely:

(b) if $F_p(\alpha)$ then $K_\alpha F_p$

together with:

(3) if $\sim F_p(\alpha)$ then $K_\alpha \sim F_p$

Looking back over the argument as we have unpacked it, (3) follows from (i) (see the last-but-one quotation), together with the account of indiscriminability stated by (ii). But (i) is not supported by any explicit argument.

As mentioned earlier, assuming that the opponent accepts Exp, the first part's characterization of perceptual experience, she is committed, as Martin notes, to:

(1) if $\sim F_p(\alpha)$ then $K_\alpha \sim V_p$

Although Martin does not say that (3) (or, equivalently, (2)) follows from (1), it is clear that this is his route, on behalf of the disjunctivist's opponent, to (3). Here's how we take the argument to go. For illustration, suppose the feature distinctive of a perceptual experience of a red tomato is the presentation of a red' sense-datum. Assume (1): if the subject is not having a perceptual experience of a red tomato (that is, if F_p —the presentation of a red' sense-datum—is absent), he is in a position to know that he is not having a *veridical* perceptual experience of a red tomato. Presumably the disjunctivist's opponent will say that one comes to know that one is not veridically perceiving a red tomato by inferring it from the absence of red' sense-data. But this implies that the subject is in a position to *know* that red' sense-data are absent—else the conclusion that he is not veridically perceiving would not itself be knowledge.³⁴ Hence, from (1):

(3) if $\sim F_p(\alpha)$ then $K_\alpha \sim F_p$

However, although Martin is quite correct in taking his opponent—as she is officially characterized—to be committed to a "substantive epistemic principle", this does not show that disjunctivism is the default view. Instead, it shows that the opponent's position has been characterized far too strongly.

In fact, the disjunctivist's opponent need not even accept (a), the commitment to necessary and sufficient conditions for perceptual experience, at least if this is

³⁴ We should read 'know that red' sense data are absent', not 'de dicto', but as 'know something that is incompatible with the presence of red' sense data'—the subject need not conceive of red' sense data as such. Similar remarks go for 'the subject knows he is undergoing an event with features $E_1 \dots E_n$ '. Compare Siegel (2004: 103–4).

understood in a reductive way, which is apparently Martin's intent.³⁵ But this is not the main complaint, so for the sake of the argument let us assume that the opponent holds (a).

(b)—one half of SEP—is just built into the opponent's position at the very start. Even granting the commitment to (a), there is no evident reason for doing this.

The commitment to (3)—the other half of SEP—is derived from the controversial Exp:

$$\text{Exp } E_p(\alpha) \text{ iff } \sim K_\alpha \sim V_p$$

Hence if the opponent rejects Exp, she is free also to reject (3). Moreover, adapting the reasoning that led us from (1) to (3), Exp all by itself seems to imply an equally substantive principle which, letting ' E_p ' abbreviate 'the subject has an experience as if p ', can be put as follows:

$$(5) \text{ if } \sim E_p(\alpha) \text{ then } K_\alpha \sim E_p$$

That is, if the subject in case α does not have a perceptual experience as if p , she is in a position to know that she is not having a perceptual experience as if p .

By Exp, whenever the subject does not have a perceptual experience of a red tomato, she is in a position to know that she is not having a *veridical* perceptual experience of a red tomato. Suppose she is not having a perceptual experience of a red tomato—she is merely imagining a red tomato, say.³⁶ It is hard to see how the subject could be in a position to know that she is not *veridically* perceiving a red tomato, if she was not also in a position to know that she is not *perceptually experiencing* a red tomato. Any indication that she is not veridically perceiving, only imagining, is presumably also an indication that she is not perceptually experiencing. Hence (5) is true. And since (according to Martin) Exp is a commitment of *disjunctivism*, that theory carries no lighter epistemological baggage than some of its rivals. The argument for disjunctivism's default status therefore fails.

7 MARTIN'S ARGUMENT FOR DISJUNCTIVISM

We now turn to the most sophisticated and detailed argument for disjunctivism in the contemporary literature, developed by Martin in a series of rich and intricate papers.³⁷

³⁵ Knowledge is a useful analogy here: according to the anti-disjunctivist about knowledge, belief is a common element to 'the good case' (knowing) and 'the bad case' (merely believing)—this does not commit her to giving reductive necessary and sufficient conditions for knowing. Compare Williamson (2000: 32–3).

³⁶ Assuming, with Martin (2002a: 413–14), that imagining is not to be classified as perceptually experiencing.

³⁷ For reasons of space, we are ignoring the argument for disjunctivism that appears in Martin (2002a: 402–19) (although not in Martin's most recent writings). This argument, which turns on issues about sensory imagining, deserves a lengthy treatment. Another equally important omission from this paper is any discussion of the pro-disjunctivist case made in Campbell (2002: ch. 6).

Unlike the previous argument, it does not merely purport to show that the disjunctivist's opponent "carries more theoretical burdens", but rather attempts to establish disjunctivism outright. As before, a division of the argument into parts will be convenient. The first part is an argument for the denial of something Martin calls the "Common Kind Assumption"; the second part is an argument from the denial of the common kind assumption to disjunctivism—specifically, *negative* disjunctivism.

7.1 Part 1: from naïve realism to the denial of the common kind assumption

"The prime reason for endorsing disjunctivism", Martin says, "is to block the rejection of a view of perception I'll label *Naïve Realism*" (2004: 38).

According to naïve realism (NR):

[I]t is an aspect of the essence of such experiential episodes [e.g. episodes of seeing tomatoes] that they have . . . experience-independent constituents [e.g. tomatoes]. (Martin 2006: 357)

Or:

Some of the objects of perception—the concrete individuals, their properties, the events these partake in—are constituents of the experience. No experience like this, no experience of *fundamentally the same kind*, could have occurred had no appropriate candidate for awareness existed. (Martin 2004: 39, emphasis ours)

"Fundamental kind" is an important term of art for Martin:

I will assume the following: entities (both objects and events) can be classified by species and genus; for all such entities there is a most specific answer to the question, 'What is it?' In relation to the mental, and to perception in particular, I will assume that for mental episodes or states there is a unique answer to this question which gives its most specific kind [its "fundamental kind" (2004: 43)]; it tells what essentially the event or episode is. (2006: 361, footnote omitted)³⁸

Thus, naïve realism holds that the experience in the good case (veridically perceiving a tomato, say) is of a fundamental (essential) kind K such that no experience in the absence of tomatoes could be of kind K.

According to Martin, "Naïve Realism is inconsistent with two assumptions which are common to much of the philosophical discussion of perception" (2004: 39).

The first is *experiential naturalism* (EN), that "our sense experiences, like other events or states within the natural world, are subject to the causal order, and in this case are thereby subject just to broadly physical causes (i.e. including neurophysiological causes and conditions) and psychological causes (if these are disjoint from physical causes)" (2004: 39–40; compare 2006: 357).

The second is the *common kind assumption* (CKA), that "whatever [fundamental] kind of mental event occurs when one is veridically perceiving some scene, such as

³⁸ There is one occurrence of 'state' in this quotation, which is *prima facie* at odds with the restriction to "entities (both objects and events)". In any case, Martin's official position is that experiences are events (2006: 354).

the street scene outside my window, that kind of event can occur whether or not one is perceiving” (2004: 40; compare 2006: 357). To illustrate this, return to the good case mentioned three paragraphs back, in which the veridical perception’s *fundamental* kind is K. According to CKA, an experience of kind K can occur even if one is not seeing a tomato, but merely hallucinating one.

The argument for the inconsistency, in brief, is (roughly) this. Let *e* be the experience occurring in the good case. By NR, *e*’s fundamental kind, K, is such that no event of kind K could occur in the absence of tomatoes. By CKA, the experience in an hallucinatory case is of kind K; somehow tomatoes must be around in hallucinatory cases, even though they are not seen. But, by EN, an hallucinatory case can obtain “through suitable manipulation of mind and brain” (Martin 2006: 358); that is to say, in the absence of tomatoes. Contradiction.

Let us grant the inconsistency; one of the assumptions therefore has to go. Not NR, because it “best articulates how sensory experience seems to us to be just through reflection” (2006: 354; see also 2004: 42). And not the anodyne EN either; hence CKA has to be thrown overboard.

Now Martin himself can be read as suggesting that not even a minimal version of disjunctivism is established at this stage: so far, we have “not yet captur[ed] a key thought behind disjunctivism” (2006: 368). However, that same paper (‘On Being Alienated’) begins by saying that “Disjunctivism . . . seeks to preserve a naïve realist conception of veridical perception . . . the disjunctivist claims that . . . the experiential episode [in the good case] is of a kind which could not be occurring were you having an hallucination” (2006: 354), which gives the impression that NR&~CKA suffices for disjunctivism. And in ‘The Limits of Self-Awareness’ Martin describes someone who “resist[s] the problem of perception in a way that retains both Naïve Realism and Experiential Naturalism” as “the disjunctivist” (2004: 53). Whatever Martin’s considered view may be, *Snowdon*’s current position is that NR&~CKA (or something close to it) is the kernel of disjunctivism.³⁹ In the article mentioned just before section 6.1 above he characterizes disjunctivism as follows:

The experience in a perceptual case reaches out to and involves the perceived external objects, not so the experience in other cases. (Snowdon 2005: 136–7)

This formulation has two parts. The first is a claim about “perceptual cases” (where one sees an object like a tomato). This is more-or-less naïve realism. The second part of Snowdon’s formulation of disjunctivism is that in the “other cases” (where one hallucinates a tomato), the experience does *not* “reach out and involve” any tomatoes, or indeed any “external object”. This is the claim that the naïve realist fundamental kind K is not instantiated in hallucinatory cases. With the harmlessly oversimplified assumption that K will be instantiated in hallucinatory cases just in case both NR and CKA are true, the second part of Snowdon’s formulation amounts

³⁹ See also Crane (2005: section 3.4); Soteriou (2005: 178). In McDowell (1986/1998: 204, emphasis ours) the “*innocuous* disjunctive conception of subjective appearances” is a watered-down version of NR&~CKA.

to $NR \rightarrow \sim CKA$. Hence the conjunction of the two parts is (approximately) equivalent to $NR \& \sim CKA$.

Although terminological squabbles are not very profitable, there is good reason for not allowing $NR \& \sim CKA$ to count as a kind of disjunctivism. As Martin points out (2004: 54, 58, 60; compare 2006: 360–1), $NR \& \sim CKA$ is consistent with the experiences in the good and bad cases falling under the *same* (psychological) kind K^\dagger —all that is ruled out is that K^\dagger is the *fundamental* kind of experience occurring in the good case. And if disjunctivism allows that there *is* a mental element common to the good and bad cases, then *Hinton's* characterization of the view is incorrect—an undesirable result.⁴⁰ And in any case, a common element makes ‘disjunctivism’ an inapposite label.⁴¹

So, we do not yet have an argument for disjunctivism.⁴² But we *do* have the first part of Martin’s overall case for disjunctivism, which we will now examine.

Although the argument from naïve realism to the denial of the common kind assumption is not entirely plain sailing, the main problem is with the premise. First, notice that naïve realism, as Martin explains it, is not the usual version, which is not a claim about the essence of some of our sensory episodes.⁴³ In fact, Martin’s opening statement of the position in ‘The Limits of Self-Awareness’ is much weaker:

The Naïve Realist thinks that some at least of our sensory episodes are presentations of an experience-independent reality. (2004: 38)

Suppose the quotation is elaborated by saying (as Martin does) that “[s]ome of the objects of perception—the concrete individuals, their properties, the events these partake in—are constituents of the experience” (2004: 39). Still nothing follows about the essence of experiences—presumably one may hold that the tomato is a “constituent” of one’s experience, and that *this very experience* could have occurred in the absence of the tomato. Why does the naïve realist feel the need to go further, and proclaim that “[n]o experience like this, no experience of fundamentally the same kind, could have occurred had no appropriate candidate for awareness existed” (2004: 39)?

⁴⁰ Snowdon, incidentally, thinks that the “common element” view opposed by Hinton is that the experiences in the good and bad cases “have the same nature and, therefore, do not reach out to, or involve as constituents, items external to the subject” (2005: 136). But Hinton went far beyond denying that the experiences have the same nature: see footnote 23.

⁴¹ The analogue of NR for knowledge is that (a) in ‘the good case’ (where the subject knows) knowing is the subject’s “fundamental” mental state. The analogue of $\sim CKA$ is that (b) in ‘the bad case’ (where the subject merely believes), the subject doesn’t know. This is consistent with believing being an element common to both cases. This overall position, including believing as a common element is, more-or-less, Williamson’s. And it doesn’t seem to be a kind of disjunctivism—at any rate, Williamson explicitly denies that it is (2000: 47–8), and Martin agrees (2006: 370–1; compare his 1997: 88–9).

⁴² It is worth pointing out that naïve realism, at least as officially explained by Martin, is not even a *necessary* component of disjunctivism (compare Martin 2002a: 395, fn. 25). That Austin’s experience in the good case lacks a “common element” does not imply that it has the tomato as a *constituent*.

⁴³ Naïve realism is often taken to be (approximately) the claim “things are what they seem” (Russell 1950: 15). See also, for example, Price (1950: 26–7); compare Langsam (1997: 53); Robinson (1994: ch. 2); Smith (2002: 44).

Put another way: suppose that one is forced to *deny* this claim about essence—why would that be bad? The denial is perfectly consistent with the claim that we see tomatoes and never see sense-data, that nothing mental or immaterial ‘intervenes’ between us and tomatoes, and so on. Let the essences look after themselves—we can still perceive the world as ‘directly’ as we would wish.

Thus, it is unclear why we need to go to the mat to rescue Martin’s industrial-strength naïve realism. But the problem is deeper. According to Martin, naïve realism is “the best articulation of how our experiences strike us as being to introspective reflection on them” (2004: 42).⁴⁴ And part of the motivation for ‘direct realist’ views (of which Martin’s naïve realism is a species) is, as he notes, that “our sense experience is transparent— . . . experientially we are presented with a mind-independent realm and not simply some array of mind-dependent qualities or entities whose existence depends on this awareness” (2004: 39). The usual transparency claim is that in undergoing a sense experience, one is never aware *of* the experience itself—if one is aware of any events at all, they are events in one’s environment, like flashes and bangs, not events in the head. Indeed—as stressed in Dretske (2003)—transparency makes it extremely hard to see how one knows that one has experiences in the first place. But at the very least, the transparency of experience fits nicely with the view that in having an experience of, say, a tomato, although one may be in a position to learn something about the essence of the *tomato*, one is not in a position to learn much of anything about the essence of the *experience*.

Something stronger may be said. We have so far played along with talk of ‘experiences’, conceived of as particular events or episodes, like collisions or cocktail parties, about which various theories may be offered. Following this section, we will continue to play along—and certainly for many purposes such talk is harmless. But are there really any such episodes? Undeniably, we *do* see tomatoes. But why think that when one sees a tomato something happens that should be labelled one’s ‘experience of the tomato’? For the naïve realist, the reason cannot be because we notice our experiences, as we might notice a flash of light or a car backfiring—the proponent of transparency should deny that we are aware of any such things. The experience cannot be one’s ‘seeing the tomato’: seeing itself is not an event, episode, or process (Vendler 1957).⁴⁵ Obviously vision involves various processes beginning at the tomato and continuing into the visual pathways. However, we can only label such a process ‘the experience’ if we can be assured that the subject is generally in a position to know that it occurs. True, sometimes accessible perceptual events *will* occur: if one scrutinizes or observes the tomato, then (if grammar is any guide) there is an event of scrutinizing or observing, which is available for one to speculate about. But what if one *doesn’t*

⁴⁴ It should be emphasized that Martin does not take naïve realism to be particularly naïve: it is a sophisticated philosophical gloss on, or elaboration of, what Martin thinks “we all pre-theoretically accept concerning the nature of our sense experience” (2006: 404). Martin himself does not take naïve realism to be obvious or uncontroversial, remarking that “[t]he commitment to Naïve Realism is probably not shared by most readers” (362). If there is a properly so-called ‘naïve’ picture of sense experience, it is probably the seventeenth-century view that we only perceive our own ‘ideas’.

⁴⁵ See also Bennett (1988: 4–9), who argues (following other work of Vendler’s) that “imperfect nominals” like ‘Austin’s seeing the tomato’ name facts, not events.

scrutinize or observe the tomato, but merely sees it? One may not be in a position to know much about one's experiences, conceived of as particulars (specifically, events, or something similar), simply because there aren't any.⁴⁶

Transparency goes naturally with *modesty about experience*—we know little of the nature of experiential episodes.⁴⁷ What's more, it goes naturally with *scepticism about experience*—there are no experiential episodes to begin with. Hence naïve realism, in so far as it is motivated by transparency, is a position with some serious internal tensions.

7.2 Part 2: the argument from explanatory redundancy

So far, let us suppose, NR&~CKA has been established. That is, if the fundamental kind of experience in the good case is K, then K is a "Naïve Realist" kind (by NR), and is not instantiated in the hallucinatory case (by ~CKA). But now, Martin argues, there is a problem. Consider an hallucination with the same proximal causes as the perception, and suppose that the hallucination falls under a mental kind, K†. Given sameness of proximal causes, the perception also falls under K†. But then K† would seem to pre-empt K in the good case, rendering it explanatorily idle. Since K is *not* explanatorily idle, we are forced to deny that the hallucination falls under K†. The most that can be said about the hallucination is that it has the negative epistemological property of being not knowably different from the perception—negative disjunctivism is true.

Sketched slightly less roughly, the second part of the argument itself divides into three parts—2(a), 2(b), and 2(c):

Part 2(a)

Assume that the experience in the hallucinatory case is "brought about through the same proximate causal conditions as [the] veridical perception [in the good case]"—that the experience is a "causally matching hallucination" (2006: 368). Let K† be a kind that subsumes the experience in the hallucinatory case. Then (by the argument of this part):

C(a): the experience in the good case is also of kind K†.⁴⁸

(As noted before, this conclusion is consistent with ~CKA.)

⁴⁶ Rather ironically, Hinton himself approaches this point in Part I of *Experiences*; he does, however, think that "my experience of seeing a flash" is an "event", albeit in some exceptionally thin sense (1973: 30).

⁴⁷ Martin describes his own view of perceptual experience as "modest", because it eschews commitment to any metaphysical characterization in terms of representational properties, awareness of sense-data, and the like (2004: 48). This contrasts with his decidedly *immodest* view of what introspection can tell us about the nature of veridical perceptions.

⁴⁸ It would be more accurate to say that the conclusion is that *some event* (perhaps not the veridical experience with K) in the good case has K† (see Martin 2004: 59–60). We are using the stronger formulation for simplicity. (This part of Martin's argument is adapted from Robinson 1995, and Robinson 1994: ch. 6.)

Part 2(b)

Let the fundamental kind of experience in the good case be K (a naïve realist kind, from Part 1). By $\sim CKA$ (also from Part 1), the experience in the hallucinatory case is not of kind K ; hence $K \neq K^\dagger$. Suppose for reductio that kind K^\dagger is a “positive mental characteristic” (2004: 73) (specified in terms of representational properties or sense-data, say). By 2(a), the experience in the good case is of kind K^\dagger . And (by the argument of this part):

C(b): K^\dagger makes “the Naïve realist aspects of the perception [i.e., K] . . . explanatorily idle” (2004: 71).

Part 2(c)

The experience in the hallucinatory case has the “negative epistemological property” (2006: 398) K^e of being indiscriminable from (not knowably distinct from) the experience in the good case: in other words, in the hallucinatory case, the subject is not in a position to know that he is not in the good case. Trivially, the experience in the good case also has this negative epistemological property. But (by the argument of this part):

C(c): Although K^e can indeed do some explaining, it does *not* make K explanatorily idle.

So, putting the three parts together, and assuming that the “Naïve realist aspects” of the good case are *not* “explanatorily idle”, the hallucinatory case does *not* fall under any such positive mental kind K^\dagger :

[T]here are certain mental events, at least those hallucinations brought about through causal conditions matching those of veridical perceptions, *whose only positive mental characteristics are negative epistemological ones*—that they cannot be told apart by the subject from veridical perception. (2004: 73–4, emphasis ours)

That is, *negative* disjunctivism is true.⁴⁹

As Martin emphasizes, the restriction to “causally matching hallucinations”—those hallucinations that share the same proximate causal conditions as veridical

⁴⁹ ‘In The Limits of Self-Awareness’, Part 2(a) is set out in the form of premises and conclusion on pp. 53–4. Parts 2(b) and (c) are more discursive. Part 2(b) starts with the second paragraph on p. 60 (“But this is not the only way . . .”), and ends with “. . . as not common to perception and hallucination” in the last paragraph on p. 64. The Part 2(c) argument for the explanatory virtues of the indiscriminable-from-the-good-case property starts at the top on p. 65 and ends (bar a few loose ends) with the first paragraph of section 7 on p. 68. The argument that the indiscriminable-from-the-good-case property does not screen off K starts with the first full paragraph on p. 69 and runs to the end of section 7 on p. 70. However, given that Martin concedes that the screening off argument of Part 2(b) doesn’t work as stated (see below), the second argument of Part 2(c) is apparently not needed.

In ‘On Being Alienated’, an abbreviated version of Part 2(a) is on pp. 368–9. On p. 370 it is stated without argument that “a threat of explanatory pre-emption of the common feature [K^\dagger] overcomes the claim of that which is peculiar to the case of veridical perception [K]”, but the explicit Part 2(b) argument only gets started at the bottom of p. 370 (“It is instructive to compare . . .”) and ends at the top of p. 372 (“ . . . is of no avail”). Part 2(c) is absent.

perceptions—is crucial. A more general version of this argument, applying to hallucinations in general, would fail at the first step (2(a)). Thus, Martin says, disjunctivists can be agnostic about how to characterize bad cases that are *not* causally matching hallucinations: “[disjunctivism] is not inconsistent with the view that there are some experiences among the non-veridical ones which fit the characterisations offered by sense-datum or intentional theories” (2004: 52).

Now one might think that if our criticisms of Part 1 are correct, then there is no need to spill more ink on Part 2, since it draws on a premise of Part 1, naïve realism, together with its conclusion, namely the denial of the common kind assumption. But this would be mistaken. Part 1 only comes into play in Part 2(b), where it is used to secure the assumption that kind K is a “Naïve realist aspect” of the experience in the good case, and is not instantiated in the bad case. No controversial claims about the essences of mental events are needed for this result: all we need is the bland assumption that the experience in the good case has K (perhaps contingently), and the experience in the bad case doesn’t. The naïve realism of Part 1 is redundant.

So, Part 2 cannot be left unaddressed. Part 2(a) deploys causal considerations to show that the good case is also of kind K^\dagger . The idea (put very roughly) is that since an experience’s having K^\dagger just depends on local neural activity (unlike an experience’s having K, which partly depends on what is in the subject’s environment), and since the relevant neural activity occurs in the good case, the experience in the good case also has K^\dagger . Even without the fine details, Part 2(a) looks promising. The role of Part 2(c) is basically to allay the suspicion that the argument of Part 2(b) proves too much, rendering K explanatorily idle no matter what. Part 2(b), then, is where the action is.

Part 2(b) opens by noting that the hallucinator’s actions can be explained:

[H]allucinations no less than perceptions are liable to coerce our beliefs and move us to action . . . (2004: 61)⁵⁰

That is, the hallucinator’s actions are (partly) explained by K^\dagger . But K^\dagger is also instantiated in the good case, so it will apparently screen off K from doing any explanatory work:

We have the same resultant phenomena in introspectively matching cases of perception and hallucination . . . the common kind of event [K^\dagger] between hallucination and perception seems better correlated with these common phenomena than the kind of event unique to perception and so seems to screen off the purely perceptual kind of event [K] from giving us an explanation. (2004: 62, endnote omitted)

Martin then says that “[t]his concludes the second step [Part 2(b)] in the argument” (2004: 63).

However, by Martin’s own lights, this does not conclude Part 2(b). In fact, it hardly commences it. For, as Martin goes on to concede, it just isn’t true that K^\dagger is better at explaining than K. As Williamson and others have (in effect) argued, the relational

⁵⁰ One might wonder how Martin’s “impersonal” notion of indiscriminability (see footnote 31), abstracting as it does from subjects’ cognitive limitations, can be used to account for this fact: compare his (2004: 68).

nature of K makes it *better* correlated with worldly outcomes than K^\dagger , and so it is better placed than K^\dagger to explain those outcomes:

Why was John able to pick up the glass that was on the table? Because he could see it, and could see where it was. (Martin 2004: 64)

An alternative explanation, appealing to the conjunction of the fact that John is facing a glass and is fortuitously hallucinating one, would not be as good, because one is much more likely to pick up a glass if one is *seeing* it. Given the complex feedback between perception and action required to pick up a glass, someone veridically hallucinating a glass is most unlikely to be able to pick it up.⁵¹

So, “one can at least rebut the challenge that the disjunctivist’s conception of sensory experience [i.e. K] is guaranteed to be explanatorily redundant” (64). In other words, the screening off argument just given doesn’t work as it stands. Moreover, once the Williamsonian point has been taken on board, surely something stronger can be said: K is, simply, *not* explanatorily idle. Taken at face value, the conclusion of the Part 2(b) argument is *false*.

However, that is evidently not how the conclusion is supposed to be interpreted. As Martin says at the start of the Part 2(b) argument, the worry about explanatory idleness is, specifically, that K plays no role in explaining “the phenomenal aspects of experience” (2004: 59). And immediately after noting the Williamsonian point about explanation, he claims that:

[t]his [Williamsonian] strategy does not address the question whether there are any common properties to the two situations which are distinctive of the subject’s conscious perspective on the world. Nor yet the question whether, if there are any, why they can only be explained by what is common to perception and hallucination rather than what is distinctive of perception.

It would be a severe limitation on the disjunctivist’s commitment to Narve Realism, if the Narve realist aspects of perception could not themselves shape the contours of the subject’s conscious experience. Yet this aim would be frustrated if we rested with the above responses [e.g., the point that *seeing the glass* is not explanatorily idle], since *so far no reason has been offered to show why we must think of the fabric of consciousness as relational, and as not common to perception and hallucination.* (2004: 64, emphasis ours)⁵²

By “the contours . . . of conscious experience” and “the fabric of consciousness”, Martin presumably means what he earlier called the “phenomenal aspects of experience”. The problem is supposed to be, then, that K^\dagger screens off K from explaining *the phenomenal aspects of the subject’s experience in the good case*. Notice that this quotation clearly implies (correctly) that such screening off is *consistent* with naïve realism. Hence, we may think of Martin’s (friendly) opponent as holding (a) naïve realism,

⁵¹ See Williamson (2000: ch. 2); Yablo (1997).

⁵² Summing up the worry at the end of the paper, Martin says that:

[I]f we assume that the causally matching hallucination is an event which represents the presence of a tree—that is, its having such representational properties are taken to explain why the experience is as it is and has the consequences that it does—then the explanation we can give of the salient features of the hallucination should equally be applicable to the case of veridical perception (2004: 71).

(b) that the naïve realist K earns its explanatory keep, and (c) that the phenomenal aspects of the experience in the good case are best explained by K^\dagger rather than K. According to Martin, this whole package is consistent. Evidently (c) is supposed to be highly problematic, but what is not clear to us is why Martin thinks this.

Pending some further clarification of why (c) is so bothersome, we may reasonably take Martin's argument to fail at exactly this point.

8 AGAINST $V \vee I/H$ METAPHYSICAL DISJUNCTIVISM

In the previous two sections we have raised some objections to Martin's argument that disjunctivism is the default view, and his direct argument for disjunctivism. It is now time to go on the offensive. In this section we will argue against $V \vee I/H$ disjunctivism; in the next section we will argue against $VI \vee H$ disjunctivism.

The argument against $V \vee I/H$ disjunctivism is extremely simple, if not simple-minded; we can give it without raising theoretical questions about the representational content of experience. According to $V \vee I/H$ disjunctivism, there is no specific mental state or event common to the good case and the illusory cases. Return to our example of Austin and the tomato: in the good case, Austin sees that the tomato is red and spherical; in an illusory case he sees the tomato, but it is not the way it looks. Consider two illusory cases: in I_1 , the tomato is red and ovoid; in I_2 , the tomato is green and spherical. In I_1 , Austin sees that the tomato is red, but misperceives the tomato as spherical; in I_2 , Austin sees that the tomato is spherical, but misperceives it as red.

We are seeking a common mental element, specific enough to be absent in any good case in which Austin sees a green ovoid tomato (see section 1). The common element cannot be *x sees that the tomato is red*, because the subject is not in that state in I_2 ; similarly, it cannot be *x sees that the tomato is spherical*. What about *x sees the tomato*? Although in every illusory case Austin sees the tomato, this state is not specific enough—Austin is in that state in a good case in which he sees a green ovoid tomato, which happens to be the very same tomato he sees in 'the' good case. But of course we have left out the obvious candidate: *the tomato looks red and spherical to x*. In every illusory case, the tomato looks red and spherical to Austin, as it does in the good case.

It might be objected that this is not to identify a *mental* state, because one can only be in the state if the tomato exists, but with the demise of the Cartesian view, this complaint has no apparent basis. The point can be reinforced by noting that if *the tomato looks red and spherical to x* is *not* a mental state, this is presumably because it is some sort of hybrid, consisting of being in an 'inner' tomato-independent mental state appropriately caused by the tomato. However, a credible hybrid analysis remains completely elusive.

Obviously the previous complaint is not one that the *disjunctivist* can accept—an example of a tomato-entailing mental state is provided by the good case itself. Instead, the disjunctivist will presumably point to the gap between a psychological locution's being non-disjunctive, and its picking out a state or event with a unified mental

nature. The disjunctivist has to exploit this gap elsewhere—for example, she cannot allow that ‘a visual experience of a red tomato’ (compare section 6.2) picks out an event of a distinctive mental kind. And if this move is acceptable for ‘a visual experience of a red tomato’, why not also for ‘looks red’, ‘looks spherical’, and the like?

This objection would have some merit if $V \vee I/H$ disjunctivism were the default view, or if there were a persuasive argument for it. But from the previous two sections, there appears to be no reason to believe this. The parallel between ‘a visual experience of a red tomato’ and ‘looks red’ is anyway not convincing: the former occurs mostly in philosophy papers, while the latter is a central component of ordinary psychological talk. Denying that *the tomato looks red to x* is not a mental state is scarcely more appealing than denying that *x believes that the tomato is red* is a mental state.

Moreover, expressions like ‘looks red’ are (arguably) needed to characterize the *good* case properly. If so, ‘looks red’ cannot be given a disjunctive analysis, since such an analysis would be partly in terms of the good case, and hence circular. In section 1, the good case was described as one in which Austin sees the tomato, and “sees that it is red (and spherical)”. In the context, that was sufficient for getting across the idea of the good case. But in fact, it is possible to see that the tomato is red—moreover, see that the tomato is red *on the basis of vision*—even if the tomato does not *look* red. Imagine that objects are red if and only if they are shiny. Tomatoes are red (and shiny); they also look shiny, but they do not look red. Imagine, further, that one knows this fact. One may know by looking that the tomato is red—thus one may see that the tomato is red (compare Hinton 1973: 31).⁵³ (Perhaps we can generate a similar example without relying on background knowledge: imagine a creature who is ‘hard wired’ to take visibly shiny objects to be red, which they are. Reliabilist sympathizers, at least, will allow that the creature can see that the tomato is red.)

Obviously the good case was not intended to be of this kind. To make that intention explicit we apparently need to stipulate that the tomato *looks* red in the good case—or else explain the good case using the philosophical jargon of ‘veridical perception’. The first strategy precludes giving a disjunctive account of *the tomato looks red to x*; the second just invites the question ‘And what is a ‘veridical perception of a red tomato’, precisely?’

9 AGAINST $VI \vee H$ METAPHYSICAL DISJUNCTIVISM

In hallucinatory cases, the subject does not see the tomato, and so it does not look red. If there is a specific mental state or event common to the good case and the hallucinatory (and illusory) cases, it cannot be the one previously identified. But what about *it is to x as if there is something which looks to him (x) to be red and spherical* (borrowing Snowdon’s phrase for anti-disjunctivist purposes), *x seems to see a red spherical thing*, or *it looks to x as if there is a red spherical thing before him*? The former locution is a piece of philosopher’s argot which is itself in need of explanation: there should be no presumption that it picks out a mental state. The problem with the latter two locutions

⁵³ For some relevant discussion, see Dretske (1969: 78–93); Jackson (1977: ch. 7).

is that their obvious readings are epistemic—as in ‘I seem to be out of gas/putting on weight/drunken’, ‘It looks as if I’m out of gas/putting on weight/drunken’, etc. A rough paraphrase of ‘I seem to see a red spherical thing’, for example, would be ‘I have evidence that supports the hypothesis that I see a red spherical thing’; hence ‘I seem to see a red spherical thing’ does not even purport to ascribe a distinctive kind of mental state.

The issue of the common element is unlikely to be settled by interrogating folk psychology. A small dose of high theory is needed, and here we need to return to the two main foils distinguished in section 5, *abstract* and *particular* intentionalism. Recall that the abstract intentionalist holds that the content of Austin’s experience in the good case is the proposition that $\exists x(x \text{ is red, spherical, and before Austin})$; the particular intentionalist holds that the content includes the proposition that Tom is red, spherical, and before Austin (where Tom is the seen tomato). As noted in section 5, if abstract intentionalism is correct, then the content of the experience in hallucinatory cases is exactly the same as the content in the good case, and hence (metaphysical) disjunctivism of any kind is false.

However, abstract intentionalism might well be doubted. First, as Martin has (in effect) argued, its motivations are quite suspect.⁵⁴ Second, empirical theories of vision sometimes take the contents of experience to be particular (in particular, the “visual indexes” of Pylyshyn 2003). Third, and more ambitiously, one might attempt to run a transcendental argument for particular intentionalism as a necessary condition for perceptually-based singular thought (compare Brewer 1999: ch. 2).

It would take us too far afield to examine the case for particular intentionalism in detail. Still, we may fairly conclude that arguing against disjunctivism by assuming abstract intentionalism is a risky business.

So suppose that particular intentionalism is true. To adapt a famous example from Strawson (1950: 333), the hallucinatory cases are then analogous to uttering ‘This is red and spherical’, while demonstrating nothing.

With this analogy in mind, one might see the common element in the *predication* of redness and sphericity—predication that fails for lack of a suitable subject. Cashing out the metaphor of ‘failed predication’ is by no means trivial.⁵⁵ Fortunately, though, there is an easier strategy.

The notion of the content of experience is usually explained in terms of the way the experience ‘presents/represents the world as being’, the experience’s ‘correctness conditions’, the conditions under which the experience would be ‘veridical’, and the like. (For some quotations, see Byrne 2001: 201; Travis 2004: 58–9.) These explanations are of course vague, but may be none the worse for it. The crucial point is that, although one may imagine various ways of injecting more precision, none will *rule out* abstract contents entirely. For instance, it might be suggested that the content

⁵⁴ Martin (1997: 92–4); see also Martin (2002b); Soteriou (2000).

⁵⁵ For some suggestions, see Hawthorne and Kovakovich (2006); Johnston (2004); Siegel (2005: section 5.1). As the Strawson analogy indicates, this issue is to a large extent a replay of the debate about the Evans/McDowell position on singular thought: see, for example, Segal (1989); for the connection with hallucination, see Smith (2002: ch. 8); Martin (2002b).

of experience should align closely with the intuitive conception of a visual illusion, that the contents should be ‘cognitively impenetrable’, that they should be confined to the output of ‘early vision’, or that they should be ‘non-conceptual’. None of these suggestions is inimical to the idea that the content of experience is (partly) abstract. That is why we said that, according to the particular intentionalist, the content of the experience in the good case *includes* the proposition that Tom is red, spherical, and before Austin. As far as we can see, no theoretically well-motivated precisification of the notion of perceptual content will have the consequence that abstract contents—for instance, the proposition that $\exists x(x \text{ is red, spherical, and before Austin})$, or something similar but fancier—are not represented in the good case.

Further, we can motivate the idea that experience has abstract content with certain examples of ‘semi-veridical illusions’. Consider the following Gricean case (modified from Martin 2002b: 10). Suppose one is looking at a scene with a green tomato on the left and a red tomato of exactly the same shape on the right. The green tomato is (very!) cleverly illuminated to look red, and the red tomato is cleverly illuminated to look green. One views the tomatoes through a prism, which has the effect of transposing their apparent locations. It appears to one that there is a green tomato on the left and a red tomato on the right, and indeed there is. There may well be a temptation to say that one’s experience is, with respect to colour and location, *inaccurate* or *illusory*—one perceives what is in fact the rightmost red tomato as being green and on the left. But on the other hand, there is also a temptation to say that one’s experience is, with respect to colour and location, *accurate* or *veridical*—there is a green tomato on the left, just as there appears to be. If the content of the experience is *entirely* particular, this last temptation must be resisted.

If veridical and illusory experiences have abstract content, then hallucinatory experiences can hardly be exempt. (Compare the analogous position for thought: if one *accepts* that, in the ‘good case’, when the subject says ‘This is red and spherical’, she believes (*inter alia*) that *there is* something red and spherical, one should *also* accept that the subject believes this abstract proposition in an hallucinatory case, where the demonstrative ‘this’ does not refer.⁵⁶) And if the experiences in the hallucinatory cases have abstract content then we have found a common element.⁵⁷

The upshot of the last six paragraphs is this: *if* experience has representational content at all, then VI v H disjunctivism is false. Now notice that the VI v H disjunctivist is under some pressure to accept the antecedent. In an illusory case, things look exactly the same to Austin as they do in the good case, where this ‘sameness in look’ is supposed (by the VI v H disjunctivist) to be a genuine respect of mental overlap. In the good case, things *are* as they look, and the experience is *accurate*; in an illusory case, things *aren’t* as they look, and the experience is *inaccurate*. There are certain propositions that jointly specify ‘the way things look’ in the good and illusory cases

⁵⁶ This (plus agreement with the antecedent) seems to be McDowell’s position (1986/1998: 236, fn. 17).

⁵⁷ In the terminology of Hawthorne and Kovakovich (2006), we are denying that “strong singularism” is tenable. In taking perceptual experience to have (at least) abstract content, we are *not* denying the importance or primacy of object-dependent content (compare, again, the analogous position for thought).

(that this is red and spherical, that there is something red and spherical before me, etc.), and which give the accuracy conditions of the experiences: these propositions are true in the good case and false in the illusory cases. But this is, in all but name, the almost orthodox view that experience has representational content. The VI v H disjunctivist plainly cannot reject out of hand the possibility that the experience in hallucinatory cases has false content, since she admits that this is the correct account of illusion. And once that option is explored a little further, its attractions become clear, or so we have argued. To adapt Russell on naïve realism: VI v H disjunctivism, if true, is false; therefore it is false.⁵⁸

10 SUMMARY

Our main points are these:

1. Epistemological disjunctivism (McDowell) should be sharply separated from metaphysical disjunctivism (Hinton et al.). (section 4)
2. Metaphysical disjunctivism leads naturally to epistemological disjunctivism, but not conversely. (section 5)
3. There's no compelling argument for (metaphysical) disjunctivism, or for taking it to be the default view of perceptual experience. (sections 6 and 7)
4. V v I/H disjunctivism is false. (section 8)
5. As is—more tentatively—VI v H disjunctivism. (section 9)

And, finally: the default position is not the Cartesian view, (metaphysical) disjunctivism, or naïve realism, but the *moderate view*. But we must admit that the appeal

⁵⁸ Alston would resist the step from accuracy conditions to the propositional (representational) content of experience, on the grounds that when a tomato looks red to one, this is a form of “non-conceptual cognition” of the tomato. According to him, this rules out views on which perceptual experience has propositional content (1999: 184). We lack the space to discuss this properly, but we think that in whatever defensible sense experience is ‘non-conceptual’, this is *compatible* with experiences having propositional content. (Johnston argues against “accounts that recognise *only* propositional acts of sensing” (2006: 279, our emphasis); this conclusion is also compatible with experiences having propositional content.)

A more radical view is defended by Travis (2004). According to him, it is a mistake to say that illusions involve *any* kind of inaccuracy at the level of experience—any inaccuracy is entirely a matter of the judgements the subject would tend to make. We cannot possibly do Travis's view justice here. Briefly, though, we find his model of illusion unconvincing. He thinks an illusion is a matter of taking feature A of the perceived scene to indicate that the scene has feature B, which in fact it does not. To use his example, the Müller-Lyer lines “have a certain look” (feature A), which one might take to “indicate that it is two lines of unequal length that one confronts” (feature B) (2004: 68). Whatever the “certain look” is, exactly, it is intended to be a feature of the lines that has nothing to do with the subject who is perceiving them; compare (Austin 1962: 43), on “petrol looks like water”. We think (but do not argue here) that this account is neither necessary nor sufficient for illusion. On Travis's view, the good case is a perception of the tomato's redness, an illusory case is a perception of some *other* feature of the tomato (a “certain look”, presumably), and an hallucinatory case is an entirely different kettle of fish. Assuming that the good case *also* involves perceiving the “certain look”, Travis is a VI v H disjunctivist; see also Travis (2005).

of the moderate view is largely due to the disjunctivists' insightful emphasis on the difference between the good and bad cases. So it is they—not forgetting, of course, Kierkegaard—who should get most of the credit.

REFERENCES

- Alston, W. (1999) 'Back to the Theory of Appearing', *Philosophical Perspectives*, 13: 181–203.
- Armstrong, D. M. (1976) 'Immediate Perception', in R. Cohen, P. Feyerabend, and W. Wartofsky (eds.), *Essays in Memory of Imre Lakatos*. Dordrecht: Reidel: 23–35.
- Austin, J. L. (1962) *Sense and Sensibilia*. Oxford: Oxford University Press.
- Bennett, J. (1988) *Events and their Names*. Indianapolis: Hackett.
- Brewer, B. (1999) *Perception and Reason*. Oxford: Oxford University Press.
- Byrne, A. (2001) 'Intentionalism Defended', *Philosophical Review*, 110: 199–240.
- (2004) 'How Hard are the Sceptical Paradoxes?', *Noûs* 38: 299–325.
- Campbell, J. (2002) *Reference and Consciousness*. Oxford: Oxford University Press.
- Child, W. (1994) *Causality, Interpretation, and the Mind*. Oxford: Oxford University Press.
- Crane, T. (2005) 'The Problem of Perception', in E. Zalta (ed.), *Stanford Encyclopedia of Philosophy* (Spring 2005 Edition). <<http://plato.stanford.edu/archives/spr2005/entries/perception-problem/>>.
- Dancy, J. (1995) 'Arguments from Illusion', *Philosophical Quarterly*, 45: 421–38.
- Dretske, F. (1969) *Seeing and Knowing*. Chicago: University of Chicago Press.
- (2003) 'How do You Know You are Not a Zombie?', in B. Gertler (ed.), *Privileged Access: Philosophical Accounts of Self-Knowledge*. Aldershot: Ashgate Publishing: 1–14.
- Evans, G. (1982) *The Varieties of Reference*. Oxford: Oxford University Press.
- Graff, D. (2001) 'Phenomenal Continua and the Sorites', *Mind*, 110: 905–35.
- Hawthorne, J., and Kovakovich, K. (2006) 'Disjunctivism', *Proceedings of the Aristotelian Society*, Supplementary Volume 80: 145–83.
- Hellie, B. (2005) 'Noise and Perceptual Indiscriminability', *Mind*, 114: 481–508.
- Hinton, J. M. (1967a) 'Visual Experiences', *Mind*, 76: 217–27.
- (1967b) 'Experiences', *Philosophical Quarterly*, 17: 1–13.
- (1967c) 'Illusions and Identity', *Analysis*, 27: 65–76.
- (1973) *Experiences*. Oxford: Clarendon Press.
- (1980) 'Phenomenological Specimenism', *Analysis*, 40: 37–41.
- Jackson, F. (1977) *Perception: A Representative Theory*. Cambridge: Cambridge University Press.
- Johnston, M. (2004) 'The Obscure Object of Hallucination', *Philosophical Studies*, 120: 113–183.
- (2006) 'Better than Mere Knowledge? The Function of Sensory Awareness', in T. Gendler and J. Hawthorne (eds.), *Perceptual Experience*. Oxford: Oxford University Press: 260–90.
- Langsam, H. (1997) 'The Theory of Appearing Defended', *Philosophical Studies* 87: 33–59.

- Locke, D. (1967) *Perception and our Knowledge of the External World*. New York: George Allen & Unwin.
- ____ (1975) 'Review of J. M. Hinton's *Experiences*', *Mind*, 84: 466–8.
- McDowell, J. (1982) 'Criteria, Defeasibility, and Knowledge', *Proceedings of the British Academy*, 68: 455–79. [Page references are to the reprint in McDowell (1998: 369–94).]
- ____ (1986) 'Singular Thought and the Extent of Inner Space', in J. McDowell and P. Pettit (eds.), *Subject, Thought and Context*. Oxford: Oxford University Press: 135–68. [Page references are to the reprint in McDowell (1998: 228–59).]
- ____ (1995) 'Knowledge and the Internal', *Philosophy and Phenomenological Research* 55: 877–93. [Page references are to the reprint in McDowell (1998: 395–413).]
- ____ (1998) *Meaning, Knowledge, and Reality*, Cambridge, MA: Harvard University Press.
- ____ (this volume) 'The Disjunctive Conception of Experience as Material for a Transcendental Argument'.
- Martin, M. G. F. (1997) 'The Reality of Appearances', in M. Sainsbury (ed.), *Thought and Ontology*. Milan: Franco Angeli: 81–106.
- ____ (2002a) 'The Transparency of Experience', *Mind and Language*, 17: 376–425.
- ____ (2002b) 'Particular Thoughts & Singular Thought', in A. O'Hear (ed.), *Logic, Thought and Language*. Cambridge University Press: 173–214.
- ____ (2004) 'The Limits of Self-Awareness', *Philosophical Studies*, 120: 37–89.
- ____ (2006) 'On Being Alienated', in T. Gendler and J. Hawthorne (eds.), *Perceptual Experience*. Oxford: Oxford University Press: 354–410.
- Millar, A. (1996) 'The Idea of Experience', *Proceedings of the Aristotelian Society*, 96: 75–90.
- Peacocke, C. (1983) *Sense and Content*. Oxford: Oxford University Press.
- Price, H. H. (1950) *Perception* (2nd edn). London: Methuen.
- Putnam, H. (1999) *The Threefold Cord: Mind, Body, and World*. New York: Columbia University Press.
- Pylyshyn, Z. (2003) *Seeing and Visualizing: It's Not What You Think*. Cambridge, MA: MIT Press.
- Robinson, H. (1985) 'The General Form of the Argument for Berkeleian Idealism', in J. Foster and H. Robinson (eds.), *Essays on Berkeley*. Oxford: Oxford University Press: 163–86.
- ____ (1994) *Perception*. London: Routledge.
- Russell, B. (1912) *The Problems of Philosophy*. London: Williams & Norgate.
- ____ (1950) *An Enquiry into Meaning and Truth*. London: Unwin.
- Segal, G. (1989) 'The Return of the Individual', *Mind*, 98: 39–57.
- Siegel, S. (2004) 'Indiscriminability and the Phenomenal', *Philosophical Studies*, 120: 91–112.
- ____ (2005) 'The Contents of Perception', in E. Zalta (ed.), *The Stanford Encyclopedia of Philosophy* (Summer 2005 Edition). <<http://plato.stanford.edu/archives/sum2005/entries/perception-contents/>>.
- Smith, A. D. (2002) *The Problem of Perception*. Cambridge, MA: Harvard University Press.

- Snowdon, P. (1980–1) 'Perception, Vision, and Causation', *Proceedings of the Aristotelian Society*, 81: 175–92.
- (1990) 'The Objects of Perceptual Experience', *Proceedings of the Aristotelian Society* Supplementary Volume 64: 121–50.
- (2005) 'The Formulation of Disjunctivism: A Response to Fish', *Proceedings of the Aristotelian Society*, 105: 129–41.
- Soteriou, M. (2000) 'The Particularity of Visual Perception', *European Journal of Philosophy*, 8: 173–89.
- (2005) 'The Subjective View of Experience and its Objective Commitments', *Proceedings of the Aristotelian Society*, 105: 177–90.
- Strawson, P. F. (1950) 'On Referring', *Mind*, 59: 320–44.
- Sturgeon, S. (2000) *Matters of Mind*. London: Routledge.
- (2006) 'Reflective Disjunctivism', *Proceedings of the Aristotelian Society* Supplementary Volume 80: 185–216.
- Thau, M. (2004) 'What is Disjunctivism?' *Philosophical Studies*, 120: 193–253.
- Travis, C. (2004) 'The Silence of the Senses', *Mind*, 113: 57–94.
- (2005) 'Frege, Father of Disjunctivism', *Philosophical Topics*, 33: 307–34.
- Tye, M. (2000) *Consciousness, Color, and Content*. Cambridge, MA: MIT Press.
- Vendler, Z. (1957) 'Verbs and Times', *Philosophical Review*, 66: 143–60.
- Williamson, T. (1990) *Identity and Discrimination*. Oxford: Blackwell.
- (2000). *Knowledge and Its Limits*. Oxford: Oxford University Press.
- Yablo, S. (1997) 'Wide Causation', *Philosophical Perspectives*, 11: 251–81.

3

Against Disjunctivism

E. J. Lowe

According to the causal theory of perception (CTP), a subject *S* perceives an object *O* if and only if *S* has a perceptual experience, *E*, of an appropriate kind, which is suitably caused by *O*.¹ This statement of the theory leaves certain important details to be filled out—most obviously, what is meant by ‘appropriate’ and ‘suitable’ in this context. Very roughly and only to a first approximation, an appropriate experience has to have a sensory modality matching that of the mode of perception concerned—thus, it must be a *visual* experience if the mode of perception concerned is *vision*—and it must have a content that matches at least some perceptible feature of the perceived object, *O*.² For example, *O* might be *red* and *E*’s content meets the

I am grateful to Adrian Haddock and Fiona Macpherson for their careful and penetrating comments on an earlier draft of this essay, and also to two anonymous referees for their helpful remarks on the penultimate version.

¹ I develop a version of this theory, very similar to the version presented here, in my (1996: ch. 4). In that book, I do not engage directly with disjunctivist opponents of such a theory, for reasons I give there (1996: 92 fn. 2). The present essay is designed to repair that omission and to extend the discussion of disjunctivism in my (2000: 145–9).

² Rather confusingly, in my view, some authors who discuss disjunctivism in the philosophy of perception use the expression ‘perceptual experience’ in a deliberately *contrastive* way, to denote experiences that are *not* hallucinatory. Thus, for example, Paul Snowdon writes: “As well as . . . perceptual experiences there are experiences which are certainly not perceptions of the environment, but which when undergone, are (or can be) easily mistaken for such perceptual experiences . . . for example, certain sorts of after-images and hallucinations” (2005: 135). I myself prefer to *introduce* the expression ‘perceptual experience’ into philosophical discussion as one that may be used in a generic and *neutral* way, to denote any experience that has a distinctive sensory modality—whether that modality be, say, visual, auditory, or olfactory—and without the presumption that the subject undergoing it must thereby actually be *seeing, hearing, or smelling* anything in his or her environment. Using the expression in this neutral or non-contrastive way does not *ipso facto* beg the question against disjunctivism, any more than using it in the contrastive way begs the question against opponents of disjunctivism: for it can still be debated whether the ‘perceptual experiences’ involved in actual perceivings of the environment are *of the same kind* as those involved in cases of hallucination which might be mistaken by their subjects for being such perceivings. (CTP holds that they *are* of the same kind and thus, when the expression ‘perceptual experience’ is used by advocates of CTP explicitly as a *theoretical* term of their own, it is used with the intention of denoting items of that putative kind: and this, accordingly, is how I shall later

matching requirement by including an element that registers *redness*. This way of expressing the matching requirement on content—by saying that, for some perceptible feature *F* of *O*, *E*'s content must 'include an element that registers *F*ness'—is thus far only a convenient form of words awaiting fuller explication, which I shall supply later. As for the requirement that the causation of *E* by *O* should be of a 'suitable' kind, this is needed to exclude so-called 'deviant' or 'wayward' causal chains. We have here at best only a skeleton of a theory of perception, but enough, perhaps, to be going on with for the time being. A more pressing issue than that of how to fill out the theory is to say what *kind* of 'theory' it is supposed to be. A *philosophical* theory, we might say. But what is meant by that? Well, negatively speaking, what is meant is that the theory is not intended to compete with anything that an empirical scientist might have to say, on the basis of observation or experimentation, about perceptual processes in animals of any kind, including human beings. Positively speaking, it might be suggested that what is on offer is an 'analysis' of the *concept* of perceiving, or an account of the meaning of the word 'perceive'. But that, I think, would be to put an excessively psychological or linguistic gloss on what it is that a philosophical theory of perception should be concerned with. I would prefer to say that what is being offered by such a theory is an account of *what it is* to perceive: in other words, an account of the *essence* of perception—remembering that 'essence' is just the word customarily used to translate a phrase of Aristotle's whose more literal translation is 'what it is (to be)', or 'what it would be (to be)'.³ Very arguably (although also, I concede, quite controversially), we have to grasp the essence of anything, *X*—that is to say, we have to understand *what X is*—before we can undertake any sort of empirical investigation of *X*. Thus, if that is correct, an empirical psychologist or physiologist must at least understand *what it is* to perceive if he or she is to be able to investigate the psychological and physiological processes involved in human perception. And supplying such an understanding is what a *philosophical* theory of perception is, or at least *should* be, designed to do.

Why should anyone suppose that a philosophical theory of perception should be a *causal* theory—that something along the lines of CTP is correct? Does CTP face or present insurmountable difficulties? Is there any more cogent alternative to it? Advocates of *disjunctive* theories of perception will, of course, reply to the latter two questions with a resounding 'Yes'.⁴ Let us review some of these issues, beginning with the question of why CTP should seem at all compelling to any philosopher. The most

use it when speaking on behalf of CTP.) The reason why I prefer my own practice of introducing the expression 'perceptual experience' in a non-contrastive way is simply that it seems to me that this reflects actual common usage of more specific expressions such as 'visual experience'. Ordinary, philosophically untutored users of everyday English commonly speak of 'visual experiences' in a way that is clearly intended to *leave open* the question as to whether the subjects of those experiences are thereby actually seeing anything in their environment.

³ See Aristotle (1928: Z, 4).

⁴ The *locus classicus* of disjunctivism is Hinton (1973). See also Snowdon (1980–1) and McDowell (1982). I should acknowledge that the claim that CTP is inconsistent with disjunctivism has not gone entirely unchallenged: see, in particular, Child (1994: ch. 5). However, the version of CTP that I am defending here is most certainly incompatible with disjunctivism, as I interpret the latter doctrine.

obvious reason for this seems to me to be the following, which we could call the *cut-off argument*. For the purposes of this argument, I assume that we are talking exclusively about the *conscious* perception of so-called ‘external’ objects—that is, physical objects occupying places outside the periphery of the perceiver’s body. Now, it would seem that whenever a subject, *S*, perceives an object, *O*, *S* is in a conscious mental state, *M*, whose intrinsic character does not depend ontologically on the state of the physical environment external to the periphery of *S*’s body. Ontological dependence is here to be understood as a *metaphysical* relation, not to be confused with any kind of *causal* relation—for, of course, if CTP is correct, *M*’s intrinsic character may indeed be expected to depend *causally* on features of *O* and thus on the state of the physical environment external to the periphery of *S*’s body.⁵ For present purposes, what it amounts to to say that *M*’s intrinsic character does not depend ontologically on the state of the physical environment external to the periphery of *S*’s body is just this: *M* *could* exist, with the same intrinsic character, even if the state of the physical environment external to the periphery of *S*’s body were quite different—and, in particular, even if neither *O* itself nor any other object perceptibly similar to *O* were part of that physical environment (the word ‘could’ here being expressive of so-called *metaphysical* possibility). Why does this claim seem plausible? Because, I suggest, we are strongly inclined to believe that if all causal interaction between *O* and *S* were *cut off* at any point between *O* and the periphery of *S*’s body, but the state of *S*’s body remained the same—if, for example, in the visual case, the pattern of neural stimulation in *S*’s retina remained the same, being produced by a different cause—then *S* would be in a conscious mental state *M** with exactly the same intrinsic character as *M*.⁶ (I shall say more a little later about what exactly I mean by ‘intrinsic character’ in this sense.) Now, of course, one way to cut off all causal interaction between *O* and *S* is simply to remove *O* altogether, in which case *S* can no longer be perceiving *O*. But this case does not seem to be relevantly different from one in which we cut off all causal interaction between *O* and *S* *without* removing *O*, whence we can conclude that in such a case, too, *S* is no longer perceiving *O*. The implication is that, in order for *S* to perceive *O*, causal interaction between *O* and *S* must *not* be cut off—and this is why the cut-off argument seems to support CTP.

⁵ For fuller discussion of the notion of ontological dependence in play here, see my (1998: ch. 6).

⁶ Some philosophers, it would seem, do *not* consider that we are strongly inclined to believe this, at least subsequent to reflecting on the nature of our perceptual experiences. Thus, for example, M. G. F. Martin characterizes ‘naïve realism’—which, presumably, is supposed to embody our intuitive, pre-theoretical beliefs about the nature of perception—as implying that “the actual objects of perception . . . such as trees, tables and rainbows . . . and the properties which they can manifest to one when perceived, partly constitute one’s conscious experience, and hence determine the phenomenal character of one’s experience . . . [so that] one could not be having the very experience one has, were the objects perceived not to exist, or were they to lack the features they are perceived to have” (Martin 1997: 83). I can only say that I find this unconvincing as an account of philosophically untutored opinion concerning the nature of perceptual experience, in the light, for example, of popular reaction to films such as *The Matrix* (in which, it will be recalled, almost the entire human population of Earth at a future time lives in a ‘virtual reality’ created by computer-programmed inputs to their brains). Audiences seem to find such films entirely believable. ‘Naïve realism’ as characterized by Martin seems to me to be a sophisticated philosophical position, rather than genuinely ‘naïve’. See also footnote 22 below.

Given that *S* would not be *perceiving O* if all causal interaction between *O* and *S* were cut off but *S* was in a conscious state, *M**, with exactly the same intrinsic character as *M*, how *should* we characterize *S*'s condition in these circumstances? Well, we could perhaps characterize it by saying that *S* was undergoing an *illusion* or a *hallucination* of perceiving *O*. But those terms are loaded in certain ways which make them not entirely appropriate in all cases. The most neutral way in which we could describe *S*'s condition is, perhaps, to say that in these circumstances *S* merely *seems* to perceive *O*. The qualification 'merely' is important, for we want to acknowledge that when *S* is perceiving *O*, then *S* *seems* to perceive *O* too—but doesn't *merely* seem to perceive *O*. However, in the light of points that I am about to make concerning the disjunctive theory of perception, I should emphasize here that the advocate of CTP only calls upon the phrase '*S* merely *seems* to perceive *O*' to describe a condition of *S*'s which can, in his view, be more perspicuously or more informatively described in terms of *S*'s being in a conscious mental state, *M**, which is such that *if it (or another state M with the same intrinsic character as M*) had been suitably caused by O*, then *S* would thereby have been perceiving *O*. In other words, for the advocate of CTP, merely seeming to perceive something is a matter of being in a mental state which, if it had been suitably caused by that thing, would have constituted the mental ingredient in an episode of perceiving that thing. Calling such a condition one in which *S* 'merely seems to perceive *O*' is appropriate because, if *S* is in such a condition *and is aware of the fact*, *S* may well say of himself 'I seem to perceive *O* but am not in fact doing so' or 'It seems to me as though I am perceiving *O*, but in fact I am not' or, indeed, 'I *merely* seem to perceive *O*'. But the advocate of CTP certainly doesn't have to maintain that whenever *S* is in a condition aptly describable as being one in which *S* 'merely seems to perceive *O*', *S* will *believe of himself* that he merely seems to perceive *O*, or even that he seems to perceive *O*. For, apart from anything else, *S* may not even be capable of entertaining such a belief and yet may still be in a condition aptly describable as being one in which *S* merely seems to perceive *O*. *S* may, for instance, be a young child at a stage of her intellectual development at which she is not yet capable of framing such judgements concerning her own mental condition.

Why should any philosopher see any reason to quarrel with any of the above? For present purposes, I shall focus on the kinds of worry that advocates of disjunctivism typically raise in this connection. For one thing, there are worries about some of the supposed consequences of CTP—in particular, its supposed *sceptical* implications.⁷ For another, but relatedly, there are worries about the supposed failure of CTP to capture allegedly central features of the *phenomenology* of perception—in particular,

⁷ I don't want to suggest that *all* advocates of disjunctivism are motivated by worries about scepticism, nor that they all think that disjunctivism could help to alleviate such worries. But considerations of this sort plainly do play a role in the appeal that disjunctivism has for some philosophers, such as John McDowell and Hilary Putnam. (Of course, I should also acknowledge that 'disjunctivism' isn't a single theory but a general approach in the philosophy of perception, which has significantly different variants.) Thus, for example, Putnam—acknowledging in this context a debt to McDowell (1994)—writes: "[T]he traditional claim that we must conceive of our sensory experiences as *intermediaries* between us and the world . . . makes it impossible to see how persons can be in genuine cognitive contact with the world at all" (Putnam 1999: 11).

its supposedly ‘diaphanous’ or ‘transparent’ character, which in turn is supposed to explain and perhaps even support the intuitions that underlie the appeal of naïve realism. For yet another, there are worries that the key notions utilized by CTP—especially that of a ‘perceptual experience’, as this expression is typically deployed by advocates of the theory—cannot be made sufficiently clear for it to qualify as a viable philosophical theory of perception. Finally, there are worries to the effect that, even if the theory can be coherently formulated, it incurs unwarranted ontological commitments in the shape of ‘perceptual experiences’ so conceived, because another theory which lacks those commitments—*disjunctivism*—is at least as satisfactory, if not more so.

But what *is* ‘disjunctivism’ in the philosophy of perception? For present purposes, I am taking it to be the doctrine that there is no *common factor*—in the form of a ‘perceptual experience’, *E*, or a ‘conscious mental state’, *M*—in the cases of *S*’s *actually* perceiving *O* and *S*’s *merely seeming* to perceive *O*.⁸ That there is such a common factor is, of course, a central tenet of CTP, which holds that in a case of veridical perception, *S* has a perceptual experience suitably caused by *O*, whereas in a case of merely seeming to perceive, *S* has a conscious mental state with the same intrinsic character as such a perceptual experience, but one that is caused in another way. The disjunctivist acknowledges, of course, that any case of seeming to perceive *O* is either a case of *actually* perceiving *O* or else a case of *merely* seeming to perceive *O*—but denies that the two kinds of case are related in the way that CTP proposes, namely, in virtue of having a common factor in the form of a conscious mental state of the same kind that is differently caused in the two cases. Where the advocate of CTP sees the disjunction ‘Either *S* actually perceives *O* or else *S* merely seems to perceive *O*’ as having the underlying form ‘Either $X + Y$ or else $X + Z$ ’, or at least as implying something of this form—with the ‘common factor’, *X*, being that *S* has a conscious mental state, *M*, with a certain intrinsic character—the disjunctivist sees the disjunction as being reducible or analysable no further and hence as being basic. Thus, the disjunctivist sees the disjunction in question as being perfectly intelligible just as it is, without needing any further analysis or explanation—or, at least, not one couched in terms of the presence of a supposed ‘common factor’. According to the disjunctivist, a

⁸ Hinton (1973: 60 ff.) speaks in this connection of a ‘common element’—as does Snowdon (1980–1). McDowell (1982) speaks of a ‘highest common factor’, and Putnam (1999) follows McDowell in this usage. However, caution must be exercised in characterizing disjunctivism in this way, so as not to imply that there need be *nothing at all* ‘in common’ to two such cases. Thus, Martin describes disjunctive theories of perception as claiming only that “it is simply mistaken to suppose that there need be anything more in common across veridical perceptions and delusive experience, other than the fact that all of these states of mind may be indistinguishable for the subject who has them, in some or all respects” (1997: 81). Tim Crane makes the point well: “[D]isjunctivism does not deny that there is some true description under which the perception, say, of a rabbit and a subjectively indistinguishable hallucination of a rabbit can fall . . . What the disjunctivist . . . rejects is what J. M. Hinton calls ‘the doctrine of the “experience” as the common element in a given perception’ and an indistinguishable hallucination” (Crane 2006: 139, quoting from Hinton 1973: 71). Crane points out that Putnam, for one, has incautiously gone further than this to claim that “there is nothing *literally* ‘in common’” in cases of perception and hallucination, and in that very strong sense denies that there is a ‘Highest Common Factor’: see Putnam (1999: 152). I certainly don’t want to construe disjunctivism as being committed to such a denial.

case of *merely* seeming to perceive can be adequately characterized simply by contrast with a *veridical* case of perceiving, which is taken to be the conceptually prior notion, but without calling upon the notion of anything like a ‘perceptual experience’, conceived as a mental state of some single kind, such that states of that very same kind could be enjoyed *both* in veridical cases of perceiving *and* in cases of merely seeming to perceive.

So how, exactly, *does* the disjunctivist want to characterize a case of merely seeming to perceive *O*, given that he or she *doesn’t* want to characterize it in terms of having a certain kind of experience that, *were it to have been suitably caused by O*, would have meant that *O* was perceived? This, unfortunately for the disjunctivist, is not altogether easy to say. It is tempting to say that *S*’s merely seeming to perceive *O* is *S*’s being in a state that, while it is *not* one of perceiving *O*, is somehow *indistinguishable* by *S* from one of perceiving *O*. But indistinguishable in what sense and by what means? By way of illustration, we might be asked to envisage the case of a ‘perfect’ hallucination, which is not detectably different, from the point of view of its subject, from a case of ordinary veridical perception. But this still leaves too much unexplained. Someone undergoing such a hallucination might still *know*—and in that sense have succeeded in ‘detecting’—that he or she was not veridically perceiving anything: for instance, the hallucination might be a self-induced one undergone for experimental purposes by a neuroscientist.⁹ The advocate of CTP can of course say, in such a case, that what the subject cannot detect is any difference between the intrinsic character of the conscious experience that he or she is undergoing and the intrinsic character of a conscious experience that he or she might be undergoing in a case of veridical perception. But the disjunctivist must surely avoid saying anything like that.¹⁰ At this point the disjunctivist might try to shrug off the apparent difficulty by saying something like this: ‘Now, we are all familiar with what it is to perceive a certain object, *O*—well, *merely seeming* to perceive *O* is *just like that*, except that it’s

⁹ M. G. F. Martin proposes that a case of *S*’s having a ‘perfect’ hallucination of perceiving *O* is one that, while it is not a case of *S*’s perceiving *O*, is *indiscriminable by S through reflection or introspection alone* from one of *S*’s actually perceiving *O*: see Martin (2004: 47 ff.). But the notions of reflection and introspection invoked here are not really explained by Martin. In any case, the example of a self-induced hallucination that I have just cited seems to present a difficulty for Martin’s account. For it is quite possible that the neuroscientist in that example would find that his self-induced hallucination was invested with an introspectively detectable aura of unreality in virtue of his knowledge that it was illusory—and yet, *ex hypothesi*, this is supposed to be a ‘perfect’ hallucination. In a recent paper, Katalin Farkas has raised other difficulties for any approach like Martin’s and argues in favour of an account of experiential sameness akin to my own, that is, one that is couched in terms of what she calls sameness of *phenomenal character*, where the latter is understood as ‘what it is like to have [the] experience’: see her (2006: 40).

¹⁰ If the disjunctivist *does* say this, then s/he must be prepared to offer an account of what *constitutes* the intrinsic character of a conscious experience in the hallucinatory case in a way which doesn’t simply *appeal* to its (supposed) subjective indistinguishability from a veridical perceptual experience—for it is this very notion of indistinguishability that we are now seeking to elucidate. But if such an account can be offered by the disjunctivist, then it is hard to see how s/he can justifiably resist allowing that this same account could be taken to apply also to the intrinsic character of *veridical* perceptual experiences. However, to concede this would be to lend support to the hypothesis, which is anathema to the disjunctivist, that hallucinations and veridical perceptual experiences are mental states of fundamentally the same kind.

not the case that *O* is perceived'. But that is really no help at all. For we are entitled to ask: 'Just like that *in what respect*?' Again, the advocate of CTP has a simple response to this question, namely, 'Just like that in respect of involving a conscious mental state with the same intrinsic character'. What the disjunctivist can say without compromising his position is much harder to determine. However, I don't want to dwell on this point, but am prepared to allow, at least for the sake of argument, that the disjunctivist can come up with an appropriate and non-compromising formula—or, failing that, that it is allowable for the disjunctivist to take the notion of 'merely seeming to perceive' as standing in need of no further explication of any kind, perhaps because its proper application is learnable simply by acquaintance with cases that fall under it. We might suppose, for instance, that as children we learn to say things such as 'I just seemed to see *O*' as a result of being corrected by others when we mistakenly say things of the form 'I saw *O*'—the others being able to correct us in these cases because they know, perhaps, that *O* was not present and in view at the time.

But we must press on. Let us first see whether disjunctivists can make good the charge that advocates of CTP cannot make the notion of a 'perceptual experience'—construed in the sort of way that their theory demands—sufficiently clear.¹¹ What I have said so far on behalf of CTP concerning perceptual experiences is that these are conscious mental states with a certain intrinsic character, that they have certain 'contents', and that they have a distinctive sensory modality—so that any perceptual experience is distinctively either a *visual* experience, or a *tactile* experience, or an *auditory* experience, or what-not. It is easy enough—the advocate of CTP may say—to *pick out* such states for illustrative purposes: we can say that they are states such as those we enjoy when we consciously *see*, *feel*, *hear*, or otherwise *perceive* some object in our environment—for instance, the state I enjoy when I consciously see a vase of red roses standing on a table in front of me. This is manifestly different from the state I enjoy when I *feel* a vase of red roses standing on a table in front of me. As for the idea that such states have 'contents', this is not peculiar to the notion of a perceptual experience as such but common to those of many if not all mental states—common to all if Brentano was right in saying that intentionality is the mark of the mental. That being so, the advocate of CTP may urge that there is no *special* burden upon him to explain the notion of content, although I shall indeed say more about it shortly. All that remains, then, is the idea that perceptual experiences are conscious mental states with a certain *intrinsic character*. But this too is a relatively familiar notion, often cashed out in terms of *what it is like* to enjoy the mental state in question.¹² *Partly* this is a matter of the distinctive sensory modality of the state, since what it is like to enjoy a *visual* experience, say, is never the same as what it is like to enjoy a

¹¹ Such a charge forms a major part of Hinton (1973: 63–100). I shall not attempt to engage directly with Hinton's arguments here, since they are all directed at the alleged difficulty of formulating what he calls "experience-reports" (1973: 63), thus giving his strategy a linguistic cast which is at odds with my own approach to these matters (however well-suited it was to the era in which Hinton was writing).

¹² Compare again Farkas (2006: 40). I hardly need document here the history of the philosophical use of the expression 'what it is like' in this connection, so widespread is it in the literature on perception and consciousness of the last thirty years.

tactile experience.¹³ But it is much more than just that, because what it is like to enjoy one visual experience may still be very different from what it is like to enjoy another visual experience. When, in the context of the cut-off argument, I said on behalf of CTP that, very plausibly, the intrinsic character of the kind of conscious mental state involved when a subject *S* perceives some external object does not depend ontologically on the state of the physical environment external to the periphery of *S*'s body, I meant, then, that *what it is like* to enjoy such a conscious mental state, *M*, is ontologically independent of the state of the physical environment external to the periphery of *S*'s body—where this in turn means that it is metaphysically possible for *S* to have, in a quite different state of the physical environment external to the periphery of *S*'s body, an experience, *M*^{*}, which is exactly like *M* in respect of what it is like to enjoy it. (It may occur to someone to ask here whether the disjunctivist might not appeal to this notion of 'what it is like' to enjoy a mental state in order to cash out the sense in which *S*'s merely seeming to perceive *O* is indistinguishable by *S* from *S*'s actually perceiving *O*. Surely not: for it is a notion bound up with the idea that certain mental states involve sensory *qualia*—and this is an idea repugnant to the disjunctivist. Indeed, the very thought that mental states have an 'intrinsic character' in my sense ought to be repugnant to the disjunctivist.¹⁴)

It might be objected that it is illegitimate for the advocate of CTP to try to illustrate what he means by a 'perceptual experience' by saying, as above, that they are 'states such as those we enjoy when we consciously *see, feel, hear*, or otherwise *perceive* some object in our environment', for two reasons. First, because this is *circular*—appealing to examples of *perceiving* in aid of the explication of a notion (that of a 'perceptual experience') which is designed to be used in formulating a theory of *perception*. And second, because it simply begs the question by assuming that when we perceive some object in our environment, we *are* in some 'state' of the kind that the advocate of CTP has in mind (the sort of state that he calls a 'perceptual experience'). Neither charge is fair, I think. The circularity charge is unfair because, apart from anything else, CTP—at least as I am understanding it—is not intended to provide an analysis of the *meaning* of the word 'perceive', capable of being understood by one who did not already grasp that meaning. Rather, CTP is intended to tell us *what perceiving is*—to reveal its essence. For that purpose, the advocate of CTP is entitled to assume that he or she is addressing an audience who are already competent users of words like 'perceive'. The second charge is unfair also, because the examples used for illustrative purposes are not intended to *establish the existence* of 'perceptual experiences', but merely to focus our attention on the sort of thing that CTP means by a 'perceptual experience' *assuming* that it is correct in holding there to be such things. As

¹³ I suppose that this might be disputed in cases of severely degraded or rudimentary experiences of the two types, such as those undergone by subjects whose eyesight is seriously damaged and whose sense of touch is similarly impoverished. Although I don't myself find this suggestion intuitively plausible, I suppose that the question is one that can ultimately be decided, if at all, only on an empirical basis.

¹⁴ See again footnote 10 above. I concede, of course, that some disjunctivists and their sympathizers may be happy to talk *in their own terms* of the 'phenomenal character' of perceptual experiences: see, for instance, the passage quoted from Martin in footnote 6 above.

to what perceptual experiences are supposed to *be*, that has already been explained: they are supposed to be conscious mental states with three distinctive aspects—an intrinsic character, a content, and a sensory modality—where, moreover, the first of these aspects is ontologically independent of the physical environment external to the periphery of the subject's body. As to the *existence* of such states, that is posited as part of what is argued to be the best available theory of *what perceiving is*, in the light of considerations such as those provided by the cut-off argument. So far, then, I find nothing objectionable in the manner of proceeding adopted by the advocate of CTP in order to clarify the supposedly dubious notion of a 'perceptual experience'.¹⁵

However, I promised to say more about the notion of perceptual *content*. At the outset, I mentioned what I called a *matching* requirement on perceptual content. When *S* perceives *O* and in so doing has a perceptual experience *E* that is 'suitably' caused by *O*, I said, the following must be the case: for some perceptible feature *F* of *O*, *E*'s content must include an element that registers *F*ness. For example, *O* might be *red* and *E*'s content includes an element that registers *redness*.¹⁶ I used this somewhat arcane terminology because I did not want to prejudge the issue of whether some or all perceptual content is purely 'external'. An externalist might want to say, for instance, that when *S* perceives a red object, *O*, and perceives *that O is red*, the 'element that registers redness' in the content of his perceptual experience is quite simply *O itself*—or if not *O* itself, then *O's redness* or (if this is preferred) *the fact that O is red*. On this sort of view, elements of perceptual content are just external objects and/or their properties, or facts containing external objects and their properties as constituents (I have tried to accommodate several different versions of the view). By contrast, on an internalist view of perceptual content, an 'element that registers redness' in the content of a perceptual experience would be some item that *represents* redness, rather than redness itself or some red object or some state of affairs containing redness as a constituent.¹⁷ It really doesn't matter, for the purposes of CTP, which view of content is adopted. No doubt it is true that on some views of content, the content of a perceptual experience is an extrinsic feature of the experience, while on others it is an intrinsic feature. But this, too, doesn't matter for the purposes of CTP.

¹⁵ I should stress that everything that I say about CTP in this paragraph is said in respect of a version of the theory that I am prepared to endorse myself, not just any theory that has been advertised by its proponents as 'the causal theory of perception'.

¹⁶ Why require only that for at least *one* perceptible feature, *F*, of the perceived object, *O*, *E*'s content includes an element that registers *F*ness? Because, for example, we want to be able to say that *S* sees a uniformly white wall that fills his field of view, even though the only perceptible feature of the wall that is registered in the content of *S*'s visual experience is the wall's whiteness. Similarly, when a house with white walls is seen from afar—so that it appears merely as a white dot on the horizon—its whiteness may be its only perceptible feature that is registered visually. In general, though, *many* perceptible features of an object may be expected to be registered in the content of the perceptual experience of a subject who perceives that object.

¹⁷ It should now be clear why I chose to use the neutral term 'register' in my account of content, for one accepted meaning of 'register' is *show*. Thus, for example, we say that the pointer on the scale of a weighing machine *registers* or *shows* the weight of the item being weighed—and 'show' in this case means 'represents'. Equally, however, we may say that a coloured object *shows* its true colour in daylight—and 'show' in this case means 'exhibits' or 'displays'. In the latter case, but not in the former, the thing that *shows* the property in question is the very thing that *has* the property.

All that matters for those purposes is that perceptual experiences have what I called an intrinsic *character*, which is a matter of *what it is like* to enjoy them. The *content* of a perceptual experience, whether or not it is taken to be external and whether or not it is taken to be intrinsic, is no part of its *character* as I understand the latter. Thus, when I say that the experiences involved in *S*'s *actually* seeing *O* and *S*'s *merely seeming* to perceive *O* may have exactly the same intrinsic character, which is ontologically independent of the physical environment external to the periphery of *S*'s body, I don't mean either to imply that these experiences must have the same content or to imply that they need not. Which view is taken with regard to the latter issue will depend on what view of perceptual content is favoured by any given advocate of CTP. An externalist will have to say that the experiences have *different* contents, at least in the case in which *S* merely seems to perceive *O* and *O* itself is not even present: for on this view, *O* itself is an element included in the content of the veridical experience, but *O* itself obviously cannot be included in the content of the non-veridical experience. An internalist, by contrast, may maintain that the experiences have the *same* content. As I, at least, am understanding CTP, an advocate of the theory has a free choice as to which view of perceptual content to adopt, for the following very simple reason: CTP maintains that the causation in perception operates only on the *intrinsic character* of perceptual experiences, not on their content. That is to say, according to CTP, what the causal processes issuing from the object of perception, *O*, are responsible for is the intrinsic character of the perceptual experience, *E*, in virtue of having which the subject, *S*, perceives *O*. It is true that there is also a matching requirement on *E*'s content, but *that E has the requisite content* is not something that need be *caused by O* in order for *S* to be perceiving *O* in virtue of having *E*, according to CTP. What *O* causes is a pattern of neural stimulation at the periphery of *S*'s body, which in turn gives rise to an experience *E* with a certain intrinsic character: and that, as far as CTP is concerned, is the end of the *causal story* to be told about the interaction between *O* and *S* in a case of perception. Why *E* has the *content* it does is another matter whose explanation may be expected to turn in considerable measure upon *S*'s individual psychological history. *O* could give rise to experiences with exactly the same intrinsic character in two different subjects, but those experiences have very different contents, as far as CTP is concerned.

Having now established, I hope, that CTP can at least be coherently formulated, let us look at some of the disadvantages that are often laid at its door by disjunctivists and their sympathizers—in particular, its supposed tendency to promote scepticism and its supposed lack of faithfulness to the phenomenology of perception.¹⁸ Both

¹⁸ Any attempt to name specific philosophers as being especially responsible for such accusations would be invidious and it would, in any case, detract from the generality of my responses to focus on particular versions of them. That accusations of these kinds are commonplace I take to be self-evident to those familiar with the literature, both that designed for consumption by professional philosophers and that designed for a student readership. But if examples are wanted, see again those cited in footnote 7 above. And let me acknowledge once more that not *all* disjunctivists worry that CTP promotes scepticism. Some just worry that it mischaracterizes our perceptual knowledge of external things as being always *indirect*. However, as I am about to point out in the main text, CTP is *not at all* committed to regarding our perception of external objects as being always 'indirect', in

of these charges are based on the idea that CTP represents a subject's perception of external objects as being somehow *indirect*. Note, first, that the claim cannot be that CTP implies that a subject perceives an external object *O* only in virtue of perceiving some other, 'internal' object, such as a 'sense datum': for although it is possible to combine CTP with a sense-datum theory of the immediate objects of perception, it is not necessary to do so—and, indeed, I am for present purposes assuming a version of CTP which maintains that, where the perception of external objects is concerned, such objects themselves may be the immediate objects of perception. So in what sense can it be maintained that this version of CTP represents a subject's perception of external objects as being somehow indirect? Only in the sense that *S*'s perception of *O* is 'mediated' by a perceptual experience, *E*. That is, only in the sense that there is, according to CTP, an entity, *E*, that must exist *in addition to S and O* when *S* perceives *O*, so that '*S* perceives *O*' does not express a simple dyadic relation between *S* and *O*, but a *triadic* relation between *S*, *O*, and some third entity, *E*. This implication can be seen from the statement of CTP that I offered at the outset: a subject *S* perceives an object *O* if and only if *S* has a perceptual experience, *E*, of an appropriate kind, which is suitably caused by *O*. Moreover, of course, CTP maintains that *S* *could* have *E* even if *E* were not caused by *O* and, indeed, even if *O* were not to exist, in the sense of 'could' in which this expresses metaphysical possibility. It is *this* aspect of CTP, I think, that motivates the thought that it promotes scepticism. For it seems to be consistent with CTP that *all* of a subject's perceptual experiences *could*—in the metaphysical sense of 'could'—lack causes in the shape of external objects. And if that *were* the case for any given subject, *S*, then *S* would never *actually* be perceiving any such object but only be *merely seeming* to. In other words, it seems that, according to CTP, it is not metaphysically impossible for a subject to be *entirely mistaken* concerning the veracity of his or her perceptual experiences.

In point of fact, I think that this conclusion is too strong. It may well be that CTP itself does not *rule out* the metaphysical possibility of such global perceptual error, but I don't think that it implies that such global error *is* metaphysically possible. After all, CTP may very plausibly be consistently combinable with other philosophical claims which imply that such global error is metaphysically impossible, in which case CTP itself cannot imply that such global error *is* metaphysically possible.¹⁹ To make this point clear, call the thesis of global perceptual error *G* and the putative philosophical claim which rules it out *P*. Then we have that *P* entails $\neg \Diamond G$. Now suppose that CTP entails $\Diamond G$ but also that CTP and *P* are consistently combinable and thus that $\Diamond (CTP \ \& \ P)$. Then, however, we have a contradiction on our hands, because if ' $\Diamond (CTP \ \& \ P)$ ' is true then so must be ' $\Diamond (A \ \& \ B)$ ', where *A* and *B* are any two propositions entailed by CTP and *P* respectively. So letting $\Diamond G$ and $\neg \Diamond G$ be two such propositions, we find ourselves committed to the truth of ' $\Diamond (\Diamond G \ \& \ \neg \Diamond G)$ ', which we know to be

the most obvious and salient sense of that expression, and so can happily allow that it affords us knowledge of them 'by acquaintance'.

¹⁹ The classic example of such a claim is Descartes's claim that a non-deceiving God necessarily exists who would not allow us to be victims of such global error. CTP is, of course, consistently combinable with Descartes's claim.

logically false. Hence, we must deny that CTP entails $\diamond G$ and make instead only the weaker claim that CTP does *not* entail $\neg \diamond G$ —or, as I put it earlier, that CTP does not *rule out* the metaphysical possibility of global perceptual error. On reflection, it should be clear in any case that CTP does not entail $\diamond G$, for the following reason. CTP at most entails that, for each object O that S actually perceives, S *could be* merely seeming to perceive O , in the metaphysical sense of ‘could’, and thus that *any one* of S ’s perceptual experiences could be non-veridical. But this clearly doesn’t imply that CTP entails that *all* of S ’s perceptual experiences could be non-veridical.²⁰ To suppose that it did imply this would be like supposing that the fact that *any one* banknote could be a forgery implies that *all* banknotes could be forgeries. In short, it is wrong to charge CTP with giving succour to global perceptual scepticism. All that CTP implies—and very reasonably—is that perception of ‘external’ objects is *fallible*. That, surely, is something that any acceptable theory of perception should imply.

To be sure, I have not denied that CTP is *compatible* with global perceptual scepticism and so doesn’t *rule it out*. But disjunctivism is no different from CTP in this regard. Nothing in disjunctivism by itself rules out the possibility that, for a given subject S —myself, for example—*all* of S ’s seemings to perceive are *merely* seemings to perceive rather than actual perceivings. It might be thought that this is not so because the disjunctivist contends that we understand what *merely seeming* to perceive is only by contrast with cases of veridical perception, which are taken to be conceptually prior. However, even if the disjunctivist were right about this, it would not imply that disjunctivism rules out the metaphysical possibility of global perceptual error. Consider the following analogy. It might be urged that we understand what a forged banknote is only by contrast with the notion of a genuine banknote—for forgery is by definition an attempt to pass something off as genuine. But this doesn’t imply that it is metaphysically impossible for all banknotes to be forgeries. To be sure, it may be necessary for some people to *believe* that some banknotes are genuine in order for there to be any forged banknotes, but it is not necessary for there actually to *be* any genuine banknotes. The concept of a genuine banknote could fulfil its role of giving sense to that of a forged banknote even if the former concept *in fact* lacked application. By the same token, the concept of veridical perception could fulfil its (putative) role of giving sense to that of merely seeming to perceive even if there were *in fact* no cases of veridical perception—that is, even if there were global perceptual error. Hence, notwithstanding disjunctivism’s contrastive account of *merely seeming to perceive*, it does not rule out global perceptual scepticism any more than CTP does.

So much for the claim that CTP promotes scepticism: it doesn’t do so any more than disjunctivism does.²¹ But what about the claim that CTP is unfaithful to the

²⁰ Formally, the difference is between ‘ $\forall x(x \text{ is a perceptual experience of } S \rightarrow \diamond x \text{ is non-veridical})$ ’ and ‘ $\diamond \forall x(x \text{ is a perceptual experience of } S \rightarrow x \text{ is non-veridical})$ ’.

²¹ I certainly don’t mean to imply here that all disjunctivists think that disjunctivism is better equipped than CTP is to combat the threat of global perceptual scepticism: see footnote 7 above. But I don’t think it can be denied that some philosophers are motivated to favour disjunctivism at least in part by a belief that this is so. My purpose in the preceding three paragraphs has been to make it clear why such a belief is quite unfounded.

phenomenology of perception, by failing to account for the intuitive appeal—or, indeed, to allow for the truth—of naïve realism? It is not entirely clear what is supposed to be at issue here. Certainly, it is true that the act of perceiving an external object, *O*, does not, in most cases, *seem to its subject*, *S*, to involve any kind of causal transaction between *O* and *S*. When I see a vase of red roses standing in front of me, it does not *seem* to me that my seeing this thing is a matter of its affecting me in a certain way. But that is because no such causal transaction is part of the *content* of my visual experience. To use the terminology introduced earlier, no element in the content of my visual experience registers such a transaction. All that is registered is the object of perception—the vase of roses—and certain of its perceptible features, such as the redness of the roses. It is perfectly understandable why this should be so, if we think of perception in evolutionary terms. Perceptual systems evolved to inform higher animals about the properties of objects in their physical environments, not to inform them about the causal relations between those objects and themselves in virtue of which (according to CTP) they perceive those objects. Human beings are sufficiently intelligent to be able to use their perceptual systems to find out about those causal relations, but that is not the purpose for which evolution ‘designed’ those systems. So, evolutionary psychology can explain why CTP is not ‘intuitively’ true—explain, that is, why CTP does not strike us *pre-theoretically* as being true simply in virtue of our knowing what it is like to be a subject of perception. But that CTP is not intuitively true in this sense cannot be held against it, given that we have a perfectly good explanation of why our intuitions in this case should be misleading.

Can it even be said that CTP necessarily conflicts with the truth of naïve realism? Well, that depends on what is meant by ‘naïve realism’. If naïve realism is simply taken to be the view that, when a subject, *S*, perceives an external object, *O*, *O itself* is a constituent of the content of *S*’s act of perceiving, then we have already seen that CTP is consistent with such a thoroughly externalist account of perceptual content.²² CTP can allow that, in this sense, perception puts us *directly in contact with* external objects. Of course, if naïve realism is held to imply more than just this, it may well not be consistent with CTP. For example, if naïve realism is held to imply that we don’t have ‘perceptual experiences’, or that causation is not essentially involved in perception, or that all talk of sensory qualia or ‘what it is like’ to enjoy certain conscious mental states is misguided, then it is not consistent with CTP, at least as I have formulated the latter. But the more that ‘naïve realism’ is held to imply in such ways, the less clearly it deserves to be called *naïve*. After all, naïve realism is supposed to be the intuitive view of the nature of perception that is instinctively adopted by philosophically untutored subjects—and even, perhaps,

²² M. G. F. Martin characterizes naïve realism as follows: “The Naïve Realist . . . claims that . . . some of the objects of perception . . . are constituents of the experience. No experience like this, no experience of fundamentally the same kind, could have occurred had no appropriate candidate for awareness existed” (2004: 39). The first sentence quoted here is relatively uncontroversial, but the second is not. It makes an essentialist claim which cannot plausibly be read into naïve realism, if the latter is construed as being genuinely *naïve*—that is, as constituting philosophically untutored opinion concerning the nature of perception. See also footnote 6 above.

by advocates of CTP when they are not being philosophically reflective. But in that case it seems improper for opponents of CTP to recruit the term 'naïve realism' to denote a position carrying very much philosophical baggage in the form of implicit denials of the kinds just cited.

What about the supposedly 'diaphanous' or 'transparent' character of perception? If this amounts, once more, to a point concerning the *content* of perception—that external objects *themselves*, rather than anything representing them, are elements of perceptual content—then we have already dealt with it. But I think that it amounts to more than just that. The suggestion, I think, is that the diaphanous character of perception is incompatible with the view, integral to CTP as I have formulated it, that in perceiving external objects we enjoy conscious mental states possessing what I have called *intrinsic characters*. The complaint is that, to be faithful to the phenomenology of perception, we must acknowledge that in perception we are conscious of *nothing but* the objects of perception and their properties—we are *not* conscious of any intrinsic feature of the mental state that we are in. However, I have said only that conscious perceptual experiences *have* intrinsic characters, not that we must be *consciously aware of those characters as such* whenever we enjoy perceptual experiences. I do indeed believe that we can become (or be made) aware of the intrinsic character of a perceptual experience: we do so whenever we consciously reflect on what it is like to enjoy the experience. But it would be absurd to suppose that we are *always* engaging in such conscious reflection whenever we undergo perceptual experiences, and hence whenever we perceive things. Certainly, it is no implication of CTP that this is the case. Largely, what is involved here is a matter to do with *attention* rather than perception. Normally, when we perceive an object it is *the object* that we are attending to, not any feature of our *perceptual experience* of that object. But we can, if need be, switch our attention to the latter. This is, in effect, what artists sometimes do in order to get the perspective right in a drawing. Similarly, when we are watching a sporting event on television, we normally attend to the scene of action—the movements of the players and so forth—but we *can* switch our attention to the array of colours on the television screen itself and ignore the players. What we can't very easily do is to attend to *both* at the same time. So it is with our perceptual experiences: if we attend to intrinsic features of *them*, we cannot easily also attend to *the objects* of those experiences, and vice versa. This is why, for the most part, we simply don't notice the intrinsic characters of our perceptual experiences, with the result that such experiences generally seem 'diaphanous'.

So far I have been trying to defend CTP against the sorts of charges that disjunctivists and their sympathizers typically raise against it, but I have not really been directly critical of disjunctivism, nor have I attempted to weigh the merits of the two positions against each other. It might be thought that, even if both theories are equally coherent and meet all of the basic desiderata of a philosophical theory of perception, disjunctivism still wins by default because, allegedly, it involves less of a departure from common sense and incurs no distinctive ontological commitments. However, I think that saying this misrepresents the true dialectical situation. Disjunctivism does not really have a rationale independent of its opposition to CTP, or at least to theories of perception committed to the existence of perceptual

experiences as the ‘common factor’ in cases of perceiving and merely seeming to perceive. It is essentially a *negative* theory, *denying* the existence of any such common factor.²³ Thus, if CTP can be defended against the various charges brought against it by disjunctivists, most of the wind can be taken out of the sails of disjunctivism. However, there does still remain the question of whether the main argument in favour of CTP—what I have called the cut-off argument—can be neutralized by the disjunctivist. In other words, can disjunctivism satisfactorily accommodate on its own terms the considerations that are drawn on by the cut-off argument and thereby show that the invocation of ‘perceptual experiences’ as these are conceived by CTP is explanatorily otiose? My own view—perhaps unsurprisingly—is that it cannot do so at all plausibly.

It will be recalled that at the heart of the cut-off argument is the claim that if, in a case of a subject *S*’s consciously perceiving an external object *O*, all causal interaction between *O* and *S* were *cut off* at any point between *O* and the periphery of *S*’s body, but the state of *S*’s body remained the same—if, for example, in the visual case, the pattern of neural stimulation in *S*’s retina remained the same, being produced by a different cause—then *S* would be in a conscious mental state with exactly the same intrinsic character as before. The proposal then is that the conscious mental state that *S* is in when perceiving *O* just *is* the same as the conscious mental state that *S* would be in if all causal interaction between *O* and *S* were cut off—and hence that this conscious mental state is a ‘common factor’ in the two cases, the first being a case of perceiving and the second a case of merely seeming to perceive. Or, if this formulation is not liked because it involves a subjunctive conditional, let us put it this way: if *S* and *S** are two subjects and *S* is perceiving *O* while *S**’s circumstances differ from those of *S* *only* inasmuch as all causal interaction between *O* and *S** is cut off, then *S* and *S** are both in a conscious mental state of exactly the same kind. Notice that one version of the latter formulation is supplied by letting *S* and *S** be the *same* subject *at two different times*: this is the case in which, for example, *S* first perceives *O* and then all causal interaction between *O* and *S* is cut off, without any alteration to the state of *S*’s body.²⁴ The disjunctivist, in denying the existence of any ‘common factor’ in these situations, is committed to saying that, for instance, in the temporal case just described, the mere cutting off of all causal interaction between *O* and *S*, without any alteration being made to *S*’s body, suffices to change the kind of conscious mental state that *S* is in. For *S* is first of all *actually* perceiving *O* and subsequently is *merely seeming* to perceive

²³ Again, the *caveat* must be issued here that disjunctivism isn’t committed to asserting that perception and hallucination have *nothing at all* ‘in common’: see footnote 8 above. When I speak in the main text of a ‘common factor’, I should always be understood to intend this in the robust sense required by CTP. I should also acknowledge that, of course, disjunctivists cannot be expected to agree with my assessment of the dialectical situation. In response, I would simply point out the incontrovertible fact that disjunctivism in the philosophy of perception first emerged only in response to some alleged difficulties with CTP and related theories, not autonomously as an independently motivated theory of perception.

²⁴ This could be arranged by, for example, deflecting the light from *O* and replacing it by light from another source (such as a hologram) in such a way as to ensure that the pattern of light-waves impinging on *S*’s retina was unchanged.

O,²⁵ both of which are conditions that involve *S* being *conscious* and hence that involve *S* being in a conscious mental state of *some* kind—and yet not, according to the disjunctivist, of the *same* kind. But who now should be accused of ontological extravagance—the advocate of CTP, who says that *S* remains in a conscious mental state of the *same* kind, or the disjunctivist, who says that *S* undergoes a *change* in the kind of conscious mental state that he is in? Both theorists may say that *S* is in a different *token* state at the two different times—although the advocate of CTP may not even need to say this—but the disjunctivist must additionally say that those token states are of different *types*. Seen in this light, CTP's claim that perceiving is really a *triadic* relation between a subject, *S*, an object, *O*, and third entity, a perceptual experience, *E*, is not as ontologically unparsimonious as it might initially seem. For the disjunctivist denies this only at the expense of having to proliferate the types—and perhaps also the tokens—of conscious mental states considerably beyond those to which the advocate of CTP is committed. Sometimes, proposing that what *seems* to be a dyadic relation is really a triadic one can effect an ontological economy: and that, I suggest, is precisely what CTP achieves and why it is superior to disjunctivism as a philosophical theory of perception. Of course, it can only succeed in this regard if it can be defended from the various criticisms that have been raised against it. But we should *not* include as a legitimate criticism the claim that CTP is ontologically extravagant in comparison with disjunctivism, for the very reverse of this is true. As for the other criticisms, I hope that I have succeeded in defeating them earlier in this essay. My conclusion is that disjunctivism can be—indeed, I hope, *has been*—comprehensively defeated.

REFERENCES

- Aristotle (1928) *Metaphysics*, ed. W. D. Ross, Oxford: Clarendon Press.
 Child, W. (1994) *Causality, Interpretation and the Mind*, Oxford: Clarendon Press.
 Crane, T. (2006) 'Is There a Perceptual Relation?', in T. Szabó Gendler and J. Hawthorne (eds.), *Perceptual Experience*, Oxford: Clarendon Press.
 Farkas, K. (2006) 'Indiscriminability and the Sameness of Appearance', *Proceedings of the Aristotelian Society*, 106: 39–59.
 Hinton, J. M. (1973) *Experiences*, Oxford: Clarendon Press.
 Lowe, E. J. (1996) *Subjects of Experience*, Cambridge: Cambridge University Press.
 — (1998) *The Possibility of Metaphysics: Substance, Identity, and Time*, Oxford: Clarendon Press.
 — (2000) *An Introduction to the Philosophy of Mind*, Cambridge: Cambridge University Press.
 McDowell, J. (1982) 'Criteria, Defeasibility, and Knowledge', *Proceedings of the British Academy*, 68: 455–79.

²⁵ I am not begging any question against the disjunctivist in saying this, for disjunctivists concede that *as a matter of scientific fact* one cannot perceive an object if one is causally isolated from it—that, for instance, one cannot *see* an object if no light waves from it reach one's retina.

- (1994) *Mind and World*, Cambridge, MA: Harvard University Press.
- Martin, M. G. F. (1997) 'The Reality of Appearances', in M. Sainsbury (ed.), *Thought and Ontology*, Milan: Franco Angeli.
- (2004) 'The Limits of Self-Awareness', *Philosophical Studies*, 120: 37–89.
- Putnam, H. (1999) *The Threefold Cord: Mind, Body, and World*, New York: Columbia University Press.
- Snowdon, P. (1980–1) 'Perception, Vision, and Causation', *Proceedings of the Aristotelian Society*, 81: 175–92.
- (2005) 'The Formulation of Disjunctivism: A Response to Fish', *Proceedings of the Aristotelian Society*, 105: 129–41.

4

Disjunctivism About Visual Experience

Scott Sturgeon

This essay explores an approach to visual experience known as ‘disjunctivism’. Section 1 lays out a conception of visual experience focused on two charged aspects of the phenomenon—its conscious portrayal of the world, and its perceptual contact with the world. Section 2 lays out the backbone of disjunctivism and relates it to that conception of visual experience. Section 3 lays out a range of disjunctive positions and shows how they share the backbone. Section 4 canvasses major motivations for disjunctivism and shows how they confer a dialectical advantage on one version of the approach. Sections 5–6 sketch that version, explore its details, and show how they can be used to rebut criticism found in the literature. Section 7 presents new worries for the view; and section 8 explores a pure kind of disjunctivism—one which avoids these worries—sketching when it counts as well motivated.

1 WORKING TO A DISJUNCTION

Visual experience is remarkable for two reasons. It seems to involve conscious portrayal of the world; and it seems to involve perceptual contact with the world. When one sees a cat in the ordinary way, for instance, that bit of consciousness seems to involve the worldly depiction of a cat, somehow; and it also seems to involve perceptual contact with a cat.

There is much debate about such depiction and perceptual contact. We needn’t commit to any story about them. For our purposes, we shall say merely that visual experience has a portrayal side; and we shall mean by the remark solely that it involves conscious depiction of reality somehow. And for our purposes, we shall say merely that visual experience has a perceptual side; and we shall mean by the remark solely that it can make for perceptual contact with reality somehow. Our starting view will

Thanks to Alex Byrne, David Chalmers, Tim Crane, Kati Farkas, John Hawthorne, Mike Martin, Fiona Macpherson, Susanna Schellenberg, Susanna Siegel, Paul Snowdon, Matt Soteriou, the Birmingham Philosophy Department, and Maja Spener most of all.

be one on which visual experience has portrayal and perceptual sides. We shall take no stand on how any of that works.

Such neutrality is fortunate, of course, for there is no serious consensus about how visual experience portrays the world; and there is none about how it makes for perceptual contact. While most agree that visual experience has portrayal and perceptual sides, none agree on the details. Save for this: everyone thinks the two sides are deeply connected. As we'll see in a moment, the specific common-sense category to which a given visual experience belongs is itself fixed by the way its perceptual and portrayal sides inter-weave.

This process can work out a number of ways. In turn those ways form into a range of cases. On its most successful end are cases in which conscious portrayal of the world is fully grounded in perceptual contact. A bit further down in the range are cases in which an experience's portrayal side is mostly grounded in its perceptual side. These are cases in which the delusive aspect of an experience is dwarfed by the degree of accuracy attaching to its perceptually grounded portrayal of the world. Further down still in the range are higher and higher levels of ungrounded portrayal, until finally there are cases in which portrayal of the world is fully unyoked to perceptual contact. The lesson is simple enough: visual experience varies in the degree to which conscious portrayal of the world is grounded in perceptual contact with that world. Best-case scenario involves fully grounded portrayal. Worst-case scenario involves fully ungrounded portrayal. In-between there are intermediate levels of groundedness, and that is why visual experience makes for a range of cases. That range is ordered by the degree to which worldly portrayal is grounded in perceptual contact.

The range can be idealized with numbers:

$$(1) \quad VE(x) \leftrightarrow (\exists n) DG[Por(x), Per(x)] = n\%.$$

Put in English: an episode is a visual experience exactly if a number measures the degree to which its portrayal side is grounded in its perceptual side. This claim idealizes in at least three spots: it presupposes that there is such a thing as *the* portrayal side of a visual episode; it presupposes that there is such a thing as *the* perceptual side of a visual episode; and it presupposes that the former can be grounded in the latter by degrees usefully measured with numbers.

All three assumptions might be wrong. Visual episodes might have multiple portrayal sides, multiple perceptual sides, or both; and even if they do have just one side of each kind, there may be no useful way to measure the degree to which one grounds the other. Having said that, the thought behind (1) seems obviously right: visual episodes seem to vary in the degree to which their portrayal of the world is grounded in their perceptual contact with it.

Common sense blunts this variation by slotting visual episodes into a threefold scheme. The idea in symbols is

$$(2) \quad VE(x) \leftrightarrow [VP(x) \vee I(x) \vee H(x)].$$

Put in English: an episode is a visual experience exactly if it is a veridical perception, illusion, or hallucination. This is true because all episodes fall into exactly one of three slots in the range of cases idealized at (1): some land on (or near enough for

present purposes to) an end of that range—being fully grounded (for present purposes) or fully ungrounded (in that way)—some land between the extremes—being neither fully grounded (for present purposes) nor fully ungrounded (in that way). Veridical perception falls on the fully grounded end of the spectrum; hallucination falls on the fully ungrounded end of the spectrum; and illusion lands between the extremes. The common-sense types of visual experience—veridical perception, illusion, and hallucination—make for a range of cases ordered by how well portrayal is grounded in perceptual contact. Common sense chunks the range into three useful categories.

This means an event is a veridical perception if its portrayal of the world is fully grounded in its perceptual contact with that world. Veridical perception is a more demanding category of consciousness than is often supposed, more than just accurate perceptual contact with the world. Veridical perception is accurate perceptual contact the accuracy of which comes through, but only through, perceptual contact as such. Veridical perception is the best kind of visual experience, occurring when the portrayal side of an experience springs fully from its perceptual side.

Or again: an event is an illusion if its conscious portrayal of the world springs partly from its perceptual contact with it. Illusion is a less demanding category of consciousness than is often supposed, less than just inaccurate perceptual contact with the world. Illusion is perceptual contact the accuracy of which comes through, but not only through, perceptual contact as such. Illusion is the second-best kind of visual experience, occurring when an experience's portrayal side comes some-but-not-fully through perceptual contact. There are veridical as well as delusive illusions. Old chestnuts involve inaccurate perceptual contact with the world—say when a partially submerged stick looks bent despite being straight—but new chestnuts involve accurate perceptual contact too—say when a red herring looks red not because it is so but because of the ambient light. What makes for illusion—what old and new chestnuts have in common—is conscious portrayal of the world yoked to, but not drawn wholly from, perceptual contact with it.

Or again: an event is an hallucination if its conscious portrayal of the world is not grounded in perceptual contact with that world. Hallucination too is a less demanding category of consciousness than is often supposed, less than just inaccurate non-perceptual portrayal of the world. Hallucination is the worst kind of visual experience, occurring when an experience manifests fully non-perceptual portrayal of the world. There are veridical as well as delusive hallucinations. Old chestnuts involve delusive portrayal of the world—say when Macbeth hallucinates his dagger—but new chestnuts involve accurate portrayal too—say when a red herring is hallucinated to be where a red herring chances to be. What makes for hallucination—what old and new chestnuts have in common—is conscious portrayal of the world unyoked to perceptual contact with it.¹

In order to simplify things, however, I shall ignore the degree-of-groundedness facts which make for the range of cases idealized at (1). I shall also ignore the

¹ Blindsight seems to involve perceptual contact with the world unyoked to conscious portrayal. If that is right, then perceptual and portrayal sides of visual experience are two-way disentangle-able.

common-sense blunting of those facts which makes for the platitude codified at (2). Instead I shall speak simply of Good and Bad visual experiences, Good and Bad episodes, Good and Bad cases. I shall understand Good ones to involve a degree of portrayal groundedness surpassing a theoretically motivated threshold, Bad ones to involve a degree of portrayal groundedness falling below such a threshold, and I shall take no stand on whether the thresholds in question are identical.²

Some disjunctivists claim that every experience is either a Good case or a Bad one; and they say this because every case involves perceptual contact with the world or fails to do so (vagueness aside). On this view: what makes for Goodness is perceptual contact with the world, and what makes for Badness is lack of it. Other disjunctivists liken Good and Bad cases to ends of the spectrum mentioned before. On their view: what makes for Goodness is full grounding of portrayal in perceptual contact, and what makes for Badness is no grounding of portrayal in such contact. This way of carving things makes out illusion to be neither Good nor Bad.

On the first approach, a single threshold is needed to make for the divide between Good and Bad cases: any grounding of portrayal at all in perception is both necessary and sufficient for Goodness; and thus a total lack of it is both necessary and sufficient for Badness. On the second approach, two thresholds are needed to make for the divide between Good and Bad cases: when the degree of portrayal groundedness surpasses one of them, that makes for Goodness; when it sinks below the other, that makes for Badness. Nothing in what follows depends on which approach is taken, so I will assume for simplicity that a single threshold is used; and I will also assume, again for simplicity, that this threshold is both exact and constant across conversational context. This will allow me to speak—doubtless contrary to fact—as if every visual experience is absolutely Good or absolutely Bad.

With these assumptions in place we can symbolize a thought prominent in work on disjunctivism:

$$(3) \quad VE(x) \leftrightarrow [G(x) \vee B(x)].$$

Put in English: an episode is a visual experience exactly if it's a Good case or a Bad one. This comes to the thought—in our terminology and relative to our assumptions—that an episode is a visual experience exactly if its conscious portrayal of the world is grounded in its perceptual contact with it to a degree that surpasses a theoretically motivated threshold or to a degree that fails to do so. Disjunctivism is most easily discussed by speaking of Good and Bad cases, so that is what I shall do; but I shall take it as read that Goodness and Badness turn on the degree to which an episode's portrayal side is grounded in its perceptual side: Good cases involve sufficient grounding, Bad ones involve insufficient grounding (relative to our assumptions).

² In the terminology of Byrne and Logue's contribution to this volume, I take no stand on whether VI-v-H disjunctivism is preferable to V-v-IH disjunctivism, or vice versa. Byrne and Logue's essay nicely locates various theorists relative to distinctions we're making here. Their essay and this one were written independently of one another.

2 THE BACKBONE OF DISJUNCTIVISM

Disjunctivism is not really a theory of visual experience but a rough approach to the topic. Its defenders include John Hinton, Mike Martin, John McDowell, and Paul Snowdon. These philosophers believe all manner of things about visual experience; and often those things conflict with one another. One must tread carefully, then, when shooting for the backbone of the approach they champion. One stands no chance of hitting it unless one aims at a high level of abstraction. That will be my strategy.

I suggest two views are jointly the backbone of disjunctivism (or at least should be). One concerns the conscious character of Good cases—what it is consciously like to enjoy Good visual experience. The other concerns the conscious character of Bad cases—what it is consciously like to enjoy Bad visual experience. Each view can be spelled out in various ways—as we’ll see in the next section—but it is particularly important that the claim about Bad character can be multiply understood. Those understandings make for a range of disjunctive positions; and the positions in that range vary in their radicalism, resources, and appeal.

The backbone claim about Good character is this:

(G) Good character derives from bits of the physical world standing in an explanatorily basic relation to percepts.

The idea here, roughly stated, is that what it’s like to enjoy Good experience is fixed by the explanatorily fundamental perceptual contact one has with the world. Needless to say, that broad idea can be understood in various ways. McDowell says that Good cases involve real-world facts manifesting themselves to percepts; and he treats the manifestation relation as explanatorily basic. Martin says that Good cases involve objects, events, and features being taken in by percepts; and he treats the taking-in relation as explanatorily basic. Other disjunctivists spell out (G) in their own way. But all claim—or at least should claim—that Good character comes from explanatorily basic conscious perceptual contact with the world. This idea is the soul of disjunctivism’s take on Good character. That character is said to be drawn directly from perceptual contact with the world, contact which cannot itself be explained in more basic mental idioms. My suggestion is that this is the first half of disjunctivism’s backbone. For reasons to be spelled out in section 4, it is the most important half of that backbone.

The other half concerns Bad character:

(B) Bad character does *not* derive from bits of the physical world standing in an explanatorily basic relation to percepts.

The idea here, roughly stated, is that what it’s like to enjoy Bad experience is not fixed by explanatorily basic perceptual contact with the world. This thought can be spelled out in multiple ways too; and that yields a range of disjunctive positions. In turn that range can be ordered by the degree of difference they see between Good and Bad cases.

3 A RANGE OF DISJUNCTIVE POSITIONS

According to the Theory of Appearing, Good character derives from real-world objects appearing to percipients; the appearing relation is explanatorily basic; and Bad character derives from a different kind of thing—a non-real-world object—appearing to percipients. The approach says Good and Bad cases turn out to be different, therefore, but not terribly so, as they both involve people standing on the catching-end of the appearing relation. The key difference between Good and Bad cases comes on the pitching-end of that relation. Good cases involve real-world pitchers; Bad cases involve unreal-world pitchers (as it were). Good and Bad cases are not radically different, as they share a basic metaphysical structure. That structure is filled out with a single relation and a single kind of thing on its catching end. Good and Bad cases differ only in what kind of thing appears to people in each sort of case. This is a mild form of disjunctivism, a view on which Good and Bad cases are different but not hugely so.³

According to one reading of McDowell, at least, Good character involves real-world facts manifesting themselves to percipients; the manifestation relation is explanatorily basic; Bad character involves a different kind of object—a mere appearance—standing in a different basic relation to percipients—itsself expressed with the word “appearing”. Call this Reading A of McDowell. On it Good and Bad cases differ more than they do on the Theory of Appearing. After all, Reading A of McDowell says that Good and Bad cases both involve people standing on the catching end of a relation—in line with the Theory of Appearing—but Reading A continues that the two kinds of case involve distinct relations as well as distinct kinds of thing on their pitching ends. Good cases involve worldly facts manifesting themselves to people. Bad cases involve unworldly things appearing to people (as it were). On this kind of view, Good and Bad cases are fairly different. They share a deep metaphysical structure; but they fill out that structure with two relations and two kinds of pitching-end object. The result is a not-so-mild form of disjunctivism, a view on which Good and Bad cases are fairly different from one another.

According to another reading of McDowell, at least, Good character involves real-world facts manifesting themselves to percipients; the manifestation relation is explanatorily basic; and Bad character involves intentionality—that ill-understood feature of propositional attitudes. Call this Reading B of McDowell. On it Good and Bad cases differ more radically than they do on either the Theory of Appearing or Reading A of McDowell. After all, Reading B claims that Good character involves people standing on the catching-end of a real relation to real things, while Bad character involves intentionality possibly linked to unreal things. Reading B thus denies that Good and Bad cases share a deep metaphysical structure. It does not generate Good and Bad character by filling out such a structure in some or mostly the same way. Rather, it sees Good cases as relational and Bad ones as intentional. This is a

³ See Alston (1999), Chisholm (1965), and Langsam (1997).

strong form of disjunctivism, a view on which Good and Bad cases are obviously very different.⁴

And then there is a view I call Reflective Disjunctivism. Like all versions of disjunctivism, the view claims that Good character derives from bits of the physical world standing in an explanatorily basic relation to percipients. Like all versions of disjunctivism, the view claims that Bad character does not derive in that way. But Reflective Disjunctivism distinguishes itself from other versions of disjunctivism—and thereby earns its name—by claiming that Bad character derives from *negative epistemics*: in particular, the view says that Bad character derives from Bad cases being indiscriminable from Good ones. As we'll see in a moment, the key idea is that visual experience as such is nothing but a state of reflective indiscriminability from Goodness.

There are many ways to fill out the picture. They all result in a view which claims that Good cases involve people standing on the catching-end of a real relation to bits of physical reality, while Bad ones involve nothing like that at all. Rather, Reflective Disjunctivism insists that Bad cases are merely situations which cannot be discriminated from Good ones. Bad character is then spelled out by appeal to negative epistemics: when Macbeth hallucinates his dagger, for instance, what it is like for him visually derives solely from his inability to discriminate his case from one of seeing a dagger properly. Reflective Disjunctivism denies that Good and Bad cases share a deep metaphysical structure. It does not generate Good and Bad character by filling out such a structure in partly or mostly the same way. The view sees Good cases as relational in nature and Bad ones as deeply epistemic. It is a radical form of disjunctivism, a view on which Good and Bad cases are very, very different.

I have drawn a chart of the views just sketched. Its first column lists positions by name. Its second column lists a view's take on where Good character comes from. Its third column lists a view's take on where Bad character comes from. And its fourth column lists key differences between Good and Bad character on each approach.

Table 4.1 Types of disjunctivism

| | Good Character | Bad Character | Good versus Bad |
|------------------------------|---|--|---|
| <i>Theory of Appearing</i> | Worldly objects appearing to persons | Unworldly objects appearing to persons | <ul style="list-style-type: none"> * One metaphysical structure * A single relation used to build both types of character * Two kinds of pitching end <i>relatum</i> |
| <i>Reading A of McDowell</i> | Worldly facts manifesting themselves to persons | Unworldly facts appearing to persons | <ul style="list-style-type: none"> * One metaphysical structure * Two relations used to build character, one for Good and one for Bad * Two kinds of pitching end <i>relatum</i> |

⁴ There is a great deal of discussion about what McDowell's position really is. I take no stand on that. See McDowell (1982), (1994), (1998); and for relevant discussion see Byrne and Logue's contribution to this volume, Crane (2001), Hawthorne and Kovakovich (2006), and Thornton (2004).

| | | | |
|---------------------------------|---|---|---|
| <i>Reading B of McDowell</i> | Worldly facts manifesting themselves to persons | Possibly non-actual states of affairs linked to persons by their intentional states | <ul style="list-style-type: none"> * Two metaphysical structures * One relation used to build Good character * Intentionality used to build Bad character * Only real objects slot into the relation * Potentially unreal objects can figure in intentionality |
| <i>Reflective Disjunctivism</i> | Worldly objects and/or events and/or features being taken in by persons | Episodes being indiscriminable from Goodness | <ul style="list-style-type: none"> * Two metaphysical structures * One relation used to build Good character * Indiscriminability used to build Bad character |

This spread of views indicates that disjunctivism can be more or less radical. All versions of the approach endorse (G) and (B). But there are very different reasons why one might go in for the latter. This leads to multiple ways a disjunctivist might endorse the idea that Good and Bad cases are different. Some of those ways are theoretically mild, others are theoretically severe.

All versions of disjunctivism account for Good character by appeal to explanatorily basic conscious perceptual contact with the physical world. Reflective Disjunctivism does this like other versions of the approach. The key difference between it and those rivals turns on their take on Bad character. After all, Reflective Disjunctivism accounts for such character by appeal to indiscriminability. Its rivals do so by appeal to non-physical objects, extra basic relations, intentionality, and suchlike. These resources are dialectically non-trivial in a way that indiscriminability is not; for every view of experience—disjunctive or otherwise—can make use of indiscriminability. The fact that experience is indiscriminable from veridical perception is a datum of everyday life. It is a non-partisan resource in the area. Every theory can use it for whatever it sees fit. By using indiscriminability to account for Bad character, therefore, Reflective Disjunctivism ensures that its resources are austere when compared to those of its charted rivals. They too could use indiscriminability to deal with Bad character; but they do not choose to do that. Instead they invoke resources like intentionality or non-physical objects. What sets Reflective Disjunctivism apart from its charted rivals—indeed what makes it a remarkably distinctive view—is its use of indiscriminability to limn the contours and character of visual experience as such.

4 MOTIVATIONS

In this section I sketch the major motivations for disjunctivism about visual experience. I mean neither to endorse nor reject any motivation to follow, only to set out

the motivations in bare outline. A key fact about them will emerge in due course; and that fact will be used, in section 8, to pin down a surprisingly interesting version of disjunctivism.

To begin, a great many lines of thought have been used to support disjunctivism. In turn those lines partition into three natural groups. Members of each group bear a striking resemblance to one another. Members across groups do not do so. For this reason, there are three *families* of motivation for disjunctivism in the literature rather than three motivations as such.

The families are sometimes run together—or at least mentioned in one breath—but they are very different from one another. They do not include every consideration which could be used to argue for disjunctivism; but they do contain every consideration which has been used for that purpose with influence. In alphabetical order I call them the *epistemic*, *phenomenal*, and *semantic* families of motivation for disjunctivism.

(a) Epistemic motivations connect a disjunctive metaphysics of experience with ideas about external-world knowledge—knowledge of picnics, knowledge of pumpkins, and so forth. The ideas in this family can be put together in various dialectical ways. One might start with our possession of external-world knowledge, for instance, and then argue that its possession requires a disjunctive metaphysics of experience. This would be to mount a straightforward argument for such a metaphysics on the basis of our knowledge. Or one might use the conceptual coherence of disjunctivism to undermine the conceptual inevitability of a sceptic's argument against external-world knowledge. This would be to mount a defence of the sheer possibility of such knowledge, with said defence in turn being thought somehow to raise the philosophical likelihood of disjunctivism as such. Or one might argue that the conceptual coherence of disjunctivism shows the conceptual incoherence of a sceptic's argument against external-world knowledge. This would be to mount a straightforward refutation of the sceptic, with that refutation also being thought somehow to raise the philosophical likelihood of disjunctivism. And so on.⁵

In each case, support for disjunctivism is drawn from reflection upon external-world knowledge. In each case, moreover, the line turns crucially on the disjunctive approach to *Good* character; for experience with such character is said to be what makes for (visually based) external-world knowledge. The key idea—time and again in this family of motivations—is that Good character derives from bits of the physical world standing in an explanatorily basic relation to percipients, and, as such, that Good cases make for (visually based) external-world knowledge. All lines of thought in the epistemic family of motivations for disjunctivism pivot on this idea about Good character. All of them turn crucially on (G): the thought that Good visual consciousness derives from explanatorily basic perceptual contact with the world.

(b) Phenomenal motivations for disjunctivism connect a disjunctive metaphysics of visual experience with what it is like to enjoy that experience. Ideas in this family can

⁵ McDowell famously emphasizes epistemic motivations for disjunctivism. See McDowell (1982), (1994), (1998) and especially his contribution to this volume.

also be put together in various dialectical ways. The most straightforward involves direct argument for disjunctivism by introspection of what visual experience is like. The line says that introspection itself is sufficient to establish that visual experience has (G)-like character in at least some cases. From this it follows, of course, that visual experience has (G)-like-character-or-something-else in all cases. Introspection is thus said to lead directly to disjunctivism about visual experience; and in the verbal tradition, at least, one often hears this line of thought accompanied by the remark that a (G)-style canvassing of visual experience is no more than an innocent gloss on how such experience strikes one introspectively. A less straightforward (but more creative) argument for (G) compares what it is like to see with what it is like to imagine seeing. The line claims that a disjunctive metaphysics of experience falls out of the comparison. By holding up what it is like to see, and comparing it to what it is like to imagine seeing, one then argues for a (G)-like approach to character in some cases.⁶

In each of these lines of thought, support for disjunctivism is drawn from reflection upon what visual experience is like. In each of these lines of thought, moreover, the crucial point turns on the disjunctive approach to *Good* character; for experience with that character is said to be what makes for the way that visual experience strikes us introspectively. The key idea—time and again in this family of motivations—is that Good character derives from bits of the physical world standing in an explanatorily basic relation to percipients, and, as such, that Good cases strike us introspectively as they do. All lines of thought in the phenomenal family of motivations pivot on this idea about Good character and introspection. All of them turn crucially on (G): the thought that Good visual consciousness derives from explanatorily basic perceptual contact with the world.

(c) Semantic motivations for disjunctivism connect a disjunctive metaphysics of experience with thought about the external world—thought about picnics, thought about pumpkins, and so on. The ideas in this family can also be put together in various dialectical ways. One might start with the possession of thought about the external world, for instance, and then argue that its possession requires a disjunctive metaphysics of experience. This would be to mount a straightforward argument for disjunctivism on the basis of actual thought. Or one might use the conceptual coherence of disjunctivism to undermine the conceptual inevitability of any argument against the possibility of external-world thought. This would be to mount an in-principle defence of thought targeting the external world, with said defence in turn being somehow meant to raise the philosophical likelihood of disjunctivism as such. Or one might argue that the conceptual coherence of disjunctivism shows the conceptual incoherence of any argument against the possibility of external-world thought. This would be to mount a straightforward refutation of nay-saying about such thought, with said refutation again being meant somehow to raise the philosophical likelihood of disjunctivism. And so on.⁷

⁶ Martin (2002) deploys a line like this in its main argument for disjunctivism.

⁷ Putnam famously emphasizes semantic motivations for disjunctivism in Putnam (1999). Those motivations are also pronounced in McDowell (1998).

In each case, support for disjunctivism is drawn from reflection upon external-world thought. In each case, moreover, the line turns crucially on the disjunctive approach to *Good* character; for experience with that character is said to be what makes for (visually based) external-world thought. The key idea—time and again in this family of motivations—is that Good character derives from bits of the physical world standing in an explanatorily basic relation to percipients, and, as such, that that is why Good cases make for (visually based) external-world thought. All members of the semantic family of motivation pivot on this idea about Good character and mental content. All of them turn crucially on (G): the thought that Good visual consciousness derives from explanatorily basic perceptual contact with the world.

In a nutshell: there are three families of motivation for disjunctivism. One reflects on knowledge of the external world to arrive at (G)’s take on Good character. One reflects on phenomenology to arrive at (G)’s take on Good character. One reflects on thought of the external world to arrive at (G)’s take on Good character. The key idea each time—the one on which every dialectic turns—is that Good character derives from explanatorily basic perceptual contact with the world. If disjunctivism is motivated at all—by epistemic, phenomenal, or semantic considerations—it is by virtue of this aspect of its story. (G)’s *bona fides* exhaust the view’s chance of being motivated. For this reason, though, nothing in the view’s motivations yields a take on Bad character: neither epistemic, phenomenal, nor semantic considerations lead us to any account of such character. Root motives for disjunctivism fail to prompt any *spiel* about what it is like to suffer delusive visual experience.

This generates a non-trivial dialectical advantage for Reflective Disjunctivism over its charted rivals. After all, the view accounts for Bad character by appeal to something in everyone’s toolkit: namely, the indiscriminability of experience—and in particular the indiscriminability of Bad experience—from veridical perception. If this non-partisan resource can be used to account for Bad character—and we’ll look more closely at that in section 8—we shall have no need for intentionality, non-physical objects, or extra basic relations to do so. We shall have no need for anything found in the ideology of Reflective Disjunctivism’s charted rivals concerning Bad character. If indiscriminability can be used to account for such character, a dialectically free resource will do the trick; and that will make for a serious head start in the debate. After all, Reflective Disjunctivism will account for Good character with (G)—like all versions of disjunctivism—but unlike its charted rivals the view will account for Bad character by appeal to a non-partisan resource, by appeal to something everyone recognizes. This will make for a substantive dialectical advantage. Our initial task, therefore, is to see whether Reflective Disjunctivism is plausible. To do that we shall need a much better grip on its details than we presently have.

5 SKETCHING THE APPROACH

Recall that Bad cases consciously portray the world as being a way in which they are not aptly grounded. Reflective Disjunctivism spells them out by appeal to reflective

indiscriminability from Goodness—by appeal, that is to say, to indiscriminability from Goodness by introspective reflection.⁸ An episode is said to be a Bad Φ -case only if it is so indiscriminable from Good Φ -cases:

$$(4) \quad B^{\Phi}(x) \rightarrow xIND_g^{\Phi}.$$

On the conception of indiscriminability used by Reflective Disjunctivism, moreover, Good Φ -cases turn out to be automatically indiscriminable from themselves. Eventually we'll see why that's so; but for now we note merely that because it is so facts of this form turn out to be conceptual truths:

$$(5) \quad G^{\Phi}(x) \rightarrow xIND_g^{\Phi}.$$

Put in English: an episode is a Good Φ -case only if it is indiscriminable from Good Φ -cases.

Now recall that (3) is the claim that an episode is a visual experience as of Φ exactly when it's either a Good Φ -case or a Bad one. This, (4) and (5) jointly entail

$$(6) \quad VE^{\Phi}(x) \rightarrow xIND_g^{\Phi}.$$

Put in English: an episode is a visual experience as of Φ only if it is indiscriminable from Φ -Good cases. A striking feature of Reflective Disjunctivism—to be vetted in section 8—is that it accepts the converse idea:

$$(7) \quad xIND_g^{\Phi} \rightarrow VE^{\Phi}(x).$$

Put in English: an episode is indiscriminable from Φ -Good cases only if it is a visual experience as of Φ .¹⁰ Reflective Disjunctivism is thus led, via (6) and (7), to its main claim about the extension of visual experience:

$$(8) \quad VE^{\Phi}(x) \leftrightarrow xIND_g^{\Phi}.$$

⁸ As Siegel's contribution to this volume notes, however, it seems possible that visual experience portrays the world as being a way it could not be—Escher-wise, so to say. If that is right, such experience will almost certainly be reflectively discriminable from veridical perception. Martin's way with the issue—sketched at the end of Martin (2004)—is to apply his reflective construction in the first instance to elements of experience rather than whole experiences. It is not obvious that this will work; for it is not obvious that the content of experience “adds up” simply from that of its elements. Escher-like conflicts might crop up as a result of the global arrangement of elemental bits of subject matter. Unfortunately Martin does not say enough to know what he thinks about this, so I shall ignore the issue entirely. Chapter 5 of my (2001) reinforces the worry.

⁹ Here I am being purposefully vague in my symbolization. I want to use English idioms common to the literature on visual experience, along with natural-but-multiply-interpretable symbolizations of them. This should raise pointed questions of interpretation as we go. The aim will be to resolve those questions into recognizably clear symbols, and recognizably clear English, as we proceed.

¹⁰ Reflective Disjunctivism must find a non-question-begging way to delineate the episodes available to introspective reflection. In a sense, after all, episodes unavailable to it are automatically indiscriminable from anything else by its means. I shall take this as read in what follows, always assuming that episodes under discussion are introspectable. For relevant discussion, see the distinction between strong and weak indiscriminability properties in Siegel's contribution to this volume, and the last note of this essay.

Put in English: an episode is a visual experience as of Φ exactly when it is indiscriminable from Good Φ -cases.

If that is right—and we'll consider whether it is in a moment—the most natural explanation would be the identity of visual experience and indiscriminability. (8)'s truth would be most cleanly explained by Reflective Disjunctivism's main claim about the nature of visual experience:

$$(9) \quad \lambda x \text{VE}^\Phi(x) = \lambda x(x \text{IND} g^\Phi).$$

Put in English: being a visual experience as of Φ is identical to being a state indiscriminable from Good Φ -cases; in other words, being a Φ -experience is the same thing as being an episode that cannot be reflectively discriminated from such cases. This is the view's core conception of visual experience. It sees it as nothing but reflective indiscriminability from Goodness.

To understand Reflective Disjunctivism properly, we need a better grip on its notion of reflective indiscriminability. There are several versions of this notion at work in the philosophical literature. Each could be used to build a version of Reflective Disjunctivism. Full treatment of the approach requires inspecting every (9)-like claim got by appeal to a coherent notion of indiscriminability. We have no time for that here, of course; so we focus instead on the most prominent type of indiscriminability notion found in philosophical discussion of perception. That notion turns on the idea of *knowledge*.

To see how note that indiscriminability is inability to be discriminated. On a knowledge-based understanding of indiscriminability, discrimination is itself construed as the activation of knowledge; in particular, it is construed as the activation of knowledge that things are not a certain way. This means that a knowledge-based take on indiscriminability sees it as the inability to be known not to be a certain way. When the root idea is applied to (8) the result is

$$(10) \quad \text{VE}^\Phi(x) \leftrightarrow \neg \diamond \text{K}[\neg(x \approx g^\Phi)];$$

and when it is applied to (9) the result is

$$(11) \quad \lambda x \text{VE}^\Phi(x) = \lambda x \neg \diamond \text{K}[\neg(x \approx g^\Phi)].$$

Put in English: the first claim is that an episode is a visual experience as of Φ exactly when it cannot be known not to be a Good Φ -case; and the second claim is that an episode's being a Φ -experience is identical to its being not possibly known not to be a Good Φ -case.

Pulling all this together, then, Reflective Disjunctivism plumps for (G), (B), (11), and so (10). The view claims that Good character derives from bits of the physical world standing in an explanatorily basic relation to percepts. It continues that Bad character does not derive in that way. And it accounts for the source of Bad character by appeal to the purely experiential nature of Bad cases. This is fleshed out via negative epistemics: visual consciousness is said to be a state of reflective indiscriminability from Goodness—a state of such indiscriminability from bits of the physical world standing in an explanatorily basic relation to percepts.

6 WORKING THROUGH SOME DETAILS

Our take on Reflective Disjunctivism should depend on how principles (10) and (11) are fleshed out. That can be well understood by seeing how proponents of them might deal with challenges put to their view. Working through several proves a useful way to see how Reflective Disjunctivism must be interpreted if it is to be credible.

To that end, note the right-hand sides of (10) and (11) each lean on three non-trivial items: modality, knowledge, and the claim that x is not a Good Φ -case. The rest of this section presents challenges to (10) and (11). Responses to them help fix our take on the modality, knowledge, and claims used by Reflective Disjunctivism.

6.1 What about everyday knowledge of Badness?

It is a datum of everyday life that we can know—in certain contexts at least—that we suffer visual illusion; and it is a datum of everyday life that we can know—in certain contexts at least—that we suffer visual hallucination. It is part of everyday life, therefore, that we can know—in certain contexts at least—that our visual experience is not Good. That seems to conflict with (10) and with (11); for if we know that a given visual episode is not Good, it must *be* not Good by the factiveness of knowledge. It follows that we *can* know of a Bad one that it is not Good; and that looks to mean that visual experience pulls apart from inability to know that one's experience is not Good. In turn that looks to conflict with the extensional claim at (10) and the identity claim at (11). Everyday knowledge of Badness puts pressure on Reflective Disjunctivism from the start.

That pressure must be alleviated by restricting the source of knowledge used in the approach. Specifically, Reflective Disjunctivism must say that its knowledge can only come by introspective reflection. For short: it can only come *by reflection*. The thought must be that when one knows in everyday circumstances that one's visual episode is not Good, that knowledge is got other than by such reflection. Martin notes in (2006), for instance, that it could be got by testimony; so Reflective Disjunctivism must stipulate that knowledge got in that way is not reflective in the experience-making sense. And more generally this is how everyday knowledge of Badness will be marked as compatible with the intended readings of (10) and (11). That knowledge will be said not to refute Reflective Disjunctivism from the start; for the view will be said not to cover every kind of knowledge in its founding principles. It will be said to cover only knowledge got by reflection.

Later we'll see whether a restriction like this can be made to work. Here we take the restriction for granted and mark it in our notation:

$$\begin{aligned} (10)_{kr} \quad & \text{VE}^\Phi(x) \leftrightarrow \neg \diamond \text{KR}[\neg(x \approx g^\Phi)]; \\ (11)_{kr} \quad & \lambda x \text{VE}^\Phi(x) = \lambda x \neg \diamond \text{KR}[\neg(x \approx g^\Phi)]. \end{aligned}$$

Put in English: the first claim is that an episode is a visual experience as of Φ exactly when it cannot be known by introspective reflection not to be a Good Φ -case; and

the second claim is that an episode's being a Φ -experience is identical to its being not possibly known by such reflection not to be a Good Φ -case.

6.2 What proposition is involved here?

Reflective Disjunctivism is grounded in a knowledge-based understanding of reflective indiscriminability. Before we can understand the view properly we must know what kind of claim is used in the knowledge said to make for indiscriminability.

The symbolization of this at (10) and (11)—along with its vocalization in English—was left purposefully vague on this issue. In particular, it was left vague what kind of claim is to be used in the knowledge-based approach to reflective indiscriminability. In this essay we focus on a version of Reflective Disjunctivism got by precisifying in line with Williamson's work in indiscriminability. A different precisification—inspired by Martin's work on the topic—is discussed in my 2006.¹¹

Williamson's *Identity and Discrimination* (1990) is a *tour de force* exploration of identity and discrimination. He investigates a great many topics concerning these notions. Along the way he looks at the relationship between the identity of conscious character and our capacity to discriminate its instances. We need only draw from that brilliant discussion our inspiration for a take on the item of knowledge involved in the indiscriminability of episodes.

Roughly, Williamson says that episodes u and v are discriminated exactly when someone activates knowledge that u isn't v . For Williamson: discrimination turns on knowledge of numeric distinctness, so discriminability turns on the possibility of such knowledge, indiscriminability turns on its impossibility. This leads to the idea, for our purposes, that *reflective* indiscriminability turns on the impossibility of knowledge of distinctness based on reflection. On the approach I have in mind:

$$(\text{ind})_w \quad u\text{IND}v \leftrightarrow \neg \text{KR}(u \neq v).$$

Put in English: u is indiscriminable from v exactly when it's impossible to know by reflection that u and v are not identical; exactly when, in other words, it is impossible to know by reflection that u and v are numerically distinct.

When this take on indiscriminability is applied to (10), the result is a faintly Williamsonesque version of Reflective Disjunctivism's main take on the extension of visual experience:

$$(10)_{\text{kr-w}} \quad \text{VE}^\Phi(x) \leftrightarrow \neg \text{KR}(x \neq G^\Phi s).$$

Put in English: an episode is a visual experience as of Φ exactly when it is not possible to know by reflection that it is numerically distinct from Φ -Good cases. Indiscriminability here is understood via knowledge got by reflection alone; and that knowledge is understood to concern plural claims of numerical distinctness. When it cannot be known by reflection that a given episode is distinct from Good Φ -cases, what cannot be known is the claim that the episode is numerically distinct from those cases. What cannot be known, on this view, is a plural claim of numerical distinctness.

¹¹ Nothing of general substance is lost by focusing our attention in this way.

When the view of indiscriminability found at $(\text{ind})_w$ is applied to (11), the result is a faintly Williamsonsque version of Reflective Disjunctivism's take on the nature of visual experience:

$$(11)_{kr-w} \quad \lambda x \text{VE}^\Phi(x) = \lambda x \neg \diamond \text{KR}(x \neq G^\Phi s).$$

Put in English: an episode's being a Φ -experience is identical to its being not possibly known by reflection to be numerically distinct from Φ -Good cases. Once again indiscriminability is understood via knowledge got by reflection; and once again knowledge is understood to concern plural claims of numerical distinctness. We shall discuss versions of Reflective Disjunctivism like this in the remainder of this essay.

6.3 What about children, puppies, and the like?

It is a fact of everyday life that some visual experiencers *lack* the capacity for reflective knowledge seemingly presupposed by (10) and (11). Small children cannot reflect introspectively yet they enjoy visual experience; and the same is true of puppies and other creatures of diminished epistemic capacity. In a nutshell: there are visual experiencers who cannot reflect introspectively and thus cannot know anything by doing so. This puts pressure on (10) and on (11).¹²

That pressure must be relieved by de-coupling the possibility of knowledge used in the approach from capacities enjoyed by visual experiencers. The idea must be to read (10) and (11) so that manifestation of visual experience does not oblige personal epistemic capacity. Rather, visual experience must be said to require *impersonal* epistemic impossibility of some kind. The key question is what kind of possibility is that?

This is a good question. I also think it is an answerable one. The challenge to Reflective Disjunctivism posed by children, puppies, and the like does not amount to refutation of the view. There is a straightforward response to the challenge. Here's how it goes.

A common thought in epistemology is that its subject matter includes impersonal norms or idealizations. When degrees of belief are said to be rational only if measured by probability functions, for instance, the idea need not be that rational agents must actually possess the personal capacity to make their degrees of belief conform to the probability calculus if they are to be subject to probabilistic norms. The idea could be—indeed normally is—that such norms are impersonal, that they limn the contours of rationality divorced from epistemic capacity. The idea is that probabilistic norms apply to an agent's degrees of belief independently of her capacities—mechanism-based or otherwise—to bring credence into line with the probability calculus. Or when belief sets are said to be rational only if logically consistent, for instance, the idea need not be that rational agents must actually possess the personal capacity to make their entire set of beliefs conform to logic if they are to be subject to logical norms. The idea could be—indeed normally is—that such norms are impersonal, that they limn the contours of rationality divorced from epistemic capacity.

¹² Hawthorne and Kovakovich (2006) and Siegel's contribution to this volume object to Reflective Disjunctivism on these grounds. Martin's response in Martin (2006) is critiqued in my (2006).

The idea is that logic-based norms apply to an agent's beliefs independently of her capacities—mechanism-based or otherwise—to bring belief into line with logic.

Impersonal idealizations are the meat and potatoes of epistemic theory. There is no reason why Reflective Disjunctivism should be disallowed from using them. They provide a good resource for dealing with challenges to the approach based on diminished epistemic capacity. On my view defenders of that approach should use them to do so. They should deal with children, puppies, and the like by appeal to idealization implicit in their theory.

The idea would be to read modality on the right-hand sides of (10) and (11) as impersonal idealized epistemic possibility. Just as norms of consistency can be seen to involve impersonal epistemic modality—namely, an impersonal epistemic impossibility of contradiction lurking within rational belief—so (10) and (11) could be seen to involve impersonal epistemic modality—namely, an impersonal epistemic impossibility of reflective knowledge that one's case is not identical to a Good one. On this interpretation of (10) and (11), their modal operator is an idealized epistemic operator abstracting from real-world capacity to know by reflection.

Norms of logic are often thought to make for an impersonal space of possibility concerning justified belief. The idea is that they fix what is impersonally possible for such belief. Similarly, norms of probability theory are often thought to make for an impersonal space of possibility concerning rational degrees of belief. The idea is that they fix what is impersonally possible for such credence. My view is that (10) and (11) should be thought of as resting on similar norms, only this time the norms should be said to apply to reflective judgement rather than belief or credence. Reflective norms—as we might call them—make for an impersonal space of possibility concerning reflective judgement. They fix what is impersonally possible for knowledge grounded in such judgement.

On this approach, reflective norms—together with logic, conceptual links, and other a priori tools, of course—make for a space of impersonal possibility. Modal operators in (10) and (11) reflect the contours of that space. Reflective Disjunctivism should make use of impersonal idealization. The resulting view should be one on which (G) and (B) are true along with these principles:

$$(10)_W \quad VE^\Phi(x) \leftrightarrow \neg \diamond_{\text{imp}} KR(x \neq G^\Phi s).$$

$$(11)_W \quad \lambda x VE^\Phi(x) = \lambda x \neg \diamond_{\text{imp}} KR(x \neq G^\Phi s).$$

Put in English: the first claim is that an episode is a visual experience as of Φ exactly when it is not impersonally possible to know by reflection that it is numerically distinct from Φ -Good cases; and the second claim is that an episode's being a Φ -experience is identical to its being not impersonally possible to know by reflection that it is numerically distinct from Φ -Good cases.

Reflective Disjunctivists should side-step the worry of this subsection by appeal to epistemic idealization. They should see visual experiencers of diminished capacity as enjoying sufficiently robust background states—beliefs, veridical perceptions, and so forth—to fall under impersonal norms of reflective judgement. They should then make use of the impersonal nature of such norms to clarify how

agents of diminished capacity can be visual experiencers despite their diminished capacity.¹³

Epistemic idealization is an impersonal matter. Reflective Disjunctivism tries to pin down the nature of visual experience in epistemic terms. When those terms are said to involve epistemic idealization, the resulting view entails that the nature of visual experience is an impersonal matter, one unfixed by capacities of those who visually experience. The resulting view de-couples the nature of visual experience from the capacities of those who enjoy it; and for this reason, the resulting view is not threatened by visual experiencers of diminished capacity. Proponents of Reflective Disjunctivism are well advised to see their resources to involve idealization. That is the lesson of children, puppies, and the like.

6.4 What about the non-transitivity of indiscriminability?

A worry for Reflective Disjunctivism grows from the putative non-transitivity of indiscriminability.¹⁴ It can be sketched as follows: there might be episodes u , v , and w so that u is indiscriminable from v , v is indiscriminable from w , while u is discriminable from w . Were that to happen, one could not tell that u was distinct from v , one could not tell that v was distinct from w , but one could tell that u was distinct from w . For reasons to be spelled out in a moment, that possibility puts pressure on the idea that a feature of episodes is pinned down by indiscriminability facts concerning them; and that, in turn, puts pressure on Reflective Disjunctivism's claim that experiential features as such are pinned down by indiscriminability.

This all takes some explaining. First we'll spend a paragraph setting out a general fact about relations and properties. Then we'll examine whether that fact makes trouble for Reflective Disjunctivism. Here are the steps in order:

A general fact. No non-transitive relation R and feature F can be so that for variable x and y :

$$(*) \quad [F(x) \ \& \ F(y)] \leftrightarrow xRy.$$

To see this, suppose a non-transitive relation R and feature F do make $(*)$ true for variable x and y . Then by R 's non-transitivity there can be a , b , and c so that a stands in R to b , b stands in R to c , but a does not stand in R to c . Since a stands in R to b , the right-to-left direction of $(*)$ entails that a and b each have feature F . Since b stands in R to c , that direction of $(*)$ entails that b and c each have feature F . It follows that a and c both have feature F . The left-to-right direction of $(*)$ then entails that a

¹³ It is no part of my task here to spell out how this could plausibly be done in detail. As will become clear in section 7, I do not think it could; but the worries that unfold as Reflective Disjunctivism is spelled out in detail do not turn on the phenomenon of diminished epistemic capacity.

¹⁴ Recall that we're restricting our attention to Williamsonsque versions of Reflective Disjunctivism. My (2006) looks at the putative non-transitivity of indiscriminability with reference to another version of Reflective Disjunctivism. The outcome of that discussion is somewhat different than it shall be here, though it turns out the worry of this subsection can be surmounted by either version of Reflective Disjunctivism.

stands in relation \mathbf{R} to c . This contradicts our assumptions about \mathbf{R} and \mathbf{F} , so those assumptions must be wrong. No non-transitive relation \mathbf{R} and feature \mathbf{F} can make (*) true for variable x and y .

The rub. Reflective Disjunctivism uses indiscriminability to build a conception of visual experience. The general fact constrains how this might be done. To see how, recall the view's basic take on the extension of visual experience:

$$(8) \quad \text{VE}^\Phi(x) \leftrightarrow x\text{IND}_g^\Phi.$$

Put in English: an episode is a visual experience as of Φ exactly when it is indiscriminable from Good Φ -cases. This claim looks a lot like (*) in our general fact; in particular, it looks to have this form:

$$\mathbf{F}(x) \leftrightarrow x\mathbf{R}y.$$

That isn't quite the general fact's (*); but we can get to it with no difficulty.

After all, Reflective Disjunctivism sees the claim that x is indiscriminable from Good Φ -cases as meaning that x cannot be known to be distinct from those cases. The factive nature of knowledge ensures that Good Φ -cases trivially cannot be known to be distinct from themselves. As flagged before, then, the view sees claims of this form as conceptual truths:

$$(5) \quad \mathbf{G}^\Phi(x) \rightarrow x\text{IND}_g^\Phi.$$

The right-to-left direction of (8) then ensures that Good Φ -cases conceptually count as visual experiences as of Φ . But that fact joins with (8) to entail

$$(12) \quad [\text{VE}^\Phi(x) \ \& \ \text{VE}^\Phi(g^\Phi)] \leftrightarrow x\text{IND}_g^\Phi.$$

And this claim looks very much like our general fact's (*). Indeed: when \mathbf{F} is set equal to the feature of being a visual experience as of Φ , \mathbf{R} is set equal to a binary relation of indiscriminability, and ' $x\text{IND}_g^\Phi$ ' is read to involve that relation, (12) is revealed as a claim of (*)'s form. The general fact then has bite; for when (12) is so read, it ensures that (8)'s relation of indiscriminability is transitive. Yet the indiscriminability of episodes can look clearly non-transitive.

For instance, suppose you see a given white picket fence on Monday, Wednesday, and Friday. Each time you see it from the same point of view; and each time the fence and its surroundings are exactly the same. Save for this: the fence changes colour each day, shifting subtly from one shade of white to another. Suppose there are five shades of white involved. The fence has the first one on Monday, the second one on Tuesday, and so forth. It looks possible that your veridical perception on Monday is indiscriminable from your veridical perception on Wednesday, your veridical perception on Wednesday is indiscriminable from your veridical perception on Friday, yet your veridical perception on Monday is discriminable from your veridical perception on Friday. After all, the shades of white may render Monday–Wednesday differences and Wednesday–Friday differences vanishing, while making for just noticeable Monday–Friday differences in colour. If the shades turn out to be indiscriminable in this

way, however, it's extremely plausible that your Good cases of them are too.¹⁵ But that would mean

$$(13) \quad g^m \text{IND} g^w \ \& \ g^w \text{IND} g^f \ \& \ \neg(g^m \text{IND} g^f).$$

This contradicts our current reading of (8). After all, (13) ensures that its relation of indiscriminability is non-transitive. The general fact then entails that there can be no feature *F* which joins with that relation to make something of (*)'s form variably true. That is just what (12) asserts under its present reading—which is itself entailed by a certain reading of (8)—so the latter must be wrong. Reflective Disjunctivism cannot employ a non-transitive notion of indiscriminability.

It is fortunate, therefore, that (the Williamsonsque version of) Reflective Disjunctivism is best thought to employ a transitive notion of indiscriminability in the first place. Seeing why that is so will expose the truly radical nature of the view; and it will set up critical discussion to follow in the next section.

To begin, recall the notion of indiscriminability in play:

$$(\text{ind})_{\mathbb{W}} \quad u \text{IND} v \leftrightarrow \neg \Diamond_{\text{imp}} \text{KR}(u \neq v);$$

Put in English: episode *u* is indiscriminable from episode *v* exactly when it is not impersonally possible to know by reflection that *u* and *v* are distinct. Reflective Disjunctivism uses this notion to build a conception of visual experience as such. To understand that conception properly we must understand reflective knowledge in detail. But the details have not yet been given. In particular, we have not yet been told what one reflects *upon* when one achieves knowledge by reflection. This proves central to the radical nature of Reflective Disjunctivism's take on discriminability.

To see this, note that the most natural thought here can be shelved right away. Reflective Disjunctivism cannot say that reflective knowledge is got solely by reflection on the *visual character* of episodes. After all, the view is in the business of spelling out a theory of such character by appeal to indiscriminability. Since Reflective Disjunctivism builds a view of visual character from facts about indiscriminability, and builds a view of indiscriminability from facts about reflective knowledge, Reflective Disjunctivism cannot spell out such knowledge by appeal to reflection on visual character. That would put things back to front, yielding no theory at all about episodes with only visual character (i.e. hallucinations). For this reason, Reflective Disjunctivism must see reflective knowledge in a broad way, as something like knowledge got by reflection on one's epistemic context.

As we'll see in section 8, it is unclear whether this idea can be spelled out properly. What matters here, though, is simply that Reflective Disjunctivism is built on the idea that visual experience is itself made from reflection on one's overall epistemic context, and *not* by reflection on the visual character of one's episode. The view's notion of indiscriminability is

$$(\text{ind}_c)_{\mathbb{W}} \quad u \text{IND}_c v \leftrightarrow \neg \Diamond_{\text{imp}} \text{KR}_c(u \neq v).$$

¹⁵ For a contrary perspective on this point see Hellie (2005).

Put in English: episode u is indiscriminable in context c from episode v exactly when it is not impersonally possible to know by reflection on c that u and v are numerically distinct. This is the full-dress notion of indiscriminability used by (our Williamsonsque version of) Reflective Disjunctivism. It concerns a contextually indexed relation built from reflection on epistemic context.

This leads to a quick-and-dirty reaction to the worry of this subsection; for $(\text{ind}_c)_w$ makes explicit that Reflective Disjunctivism's notion of indiscriminability is at bottom—from a God's-eye metaphysical point of view, as it were—a *tertiary* relation between two episodes and a context. It is not at bottom a binary relation at all. For this reason, Reflective Disjunctivism can formally side-step trouble posed by our general fact. Once its contexts are made explicit, after all, the view's notion of indiscriminability is revealed not to be a binary relation on episodes. Rather—and intuitively put—it is revealed to be a binary relation on episodes relative to a context; and so Reflective Disjunctivism can side-step difficulty posed by the putative non-transitivity of a binary notion of indiscriminability.

But this response is too quick and too dirty. The general fact can be easily extended to binary relations which hold relative to contexts; and once that's done the problem of this subsection crops up as before. Reflective Disjunctivism is then forced to admit that its notion of indiscriminability is transitive within a context; and bringing that out reveals the view's truly radical nature.

Some terminology. Call a relation R *contextually transitive* when there can be no u , v , and w and context c so that u stands in R to v at c , v stands in R to w at c , but u does not stand in R to w at c . When a contextually indexed relation R is contextually transitive, u standing in R to v at c plus v standing in R to w at c guarantees that u stands in R to w at c . When a contextually indexed relation R permits there to be u , v , w , and context c so that u stands in R to v at c , v stands in R to w at c , but u does *not* stand in R to w at c , R is a contextually *non-transitive* relation. Only contextually indexed relations like this are in the business of being contextually transitive or non-transitive. Only they can be preserved at a context—or fail to be so—by iteration across pairs of things.

A new general fact. No contextually non-transitive relation R and feature F can be so that for variables x , y , and c :

$$(**) \quad [F(x) \ \& \ F(y)] \leftrightarrow [xRy \text{ at } c].$$

To see this, suppose a relation R and feature F do make $(**)$ true for variables x , y , and c . Then by R 's contextual non-transitivity there can be u , v , w , and c so that u stands in R to v at c , v stands in R to w at c , but u does not stand in R to w at c . Since u stands in R to w at c , the right-to-left direction of $(**)$ entails that u and v each have feature F . Since v stands in R to w at c , that direction of $(**)$ entails that v and w each have feature F . It follows that u and w both have feature F . The left-to-right direction of $(**)$ then entails that u stands in relation R to w at c . This contradicts our initial assumptions about R and F , so those assumptions must be wrong. No contextually non-transitive relation R and feature F can make $(**)$ true for variables x , y , and c .

The extra rub. Now things go as before, only this time the problem springs from a new general fact. We appeal to multiple-episode contexts which look to involve non-transitive indiscriminability across episodes within them. Borrowing an idea from Richard Heck,¹⁶ suppose you look at a pie diagram; and suppose the pie is divided into white sections, with the shades of white in play being just those of our picket-fence example (leading up to (13)). In the event, arguably, when you look at the pie you occupy a single epistemic context with multiple veridical episodes, one for each section of the pie. Those episodes look to be pair-wise indiscriminable when they are of adjacent slices of the pie, though you can discriminate episodes of non-adjacent slices. Your context c looks to contain good episodes g^1 , g^2 , and g^3 —corresponding to slices 1–3 of the pie—so that

$$(14) \quad g^1 \text{IND} g^2 \ \& \ g^2 \text{IND} g^3 \ \& \ \neg(g^1 \text{IND} g^3).$$

In our present notation:

$$(15) \quad \neg \diamond_{\text{imp}} \text{KR}_c(g^1 \neq g^2) \ \& \ \neg \diamond_{\text{imp}} \text{KR}_c(g^2 \neq g^3) \ \& \ \diamond_{\text{imp}} \text{KR}_c(g^1 \neq g^3).$$

And put in English: it is not impersonally possible to know by reflection on c that g^1 is distinct from g^2 , and it is not impersonally possible to know by reflection on c that g^2 is distinct from g^3 , but it is impersonally possible to know by reflection on c that g^1 is distinct from g^3 .

We have already seen, however, that Reflective Disjunctivism entails this kind of claim about an episode x and a Good Φ -case g^Φ :

$$(16) \quad [\text{VE}^\Phi(g^\Phi) \ \& \ \text{VE}^\Phi(x)] \leftrightarrow \neg \diamond_{\text{imp}} \text{KR}_c(x \neq g^\Phi).$$

This has the form of our contextually indexed (**); and that looks to show that Reflective Disjunctivism conflicts with the contextually indexed general fact.

But notice: (14) and (15) look plausible only if their indiscriminability is thought to involve *discrimination on the basis of character alone*. After all, reflection on your overall epistemic context will render it obvious that g^1 is not identical to g^2 , g^2 is not identical to g^3 , etc. That is because reflection on your overall epistemic context will involve reflection on more than the intrinsic character of episodes within it. Such reflection will also have access to extrinsic aspects of episodes—their temporal features and interrelations, their phenomenological interrelations in the visual field, and so forth. The only way g^1 could begin to look indiscriminable from g^2 is if reflection upon them abstracted from such information, paying attention only to the character of each episode.

That is precisely not the kind of reflection used by Reflective Disjunctivism. As we have seen, the view uses a more general kind of reflection, one which takes full epistemic contexts as input. This undermines the plausibility of key claims in the pie-diagram and picket-fence cases—i.e. (13), (14), and (15). After all, reflection on your full epistemic context in the former renders it clear that g^1 is not identical to g^2 , clear that g^2 is not identical to g^3 , and clear that g^1 is not identical to g^3 ; and reflection on

¹⁶ Who in turn offered it to Delia Graff: see footnote 11 of Graff (2001).

your full epistemic context in the latter—say on Wednesday—renders it clear that g^m is not identical to g^w , clear that g^w is not identical to g^f , and clear that g^m is not identical to g^f .

Reflective Disjunctivism is simply untroubled by the putative non-transitivity of indiscriminability. That non-transitivity turns on a reflection-on-character notion of indiscriminability. Reflective Disjunctivism uses a reflection-on-context notion of indiscriminability. Intuition may show that the former is non-transitive in the epistemic ideal; but it does not show—indeed it does not even suggest—that the latter is non-transitive there. Rather, epistemic reflection on contexts seems obviously to generate a transitive relation of indiscriminability on episodes; and so Reflective Disjunctivism is not touched by the worry of this subsection.

7 TWO WORRIES

Reflective Disjunctivism sees visual experience as the impersonal limit of reflective knowledge. In this section I probe that conception. First I present an obvious worry for the approach; then I present a less-obvious worry. The next section sketches a new version of disjunctivism, one which avoids these worries.

7.1 (a) The zombie problem

Here is Reflective Disjunctivism's take on the nature of visual experience:

$$(Nat) \quad \lambda x VE_c \Phi(x) = \lambda x \neg \diamond_{imp} KR_c(x \neq G^\Phi s).$$

Put in English: being a visual experience as of Φ in context c is identical to being not impersonally possibly known by reflection on c to be distinct from Φ -Good cases. This entails a striking claim, one flagged as such in our first gloss of Reflective Disjunctivism. In our present terminology:

$$(17) \quad \neg \diamond_{imp} KR_c(x \neq G^\Phi s) \rightarrow VE_c \Phi(x).$$

Put in English: if it's impersonally impossible to know by reflection on c that x is distinct from Φ -Good cases, then x is a c -bound visual experience as of Φ .

(17) strikes forcefully because it is meant to codify the idea that a purely epistemic condition is sufficient for a robustly phenomenal one. (17)'s antecedent concerns an idealized limit on knowledge got by reflection, while its consequent concerns a state with Φ -indexed visual phenomenology. How *could* the former be sufficient for the latter? What about physical-functional duplicates of us who lack consciousness? In the vernacular, what about *zombie twins*? Don't they make trouble for (17) and (Nat)? After all, their episodes look to satisfy (17)'s antecedent but not its consequent; their episodes look to satisfy (Nat)'s right-hand side but not its left-hand side. The mere possibility of zombies indicates that no cognitive condition can capture visual experience as such; for none can capture what it's like.

This is the zombie objection. Serious reaction to it will involve two commitments: one proceeds from indiscriminability to a certain kind of cognitive awareness

- (i) $\neg \diamond_{\text{imp}} \text{KR}_c(x \neq G^\Phi S) \rightarrow [\text{it seems to one in } c \text{ that } x \text{ is a Good } \Phi\text{-case}];$

the other proceeds from that kind of awareness to visual consciousness

- (ii) $[\text{it seems to one in } c \text{ that } x \text{ is a Good } \Phi\text{-case}] \rightarrow \text{VE}_c^\Phi(x).$

These steps forge an intellectual path *from* something cognitive-cum-epistemological—namely indiscriminability from Goodness—to something meant to be phenomenal—namely visual experience as of Φ .¹⁷ Both steps are highly non-trivial. If either is defended as conceptual truth, in fact, the defender's theory will have slipped from non-triviality to full implausibility.

After all, (ii)'s conceptual necessity looks to entail the conceptual impossibility of zombies. Their genuine possibility is a matter of serious debate within the profession. Their conceptual possibility is not. Zombies look conceptually coherent. They may be impossible, at the end of the day; but it does not look as if this is a conceptual truth. Zombies fail to satisfy the intended reading of (ii)'s consequent, by hypothesis; but they seem perfectly capable of satisfying the intended reading of (ii)'s antecedent. For this reason, the conceptual coherence of zombies looks to show that (ii) is not a conceptual truth at all. (ii) looks to be an implausible claim in need of serious defence.¹⁸

Moreover, there are strong reasons to think (i) is plain false; for its general form is

$$\neg \diamond_{\text{imp}} \text{KR}_c \neg \Psi \rightarrow [\text{it seems to one in } c \text{ that } \Psi].$$

The item known in the antecedent of this conditional is the negation of the embedded claim in its consequent. Shifting negations yields another expression of the idea

$$\neg \diamond_{\text{imp}} \text{KR}_c \Psi \rightarrow [\text{it seems to one in } c \text{ that } \neg \Psi].$$

Put in English: the first claim is that if it cannot impersonally be known that $\neg \Psi$ —by reflection on context c , of course—then in c it seems intellectually that Ψ ; and the second claim is that if it cannot impersonally be known that Ψ —again by reflection on c —then in c it seems intellectually that $\neg \Psi$. These claims amount to a single idea: when something cannot be ruled *out* by reflection on context it thereby seems intellectually to be ruled *in*.

That seems wrong for a number of reasons.

First, contexts lacking information about what you see make trouble for (i). I call them *mute* epistemic contexts. Suppose you are in one and do not know, for instance, whether you see a white picket fence. Idealized impersonal reflection on your context will not rule out that you see such a fence; but nor will it rule out that you do not see in that way. Whenever you're in a mute epistemic context of this sort, Reflective

¹⁷ In her contribution to this volume, Siegel diagnoses Martin's (2004) response to a zombie-like worry for his view in basically these terms.

¹⁸ I know of only one attempt. The latter stages of Martin (2006) sketch a conception of visual phenomenology—and our awareness of it—meant to entail something like (ii). I have not yet been able to understand the conception well enough to comment on it here.

Disjunctivism entails that it seems to you that you see a white picket fence; but the view also entails that it seems to you that it's not the case that you see such a fence. Whenever you're in a mute epistemic context of this sort, Reflective Disjunctivism entails that your intellectual perspective on things is explicitly contradictory. But that seems wrong. No contradictory perspective follows from being in an epistemic context which does not decide—because it does not speak—to whether your episode is a veridical perception of a white picket fence. Being in such a context does not cripple your intellectual perspective on things. It only limits reflective knowledge. That seems to show that intellectual perspective is not fixed by such knowledge, from which it follows that (i) is false.

Second, contexts with conflicting information about what you see make trouble for (i). I call them *mixed* epistemic contexts. Suppose you are in one and do not know, as a result, whether you see a white picket fence. Idealized impersonal reflection on your context will not rule out that you see such a fence; but nor will it rule out that you do not see in that way. Whenever you're in a mixed epistemic context of this sort, Reflective Disjunctivism entails that it seems to you that you see a white picket fence; but the view also entails that it seems to you that it's not the case that you see such a fence. Whenever you're in a mute epistemic context of this sort, Reflective Disjunctivism entails that your intellectual perspective on things is explicitly contradictory. But that too seems wrong. No contradictory perspective follows from being in an epistemic context which does not decide—because it speaks with a mixed voice—to whether your episode is a veridical perception of a white picket fence. Being in such a context does not cripple intellectual perspective on things. It only limits reflective knowledge. That seems to show that intellectual perspective is not fixed by such knowledge, from which it follows that (i) is false.

Third, contexts with unreliable belief-forming mechanisms, ill-designed belief forming mechanisms, causally deviant chains, evidence just out of sight, and all manner of external-to-perspective knowledge precluders make trouble for (i).¹⁹ I call these *corrupt* epistemic contexts. Intuitively, you can't know anything by reflection on them: their corrupt nature precludes knowledge. That fact and (i) jointly entail that when you're in a corrupt epistemic context, it seems from your perspective as if you see all kinds of things to be so, that your intellectual take on things is again explicitly contradictory. But that seems wrong as well. No contradictory perspective follows from being in an epistemic context which does not decide—due to corruption—whether your episode is a veridical perception of a white picket fence. Being in such a context does not cripple intellectual perspective on things. It only limits reflective knowledge. This is yet another reason to think that intellectual perspective is not fixed by such knowledge, from which it follows that (i) is false.²⁰

¹⁹ Arguably: like the veridical perception two days hence (or two slices over in a pie being seen) of a subtly different version of what one sees now—as in the picket-fence and pie examples of the last section.

²⁰ Martin (2004) defends the idea that in a certain range of cases, at least, whenever something cannot be ruled *out* by reflection on context it thereby seems intellectually to be ruled *in*. That defence is worrying; see footnote 13 of my 2006.

7.2 (b) The alignment problem

It does not seem that reflection on context can yield knowledge in line with the demands of Reflective Disjunctivism.²¹ To see why, note the view's main take on the extension of visual experience:

$$(Ext) \quad VE_c^\Phi(x) \leftrightarrow \neg \diamond_{imp} KR_c(x \neq G^\Phi S).$$

Put in English: an episode in c is a visual experience as of Φ exactly when it's not impersonally possible to know by reflection on c that the episode is distinct from Φ -Good cases. Negating each bit of (Ext)—and dropping a double negation—yields an equivalent idea:

$$(Ext)^* \quad \neg VE_c^\Phi(x) \leftrightarrow \diamond_{imp} KR_c(x \neq G^\Phi S).$$

Put in English: it's *not* the case that an episode in c is a visual experience as of Φ exactly when it *is* impersonally possible to know by reflection on c that the episode is distinct from Φ -Good cases.

But it's a datum of everyday life, recall, that we can sometimes know that we suffer visual illusion as of Φ ; and it is a datum of everyday life that we can sometimes know that we suffer visual hallucination as of Φ . We can sometimes know, for this reason, that our visual episodes are distinct from Φ -Good cases. Reflective Disjunctivists must say that such knowledge does not make for visual experience—it is not reflective, in their terms. But Reflective Disjunctivists must tell us what *does* make for such knowledge, what *does* go into reflection on context. This they have never done.

Further, it does not seem that any specification of detail will work out; for there are all kinds of intellectual routes to everyday knowledge of Badness. Here is a case drawn from real life:

During the summer of 2005 my wife and I used a big fan in our flat. We then stopped monitoring our toddler's sleep for the first time. The result was a distinctive pattern of auditory hallucination. Every night I would 'hear' Sascha, our daughter, crying; but after several trips down the hall—over several nights—the hallucinatory nature of the set up became known. Then I had everyday knowledge of Badness.

Reflection on information like this applies to visual episodes as well as auditory ones. It too can yield everyday knowledge of visual Badness. For this reason, Reflective Disjunctivism must bar any information which could be used in a normal route to knowledge got by reflection. But there are all manner of such routes available—situations well within the sphere of everyday possibility which permit the visually deceived to work out, by reflection on circumstance, that they are so deceived. This means Reflective Disjunctivism must say

- (iv) Background beliefs do not generally make for reflective knowledge.

²¹ The problem generalizes to other forms of Reflective Disjunctivism, as the main text is meant to make clear. See my (2006) for its application to another form of the view.

The information involved in background beliefs cannot be generally available to reflection on circumstance. Otherwise the possibility of everyday knowledge of Badness will slip through the net, count as knowledge got by reflection, and falsify the left-to-right direction of (Ext).

But now let h be an hallucination as of Ψ —say as of a white picket fence. Then for a huge number of Φ ,²² h is *not* a visual experience as of Φ . When Φ concerns portrayal of an empty green field, for instance, h is not a visual experience as of Φ ; for h is not a conscious portrayal as of an empty green field. When Φ concerns portrayal of a big blue sky, for instance, h is not a visual experience as of Φ ; for h is not a conscious portrayal as of a big blue sky. And so on. For a huge number of Φ , h is not a visual experience as of Φ since it is not a conscious portrayal as of Φ . The left-to-right direction of (Ext)* entails that for each of the Φ s one *can* know, by reflection alone, that h is not a Good Φ -case.

That is a huge amount of knowledge to be got solely by reflection. In layman's terms, one can know by reflection alone that one is not seeing an empty green field, that one is not seeing a big blue sky, that one is not seeing all kinds of other things. For each of these Φ s, Reflective Disjunctivism entails that one can know by reflection alone—and *not* by reflection on the visual character of h , recall—that one is not seeing Φ . The only way that can be true is if background beliefs are generally available to reflection on context. Otherwise one could not know all these things by reflection. Reflective Disjunctivism must also endorse

(iv) Background beliefs generally make for reflective knowledge.

(iii) and (iv) jointly entail that no source of knowledge can do what Reflective Disjunctivism obliges reflective knowledge to do: namely, square with the datum that ordinary background information can make for everyday knowledge of Badness yet secure the extensional adequacy of (Ext). There is a sharp misalignment between the datum and the view's take on the extension of visual experience.

8 A MODEST PROPOSAL

Reflective Disjunctivism accounts for Bad character with a non-partisan resource: reflective knowledge. Section 3 charted beefier versions of disjunctivism. They use partisan resources to deal with Bad character. Sometimes those resources involve mere appearance, other times sense-data, still other times intentionality; and so on. Every time the extra resource turns out to be partisan, something denied by many who work in the area. Reflective knowledge is in everyone's toolkit, however, and that means the ideology of Reflective Disjunctivism is pleasing when compared to that of its charted rivals. It is minimalist in a certain way; and for that reason Reflective Disjunctivism enjoys a dialectical advantage *ceteris paribus*.

²² Perhaps all Φ not equal to Ψ ; perhaps only Φ 's not entailed by Ψ ; perhaps something else again. We won't bother with this, as it turns on the exact way that visual experience portrays the world. We continue to abstract from that issue in this essay.

As we have seen, moreover, the view has internal resources to deal with worries based on diminished epistemic capacity and the putative non-transitivity of indiscriminability. Reflective Disjunctivism is a surprisingly resilient view. But it does face theoretical pressure: most importantly, its details need canvassing more effectively than they have been so far. No matter how that is done, however, it is hard to see how indiscriminability alone can make for an epistemic perspective sufficient for visual phenomenology; and it is hard to see how everyday knowledge of Badness can be reconciled with the approach's take on the extension of visual episodes.

Pressures like these call into question the viability of Reflective Disjunctivism. They raise serious doubt as to whether the view can be made to work. But what is the choice point at this stage of discussion? What should disjunctivists begin to contemplate if they are struck by the force of what's gone before? I close with a modest proposal about that.

It is natural to see the relevant choice point in the following way: one should either drop disjunctivism about visual experience, or plump for a beefier version which uses partisan resources to deal with Bad character. But that does not truly capture the stage of our dialectic; for there is a version of disjunctivism which avoids theoretical pressures faced by its reflective variant but also avoids partisan resources. The relevant choice point here is this: one should either drop disjunctivism about visual experience, plump for a beefier version, or go in for the view now to be sketched.

The view I have in mind hears the second backbone claim of disjunctivism in a novel way. Recall that claim:

(B) Bad character does *not* derive from bits of the physical world standing in an explanatorily basic relation to percepts.

To this point, we've heard (B) to involve narrow-scope negation. We've assumed—with common sense, of course—that there is such a thing as Bad character; and we've heard (B) state of it that it does not derive like Good character. Section 3 made clear that the main motivations for disjunctivism neither entail nor suggest that there is such a thing as Bad character. Epistemic ones turn on the putative (G)-like nature of Good character and its relation to knowledge; phenomenal ones turn on the putative (G)-like nature of Good character and its relation to phenomenology; and semantic ones turn on the putative (G)-like nature of Good character and its relation to thought. Each time the key issue is the nature of Good character. None of the main motives for disjunctivism has anything to do with Bad character. None even requires that there be such a thing. When it comes to disjunctivism and its motivations, therefore, Bad character doesn't come into it.

This makes room for a pure form of disjunctivism, one in line with the approach's main motivations. The form I have in mind—Pure Disjunctivism—avoids partisan resources while side-stepping theoretical pressures faced by its reflective cousin. The view accepts (G) and (B), of course—and thereby counts as a version of disjunctivism—but it accepts (B) because it denies that there is such thing as Bad character. Pure Disjunctivism rejects the very idea of delusive phenomenology.

This rejection entails that when one suffers an hallucination as of an empty green field, for instance, there is nothing it is like to enjoy the episode. And more generally:

Pure Disjunctivism says that the delusive side of experience brings no phenomenology with it at all. In saying this, the view disagrees with its reflective cousin, with beefier versions of disjunctivism, with common sense, and with every other school of thought about visual experience. This makes it initially look as if Pure Disjunctivism is a drastic view, a port of last resort for disjunctivists.²³ But that is not so. In many dialectical settings, in fact, Pure Disjunctivism is the theory of choice, the best-motivated view. I close by explaining why.

To begin, note that everyone should agree—including Pure Disjunctivists—that it *seems* to one introspectively as if there is something it is like to enjoy Bad experience. After all, everyone thinks there *is* something it is like to enjoy Good experience; everyone thinks such experience introspectively seems to be like something because it is like something; and everyone thinks Bad experience is introspectively indistinguishable from Good experience (in some sense). Since the latter introspectively seems to possess phenomenology—on everyone's view, to repeat, because it does possess phenomenology—and since the former introspectively seems to be like the latter—again by universal assent—all should agree that Bad experience introspectively seems to be like something. All should accept that Bad experience introspectively seems to possess phenomenology.

Pure Disjunctivism denies that it does, of course; but that denial coheres perfectly with disjunctivism's main motivations; for none of them suggest that there is such a thing as Bad character. They all turn exclusively on the nature of Good character, pivoting on the role such character plays in our theorizing about knowledge, phenomenology, or semantics. If we're sensitive to motivations for disjunctivism alone, therefore, we'll feel no need to accept that Bad experience is like anything at all, no need to accept Bad character. Our commitment to such character must come from some place other than our motive(s) for disjunctivism.

Where does it come from?

Introspection, of course. We have no reason to think that Bad cases are like anything—no reason to accept Bad character—save for the fact that they seem to be like something introspectively. When deciding whether a beefy version of disjunctivism is better motivated than its pure cousin, therefore, we must first reckon whether that beefy version is in greater tension with the deliverance of introspection than is that cousin. In turn that means we must do three things:

- (i) discern the level of tension between the deliverance of introspection and the beefy version of disjunctivism's story about experience;
- (ii) discern the level of tension between the deliverance of introspection and Pure Disjunctivism's story about experience;
- (iii) compare the levels of tension.

If the level of tension between introspection and Pure Disjunctivism does not exceed that between introspection and a beefy version of the approach, then Pure

²³ As one philosopher reacted when I first spoke of the view—someone whose work has been central to disjunctivism for decades, someone who shall remain namelessly Paul Snowdon—“what in the name of all that is holy are you *talking* about?!”

Disjunctivism will be *better* motivated than its beefier cousin. After all, our only reason to believe in Bad character is the deliverance of introspection. On present assumptions, Pure Disjunctivism coheres with that deliverance at least as well as its beefier cousin. The extra resources involved in the beef are thereby unmotivated.

On the other hand: if the tension between introspection and Pure Disjunctivism exceeds that between introspection and a beefy version of the approach, then Pure Disjunctivism may or may not be better motivated than its beefier cousin. It will all depend on whether the beefier view's extra resources are worth buying in order to save the introspective phenomena. If those resources are exotic by one's lights—say when the view appeals to immaterial things and one is a materialist—then one should *still* find Pure Disjunctivism better motivated than its beefier cousin. If the beefy view's resources are not too exotic by one's lights, however, then one should find Pure Disjunctivism less well motivated than its beefier cousin.

How do we discern the level of tension between a given version of disjunctivism and the deliverance of introspection? That is a complicated matter. It can be sketched without pain, though, by letting **D** be an arbitrary version of disjunctivism and **INT** be introspection. The tension between them will then be a function of two things: the level of detail found in **INT**'s story about visual experience, and the degree of difference **D** sees between Good and Bad cases. The tension will look something like this

$$\text{Tension}(\mathbf{D}, \mathbf{INT}) = f\{\text{detail}(\mathbf{INT}), \text{Good-Bad-diff}(\mathbf{D})\}.$$

To see how this works, consider a limit case. Suppose introspection says only that Good and Bad cases are like something; suppose it says nothing about what they are like. In the event, the level of detail in **INT**'s story about experience is very low indeed. It is so low, in fact, that every version of disjunctivism can square with it save for the pure version. The level of tension between **D**'s story about experience and **INT**'s story is nil, that is to say, unless **D** happens to be Pure Disjunctivism.²⁴

Now consider the other extreme: suppose introspection says a great deal about visual experience. To fix ideas, suppose it says experience is always (G)-like. This is to suppose that introspection says that Good and Bad cases involve basic conscious perceptual contact with the physical world, say because it says that facts in that world manifest themselves to us visually. In the event, **D**'s story about experience will agree with **INT** when it comes to Good cases. **D**'s story will conflict with **INT**, therefore, to just the extent that **D** sees a big difference between Good and Bad cases. If it sees very little difference—say because it claims that both Good and Bad cases involve a single metaphysical structure filled out mostly the same way—then there will be very little conflict between the deliverance of introspection and **D**'s story about experience. If **D** sees a big difference between Good and Bad cases, however, there will be real tension between **D** and **INT**.

²⁴ Even then it seems that the conflict is of little theoretical moment. In these circumstances to save the introspective phenomena is to do no more than to save the claim that Bad cases are like something or other. Without that bringing with it some idea concerning what they are like, however, it's hard to see what difference it makes if we give up that claim and maintain instead that introspection goes wrong when it comes to Bad cases.

We can say in rough outline, therefore, that the level of tension that exists between the deliverance of introspection and a particular version of disjunctivism is itself a function of two things: first, the level of detail found in the introspection-based story about experience; and second, the degree of difference acknowledged between Good and Bad cases on the version of disjunctivism in question.²⁵ Moreover, the relevant function can be roughly characterized as follows: as the level of detail goes up in the deliverance of introspection, the potential for tension between it and disjunctivism goes up as well; and that tension will increase in proportion to the level of difference that is said to exist between Good and Bad cases.

Pure Disjunctivism is no port of last resort. It is not even a *disjunctivism* of last resort. Whether the view is motivated—say in relation to other versions of the approach—will itself turn on a number of things. It will depend on one's take concerning the deliverance of introspection. It will depend on the degree of difference acknowledged between Good and Bad cases on the impure version of disjunctivism being considered. It will also depend on one's background commitments and their relation to extra resources found in the impure view's ideology. Pure Disjunctivism does not always lose out in the dialectical competition. There are plenty of dialectical settings in which it is the best-motivated version of disjunctivism. There may even be some in which it is the best motivated view of visual experience full stop.

Pure Disjunctivism accepts that there are visually delusive episodes. It just denies that their delusive nature is like anything. The view needn't contain any commitment about the extension or nature of visual experience. It need only involve a take on what that experience is like, claiming that it's (G)-like through and through. For this reason, the view can avoid partisan resources needed to deal with Bad character; it can avoid theoretical pressures faced by its reflective cousin—as it does no serious work with the notion of indiscriminability—and the view can stay faithful to core motivations for disjunctivism.

There are several ways to spell out the approach. At bottom they all defend a disjunctivism about visual episodes, splitting their class into two groups. One of them is said to contain episodes with visual character; the other is said to contain episodes with no such character. Some versions of the approach say that visual experience as such has a nature specifiable in other terms—it is compatible with Pure Disjunctivism as sketched, after all, that an episode is a visual experience as of Φ exactly when it plays a certain functional role. But other versions of the approach say that visual experience as such has no nature at all—it is compatible with the picture as sketched, after all, that visual experience makes for a hodgepodge class of event tied together by nothing of theoretical moment.

To use the terminology with which we began: Pure Disjunctivism agrees with common sense that Bad cases lack perceptual sides; and it agrees with common sense that introspection misses this fact; Pure Disjunctivism continues that Bad cases lack conscious portrayal sides too; and it claims that introspection misses that fact as well. The

²⁵ I used to think it obvious that introspective reflection on visual phenomenology leads directly to non-trivial commitments about visual experience. Maja Spener's work has disabused me of that naiveté. See Spener (2003) and (forthcoming).

view sees Good and Bad cases as polar opposites, each at the end of a range of possibilities concerning how an episode's portrayal of the world might spring from its perceptual contact with it.²⁶

REFERENCES

- Alston, W. (1999) 'Back to the Theory of Appearing', *Philosophical Perspectives*, 13: 181–203.
- Byrne, A. and Logue, H. (this volume) 'Either/Or'.
- Chisholm, R. (1965) 'The Theory of Appearing', in R. Schwarz (ed.) *Perceiving, Sensing and Knowing*, Berkeley: University of California Press.
- Crane, T. (2001) *Elements of Mind*, Oxford: Oxford University Press.
- Graff, D. (2001) 'Phenomenal Continua and the Sorites', *Mind*, 110: 905–36.
- Hawthorne, J. and Kovakovich, K. (2006). 'Disjunctivism', *Proceedings of the Aristotelian Society*, 106: 185–216.
- Hellie, B. (2005) 'Noise and Perceptual Indiscriminability', *Mind*, 114: 481–508.
- Langsam, H. (1997) 'The Theory of Appearing Defended', *Philosophical Studies*, 87: 33–59.
- McDowell, J. (1982) 'Criteria, Defeasibility and Knowledge', *Proceedings of the British Academy*, 68: 455–79.
- (1994) *Mind and World*, Cambridge, MA: Harvard University Press.
- (1998) 'Having the World in View: Sellars, Kant and Intentionality', *Journal of Philosophy*, 95: 431–91.
- Martin, M. G. F. (2002) 'The Transparency of Experience', *Mind and Language*, 17: 376–425.
- (2004) 'The Limits of Self-Awareness', *Philosophical Studies*, 120: 37–89.
- (2006) 'On Being Alienated', in T. Szabó Gendler and J. Hawthorne (eds.) *Perceptual Experience*, Oxford: Oxford University Press.
- Putnam, H. (1999) *The Threefold Chord*, New York: Columbia University Press.
- Siegel, S. (this volume) 'The Epistemic Conception of Hallucination'.
- Spener, S. (2003) *Gilding or Staining the Mind: Introspection and the Metaphysics of Visual Phenomenology*, PhD Thesis: King's College London.
- (forthcoming) 'Mind-Independence and the Spatial Character of Visual Experience'.
- Sturgeon, S. (2001) *Matters of Mind*, London: Routledge.
- (2006) 'Reflective Disjunctivism', *Proceedings of the Aristotelian Society*, 106: 185–216.
- Thornton, T. (2004) *John McDowell*, Chesham: Acumen Press.
- Williamson, T. (1990) *Identity and Discrimination*, Oxford: Basil Blackwell.

²⁶ One initially bitter aspect of the view—by my lights anyway—is that it is committed to the possibility of visual episodes which are introspect-able but unlike anything. One might reasonably query the coherence of that. Matt Soteriou did in discussion; but Maja Spener reminded me that we do not worry much about the idea of non-occurrent beliefs being introspectable. There's certainly room for debate, then, about the bitterness of the possibility in question here. For further discussion see the last section of my (2006).

5

Disjunctivism, Indistinguishability, and the Nature of Hallucination

William Fish

In the eyes of some of its critics, disjunctivism fails to support adequately the key claim that a particular hallucination might be indistinguishable from a certain kind of veridical perception despite the two states having nothing other than this in common. Scott Sturgeon, for example, has complained that disjunctivism “offers no positive story about hallucination at all” (2000: 11) and therefore “simply takes [indistinguishability] for granted” (2000: 12). So according to Sturgeon, what the disjunctivist needs to provide is a plausible explanation of just how two mental states which have no common component might be indistinguishable for their subject and this in turn will require the telling of a positive story about hallucination. This is the goal of the present essay.

1

Our datum then is that some mental states that are not veridical perceptions—hallucinations—can nonetheless be indistinguishable from veridical perceptions. The evidence for this lies in the annals of literature, psychopathology, and in some cases even personal experience. There we find numerous cases in which subjects have

I would like to thank audiences from the Institut Jean Nicod, the ANU, and the universities of Nottingham, Victoria, and Waikato for discussions of previous incarnations of this essay. Special thanks go to David Chalmers, Pascal Engel, and Susanna Siegel, as well as the editors and readers of the present volume, for their insightful and helpful comments and criticisms. I would also like to take this opportunity to note that, although Scott Sturgeon appears in the present essay as the critic the disjunctivist is aiming to satisfy, in more recent work his opposition to disjunctivism appears to be waning. A feature of particular interest is that, in his contribution to the present volume, Sturgeon (this volume) indicates the availability of a position that bears many similarities to the one developed here, a position he calls “Pure Disjunctivism”.

sincerely taken themselves to be having an experience just like that of seeing something (we will focus only on the visual case) in situations where the thing apparently encountered was not there to be seen. In many of these cases, the subject of the anomalous experience may not suspect that her experience is anything other than veridical and will thereby come to believe that things really are as they seem to be. Purely for the purposes of marking a distinction, I will call hallucinations which result in belief in this way 'true' hallucinations. Yet it is not always the case that subjects are taken in by their hallucinations; there is plenty of evidence that there can be situations in which the subject undergoing the hallucinatory experience is aware of its non-veridicality and I will call those hallucinations which are recognized as such 'resisted hallucinations'. Now a fully adequate disjunctive analysis of hallucination will need to account for both true and resisted hallucinations and the theory to be presented here is intended to do so. However, because of the differences between the two kinds of hallucination, there will also be important differences between the accounts offered of each. So with this in mind, I will begin by discussing and offering an account of the paradigmatic cases of true hallucination before extending this analysis to resisted hallucination in section 10.

Before we get on to developing the account of true hallucination, we need to first say something about the kind of mental state hallucinations are indistinguishable from—veridical perceptual experiences. Thankfully, we can do without a substantial analysis of veridical experience for present purposes, and work with the minimal (and hopefully untendentious) claim that, whatever else we may ultimately want to say about them, veridical perceptual experiences are mental states/events which have phenomenal (sensory/qualitative) character, and that it is in virtue of an experience's having phenomenal character that there is something it is like to be in that mental state. When it comes to substantive explanatory questions—such as questions about whether, in order to account for the phenomenal character a veridical experience has, we need to appeal to intrinsic, non-representational qualia, representational content, a combination of these, or something else entirely—we can for the time being remain silent. I will assume that there is a positive, realist explanatory story to tell about the phenomenal character of veridical perceptual experience—given the motivations for disjunctivism, the disjunctivist is likely to want to offer some kind of naïve realist account—but the detail need not concern us here. All we need to focus on for the current project is the idea that veridical perceptual experiences have a phenomenal character, and that it is in virtue of this that there is something it is like to perceive veridically.

With this minimal claim in hand, we can see that, as hallucinations are mental states that are indistinguishable from veridical perceptual experiences, hallucinations are mental states that are indistinguishable from states with phenomenal character. The challenge for the disjunctivist is therefore to explain this indistinguishability without positing some common component shared by both hallucinations and veridical perceptions.

In order to do this, we need to begin by getting clear about what it is for two things to be indistinguishable in order that we can thereby focus precisely on what it is that needs to be explained. In his recent book on the epistemology of identity, Timothy Williamson defines indiscriminability—a notion I treat as equivalent to indistinguishability—as follows:

a is indiscriminable from *b* for a subject at a time if and only if at that time the subject is not able to discriminate between *a* and *b*, that is, if and only if at that time the subject is not able to activate (acquire or employ) the relevant kind of knowledge that *a* and *b* are distinct. (1990: 8)

I will say more about the definiens of this definition in due course, but for now, I just want to focus on a feature of the definiendum: Williamson is explicitly defining a notion of indiscriminability which is doubly relative—it is relative to a subject and it is relative to a time. The subject relativity is there to account for the fact that you may be able to distinguish between two things where I cannot—perhaps because you have sharper eyesight than I do. The temporal relativity is there to account for the possibility that I may be unable to distinguish between two things at one time, yet be able to distinguish between them at a different time. For example, if I am looking at two people at a distance or in half-light, I may well be unable to make a discrimination. According to this definition, those individuals are therefore indistinguishable (for me, at that time). But bring them closer to me, or turn up the lights, and I may then be able to tell one from the other. So at this later time when the conditions under which I am attempting the discrimination have improved, the individuals are now distinguishable (again, for me, at that time).

With these sources of relativity made explicit, we can see that there are different factors which can go to determine whether or not two things are found to be indistinguishable. The most conspicuous factor lies, of course, in the similarity of the things to be distinguished: as a general rule, the more similar two things are, the more difficult they will be to tell apart; indeed, in the limiting case, two qualitatively identical objects will simply be impossible to tell apart. Yet in addition to the similarity of the two objects, the double relativity of the notion of indistinguishability makes it clear that there are a range of other relevant factors relating to what we might loosely call the ‘discriminatory context’—the context under which the discrimination is attempted—which in everyday cases would include such things as the subject’s discriminatory capacities and the observation conditions under which the discrimination is attempted. When the contribution of the discriminatory context is taken into account, we can see that it would be possible for two things, intrinsically quite distinct, nonetheless to be indistinguishable for a subject at a time.

Now when we are faced with the datum that hallucinations can be indistinguishable from veridical perceptions, a characteristic initial reaction is to think that this indistinguishability must be due to the similarity of the two states: that hallucinations must therefore be a class of mental states which have a rich and detailed phenomenal

character that is similar enough to the phenomenal character of a veridical perception to mislead the experiencing subject. Yet as Williamson's definition shows, the very fact of indistinguishability alone does not entail such a conception of the hallucinatory mental state. As is true of more everyday attempts at discrimination, if the contribution of the discriminatory context is taken into account, it is conceivable that one may fail to distinguish between two things which are intrinsically quite distinct. Indeed, this suggests that the very claim that *hallucinations* are indistinguishable from veridical perceptions may be a subtly, yet importantly, leading formulation of our datum. This way of putting the point implies that there is a mental state, antecedently identifiable as a hallucination (in virtue of having perception-like phenomenal character), which is liable to be misidentified as a veridical perception. Yet once we take note of the potential contribution of some analogue of the discriminatory context, we can see that the very fact of indistinguishability does not require that a mental state be identified as a hallucination *prior* to its being misidentified as a veridical perception. Instead, it would be consistent with the very fact of indistinguishability for the class of hallucinations to comprise a collection of disparate mental states that are united only by the fact that, for some reason, the subject of those mental states finds them to be indistinguishable from veridical perceptions of certain kinds at certain times. If this were the case, these mental states would not be antecedently identifiable as hallucinations, but would rather only *qualify* as hallucinations inasmuch as the subject finds them indistinguishable from veridical perceptions.

Of course, this does not show that our initial intuition that hallucinations are phenomenally similar to veridical perceptions is *false*, merely that we do not have to assume its truth in order to account for the very fact of indistinguishability. This suggests that, when it comes to providing the analysis of hallucination, there are at least two options open to the disjunctivist. One would be to offer a broadly similarity-based account of hallucinatory indistinguishability: to try and show how hallucinations might be phenomenally similar to veridical perceptions despite having nothing in common. The other would be to offer an account of hallucinatory indistinguishability that does not assume that hallucinations are phenomenally similar, but which instead shows how a mental state, which may be intrinsically quite different from a veridical perception, might come to be mistaken for such an experience and therein acquire the status of hallucination. In this essay, I will examine the possibilities for developing an account of hallucination in accordance with the latter of these alternatives.

3

Some philosophers would argue that this approach is misguided, and that an adequate philosophical analysis of hallucination must be of the former kind. A. D. Smith, for example, insists that as far as the philosophical problems of perception are concerned, what is important is "that there *could possibly be* genuine sensory experiences of the same intrinsic, qualitative character as veridical perceptions, but which are not veridical" (2002: 202 emphasis in original). Now Smith goes on to argue that, once this

concession has been extracted, the disjunctivist will be sunk, but this just highlights the fact that the disjunctivist ought to ask why we should grant that such ‘phenomenally identical’ hallucinations are indeed possible in the first place. And when Smith does offer a defence of this claim of possibility, it runs as follows: “it is surely not open to serious question that [the principle of same proximate cause, same effect] does apply with respect to the merely sensory character of conscious states” (2002: 202). Inasmuch as there is an argument here, the case for the bare possibility of phenomenally identical hallucinations turns on a particular theory of the metaphysical source or explanation of the phenomenal character of experience. To elaborate, we can see Smith as defending the possibility in principle of such hallucinations—non-veridical mental states with the same phenomenal character as veridical perceptions—by arguing (1) that it is at least theoretically possible to replicate precisely the neural conditions which occur when a subject has a veridical perceptual experience in the absence of relevant objects, and (2) that “it is surely not open to serious question” that in replicating these neural conditions, you would thereby produce a mental state with the same phenomenal character as that veridical perception.

Yet despite the fact that many of us find it very natural to suppose that something like (2) is true, it would be philosophically irresponsible not to subject the claim to “serious question”. And interestingly, when we do inquire why we are inclined to think that neural replication alone would suffice to replicate the kind of full-blown phenomenal character possessed by veridical perceptions, we discover that things are not all that straightforward after all. For example, one central motivation for thinking that (2) is true is that *actual* hallucinations seem, on the face of it, to be states with perception-like phenomenal character whose origination lies solely in the brain (for example, Broad 1951: 39 and Robinson 1994: 152). But this is effectively an appeal to the very fact of indistinguishability and, as we have just seen, this does not entail that hallucinations are states with perception-like phenomenal character. Indeed, the analysis of hallucination to be presented in this chapter will explain why we might think hallucinations are like this, even though they are not. So with the dialectic poised as it is, such an appeal can carry no weight as an argument for a claim like (2).

Another defence of (2) appeals to empirical results from neuroscience but, as Alva Noë forcefully points out (2004: 7.2), whilst something like (2) may indeed function as a background assumption in much contemporary cognitive science, it does so despite a notable lack of any compelling empirical evidence that anything resembling full-blown experience has been produced artificially. Take, for example, the widely cited findings of Wilder Penfield (for example Robinson 1994: 247 and Lowe 1996: 103). While operating on the brains of epileptic patients, Penfield discovered that if the part of the cortex which lies between the auditory sensory and visual sensory areas is directly stimulated with an electric probe, subjects will occasionally report detailed and even multi-modal experiences (subject reports collated in Penfield and Perot 1963). Yet whilst Penfield’s findings do indeed suggest that direct electrical stimulation of the brain can induce experiences in subjects, it is far from clear that they offer much support to the claim that *phenomenally identical* experiences can be produced by direct stimulation of the brain. When Mahl et al. revisited Penfield’s transcripts of his patients’ responses, they concluded that the experiences he induced

are “essentially new creations based on memories, analogous to dreams” (1964: 357). This conclusion is supported by further results from Horowitz et al. who conclude that the “relationship between the prestimulation mental content and the sensory experience resemble[s] the process of dream construction” (1968: 471; see also Halgren 1982: 251).

So whilst Penfield’s findings do show that some kind of mental state can be produced by direct electrical stimulation of the brain, the results from Mahl et al. and Horowitz et al. give us reason to be sceptical that what is produced is anything more than a regular dream or hallucination. So this defence of (2) can be seen to be little more than an appeal to the very fact that (artificially induced) hallucinations are possible. But once again, an appeal to the very fact of indistinguishability can carry no weight at this point in the dialectic. As things stand, then, I suggest the disjunctivist need not assume that the theoretical possibility of precise neural replication alone is enough to establish the possibility in principle of a special kind of phenomenally identical hallucination. And if we do not, then the way remains open for the disjunctivist to offer instead an explanation of the very fact of indistinguishability that does not appeal to the similarity of the two states.

4

We now have some idea of the nature of the task before us. However, we are not yet ready to leave the issue of indistinguishability behind. Before we do, there are two other unusual features of our use of the term ‘indistinguishable’ in this context which need to be discussed. First is that, where our everyday concept of discrimination is concerned, discrimination is essentially a comparative act which requires the discriminator to consider two things—which may be presented either simultaneously or consecutively—and reach a judgement about whether they are the same or different. However, when we endorse the claim that a hallucinatory episode is indistinguishable from a veridical perceptual experience, we do not thereby intend to commit ourselves to the claim that a hallucinating subject will have actually had a veridical perceptual experience of the kind that the hallucination is claimed to be indistinguishable from. And even in cases where the subject can be taken to have had such an experience, we are certainly not intending to claim that she has this experience at the same time as the hallucination, nor that the experiences are had in anything like the kind of temporal proximity which might underpin an attempt at consecutive discrimination as we usually understand it.

This raises the question of exactly how we are to interpret the datum that hallucinations can be indistinguishable from veridical perceptions if we are not thereby intending to assert that there are two events—one perceptual and one hallucinatory—which the subject undergoes and then fails to tell apart. Now it might be tempting to read the claim counterfactually: as claiming that *if* the subject were to have both the perceptual and the hallucinatory experience (presumably in close temporal succession) *then* he would be unable to tell one from the other. But once again, if we bracket the assumption that neural replication will bring about phenomenally identical

experiences, it is not at all clear that this counterfactual would be true. When you consider a particular real-world case of a hallucination of a certain kind of object, what other reasons do we have for assuming that if *this* experience either followed, or was followed by, a veridical perception of an object of that very kind, the subject would not notice any difference? This isn't the case for dreams, for example. Even though dreams can seem quite real while one is dreaming, the vast majority of people have no difficulty in distinguishing between dream experiences and real experiences once they have woken up. So it is far from clear that we are entitled simply to take it for granted that the subject would always be unable to tell the difference between consecutive presentations of a hallucination and a veridical perception as the counterfactual reading of the indistinguishability claim would require.

If we do not read the indistinguishability claim counterfactually, then we need to say something about what it means to say that a hallucination is indistinguishable from a veridical perception when (1) there is no implication that the subject of the hallucination has actually *had* a perceptual experience of the kind the hallucination is claimed to be indistinguishable from, and (2) that there is no implication that, *if* the subject were to have (had) such a perceptual experience, she *would* find it indistinguishable from the hallucination she is actually having. A suggestive way of doing so is provided by Susanna Siegel, who suggests we might understand the indistinguishability claim in the following way: that purely on the basis of being in the hallucinatory state (in other words, ignoring any other relevant knowledge), the subject could not know that he or she was not having a veridical perception (2004: 98). This suggestion can be seen to be a development of Williamson's claim that two things are indistinguishable if and only if, at the time the discrimination is attempted, the subject cannot activate knowledge that the two things are distinct. Once we assume that the two states are not, in any sense, co-presented, it is a natural step to reinterpret the clause that the subject could not activate the knowledge that the two states are *distinct* as requiring that the subject could not activate knowledge that he or she was in a state of one kind rather than a state of the other kind.

This seems to be a promising way of interpreting the indistinguishability claim, but it leaves us squarely facing the second peculiarity in the current employment of the term 'indistinguishable'. This issue is that, as Williamson points out, discrimination is paradigmatically a "cognitive act" (5) which is "concerned with the way things seem to a pre-reflective observer" (13). So as far as our everyday use of the term goes, to say that two things are distinguishable or indistinguishable implies the existence of a *discriminator*: an agent who is engaged in the process of trying to tell the two things apart. Yet it is far from clear that the process of distinguishing between two mental states involves anything remotely like a "pre-reflective observer" engaged in the act of trying to tell them apart. So just as our use of indistinguishability in this context should not be taken to imply that the two mental states are experienced in succession, it also ought not be taken to imply that, in making the discrimination (whatever that may involve), we take up anything like an 'observational' stance towards our own mental states. Although there is a philosophical tradition according to which introspection involves an additional 'inner sense' (for example Locke (1690: II, I, 4), Armstrong (1968: ch. 15), and Lycan (1996: ch. 2)) which might therefore analyse this

kind of discrimination as involving something like a pre-reflective observer engaged in the task of discriminating between mental states, there are both significant philosophical concerns with such a claim (see, for example, the essays in parts I and IV of Shoemaker 1996), as well as alternative philosophical analyses of introspection around. Given this, it would be odd if, merely by agreeing to the claim that hallucinations and veridical perceptions can be indistinguishable, we were thereby taking a stand on this substantive philosophical issue.

So the claim that perception and hallucination may be indistinguishable should not be taken to imply that both a veridical perceptual experience and a hallucination must both be experienced in order for a discrimination between them to be attempted. We have seen that this can be avoided if we interpret the indistinguishability claim as requiring it to be the case that, solely on the basis of being in the hallucinatory state, the subject could not come to know that he or she was not having a veridical perception. Yet in explaining how this might be the case—how it might be the case that a subject who was not having a veridical perception might fail to realize this—we should not assume that the subject is in any sense ‘observing’ his or her own mental states. So how might we account for the subject’s failure to realize that he is not enjoying a veridical perception if we *don’t* think of indistinguishability in this way? Sturgeon offers a suggestive alternative. He claims that:

When introspecting what a visual state is like, we form judgments about the visual state directly on its basis. Introspection is one type of belief formation. The beliefs so formed are about visual states. But their formation does not spring from inner visual impressions of those visual states . . . Introspective beliefs about visual states spring directly from visual states themselves. To say two such states are Indistinguishable, then, is to say that they register equivalently in introspection. (2000: 12)

This approach to introspection and indistinguishability provides a plausible way in which we can make sense of the indistinguishability claim which does not require us to adopt an observation-based model of introspection. On this account, to say that two mental states are indistinguishable would be to say that they register equivalently in introspection, where this is understood as requiring the two states to generate equivalent introspective beliefs. Not only does this approach avoid appealing to an observational model of introspection (although it does not preclude such a model either), it also offers a neat explanation of how a hallucinating subject might fail to realize that he or she was not having a veridical perception. If the introspective beliefs a mental state were to generate were equivalent to the introspective beliefs a certain kind of veridical perception *would have generated*, then that mental state could be said to register equivalently in introspection and therein be indistinguishable from a veridical perception of that kind.

We need to say something more about what would be required for the introspective beliefs generated by a hallucination to qualify as *equivalent* to the introspective

beliefs that would have been generated by the associated veridical perception. Now, we might think that, to be equivalent, the introspective beliefs would have to be identical, but given that the mental state in question need not occur in close temporal proximity to the relevant kind of perception (indeed, the relevant kind of perception need not occur *at all*), then the introspective beliefs generated by the hallucination need not necessarily include each and every belief that the relevant perception would have generated, but would only need to be what we might call 'sufficiently similar'—similar in certain important respects. This does raise the question of what these important respects *are*—what *kind* of introspective beliefs does a hallucination need to yield to qualify as such?—but it is not clear that we can give a perfectly precise answer to this question. For one thing, there may be a number of subtly different ways in which a veridical experience, and in turn a hallucination, might seem 'real' (e.g. see Aggernaes 1972), which raises the possibility that a mental state might generate the same introspective beliefs as a certain kind of veridical perception in some, but not all of these dimensions, yet still do enough to qualify as a hallucination. What is more, the specific introspective beliefs a certain kind of *perception* might yield at any one time would be highly variable according to various other factors such as the subject's levels of attention and concentration, her current concerns and intentions, her background beliefs, and so on. But nevertheless, inasmuch as we are dealing with paradigmatic cases of hallucination here—situations in which a subject takes himself to be having a certain kind of veridical perception when he is not—the most important introspective beliefs will be those that ground this critical error.

To come at this from a different perspective, consider the kinds of effects a particular veridical perception would have. Say I veridically perceive a pink ice cube: what kind of effects would such a mental state have? Of course, for reasons given above, there would be a vast array of effects such a state might have—it might make me want a cold drink, or remember the last occasion on which I saw such an ice cube, or think of Wilfrid Sellars, and so on. But there would also be more standard effects than these. In particular, where subjects such as ourselves are concerned, someone who veridically perceives a pink ice cube would, unless something very curious was going on, acquire a range of fairly predictable beliefs. We would expect these beliefs to include, for example, the belief that a pink ice cube exists and the belief that she sees the pink ice cube. Now note that, if a mental state which was *not* a veridical perception of a pink ice cube somehow led the subject to form beliefs of this kind, then these beliefs could be appealed to in explaining why the subject takes herself to be having a veridical perception of a pink ice cube even though she is not. So long as I believe that I am seeing something, I will take myself to be enjoying a certain kind of veridical perception. So the thought is this: for subjects such as ourselves, a certain kind of veridical perception will have an array of cognitive effects. Among these cognitive effects will usually be higher-order beliefs about our own experiences. If a mental state which was not a veridical perception were somehow to come to have effects which included such higher-order beliefs, then this would look to be sufficient to explain why that mental state was mistaken for a veridical perception of that kind. If we therefore stipulate that the appearance of such higher-order beliefs is indicative of sufficient similarity in effects, and hence indicative of registering equivalently in

introspection, then given the account of indistinguishability developed earlier, this would ground the claim that this mental state is indistinguishable from that kind of veridical perception. The mental state would thereby qualify as a hallucination.

6

At this point, we can start to see a familiar piece of philosophical landscape hovering into view. In *Perception and the Physical World*, David Armstrong also emphasizes the fact that some of the beliefs we acquire in the hallucinatory case will be the same as those we acquire (or would have acquired) in veridical perception. To use one of his examples, suppose I walk into a room and see that there is a cat on the mat. What would I (normally) come to believe? Well, suggests Armstrong, I would normally come to believe both that there is a cat on the mat, and that I am seeing the cat (1961: 83). Now suppose I come into a room and hallucinate a cat on the mat. In such situations, Armstrong suggests, I would acquire a pair of beliefs of the very same kind—I would believe that there was a cat there, and I would believe that I see that there is a cat there—it is just that in this case, the beliefs would be false.

As Armstrong is keen to point out, if beliefs of this kind are common to the perceptual and hallucinatory cases, they can then be employed to explain why the hallucinating subject takes their hallucination to be a veridical perception. Here we can see Armstrong making the *Explanatory Claim*—the claim that we can explain everything we need to explain (in other words, everything a hallucinating subject thinks, does, and says) by appeal to the beliefs that the hallucinating subject holds. Now Armstrong's official line appeals to two kinds of false belief: false first-order beliefs about the world and false higher-order beliefs about experiences. But inasmuch as first-order beliefs are needed to explain the subject's erroneous claims about the world, they can be derived from the relevant higher-order beliefs anyway. If I believe that I see that there is a cat on the mat then I will, all other things being equal, believe that there is a cat on the mat. So in order to explain everything that needs explaining, the Explanatory Claim need in fact only appeal to the higher-order beliefs that a subject holds about his experiences. As long as I sincerely believe that I see that there is a cat on the mat before me, I thereby take myself to be seeing that there is a cat on the mat before me—I take myself to be having a veridical perceptual experience of a cat on the mat—even if I am not.

Now, on the face of it, beliefs of this kind look like common effects of both veridical perceptions (of cats on mats) and hallucinations of such. But in addition to the Explanatory Claim, Armstrong also endorses the *Constitutive Claim*: "The suggestion being put forward now is . . . that the occurrence of these two false beliefs constitutes [hallucination], that this is all that [hallucination] is" (1961: 83, emphasis in original). This yields Armstrong's official theory of hallucination, which (assuming the minor restriction to higher-order beliefs) consists of two claims: the Explanatory Claim, according to which we need appeal to no more than our (false) higher-order beliefs about our experiences in order to explain everything we need to explain, and the Constitutive Claim, according to which the state of hallucination is actually

constituted by the occurrence of these higher-order beliefs about experiences. These two claims mesh together as follows: the Constitutive Claim ensures that, whenever we have a hallucination, we will have higher-order beliefs that we see something. Then, by appealing to the higher-order beliefs guaranteed by the Constitutive Claim, the Explanatory Claim ensures that we can explain the hallucinating subject's reports and other behaviour.

Clearly, endorsing the Constitutive Claim would preclude us from viewing these higher-order beliefs as the registering of a hallucination in introspection: according to the Constitutive Claim, these beliefs *are* the hallucination, they are not the introspective consequences of the hallucination. Yet as critical discussions of Armstrong's theory have shown, there are reasons to be sceptical of the Constitutive Claim. A familiar objection to Armstrong's theory is that, in attempting to reduce hallucination to beliefs about experiences, we thereby over-intellectualize hallucinations. On one reading of this objection, we might see it as questioning whether hallucinating human subjects really do have these kinds of higher-order beliefs. However, I think that if this is how the objection is intended to be understood, it would be unproblematic for Armstrong. When such a subject hallucinates, we should always be able to reveal that she holds, or at least held, such a belief. To paraphrase Armstrong, this is shown by the fact that if I am asked why I am under the apprehension that there is a cat there, I would reply: "I can *see* it" (1961: 83). And whenever a subject reports the content and nature of her hallucination—reports that she sees something which isn't there—she simultaneously *expresses* her higher-order belief that she sees something (Rosenthal 1990: 747). For example, if I report that "I see that there is a cat in front of me", I thereby express the higher-order belief that I see that there is a cat in front of me. So any reports that hallucinating subjects provide about what they see are at the same time expressions of just the kinds of higher-order beliefs Armstrong's analysis appeals to.

If we read this objection as applying to those creatures which fundamentally lack the cognitive capacity for higher-order beliefs, however, it acquires a great deal more bite. Even given the obvious conceptual difficulties that attend this question, there is some fairly compelling empirical evidence that lower animals (such as rodents) can have hallucinations (for example, Corne and Pickering 1967 and Yamamoto and Ueki 1980), as well as strong intuitions that such creatures would lack the kind of cognitive sophistication required to have beliefs about their own experiences. Given that such animals are incapable of describing their experiences to us, the evidence for animal hallucinations will never be demonstrative, so it will always be possible to maintain the Constitutive Claim by simply denying that animal hallucinations occur. But if we do accept that some animals incapable of forming and entertaining higher-order beliefs can nevertheless suffer from hallucinations—and I suggest we should—then as the Constitutive Claim defines hallucination as no more than the occurrence of a suitable second-order belief, it must be rejected.

Yet if we reject the Constitutive Claim that higher-order beliefs just are the hallucination, the most natural position for someone broadly sympathetic to the Armstrongian project to draw back to is one on which, at least in the case of suitably sophisticated subjects, the higher-order beliefs still occur, it is just that they

are effects of the hallucination. This yields the position recommended earlier: that a hallucination is a mental state which, whilst not being a veridical perception of a particular kind, has effects which are sufficiently similar to those a veridical perception of that kind would have had. And of course, where suitably sophisticated subjects are concerned, to qualify as sufficiently similar these effects would need to include the higher-order belief that the subject was enjoying a veridical experience of that kind. Taken together, these two claims ensure, at least in those cases of true hallucination we are currently considering, that Armstrong's Explanatory Claim will always be available to be called upon when required, despite our having rejected his Constitutive Claim.

Now when it comes to a creature—a cat, for example—which lacks the conceptual sophistication required to entertain higher-order beliefs, a veridical perception of, say, a butterfly will still have certain kinds of effects. Of course, given that we are assuming that the cat is lacking in higher conceptual capacities, these effects will not include cognitive effects such as higher-order beliefs. Instead, they will be primarily behavioural, possibly including trying to strike the butterfly with its paw, and so on. But then our evidence for animal hallucinations is essentially a set of behavioural observations too. If we were to find a cat “striking at imaginary objects in the air” (Slade and Bentall 1990: 12)—in other words, behaving as though it perceived a butterfly when there was not one there—then we might have good reason to say that, in such a case, the animal is not perceiving a butterfly but is nevertheless in a mental state which has the same effects as such a perception. As the two states have sufficiently similar effects (given the complexity of the system in question), we can therefore allow that the two states are indistinguishable for the cat, and that the animal is having a hallucination of a butterfly (or other flying insect). Of course, what the animal *lacks* is a belief that it is actually *seeing* a butterfly, but that should be expected as we are explicitly dealing with those creatures which lack the conceptual sophistication required to entertain such beliefs.

7

The aim of this chapter is to outline a positive theory of hallucination in order to explain how two mental states which share no common component can be indistinguishable for their subject. How do things stand with respect to this goal? Well, the tentative analysis of what is going on in cases of true hallucination is this: hallucinations occur when a mental state that is not a veridical perception of a certain kind comes to have effects that are sufficiently similar to those that that kind of veridical perception would have had. The mental state will therefore qualify as indistinguishable from such a veridical perception which in turn elevates the mental state to the status of hallucination. Because sufficiently similar effects are defined as including, in a subject capable of such cognitive states, higher-order beliefs that he is having a certain kind of veridical perception, the subject will therefore take himself to be having a perceptual experience of that kind, which in turn enables us to explain everything that the hallucinating subject thinks, does, and says.

It would be fair to say that the positive story about hallucination embodied in these claims is rather minimal. But this is a consequence of the nature of the dialectic, not a result of quietist tendencies on the disjunctivist's part. Any mental state, whatever its intrinsic nature, that has effects that are sufficiently similar to those a veridical perception would have had will thereby qualify as (a) indistinguishable from that veridical perception and (b) as a hallucination. So the essence of hallucination—what distinguishes hallucinations as a class from other mental states—lies in their being indistinguishable from veridical perceptions, not in some antecedently identifiable feature of the state. This is why, “when it comes to a mental characterisation of the hallucinatory experience, nothing more can be said than the relational and epistemological claim that it is indiscriminable from the perception” (Martin 2004: 72). Of course, we can say more about what it is for a hallucination to be indistinguishable from a veridical perception: we can say that the indistinguishability of the two states is a matter of the hallucination having sufficiently similar effects to those that the veridical perception would have had. But the demand for a more substantial intrinsic characterization of the hallucinatory mental state is misguided. For a similar reason, it would also be misguided to demand any further explanation of why hallucinations, considered as a mental kind, have similar effects to veridical perceptions. As a mental state only qualifies as a hallucination inasmuch as it has the same kinds of effects as a certain kind of veridical perception, there is therefore something wrong with asking why *hallucinations* have these kinds of effects—it would be akin to asking what bachelors have in common in virtue of which none of them are married. Hallucinations are simply defined as states which have the same kinds of effects as, and are therefore indistinguishable from, certain kinds of veridical perceptions.

Yet whilst this positive story about hallucination is indeed somewhat minimal, it is a positive story nonetheless. And what is more, it is sufficient to explain everything we need our account of hallucination to explain, at least for those cases we are currently considering in which the subject is taken in. A hallucination is a mental state which is indistinguishable from a certain kind of veridical perception, and this indistinguishability is a matter of the state's having effects that are sufficiently similar to those that veridical perceptions of that kind would have had. Where subjects capable of higher-order mental states are concerned, these effects will include the forming of beliefs about their experiences. We can therefore be sure that whenever we find a sophisticated subject undergoing a (true) hallucination, we will find the kinds of higher-order beliefs which Armstrong's Explanatory Claim appeals to in order to explain everything that subject thinks, does, and says. But inasmuch as this story about hallucination appeals to higher-order beliefs, they are appealed to solely to explain a sophisticated subject's claims about her hallucinatory experiences—to explain why the subject claims to see something which is not there—not to provide a reductive base for the hallucinatory state. Inasmuch as we have a mentalistic *analysis* of hallucination—a definition of the hallucinatory mental state—it is merely as a mental state which is not a veridical perception of a certain kind, but which is, at a particular time, indistinguishable from a veridical perception of that kind.

Although it would be inappropriate to ask why hallucinations in general have the same kinds of effects as certain kinds of veridical perceptions, there is a reasonable question why, on a certain occasion, a particular mental state comes to have the same kinds of effects as a certain kind of veridical perception. Whilst there is not space to address this question in the detail it deserves here, I will offer a sketch of a possible explanation that is based on current empirical work on the topic. In a recent paper, Collerton et al. suggest that “within scene perception [i.e. against the background of veridical perceptual experience of a scene], a hallucination is experienced when . . . the visual system is constrained by a combination of impaired attentional binding and . . . a relatively intact scene representation biases perception towards an incorrect image” (2005: 749). In our terms, the suggestion is that the explanation of why a particular mental state has anomalous effects has two components: one that appeals to some kind of cognitive impairment (roughly corresponding to the discriminatory context) and another that appeals to some component of the experience which biases the subject to make the error.

As to the nature of the cognitive impairment, recent psychological research suggests that hallucination-prone individuals (as determined by high scores on the Launey-Slade Hallucination Scale (LSHS)) may be deficient in the important meta-cognitive skill of “reality discrimination” (see Bentall (1990: 88), Slade and Bentall (1990: 125), and Rankin and O’Carroll (1995: 518)). Reality discrimination is characterized as the ability we have of telling mental episodes that are internally generated apart from real veridical experiences. And as Rankin and O’Carroll point out, “a dramatic failure in an individual’s meta-cognitive skill of discriminating internally and externally generated events might cause a person to misattribute internal events to external sources and thus bring about hallucinations” (1995: 518). If this were right, then a non-perceptual mental state might come to have effects equivalent to those a certain kind of veridical perception would have had—and would thus qualify as a hallucination—when the subject’s reality discrimination capacity breaks down.

As such a cognitive disorder will most often not be permanent—reasons for a temporary disturbance may include such things as high fever, drug intake, the onset of a schizophrenic episode, or simply being half-asleep—it would likely be the case that, at those times when reality discrimination is functioning effectively, the subject would not have mis-identified the internally generated event and hence would not have suffered from a hallucination. So this explanation of the occurrence of hallucinations does not entail that a mental state that generates the same kind of introspective beliefs as a certain kind of veridical perception on a particular occasion would *always* generate such beliefs. It is compatible with the possibility of such a hallucination that, for the most part, the kind of mental state in question would not yield introspective beliefs that were equivalent to those *any* kind of veridical perception would have generated. However, if it were the case that, when cognitive disturbances caused the subject’s reality discrimination ability to break down, a certain mental state *did* come to generate introspective beliefs equivalent to those a certain kind of veridical perception

would have yielded then, for that subject at that time, the mental state would be indistinguishable from a veridical perception of that kind, and would thereby acquire the status of a hallucination.

Where the second, experiential, component is concerned, we can begin by noting that, in many cases, hallucinations actually occur against a background of veridical perceptual experiences of the world and indeed that in some cases, stopping the perceptual component of the experience (by shutting one's eyes) will actually make the hallucination cease (Manford and Anderman 1998: 1820). This suggests that the mental state which has the effects of a veridical perception of an x may itself be a perceptual experience of a different kind. Why would a perceptual state of one kind be mistaken for another? Well, in addition to the relevance of any cognitive impairments suffered by the subject, Collerton et al. also draw our attention to the possibility that there might be features of the experience that bias the subject to form false beliefs. For example, consider a situation in which a subject is faced with a suitably shaped shadow in the corner of the room yet, because of a short-lived cognitive disturbance leading to reality discrimination breakdown, comes to think that she sees a person in the corner of the room. The suggestion is that what happens here is that the shape and location of the shadow, in conjunction with the nature of the subject's cognitive disturbance, biases that subject to form the belief that she sees a person in the corner of the room. In such a case, the experience of the shadow comes to have the cognitive effects of a perception of a person in the corner of the room, and it thereby qualifies as indistinguishable from such a perception and, in turn, as a hallucination of a person in the corner of the room. So our explanation of *why* the mental state has these particular anomalous effects will have two components: one which appeals to the reality-discrimination breakdown and one which appeals to the particular shape and location of the shadow experienced.

In other cases of hallucination, however, the explanatory relevance of what in the world is seen can be much reduced. Horowitz discusses a case in which dragons and demons are hallucinated in response to flushing water in a toilet bowl (1964: 515). In such a case, an appeal to features of what is seen—to the patterns formed by the water in the bowl—doesn't really help to explain why the subject's experience of those features ended up having the particular effects it did. In order to explain why the experience had those specific anomalous effects, we can only really point to the nature of the subject's cognitive disorder, perhaps together with the contents of other of the subject's current mental states (see Mahl et al. (1964: 358) and Horowitz et al. (1968: 484)). Indeed, in the limiting cases, it is conceivable that the mental state that has the relevant effects need not antecedently be an experience at all. For example, Collerton et al. report that drug induced visual hallucinations can occur with the eyes closed (2005: 740). This suggests that the mental state that has the effects of a veridical perception of an x need not itself be a visual experience at all. Nevertheless, as long as it comes to have effects which are sufficiently similar to such a veridical perception, it will qualify as a hallucination of an x . And when asked why the mental state has the anomalous effects it does, in such a situation all we can usefully do is point to the subject's acute cognitive disorder, together with other features of the subject's mental make-up that explain why it is those particular beliefs that are formed.

This account of how a particular mental state comes to have the anomalous effects it does highlights what is probably the most unconventional aspect of our analysis. When we discussed the nature of veridical perception at the outset, it was claimed that the reason there is something it is like to have a veridical perceptual experience is explained by the fact that the mental state in question has phenomenal character. But in the limiting cases of hallucination we have just discussed, the mental state which is mistaken for a veridical perception and thereby qualifies as a hallucination may have no phenomenal character whatsoever. And even in those cases in which the mental state does have phenomenal character—as in the dragons and demons case—the phenomenal character it has may well not account for the particular content of the hallucination. Yet a hallucinator who possessed the relevant conceptual resources would probably *claim* that there was something it is like for them to hallucinate and that their hallucination has appropriate phenomenal character. What should we say about such cases?

The most straightforward approach would be to simply bite the bullet and claim that, although such subjects think/believe/judge that hallucinatory states have phenomenal character, they are wrong. This would be to adopt a circumscribed Dennettian irrealism about the phenomenal character of hallucinations (or at least, the hallucinatory elements of the phenomenal character). When talking about qualia, Dennett says that “I [seem] to be denying that there are *any* such properties, and for once what seems so *is* so. I *am* denying that there are any such properties. But . . . I agree wholeheartedly that there seem to be qualia” (1991: 372). And whilst the position being advocated here is far less radical than Dennett’s position—the disjunctive nature of the overall picture allows, and indeed I am assuming, a realist theory of the phenomenal character of veridical experiences—it is similarly irrealist about the phenomenal character of hallucinations: it denies that hallucinations have phenomenal character whilst agreeing wholeheartedly that they certainly seem to.

Interestingly, by limiting the irrealist aspects of the picture in this way, a number of objections faced by Dennett’s unrestricted irrealism can be either avoided or significantly weakened. Take, for example, the objections in (Tye 1993: 896) that Dennett’s theory is viciously circular, that it has the unwelcome consequence that one’s judgements about how things seem will come out as incorrigible, and that it precludes animals and children from being in states in which things seem a certain way to them. A version of Dennettian irrealism which is limited in the way suggested is not so adversely affected by these points: such a theory need not be circular as it leaves scope for there to be ‘real seemings’ in the veridical cases to provide the conceptual ingredients for the false judgements in hallucination; the suggestion that judgements about one’s own hallucinations possess some kind of incorrigibility is not subject to the kinds of worries that attend the corresponding claim in the perceptual case; and the circumscribed position not only allows for animals and children to be in states in which things seem a certain way to them (the veridical cases), it also allows that such conceptually unsophisticated creatures can suffer from hallucinations as we have seen.

Despite this, however, some people might think that the claim that hallucinations have no phenomenal character entails that there is nothing it is like to hallucinate and that this is just absurd. For example, one might think that it is simply a conceptual truth that, if a subject honestly judges that there is something it is like to be in a particular state, then there *is* something it is like to be in that state. So as the view presented here allows that subjects may indeed make such judgements, if it also entails that there isn't anything it is like for them when they do so, it will thereby be revealed as unsatisfactory.

Be that as it may, the position we have outlined can quite easily be adapted to accommodate the thought that it is conceptually necessary that an honest, first-person claim that there is something it is like to be in a particular mental state will be true. As discussed in section 1, in the veridical cases the truthmaker of such an honestly made claim will be the fact that the perceptual state itself has phenomenal character. The purported problem is that, if we accept both the limited irrealism and the conceptual thesis, then in the case of hallucination such honestly made claims will still be true (from the conceptual thesis), but the truthmaker of such claims cannot be found in the hallucinatory state's phenomenal character (from the irrealism). So if we do suppose that it is conceptually necessary that such claims will be true, we will need either to accept that the hallucinatory state does have phenomenal character after all—essentially, to reject this analysis of hallucinatory indistinguishability in favour of a broadly similarity-based analysis—or to provide an alternative truthmaker for something-it-is-like claims.

We can take the latter option if we note first that it is crucial that the first-person claim is honestly made: if the claim weren't honestly made—if the subject was lying about what it was like for them—then the guarantee of truth would not follow. So perhaps we can fix on what makes such claims true by looking at what makes it the case that they are honestly made. According to the analysis of indistinguishability we have given, a hallucination of an *x* will lead a suitably sophisticated subject to believe that they see an *x*. Whilst such a belief would of course be false, *having* such a false belief would nevertheless put the subject in a position to claim *honestly* that there is something it is like for them to be in that state. Having this kind of false belief guarantees that the subject will be in a position to make the honest claim that there is something it is like for them because, if you believe that you are seeing something then, as there is something it is like to see something, you will be disposed to judge honestly, and hence claim, that there is something it is like for you.

So if a subject enters a mental state which is not a veridical perception of an *x* but which, for whatever reason, she takes to be a veridical perception of an *x*, she will thereby be placed in a position to claim honestly that there is something it is like for her to be in that mental state. This suggests that the truthmaker of the claim in cases of hallucination is therefore the fact that the subject is in a mental state which is indistinguishable from a certain kind of veridical perception, as it is in virtue of this indistinguishability that the subject is thereby put into a position to claim honestly that there is something it is like for her. In order for this to work, however, it is crucial that there is something it is like to perceive veridically: it is only because there is

antecedently something it is like to see an x that merely believing that one sees an x suffices to put you in a position to make the claim honestly.

This does reverse the standard order of explanation—normally, the indistinguishability of perception and hallucination is held to derive from their similarity in terms of phenomenal character, whereas I am suggesting that the similarity in terms of what it is like to perceive an x and hallucinate an x derives from the indistinguishability of the two states—but it still allows us to rebuff the charge that there is nothing it is like to have a hallucination. There is indeed something it is like. And what it is like is given by our analysis—it is like being in a certain kind of perceptual state (with a certain kind of phenomenal character), even though you are not. In a sense, then, I am recommending the endorsement of a *disjunctive* account of the truthmakers of such claims: claims that there is something it is like to be in a particular mental state are made true *either* because the state in question has phenomenal character of a certain kind *or* because the state in question is indistinguishable from a mental state with such a phenomenal character. So even if you are inclined to think that the claim that you could be wrong when you think that there is something it is like to have a hallucination is simply absurd, this position can adapt to accommodate your intuition—it can allow that your judgement is true despite the hallucination lacking phenomenal character.

Given the claim that there is a disjunctive truthmaker for something-it-is-like claims—that the claim that there is something it is like to be in a mental state is true *either* because it possesses phenomenal character, *or* because it is indistinguishable from a state which has phenomenal character—there may appear to be a threat of overdetermining there being something it is like in the veridical cases, and a consequent threat of rendering the first disjunct's phenomenal-character explanation of there being something it is like superfluous. Robinson, for example, argues that:

If the experience of seeming to see something red when nothing relevant is red is adequately analysed in terms of acquiring some kind of belief state, then a similar analysis must also be adequate for the experience of seeming to see something red when that experience is veridical. (1994: 51)

The thought would be that, if we can fully account for its being true that there is something it is like to hallucinate in terms of the belief states a subject acquires, and if what it is like to hallucinate is the same as what it is like to perceive veridically (as the two are indistinguishable), then we ought to be able to account fully for its being true that there is something it is like to perceive veridically in terms of the belief states a subject acquires. But then it would look as though the first disjunct's explanation of there being something it is like—in terms of the phenomenal character of the experience—will drop out as redundant.

The threat posed by these objections is a metaphysical one. It is that, in offering a disjunctive explanation of there being something it is like, Occam's razor thereby recommends the elimination of the very thing that made this general approach attractive in the first place: a perceptual state with phenomenal character that can be given a broadly naïve realist explanation. However, we can draw the sting of this objection by clarifying the crucial dependencies in play. The reason there is something it is like to

have a hallucination—*despite* the fact that the hallucinatory state itself doesn't have a phenomenal character—is because it is indistinguishable from a veridical perception. Now this relational property—the property of being indistinguishable from another kind of state—is only sufficient to make it the case that there is something it is like to be in a state instantiating the relational property if there is *already* something it is like to be in the state that the relational property relates it *to*. The only reason we are disposed to judge that there is something it is like when we hallucinate is because we take ourselves to be perceiving *and there is something it is like to perceive*. And of course, the explanation of there being something it is like to perceive cannot be an explanation of the same kind, or we would be locked in a regress. So if veridical perceptual experiences did not have phenomenal character in the first place—phenomenal character that makes it the case that there is something it is like to be in such states—then the relational property of being indistinguishable from such an experience would not suffice to make it the case that there was something it is like to be in a state which instantiated that relational property. So the very possibility of there being something it is like to have a hallucination *requires* there to be other states—veridical perceptual experiences according to the current theory—that possess phenomenal character.

10

Throughout the foregoing presentations, we have been focusing on those cases of hallucination in which subjects, on the basis of their hallucination, really do come to believe that they see something which is not there. The time has now come to address a significant lacuna in this analysis: that not all cases that would naturally be brought under the banner of hallucination are cases in which the subject is fooled in this way. There are also resisted hallucinations to be accounted for—cases in which hallucinating subjects do not believe that things are as they are experienced to be, but where the subjects still claim to have the hallucinatory experience as of seeing something. As things stand, an adequate explanation of such resisted hallucinations eludes our analysis. To illustrate why, consider the following passage from Armstrong:

Suppose I am regularly subject to hallucinations, such as those of seeming to see cats, under certain conditions which I know and can recognise. Under these circumstances, I may well come to recognise my hallucination for what it is, an hallucination. I will not believe that there is a cat on the mat, and, a fortiori, I will not believe that I am seeing a cat. Yet there certainly is a black cat-like shape 'in my visual field'. (1961: 84)

As Armstrong points out, if I know that I am prone to hallucinations, that may provide enough countervailing evidence such that, when I hallucinate, I don't come to believe that I see a cat. But nevertheless, the failure to form a belief that I see a cat does not mean that I do not have a cat-like experience.

This is the problem of resisted hallucination. In fact, the problem has two related aspects. The first concerns those relationships of indistinguishability we intuitively take to hold between different experiences. For example, if you consider a hallucination of a cat on a mat in which the subject is taken in and the corresponding resisted

hallucination in which the subject is not, our natural inclination is to think that both of these experiences bear the relationship of indistinguishability to the same kind of veridical perception. But whereas the true hallucination and the veridical perception appear to have effects of the same kind—in particular, they both result in the subject believing that she sees that there is a cat on the mat—the resisted hallucination does not. So it looks as though, on the effects-based conception of indistinguishability, the appropriate indistinguishability relations do not hold where resisted hallucinations are concerned. The second aspect is that, as the resisted hallucinating subject lacks the higher-order belief that she sees something, it looks as though Armstrong's Explanatory Claim will be of no use to us when it comes to explaining why the resisted hallucinating subject claims to have an experience as of seeing something. Given that the resisted hallucinator *does not* believe that she sees a cat on the mat, we will be unable to appeal to such a belief to explain the kinds of things the subject thinks, does, and says.

It will help the clarity of subsequent discussions if we formalize things slightly here. Consider the following three scenarios: a veridical perception of an x , a true hallucination of an x , and a resisted hallucination of an x (an x -ish resisted hallucination). In the first two cases, the veridical perception and the true hallucination will have effects of the same kind, including particularly a belief of the form:

[B1] I believe that [I see that $\{x\}$]

The fact that hallucinations and veridical perceptions both have similar effects then gives us a way of understanding the claim that they are indistinguishable—because the effects the two mental states have on their subject are similar in this way, the subject will therefore be unable to activate information that they are distinct, and this will guarantee their indistinguishability.

Resisted hallucination, however, creates problems for our analysis because it concerns situations in which it is false that the subject comes to hold beliefs like [B1]. As we have seen, this is problematic for two reasons. First, because despite the fact that resisted hallucinations have different effects from veridical perceptions/true hallucinations, we are nevertheless inclined to treat them all as indistinguishable from one another. After all, as Armstrong suggests, our failure to form beliefs need not be for reasons of how things seem to us, but may instead be due to our awareness of countervailing considerations such as our knowledge that we are prone to hallucinations in this kind of context. The second reason why resisted hallucinations cause problems is because, despite the fact that the resisted hallucinator does not hold a belief like [B1], he will nevertheless still claim to have an experience *as of* an x . But in such a case, we cannot explain the subject's claims about his x -ish experiences by appealing to a higher-order belief that he sees an x as, *ex hypothesi*, he has no such belief.

In response to the first worry, note that we are explicitly considering resisted hallucinations that are purportedly indistinguishable from certain kinds of veridical perceptions. So inasmuch as these states have different effects from veridical perceptions, this cannot be because the subject can just introspectively tell that the episode is non-veridical as it is, *ex hypothesi*, indistinguishable from a veridical perception. In the truly indistinguishable cases which concern us, it is rather because the subject is aware that there are countervailing considerations to be taken into account which give her

reason to doubt that things are as they seem and thereby block her from forming the relevant higher-order beliefs. With this in mind, consider the following situation: once again, suppose that I am regularly subject to hallucinations of cats under certain conditions; perhaps in a particular armchair in my living room. But suppose that, unbeknownst to me and unexpectedly, there actually is a cat in that armchair in my living room. Given that I am aware that there are countervailing considerations in play and given that I have no reason to expect there to be a cat there, when I walk into the room and have a veridical experience of the cat, I would still have good reasons to doubt that things are as they seem. As I have often suffered hallucinations in just this context and have no reason to suppose that my current situation is relevantly different, I would therefore be highly likely to fail to form either the belief that there is a cat on the mat, or that I am seeing a cat. So if we hold the awareness of countervailing considerations constant, we can see that in fact, veridical perceptions and resisted hallucinations may not be as dissimilar in their effects as we might have thought. However, this is not yet to show that veridical perceptions and resisted hallucinations have effects of the same kind. All this shows is that, in some circumstances, veridical perceptions can fail to have certain effects for the same reason that resisted hallucinations fail to have effects—because the subject takes countervailing considerations into account. So what positive claims can we make about these situations?

Let's begin by focusing on the veridical case. Imagine we have a naïve subject who, for the first time, sees a *trompe l'œil* picture of a banjo hanging on the back of a door. When first seeing the picture, in the absence of any reason to think that this experience is in some way anomalous, our subject may well come to think that he really does see a banjo hanging on the back of the door. If we ask him about his experience, he may well, therefore, assert that:

[A2] I see that {there is a banjo hanging on the door}

And in saying this, he expresses a higher-order belief of the following form:

[B2] I believe that [I see that {there is a banjo hanging on the door}]

Now imagine that we take the subject to the door, and show him that the banjo is really just a craftily constructed picture, explain the history of *trompe l'œil* paintings, and so on. Once he has grasped what we are telling him, and understood that in situations like these there may be countervailing considerations to be taken into account, the next time the subject has a similar experience, he would be likely to hedge his bets, and make a much more cautious assertion, such as:

[A3] It is as if I see that {there is a banjo hanging on the door}

And once again, in making this statement, the subject expresses a higher-order belief. In this case, the belief has the form

[B3] I believe that [it is as if I see that {there is a banjo hanging on the door}]

In this case, the subject's experience hasn't changed—it is no more or less veridical than it was when he was naïve—but now the subject is taking the relevant countervailing considerations into account, whilst he is still acquiring a belief, it is

a different belief from that which he acquired in the naïve case. Where the naïve subject claims/believes that he sees, the more cautious subject merely claims/believes that it is as if he sees. So the precise effects a particular veridical perception has on any given occasion will depend upon other aspects of the subject's overall cognitive state.

With this in mind, consider once again the situation with regard to the different forms of hallucination. As we have seen, a veridical perception of an x would, in the absence of any countervailing considerations, lead us to form beliefs of the form [B1]. However, if the subject was aware of countervailing considerations to take into account, as in the *trompe l'œil* case we have just considered, she may well form a more cautious belief in place of [B1]. Such a belief might have the following form:

[B4] I believe that [it is as if I see that { x }]

Now consider a mental state that is not a veridical perception of an x but which comes to have effects that are sufficiently similar to those which such an experience would have had. In those situations where the subject is not aware of any countervailing considerations then, as we have seen, the effects of both that mental state and the original veridical perception will, as they are sufficiently similar, include higher-order beliefs of the form [B1]. The presence of such a belief in the hallucinatory case then provides the ammunition for the Explanatory Claim to account for everything the hallucinating subject thinks, does, and says.

If, however, the subject *were* taking countervailing considerations into account then, as we have seen, the effects of a veridical perception of an x would be different: they would not include beliefs of the form [B1], but rather more cautious beliefs of the form [B4]. So, if a mental state were to have the same kinds of effects as this veridical perception in the same kind of overall cognitive context (i.e. when countervailing considerations are being taken into account) we could expect it to also yield beliefs of the form [B4]. This enables us to resolve the first aspect of the problem of resisted hallucination as it gives us a way of constructing the relations of indistinguishability that intuitively hold between veridical perceptions, true hallucinations, and resisted hallucinations. Standard deceptive hallucinations have the kinds of effects a certain kind of veridical perception would have if the subject was trusting and did not take any countervailing considerations into account. As outlined earlier, this can ground the claim that hallucinations are indistinguishable from that kind of veridical perception. Resisted hallucinations can also have the same kinds of effects as a veridical perception of this kind, only in this case the effects are similar to the effects the perception would have had if the subject were more guarded. This can then ground the claim that resisted hallucinations are also indistinguishable from that kind of veridical perception. This enables us to see how both a true hallucination and a resisted hallucination could qualify as indistinguishable from a veridical perception of the same kind, and hence, by transitivity, how a true hallucination and a resisted hallucination could qualify as indistinguishable from one another, despite having different effects.

What is more, the fact that resisted hallucinators form beliefs of the form [B4] also provides a way of resolving the second aspect of the problem: that of explaining why a subject claims to have a certain kind of experience despite lacking the belief

that they see something. Higher-order beliefs such as [B4] would provide the ammunition for the Explanatory Claim to explain everything the *resisted* hallucinator thinks, does, and says. If a resisted hallucinating subject believed that it was as if she saw an x , then the Explanatory Claim would be able to appeal to this belief in order to explain her claim that it is as if she sees an x —she claims that it is as if she sees something because, like the cautious perceiving subject, she believes that it is as if she sees something. What the resisted hallucinator lacks, of course, is the higher-order belief that she *actually does* see something. But that presents us with no problems as it is only the more limited and non-committal claims of resisted hallucinators that need to be explained.

This concludes our positive story about hallucination. It recommends that a hallucination of an x is a mental state that, whilst not being a veridical perception of an x , nevertheless comes to have effects that are sufficiently similar to those such a veridical perception would have had. If the subject was conceptually sophisticated and was not aware of any countervailing considerations to be taken into account, these effects would include the higher-order belief that the subject was seeing an x ; if the subject were taking countervailing considerations into account, the effects would instead include the more cautious belief that it is as if the subject was seeing an x ; and if the subject were not sophisticated enough to have such beliefs at all, the effects would be solely behavioural. Taken together, these claims enable us to explain everything that hallucinating subjects think, do, and say, and thereby provides theoretical substance to the disjunctivist's assertion that two states which may potentially have nothing in common can nonetheless qualify as indistinguishable for their subject.

REFERENCES

- Aggernaes, A. (1972) 'The Experienced Reality of Hallucinations and other Psychological Phenomena: An Empirical Analysis', *Acta Psychiatrica Scandinavica*, 48: 220–38.
- Armstrong, D. M. (1961) *Perception and the Physical World*, London: Routledge and Kegan Paul.
- (1968) *A Materialist Theory of the Mind*, London: Routledge and Kegan Paul.
- Bentall, R. P. (1990) 'The Illusion of Reality: A Review and Integration of Psychological Research on Hallucinations', *Psychological Bulletin*, 107: 82–95.
- Broad, C. D. (1951) 'Some Elementary Reflexions on Sense-Perception', repr. in R. J. Swartz (ed.) *Perceiving, Sensing and Knowing*, Berkeley: University of California Press, 1965.
- Collerton, D., Perry, E., and McKeith, I. (2005) 'Why People See Things That Are Not There: A Novel Perception and Attention Deficit Model for Recurrent Complex Visual Hallucinations', *Behavioral and Brain Sciences*, 28: 737–57.
- Corne, S. J. and Pickering, R. W. (1967) 'A Possible Correlation between Drug-Induced Hallucinations in Man and a Behavioural Response in Mice', *Psychopharmacologia*, 11: 65–78.
- Dennett, D. C. (1991) *Consciousness Explained*, New York: Penguin.

- Halgren, E. (1982) 'Mental Phenomena Induced by Stimulation in the Limbic System', *Human Neurobiology*, 1: 251–60.
- Horowitz, M. J. (1964) 'The Imagery of Visual Hallucinations', *Journal of Nervous and Mental Disease*, 138: 513–23.
- , Adams, J. E., and Rutkin, B. B. (1968) 'Visual Imagery on Brain Stimulation', *Archives of General Psychiatry*, 19: 469–86.
- Locke, J. (1690/1961) *An Essay Concerning Human Understanding* London: Everyman's Library.
- Lowe, E. J. (1996) *Subjects of Experience*, Cambridge: Cambridge University Press.
- Lycan, W. (1996) *Consciousness and Experience*, Cambridge, MA: MIT Press.
- Mahl, G. F., Rothenburg, A., Delgado, J. M. R., and Hamlin, H. (1964) 'Psychological Responses in the Human to Intracerebral Electrical Stimulation', *Psychosomatic Medicine*, 26: 337–68.
- Manford, M. and Anderman, F. (1998) 'Complex Visual Hallucinations: Clinical and Neurobiological Insights', *Brain* 121: 1819–40.
- Martin, M. G. F. (2004) 'The Limits of Self-Awareness', *Philosophical Studies*, 120: 37–89.
- Noë, A. (2004) *Action in Perception*, Cambridge, MA: MIT Press.
- Penfield, W. and Perot, P. (1963) 'The Brain's Record of Auditory and Visual Experience: A Final Summary and Discussion', *Brain*, 86: 595–696.
- Rankin, P. M. and O'Carroll, P. J. (1995) 'Reality Discrimination, Reality Monitoring and Disposition towards Hallucination', *British Journal of Clinical Psychology*, 34: 517–28.
- Robinson, H. (1994) *Perception*, London: Routledge.
- Rosenthal, D. (1990) 'A Theory of Consciousness', in N. Block, O. Flanagan, and G. Güzeldere (eds.) *The Nature of Consciousness*, Cambridge, MA: MIT Press.
- Siegel, S. (2004) 'Indiscriminability and the Phenomenal', *Philosophical Studies*, 120: 91–112.
- Shoemaker, S. (1996) *The First Person Perspective and Other Essays*, Cambridge: Cambridge University Press.
- Slade, P. D. and Bentall, R. P. (1990) *Sensory Deception: A Scientific Analysis of Hallucination*, London: Croom Helm.
- Smith, A. D. (2002) *The Problem of Perception*, Cambridge, MA: Harvard University Press.
- Sturgeon, S. (2000) *Matters of Mind: Consciousness, Reason and Nature*, London: Routledge.
- (this volume) 'Disjunctivism about Visual Experience'.
- Tye, M. (1993) 'Reflections on Dennett and Consciousness', *Philosophy and Phenomenological Research*, 53: 893–8.
- Williamson, T. (1990) *Identity and Discrimination*, Oxford: Basil Blackwell.
- Yamamoto, T. and Ueki, S. (1980) 'The Role of Central Serotonergic Mechanisms on Head-Twitch and Backward Locomotion Induced by Hallucinogenic Drugs', *Pharmacology Biochemistry and Behavior*, 14: 89–95.

6

How to Account for Illusion

Bill Brewer

The question how to account for illusion has had a prominent role in shaping theories of perception throughout the history of philosophy. Prevailing philosophical wisdom today has it that phenomena of illusion force us to choose between the following two options. First, reject altogether the early modern empiricist idea that the core subjective character of perceptual experience is to be given simply by citing the object presented in that experience. Instead we must characterize perceptual experience entirely in terms of its representational content. Second, retain the early modern idea that the core subjective character of experience is simply constituted by the identity of its direct objects, but admit that these must be mind-dependent entities, distinct from the mind-independent physical objects we all know and love. I argue here that the early modern empiricists had an indispensable insight. The idea that the core subjective character of perceptual experience is to be given simply by citing the object presented in that experience is more fundamental than any appeal to perceptual content, and can account for illusion, and indeed hallucination, without resorting to the problematic postulation of any such mind-dependent objects.

This return from the current preoccupation with the representational *content* of perceptual experience to the more traditional idea of the direct *objects* presented in it also has promising consequences for the central disjunctivist contention, that perceptual experiences in which a person's subjective condition constitutes a simple openness to the mind-independent physical world are both explanatorily and metaphysically basic, only by derivative reference to which various failures of direct perceptual contact are to be characterized and understood.

Many thanks to Steve Butterfill, John Campbell, David Charles, Bill Child, Tim Crane, Imogen Dickie, Naomi Eilan, Anil Gupta, Adrian Haddock, John Hawthorne, Christoph Hoerl, Hemdat Lerman, John McDowell, Fiona Macpherson, Alan Millar, Jennifer Nagel, Johannes Roessler, Nick Shea, Alison Simmons, Paul Snowdon, Matt Soteriou, Helen Steward, Peter Sullivan, Charles Travis, Ralph Wedgwood, Michael Williams, and Tim Williamson, for comments on previous versions of this material. Special thanks to the editors and two additional referees for their excellent extended critical suggestions.

Let me begin by characterizing a *visual illusion* as an experience in which a physical object, *o*, looks *F*, although *o* is not actually *F*.¹ According to the early modern empiricists, especially Locke (1975) and Berkeley (1975a, 1975b), the way to account for the fact that something looks *F* in an experience is to construe that experience as the presentation to the subject of a *direct object*, which constitutes the core subjective character of the experience, and which must, supposedly, therefore itself *be F*. In cases of illusion, then, this direct object is distinct from the physical object *o* which is not *F*. On Locke's *materialist* view, the direct object of an illusion is a mind-dependent entity, which is *F*, which nevertheless sufficiently resembles a non-*F*, mind-independent object, *o*, for the latter to be the physical object which illusorily looks *F*.² According to the most straightforward version of Berkeley's *mentalism*, on the other hand, the direct object of an illusion is a mind-dependent entity, which is a part of an equally mind-dependent composite physical object *o.O* is not *F*, very roughly, because most of its parts are not *F*, and it does not behave, in general, in ways characteristic of *F*s: in particular we cannot use it as we can paradigm *F*s. Nevertheless, it looks *F*, on this occasion, because the part presented in the relevant illusory experience is *F*.³ Both approaches are widely regarded as unsatisfactory today, absolutely rightly in my view.⁴

According to current orthodoxy, the mistake goes right back to the idea that perceptual experience has its core subjective character given simply by citing its direct object, which must apparently therefore *be F* in a case of illusion, and hence must be distinct from the physical object, *o*, which illusorily looks *F* although it is not. The subjective character of perceptual experience is to be given instead by its *representational content*: how it represents things as being in the physical world around the subject. In an illusion, perception has the *false* representational content that *o* is *F*. In general, *o* looks *F* iff *o* is the referent of a perceptual content in which *F* is predicated. I call this the *Content View* (CV), since it characterizes perceptual experience by its

¹ I concentrate throughout on the case of vision. I believe that much of what I say applies equally to the other modalities, although I do not address this here. It will turn out later that the characterization given in the text is not sufficient for visual illusion. It is adequate to be getting on with, though, and its insufficiency highlights some interesting issues in the development of my argument.

² Things are of course more complicated in the case of secondary qualities, according to Locke. For, in one sense, all secondary-quality perception is illusory: nothing in the mind-independent physical world is red, *in the basic sense in which mind-dependent ideas are red*. Still, in having such an idea before the mind, a physical thing may look red*, that is, either disposed to produce red ideas in normal observers in normal conditions, or microscopically constituted in whichever way actually grounds that disposition. Some, but not all, such perceptions may then be illusory in a derived sense. None of these details are relevant for present purposes.

³ See Stoneham (2002) for a compelling presentation of this account of Berkeley. Note, as with Locke's account of the secondary qualities, predicates apply to persisting physical objects, according to Berkeley, in a way which is derivative of their more basic application to our fleeting ideas, which are their temporal, and 'personal', parts.

⁴ Having said that, I believe that there are significant, and illuminating, structural similarities between the latter Berkeleian view, and Lewis's (1998) account of the metaphysics of persisting (that is, perduring, rather than enduring) macroscopic physical objects, especially in the presence of his Ramseyan humility (Lewis 2002). My forthcoming book, *Perception and its Objects*, contains a detailed development of this suggestion.

representational content, and identifies the objects of perception as those to which reference is made by such contents.⁵

The debate between *disjunctivism* and its opponents is often characterized at this point as follows. It is agreed on both sides that cases in which it appears visually to *S* that *p* are either cases in which *S* (veridically) *sees* that *p*, or cases in which it *merely* (that is, non-veridically) appears visually to *S* that *p*.

Disjunctivism insists upon the explanatory priority and metaphysical distinction of the veridical case. '*S* sees that *p*' picks out a basic, unanalysable metaphysical condition of *S*'s visually apprehending the fact that *p*. The fact that *p* is, in this case, "within the reach of... [*S*'s] subjectivity" (McDowell 1986: 150), and is therefore essential to that very experiential condition itself. It *merely* appears visually to *S* that *p* whenever he is in an experiential condition which, although not a case of seeing that *p*, he cannot distinguish introspectively from this basic one. The mere appearance that *p* is therefore explanatorily parasitic upon veridically seeing that *p*. For it is defined by essential reference to such seeing, as anything else appropriately indistinguishable from it. Its merely appearing visually that *p* is also a metaphysically distinct condition from seeing that *p*. For the former does not, whereas the latter does, *essentially* involve the fact that *p*. Unlike the corresponding mere appearance, seeing that *p* is not a possible subjective experiential condition in the absence of the fact that *p*. The notion of a visual appearance to *S* that *p simpliciter*, is, therefore, both explanatorily and metaphysically, disjunctive: such things are *either* cases in which *S* sees that *p*, *or*, something derivative and quite distinct, cases in which it merely appears visually to him that *p*. 'It appears visually to *S* that *p*' is a locution which applies when one or other of two fundamentally quite distinct conditions obtains.

Opponents of disjunctivism, on the other hand, regard the general condition of its appearing visually to *S* that *p*, entirely non-committal as it is on whether or not *p*, as explanatorily and metaphysically basic. This is a unified single condition of *S*, which is the common experiential core to veridical and non-veridical appearances that *p*. When it is appropriately caused by the fact that *p*, or its relevant worldly constituents or correlates, we say that *S* sees that *p*. The visual appearance that *p* constitutes her seeing that *p*. In this way, the common experiential element is involved in an explanatory and metaphysical analysis of seeing that *p*. Otherwise, we say that it *merely* appears visually to *S* that *p*. The visual appearance that *p* constitutes some kind of illusion or hallucination. Again, the common experiential element is explanatorily and metaphysically basic.

⁵ I elucidate both the early modern empiricist insight which I aim to preserve and the CV alternative to it in terms of the most basic theoretical characterization of perceptual experience: does this proceed by the identification of an object presented in the experience, or a state of affairs represented by the experience as obtaining? This is certainly intended to indicate the metaphysical commitments of the two views—in the most abstract terms, that experiences are relations to object and relations to contents respectively—but a great deal is clearly left metaphysically open by my characterizations: which objects, what are contents, which relations, are there other relata, and so on? Note that my own development of the early modern empiricist insight differs significantly from those of its initiators in answering affirmatively to the last of these questions.

According to the disjunctivist proponent of CV, then, a visual illusion is *illusory* because its subjective character is given by a *false* representational content. It is therefore an experience of an explanatorily derivative kind. It has to be understood as something akin to a failed attempt at seeing, that is, by reference to a more basic, and quite distinct, experiential condition: seeing that *o* (or something like it) is *F*. This condition would only have been possible, though, had the world been quite different from the way it actually is: had *o*, or something else suitably placed and appropriately like it, been *F*, rather than not.

Now, every visual experience is plausibly illusory in some respect: something visible looks in some way, perhaps only slightly, different from the way it actually is. So this disjunctivist—the disjunctivist who is also a proponent of CV—appears committed to the idea that all actual perception is explanatorily and metaphysically parasitic upon some non-actual ideal. Full direct contact between mind and world is never actually established in human perceptual experience. This is a strange result; but I will not consider it further here. I think that the problem lies, not in the disjunctivist idea itself, that perceptual experiences in which a person's subjective condition constitutes a simple openness to the mind-independent physical world are both explanatorily and metaphysically fundamental, only by derivative reference to which various failures of direct perceptual contact are to be explained, but, rather, in the background commitment to CV. What I aim to do here, is to articulate and defend an alternative *Object View* (OV), which retains the early modern conviction that the core subjective character of perceptual experience is to be given simply by citing its direct object, but which is able to resist any inference from the existence of illusion to the identification of such direct objects with mind-independent entities distinct from the persisting (and, indeed, enduring)⁶ mind-dependent physical things we know and love, and which also provides a more conducive setting for the basic disjunctivist idea.

So, what is the *Object View*, and how does it account for illusion? The basic idea is that the core subjective character of perceptual experience is given simply by citing the physical object which is its mind-independent direct object. This is what I earlier called the early modern empiricist insight, that perceptual experience should be conceived as a relation between a perceiving subject and the object presented. From various points of view, and in various circumstances of perception, physical objects have *visually relevant similarities* with paradigms of various kinds, or types, of such things. These may intelligibly lead us to take them as instances of such kinds when seen from the relevant points of view in the circumstances in question. Thus, they look various ways to us. So OV supplements the early modern insight with the insistence that perceptual experience should be conceived as a *three-place* relation, in which the third relatum is an index of the conditions of perception, which involve the subject's spatio-temporal point of view, and other relevant circumstances, such as lighting, and

⁶ Any serious engagement with the debate between endurance and perdurance accounts of the persistence of physical objects is of course well beyond the scope of the present essay. I simply record my conviction that OV and the endurance view are mutually supporting. See Lewis (1998) for the launch of the perdurantist's case. Sider (2001) constitutes an excellent recent discussion with the same sympathies. My forthcoming book begins to make the connection which I see between OV and the defence of endurantism.

so on. Illusions are simply cases in which the direct object of experience has such visually relevant similarities with paradigms of a kind of which it is not in fact an instance.

Consider, for example, the Müller-Lyer illusion (ML), in which two lines which are actually identical in length are made to look different in length by the addition of misleading hashes. The ML diagram is visually-relevantly similar to a pair of lines, one longer and more distant than its plane, one shorter and less distant—a paradigm of inequality in length. It is therefore perfectly intelligible how someone seeing it might take that very diagram as consisting of unequal lines, regardless of whether she does or not. In this sense: they look unequal in length.

Which similarities are *visually relevant*, though; for anything has unrestricted similarities with everything? Clearly, the visually relevant similarities cannot be defined as identities in the ways the relata are visually *represented* as being, or else OV depends upon CV. That is to say, we cannot simply say that two objects have visually relevant similarities just when there are sufficiently many common properties amongst those which each is visually represented as having. Rather, they are, according to OV, to be what ground and explain the ways in which the relata may intelligibly be taken to be when seen. That is to say, two objects have visually relevant similarities when they share sufficiently many common properties amongst those which evidently have a significant involvement in the external physical processes underlying vision. Thus, and very crudely, visually relevant similarities are identities in such things as, the way in which light is reflected and transmitted from the objects in question, and the way in which stimuli are handled by the visual system, given its evolutionary history and our shared training during development.⁷

Furthermore, what are the *paradigms* of physical kinds supposed to be? Again very roughly, these are instances of the kinds in question, whose association with the terms for those kinds partially constitutes our understanding of them, given our training in the acquisition of the relevant concepts: paradigm exemplars of the kinds in question relative to our grasp of the concepts for those kinds.⁸

It may be objected at this point that similarity is symmetrical. So OV has the unacceptable consequence, in connection with the ML, for example, that the relevant paradigm pair of lines of unequal lengths at different depths look equal in length, for the very same reason.⁹

I would make two points in reply to this objection. First, *a* looks *F*, to a first approximation, according to OV, iff *a* may intelligibly be taken to be *F*, when seen from the point of view in question in the relevant circumstances of perception, in virtue of its visually relevant similarities with certain *paradigm F*s. The ML diagram does not constitute a paradigm case of lines which are equal in length. So, although plain similarity is symmetrical, the relevant condition of similarity to a paradigm is

⁷ I entirely acknowledge that this is a very rough placeholder for what must in the end be a far more developed account of visually relevant similarities.

⁸ This idea is clearly in need of far more extended discussion. It also involves a controversial account of concepts and their possession. To make progress with my primary argument here I will have to leave further elucidation and defence for another occasion; but see Fodor (1998) for strong opposition.

⁹ Thanks to Tim Williamson for this objection.

not. Second, misleading cues *could*, no doubt, be added to unequal lines at different depths, to bring about an inverse to the ML illusion. Notice, though, that *which* such cues should be added, would be ascertained precisely on the basis of knowledge of the physical processes involved in vision: those which would intelligibly ground a mistaken judgement as to the relative lengths of the lines if the composite were seen from the point of view in question, in the relevant circumstances, in virtue of its visually relevant similarities with an appropriate paradigm of lines of equal length.

I claim that the same account covers many of the most standard cases of visual illusion. Here are two further examples for illustration.

First, a partially submerged straight stick looks bent. Here, the direct object which constitutes the core subjective character of the experience is that very (straight) stick itself. Nevertheless, it looks bent, in virtue of its visually relevant similarities with a bent stick, with a coincident unsubmerged portion, and its bottom half in the position of the relevant virtual image, from the subject's point of view and given the refractive index of the liquid in question. These similarities exist in virtue of the refraction of the light from the submerged portion of the actual stick at the surface of the liquid. The illusion consists in the fact that those very similarities may intelligibly be taken for qualitative identities. The actual stick, presented as it is in experience, from the point of view in question and in the relevant circumstances—especially given the presence of a refracting liquid—in this sense *looks bent*.

Second, a white piece of chalk illuminated with red light looks red. Again, the OV proposal is that the core of the subjective character of such illusory experience is constituted by that very piece of chalk itself: a particular persisting (by enduring) mind-independent physical object. From the viewpoint in question, and given the relevant perceptual circumstances—especially, of course, the abnormally red illumination—it looks red. This consists in the fact that it has visually relevant similarities with paradigm red objects: the light reflected from it is like that reflected from such paradigms in normal viewing conditions.

Of course there are very many quite different kinds of visual illusion. I cannot possibly consider representatives of all kinds. Further materials for the overall OV account will emerge in what follows; but I leave the direct enumeration of examples for now, and move on to introduce further key distinctions of principle in the OV approach to illusion. To set the scene, we should ask how OV accounts for hallucination.

The intuitive category I have in mind consists of *purely inner* phenomena, in the following sense. Nothing in the mind-independent world is subjectively presented in hallucinatory experiences. They have no mind-independent direct object. OV also rejects any characterization of hallucination in terms of purportedly mind-independent direct objects. Rather, hallucinatory experiences have to be characterized by giving a qualitative description of a mind-independent scene, and saying that the subject is having an experience which she cannot distinguish by introspection alone from one in which the constituents of such a scene are the direct objects. No more positive characterization of the experience may be given.¹⁰ Thus, for

¹⁰ I rely heavily here upon Martin (2004).

example, I once had an experience which I could not distinguish by introspection alone from one in which a large pink elephant in a desert was the direct object of my perception.

Now, hallucinations may be caused in many and varied ways, such as by taking certain drugs or getting a firm knock to the head. Other ways of bringing about hallucination may also involve distal external objects, sometimes relatively systematically. Indeed, this may even occur in cases in which the relevant mind-independent objects are also presented as direct objects of vision, supplemented, as it were, by their hallucinatory products. OV has the resources to account for some cases which may pretheoretically be classified as illusions in this way. Hermann's Grid, in which pale grey patches appear at the intersections of the white channels formed by a grid of closely spaced black squares, is plausibly a case in point.¹¹

OV also has a relatively straightforward account of the distinction between illusion and robust delusion. A person may be so-conditioned that sight of a cat always makes him think that there is a pouncing tiger before him. Indeed, it may cause him to think that that very thing is a pouncing tiger. But this is clearly not a case of visual illusion. The obvious explanation is that there are insufficient *visually relevant* similarities between any given such cat and any paradigm tiger for this to be an *intelligible* way in which to take the cat which is presented to him.

This account does raise an interesting issue, though, of how OV should best classify a case in which a person takes a life-like papier-mâché model before her to be a pouncing tiger. Surely this may well be precisely in virtue of its visually relevant similarities with paradigms of tigerhood. Yet this is intuitively no *illusion*.¹² My inclination is absolutely to acknowledge that its classification as an *illusion* is intuitively forced; but to explain this by pointing out significant dissimilarities with paradigm cases of illusion, even in light of the similarity in explanation given for the error: the model is indeed visually-relevantly similar to paradigm tigers, which explains why it may intelligibly be taken to be a tiger.

The obvious point of disanalogy lies in the relation between paradigms and instances of the kinds in question: *unequal in length*, on the one hand, and *tiger*, on the other. Locke's (1975: II. xiii–xxiii) distinction between *mode* and *substance* is surely the key here, especially his idea that the real and nominal essence are identical in the former case, whereas they are distinct in the latter. Very roughly, I want to say, when considered from every point of view and in all relevant circumstances of perception, being identical in visually relevant respects to a paradigm of unequal lines is sufficient for actually being unequal in length, whereas being identical in visually relevant respects to a paradigm tiger is not sufficient for being a tiger. This, I suggest

¹¹ Illusions due to sudden faults in the visual system, which are not intelligible in the light of any evident visual similarities between the object in question and paradigms of the kind illusorily involved, also belong in this class.

¹² This is the point at which my earlier characterization of a visual illusion emerges as insufficient. For, although intuitively not a visual illusion, it *is* a case in which a physical object, *o*, looks *F*, although *o* is not actually *F*—the papier-mâché tiger looks 'tiger'-instantiating, although it is not actually 'tiger'-instantiating, although the roundabout way in which this has to be put is in itself revealing.

is why we classify ML as a visual illusion, whereas, although the papier-mâché model looks just like a tiger, we do not regard visual experience of it as illusory.¹³

You might object at this point that, in rejecting CV, and insisting that the core subjective character of perceptual experience is to be given simply by citing its mind-independent direct object, given the relevant point of view and circumstances of perception, OV misses entirely the crucial point: illusions like ML are *experiential*. The ML lines look, *phenomenologically*, unequal in length!¹⁴ I agree with the datum, but I disagree that only CV, and not OV, may accommodate it.

This can be illustrated by reflection on the phenomenon of aspect-seeing. In the most basic case, the concept 'duck' is intelligibly applicable on being presented with a particular mind-independent animal as the direct object of perception, in virtue of its visually relevant similarities with paradigms of duckhood. In this sense, it looks like a duck. Given the actual direct object involved, with its visually relevant similarities to what we take to be paradigms of duckhood, we can even apply this characterization of the experience to a child without the concept of a duck, if we wish, although the characterization makes essential reference to the paradigms constitutive of our grasp of that concept. All that is involved in her having the experience, though, is that that very animal is presented, with the similarities it actually has with various paradigms *of ours*, not that her experience in any way *represents* it as being any such way. Reference to that object, given her viewpoint and the relevant circumstances, entirely captures this phenomenology. Noting the intelligible applicability of our concept of a 'duck', in virtue of the *de facto* similarities with our paradigm, *we* may see it *as* a duck. This is a further genuinely phenomenological affair; but the difference in how things are for us phenomenologically is no change in the core subjective character of the experience; it rather concerns our classificatory engagement with what is presented to us in it: that *duck*, as we would now say. This is the further phenomenology of actual and intelligible conceptual categorization, or recognition, not that of basic experiential presentation, which is common throughout. Still, it is aptly titled *phenomenology*, all the same.

Conceptual phenomenology of this kind is not simply a matter of being caused to make a judgement employing the concept in question. It is a matter of actively and intelligibly subsuming the particular presented in the core subjective character of the experience under that concept, in virtue of its evident similarities with the paradigms central to our understanding of that concept. We may simply find ourselves with that concept in mind, but, in cases of seeing as, it is evidently appropriate to that particular in virtue of the *de facto* existence and attentional salience of such visually relevant similarities.

Similarly, in connection with Jastrow's (1900) Duck-Rabbit (see also Wittgenstein 1958: II, ii). Suppose that I am simply presented with the diagram. According to OV the core subjective character of my experience is given simply by that diagram itself.

¹³ My discussion of the issues raised by this kind of case is clearly inadequate as it stands. See my forthcoming book for further details.

¹⁴ Thanks to Ian Philipps (draft) for pressing this objection very forcefully in his paper at the 2005 Warwick University MindGrad conference.

It actually has visually relevant similarities with my paradigms of both a duck and a rabbit. In this sense both concepts are intelligibly applicable, and it looks like both a duck and a rabbit, regardless of whether I notice either resemblance: perhaps it is presented at a peculiar angle.¹⁵ When I see it *as* a duck, say, this is again a phenomenological change, but one of actual conceptual classificatory engagement with the very diagram presented to me. Similarly, when I shift aspects, and see it as a rabbit, there is an alteration in this phenomenology of the categorization of what is presented.¹⁶

Consider now the case of ML. Suppose that someone has the diagram visually presented to her, from head-on, and in good lighting conditions, with eyes open and a normally functioning visual system. According to OV, the core subjective character of her visual experience is simply constituted by that diagram itself. From that viewpoint, and given the circumstances of perception, it has visually relevant similarities with a paradigm pair of unequal lines at different depths. In this sense, the concept of inequality in length is intelligibly applicable to its main lines: the lines look unequal in length. Once again, we may even mark the *de facto* existence of these visually relevant similarities in this way in connection with children, without the relevant concepts. This is genuine phenomenology, which flows directly from the identification of the direct object of the experience, given the viewpoint and relevant circumstances. It is fully captured by OV without any need for CV. Possessing the concept, as I do, I may notice the intelligible applicability of 'unequal in length' to the direct object of my ML experience. This may be, either because the question of the relative length of its main lines becomes relevant, and I attend accordingly, or simply because the intelligible applicability of this concept jumps out at me, or *captures* my attention. In this more robust sense, the lines now look unequal in length to me, regardless of whether I actually judge them to be so. This is a perfectly genuinely phenomenological matter, but one which is again captured entirely by OV, along with my deployment of attention and active conceptual endowment, without any need for CV.

A very closely related worry about OV can be sharpened, though, in the form of a dilemma.¹⁷ Is the concept of inequality in length essential to the subjective character of my ML experience, or not? If not, then OV fails to capture the robustness of the illusion: someone might have just that experience, and yet the lines not look unequal in length, which seems wrong. If so, if inequality is essential to the subjective

¹⁵ It has visually relevant similarities with paradigms of both a duck and rabbit even if I don't actually have either concept myself. In this weaker sense, even then, it may be said to look like both, although I cannot classify it as (like) either.

¹⁶ Having granted that core subjective character is given by citing the object presented, and marking, as I do, a key distinction between presentation in experience and representation in thought, how can it then be said that conceptual classification induces a phenomenological alteration? This is a good question, deserving extended discussion. I can only make two points here. First, I reject any assumption behind the question, that there is a single, simple, and exhaustive account of the phenomenology of a given experience: this notion has layers in my view, the object at the core, and various levels of actual or intelligible classification based upon that. Second, there is surely something quite familiar and harmless in the idea that recognition—of a cloud as shaped like a bull, or of a doodle as a distorted name, say—is *both* classificatory and phenomenological, although this may indeed be difficult fully to explain. Given my rejection of the myth of a single simple level of phenomenology, this difficulty is no more acute on OV than on any other account.

¹⁷ Many thanks to Matt Soteriou for presenting the objection to me in this form.

character, that is, since the lines which are its direct objects are not themselves unequal in length, representational properties are surely also essential to its subjective characterization, and CV is back in business.

This is a very helpful challenge. For it enables us to appreciate the interesting and subtle division of labour between ourselves and the world in the ML illusion. According to OV, the concept of inequality in length is appropriate to the subjective character of ML experience owing to the visually relevant similarities between its direct object—the ML diagram—and certain paradigm cases of unequal lines. These obtain, relative to the viewpoint and circumstances of perception in question, in virtue of the intrinsic nature of the direct object itself, *given the normal operation of our perceptual systems, their historical evolution and our developmental training, and what we take to be paradigms of inequality in length*. Thus, were these factors concerning us and our perceptual systems to have been sufficiently different, the ML lines would not look unequal in length: an experience with them as direct object would not be misleading. This is right, in my view: the illusion is not unrestrictedly robust.¹⁸ Still, given these deep contingencies about our evolution and development, intelligible applicability of the concept of inequality in length follows simply from the intrinsic nature of the direct object of the ML illusion and the subject's point of view and other relevant perceptual circumstances, without any need for appeal to independent representational properties in the subjective characterization of the experience of it. Thus, OV is entirely in the clear.

Still, one might worry now that, since the apparent inequality of the lines is supposed to be a genuine feature of the phenomenology of our experience of the ML diagram, as I explained in discussing the phenomenon of aspect-seeing above, any significant *distinction* between OV and CV has been lost, or, at the very least, that the final view which I am proposing is some combination of the two. This would be a mistake. I certainly argued earlier that CV alone is bound to leave something essential out of its account of the subjective character of perceptual experience. I also believe that the position at which we have arrived is firmly to be distinguished from any simple combination of OV and CV.

An analogy with Grice on conversational implicature may be illuminating (Grice 1989a, 1989b). We communicate a complex message to our audience when we speak a language they understand. On one view, every aspect of this message is to be regarded as part of an undifferentiated notion of the meaning of what is said. This effectively undermines the possibility of any fruitful and illuminating systematic theory of meaning (Baker and Hacker 1984: ch. 4). On the Gricean alternative, we get a more satisfying and complete picture of the situation by regarding the whole message communicated as the product of at least the following two factors. First, the core *semantic meaning*: what is strictly and literally said to be the case. Second, any pragmatic implicature, which may be conveyed by choosing to say something with just those core truth-conditions in the circumstances, given the conventions

¹⁸ See McCauley and Henrich (2006) for empirical confirmation from results which suggest that susceptibility to the ML illusion is dependent upon being in a carpentered world, whose orthogonal joints invest the diagram's hashes with their misleading association with depth.

governing good communication. Notice, in particular, that the pragmatic elements of the message are systematically to be *explained* by appeal to the core semantic element, plus a further, independently plausible, theory, concerning the etiquette of good conversation.

Thus, I may say ‘the vice-chancellor was sober at the party today’, and thereby communicate that he had not drunk very much and was behaving respectably and so forth, but also that this is not normally the case—that he normally gets inappropriately inebriated. On the Gricean view, the strict and literal meaning of my words is simply that he was sober at the party: not drunk, period. Still, since that is only a remotely interesting or informative thing to say when it is not normally so, and there is a standing convention of good conversation only to say something when it is interesting, relevant, and informative, I thereby also communicate that he is normally embarrassingly inebriated at such events.

Similarly, I claim that, although ‘unequal in length’ really is part of how the ML lines look, *phenomenologically*, it is best to regard this fact as the product of a more basic phenomenal presentation of those very lines themselves, from the viewpoint and in the circumstances of perception in question, along with a further, independently motivated, theory, of how they may therefore intelligibly strike us, given our evolutionary niche, developmental training, conceptual endowment, and our attention, interests, and concerns at the time of viewing. Inequality in length is not an independent part of the phenomenology of the situation, to be explained by an appeal to CV, entirely orthogonal to the crucial application of OV in giving the core subjective character of perceptual experience simply by citing the relevant mind-independent physical object presented, from the point of view and in the perceptual circumstances in question. The relevant phenomenological ‘looks’ phenomena flow directly from the core early modern empiricist insight at the heart of OV, in the context of appropriate background theoretical materials. So it is not even accurate to portray CV as a correct account of part of the illusory phenomenon. CV is not even acceptable as an autonomous account of part of the phenomenology of perception.

I remarked earlier that the modern CV disjunctivist faces the implausible prospect that all actual perception is parasitic, both in its metaphysical nature, and in the order of correct philosophical explanation, on a non-actual ideal of perfect subjective openness to the mind-independent facts. OV provides a far less hostile environment for the disjunctivist insight, that perceptual experiences in which a person’s subjective condition constitutes a simple openness to the mind-independent physical world are both explanatorily and metaphysically basic, only by derivative reference to which various failures of direct perceptual contact are to be characterized and understood. Hallucination is to be characterized and understood only negatively, as it were, as a condition introspectively indistinguishable from a non-actual presentation of worldly objects of certain kinds; and this category may even expand to include certain cases pretheoretically classed as illusions. Still, illusion proper consists in a genuine subjective presentation of particular mind-independent objects themselves, although from a viewpoint and in circumstances of perception in which they provide the perfectly intelligible ground for some misclassification or miscategorization of those very

things, in virtue of their visually relevant similarities with paradigms of kinds of which they are nevertheless not actually instances.¹⁹

In conclusion, I contend that the early modern empiricist insight, that the subjective character of perceptual experience is to be given simply by citing the objects presented in such experience, in a given sense modality, from the point of view and in the circumstances of perception in question, is perfectly capable of accounting for the phenomena of illusion, and, indeed, of hallucination too, without any pressure towards regarding such direct objects as mind-dependent entities, distinct from the mind-independent physical objects we all know and love. Perhaps ironically, this early modern empiricist idea also provides a far more hospitable setting for disjunctivism than the more recent CV, within the context of which disjunctivism is normally formulated and discussed. In any case, OV provides the best account of illusion consistent with the empirical realist conviction that the objects which are subjectively presented to us in perceptual experience are the genuinely mind-independent persisting (by enduring) physical things we all know and love.

REFERENCES

- Baker, G. P. and Hacker, P. M. S. (1984) *Wittgenstein: Meaning and Understanding*, Oxford: Blackwell.
- Berkeley, G. (1975a) *Three Dialogues between Hylas and Philonous*, in *George Berkeley: Philosophical Works*, ed. M. Ayers, London: Everyman.
- (1975b) *A Treatise Concerning the Principles of Human Knowledge*, in *George Berkeley: Philosophical Works*, ed. M. Ayers, London: Everyman.
- Brewer, B. (forthcoming) *Perception and its Objects*, Oxford: Oxford University Press.
- Fodor, J. (1998) *Concepts: Where Cognitive Science Went Wrong*, Oxford: Oxford University Press.
- Grice, H. P. (1989a) 'Logic and Conversation', in his *Studies in the Way of Words*, Cambridge, MA: Harvard University Press.
- (1989b) 'Further Notes on Logic and Conversation', in his *Studies in the Way of Words*, Cambridge, MA: Harvard University Press.
- Jastrow, J. (1900) *Fact and Fable in Psychology*, New York: Houghton Mifflin.
- Lewis, D. (1998) 'The Problem of Temporary Intrinsic: An Excerpt from *On the Plurality of Worlds*', in P. Van Inwagen and D. W. Zimmerman (eds.) *Metaphysics: The Big Questions*, Oxford: Blackwell.
- (2002) 'Ramseyan Humility', Gareth Evans Memorial Lecture to Oxford University, delivered posthumously by Stephanie Lewis, 2 March 2002.
- Locke, J. (1975) *An Essay Concerning Human Understanding*, ed. P. H. Nidditch, Oxford: Oxford University Press.

¹⁹ It is also worth pointing out in this connection that OV entails the disjunctivist insight that perceptual experiences in which a person's subjective condition constitutes a simple openness to the mind-independent physical world are both explanatorily and metaphysically basic, whereas there are certainly *prima facie* consistent forms of CV which reject this.

- McCauley, R. N. and Henrich, J. (2006) 'Susceptibility to the Müller-Lyer Illusion, Theory-Neutral Observation, and the Diachronic Penetrability of the Visual Input System', *Philosophical Psychology*, 19: 79–101.
- McDowell, J. (1986) 'Singular Thought and the Extent of inner Space', in P. Pettit and J. McDowell (eds.), *Subject, Thought and Context*, Oxford: Oxford University Press.
- Martin, M. (2004) 'The Limits of Self-Awareness', *Philosophical Studies*, 120: 37–89.
- Philipps, I. (draft) 'Illusion and Content', paper presented at the 2005 Warwick University MindGrad Conference.
- Sider, T. (2001) *Four-Dimensionalism: An Ontology of Persistence and Time*, Oxford: Oxford University Press.
- Stoneham, T. (2002) *Berkeley's World*, Oxford: Oxford University Press.
- Wittgenstein, L. (1958) *Philosophical Investigations*, Oxford: Blackwell.

7

Disjunctivism and Discriminability

A. D. Smith

Disjunctivism is defined by the claim that when one genuinely perceives some object in the real world, one thereby has an experience of a kind that it is impossible to have unless one is thus genuinely perceiving something. Even if it is possible, as is generally accepted, for there to be experiences that are not perceptions of any real object and yet are subjectively indistinguishable from such perceptions—totally convincing hallucinations—these cannot possibly be of the same psychological kind as the experiences had when genuinely perceiving. This is because perceptual experiences, even considered just as experiences, are essentially *of* some real object. Disjunctivists are, thereby, united in rejecting what J. M. Hinton, the father of disjunctivism, called “the doctrine of the ‘experience’” as something supposedly to be found both when one perceives and when one hallucinates (Hinton 1973: 71). These days the rejected view is commonly referred to as the ‘highest common factor’ view of experience (following McDowell 1982).¹ Although this is the essential core of the position, disjunctivism is also commonly associated with a further thesis. This thesis is, in the words of Bill Child, that “a hallucination is simply a state of affairs in which the subject is not seeing anything, but which is for her just like a case of vision” (1994: 144).² An experience is ‘just like’ a perception ‘for’ a subject if it is, as it is now commonly put, ‘subjectively indistinguishable’ or ‘subjectively indiscriminable’ for the subject from such a perception. The notions of discriminability and indiscriminability are standardly spelled out in terms of possibilities of a certain kind of knowledge. As M. G. F. Martin writes, for example, “For one situation to be indiscriminable from another requires only that it not be possible to know that it is distinct in kind” (Martin 2004: 50). The specifically

¹ There is some lack of uniformity among disjunctivists concerning how to treat illusions. Some would strengthen the claim stated in my first sentence by replacing ‘genuinely perceive’ by ‘genuinely and veridically perceive’. Although I have views on this matter, and shall treat it in some detail elsewhere, in this essay I ignore the issue of illusion entirely. When I contrast hallucination and perception, some may wish to substitute ‘veridical perception’ for ‘perception’. They are fully at liberty to do so.

² Obviously, the claim is meant to extend beyond the visual case.

'subjective' indiscriminability that is supposed to attach to a hallucination is the impossibility of knowing, just on the basis of the experience itself, that one is not perceiving.³ Martin expresses this by saying that the discrimination in question should be just on the basis of "introspection and reflection" (Martin 2004: 47).⁴ Meeting this 'negative epistemic criterion', as it is sometimes called, while yet not being a genuinely perceptual state, is therefore presented as a sufficient condition for being a hallucination.

Why disjunctivism has come to be associated with this second thesis is an interesting question, which I hope to explore elsewhere. Here, however, I shall argue that the thesis is false: meeting the negative epistemic criterion is insufficient for a non-perceptual state to be a hallucination. There are a number of experiences that are not hallucinations (as the term is generally employed in the philosophical literature), but that, equally with hallucinations, cannot be subjectively discriminated from perceptual experiences. They are not hallucinations, since they are not *sensory* experiences at all, whereas hallucinations are. In order that meeting the negative epistemic criterion should be sufficient for a non-perceptual experience to be a hallucination, meeting it must be sufficient for a state to be a sensory state. Martin recognizes this implication of the position, and explicitly presents the negative epistemic criterion as a sufficient condition for a state to be sensory (or, as he sometimes puts it, perceptual) in character: "The concept of perceptual experience in general is that of situations indiscriminable from veridical perception . . . Indiscriminability provides sufficient conditions for an event's being a sensory experience" (2004: 65 and 74). It is precisely this, given that indiscriminability is explicated in terms of the impossibility of a certain kind of knowledge, that is not true.

If the negative epistemic criterion is inadequate to pick out sensory states, and hence inadequate to pick out hallucinations among non-perceptual states, disjunctivists will have either to abandon or modify the criterion. In the penultimate section of this essay I consider a modification that at least improves matters. It is, however, a modification that one particular form of disjunctivism cannot endorse. The form of disjunctivism in question I shall term, not unfairly I think, 'extreme disjunctivism'. According to this view, at least some possible hallucinations—specifically those that have the same kind of proximal cause as some (possible) veridical perception—have no positive psychological or mental features at all that account for their phenomenal character. As M. G. F. Martin, the leading advocate of such extreme disjunctivism, writes concerning a hallucination of this type, "Nothing more can be said than the

³ The restriction that such discrimination should be solely on the basis of the experience itself, or 'what it is like' to have the experience, is obviously required, since hallucinations can certainly be discriminated from perceptions in other ways. I might, for example, learn that I am hallucinating by being informed of the fact by someone who knows; or I might infer that I am hallucinating because I regard the apparently presented scene as physically impossible.

⁴ When the criterion is expressed in this way, it is presupposed, of course, that the subject does have reflective access to the state in question. Otherwise merely physical states of a person would be hallucinations. The state of my pancreas, for example, is not reflectively discriminable by me from a veridical perception. A problem may, therefore, lurk here when it comes to applying the criterion to non-reflective perceiving subjects. Martin (2004 and 2006) has addressed this worry at some length, and I shall not press it here.

relational and epistemological claim that it is indiscriminable from . . . perception” (Martin 2004: 72). This inability to say more is not merely a reflection of our necessary ignorance. We can say no more because there is no more to the mental nature of such a hallucination: “There are certain mental events, at least those hallucinations brought about through causal conditions matching those of veridical perceptions, whose only positive mental characteristics are negative epistemological ones—that they cannot be told apart by the subject from veridical perception” (2004: 73–4). This extreme position goes beyond disjunctivists’ general acceptance of the negative epistemic criterion. One could accept that being subjectively indistinguishable from some (possible) veridical perception, while not being a perception, is sufficient for an experience to be a hallucination, and yet hold that all hallucinations have some positive psychological character or other. It is, it might be thought, precisely *in virtue of* such a character, whatever it is, that hallucinations are indeed subjectively indiscriminable from perception. John McDowell, for example, claims that all “ostensible seeings”—a term that covers both visual perceptions and visual hallucinations—involve “representations” and “conceptual shapings of sensory consciousness” (McDowell 1998: 476). It is only because an experience has at least some representational content in common with some possible perception, and because it, too, involves ‘sensory consciousness’, that it is a hallucination, and a specific sort of hallucination. So, in explicitly ruling out such positive character, at least for a certain class of hallucination, the extreme view goes beyond the negative epistemic criterion. It is precisely in virtue of this extra step that, I shall argue, extreme disjunctivism is incapable of making the modification to the negative epistemic criterion that I shall introduce towards the end of this essay. Moreover, such an extreme position cannot abandon this criterion altogether, as other forms of disjunctivism perhaps can, and hold that having some specific psychological feature, not entailed by meeting the negative epistemic criterion, is necessary in order for an experience to be a hallucination. Since, according to the extreme view, at least some possible hallucinations have no positive psychological character at all, in their case there cannot be anything to make good the failings of the negative epistemic criterion itself. Extreme disjunctivism is, therefore, false.

Since much in this essay turns on the distinction between states that are and states that are not ‘sensory’ in character, I should say something about this distinction. Rather than presenting a positive account of what it takes for a state to count as sensory—something that could hardly avoid being question-begging at this stage—I shall offer a grouping of various kinds of psychological state into the sensory and non-sensory kinds. This should at least serve to afford an intuitive grasp of the distinction as it is employed here. First, and obviously, perceptual experiences are sensory in nature. Secondly, and hardly less obviously, merely thinking about something, having a belief about something, knowing something (even ‘occurrently’), and remembering something are not sensory states.⁵ Thirdly, and more significantly, mental imaging and concrete imagining—such as visualizing—are not sensory states. When one is in a sensory state (and one is not inattentive, confused, or otherwise psychologically impeded), one seems to be confronted with an actual instance of a perceptible

⁵ ‘Merely’ is meant to exclude cases that fall into the preceding and the following category.

quality.⁶ This is the case when one perceives. If, for example, I see a red object in front of me, I seem to be acquainted with an instance of redness. This is not the case when I merely think, or know, that there is a red object behind my back. Nor is it true when I visualize such an object: not even when I visualize it, with eyes open, in front of me. As we know from the classic experiments of Külpe and Perky, it is possible for mental imagery to be mistaken for perception, and conversely.⁷ This, however, is precisely a mistake. This possibility does not affect the intrinsic differences between the kinds of state in question.⁸ By contrast, hallucinations are sensory states. Here, however attentive and mentally focused we are, we ineluctably seem to be presented with instances of perceptible qualities. Hallucinations are, therefore, at least on one level of classification, to be grouped with genuine perceptual experiences in contrast to all the others just mentioned. In particular, hallucination is not a matter of mental imagery being mistaken for perception. Perhaps most—perhaps all—actual cases of what are called hallucinations are in fact cases of visual imagery being mistaken for perceptual experience. However, as the term is understood here, and indeed widely within the philosophical community, hallucination is a possible kind of experience that is closer to perceptual experience than this.⁹ Paradigmatically for modern philosophy, a hallucination is the sort of state that could in principle be induced in a subject by stimulating a sense-organ in precisely the way it is stimulated when the subject actually perceives something. Whatever we may wish to say about the intrinsic nature of the resultant experience, it is not going to be a matter of having mental images. That is a different sort of psychological state, with a different aetiology.¹⁰

1

Did I just hear a click? Did I just see a flash? Did I just see something move? Sometimes we find ourselves asking such questions; and we may be quite unable to answer them with any conviction. If we do venture an answer, we may be wrong. There are, I contend, three possibilities in such a situation. First, I really did perceive something, but because it was barely detectable, at the very threshold of my discriminatory ability, I am not sure that I did. Second, I did not perceive anything, but, as we say,

⁶ There is a question whether bodily sensations should be regarded as 'sensory' in nature. I shall simply skirt this issue, since we are here concerned with external perceptions and states that are not subjectively discriminable from them.

⁷ The findings of Perky in particular have been the subject of considerable controversy. The basic claim that it is possible to mistake perceived objects for mental images is not in serious dispute, however. See, for example Segal (1972).

⁸ This claim does not, in case it needs saying, prejudice the issue against the extreme disjunctivists. The latter do not claim that states other than hallucinations cannot in fact be mistaken for perceptions, but only that any such state can in principle be discriminated from perceptions.

⁹ It would be to no purpose to suggest that the negative epistemic criterion is put forward as a sufficient condition for being a hallucination in the wider sense that perhaps corresponds to everyday usage. For then the question arises of what to say about hallucinations in the narrower 'philosophical' sense.

¹⁰ For a little more on this issue, see Smith (2002: 201–3).

I simply ‘thought’ that I did—or that I may have. A third possibility, though one that would be quite exceptional, is that I very briefly hallucinated a click or flash or movement. In this last case, although I did not actually perceive anything, I did have a very brief sensory experience. By contrast, in the second case I had no (relevant) sensory experience at all, however brief. In such a case I am unsure, or mistaken, even about the character of my own experience. Such everyday occurrences illustrate something that is found in psychological experiments involving a tachistoscope: a device that flashes images on a screen for very brief periods. Subjects are requested to report what they see, if anything. What is relevant for us in such experiments is the so-called ‘false positive’ response that such subjects sometimes give: saying (and believing) that they saw something when nothing actually flashed on the screen. Here, again, it would be unusual for such subjects to hallucinate a flash. If they did, they would make only one of two possible mistakes. Although they would be wrong in thinking that a flash had actually occurred, they would be right in thinking they did have a brief visual experience as of a flash. In the usual false positive case, however, they would be wrong on both counts. There are here three psychological states that need to be distinguished from each other: having a momentary perception, having a momentary hallucination, and having neither, but merely ‘thinking’ that one has, or may have, just perceived something. Disjunctivists who adhere to the negative epistemic criterion are incapable of distinguishing between the latter two, since they are both reflectively indiscriminable from the first, perceptual case.

Now consider a very different sort of case: dreaming. As Descartes famously pointed out, we cannot, when we are dreaming, and simply on the basis of the experience itself, tell that we are not awake and perceiving. Dreams are not, however, hallucinations that we have while asleep: they are non-sensory episodes that merely seem sensory to the subject. In their actual nature they are akin to the sort of mental imagery that is involved in daydreaming and visualizing. We get so caught up in the imaginary scene that we cannot tell it apart from reality; but the experiences themselves remain merely imagistic, not sensory, in character. They are not the sorts of experiences that we would have if our brains were stimulated in just the way that they are when we actually see things of the sort that we seem to see in dreams. Or so I think. However, even if I am wrong about this, and dreams are indeed sensory and not merely imagistic in nature, it is surely *possible* (metaphysically possible) that there should be dreams of the kind I have described. The negative epistemic criterion would rule out even the possibility.

Finally, consider the experiences of subjects who are hypnotically induced to ‘see’ and ‘hear’ things that are not there. I think it likely, and certainly possible, that although it seems to these subjects that they are seeing or hearing certain things, they are not actually having relevant visual or auditory experiences at all. Although such experiences are commonly referred to in the literature as ‘hallucinations’, they are, I suggest, not hallucinations in the narrow sense in which I and most people in the philosophical community employ the term. They are not the sort of experiences a subject would have whose brain was stimulated in exactly the way that happens when a subject really does see and hear things. Nevertheless, these experiences are, like genuine hallucinations, wholly indiscriminable for their subjects from perceptions.

Similar remarks apply, I also suggest, to the voices that schizophrenics sometimes seem to hear.

Disjunctivists commonly write of perception and hallucination in terms of the former being the 'good' and the latter being the 'bad' disjunct. What the previous cases show is the need to recognize what we may call a 'very bad' disjunct. There is a clear sense in which perceptions and hallucinations, in contrast to such 'very bad' cases, fall into one kind of experiential state: the sensory kind. For any disjunctivist, this sensory kind will not be a fundamental kind, since it will comprise the fundamentally different states that are perceptions and hallucinations. Nevertheless, this class of experiences needs to be delimited and distinguished from the class of very bad experiences. The class of 'sensory experiences' needs to be specified. My charge against disjunctivism is that it cannot do this if it remains wedded to the negative epistemic criterion. In virtue of meeting the negative epistemic criterion very bad cases are lumped together with hallucinations, and their non-sensory character is not recognized. Alternatively, the sensory character of hallucinations is not recognized, since they are characterized in a way that applies to the non-sensory very bad cases.

2

How might a disjunctivist who accepts the negative epistemic criterion respond to the above charge? M. G. F. Martin has recently offered a response (in the interests of upholding extreme disjunctivism); and it seems to be the only response that is possible.¹¹ Before turning to the substance of his reply, it is perhaps worth setting aside a possible misconstrual of my very bad cases, if only because it is a misconstrual that partly informs Martin's own reply. At one point he writes of the "epistemological irrelevance" of my very bad cases. They are irrelevant, he suggests, because the subject in such cases "need not be rationally responding to how things sensorily seem to him" (Martin 2006: 389). He construes my very bad cases as merely involving what he calls "cognitive concomitants" and "cognitive consequences" of experience (Martin 2006: 374). In other words, my very bad cases are construed as simply a matter of subjects acquiring false (and irrational) beliefs about their experiences. The point is, however, that my very bad cases involve the occurrence of certain *experiences*: experiences that, I claim, are indiscriminable from perceptual experiences. A hypnotized subject's visual experience may, we may suppose, be constantly as of an empty room. When, however, he starts to 'see' a pig before him, he does not just start to have false beliefs that are an irrational response to this continuing perceptual experience of an empty room; he begins to experience in a certain way. Dreams, merely apparent momentary experiences, and the schizophrenic's voices are all equally experiential episodes. Because

¹¹ I have not published my criticisms of the negative epistemic criterion or of extreme disjunctivism before. Martin is explicitly responding to some brief remarks of mine (Smith 2002: 224) that were not explicitly directed against these theses, but which can reasonably be interpreted in this light. I did, however, explicitly argue against the negative epistemic criterion and extreme disjunctivism at a conference on disjunctivism in Frankfurt in 2004 at which Martin was present.

of this misconstrual of the very bad cases, Martin takes my criticism of the negative epistemic criterion to be that it allows mere beliefs into the class of sensory states. He then replies, in part, by insisting that there necessarily is “something it is like” to be in a state that meets the negative epistemic criterion (Martin 2006: 378). If there is, the suggestion seems to be, the criterion will be adequate to pick out the class of sensory experiences. This is not true. For any experience at all, even a ‘very bad’ one, there is, no doubt, something it is like to have it. But not every experience is sensory in nature: mental imaging and sudden realizations are not, for example. It is specifically sensory experiences that the negative epistemic criterion is meant to capture; but it does not, as the very bad cases show. That said, let us turn to the substance of Martin’s response.

This response invokes a distinction between two notions of indiscriminability. A subject may be unable to discriminate between two things because of some particular limitation that attaches to the subject or to the situation the subject is in. This is, we may say, a case of ‘local’ indiscriminability. It does not follow from this that the two things in question cannot be told apart from one another at all: that they are, as Martin terms it, ‘impersonally’ indiscriminable. To use one of Martin’s own examples: that I, with my somewhat poor eyesight, cannot spot a tailor’s ‘invisible mend’ to my jacket does not mean that the mend cannot be discerned—discriminated from the unmended fabric—at all. If someone with better eyesight than I can detect the mend, it is not invisible. Here we have a clear case of local indiscriminability (I can’t do it) and yet impersonal discriminability (someone else can). We can make out this contrast even if no human being is able to spot the mend. If some (perhaps merely possible) sighted creature could visually detect the mend, then the mend is not impersonally undetectable: it is not absolutely undetectable for sight as such. The visibility of the mend can also be made relative to a set of circumstances: the mend may not be visible to the casual glance, or in subdued lighting, and yet be visible impersonally, since it could be spotted on closer inspection or in a clearer light. Such notions of visibility and discriminability apply not just to individual objects and features, but also to kinds. We can sensibly ask whether this sort of mend is visible in either a local or an impersonal sense.

Martin’s proposed response to the challenge presented by my very bad cases is that they involve only the local sort of indiscriminability, whereas it is the notion of impersonal indiscriminability that the disjunctivist should employ in giving an account of hallucination. Only hallucinations and perceptions are impersonally indiscriminable from perception. My very bad cases, by contrast, are indiscriminable from perception only because of some local impediment.

3

This distinction between two notions of indiscriminability will not help in the face of my very bad cases, however, since they are absolutely, “impersonally”, indiscriminable from perceptions. There is *no possible* situation in which they could be discriminated: no possible ‘optimal’ situation, as we may put it. Consider dreams, for example. Dreamers, it is allowed, cannot tell that they are not perceiving real things.

But that is, the suggestion goes, simply because they are unfavourably placed to effect the discrimination. Someone better placed would be able to effect it. What, however, is the local impediment to reflection's discriminative power in this case? It is simply the fact that the subjects are asleep. So the suggestion would seem to be that dreams are (impersonally) discriminable from waking states of perception, because subjects who are dreaming could tell that they were dreaming *if they were awake*. But this, it may be thought, is absurd. Is it not analytic that if you are dreaming, you are asleep? In fact, the issue is not straightforward. For one thing, subjects under hypnosis can be induced to dream. 'Hypnotic dreams', as they are called, are had by subjects who are not asleep. (See, for example, Hilgard and Nowlis (1972) and Barrett (1979).) One could argue over whether these induced experiences are 'really' dreams; and in any case, the subjects in question are in a hypnotized, and hence non-optimal, condition. But even if there is some necessity that normally awake subjects do not dream, the necessity may be a relatively superficial one. The substantive question is whether someone who is awake and alert could have the very same kind of experience as a dreamer. Such waking experiences would, perhaps, not properly be regarded as dreams. If so, let them be 'dreams'. If it is indeed not possible for such experiences to occur during waking life, this has to be shown. It is not shown simply by pointing to some supposed analytic connection between dreaming and sleeping. What might it be like for an awake person to have an experience that was identical in character to a dream? Well, it might be somewhat like a daydream: something that we can certainly distinguish from a perception while we are awake. A sleeping person who has such an experience, however, is unable to make this discrimination.

Certain facts tell against such a suggestion, however. The epistemic problem with typical dreaming subjects is what has been called their 'single-mindedness' (Rechtschaffen 1978). We are generally 'locked into' our dreams in a way that allows little or no critical, reflective distance between us and the experience. However, although 'single-mindedness' does generally characterize the dreaming subject, it does not do so universally. There is considerable evidence that it is possible to be in the sort of reflective, critical frame of mind that is characteristic of being awake while one is asleep and dreaming. Indeed, because of the absence of 'single-mindedness', it is possible for dreamers to be aware that they are dreaming. Such dreams are termed 'lucid dreams'. This, however, may be thought only to strengthen the case for the negative epistemic criterion. For now it may seem that all that needs to be claimed is that it is possible that any dream should have been, or should become, lucid. We should not, then, have to rely on the perhaps contestable possibility of having 'dreams' while awake. In fact, however, the reports of lucid dreamers lend no support to the epistemic criterion. Moreover, they cast doubt on the suggestion that an awake subject would both be able to have 'dreams' and discriminate them from perceptions.

Lucid dreamers generally become aware that they are dreaming because of something about the content of the dream. For example, subjects often come to such awareness when they realize that they are apparently witnessing something that is impossible or extremely improbable in reality. The following is typical: "Dreamed that my wife and I were in bed in a strange room . . . Suddenly she disappeared from my sight; she seemed to dissolve into a cloud and vanish. This told me that I was

dreaming" (cited in Green 1968: 32). Lucid dreams are often preceded by a period of so-called 'pre-lucid' dreaming, when subjects merely wonder whether they may be dreaming. Manuals on lucid dreaming direct subjects in such situations to search for some such telltale impossibility, so as to bring about the realization they are indeed dreaming, and thereby effect the transition to fully 'lucid' dreaming.¹² One subject reports on his test as follows: "At one point I'm sitting on a bed. Have two toothbrushes. Wonder if this is a dream. Easily pass them through one another and realise that it is" (quoted in Green and McCreery 1994: 20). Sometimes, also, the familiarity of a recurrent dream suggests to subjects that they are dreaming once again (Green 1968: 45). Such ways of determining that one is dreaming are, however, not at all what defenders of the negative epistemic criterion mean by discriminating an experience from perception 'through reflection' or 'just on the basis of the experience'. The latter must be restricted to reflection on the *experiential character* of the experience. That the notion of discrimination by reflection must be restricted in this way is clear, because otherwise the same indicators that lucid dreamers rely upon could be used to tell hallucinations apart from perceptions. If I am sure that I am awake, and I hallucinate my wife disappearing in a cloud, I could come to realize that I am hallucinating, because I believe that such a thing cannot really happen. And a recurrent hallucination could be recognized as such.

When we turn to what lucid dreamers say specifically about the experiential character of their lucid dreams, what we find repeatedly stressed is the perception-like clarity and lifelikeness of such experiences. The following is typical: "When I see the blue sky in a lucid dream, I see it as clearly, as brightly, as consciously as I now see this paper" (quoted in Green and McCreery 1994).¹³ Also significant are the so-called 'pre-lucid' dreams that do not develop into lucid dreams: that are not followed, that is to say, by a realization that one is indeed dreaming. Pre-lucid dreams feature the same ability to reflect on the experience as do lucid dreams; and very often, in the absence of any telltale signs, and after the closest reflective scrutiny, pre-lucid dreamers resolve their doubt as to whether they are dreaming or not by deciding that they are definitely awake and perceiving!

On the other hand, not all lucid dreams are triggered by some impossible or familiar element in the content of the dream; nor are they always preceded by pre-lucid dreams. Sometimes, it seems, lucidity simply dawns on the subject. One investigator reports that out of 100 subjects he interviewed about how they determined that they were dreaming, 16 per cent said that they 'just knew' that they were dreaming (Hearne 1982: 14). So a defender of the negative epistemic criterion might suggest that such a realization, not based on any irrelevant telltale content, is always a possibility for a dreamer. Even if it is, this will not provide support for the criterion,

¹² Although we are probably now stuck with the terminology of 'lucid' and 'pre-lucid', it is somewhat misleading. Lucidity connotes reflective alertness and clarity, which is present in both the 'pre-lucid' and the 'lucid' stages. In the technical sense, a dream is 'lucid' only if, in addition, the subject is aware that he or she is dreaming.

¹³ Sometimes, it is true, a lucid dreamer reports that the dream experience was unlike a perceptual experience (for example, Green 1968: 35). This, however, is unusual. It is the more typical cases that constitute a counter-example to the negative epistemic criterion.

since this would, yet again, not be a relevant form of discrimination. Consider an analogy. Suppose that red objects come in two varieties, depending on whether the amount of invisible ultra-violet light they reflect attains a certain level or not; and suppose that there are certain subjects who are sensitive to the amount of this light in such a way that they 'just know' which sort of red object they are seeing. Both sorts of object may, however, *look* exactly the same to these subjects. It may be that ultra-violet light is truly invisible to them, and the different beliefs and discriminations are caused by some 'sub-personal' mechanism that operates independently of the character of the subjects' visual experiences.¹⁴ Since the two sorts of red object look absolutely the same to these subjects, they cannot visualize two red objects as being different in the relevant respect, since it is not a visible difference; and were they to hallucinate a red object, they could not hallucinate one kind of red object rather than another. This story is surely intelligible; and yet it would not be for anyone who accepts the negative epistemic criterion if these subjects' ability to discriminate between their experiences of one sort of red object and their experiences of the other counted as a relevant sort of discrimination. For if it did, the fact that both sorts of red objects look exactly the same to them, that the visual character of their respective experiences is identical, could not be represented. This shows that the notion of indiscriminability as it features in the negative epistemic criterion must be restricted yet further. We now see that not only must such reflection focus only on the experiential character of an experience; it must, if the experience in question is visual, focus on the purely visual, presentative aspects of the experience, for these alone will correspond to the genuinely sensory aspects of a perceptual situation.¹⁵ If, however, this is the notion of indiscriminability that the criterion must involve, then lucid dreams pass the test for being sensory states. For, as I have pointed out, lucid dreamers constantly stress that as far as the purely visual aspects of their dreams are concerned, they are strikingly lifelike and 'perceptual' in character; and the same is true of the other sense modalities when they feature in lucid dreams. The fact that lucid dreamers exhibit the focused alertness that is characteristic of an awake state, and yet still cannot, in the relevant way, discriminate their experiences from perceptions, undermines the suggestion that an awake person who had a 'dream' experience would be able to tell that the experience was not a perception.

If, however, we take at face value what most lucid dreamers say about their experiences, perhaps we should doubt that their experiences are, after all, non-sensory in nature. Perhaps lucid dreams, unlike ordinary dreams, really are hallucinations that we have while asleep. If this were the case, the reports of lucid dreamers would not count against the suggestion that an awake person could have an experience that was experientially identical to an *ordinary* dream, and tell it apart from perception. If, however, it is impossible for a lucid sleeper to have an ordinary, non-sensory

¹⁴ Indeed, we can suppose that blue objects similarly divide into two for these subjects, so that they find some indefinable 'something' in common between certain red and certain blue objects.

¹⁵ Even when our subjects 'just know' that they are seeing one sort of red object, this will not involve them being apparently confronted by an instance of a perceptible quality that defines the kind of red object in question. The very language of 'just knowing' rules this out.

dream—that is, an experience that is of the same kind as an ordinary dream, differing only with respect to the reflective clarity with which it is appreciated—but only a hallucination, why on earth should it be possible for an awake person to have a ‘dream’? Lucid dreamers, after all, are closer to ordinary dreamers than people who are awake. At least they are asleep.

4

One way for a defender of the negative epistemic criterion to attempt to rebut the suggestion that the reports of lucid dreamers indicate that dreams are indiscriminable from perceptual experiences even for focused, attentive consciousness is to claim that lucid dreamers, despite what they themselves report, are not really as reflectively alert and attentive as an awake person can be. If this were so, their reports would not undermine the original suggestion that an awake person would be able to discriminate a ‘dream’ from a perception. Given the reports, I regard this suggestion as at best speculative. The subjects of such experiences give every indication of an attentive, reflective interest in their experiences while they are occurring that is as great as anything that awake people achieve. One should not be over-impressed by the fact that lucid dreamers are fast asleep and hence out of touch with ‘reality’. For there is strong evidence that lucid dreamers, while asleep and dreaming, can communicate, through prearranged signals made by eye movements, with outside observers—a technique originally devised by Keith Hearne. (See Schatzman et al. 1988.) Moreover, as several writers have pointed out, lucid dreams are similar in many respects to another sort of experience, which can be had by fully awake subjects: out-of-body experiences. I regard it as a genuine possibility that these experiences are also not genuinely sensory states (hallucinations), but ‘dreams’, or something analogous to them; and yet subjects once again attest to the ‘perceptual reality’ of such experiences. Indeed, the subjects of such experiences actually take themselves to be perceiving the world, though from a point of view other than that occupied by their body. Such experiences furnish yet another type of very bad case, and one that directly casts doubt on the suggestion that an awake person would be able to discriminate a ‘dream’ from a perception. The reply to this observation will doubtless be that, although the subjects of out-of-body experiences can be awake, they are still not in an ‘optimal’ state for appreciating the nature of their experiences. A similar line will also, no doubt, be taken in connection with various other very bad cases that can befall an awake subject, such as the apparent perceptions of hypnotized and schizophrenic subjects.¹⁶

¹⁶ Schizophrenia is obviously a non-optimal condition. The hypnotic state is commonly regarded as involving something like what Ronald Shor (1959) characterized as a diminution in ‘the generalized reality-orientation’ that governs our normal cognitive dealings with the world. Perhaps something similar holds of out-of-body experiences. Hypnotized subjects are also able to accept the reality of their induced ‘perceptions’ together with that of their conflicting real perceptions: a phenomenon that Martin Orne (1959) dubbed ‘trance logic’. This would seem to indicate a less than optimal condition.

Although the case of lucid dreams remains problematic, let us suppose that, necessarily, all subjects of all the very bad cases I have just discussed are in a less-than-optimal situation for appreciating their experiences. Even if this is true, the upholder of the negative epistemic criterion now faces the worrying possibility that there should be kinds of non-sensory experience that can be had only under certain non-optimal conditions that preclude relevant discrimination. If this were the case, it would be absolutely impossible for such experiences to be discriminated from perceptual experiences, and the negative epistemic criterion will wrongly count them as sensory in nature. It is not at all clear, after all, that a normal awake person could indeed have a 'dream'. I suggested earlier that dreams are perhaps something like daydreams. Could, they, however, be *identical* in character to daydreams—or to any other experience that an optimal subject could have? That such a subject could have an experience that was merely something like a dream experience is hardly to the point. To take another example, consider the induced experiences of hypnotized subjects. Is it at all clear that such experiences can be had by a subject who is not hypnotized, and who is otherwise normal? If all that happens in such cases is that the subjects simply imagine or visualize certain objects being present while being unaware that this is what they are doing, then we can certainly suppose that if they did the same thing when not hypnotized, they would realize it. It is somewhat doubtful, however, that hypnotized subjects' experiences consist in their having mental images. The formerly prevalent view that the hypnotic state involves an 'absorption' in imagining what has been suggested to the subject is now in serious doubt. (See, for example, Zamansky (1977) and Bowers (1992).) Moreover, scans of the brains of subjects who are 'seeing' things under hypnosis do not support the idea that such experiences are just a matter of mental imagery. (See for example, Szechtman et al. (1998) and Kosslyn et al. (2000).) So perhaps the hypnotic state can produce experiences of a kind that are simply not to be had without being in that state. A similar possibility exists for the voices that are 'heard' by the schizophrenic, and for out-of-body experiences.¹⁷ I do not see how such possibilities can be discounted a priori, as they must be by a defender of the negative epistemic criterion.

5

Let us, however, suppose (over-generously, in my view) that the upholder of the negative epistemic criterion can provide an adequate response to the sorts of very bad cases we have just been considering. There remains one type of very bad case where the claim that it could be both experienced optimally and discriminated from perception is demonstrably untenable: the momentary or 'tachistoscopic' case. Moreover, this type of case presents a peculiar and insurmountable problem for extreme disjunctivism. In this section I shall explain this peculiar difficulty, and in the next explain why at least certain possible very bad tachistoscopic cases refute even the weaker negative epistemic criterion.

¹⁷ Here, too, certain brain scanning results are suggestive. See Cleghorn et al. (1992).

Suppose that it seems to you that you have, or may have, just seen a yellow patch flash very briefly on a screen that you are looking at; but, because it was so quick, you are unsure. Let us also suppose that nothing actually did flash on the screen, and that you did not hallucinate such a flash. In this sort of situation you just cannot tell, by reflection on your experience, that you did not see a flash. It may be thought that Martin's distinction between two sorts of discriminability applies straightforwardly in such a case. For here at least, surely, we are dealing with but a local impediment to reflective knowledge. We are unsure about the character of our visual experience when it comes to very brief experiences simply because of their brevity. So we simply think away this local limitation. The most obvious way of doing this would be to lift the temporal constraint. By analogy with the tailor's mend that is not visible to the casual glance, but is so on a closer inspection, we can suppose that any momentary perception or hallucination we may have could have been longer in duration than it actually was. If it had been (sufficiently) longer, we would have been able to tell (other things being optimal) that we were at least in a genuine sensory state. For, surely, it is not possible for it to seem to us that we are in a sensory state for any prolonged period when actually we are not, unless we are deranged or dissociated or otherwise not properly placed to make the discrimination—in which case, we simply think away such local imperfections. There is also another way in which a very bad tachistoscopic case could be referred to an optimal situation. By analogy with Martin's example where the tailor's mend is not visible to humans with their limited visual acuity, but is yet impersonally visible because a subject with more acute vision could detect it, we can suppose that, although for us the very bad momentary case is indiscriminable from perception, this is but a human limitation. There is a possible subject with more acute powers who is not over-taxed by such situations, and who would be able to make the discrimination. Let us call him 'Superman'. The defender of the negative epistemic criterion can then say that a very bad momentary case is impersonally discriminable from any perception because Superman could (indeed, would) so discriminate it. Even he, however, would not be able to discriminate a hallucination in this way. These are the only two ways in which a momentary very bad situation can be rendered optimal. For the only thing about such cases that makes them non-optimal is their brevity. So, either we consider a situation where such brevity is absent, or we consider a subject for whom brevity is no problem. There are no other possibilities.

Before proceeding, we should consider in a little more detail the nature of momentary very bad cases. I have so far characterized them in terms of its seeming to subjects that they *have* (or may *have*) seen something. It is, however, not such retrospective seemings that are supposed to be indiscriminable from perceptions. A psychological state that concerned only the past could hardly be mistaken for a perception (or a hallucination), since the latter concerns what is apparently presented to the subject *now*. When it seems to me that a flash has just occurred, or that I have just seen a flash, it is not this present seeming that I take to be a perception of a flash. The perception, if there was one, was in the immediate past, and seems to be. The point, of course, is that when things happen so fast, the experience we are interested in has gone before we are in a position fully to appreciate it—even its existence. If there was indeed a fleeting sensory experience, I am reflectively aware of it only as just having

been. When, as in a very bad case, there was no such prior sensory experience, what exactly happens?

One possibility, perhaps, is that a very bad momentary case involves the very brief occurrence of a mental image. No doubt we could mistake such a thing for an actual perception of the thing imaged. If, however, the mental image were to persist, we would surely know, other things being optimal, that it was a mere image, and no perception; and Superman would know straight off that he was having merely a mental image. Extreme disjunctivism is, however, incapable of exploiting these facts, if they are facts, without introducing a vicious circularity into the overall position.

Consider, first, the appeal to an extension. When we consider such an extension, we are considering something different from the postulated momentary experience itself whose discriminability from perceptual experience is at issue. Our question is whether a certain momentary experience could possibly be told apart from any perception, not whether some other (longer) experience could be. So the case is quite disanalogous to that of the tailor's mend that is not visible to a cursory glance, though it is on a more extended inspection. In the latter case we move from a limited to an optimal situation in relation to the same object. Where the object of discrimination is itself a brief experience, however, the consideration of a more extended experience involves a change of object. Why should these other objects, these other, longer experiences, be thought to be at all relevant? The answer, of course, is that the longer experiences are supposed to be relevant because, being extensions, they include the original momentary experiences as proper parts. Such an extension is going to be relevant, however, only if it is what we may call a 'homogeneous' extension: one that preserves the character of the initial experience. An extension should simply involve *more of the same*, not a continuation by a different sort of experience. This may seem obvious, but it presents the extreme disjunctivist with an insuperable problem when it comes to giving an account of momentary hallucinations that have the same proximal causes as possible momentary perceptions. My optic nerve could theoretically be stimulated in the way that it is when I see a yellow square, but so briefly that I am unsure whether I had a visual experience at all. Even if I incline to the view that I did, I may be wholly unable to specify the precise phenomenal character of the experience. I think something flashed on the screen; but as to whether it was a circle, or a triangle, or a square, I am completely at a loss. I think that what I saw (if I did) was yellow; but it may have been green. In other words, this hallucination is indiscriminable for me from a perception of a yellow triangle, but also indiscriminable from a perception of a yellow circle, a green square, and so on. If we are to make sense of a homogeneous extension of this momentary experience, one that preserves its phenomenal character, there must be a fact about what its phenomenal character is. Since it is a causally matching hallucination, one that, therefore, according to the extreme disjunctivist, lacks any distinctive positive psychological feature that determines or constitutes its phenomenal character, this character must be determined solely in terms of subjective discriminability. Since I am not in a position to effect the relevant discriminations, we are now considering an 'optimization' of my situation as a result of homogeneously extending my experience. But now we have a circularity in the account. A homogeneous extension presupposes a fact of the matter concerning the

phenomenal character of the experience to be extended, when this fact itself consists of possibilities of discrimination in relation to just such a homogeneous extension. The optimal situation can be specified only by relying on a notion of sameness of phenomenal kind that itself can be explicated only by reference to this optimal situation.

One possible response to this problem of circularity would be to hold that the phenomenal character of my momentary hallucination is in fact as indeterminate as my discriminatory powers. This would, however, render impossible the application of the notion of a homogeneous extension to the case. This is because the extended experience that has the momentary hallucination as its initial phase is not itself going to be phenomenally indeterminate, since such supposed indeterminacy arises simply from my inability to discriminate properly, and this inability arises simply from the brevity of the hallucination. An *extended* hallucination of a yellow square is, other things being optimal, going to be manifestly as of a yellow square. But this means that a homogeneous extension of a phenomenally indeterminate hallucination will be impossible, since the experience will have changed from being phenomenally indeterminate to being phenomenally determinate (or at least more so).

A circularity in the extreme disjunctivist's account of hallucinations will also arise if appeal is made to an ideal superhuman subject. The circularity this time does not arise in attempting to make sense of a homogeneous extension of a momentary experience, since superhuman reflective powers preclude the need for any such extension. The circularity arises, rather, in referring *our* momentary very bad experiences to *another* subject for adjudication. The extreme disjunctivist must be able to say of some very bad momentary experience that you may actually have that it—that very experience of yours—is impersonally discriminable from any perception. Clearly, however, it makes no sense to suppose that some other subject, such as Superman, should have that experience of yours. The most we can do is to suppose that he should have the same *kind* of experience as you. Once again, however, as soon as the issue of impersonal indiscriminability is applied to kinds of experiences a circularity will emerge in the extreme disjunctivist account when it comes to dealing with causally matching hallucinations. For when such indiscriminability is determined, as in the present case, by what Superman can and cannot discriminate, we shall have to say that a certain momentary experience is impersonally discriminable from any perception if and only if it is of a kind such that, if Superman were to have an experience of that kind, he would be able to discriminate it from any perceptual experience. The problem is that, according to extreme disjunctivism, causally matching hallucinations are assigned to a kind solely on the basis of their being impersonally indiscriminable from some perception, whereas we are now spelling out what impersonal indiscriminability amounts to by reference to what Superman can discriminate in relation to just such a kind of experience. Once again, therefore, the account is empty circular. The extreme disjunctivist's negative epistemic criterion cannot, therefore, operate at the level of types of experience, but only at the level of individual experiences.¹⁸ However,

¹⁸ Martin's own formulations of the criterion always make reference to individual experience. See, for example, Martin (2006: 364), where he emphasizes the point.

each of the only two ways in which momentary experiences can be referred to an optimal situation makes essential reference to a kind of experience.¹⁹

6

Let us now consider how the negative epistemic criterion stands up in the face of certain very bad momentary cases. Being weaker than extreme disjunctivism, it avoids the problems detailed in the previous section. And, indeed, the kind of very bad case considered in that section—one that involves a fleeting mental image being mistaken for a perception—poses no problem for the criterion. As I have said, if the image were homogeneously prolonged, doubtless we should be able to tell that it was a mere mental image; and Superman would be able to tell this, we may suppose, straight off. It is, however, not necessary for a very bad momentary case to feature a mental image in this way.²⁰ Indeed, I am not sure that the having of a mental image for *such* a brief moment is even possible. But let us suppose that it is, and that the occurrence of such a fleeting image may explain some very bad cases. There is, however, a good reason for supposing that mental images cannot explain all possible such cases. For if a momentary mental image is indeed possible, it is surely possible that I should think that I have just had such an image when I have not. We cannot, however, explain this mistake in terms of the occurrence of a mental image, for that would mean that we could never be in error in such cases: whenever it seems to us that we have just had a fleeting mental image, we would indeed have had one, because such an image is postulated to explain why it seems to us that we had one in the first place. If we must in this case explain a mistake about our own fleeting experience without postulating a mental image, why should it not be possible when apparent perceptual experience is in question?

A possible alternative sort of very bad momentary case, and the one I offer as a counter-example to the negative epistemic criterion, is that in such a case it suddenly seems to us that, say, a flash has (perhaps) occurred—and, if we are reflective beings, that we have (perhaps) just seen a flash—and that is the whole story.²¹ Here we would have a wholly groundless episode of retentional consciousness: one that offers us ‘a past that has never been present’, to employ a trope from recent continental philosophy. In this sort of momentary very bad case there is not an identifiable non-sensory experience that subjects mistake for a sensory experience. The subjects in such situations are even more radically mistaken about their own experience than that: they

¹⁹ This problem also arises for extreme disjunctivism in relation to the other very bad cases discussed earlier. Since any reference to a *type* of experience is ruled out, the extreme disjunctivist will have to claim, for example, not just that dreams or hypnotic experiences could be experienced while the subject is awake or not hypnotized, and be discriminated from perceptions, but that any individual dream or hypnotic experience could *itself* have been experienced while in an optimal condition. This is even less plausible than the original claim.

²⁰ One will typically have a mental image in retrospect, when one wonders if one has just perceived *this* sort of thing. The question is, however, whether a prior mental image lies behind this retrospective episode.

²¹ At least on the psychological level. There may be some neurological story to tell.

do not have a relevant experience at all, though it seems to them that they do, and that it is a sensory experience. What they mistake for a flash is the non-existent intentional object of a state of retentional consciousness.²² Despite the fact that there is, on this account, no very bad experience that is itself mistaken for a perceptual experience, these cases still challenge the adequacy of the negative epistemic criterion. For what they show is that there could be a stretch of experience that does not include a certain sort of sensory occurrence—but only an ungrounded ostensible retention of one—that is, as a whole, reflectively indiscriminable from a stretch of experience that includes such a sensory episode (and a retention properly grounded on it). This, I submit, is a characterization of one possible sort of very bad momentary case. Such retentional consciousness is not, perhaps, what upholders of the negative epistemic criterion have in mind when they speak of ‘reflection’. It is, however, the best that is available when it comes to momentary experiences. ‘What it is like’ to have such a momentary experience is exhausted by this retentional consciousness we have of it.

The important point now is that the above sort of very bad momentary case will remain indiscernible from a perceptual case even if it is extended, and even for Superman. Let us first consider the extension. Retentional consciousness itself can, of course, be extended, or persist. When it seems to us that we have just seen a flash, it typically continues to seem to us that we saw a flash back then. This kind of continuation is not relevant, however, for it will not facilitate discrimination. When such retentional consciousness is extended, it is simply that the ‘then’ of the putative experience recedes further and further into the past as our conscious life continues its course. This will in no way render what happened ‘back then’ in a very bad case—that is, nothing—any more discriminable from a perception than it was before. Indeed, as time passes, we become more and more removed from the putative experience itself, so that our grasp on what happened will, if anything, be weakened. Extending the period during which the subject did not perceive a flash, or even seem to, is also irrelevant. Indeed, we can suppose this period to be as long as we like in the original scenario. However long it is, if an ungrounded retention then occurs, the subject will have the sense that a flash has (or may have) just occurred. The only other candidate for being extended is the non-existent very brief experience that is the intentional object of the retention. There is only one way to make sense of this suggestion. We can suppose that the ungrounded retentional episode gave us the sense that we had been experiencing the flash for a longer time. ‘Very bad retentions’, if I can so term them, need not concern putative *momentary* experiences. Sometimes, when an unnoticed light in our visual field goes out, we can have a sense, for the first time, that we were seeing it for some time before. There is a ‘very bad’ version of such an experience, where there was no light, nor even a hallucination of a light,

²² I borrow the term ‘retention’ from Husserl, who distinguishes it from ‘recollection’. The latter is a mental act that (apparently) ‘revives’ a past experience that one has had. Retention, by contrast, is the (apparent) ‘grip’ that any phase of one’s experience has on the immediately preceding phase. Husserl himself thought, mistakenly in my view, that such retentional consciousness is infallible. For an account of these Husserlian matters, see Smith (2003: 86–100).

but only the sense that one has (or may have) gone out after shining for a while. Such backward extensions are, however, of no significance in the present connection, since we have no firmer an epistemic grip on them than we do in the momentary case. Moreover, not even Superman will be able to distinguish an ungrounded retention from a grounded one. The former, equally with the latter, is necessarily revealed to reflective consciousness as its seeming to one that one has just perceived something. Such a seeming is necessarily false; and even Superman cannot jump back in time and determine its falsity. The only way in which a superhuman subject can avoid the mistakes that we humans make is by not having such ungrounded retentions at all. But then we are not placing such an ideal subject in the same very bad situation that we can be in, and so he will be useless as an ideal discriminator of such situations from perceptual situations.

7

The negative epistemic criterion is invariably expressed in terms of the impossibility of a certain sort of knowledge. Martin's most precise formulation of the criterion, for example, explicitly employs the notion of its being "not possible to know through reflection" that an experience is not a perception, or a perception of a certain sort. The criterion adverts to "modal facts concerning the possibility or impossibility of certain knowledge" (Martin 2006: 364). If the criterion is interpreted in this way, it is, I have argued, not adequate. The distinction between local and impersonal impossibility is of no avail, since it is *absolutely* impossible for my very bad cases, or at least some of them, to be discriminated from perceptions. Perhaps, however, the negative epistemic criterion can be reformulated in such a way that it does not deal simply with 'modal facts'. The core notion, after all, is (impersonal) *discriminability*. Perhaps we should take this notion as basic, and not to be analysed in terms of possibilities for discrimination.²³ This modification is certainly worth taking seriously, since it is indeed possible for statements attributing dispositions and abilities to be true even when it is not possible for the disposition or ability to be manifested or exercised.

Suppose that a certain substance is soluble in water. This substance has a certain disposition—it is soluble—and water has a correlative power or ability—to dissolve the substance. These facts are compatible with the total impossibility of the substance dissolving in water, since there may be an explanation, compatible with the thing's solubility and water's power to dissolve it, of why it cannot dissolve. There are, in fact, two possible sorts of explanation. To illustrate the first, suppose that the substance and water exert an insurmountable repulsive force on one another, so that they cannot be brought into contact; or that if the substance is brought within a foot of any water, it explodes into molecular fragments, or causes the water to evaporate. In such cases, the situation—the 'relevant situation' as I shall call it—in which the disposition would be manifested and the correlative power exercised simply cannot

²³ Martin himself indicates that this may well be his position. See Martin (2006: 396, fn. 44), and, perhaps, the last few lines of section 6 of the same paper.

be realized. A variant of this first type of explanation is where the relevant situation is in some sense realized, but not adequately. Even when soluble substances do dissolve when put in water, they do not do so immediately. A water-soluble substance might be such that something inevitably happens to it when put in water before it has a chance to dissolve: something that pre-empts dissolution. The substance might, for example, cause any water it comes into contact with to evaporate, a reaction that always and necessarily takes place before dissolution has a chance to start. Another possibility is that a soluble substance should only be able to exist for such a brief moment that it has not time to dissolve, even if it comes into existence surrounded by water. The second sort of explanation of the necessary non-manifestability of a disposition is furnished by C. B. Martin's 'finkish dispositions' (Martin 1994). Here the idea is that it is possible that bringing something with a certain disposition into the situation relevant to its being manifested may cause it to lose that disposition. A substance can be water-soluble, for example, even though it necessarily loses this disposition whenever it is brought within a foot of any water, or immersed in water. In a water-free environment this substance is actually water-soluble, and water has the ability to dissolve it, even though it is quite impossible for it to dissolve in water and for water to dissolve it. Correlatively, water itself may have a 'finkish' ability in relation to this substance: namely, one that disappears upon the approach of the substance, the latter remaining unchanged and, therefore, retaining its own disposition to dissolve. (See Lewis 1997: 144–5.) These are the only possible explanations of how it can be impossible for an object with a certain disposition or ability to manifest or actualize it.

The application of the preceding remarks to our own concern with discriminability is straightforward. For we can say that something is discriminable for reflection if some (perhaps merely possible) subject has the power or ability to discriminate it. It may true that subjects—even actual subjects—have this ability even when it is absolutely impossible for it to be exercised. This will be true if there is an explanation of such impossibility. The important point now is that some such explanation *must* be available whenever it is true that an ability (or a disposition) is possessed but cannot possibly be displayed. This requirement operates on two levels. First, in the absence of any such explanation, it will be entirely unjustified to claim that a disposition or ability is indeed present. In the case in hand, disjunctivists who employ the modified negative epistemic criterion claim that my very bad cases are impersonally discriminable from any perception, even though they cannot possibly be so discriminated. But why, in the absence of an explanation for this impossibility, should we believe any such thing? The simplest way to see the unjustifiability here is to turn the tables on such disjunctivists. They claim that a hallucination is impersonally indiscriminable from some perception. I turn round and claim that every such hallucination is, in fact, impersonally discriminable from any perception; it is just that they cannot possibly be so discriminated. In the absence of an explanation, the defenders of the negative epistemic criterion would quite rightly be unimpressed by this suggestion.

Secondly, it is unintelligible that there should be an unmanifestable disposition or ability without an explanation of its unmanifestability. There is an a priori link between (concepts of) dispositions and abilities on the one hand and manifestations

on the other. Indeed, the latter are, as Aristotle stressed, conceptually prior. I can understand what it is for something to dissolve without possessing the concept of solubility, or the power to dissolve something; but not conversely. Similarly, I can understand what it is for something to be discriminated without possessing the concept of discriminability; but not conversely. My very understanding of a disposition or ability presupposes and depends upon my understanding of a kind of occurrence that is its manifestation. Moreover, dispositions and abilities are differentiated from one another by differentiating between their manifestations. Solubility differs from the disposition to explode when immersed in water only because dissolving differs from exploding. Since, however, a disposition or an ability can be possessed by a thing without it ever being displayed, this *a priori* link can be to nothing more than the *possibility* of such a display. Even this, as we have seen, is too strong a link. However, without the requirement of some explanation of the impossibility of a display, where this is indeed impossible, the conceptual link would be entirely cut. To say that something is soluble in water, but that it cannot dissolve in water, or that water has the power to dissolve a certain substance, but that it cannot dissolve it, and to be unable to say more, would be incomprehensible. And if there is actually nothing more to be said, such statements are simply false. In the absence of some such explanatory facts, there is no difference, when a suitable manifestation is not possible, between a thing's possessing a certain disposition or ability and its not.

Our final question, therefore, is whether either of the two possible forms of explanation outlined above apply to the impossibility of my 'very bad' cases being discriminated from perceptual experiences, despite their supposed discriminability. We can immediately rule out the suggestion that the very bad cases involve finkish abilities, or abilities in relation to finkishly disposed objects. The suggestion that reflection loses its own discriminating power when it comes to very bad cases, or that very bad cases cease to be discriminable from perception when they are reflectively experienced, is a clear non-starter. It may, however, seem plausible that the other possible explanation, in terms of the relevant situation not being realizable, or not adequately realizable, does apply to at least a number of my very bad cases. Even if it is impossible, for example, for subjects in an optimal, or even normal, condition to have the sorts of experiences that people who are under hypnosis, or who suffer from schizophrenia, can have, this does not count against the power of reflection as such to discriminate such experiences from perceptual experiences. It is simply that the normal discriminating power of reflection cannot have an opportunity to be applied to such experiences.

So the modified negative epistemic criterion may be thought to work. It is at least a great improvement on the original 'modal' formulation. It is not, however, a version of the criterion that the extreme disjunctivist can adopt. For we have seen that abilities and dispositions can be possessed in the absence of any possible manifestation only when an explanation of this absence is possible. The required explanation is itself possible and intelligible, however, only by virtue of the possessors of such dispositions or abilities having an *intrinsic nature*. It is only by reference to such a nature that the explanations themselves can work. Why, for example, do we suppose

that a certain substance that cannot possibly dissolve in water is yet water-soluble? Only because we suppose it to have a nature in virtue of which dissolution follows immersion in water in certain specified circumstances according to the laws of nature. In the absence of a belief in such a nature, the claim that a disposition is present in such a case would not only be unjustified, it would be unintelligible. And in the absence of such a nature, the claim would be false. Compare a substance that is insoluble in water with one that is soluble but cannot possibly dissolve in any water for one or other of the two reasons outlined above. Given that the facts, even the 'modal facts', concerning their dissolution in water are identical, the difference between them must consist in some actual difference in their intrinsic nature. Or consider a soluble substance that cannot possibly dissolve in water *becoming insoluble*. We can make sense of this event only in terms of an actual change in the thing's nature. Extreme disjunctivists can, therefore, avail themselves of the reformulated negative epistemic criterion only if they attribute a positive nature to the states to which the criterion applies—a nature that underwrites and explains the modal facts concerning discrimination. This, however, is just what the extreme disjunctivist is unable to do with one class of experiences to which the criterion is to be applied: hallucinations with the same proximal cause as some possible perception. These states supposedly *have no positive nature*. Or, as Martin puts it, their "only positive mental characteristics are negative epistemological ones—that they cannot be told apart by the subject from veridical perception" (2004: 73–4). Leaving aside the puzzling suggestion that *positive* characteristics are *negative* epistemological ones, such supposedly positive characteristics cannot amount to an intrinsic nature that serves to ground a difference between a discriminable state and an indiscriminable one when neither can possibly be discriminated. There is, Martin writes, "only a negative characterisation of causally matching hallucinatory experience: it is nothing but a situation which could not be told apart from veridical perception" (Martin 2004: 72). There are, in the case of such hallucinations, Martin seems to suggest, no "relevant non-epistemological mental features which act as a ground to the facts of indiscriminability" (Martin 2004: 82). Such states have, he implies, no "non-relational" features (Martin 2004: 46 and 52). It is this very denial, indeed, which even Martin himself characterizes as "striking and implausible" and even as seeming "incredible" (Martin 2004: 38 and 85), that defines extreme disjunctivism and distinguishes it from a simple acceptance of the negative epistemic criterion.

There are, I should say, certain passages in Martin's writings that may suggest that his position is less extreme than I have been taking it to be. Such passages may suggest that what he is denying is not that the hallucinations in question have no positive nature at all, but only that they have no positive nature that the subject can detect. (See, for example, Martin (2004: 82).) On this view, such hallucinations may have a psychological nature that grounds their indiscriminability from perceptions; what is denied is that it is the subject's *awareness* of that nature that serves as this ground. In fact, however, Martin needs the stronger position. This is because the reason for embracing extreme disjunctivism is to allow us to say that veridical perceptions and causally matching hallucinations have the same psychological nature compatibly with

giving a disjunctivist (and, in Martin's case, a naïve realist) construal of veridical perception as having a distinctive and positive psychological nature.²⁴ The worry with allowing such a commonality of nature—a commonality that Martin believes, quite rightly in my view, is forced on us by a certain version of the 'argument from hallucination'—is that it will undercut the explanatory power of the positive psychological nature possessed by veridical perceptions. For veridical perceptions themselves are, of course, subjectively indiscriminable from veridical perceptions of a certain sort. Here the indiscriminability is grounded on a positive psychological nature. It is precisely because a veridical perception of, say, a red ball has a positive psychological nature (in virtue of which, such a kind of experience cannot be had in the absence of a red ball) that it is indiscriminable from any other subjectively identical veridical perception of a red ball—since all of the latter, too, will have just such a nature. If, however, we were to attribute a positive psychological nature, whether detectable by the subject or not, to a causally matching hallucination of a red ball as the ground of its indiscriminability from a veridical perception of a red ball, this nature, given the commonality mentioned above, will have to be attributed to the veridical perceptions also as the ground of *their* indiscriminability from veridical perception of a red ball. Since, however, this nature can be possessed by an experience when a red ball is not being perceived (since hallucinations have it), attributing such a nature to veridical perceptions would undercut the explanation of their indiscriminability from perceptions of a red ball in terms of their own positive psychological nature. It is irrelevant whether the nature in question is detectable by the subject or not. Whereas the indiscriminability of veridical perception from veridical perception is grounded on a positive psychological nature—so that, as Martin says, such indiscriminability does not constitute "the most fundamental kind that the perceptual event is"—the indiscriminability of at least a causally matching hallucination must be grounded on *nothing* psychological at all, so that such indiscriminability is indeed "*faute de mieux* . . . the fundamental mental character of the event" (Martin 2004: 60–1 and 72).

It is true that extreme disjunctivists cannot be supposing that hallucinations have no positive nature whatsoever. Nothing has no nature whatsoever. Their claim is only that causally matching hallucinations have no positive *psychological* nature. They will have some other nature, however: perhaps a neurological one. Disjunctivists in general, however, refuse to descend to this neurological level in giving their accounts of psychological states. Indeed, their very position depends upon denying that the psychological, at least in this area, supervenes on the neurological. According to the disjunctivist, a perception and a matching hallucination may involve the very same kind of brain state, and yet they are psychologically different kinds of state. Perhaps, however, the nature of the type of hallucination in question could be spelled out not in 'narrowly' neurological terms, but in more 'broadly' physical terms that concern both the individual perceiving subject and the environment.²⁵ And perhaps disjunctivism can sit comfortably with a supervenience of the psychological on the physical thus

²⁴ In Martin's particular version of disjunctivism—as in others, such as Paul Snowdon's—this nature is explained in terms of the perceived object being a *constituent* of the veridical experience.

²⁵ Thanks to Adrian Haddock for pointing out to me the need to address this.

broadly conceived. However, leaving aside the rather surprising finding that an analysis of sensory consciousness entails a form of physicalism, what would be required to answer the present problem is something much stronger than such supervenience. Since what is required is an account of a certain *kind* of psychological state's *nature* in physical terms, nothing short of a *reduction* of the psychological to the physical will serve. Martin for one, I think, would find this unattractive. And for good reason.

8

I wrote in the previous section that the negative epistemic criterion for hallucination and sensory states “may be thought to work” if it treats discriminability as not entailing facts about possible discrimination. In fact, however, not only is there a residual doubt about lucid dreams (since the lucidity of the subjects perhaps already constitutes an optimal condition for discrimination), there remains one type of non-optimal very bad case that is recalcitrant to employment of the criterion even as recently reinterpreted: the momentary very bad case that consists in the occurrence of an ungrounded retention. In such a case there is no possible explanation at all of why reflection is unable to come to grips adequately with the experiences involved: of why there is no possible discrimination of such cases from perceptions despite their supposed discriminability. It is not just that such cases cannot possibly be discriminated from perceptions: they are simply *indiscriminable* from perceptions. Here there is nothing more for reflection to discover, even *per impossibile*.

REFERENCES

- Barrett, D. (1979) ‘The Hypnotic Dream: Its Relation to Nocturnal Dreams and Waking Fantasies’, *Journal of Abnormal Psychology*, 88: 584–91.
- Bowers, K. S. (1992) ‘Imagination and Dissociation in Hypnotic Responding’, *The International Journal of Clinical and Experimental Hypnosis*, 40: 253–75.
- Child, W. (1994) *Causality, Interpretation, and the Mind*, Oxford: Clarendon Press.
- Cleghorn, J. M., Franco, S., Szechtman, H., Brown, G. M., Nahmias, C., and Gannett, E. S. (1992) ‘Toward a Brain Map of Auditory Hallucinations’, *American Journal of Psychiatry*, 149: 1062–9.
- Green, C. E. (1968). *Lucid Dreams*, Oxford: Institute of Psychophysical Research.
- Green, C. and McCreery, C. (1994) *Lucid Dreaming*, London: Routledge.
- Hearne, K. (1982) ‘Settings and Causes of Lucidity’, *Lucidity Letter*, 1: 14.
- Hilgard, E. R. and Nowlis, D. P. (1972) ‘The Contents of Hypnotic Dreams and Night Dreams: An Exercise in Methodology’, in E. Fromm and R. E. Shor (eds.), *Hypnosis: Research Developments and Perspectives*, Chicago: Aldine Atherton.
- Hinton, J. M. (1973) *Experiences*, Oxford: Clarendon Press.
- Kosslyn, S. M., Thompson, W. L., Costantini-Ferrando, M. F., Alpert, N. M., and Spiegel, D. (2000) ‘Hypnotic Visual Illusion Alters Color Processing in the Brain’, *American Journal of Psychiatry*, 157: 1279–84.

- Lewis, D. (1997) 'Finkish Dispositions', *The Philosophical Quarterly*, 47: 143–55.
- McDowell, J. (1982) 'Criteria, Defeasibility, and Knowledge', *Proceedings of the British Academy*, 68: 455–79.
- (1998) 'Having the World in View: Sellars, Kant and Intentionality', *Journal of Philosophy*, 95: 431–91.
- Martin, C. B. (1994) 'Dispositions and Conditionals', *The Philosophical Quarterly*, 44: 1–8.
- Martin, M. G. F. (2004) 'The Limits of Self-Awareness', *Philosophical Studies*, 120: 37–89.
- (2006) 'On Being Alienated', in T. S. Gendler and J. Hawthorne (eds.), *Perceptual Experience*, Oxford: Clarendon Press.
- Orne, M. T. (1959) 'The Nature of Hypnosis: Artifact and Essence', *Journal of Abnormal and Social Psychology*, 58: 277–99.
- Rechtschaffen, A. (1978) 'The Single-Mindedness and Isolation of Dreams', *Sleep*, 1: 97–109.
- Schatzman, M., Worsley, A., and Fenwick, P. (1988) 'Correspondence during Lucid Dreams Between Dreamed and Actual Events', in J. L. Gackenback and S. LaBerge (eds.), *Conscious Mind, Sleeping Brain: Perspectives on Lucid Dreaming*, New York: Plenum.
- Segal, S. J. (1972) 'Assimilation of a Stimulus in the Construction of an Image: The Perky Effect Revisited', in P. W. Sheehan (ed.), *The Function and Nature of Imagery*, New York: Academic Press.
- Shor, R. E. (1959) 'Hypnosis and the Concept of the Generalized Reality-Orientation' *American Journal of Psychotherapy*, 13: 582–602.
- Smith, A. D. (2002) *The Problem of Perception*, Cambridge, MA: Harvard University Press.
- (2003) *Husserl and the Cartesian Meditations*, London: Routledge.
- Szechtman, H., Woody, E., Bowers, K. S., and Nahmias, C. (1998) 'Where the Imaginal Appears Real: A Positron Emission Tomography Study of Auditory Hallucinations', *Proceedings of the National Academy of Sciences of the United States of America*, 95: 1956–60.
- Zamansky, H. S. (1977) 'Suggestion and Countersuggestion in Hypnotic Behavior', *Journal of Abnormal Psychology*, 86: 346–51.

8

The Epistemic Conception of Hallucination

Susanna Siegel

Disjunctivism about perception is a view about the relation between veridical experiences and hallucinations. Suppose that I see a green cube, and my experience is veridical—no illusion or hallucination is involved.¹ The veridical experience, according to disjunctivists, includes as constituents the bit of the world that is perceived and the perceptual relation that the subject bears to it, and these constituents characterize what it is like to see the green cube. The phenomenal character of my veridical experience, they hold, is thus constituted by facts that could not obtain if I were hallucinating. Disjunctivists typically grant that there is a corresponding hallucination in which I could not know just by introspection that I was not veridically seeing a green cube. But they hold that the most basic characterization of what it is like to have such a hallucination—its most basic phenomenal character—must differ from that of the veridical green-cube experience. The main opponents of disjunctivism are *common-kind theorists*, who hold that the most basic phenomenal character in both experiences is the same.

Although the primary commitment of disjunctivism concerns veridical perception, rather than hallucination, disjunctivists aim to respect two central assumptions about hallucination: that they can ‘match’ veridical perceptions, and that they can differ phenomenally from one another. Both assumptions are commonplaces in their own right, and the first is presupposed by the debate between disjunctivism and the common-kind theory. Common-kind theories accommodate both assumptions straightforwardly. ‘Matching’ hallucinations and non-hallucinations share the same most basic sort of phenomenal character, and variation among hallucinations is variation in that same sort of phenomenal character. Disjunctivists, in contrast, cannot

For helpful discussion and criticism, many thanks to Matthew Boyle, Alex Byrne, Bill Fish, Adrian Haddock, Benj Hellie, Heather Logue, Fiona Macpherson, M. G. F. Martin, Ram Neta, Bernhard Nickel, and Stephen White. Special thanks to David Chalmers, Matt Soteriou, and Scott Sturgeon for so many helpful conversations and for comments on multiple drafts.

¹ Throughout this essay, I’ll use ‘veridical experience’ to pick out instances of perception that are not hallucinations or illusions. So there will be no such thing as a veridical hallucination in this sense of ‘veridical’.

accommodate the assumptions in this way, since they deny that matching experiences share the same basic phenomenal character. To respect these assumptions consistently with disjunctivism, some other conception of hallucination is needed.

Recently, some disjunctivists, including M. G. F. Martin and William Fish, have developed an epistemic conception of hallucination. According to this conception, there is nothing more to the phenomenal character of hallucination of a green cube, besides the fact that the hallucination is indiscriminable from a veridical perception of a green cube as such. No further mental properties underlie its being so indiscriminable. Fish (this volume) describes what he takes to be the essence of hallucination, quoting Martin (2004: 72) at the end:

So the essence of hallucination—what distinguishes hallucinations as a class from other mental states—lies in their being indistinguishable from veridical perceptions, not in some antecedently identifiable feature of the state. This is why, “when it comes to a mental characterization of the hallucinatory experience, nothing more can be said than the relational and epistemological claim that it is indiscriminable from the perception”.²

Martin and Fish offer more detailed treatment of hallucination than earlier defences of disjunctivism offered. The main idea, however, that ‘matching’ should be given a purely epistemic construal is central to disjunctivism itself. Disjunctivism takes facts about matching experiences to be facts about how those experiences seem to the subject, rather than taking those facts to reflect any deeper similarity between the matching experiences themselves. We can thus assess the plausibility of disjunctivism by assessing the plausibility of the conception of hallucination on which it depends.

In this essay, I argue against the epistemic conception of hallucination. In section 1, I state some central theses about hallucination put forward by Martin, who has done the most to develop the epistemic conception. In section 2 I introduce a notion of indiscriminability that figures in the epistemic conception. In section 3, I argue that the epistemic conception falters with its treatment of cognitively unsophisticated hallucinators. In section 4, I introduce Fish’s version of the epistemic conception and raise two objections to it. In section 5, I argue that neither version of the epistemic conception has a promising account of what hallucinators can know on the basis of introspection. I conclude that the prospects for a plausible disjunctivist theory of hallucination are not promising. Although much of the discussion focuses on commitments made explicitly by Martin and Fish, my main aim throughout is to explore the space of possible disjunctivist proposals about hallucination and their pitfalls, rather than to locate Martin’s or Fish’s exact positions in it.

1 DISJUNCTIVISM

Return to the case in which I veridically see the green cube. My visual experience in this case is indiscriminable from a veridical perception of a green cube. This follows

² See also Martin (2006).

trivially, because the experience is a veridical perception of the green cube. On this much disjunctivists of Martin's and Fish's stripe (from now on I'll simply call them 'disjunctivists')³ and their opponents can agree. They can also agree that a hallucinatory experience can be indiscriminable from a veridical perception of a green cube.

The central theses of disjunctivism concern the status of specific indiscriminability properties of experiences, such as the property of being indiscriminable from a veridical perception of a green cube. For our purposes, the most important theses concern hallucinations. According to disjunctivism, my hallucination of the green cube is indiscriminable from a veridical perception of a green cube, even though there is no robust property the hallucination has in virtue of which it is so indiscriminable. This is Martin's line on *causally matching hallucinations* generally: hallucinations with the same proximate causal antecedents as veridical perceptions. (So let us retroactively suppose that my hallucination of a green cube causally matched a normal veridical perception of a green cube.) The most basic mental kind that hallucinations belong to, disjunctivists hold, is: being indiscriminable from a veridical perception. Veridical perceptions also belong to this kind, they think, but it is not their most basic mental kind. That kind is: being veridical perceptions; whereas the most basic mental kind of causally matching hallucinations is: being indiscriminable from a veridical perception. This is the most basic mental kind (or in Martin's terminology, the fundamental kind) of hallucinations in the sense that there is no further mental property hallucinations have, in virtue of which they are so indiscriminable.

Veridical perceptions are, trivially, indiscriminable from veridical perceptions; but according to disjunctivists further properties underlie their indiscriminability. For instance, my veridical perception of the green cube, they think, is indiscriminable from a veridical perception of the green cube in virtue of the perceptual relation that holds between the perceiver (me), on the one hand, and the cube and the properties of it that appear to me, on the other. These further properties, on their view, constitute the most basic phenomenal character of the veridical experience. This phenomenal character is thus robust in the sense that it underlies the indiscriminability.

So disjunctivists agree with the common-kind theorist that there is a common element between causally matching hallucinations and the veridical perceptions that they causally match;⁴ but they disagree about the depth and significance of the

³ This use of 'disjunctivism' may be sociologically misleading since many theorists doing business under that label say very little about the nature of hallucination, and do not even invoke the notion of indiscriminability. (Cf. McDowell (1982) and Putnam (1999). Campbell (2002) repudiates the label but holds that the natures of non-hallucinations and hallucinations fundamentally differ.) To uphold their views, however, these theorists need some conception of hallucination that is consistent with disjunctivism, and they could sign on to the epistemic conception if they wanted to. So the discussion here bears on possible developments of their views.

⁴ Martin grants the conclusion of the Causal Argument from Hallucination (see section 4), which he formulates as follows: "whatever kind of experience does occur in situations like *h* [hallucination], it is possible that such a kind of experience occurs when one is veridically perceiving" (Martin 2004: 38). Of course, "whatever kind..." is restricted to exclude the kind: hallucination.

common element. For disjunctivists, it goes no deeper than the indiscriminability property, and it does not constitute the fundamental kind to which both experiences belong.⁵

We can now draw out the central claim that disjunctivists defend about hallucinations. Hallucinations, according to them, are events or states⁶ whose phenomenal character consists exclusively in the having of indiscriminability properties. Let us formulate the thesis for a specific case and call it (H):

(H) If S hallucinates a sausage as such, then the hallucination is indiscriminable from a veridical perception of a sausage as such, and there is no further mental property that constitutes the phenomenal character of the hallucination.

I am going to let 'S hallucinates (/veridically perceives) a sausage' abbreviate 'S hallucinates (/veridically perceives) a sausage as a sausage' (and similarly for things other than sausages). The additional clause ('as a sausage') makes a difference because in principle one could veridically perceive a sausage that looked nothing like a sausage (for example, if it was disguised as a carrot), yet the experience could still be veridical: the experience might take a stand on only the colour and shape of the disguise (if these were the only properties 'manifested' to the subject, and these were properties the disguise actually had). This would be a case in which a sausage was veridically perceived, but wasn't perceived to be a sausage.⁷ For the sake of brevity I'll sometimes leave out the extra clause, but the result should always be taken as an abbreviation.

The indiscriminability property that figures in (H) is both epistemic and purely negative. If your experience has it, then a certain fact about your experience is not introspectively knowable by you—namely, the fact that it is not a veridical perception of a sausage. Since a hallucination is a kind of visual experience, and since nothing is a visual experience if it lacks phenomenal character, disjunctivists who accept the generalization of (H) are also committed to the view that an event or state's having the relevant sort of indiscriminability property suffices for it to be a visual experience. Let us call this claim $I \rightarrow E$:

$I \rightarrow E$: If a mental state or event has the property of being indiscriminable from a veridical perception of an F as an F, then it is an experience.

The conditional $I \rightarrow E$ is a schema. Once a more specific indiscriminability property is filled in, such as the property of being indiscriminable from a veridical perception of a sausage as such, the result is a claim that having that property suffices for having

⁵ "...while the perceptual event is of a fundamental kind which could not occur when hallucinating, nonetheless this very same event is also of some other psychological kind or kinds which a causally matching hallucinatory event . . . belongs to" (Martin 2004: 60).

⁶ Martin speaks of events rather than states. My discussion is indifferent to whether disjunctivism is formulated in terms of events or states.

⁷ Some might deny that anything can be visually perceived as a sausage, on the grounds that only colours, shapes, and other 'low-level' properties can be visually perceived. Nothing turns on whether this denial is correct. The sausage hallucination example comes from Martin (2006), but the green cube hallucination would work just as well throughout.

a correspondingly specific experience, such as an experience as of a sausage. If $I \rightarrow E$ is true, then there are also strong links between the notions of experience, on the one hand, and veridical perception, on the other. In Siegel (2004) I criticize the claim that the notion of experience is as closely tied to the notion of veridicality as $I \rightarrow E$ suggests, but I will set that issue aside here.

Let us now consider more closely what the notion of indiscriminability is that figures in H and $I \rightarrow E$.

2 INDISCRIMINABILITY

So far our discussion has allowed that a particular, unrepeatable experience could be indiscriminable from another (perhaps repeatable) event. It's not entirely clear what this means. The relevant notion of indiscriminability can't be a statistical notion, since the event said to be indiscriminable is unrepeatable. Nor can the notion be reasonably understood by considering what would happen if the subject had two simultaneous experiences, compared them, and found that they were the same in the relevant respect (as one might be able to do with two physical objects). Perhaps one could think of how the subject would regard the pair of experiences, if she had them successively; but that introduces complications about memory that seem extraneous. Nevertheless, there is some intuitive sense in which certain pairs of experiences seem the same to the subject. The disjunctivist (and anyone else applying the notion of indiscriminability to experiences) thus faces the task of specifying what it is for one experience to be indiscriminable from another. And the disjunctivist alone faces the task of specifying what indiscriminability is, without bringing in robust phenomenal character.

As Martin thinks of it, indiscriminability is a notion defined in terms of judgement. "To discriminate two things" Martin writes, "is to judge them non-identical" (2004: 62). This suggests that when A and B are *indiscriminable* for a subject, the subject cannot tell them apart *in judgement*. This suggests the following notion of indiscriminability:

X is indiscriminable from Y by a subject S at time t iff S cannot know at time t by introspection alone that X is not Y.⁸

Here indiscriminability is a relation that holds between tokens. According to a closely related notion, it is a relation that holds between a token and a type:

X is indiscriminable from Fs by a subject S at time t iff S cannot know at time t by introspection alone that X is not an F.

Of the two relations, the latter is closer to the one endorsed by Martin. According to him, if a subject S hallucinates a green cube, and if her hallucination—call it h—is indiscriminable from a veridical perception of a green cube, then applying this notion

⁸ This formulation of indiscriminability comes from Williamson (1990).

of indiscriminability yields the result that *S* cannot know by introspection alone that *h* is not a veridical perception of a green cube.⁹

We are now in a position to see what role in (H) is played by the notion of being indiscriminable from a veridical perception of a sausage, when it is cashed out in terms of such knowability. If (H) is true, then in the case of hallucination, a fact about what can't be known introspectively about the hallucination is supposed to characterize how things are visually from the subject's perspective. For instance, if you are hallucinating a sausage, then by Martin's lights, the impossibility of knowing introspectively that you are not veridically perceiving a sausage is all there is to the phenomenal character of your experience.

3 COGNITIVELY UNSOPHISTICATED HALLUCINATORS

When (H) and its generalization are combined with the interpretation of indiscriminability discussed above, they seem to have no application to cognitively unsophisticated creatures such as dogs, even though it seems plain that dogs can hallucinate.¹⁰ Consider a dog that lacks the cognitive equipment needed to form judgements, including ones that count as knowledge. For such a dog, there will be no perceptual experience such that he can know that it is distinct from a veridical perception. So all of the dog's perceptual experiences will count as indiscriminable from every kind of veridical perception. If the fact that two experiences share a specific indiscriminability property suffices for them to seem the same to the subject, then by accepting this interpretation of $I \rightarrow E$, disjunctivists predict that all the dog's experiences will seem the same to the dog. That seems plainly wrong.

This consequence hits hardest in the case of hallucinations. In the case of dog perceptions, the disjunctivist thinks there are robust world-involving properties that constitute the experiences, and that the dog's perception of a sausage will differ phenomenally from his perception of a carrot because of the difference between the corresponding robust properties. Although these experiences also differ in their indiscriminability properties, the disjunctivist need not rely upon these properties to make it the case that the experiences differ from the dog's point of view. In contrast, in the case of hallucinations, there is supposed to be nothing to the experiences (by disjunctivists' lights) besides the having of indiscriminability properties.

Martin considers and responds to this worry. His response introduces the idea that the relevant notion of discriminability should be impersonal:

when we turn to the experiences of sentient but unselfconscious creatures, to the extent that we do have a positive grip on the kinds of experience that they can have, and which can differ

⁹ It is possible to define other notions of indiscriminability. Some of these are discussed and criticized in Siegel (2004) and Farkas (2006). The most extended discussion of the notion is in Williamson (1990).

¹⁰ For brief discussion of empirical evidence of hallucination on the part of non-human animals, see Fish, this volume.

one from another, we also have a grip on how such experience would be discriminable through reflection or not . . . a dog might fail to discriminate one experience from another, making no judgment about them as identical or distinct at all, [but] that is not to say that we cannot judge, in ascribing them such experience, that there is an event which would or would not be judgably different from another experience. (2004: 54)

[W]e are interested in the impersonal notion of inability or incapability here. That is we are interested in the claim that John [or the dog] is in a situation for which it is impossible *simpliciter* and not just impossible for John [or the dog] to tell apart through introspective reflection from a veridical perception [of a sausage]. (2006: 381)

In considering this worry, Martin here is discussing John, an inattentive, hasty subject who treats samples of scarlet and vermillion indifferently. Intuitively, his experiences of each should count as distinct; yet, as Martin points out, it seems appropriate to say that inattentive John can't discriminate scarlet from vermillion. Martin takes the moral to be that the disjunctivist should adopt an 'impersonal' notion of indiscriminability, where this impersonal notion applies equally to the dog, to inattentive John, and to perfectly attentive cognitively sophisticated subjects.

Let's say that an indiscriminability property I of an experience is *impersonal* just in case there are no cognitive abilities that the subject must actually have, beyond any that may be needed to have conscious experiences, in order for her to have an experience with I. Martin proposes that we should understand these properties by comparison with the ordinary notion of invisibility. We can make sense of the idea that a bit of mending on a jacket is invisible, by considering whether someone ideally placed to see things—an ideal visual perceiver in ideal circumstances—could see the mending or not. If not, then the mending is invisible (Martin 2006: 381). By analogy, he suggests, we can make sense of the idea that the dog's experience is objectively indiscriminable from a veridical perception of a sausage while being objectively discriminable from a veridical perception of carrots, by considering whether by some means not employable by the dog himself, someone could know that the dog's experience is not a veridical perception of a sausage. If she could not know this, and if she could know that it's not a veridical perception of a carrot, then the experience is (non-trivially) indiscriminable from a veridical perception of a sausage.

The crucial question for the disjunctivist is what it could be for a dog's experience to have such an impersonal indiscriminability property. The disjunctivist must deny that an experience's having it involves its having underlying robust properties, in virtue of which the experience is indiscriminable from a sausage-perception. If the experience had such robust properties, that would violate the thesis that the hallucinator's experience consists merely in its having a negative epistemic property.

The fact that on pain of violating (H), impersonal indiscriminability properties cannot involve underlying robust properties suggests that they should be understood as some sort of counterfactual about the kind of knowledge that figures in the epistemic notion of indiscriminability endorsed by Martin and Fish. Here are three

counterfactuals that might be proposed as making explicit what it consists in, for the case of the dog:

- (1) If the dog could ideally reflect on his situation, he would not be able to know by reflection that he was not veridically perceiving a sausage.
- (2) If I could ideally reflect on the dog's situation, I would not be able to know by reflection that he was not veridically perceiving a sausage.
- (3) If an ideal introspector were in the dog's situation, such an introspector would not be able to know by reflection that she was not veridically perceiving a sausage.

If any of these explications is a way to understand what impersonal indiscriminability properties are, then given the claim that $I \rightarrow E$, it must support this conclusion:

C. The dog has an experience as of a sausage.

These counterfactuals all mention the dog's 'situation'. What is it to be in the [hallucinating] dog's situation? To be in this situation can't be to have a robust phenomenal character, by disjunctivists' lights. It cannot be a subpersonal physical or functional state that was identified with the experience, since that too would introduce an underlying common phenomenal character. It cannot be a subpersonal physical or functional state that normally is causally antecedent to the experience, because then the account of what the hallucinating dog's experience consists in will undergeneralize if there are other causal routes to having the experience, and overgeneralize in cases where the subpersonal state does not cause any experience.¹¹ And it can't be simply to have the property of being impersonally indiscriminable from a veridical perception of a sausage—or else 3, for example, would mean the same as the trivially true 3*:

- (3*) If an ideal introspector couldn't know by reflection that she wasn't veridically perceiving a sausage, then she couldn't know by reflection that she wasn't veridically perceiving a sausage.

A similar objection applies to counterfactuals 1 and 2. 'The dog's situation' in each case cannot pick out a robust phenomenal character or a subpersonal state that is or is causally antecedent to the experience. An appeal to an indiscriminability property moves us in a circle, since that is the property that the counterfactuals are attempts to explain. Finally, it seems ill-advised for the disjunctivist to get rid of mention of a situation altogether in 1, so that the antecedent read 'If the dog could ideally

¹¹ To elaborate the counterfactual is supposed to provide an account of what it is for the hallucinating dog to have an experience as of a sausage. It accounts for this only if the cases where the counterfactual holds are all and only cases where the dog is in the situation mentioned in the antecedent. If that situation is characterized in terms of the normal causal antecedents of the experience, then the account will be adequate only if those causal antecedents occur when and only when the dog has the relevant kind of experience. There will thus be a problem of undergeneralization if there are other causal routes to the same effect (namely, the experience), and a problem of overgeneralization if there are cases where the normal causal antecedents don't produce the effect (say because something intervenes, or because of a chance occurrence). Since effects typically can have more than one cause, and since causal antecedents are not guaranteed to have their normal effects, both problems arise.

reflect . . .'.¹² Substituting this for I in $I \rightarrow E$ would have the consequence that the dog was always having an experience, assuming that the dog could ideally reflect.

Disjunctivists must find a way to talk about the dog's situation—that is, his being in a state or event that $I \rightarrow E$ says shall count as an experience. Here all three counterfactual proposals face the same problem. If this situation is not a specifically experiential situation, then there is no reason to think that the counterfactual in which it appears explicates what it is to have an experience. Once the situation in the antecedent is an experiential situation, however, the account becomes trivial, since the disjunctivist's view is precisely that there is nothing else to the experience besides the having of the indiscriminability property itself. These doubts are general, and seem to apply to any attempt to explicate what impersonal indiscriminability properties could be. Given the assumption that some creatures who can have experiences are nonetheless not capable of having the kind of knowledge invoked by the impersonal indiscriminability properties, the prospects for developing the modalized epistemic notion of impersonal indiscriminability consistently with disjunctivism thus seem dim.¹³

Disjunctivists might try responding that dogs do not after all lack the kind of knowledge involved in the indiscriminability properties. To follow this strategy, the disjunctivist would first have to identify a general cognitive ability that dogs as well as cognitive sophisticates actually have, so that the notion of indiscriminability could be formulated using that notion. A starting point might be Sosa's notion of animal knowledge in his (1991). So long as there could be a perceiver who lacked the ability to have the suitably undemanding kind of knowledge, however, the dog problem will return. To follow this strategy, disjunctivists would thus have to argue that *any* creature capable of having hallucinatory experiences is also capable of having this kind of knowledge.

This last key claim seems difficult to establish. Suppose that the animal knowledge in question is a reliable connection of some sort between putative perceptual states and the environment (as suggested by Sosa's notion of animal knowledge). Then the key claim says that any creature capable of hallucinating is also capable of being in states that are reliably connected to the environment. It seems easy to imagine a creature with separate bits of mental machinery dedicated to producing conscious experience on the one hand and to receiving inputs from the environment on the other, where the two bits could operate independently. Suppose that in such a creature, the machinery dedicated to receiving inputs from the environment was so defective that no correlation ever existed between states of the environment and states of the machinery. It seems coherent to suppose that the other bit of machinery might produce conscious experiences perfectly well, including those phenomenally like the ones had by normal animals when they perceive. The result would be a creature capable of conscious hallucinatory experiences, yet incapable of animal knowledge, on

¹² This was suggested by a referee.

¹³ Some of these difficulties are acknowledged by Martin (2006: 383 and fn. 44) where he disowns some of these ways of developing the notion of impersonal indiscriminability. Martin retains a notion of impersonal indiscriminability. His preferred version of impersonal indiscriminability seems to be implicit in his view of introspection, which is discussed further in section 5.

this construal of it. More generally, many ways of developing the notion of animal knowledge seem to leave it an open possibility that the features enabling animals to have that kind of knowledge need not be the very features that enable them to hallucinate. The disjunctivist would have to show that despite appearances, these two capacities (for animal knowledge on the one hand, and for hallucination on the other) could not each operate separately.

There is a notion of animal knowledge, however, on which these two capacities could not operate separately. Suppose that the kind of knowledge that animals could have is a sort that one is automatically capable of having as a result of having conscious experiences of the right sort, where the right sort includes hallucinations. Then it would not be possible for a creature to hallucinate without the requisite kind of knowability facts obtaining. In section 5 (under the heading 'brute fact proposal') I criticize a general version of this proposal, one that is indifferent to whether the kind of knowledge in question is available to cognitively unsophisticated creatures or not.

Sympathizers with disjunctivism might try to dismiss the entire worry about cognitively unsophisticated creatures, by responding that the view as a whole should not stand or fall with its treatment of hallucinating animals. This dismissal is inconsistent with a commitment of Martin's: that being indiscriminable from veridical perception "is the most inclusive conception we have of what sensory experience is" (Martin 2004: 56). If disjunctivists propounding (H) take on this commitment, then they are stymied by the dog objection. If they deny this commitment, opting out of providing a unified account of perceptual experience as such, then the dog objection will not arise, but other objections will still arise. Some of these objections target Fish's version of epistemic conception, which will be discussed next. Others target both Fish and Martin's versions of the epistemic conception, and they are presented in section 5.

4 FISH'S EFFECTS-BASED VERSION OF THE EPISTEMIC CONCEPTION

So far we've been discussing Martin's knowability-based notion of indiscriminability. In contrast to Martin, Fish develops a notion of indiscriminability of hallucinations from veridical perception that fixes on their effects. In this section I argue against Fish's version of the epistemic conception.

Fish takes it as analytic that hallucinations have certain effects:

[W]e can say that the indistinguishability of the two states [hallucination and veridical perception] is manifested by the hallucination having sufficiently similar effects to those that the veridical perception would have had. But the demand for a more substantial intrinsic characterization of the hallucinatory mental state is misguided. For a similar reason, it would also be misguided to demand any further explanation of why hallucinations, considered as a mental kind, have similar effects to veridical perceptions. As a mental state only qualifies as a hallucination inasmuch as it has the same kinds of effects as a certain kind of veridical perception, there is therefore something wrong with asking why *hallucinations* have these kinds

of effects—it would be akin to asking what bachelors have in common in virtue of which none of them are married. (Fish, this volume: section 7)

Which effects are sufficiently similar? According to Fish, the relevant effects in creatures like us are beliefs, including higher-order beliefs, such as might be expressed by saying ‘I am veridically perceiving a sausage’:

Say I veridically perceive a pink ice cube: what kind of effects would such a mental state have? . . . [T]here would be a vast array of effects such a state might have—it might make me want a cold drink, or remember the last occasion on which I saw such an ice cube, or think of Wilfrid Sellars, and so on. But there would also be more standard effects than these. In particular, where subjects such as ourselves are concerned, someone who veridically perceives a pink ice cube would, unless something very curious was going on, acquire a range of fairly predictable beliefs. We would expect these beliefs to include, for example, the belief that a pink ice cube exists and the belief that she sees the pink ice cube. (Fish, this volume: section 5)

So according to Fish, the property of being indiscriminable from a veridical perception of a sausage is constituted by the property of having sufficiently similar effects as a veridical perception of a sausage.

On the face of it, Fish’s effects-based notion of indiscriminability makes it easier for a proponent of (H) to solve the dog problem. According to Fish, in the case of creatures who are not sophisticated enough to have higher-order beliefs, hallucinations will not include among their effects the belief on the part of the subject that she is veridically perceiving. Instead, the effects that constitute their hallucinations will be limited to behavioural effects:

[W]hen it comes to a creature—a cat, for example—which lacks the conceptual sophistication required to entertain higher-order beliefs, a veridical perception of, say, a butterfly will still have certain kinds of effects. Of course, given that we are assuming that the cat is lacking in higher conceptual capacities, these effects will not include cognitive effects such as higher-order beliefs. Instead, they will be primarily behavioural, possibly including trying to strike the butterfly with its paw, and so on. But then our evidence for animal hallucinations is essentially a set of behavioural observations too. If we were to find a cat “striking at imaginary objects in the air” (Slade and Bentall 1990: 12)—in other words, behaving as though it perceived a butterfly when there was not one there—then we might have good reason to say that, in such a case, the animal is not perceiving a butterfly but is nevertheless in a mental state which has the same effects as such a perception. As the two states have sufficiently similar effects . . . we can therefore allow that the two states are indistinguishable for the cat, and that the animal is having a hallucination of a butterfly . . . (Fish, this volume: section 6)

With the effects-based notion of indiscriminability in place, worries about how the cat could non-trivially fail to know that she is not veridically perceiving a butterfly disappear. If effects-based notion is correct, then it avoids the dog problem.

There are problems facing this notion, however. The theory does not ensure that hallucinations have any felt reality from the point of view of the hallucinator—not even a deflationary sort of felt reality that Fish himself invokes. Fish considers his theory to be an irrealist account of hallucination that “denies that hallucinations have phenomenal character whilst agreeing wholeheartedly that they certainly seem to”.

Their seeming to have phenomenal character, he says, consists in part in the subject's believing that they are veridically perceiving.¹⁴

[I]f a mental state which was *not* a veridical perception of a pink ice cube somehow led the subject to form beliefs of this kind, then these beliefs could be appealed to in explaining why the subject takes herself to be having a veridical perception of a pink ice cube even though she is not. So long as I believe that I am seeing something, I will take myself to be enjoying a certain kind of veridical perception If a mental state which was not a veridical perception were somehow to come to have effects which included such higher-order beliefs, then this would look to be sufficient to explain why that mental state was mistaken for a veridical perception of that kind. (Fish, this volume: section 5)

The appeal to actual higher-order beliefs cannot account for what makes all hallucinators take themselves to be enjoying a veridical perception when they are hallucinating, because some of these hallucinators will not actually form the higher-order beliefs. Cognitively unsophisticated subjects may be one example. We can find other examples if we unpack the effects-based theory a bit more.

Fish's position seems best construed as a disjunctivist version of analytical functionalism. On this view, hallucinations are defined using an *a priori* theory such as folk psychology (or some analogue for animals). Hallucinatory experiences are whatever states have the same effects as veridical perception, with the stipulation that they are not themselves veridical perceptions.¹⁵ Veridical perceptual experiences, in contrast, would have a functional role that included these effects but also included worldly causes as well.

Suppose, with Fish, that there is a set of effects that veridical perceptions of a butterfly standardly have—call this set *E*. These effects are not effects of the veridical perception considered in isolation, but rather they are effects of the veridical perception together with other mental states. As Geach (1957: 8) pointed out in criticizing logical behaviourism, mental states do not each have their own individual fund of behavioural effects. A cat veridically perceiving a butterfly standardly paws at it partly because she wants to make contact with it. If the cat is lethargic from illness and so lacks that desire, even if she veridically perceives the butterfly, she won't paw at it or even be disposed to paw at it. In identifying a set *E* of standard effects of veridically perceiving a butterfly, we are identifying effects of the veridical perception together with other mental states. Fish does not present his view this way, but Geach's point seems correct, so let us proceed by considering a Fish-inspired theory of hallucinations according to which they are states that have the certain effects, given the other mental states of the creature. What necessary conditions does (H) impose on an event or state's being a hallucination, given an effects-based notion of indiscriminability?

¹⁴ Fish offers separate treatment of 'resisted hallucinations', which are hallucinations in which the subject is 'aware of the non-veridicality' of the experience. His treatment of these hallucinations appeals to different higher-order beliefs. See Fish (this volume: section 10).

¹⁵ According to Fish's definition of hallucination, "a hallucination is a mental state which, whilst not being a veridical perception of a particular kind, has effects which are sufficiently similar to those a veridical perception of that kind would have had" (section 5). This definition includes illusion too—though to Fish this is not an unwelcome result (Fish MS).

Where E is still the set of effects standardly brought about by veridically perceiving a butterfly, the condition is this:

(*) An event or state e is a hallucination of a butterfly only if: if the subject has e and certain other mental states, then e will ground a disposition to produce E.¹⁶

By itself, the necessary condition on hallucinating butterflies laid down by (*) does not reflect any specifically disjunctivist treatment of hallucination. For all (*) says, e could ground a disposition to produce E in virtue of its phenomenal character, and that phenomenal character could be shared with a veridical perception of a butterfly. Condition (*), however, is compatible with the conditions that (a) butterfly-hallucinations lack robust phenomenal character altogether, and in particular (b) butterfly-hallucinations lack robust phenomenal character that veridical perceptions of butterflies have. In the conjunction of (*), (a) and (b), we can see the structural similarity between this treatment of hallucination and Martin's treatment of them. Both hold that an event or state is a hallucination only if it is indiscriminable from a veridical perception. And both hold that there is no common phenomenal character underlying the indiscriminability.

Let us now return to the objection that Fish's theory of hallucination does not ensure that they have any felt reality. According to Fish, having a higher-order belief that you are veridically perceiving a butterfly is supposed to account for why you take yourself to be having an experience with phenomenal character, even though (by his lights) you are not. But even some hallucinators will not form them—even if they are not 'resisting' their hallucinations (that is, even if they do not believe that they are hallucinating—Fish offers a separate treatment of 'resisted' hallucinations). We can leave aside the question whether failing to form these higher-order beliefs would be irrational—rationality is not a prerequisite of hallucinating. For subjects who fail to form these higher-order beliefs, the effects-based theory lacks a crucial resource for accounting for the felt reality of the hallucination. Neither the hallucination's actual effects nor any phenomenal basis of them is available to Fish to account for the felt reality of the hallucination. It is hard to see what else could play this role.

In response to this problem, the effects-based theory could be reformulated to make it analytic that hallucinations actually have the effects in set E, rather than making it analytic that they have those effects given appropriate background conditions. But this refinement would have implausible consequences. It should be possible to have a hallucination that is indiscriminable from a veridical perception of a butterfly, even if the standard effects E don't actually come about. And given a veridical perception of a butterfly that itself lacks the standard effects E, it should be possible for there to be a hallucination that is indiscriminable from it. A subject might suddenly expire just after veridically perceiving or hallucinating a butterfly, in which case the veridical perception or hallucination would not have E. So to ensure that such hallucinations are possible, the effects-based theory should not formulate the definition of hallucination in terms of actual effects.

¹⁶ The Fish-inspired theory would arguably increase in plausibility by adding "if the subject is rational . . ." to the necessary condition in (*). But adding this would not change the dialectic.

5 GROUNDING EPISTEMIC FACTS

My next objection to (H) focuses on what is introspectively knowable in cases of hallucination. First, I will argue that in addition to countenancing negative epistemic facts about the indiscriminability from other states, proponents of (H) must countenance certain positive epistemic facts about the discriminability of hallucinations from other states. Second, I will suggest, by considering and rejecting four proposals, that the disjunctivist cannot give a good account of the epistemic ground of these positive facts about discriminability.

Earlier, we mentioned two assumptions about hallucinations that the disjunctivist aims to respect. First, there can be ‘matching’ hallucinations; second, hallucinations can vary in their phenomenal character (the ‘difference datum’). According to the difference datum, the character of hallucination of a sausage differs from that of a hallucination of a pyramid or of a butterfly. So for any hallucination, there will be a set of possible experiences with contrasting characters. Let a butterfly-experience be an experience with the phenomenal character of veridically perceiving a butterfly as such, or with the phenomenal character of hallucinating a butterfly as such. (Disjunctivists will cash out the latter kind of phenomenal character in terms of indiscriminability, whereas common-kind theorists won’t see the need for the disjunction in the definition).

Given the contrast in character between a sausage-hallucination and the butterfly-experiences, it seems plausible to suppose that when you hallucinate a sausage, you can know introspectively that you’re not veridically perceiving a butterfly as such. Generalizing from this, it is equally plausible to suppose that if you have a butterfly-experience, you can know introspectively that you’re not veridically perceiving a sausage as such. For such facts to be knowable, it need not be the case that every detail of the character of the hallucination is introspectively accessible. Introspection just has to be able to be sensitive to the phenomenal contrast between sausage-hallucinations and q-experiences. We can abbreviate the last claim that I said was plausible to accept as claim (i) below, where $K \neg VP$ (sausage) is an abbreviation for the claim that it is introspectively knowable that you are not veridically perceiving a sausage as such. (So note that ‘K’ has a modal element built into it.)

(i) S has a butterfly-experience $\rightarrow K \neg VP$ (sausage).

The consequent of (i) is a fact (some may consider it a merely putative fact) about what is introspectively knowable. Let’s call this the Positive Epistemic Fact. (It is capitalized because it is specific to sausages.) My objection to (H) concerns what epistemic ground disjunctivists can offer for the Positive Epistemic Fact and others analogous to it.

Someone might challenge (i), on the grounds that it overestimates introspection’s actual epistemic capacities, inappropriately crediting it with being a source of knowledge about what one is not veridically perceiving. Could a disjunctivist avoid discussing the objection by denying (i) on these grounds?

The position on introspection is not open to the proponent of (H) who construes indiscriminability as unknowability. According to this construal of (H), when you hallucinate a sausage, you can't know by introspection that you're not veridically perceiving a sausage, and this introspective unknowability is all there is to the phenomenal character of your hallucination. The difference in character between your sausage-hallucination and, say, a hallucination of a butterfly is supposed to be reflected in the fact that each hallucination consists in a different fact about introspective unknowability. The fact that you can't know by introspection that you're not veridically perceiving a sausage is thus supposed to suffice to make it the case that you are not having a butterfly-experience: for example, you are not hallucinating a butterfly, and you are not having an experience with the character of veridically perceiving a butterfly. We can abbreviate this claim as claim (ii):

(ii) $\neg K \neg VP(\text{sausage}) \rightarrow S$ is not having a butterfly-experience.

The proponent of thesis (H) is thus committed to (ii), given that the unknowability fact in its antecedent is supposed to respect the datum that hallucinations can differ in their character from other hallucinations (as well as from other non-hallucinations). But since claim (ii) is the contrapositive of claim (i), the proponent of thesis (H) is committed to claim (i) as well. So there is no escaping a discussion of what grounds the Positive Epistemic Fact, if one accepts thesis (H) with the unknowability construal of indiscriminability.

What grounds the fact, illustrated by the consequent of (i), about what is introspectively knowable in the case of hallucination? It is of course not open to the proponent of (H) to appeal to the robust phenomenal character of the hallucination here. So let us consider the other options that are open to them and see whether they are plausible.

Is the epistemic fact grounded by the indiscriminability property? According to the first proposal, the knowability in the Positive Epistemic Fact can be derived from the unknowability that constitutes the sausage-hallucination—that is, the unknowability by introspection that one is not veridically perceiving a sausage. Call this unknowability the Negative Epistemic Fact. This proposal may appeal to a proponent of (H) because such unknowability is the most basic phenomenal character of hallucination. So the proposal preserves a sense in which the experience itself grounds the knowability.

However, it is not possible to derive the positive epistemic fact from the negative epistemic fact in this way. Suppose that I cannot know by introspection that I am not veridically perceiving a sausage (as such). Then it is an open possibility for me, for all introspection tells me, that I am veridically perceiving a sausage (as such). But this open epistemic possibility for me may be just one among many. For all introspection says, I may alternatively be veridically perceiving a pyramid as such, or a bridge as such, or carrots as such. All that follows from the Negative Epistemic Fact is that it is an open possibility that I am veridically perceiving a sausage as such, but it is compatible with the openness of that possibility that I am instead veridically perceiving (say) a pyramid as such.

In contrast, these other possibilities (that I am veridically perceiving a pyramid as such, and so on) are ruled out if the Positive Epistemic Fact obtains. According to the Positive Epistemic Fact, it is knowable to me introspectively that I am not veridically perceiving a pyramid. So it is not possible to derive the Positive Epistemic Fact from the Negative Epistemic Fact. To make that transition, something else is needed to rule out these other possibilities. The first proposal thus does not pin-point what makes the proposition that I am not veridically perceiving a non-sausage introspectively knowable, in the case where I am hallucinating a sausage (as such).

Is the epistemic fact grounded by introspection construed as a procedure? Suppose that there is some feature of introspection in virtue of which the Positive Epistemic Fact (and others like it) obtains. Versions of this proposal will vary depending on what the feature of introspection is taken to be. For example, according to a reliabilist version, the proposition that I'm not veridically perceiving a pyramid is introspectively knowable to me when I'm hallucinating a sausage, in virtue of the fact that I could employ a type of introspective process that reliably generates true beliefs that I'm not veridically perceiving a pyramid. Reliabilism is just one example. Let us focus more generally on the idea that there are dedicated mechanisms or processes of introspection whose employment is an optional addition to experience itself, as opposed to being constitutively linked to experience. Call this the procedural model of introspection.

In our discussion so far we have been focusing on cases in which the Negative and Positive facts obtain, and in which by hypothesis I am hallucinating a sausage. But the proponent of (H) holds that these facts suffice for me to have an experience. They accept the conditional: $I \rightarrow E$. Once we fix on the procedural model of introspection, it is coherent to suppose that a procedure or mechanism of introspection could operate in a way that resulted in the subject's knowing that she was not veridically perceiving a pyramid, even if she was not having any experience at all. For instance, suppose I have an introspective process or mechanism that reliably generates true beliefs to the effect that I am not veridically perceiving a pyramid. Suppose it operates and generates such a belief. By reliabilist lights, I will know introspectively that I am not veridically perceiving a pyramid. The Positive Epistemic Fact will thus obtain. Suppose further that the mechanism cannot in this situation generate the belief that I'm not veridically perceiving a sausage. Then I will not be able to know introspectively that I am not veridically perceiving a sausage. So the Negative Epistemic Fact will thus obtain. But it is compatible with the case as it has been described that I might lack experience altogether. If the introspective procedure is not constitutively linked to my having an experience, then it could operate in the absence of any experience. So the procedural model itself undermines $I \rightarrow E$, because of the lack of a constitutive connection between the operation of introspection and the experience itself. So if the proponent of (H) appealed to some feature of introspection, procedurally construed, to ground the Positive Epistemic Fact, she would run afoul of another disjunctivist commitment, namely $I \rightarrow E$.

Is the epistemic fact a brute fact? I'll call the third proposal for what grounds the Positive Epistemic Fact the brute-fact proposal, because it gives the Positive Epistemic Fact the status of brute fact. Giving the Positive Epistemic Fact this status may seem attractive for disjunctivists. Prima facie, the natural place to look for the aetiology of

such knowledge or for its ground would be the specific character of the experience. The epistemic conception of hallucination replaces talk of robust phenomenal character of hallucinations with facts like the Negative Epistemic Fact. The whole point of the epistemic approach to hallucination is that epistemic facts suffice to reflect that specific character. In effect, they are substitutes in the theory of hallucination for robust phenomenal character. But as we have seen, such facts cannot ground the Positive Epistemic Fact.

According to the brute-fact proposal, when I form the belief (while hallucinating a sausage) that I am not veridically perceiving a non-sausage, the status of my belief as knowledge, and the introspective knowability of what I believe, are both just brute epistemic facts, with no further explanation. In cases where I actually know introspectively that I am not veridically perceiving a non-sausage, the proposal says that there is no feature of introspection that makes my belief count as knowledge. And in cases where I have no non-introspective basis for this belief, the proposal says there is nothing at all that makes the belief count as knowledge. Its status as knowledge will be a brute fact.

The brute-fact proposal puts pressure on the idea that the beliefs in question can have the status of knowledge at all. If there is no feature of introspection that gives the belief its status as knowledge, then introspection contributes nothing to giving the belief its putative epistemic status as knowledge. In cases where there is no other basis for the belief, there will be nothing else that gives it this status either. So if the belief had the status of knowledge, it would be knowledge without any grounding or basis, evidential or otherwise. At that point is hard to see what makes the belief deserve its status as knowledge, as opposed to being merely true. Yet it seems plain that we can know introspectively, while hallucinating a sausage, that we are not veridically perceiving a pyramid—and moreover, proponents of (H) are committed to this. The brute-fact proposal is not clearly compatible with this.

Even if the brute-fact proposal could allow that the belief has the status of knowledge, it closes off the possibility of explaining a fact for which we would like an explanation anyway. Independently of assessing disjunctivism, we would like to know how it is possible for us to know what we do from a first-person perspective about our experiences. According to the brute-fact proposal, we will be barred from learning anything about this, because there is nothing to learn.

We can distinguish between an extreme and a less extreme version of the brute-fact proposal. The extreme version embodies a general view about introspection, holding that there is never any such thing as introspective evidence for knowledge claims, and no other feature of introspection that can make it the case that a belief formed on the basis of introspection counts as knowledge (when it does). A less extreme version holds this only for the relevant introspectively anchored knowledge claims in the case of hallucination (for example, the knowledge claim that can be derived from the Positive Epistemic Fact).

The extreme claim is pleasingly uniform, but not all disjunctivists can accept it. Suppose that introspecting a veridical experience of seeing a sausage as such gives us access to some aspect of its nature. If the extreme claim is true, such access will not have the status of knowledge. But some disjunctivists, including Martin, motivate

disjunctivism on the grounds that its treatment of veridical experience best accords with the way introspection presents those experiences as being.¹⁷ They thus seem to treat introspection as a source of knowledge, or at least justified belief, about the nature of experience. This claim seems weakened if it turns out that there are simply no answers to the question of what introspection is such that it provides such knowledge, or what in the nature of experience enables us know about it introspectively. So disjunctivists opting for the brute-fact proposal will have to adopt a version of it according to which introspection sometimes contributes to making a belief count as knowledge, and sometimes does not.

Are the epistemic facts grounded in the experiential perspective? According to Martin, the fact (as he sees it) that I cannot know by introspection that I'm not veridically perceiving a sausage constitutes the phenomenal character of my hallucination. If so, whatever introspection can tell me about my situation, it will not include that I am not veridically perceiving a sausage. The Negative Epistemic Fact that constitutes my sausage-hallucination thus gives me a 'perspective' on my situation. (Let us set aside the tricky question of what exactly 'my situation' is supposed to be—this point arose in our discussion of the dog objection.) Martin anticipates that someone might challenge (H) by asking: why think that any perspective on my situation delivered by introspection is accurate? Why couldn't the deliverances of introspection be inaccurate about the specific character of my experience—or even be wrong in supposing that I am having any experience at all? (Compare the putative case of zombies who falsely self-ascribe experiences.)

Martin responds to this challenge by holding that introspection's deliverances determine the existence and specific character of experiences, including hallucinations. More specifically, they constitute experiences—so 'deliverances' may be misleading, to the extent that the word suggests deliverances of verdicts about experiences. Martin's proposal links the perspective of introspection constitutively to experience itself:

[T]he perspective we have on our own phenomenal consciousness cannot be grounded in some specific mode or source of knowing about something independent of that perspective. If it is true of someone that it seems to them as if things seem a certain way, as if they are having a certain sense experience, then they are thereby having that experience. Our reflective standpoint on our own experience cannot stand outside of it. (Martin 2006: 397)

If the suggestion about the link between introspection and hallucination applies to the entirety of introspection's perspective, so that it is the entire perspective that constitutively determines the facts about one's experience, then there does not seem to be anything special about the Negative Epistemic Fact as opposed to the Positive Epistemic Fact. When the sausage-hallucinator can know by introspection that she is not veridically perceiving a non-sausage, it seems, this deliverance of introspection should also be constitutively linked to the hallucination. This in itself suggests a revision of

¹⁷ For example, Martin writes: "the disjunctivist advocates Naive Realism [the disjunctivist's view about the nature of veridical experience] because they think that this position best articulates how sensory experience seems to us to be just through reflection" (2006: 354).

thesis (H), according to which both Epistemic Facts constitute the sausage-hallucination, not just the Negative Epistemic Fact.

Suppose the disjunctivist accepts this revision of thesis (H). Can this position on the status of introspection help account for what grounds the Positive Epistemic Fact? It does not seem to help. We cannot look to the experience as the ground, since the experience is identified with introspective perspective. Alternatively, a proponent of identifying experience with the introspective perspective might try to dismiss the question as senseless, on the grounds that experience (unlike belief) has no ground, hence no question of the ground of introspective perspective can arise. But this response will not work either, since the proponent of (H) who relies on the knowledge-based notion of indiscriminability—and Martin is just such a proponent—is committed to cashing out the introspective perspective in terms of knowability claims. Such claims have grounds, whatever identity claims they may figure in. So the approach that identifies introspection's perspective with hallucination does not seem to do any better at accounting for the Positive Epistemic Fact.

Martin addresses something like the objection just made in a footnote, where he considers a comparison of his proposal about introspection with theories that posit constitutive relations between self-ascriptions of thoughts and the thoughts self-ascribed. He notes that some philosophers object to these theories that such self-ascriptions could have no rational ground. But he says no analogous objection applies to “the constitutive connection between the subject’s perspective on his or her own mind, how it seems to be, and how his or her mind then is” because the perspective “need not be identified with judgments he or she actually makes” (Martin 2006: n. 42). But even if the perspective is identified with judgements that the subject could make if she were rational, the question of what grounds them will still arise. And given the notion of indiscriminability Martin invokes, it is not in the cards to identify the perspective with something other than judgements. So the ground of the Positive Epistemic Fact still seems unaccounted for.

Given the difficulties with the four proposals for grounding the positive epistemic facts, I conclude that the chances of developing a plausible account of this grounding consistently with disjunctivism are low. When combined with the difficulties generated by the case of cognitively unsophisticated hallucinators, this result makes the prospects seem dim for the epistemic conception of hallucination.

REFERENCES

- Campbell, J. (2002) *Reference and Consciousness*, Oxford: Oxford University Press.
 Farkas, K. (2006) ‘Indiscriminability and Sameness of Appearance’, *Proceedings of the Aristotelian Society*, 106: 205–25.
 Fish, W. (this volume) ‘Disjunctivism, Indistinguishability and the Nature of Hallucination’.
 ——— (MS) *Perception, Hallucination and Illusion*.
 Geach, P. T. (1957) *Mental Acts*, London: Routledge and Kegan Paul.

- McDowell, J. (1982) 'Criteria, Defeasibility, and Knowledge', *Proceedings of the British Academy*, 68: 455–79.
- Martin, M. G. F. (2004) 'The Limits of Self-Awareness', *Philosophical Studies*, 120: 37–89.
- (2006) 'On Being Alienated', in T. S. Gendler and J. Hawthorne (eds.), *Perceptual Experience*, Oxford: Oxford University Press.
- Putnam, H. (1999) *The Threefold Cord*, New York: Columbia University Press.
- Siegel, S. (2004) 'Indiscriminability and the Phenomenal', *Philosophical Studies*, 120: 91–112.
- Sosa, E. (1991) 'Intellectual Virtue in Perspective', in his *Knowledge in Perspective: Selected Essays in Epistemology*, Cambridge: Cambridge University Press.
- Williamson, T. (1990) *Identity and Discrimination*, Oxford: Blackwell.

PART II
ACTION

9

Disjunctive Theories of Perception and Action

David-Hillel Ruben

There are many different ways in which to state something that might accurately be called a disjunctive theory of perception, but by the disjunctive theory of perception, I will mean this biconditional:¹

(1) It looks to S that p iff *either* (a) S sees or perceives that p *or* (b) it merely looks to S as if that p.²

I have expressed the theory in terms of fact-perception rather than object-perception, but nothing in what follows depends on the details of this particular statement of the theory. In truth, the theory might better be described as a disjunctive theory of how things look, or of experience: experience is either perceptual or hallucinatory (I disregard cases of illusion). But the theory is often described in the literature as a disjunctive theory of perception, although perception is merely one of the disjuncts. The theory addresses a traditional problem of perception; hence its more common name.

The idea of the theory is that its looking to a person that there is something that is so is made true either when he perceives that it is so or when it merely looks to him that it is so. The theory is not about conceptual priority, for 'the left-hand side' concept of how things look is conceptually prior to the 'right-hand side' concept of how they merely look. The point of the theory is to deny that 'left-hand side' statements about how things look to a person or about a person's experiences are to be understood as introducing reports of a single type of mental state common to both

I wish to thank the editors of this volume, but especially Adrian Haddock, who caused this chapter to be much better than it otherwise would have been. Haddock suggests another way in which I might express my view: the idea of an event embraces items of three different kinds: actions, mere events, and intrinsic events. We might say that, on my view, the kind EVENT is a determinable, and the kinds ACTION, MERE EVENT, and INTRINSIC EVENT are its determinations. In this terminology, event in the wide sense is the determinable, EVENT; event in the narrow sense includes two of its determinations.

¹ My statement of the theory is an amalgam of several such statements. See Child (1992: 145–9); Martin (2004: 37–89); Hinton (1973); Snowdon (1980–1: 175–92); and McDowell (1982: 455–79). The last two are reprinted in Dancy (1988).

² The left-hand side of the biconditional could also be expressed as: 'S experiences that p'.

perception and hallucination. If statements about how things look introduce anything about mental states at all, they introduce a disjunction of two types of mental states, each one introduced by one of the 'right-hand side' disjuncts.

It is, of course, too strong to say that these two disjunct states, perception on the one hand and hallucination on the other, share nothing in common. Since they are, *ex hypothesi*, perceptually indiscriminable, they must share at least these two features in common: the features of being perceptually indiscriminable from a veridical perception and of being perceptually indiscriminable from a hallucination (since each is also perceptually indiscriminable from itself). But what is meant to be true is that there is no significant or fundamental type (various statements of the theory express this somewhat differently) to which they both belong, and in virtue of which they are perceptually indiscriminable.³

Statements of the disjunctive theory of perception move rather too quickly between two ideas, both of which appeared above, which we might call the ideas of 'disjunctiveness' and 'lack-of-commonality'. Here are just two examples:

Disjunctive theorists . . . urge that there is nothing literally in common between . . . a case of veridical perception and . . . a case of hallucination . . . (Lowe 2000: 145)

The disjunctive theory of perception claims that we should understand statements of how things appear to a perceiver to be equivalent to statements of a disjunction . . . and that such statements are not to be viewed as introducing a report of a distinctive mental event or state common to these various disjoint statements. (Martin 2004: 37)

The ideas of disjunctiveness and lack-of-commonality are put forward as if they are almost equivalent ideas, or anyway as if each one logically follows from the other.

Is the disjunction that lends its name to the theory intended to be an inclusive or exclusive disjunction? It seems clear that the intention is for the disjunction to be exclusive: one is either having a veridical perception or merely a hallucination, but not both. Indeed, the meaning of each disjunct excludes the applicability of the other disjunct. To have a veridical perception is not to be having a hallucination and conversely. To that extent, the exclusivity of the disjunction is trivially true; it follows from the meaning of the words in the disjuncts.

On the other hand, Martin's minimalist characterization of the hallucination disjunct—that all that can be said about a hallucination is (a) that it is something that is perceptually indistinguishable from a veridical perception—would seem to make the disjunction inclusive, since a veridical perception is also perceptually indistinguishable from itself (see Martin 2004: 48). The minimalist account of the second disjunct

³ Although there is a lot of work to be done in clarifying this idea; see Martin (2004). The above formulation denies that there is some significant feature of both of the token experiences, the veridical perception and the hallucination, a significant feature that they both share. Here is another, similar, although not equivalent, way in which one might make the point, but which relies on the idea of numerical identity of a single token across possible situations: there is no veridical perception which could have been a hallucination, and there is no hallucination which could have been a veridical perception. But in discussion of the disjunctive theory of perception and that of action, I will use the first formulation, in terms of two tokens and the existence or not of a single significant mental type true of them both.

that Martin defends needs an additional idea in order to be adequate: (b) the idea of non-identity to a veridical perception as well. A hallucination is something that is perceptually indistinguishable from but not the same as a veridical perception. With (b) added, the disjunction will be exclusive. However, if Martin requires that the minimalist characterization of the second disjunct experience be one that is available to the subject's awareness, that adequate characterization is not available, since the very point of perceptual indiscriminability is that one may not be able to tell by introspection whether or not his experience is or is not identical to a veridical perception.

Just what is the relationship between these two ideas, that of disjunctiveness and lack-of-commonality? They certainly are not equivalent. It is obvious that there are exclusive (and a fortiori, inclusive) disjunctions which do have a common element. It is true that ' $p \vee \neg p$ ' is an exclusive disjunction which has no common element, but on the other hand ' $(p \& q) \vee (\neg p \& q)$ ' is an exclusive disjunction whose disjuncts do have a common element, viz. ' q '. So even if 'hallucination or veridical perception' is read as an exclusive disjunction, one cannot infer from that fact to the failure of a common element between them. Equally, and even more obviously, failure of a common element does not argue to exclusivity of disjunction: ' $p \vee q$ ' is a disjunction whose disjuncts have no common element, but it is not an exclusive disjunction—both disjuncts could be true.

When one reflects on the point of the disjunctive theory of perception, it seems that it is the idea of the lack-of-commonality and not the idea of exclusive disjunctiveness which bears the weight of the theory. The core thought of the disjunctive theory depends on the idea of no-feature-in-common-to-both. There is no single fundamental type of state, for example how things look to one, such that being in that state is or is part of the truth-maker for both veridically perceiving and hallucinating. (The biconditional should be read from right to left.) Disjunctiveness is, in any event, only a surface, syntactic feature of a sentence, as surely as we have known since Nelson Goodman's (2005) work on *grue*.

Sometimes a perceiver cannot tell whether he is the subject of a veridical perception or of a hallucination. That this is so is the basis for many formulations of scepticism. But even if a theory of perception cannot answer scepticism, it ought to be able to offer some account of how confusion, doubt, or error in perception can arise. Other theories of perception claim to be able to do this. If there were something common to veridical perception and hallucination, and if the only difference between them were the absence or presence of some yet further feature *F* (suppose *F* is not an observable property), then doubt, error, or confusion can be given this explanation: the subject is experiencing whatever is common to both veridical perception and hallucination, so that can't help him distinguish which sort of experience he is having, and he can't tell whether the experience he is having has *F* or not, since *F* is not an observable property. So the perceiver might be in doubt or confused about which kind of experience it is of which he is the subject.

But how can a theory that denies the commonality assumption explain perceptual error, doubt, or confusion? Martin says that the only mental characterization of a hallucinatory experience (beyond what might be said about experience in general) is that it is an experience that is perceptually indistinguishable from

a veridical perception. That is not to explain doubt or confusion, merely to redescribe it. There is, on this view, no explanation at all for perceptual doubt and error at the epistemological level, although surely there will be some straightforward physical account for them. In particular, this answer is unavailable: they are perceptually indistinguishable because they share the same phenomenal look. That answer is unavailable if it presupposes there is a common phenomenal character or feature that both kinds of experience share.

With this in mind, let's look at what we might call a disjunctive theory of action. What is the view that the disjunctive theory of action is opposing? That rejected (by me, anyway) theory, which for brevity I shall simply call the non-disjunctive theory of action, also assumes a type of commonality, but in this case a certain sort of commonality between actions, intrinsic events, and mere events (compare: veridical perceptions and mere hallucinations). First, we need two pieces of terminology: an intrinsic event and a mere event.

In what follows, I intend for what I say to be neutral between so-called austere and prolific theories of act identity. My own prolific terminology is based on the assumption that plurality is often in the actions themselves and not just in the descriptions, so that when I kill the Queen by shooting her, I count my killing of her and my shooting of her as two actions, not two descriptions of one action. But for those who see the plurality only in the description and not in the action, I believe that everything I have to say in what follows can be rewritten using the austere terminology.

1 INTRINSIC EVENTS AND MERE EVENTS

What is an intrinsic event? Suppose an agent *a*-ed, where his *a*-ing is a token action (of type *A*). On one view of the matter, but not one that I accept, at least not for all cases of action (and not for all senses of the word 'event'):

- (2) 'Someone *a*-ed' logically entails 'there was a token event e_a of type E_A '.

Jennifer Hornsby (1980), limiting her claim to actions described by transitive verbs, holds that there is an inference from ' $aV_t b$ ' ('*a* moved his finger') to ' bV_i ' ('his finger moved').⁴ David Hamlyn (1990: 130), evidently without limiting the claim to cases of action described by transitive verbs, says that "it is undeniable that, when we make a bodily movement [that is, when we act], a bodily movement in the intransitive sense [an event] occurs; when we move an arm certain arm movements take place." So, for example, if someone bent his finger, then ' e_a ' stands for a token occurrence of the event type, a bending of his finger. I shall sometimes refer to the token action as 'his

⁴ Jennifer Hornsby (1980: 2) thinks of an action as what makes a sentence with the form ' $aVTb$ ' true ('Simon moved his hand'), and an event as what makes a sentence with the form ' bVi ' true ('Simon's hand moved'). '*a*' names a person, '*b*' a bodily part, '*V*' stands for a verb, and the subscripts '*T*' and '*I*' for intransitive and transitive occurrences of the verb respectively.

a_e -ing', to mark the fact that token event e is meant to be the event whose existence is supposedly entailed by the a -ing.

Different writers use different terms for this event, ' e_a '. Von Wright (1977: 39–40) calls it 'an event-result', and Maria Alvarez and John Hyman (1998: 233) adopt his terminology. Davis (1979: 5) calls it 'a doing-related event'. Bishop (1989: 105) says that it is the event that is 'intrinsic' to the action. Others have called it 'the associated event'. The intransitive sense of the verb describes the intrinsic or associated event. I shall follow Bishop's terminology. Actions, including basic actions on this view (more on basic versus non-basic actions below), can be 'paired' with intrinsic events. On this view, if an action occurs, it is necessary that the intrinsic event does.

Of course, Hornsby does not say that all intuitively action-verbs are transitive, and indeed many are not. The verb 'ran', for example, has both a transitive and an intransitive sense, and both senses indicate action. In the intransitive sense, in reply to 'how did you get there?', I can reply: 'I ran'. This running of mine is an action, even though 'ran' is intransitive in this usage.

For any specific action, how can one determine which is its intrinsic event? In the cases that I will discuss, the pairing of action and intrinsic event is obvious and unproblematic, since the act-description and the event-description are closely grammatically related: my moving of my finger and my finger's moving; my raising of my hand and my hand's rising. But there are cases which are not so obvious, and this is especially true in the case of action verbs like running, speaking, or spitting, which are intransitive or have an intransitive sense, because the obvious grammatical transformations are unavailable. If I run, the event intrinsic to this running is something like: my legs moving in such-and-such characteristic way.⁵

If this logical or conceptual entailment were to hold, as (2) asserts, the metaphysical nature of the relationship between basic action and intrinsic event is still open in many ways, and my use above of 'pairing' was meant to be vague and hence neutral in this regard. The fact of that entailment is compatible, for example, with the relationship between token action and its intrinsic event being metaphysically contingent. Here are some possibilities for that relationship:

- (3) The action, his a_e -ing, causes its intrinsic event, e_a .
- (4) The event, e_a , is metaphysically or de re necessary but insufficient for the action, his a_e -ing. For instance, the event might be a proper part of the action.
- (5) The action, his a_e -ing, supervenes on, but is not identical to, the event e_a .
- (6) The action, his a_e -ing, is identical to the event, e_a .

It is easy to miss the fact that the entailment mentioned for example by Hornsby and Hamlyn is consistent with there being only a causal and hence contingent relation between basic action and its intrinsic event.⁶ The necessity required by the entailment might only be de dicto, arising from the ways in which the cause and effect are described. Just as 'the cause of b caused b ' can be a metaphysically contingent causal

⁵ Davidson noted this fact about 'walks'. Quoted in Bach (1980: 114–20).

⁶ I think it is easy to miss, because I once missed it. See Ruben (2003: ch. 2).

claim ('the cause of *b* might not have caused *b*'), despite the de dicto necessity that arises from the descriptions used for cause and effect, so too 'his *a_c*-ing caused the event *e_a*' ('His bending of his finger caused his finger to bend') could be a genuine causal claim, on (3) at least, for the same reason.

(3) is dualistic, in the sense that it counts the action and the intrinsic event as two distinct existences. (6) is monistic, in the sense that for it there is only one item, describable as either an action or an intrinsic event. (4) and (5) represent a sort of compromise, since for these views there are not two distinct existences and yet there are two non-identical ones.

(3) is the view of Jennifer Hornsby (1980) and Paul Pietroski (2002), and echoes an earlier view of Pritchard's (1949). For Hornsby, this is summed up in the claim that "movements_{*t*} cause the body to move_{*t*}. And actions are movements_{*t*}" (1980: 13). Pietroski says of basic actions:

[I]f Nora moved_{*t*} her finger, then Nora's action caused the moving_{*i*} of her finger. Her action caused the finger motion, just as Nora's action causes the melting_{*i*} of the chocolate. A similar treatment of 'raised_{*t*}' suggests that if Nora raised_{*t*} her hand, her action caused the rising_{*i*} of her hand. (2002: 35)

Actions, for Hornsby and Pietroski, are 'internal' tryings (or, for Pritchard, willings) which cause external bodily movements and, via those bodily movements, other, non-bodily motions in the external world. 'Nora's moving of her finger' is the name of a basic action, a trying. (Of course, on their 'austere' view of act identity, lots of other descriptions might also be true of the trying, typically, physical action descriptions.) 'The moving of Nora's finger' is the name of an external bodily event. As they sometimes say, actions like Nora's moving of her finger are inner causes of external motions like the moving of Nora's finger.

(6), on the other hand, is claimed by many causal theories of action. If one holds that actions just are events non-deviantly caused by prior rationalizing mental states, then the action, his bending of his finger, just is the event, the bending of his finger, if the latter is so caused. The view is common and well-known. There are also many objections to it, some of which I have set out at length elsewhere (Ruben 2003), but which I do not rehearse here.

The simplest way in which to understand (4) is to construe the action to be the combination or mereological sum of his *a_c*-ing and its intrinsic event, *e_a*. If, for example, his *a_c*-ing is a mental item of some sort to be determined (but not itself the whole action), the action commences when that mental item occurs and terminates when the intrinsic event, *e_a*, finishes. In that sense, the intrinsic event would be a part of the larger action.

There is an assumption common to all four positions, (3)–(6):

(7) All physical actions, including actions like my moving or raising of my hand or my bending of my finger, have events intrinsic to them.

Once (7) is accepted, the question addressed by (3)–(6) arises, namely, the relationship between actions and their intrinsic events.

Now for the second piece of terminology. What is a mere event? A mere event is a token event which is not intrinsic to any action. So, from the fact that a token mere event of type *E* occurred at *t* involving an agent's body, and no other token event of type *E* involving that agent's body occurred at the same time, it follows that that agent performed at *t* no token action of type *A_E*. Note that the concept of a mere event (like the idea of an intrinsic event) requires the prior concept of an action, since it is characterized by being an event that is not intrinsic to one.

Mere events include events like the explosion of a volcano, the collapse of a star, the gravitational pull of the moon, or the lapping of the waves. No tokens of those event types are going to count as events intrinsic to any action. One might dispute whether it is necessarily true or only contingently true that no tokens of these types are intrinsic events, but on either view, it is at least a truth that none are.

But the tokens of other event types sometimes are mere events and sometimes are events intrinsic to an action. The death of the Queen can be a mere event when she dies of natural causes, but can be an event intrinsic to my action if I kill her.

So too it is with the waving of my hand or the moving of my left arm. When my left arm moves or my left hand waves as a result of a neurological spasm, these are mere events. On the other hand, I may wave my left hand by moving it with my right hand. When I bring about the waving of my left hand or the moving of my left arm by manipulating my left hand or arm by my right hand, then the waving of my left hand and the moving of my left arm are events intrinsic to my (non-basic) action, my waving of my left hand or my moving of my left arm. (I use the idea of basic and non-basic action at this point, relying on readers' intuitions: an action of mine is not basic if I do it by doing something else.)

2 A DISJUNCTIVE THEORY OF ACTION

Let's accept that there are events, although even this is itself hardly beyond dispute.⁷ Both mere events and intrinsic events belong to the same fundamental type, being an event; or, (modulo some controversial assumptions about event identity) what is an intrinsic event could have been a mere event, even if the converse is not always true. There are many competing accounts in the literature of what it is for something to be an event; what I say here is intended to be neutral between all such accounts. Jonathan Bennett (Bennett 1988: chs. 4–8) offers an excellent survey of some of those options. I make commitments about the concept's extension, but certainly not about its intension. I have no, and hence offer no, analysis of eventhood.

If (2) were true, whenever there is an action, there would be an event intrinsic to that action. So, just as on a non-disjunctive theory of perception, in the cases of both a veridical perception and a hallucination, the appearance or look is meant to be the same and common to both, so too on a non-disjunctive theory of action, we could say

⁷ See for instance Horgan (1978: 28–47).

that in the case of both a mere event and an action, something belonging to the same fundamental type occurs and is common to both. If a mere event occurs, trivially, an event occurs; if an action occurs, then, on (2), an event intrinsic to that action occurs, so an event occurs (even if that last event is not identical to the action). There is a kind of commonality between actions and mere events, since in either case something of the same fundamental type occurs, namely an event.

Let me now prepare the grounds for introducing a disjunctive theory of action. We broke off the discussion of (2) above, wondering about the relationship between actions and intrinsic events. If these actions and their intrinsic events are neither identical nor non-identical (because I think all the options canvassed thus far are unattractive), what other option might remain open to us, for understanding the relationship between actions and their intrinsic events?

I think a possible suggestion is this: to deny the inference claimed by (2), and thus to deny (7). So, pace (2) and (7), there is no element common to the occurrence of both mere events and basic physical actions, via the events allegedly intrinsic to the latter, because there are no such events in this case. That is, I claim that in an important sense there is no such logical or conceptual entailment as (2) claims. We should deny the existence of any such intrinsic events, in certain (but not all) cases of physical action. Unlike (6), which is a reductionist theory, I am proposing what might be thought of as an eliminativist one. For cases of basic physical action, I want to eliminate intrinsic events. By denying their existence, I thereby deny the commonality assumption for basic physical actions and events.

For these cases of physical action, namely, basic ones, nothing occurs that is fundamental and in common to both these actions and events. The question of the nature of the relationship between basic physical action and intrinsic event would then never even arise, just as: if we deny the existence of a look common to veridical perception and hallucination, the question of the relationship between appearance and object would never even arise. I am not just denying that the basic action and its intrinsic event are identical. I deny that any such intrinsic event occurs at all, even one necessary but insufficient for the basic action; a fortiori, if there is no intrinsic event, it can't be identical to anything.⁸ No identity of basic action and intrinsic event, and no duality either, because no intrinsic event.

Other philosophers have also denied (7). For example, Hugh McCann (1998) says that the denial of events intrinsic to basic actions is essential to halt a vicious regress. Wittgenstein asked: "What is left over if I subtract the fact that my arm goes up from the fact that I raise my arm?"⁹ Wittgenstein's question can arise, says McCann, whenever an action has an event intrinsic to it. If an action (my killing of the Queen) has an event intrinsic to it (namely, the Queen's death), which intrinsic event is itself

⁸ Someone might confuse my thesis with the question of inactions, omissions, and forbearances. Most theorists count these as actions. I suppose that, if it is an action, my intentionally keeping quiet is a basic action, since it is not an action I do by doing something else. Whatever we might decide about such cases of 'negative' action or inactivity, my claim is about 'positive' action that evinces activity.

⁹ Wittgenstein's question is quoted in McCann (1998: 75). All my page references are to McCann (1998).

brought about by a more basic action of the agent (my pulling of the trigger), or by an event intrinsic to that more basic action, we shall finally get to an action which can't be given this sort of analysis, on pain of infinite regress. That action we might call a basic action. That action will fail to have an event intrinsic to it. (Of course, there are in principle other ways to halt the regress that perplexed Wittgenstein, which McCann does not consider.)

As McCann (1998: 87) says: "If there is no result [intrinsic event] to be distinguished from the action, there can be no question as to what makes it a result." So McCann and I would both deny that basic actions have events intrinsic to them. Call such basic actions 'pure'. Indeed, this provides part of the account I would give for the basicness of an action: an action is basic only if it is pure, and a physical action is basic iff it is pure. McCann claims to have found actions that have no events intrinsic to them, and hence which stop Wittgenstein's regress, but only in the case of certain mental acts: for example, acts of thinking and willing (to which we might add tryings, on some views). "That acts of thinking do not have results [events intrinsic to them] means there can be no action-result problem about thinking" (McCann 1998: 87).

McCann and I agree that (7) is false. I could also agree that there are such things as pure mental actions (although nothing in this essay requires that view). My disagreement with McCann is that I do not think that we have to hold that there is pure mental action, in order to find cases of basic action. Some physical actions are pure as well. I think that there are basic physical actions that are pure, that have no events intrinsic to them. (The fact that there may be some pure mental actions, but which may not count as basic actions, is why purity of an action, for me, cannot be a sufficient condition for basicness *simpliciter*).

Since for McCann these willings and thinkings (or other mental actions like them) are the only pure actions there are, he must claim that every action is done by or is caused by (or, on an 'austere' view of act individuation, is identical to) an act of willing or thinking or similar. On this view, it is from willing or thinking that all other cases of agency stem. It locates basic action 'inside' the agent, in his mind. McCann's account, like Pritchard's, is Cartesian. If Hornsby locates actions inside the brain or nervous system, her view would be best understood, I think, as a physicalist solution to a Cartesian problem. My view locates (at least some) pure, basic action at the bodily surface of the agent. We need not, on my theory, be driven inside the agent to find an act of will or a brain event as a stopping place to account for action.¹⁰ My account has no need to solve the apparently Cartesian location problem for action. My account is Aristotelian in its inspiration.¹¹ I find purity at the surface of an agent's body.

Presumably, assuming that there are lots of things an agent does, there will have to be, on McCann's theory (and on Pritchard's), an awful lot of willings or other pure mental acts around as well: "bodily actions typically involve the mental activity of

¹⁰ The only discussion I know of this occurs in Candlish (1984: 83–102).

¹¹ Or even Marxian. See Marx (1965) for a view of action which, at least under one interpretation of those cryptic statements, I find congenial. The influence of Aristotle on Marx, probably via Hegel, is profound.

volition or willing and . . . this activity is . . . the basic activity we engage in when” we physically act (McCann 1998: 75). Other things being equal, it is unattractive to have to hold that, on every occasion on which I physically act, on every occasion on which, for example, I bend a finger, raise my arm, or open my hand, I am engaging in some mental act, as this view requires. Or, to word substantially the same criticism differently so that it will fit the Hornsby–Pietroski ‘austere’ view of act individuation: it is unattractive (so I say, anyway) to have to hold that on every occasion on which I physically act, it is also true that my act is a trying.

So I hold, contrary to McCann:

- (8) Some physical actions are pure, and have no event intrinsic to them.

What do I mean when I say that when a basic physical action occurs, no event intrinsic to that action occurs? To repeat, my view is that, when Nora moves her hand (and not by doing something else), there is *in one sense* no such event as her hand’s moving that occurs. All that there is, is the action. What do I mean by: “there is in one sense no such event as her hand’s moving that occurs”? Why the qualification, “in one sense”? I think that the philosophical literature uses the idea of an event, and specific event type names, like the moving of a hand, ambiguously and we need to take account of this ambiguity in order to see the force of the disjunctive theory of action.

If we agree that when a basic physical action occurs, no event intrinsic to it occurs, we can still ask this question: does any event occur when a basic physical action does? Even supposing (8) to be true, the answer to that question could still be in the affirmative, in the following way. One might hold that actions are a type of event as a result of what is in effect a classificatory decision. Let’s call this the classificatory or wide concept of an event. Actions may simply be classified as a subclass of events, by stipulation of the concept’s extension. The wide concept of an event has, in its extension, mere events (like the eruption of Vesuvius), events intrinsic to non-basic actions (like the death of the Queen when I kill her by shooting her), and actions themselves, basic and non-basic. This concept of event is compatible with a disjunctive theory of action.

Similarly, the event type, the movings of a hand, taken in this wide or classificatory sense, will have the movings of a hand which are mere events (for example, my hand moves as a result of your pulling it), the movings of a hand which are intrinsic to non-basic actions (for example, my hand moves because I move it by twisting it with my other hand), and basic actions (for example, I move my hand but not by my doing something else), all in its extension. I do not think that this wide or classificatory sense of event is a sense merely invented or stipulated by me. Many philosophers rather uncritically take actions to be a subset of events, without argument, and when they do this, they are best interpreted as using the wide or classificatory sense of event in making that claim.

When they use the concept of an event in this wide way, I don’t think it is (2) and its likes that they have in mind. They are not arguing that actions are events because they are identical to their intrinsic events; that would require a substantive

philosophical argument. They are simply, and often uncritically, classifying actions as a type of event in their own right. There would be no philosophical ‘finding out’ if actions are events on the wide concept. They would be, by ‘definition’, events. Since this is only a question of classification, and not one of substantive philosophy, no honest philosophical labour is required for this result. On this wide concept of event, if an action occurs, even a basic physical action, it follows trivially that an event occurs—namely the action itself—but it does not follow that an event intrinsic to the action does.

So this introduces a complication in the comparison between disjunctive theories of perception and action. One disanalogy between the disjunctive theory of perception and that of action is that in the latter, there is an ambiguity for which there is no counterpart in the former: the concept of an event and the concept of (some) specific event types. In contrast to the wide or classificatory concept of an event described above, there is a narrower, more philosophically specialized sense of that concept. In that second sense, only mere events and events intrinsic to actions (if and when there are such) are events. In this narrow sense, actions would count as events only if identical to their intrinsic events.

As I said earlier, mere and intrinsic events are of the same fundamental kind. The extension of the narrow concept of event includes these two kinds of items. The various different accounts of eventhood in the literature are well-known and the reader can plug in at this point, for this sense, whatever might be his favourite account. It seems to be that many of the philosophers who have held that actions are not events mean that they are not events in this second sense.¹² Both concepts of event, the wide or classificatory and the narrow, seem to be legitimate and useful.

On the non-disjunctive view of action that I reject, what would be ‘common’ to both mere events on the one hand and to actions (which on this view would have intrinsic events) on the other would be the occurrence of an event in this narrow sense. If basic physical actions are to be identified with their intrinsic events, as (6) claims, then the former after all do fall into the extension of the narrow concept of an event since the latter do. If basic physical actions have intrinsic events but have a relationship with them other than one of identity (for example, as (3)–(5) propose), the commonality is different, but still a type of commonality: when a basic physical action occurs, an event in the narrow sense, intrinsic to that action, also occurs, an event to which the basic action is related but with which it is not identical.

¹² See for example Bach (1980: 114–20). I find Bach’s view hard to understand. If actions were causings, and causings are ‘instances’ (his word) of a relation, I do not see why causings would not count as individuals. Bach uses as illustration the idea of a round, red ball, and says that it would be ludicrous to ask how many instances the ball displayed. I am not so sure. If Iago loves Desdemona and Romeo loves Juliet, there is a perfectly good answer to how many lovings, or instances of love, are under discussion: two. Similarly, if there are two red balls, there are two instances of the colour, red. I see no objection to taking instances seriously, as individuals. Bach thinks that actions are causings, and that causings are not events. Causings, for him, are instances of a relation, and can neither be quantified over nor be individuated. For Alvarez and Hyman (1998), for various reasons they give, actions are not events but there is no suggestion on their view that actions do not count as a kind of individual, parallel to but different in kind from events.

On the other hand, one can express my disjunctive view of action as an exclusive disjunction, with three disjuncts (and, in terms of truth-making, the disjunction is to be read right-to-left):

(9) *i* is an event (in the wide, classificatory sense) iff *i* is an action or *i* is a mere event (in the narrow sense) or *i* is an event (in the narrow sense) intrinsic to an action.

Since I have characterized mere events simply as the complement class of intrinsic events, there is good reason to take (9) as exhaustive as well as exclusive. But if there is someone who thinks there is another disjunct that can be added to (9), I have no reason to resist that move.

On the disjunctive view represented by (9), from the fact that no token mere event of type *E* occurs to an agent's body at *t*, it does not follow that the agent has not performed at *t* a token action of type *A_E* (recall that the occurrence of suitable mere events actually excludes corresponding action). And from the fact that no mere event of type *E* and no intrinsic event of type *E_A* has occurred to an agent's body at *t*, it does not follow the agent has not performed a basic action of type *A* at *t*.

One might describe (9) as a disjunctive theory of events, rather than of action, and surely it is at least that. But just as the disjunctive theory expressed by (1) can be called a disjunctive theory of perception rather than of experience, because it is a response to what is often called the problem of perception, so too (9) might be called a disjunctive theory of action rather than of events, because it is a response to what is often called the problem of action. Instantiations of (9) will apply for some specific event types as well as for the general concept of an event. Since (9) is intended as an exclusive disjunction, no event in the narrow sense occurs if a basic action does, and hence a fortiori no such event is common to both the case of basic action and the occurrence of an event in the narrow sense.

Why should we accept (9) and deny the existence of events intrinsic to basic physical actions? What counts in favour of this option, relative to the options sketched in (3)–(6)? Suppose (9) were not superior in any way to the other choices. All the same, I believe that it is important to show that this view can occupy a perfectly acceptable position in the logical space of alternatives. I know of no other place in which the purity, in my sense, of some physical actions is put forward as an alternative position.

But I do think that the main arguments for (9) are twofold: (a) I know of no consequences of (9) that raise insoluble difficulties, and (b) the greater unacceptability of the alternatives—either the identity or the duality of basic action and intrinsic event. All the alternative theories of action seem to have some unintuitive consequences. The location of action inside the agent's brain or mind is *prima facie* implausible. On the other hand, the identification of actions with bodily events robs us of the very activity of actions; it is hard to see how action can be constructed from the passivity of what just happens. (9) has neither of those disadvantages. Coupled with my disarming of what might be thought to be its implausible consequences, I regard the case I put forward, for the acceptance of (9) and the denial that there are events intrinsic to basic physical actions, as a strong one. I will not, in this essay, rehearse in any detail the points that seem to me to constitute a strong case for (b), points that I have made elsewhere (Ruben 2003).

So either a person basically moves his hand or his hand moves (and this is an exclusive disjunction). The first is a basic action with no event intrinsic to it; the second is a mere event or an event intrinsic to his non-basic action (he may move his one hand by using the other to move it). So if I move my hand (as a basic action), it is false that my hand moved or that it changed place (in the narrow sense), only true that I moved it or that I changed its place.¹³ On the flip side, if the event of my hand's moving or changing place does occur (in the narrow sense), then it is false that I basically move my hand. When *p* moves his hand, or when he changes its place, and these are basic actions, there is no event that occurs intrinsic to *that* basic action. Other events may occur, and it may be that his basic moving of his hand is itself an event as well as an action, in the wide, classificatory sense, but there is no hand moving which is intrinsic to my moving of it, and hence no hand moving in the narrow sense.

Does this view have any unacceptable consequences? There are two pieces of data, (10) and the inferences between (11), (12), and (13), which I want to try to account for and which appear to be inimical to my view.

The first piece of data is this. Is it really *false* that when I move my hand, my hand moves? Surely an observer, who sees me move my hand, may say:

(10) "I am sure that I saw his hand move. Whether he moved it, or someone or something else did, I don't know."

What might the thought be behind (10)? The observer may be in doubt whether he has observed an action, even a basic action, or a mere event (say, the hand's moving as a result of a nervous spasm). But if the observer can be in doubt, then the two possibilities that (10) canvasses must be perceptually indiscriminable, and that must be because whatever it is that is observed is 'common' to both possibilities. In either case, whether it be his hand moving or his moving of his hand, what the observer would have observed is the same item, an event, and that item is 'neutral' between the cases of basic action, intrinsic event, and mere event. If what is observed is a mere event, then an event is observed. If what is observed is a basic action, then what is observed is also an event—but in this case, an event intrinsic to the basic action. Construing the data in this way will certainly refute my disjunctive theory by finding an element common to basic action, intrinsic event, and mere event.

The disjunctive theory of action is better placed to deal with this issue than is the disjunctive theory of perception. What the disjunctive theory of action must say is this: from the fact that a basic physical action, an intrinsic event, and a mere event

¹³ Fiona Macpherson has suggested the following. There might be a straightforward case in which this is the natural description: suppose I am on a train travelling in direction A at 10 mph and that I move my hand in the direction opposite to A at 10 mph. In terms of objective location my hand has not moved. However, it does seem right to say that I moved my hand. This seems not really to be especially relevant to action and event, for the same thought can arise in the case of two mere events: can there be a mere moving of my hand without a mere changing of its place (in the same circumstances she describes, for example)? The thought is an interesting one and surely much about its answer depends on questions concerning absolute versus relative space. But as it is not particularly a thought experiment relating to the relationship between actions and events, I don't think it will help me here.

might all be perceptually indistinguishable, it does not follow that they must have something philosophically significant in common to explain that indistinguishability. Of course, they might have these three features in common: being (or entailing the existence of something which is) perceptually indistinguishable from an action, from an intrinsic event, and from a mere event. But, as in the case of being perceptually indiscriminable from a veridical perception and from a hallucination, these commonalities are trivial.

Is there anything else that a disjunctive theory of action could say? Could it explain why one of these items might be mistaken for another, in the way in which hallucinations might get mistaken for veridical perceptions? A disjunctive theory of action can say more, because it can say exactly what a disjunctive theory of perception cannot say: a person may be in error or doubt or confusion about whether he is observing a mere event, intrinsic event, or action, because of the common look shared by all three sorts of occurrences. They may have nothing philosophically significant in common that would refute (9), but they can all have a certain look in common. They look alike.

To that degree, actions and events are like gold and fool's gold, or water and twin-earth water. Gold and fool's gold look alike, and water and twater look alike, but phenomenological appearance is no guide to philosophically or scientifically significant classification. This is exactly what a disjunctive theory of perception cannot say, since it is dealing precisely with phenomenal character itself, and not the significant features of things beyond the appearances. Phenomenology is a bad guide to kind-hood, at least for minerals, compounds, events, and actions. Two observationally or perceptually indiscriminable items, for example an action and a mere event, may be perceptually indiscriminable because they look alike, yet be of fundamentally different kinds.

We have been speaking until now of how an action, an intrinsic event, and a mere event might all be perceptually indiscriminable to an external perceiver of actions and events. But the perceptual indiscriminability for the disjunctive theory of perception is indiscriminability by the perceiver of his own experiences or states. Is there an analogue to this for a disjunctive theory of action? Various things may happen to, or be done by, an actor: he might raise his arm as a basic action; he might raise his arm by his doing something else; his arm might rise without him raising it at all. Might he himself be unable to discriminate amongst those possibilities, or anyway between the first and third, whatever the case might be for an external observer?

William James (1890: 489–92) cites two examples of just such a case. (a) Professor Strümpell's description of his "wonderful anaesthetic boy": "Passively holding still his fingers did not affect him. He thought constantly that he opened and shut his hand, whereas it was really fixed." (b) Dr Landry's account of a blindfolded patient: "If, having the intention of executing a certain movement, *I prevent him*, he does not perceive it, and supposes the limb to have taken the position he intended to give it."

If opening and closing one's hand or "executing a certain movement with a limb", at least on the occasions being considered, are basic physical actions, both Strümpell's and Landry's cases are ones that are or can be made pertinent to our discussion. In both cases, the person has done nothing, or anyway executed no physical action, but

falsely believes that he has done something, and indeed falsely believes that he has a-ed, where 'a-ing' is a basic physical action description. That is, the agent believes that 'my a-ing' refers to his token basic physical action when it in fact refers to nothing. The wonderful anaesthetic boy falsely believed that he opened his hand; Landry's patient falsely believed that he had executed a certain movement with a limb. In both cases, no physical action occurred, basic or otherwise.

The examples would make the same point if, for example, the wonderful anaesthetic boy's hand had been opened and shut after all, not by him, but rather by some force made invisible to him. The boy, as long as it is not observable by him that he was not responsible for his body's movement, might falsely come to believe concerning some mere physical occurrence involving his body, perhaps the opening and closing of his hand, that it was a basic action of his, when it in fact was not. James's counterexamples do not have to concern beliefs about action rendered false by total lack of movement; the counterexamples could also concern beliefs about action rendered false by the occurrence of mere bodily movements that are not actions.

So not only can the agent himself falsely believe that he is basically physically acting when nothing relevant occurs at all. Concerning some mere event or happening, perhaps the opening of his hand, a person could falsely believe that that mere event is his basic physical action (his opening of his hand) when it is not an action at all. So an agent might not know which of the two disjunctive situations, action or mere event, that he is in, just as a perceiver might not know which of the two disjunctive states, veridical perception or hallucination, that he is experiencing.

In perception, a perceiver might mistake hallucinations for veridical perceptions but could also mistake veridical perceptions for hallucinations. The cases cited by James show or can be made to show that an agent might mistake mere events for his actions. Could an agent also mistake his actions for mere events? I believe that this is also possible, and have described such cases elsewhere (Ruben 2003: 218–19).

The second piece of seemingly recalcitrant data is this:

(11) 'Simon moved his hand' entails

(12) 'Simon's hand was moved by Simon.'

And (12) 'Simon's hand was moved by Simon' surely entails

(13) 'Simon's hand moved.'¹⁴

So doesn't the action, his moving of his hand, entail that there was an event, the moving of his hand, after all? The second inference looks as if it is merely a case of simplification. If so, there is surely no exclusive disjunction, since from the truth that Simon moved his hand, it follows that there was a moving of his hand.

I want to argue that there is a sense of 'moving' in which the inference goes through and a sense in which it fails, but in that first sense the success of the inference does not damage my view. The sense of 'moving' for which it fails is the narrow sense; the sense in which it goes through is the wide or classificatory sense.

¹⁴ Berel Lerner pointed these inferences out to me.

In fact, there is a missing step in the argument above. It is this: (12) entails (12a) 'Simon's hand was moved', by simplification. (12) and (12a) are both about action, even though they are in the passive voice. (12a) is of course elliptical, since anything that was moved was moved by something or someone.

That there is a sense in which the inference is valid hardly needs arguing. It will surely strike the reader as obvious. We can allow that the inference is valid if we take movement or event in the wide or classificatory sense, but in that sense the admission will be harmless to my position. In the wide sense, (11), (12), (12a), and (13) are all made true by actions, even though (12) and (12a) are in the passive voice; if (13) uses 'moves' in the wide sense, what makes (13) true is the truth of the action disjunct in that disjunction which gives the extension of event or movement in the wide sense.

But let me try to convince the reader that there is a sense of 'moving' in which the inference fails and is invalid. If (13) uses 'move' in the narrow sense, I would say that in the case in which his hand was moved by the agent, (12), and hence in which his hand was (elliptically) moved (12a), (13) does *not* follow. It does not follow that it moved (in the narrow sense), for there is no event, a moving of the hand, in the narrow sense, which could make (13) true in that sense.

So, if Simon's hand was moved by him, it did not move in the narrow sense. (12a) does not entail (13), in the narrow sense of event or the narrow sense of the moving of his hand.

REFERENCES

- Alvarez, M. and Hyman, J. (1998) 'Agents and Their Actions', *Philosophy*, 73: 219–45.
- Bach, K. (1980) 'Actions Are Not Events', *Mind*, 89: 114–20.
- Bennett, J. (1988) *Events and Their Names*, Oxford: Oxford University Press.
- Bishop, J. (1989) *Natural Agency*, Cambridge: Cambridge University Press.
- Candlish, S. (1984) 'Inner and Outer Basic Actions', *Proceedings of the Aristotelian Society*, 84: 83–102.
- Child, W. (1992) 'Vision and Experience: The Causal Theory and the Disjunctive Conception', *Philosophical Quarterly*, 42: 297–316.
- Dancy, J. (1988) (ed.) *Perceptual Knowledge*, Oxford: Oxford University Press.
- Davis, L. (1979) *Theory of Action*, Englewood Cliffs, NJ: Prentice-Hall.
- Goodman, N. (2005) *Fact, Fiction, and Forecast*, Cambridge, MA: Harvard University Press.
- Hamlyn, D. (1990) *In and Out of the Black Box: On the Philosophy of Cognition*, Oxford: Blackwell.
- Hinton, J. M. (1973) *Experiences*, Oxford: Oxford University Press.
- Horgan, T. (1978) 'The Case against Events', *The Philosophical Review*, 88: 28–47.
- Hornsby, J. (1980) *Actions*, London and Henley: Routledge & Kegan Paul.
- James, W. (1890) *The Principles of Psychology*, vol. 2, London: Macmillan & Co.
- Lowe, J. (2000) *An Introduction to the Philosophy of Mind*, Cambridge: Cambridge University Press.

- McCann, H. (1998) 'Volition and Basic Action', in his *The Works of Agency*, Ithaca, NY: Cornell University Press, 1998: 75–93. [First published in *Philosophical Review*, 20 (1974): 451–73.]
- McDowell, J. (1982) 'Criteria, Defeasibility, and Knowledge', *Proceedings of the British Academy*, 68: 455–79
- Martin, M. G. F. (2004) 'The Limits of Self-Awareness', *Philosophical Studies*, 120: 37–89.
- Marx, K. (1965) *The German Ideology (& Addenda)*, London: Lawrence & Wishart.
- Pietroski, P. (2002) *Causing Actions*, Oxford: Oxford University Press.
- Pritchard, H. A. (1949) 'Acting, Willing, Desiring', in his *Moral Obligation*, Oxford: Clarendon Press.
- Ruben, D.-H. (2003) *Action and its Explanation*, Oxford: Oxford University Press.
- Snowdon, P. (1980–1) 'Perception, Vision, and Causation', *Proceedings of the Aristotelian Society*, 81:175–92.
- Wright, G. K. von (1977) *Norm and Action*, London and Henley: Routledge & Kegan Paul.

A Disjunctive Conception of Acting for Reasons

Jennifer Hornsby

Bertrand Russell once said “There are reading parties and reading parties, and they are mainly the latter.” I take it that Russell wanted to convey that not all of the things that pass as reading parties conform to a certain inflated conception that one might have of them. Someone of disjunctivist persuasion says that there are experiences and experiences. The message then is that not all experiences conform to the deflated conception that a philosopher might have of them. Some reading parties are constituted of tracts of intellectual excitement, but Russell wanted us to know that plenty are not. A proponent of disjunctivism wants us to recognize that, despite the fact that some of our experiences might not be constituted of facts being manifest to us, plenty actually are. This disjunctivism evidently makes a claim in the philosophy of perception.

Here I want to promote a different disjunctivism, which belongs in the philosophy of action: it concerns acting for reasons. I shall suggest that it can play a role analogous to that of the disjunctive conception that John McDowell recommends in thinking about perception.¹ So I say that there are cases and cases of acting for reasons; and not all of them conform to the conception that philosophers have of them when they think of actions as explained by beliefs and desires. When this disjunctivist conception is treated as a sort of counterpart to the disjunctive conception of appearances, the two can be shown to have work to do in combination. And my hope is that McDowell’s disjunctive conception emerges in good standing when it is thought about alongside the disjunctive conception of acting for reasons.

I acknowledge the support of the AHRC: I was a beneficiary of their Research Leave Scheme when I completed this essay.

To Adrian Haddock, Guy Longworth, and Scott Sturgeon, I owe thanks. All of them gave me useful comments on, and discussed with me, drafts of material here. To Adrian, I owe double thanks: he was the commentator for my paper at the Glasgow conference when he set me straight on something crucial, and he has helped at every stage since.

To Jonathan Dancy, I owe an apology. I’ve rewritten my essay even since I gave him what I then took to be the final version. I hope that there will be another occasion when my delays won’t prevent us from engaging with one another.

¹ McDowell (1982), and subsequent of his writings referred to in the bibliographies of other chapters in the present volume. (In this context and for my purposes, there will be no need to cite chapter and verse.)

The first thing I have to do is to show that a disjunctive claim about acting for reasons makes a point in its own right. I think that we need a view which gives a role to *knowledge*. In section 1 of what follows, I motivate and demonstrate the role of knowledge in acting for reasons, and I show that a disjunction can record its role. In section 2, I show why this might be a fitting way to record it. It is here that I make use of parallels between the two disjunctive conceptions. With the parallels in place, I return to questions about acting for reasons. In section 3, I find fault with various accounts of acting for reasons which lose sight of the connection that there is between (on the one hand) considerations that favour our acting in particular ways and (on the other hand) considerations that move us to act in particular ways. I suggest that a disjunctive conception safeguards the connection.

1

If you present a philosopher with the question “What is a reason for acting?”, then you may find you get different answers according as you ask it of a philosopher who specializes in ethics or of one who specializes in philosophy of mind. The two different answers need to be brought into harmony with one another. A satisfactory account of human agency ought to reveal human beings’ capacity to act for reasons as an aspect of both their ethical and their psychological nature. And that means that we need an understanding of reasons for acting which suits both ethics and philosophy of mind. My claim will be that we start to achieve such an understanding when we see human agents as knowledgeable, when we take ourselves to know things.

The two different answers to ‘What is a reason for acting?’ correspond to two conceptions of such reasons (section 1.1). I shall say something about how the two conceptions *cannot* be distinguished (section 1.2) and something about how they are connected (section 1.3). The aim here is to show that a disjunctivist claim illuminates a connection between them (section 1.4).

1.1 The two conceptions of reasons for acting

Jonathan Dancy (2000) gives expression to the idea of a reason for acting as it plays a part in ethics when he says “To be a reason for action is to stand in a certain relation to action, and the relation . . . is that of favouring.”²

The idea evidently has application beyond such reasons as one might be inclined to think of as ethical ones. Here are some examples. ‘Why is Jill heading south? She’s agreed to meet Jack down the road.’ ‘Why are you walking on the common? The dog needs a walk.’ ‘Why is Edna skating at the edge of the pond? The ice in the middle is very thin.’ In each case, we learn something which can be expressed in the language of reasons. The fact that Jill has agreed to meet Jack down the road is a reason she

² Like Dancy, I want an account of acting for reasons which allows that people may act for reasons in this sense. Unlike Dancy (as will become clear), I don’t think that such an account requires us to find any fault with psychological reason-explanations. In what follows, I confine mention of specific disagreements with Dancy to footnotes.

has for heading south. The fact that the dog needs a walk is a reason you have for walking on the common. And so on. In the presence of the relevant facts, one sees that there's something to be said for some agent's acting in a certain way—for Jill's heading south, or whatever. One sees that a person's acting in that way may be favoured. This conception of reasons can be set out schematically, thus:

(F) A reason for X to Φ was [the fact] that p .

What about the philosophers of mind's conception of a reason for acting? Well, many of them endorse a qualified version of a principle sometimes called the belief-desire-action principle—a principle which connects psychological states that lead to acting in a particular way with acting in that way. It can be stated as follows.

(BDA) If X desires something and believes that Φ -ing will help secure it, then X will Φ .

Nobody actually accepts (BDA) as it stands; it obviously needs massive qualification. But it can be used to convey the idea that underlies the claims of philosophers whom I shall call belief-desire theorists. They say that a belief-desire pair which is related to Φ -ing in accordance with (BDA) *is* a reason for Φ -ing.³ And they think that such reasons have a place in causal explanations. For illustration, take the case of Jill who was heading down the road. The reason Jill had for doing so, according to these theorists, is her belief that she'll find John at the bottom of the road combined with a desire to meet John, as agreed; and Jill heads where she does *because* she has this reason.

Often we mention *only* someone's belief in stating a reason they had for acting.⁴ (We might explain Jill's heading down the road by saying just that she thinks that she'll find John.) Here I'm going to confine myself to what can be brought under the head of *belief*: I shall be concerned only with the cognitive side of things.⁵ So I make use of a notion of reasons which can be summed up with schema (B).

(B) X had a reason to Φ : X believed that p .

³ This conception of reasons is familiar from Davidson (1963): Davidson said there that a primary reason for action consists of a belief and a pro-attitude.

When a reason is thought of as a *pair* of a belief and a desire, reasons are individuated in relation to the goals or ends of action (to desired things). According to this mode of individuation, a person can properly be said to have 'more than one' reason to Φ only if they are in a position to satisfy two different desires simultaneously by Φ -ing (so that someone who had many reasons to Φ could kill many birds with a single stone). I use '*a* reason' differently: see next note.

⁴ When we mention only a belief and speak of this as '*a* reason', the use of '*a* reason' is such that various things are reasons for an agent to Φ on occasion, and something counts as a reason by virtue of being part of a reason in the sense of the previous note. I adopt this 'fine-grained' way of speaking of reasons throughout. What we speak of as a reason is much dependent upon contextual factors; and when we speak of '*the* reason', I think we usually mean something which, in context, is the most salient reason.

⁵ Questions about how much belongs 'on the cognitive side of things' would need to be addressed in a full account. But I needn't take sides on this here. In the hopes of showing that I don't prejudge those questions, despite my using (BDA) to arrive at one notion of a reason, I note (a) that the '*desire*' of (BDA) can be understood as a generic pro-attitude, and (b) that facts (or beliefs) may themselves be reasons for having pro-attitudes of various sorts.

(B) can be glossed by saying that *beliefs* are reasons for acting. And it has to be so glossed if the conception of reasons for action which is at work in philosophy of mind is to be brought in. But if we do speak of beliefs as reasons, then we must be alert to what ‘beliefs’ now means. Many belief-desire theorists have a habit of talking as if the beliefs that are reasons for acting were states of people’s minds. But in fact it is only relatively rarely that a belief in this sense constitutes a reason that someone has. Consider someone whose calling of the police is explained when it is said that she believes that she is constantly being followed. Her reason for calling the police is *not* her state of believing that she is being followed. The reason she has for calling the police is, as she would tell us, *that she is being followed*—which is something she believes. That this is so becomes amply clear if we contrast hers with a different sort of case, now of the relatively rare sort. It might be that a man has a reason to go to a *doctor* because he believes that he is constantly being followed: his believing this may be a symptom of paranoia. What now favours visiting the doctor—a reason he has for doing so—is *that he believes that he is being followed*. But by and large when someone has a reason by dint of believing something, their reason itself is not something psychological. Rather, their reason is given, as it is in the more usual case of the woman who called the police, when it is said what they believe.⁶

We now have the two answers to the question *What is a reason for acting?* Reasons for acting are given when *facts* are stated: let us call these ‘(F)-type reasons’. Reasons for acting are given when it is said what an agent *believes*: let us call these ‘(B)-type reasons’. No special reification of reasons is involved in using these labels: there is commitment only to there being correct instances of the schemas, and to ordinary talk of things that are reasons and of things that people believe. In what follows I use the labels ‘(F)-type’ and ‘(B)-type’ regularly. I don’t mean to commend the labels except in so far as they can be of use in keeping track of the occurrences of ‘reason’ deriving from the two schemas. Usually, of course, we just say ‘reason’, and usually it is clear enough what we mean.

Although (F)-type and (B)-type reasons are introduced as belonging to two different branches of philosophy, each corresponds to a perfectly everyday conception. We saw various examples in which we state facts in saying what reasons there were. At least as plentiful can be examples in which we say what someone believed in saying what reason they had. Suppose that Ann told Sam that the café she would meet him at was on Beech Street, although it isn’t there in fact. It dawns on Ann that she misinformed Sam. She takes steps to ensure that he knows where the café actually is. (Maybe she can reach him on his mobile phone.) Her thinking might be that Sam will soon be on his way to Beech Street: inasmuch as Sam believes that Ann will be there, he has a reason to be going there. In this example, Sam temporarily has a false belief, and Ann takes him therein to have a reason. But one’s belief does not have to be false for one to have a (B)-type reason. Even if Sam was right about where the café is, we could still adduce his belief about its location in saying what reason he had for

⁶ Although this point about what belongs in the category of reasons was made by Parfit (1997) (and see also, for example, Dancy (2000) and Hyman (1999)), many philosophers of mind pay no heed to it.

being on his way to Beech Street. Indeed an onlooker who was ignorant of where the café actually is could explain why Sam is headed towards Beech Street by supplying Sam's reason for heading there: Sam believes that the café at which he is meeting Ann is on Beech Street. This can be correct whether Sam is right or wrong about where the café is.

When true beliefs are seen to play the same sort of role in explanation as false ones do, we realize that, according to our ordinary understanding of reasons, (B)-type reasons are at least as ubiquitous as (F)-type ones. It is not at all surprising then that philosophers of mind should work with (B)-type reasons: one seems bound to think of beliefs as reasons if one's aim is to cover the ground with a single sweep and to account for action which springs from true and false belief alike.

The ubiquitousness of (B)-type reasons has led some people to speak sometimes as if people's having reasons for acting were (on the cognitive side) only ever a matter of their having beliefs. But we could not think of reasons for acting as all of the (B)-type. If you believe that the dog needs a walk, there may, or there may not, be an (F)-type reason for you to walk on the common. Perhaps unbeknownst to you, someone just took the dog for a long walk in the park, so that he doesn't need a walk and there is no reason for you to take him to the common now. Here there is a question about what reasons you have which isn't addressed when your (B)-type reasons are all taken into account.

1.2 Distinguishing between conceptions of reasons

Endorsing both conceptions of reasons for acting might appear to lead to a puzzle. For if we accept that whenever someone acts for an (F)-type reason, they act on a true belief, then we have to say that whenever there is an (F)-type reason for which a person acts, there is also a (B)-type reason for which they act. It could then seem as if someone who acted for an (F)-type reason had double the reason to do that which they did for that reason. And that might seem puzzling: are our actions overdetermined somehow when we act for reasons?

Well, we must remember that the label '(B)-type reason' attaches to something only inasmuch as an instance of schema (B) is correct. And we have seen that, except in the relatively rare cases, someone's believing something does not *provide* them with a reason. (The belief that he was being followed did provide the possibly paranoid man with a reason to go to the doctor. But it was not among the reasons of the woman who called the police that she believed that she was being followed; what provided her with a reason was the fact that she was being followed.) Thus (B)-type reasons need not to be thought of as, as it were, additional reasons—reasons beyond those that might be provided by the facts. Indeed if your belief is that things are a certain way, and in fact that things are that way, then what you believe (that *p*, say) can be said to be a fact (it's a fact that *p*). Still, this may not completely remove the puzzle. For even if there is no need to suppose that you might have two reasons for action just in truly believing something, there are nonetheless two different accounts of why you act when it is said that you acted for an (F)-type reason and that you acted for a (B)-type reason: you acted because *p*, and you acted because you believed that *p*.

In due course, we shall see that there really need be no puzzle here. But to be going on with, we need to rule out a way of avoiding the supposed puzzle. Someone might try to avoid saying that agents have ‘double the reason’ by denying that the agent who acts on a belief really has any reason at all for acting. Reasons for acting are all of them facts, this person says: when we know of an agent’s (B)-type reasons, we simply know *reasons why* they act as they do; but (B)-type reasons are not reasons *for which* people act. The ‘reason’ of ‘(B)-type reason’ is just the ‘reason’ of *explanation*.

But this is wrong. Of course (B)-type reasons do figure in reason-*why* explanations: an explanation is given, for instance, when it is said why Sam is going to Beech Street. And it is true that if Ann is not in the café on Beech Street, then there is actually no reason for Sam to be going there. But that doesn’t show that Sam fails to *have* a reason to go there. The reason Sam has for going there is known by the ignorant onlooker, who tells us what reason Sam has in telling us what Sam believes. And the explanation given of Sam’s going to Beech Street, when it is said what Sam believes, cannot be assimilated to any old reason-*why* explanation. Perhaps the reason why the bridge collapsed was that it had a structural flaw. Still, the bridge didn’t have a reason to collapse. The explanation of Sam’s behaviour and of the bridge’s collapse are of different kinds: in the sense of ‘have a reason’ in which it could be made clear to all that Sam *had a reason* to go to Beech Street, bridges never have reasons to do anything. (Even within the class of explanations that adduce psychological facts about *X* in answering the question ‘Why did *X* Φ ?’, we distinguish between those in which *X*’s *reasons* for Φ -ing are given and those in which they aren’t. ‘Because he believed that she had noticed him’ might tell us why he had blushed, for instance, even though there is no reason-explanation here: he had no reason to blush.)

We need then to think of a (B)-type reason as a reason an agent *has* for Φ -ing, and *not* as a reason that there is for Φ -ing. This is not to say, however, that a reason that there is for Φ -ing cannot be a reason an agent has, and indeed a reason for which they Φ . The crucial point is that the inference from ‘*x* has a reason’ to ‘There is a reason *x* has’ fails when (B)-type reasons are in question but not when (F)-type reasons are. (This can be the key to understanding why one can wind up with the *prima facie* paradoxical ‘There was no reason to do what he did, even though he did it for a reason’. In my terms, this means that there was no (F)-type reason to do what he did even though his having a (B)-type reason explains his doing it.⁷)

Sometimes philosophers talk about ‘normative’ reasons and contrast them with ‘motivating’ reasons in order to separate the (F)-type from the (B)-type. But the label ‘normative’ cannot be an apt one if it is supposed to put reasons of the two types into different categories. For normative questions, about the justification of

⁷ Compare Dancy (2000: 3), whose own diagnosis is different. Dancy thinks that the key is to recognize that “Some motivating reasons are not good reasons.” He says “It is only when we have our eyes on the question . . . whether there were any good reasons for so acting . . . that we want to allow that a motivating reason can be no reason at all.” We can grant to Dancy that “There was no reason for him to do it” can be correct when he lacked any remotely good reason to do it (cf. next note). But Dancy’s diagnosis does not speak to the case in which (a) he did it because he believed that *p*, (b) it was false that *p*, (c) it is not in question that there would have been a perfectly good reason for him to do the thing if it had been that *p*.

action, don't lapse when one adverts to what the agent believes in saying what reason they had. When Ann saw Sam as having the reason which he had by dint of his having a false belief about her whereabouts, she treated him as someone who acts for reasons. Indeed the normativity which (B)-type reasons share with (F)-type ones shows why we do not think of the bridge as having had a reason to collapse: the behaviour of bridges is not subject to the sort of assessment to which people's conduct is.

When (B)-type reasons are seen to be normative in the manner of (F)-type reasons, it seems that they inherit the normativity of (F)-type reasons. We might say that if X has a reason to Φ in believing that p , then—at least from X 's perspective—there would be an (F)-type reason for X to Φ if p were true.⁸ So, for example, Sam believed that Ann was in a café on Beech Street, and he would have had an actual reason to go to Beech Street if that had been so. Here Sam's having a *reason* in believing what he does is explained by allusion to (F)-type reasons: Sam's having a reason is understood by reference to the idea that a fact can favour a way of acting.

1.3 Connecting (F)'s instances with (B)'s

In order to see how the two types of reasons for acting are related, we need to think about acting for reasons. We need to look into the two different accounts we saw of why someone acts—at two types of *reason-explanations*, that is. These can be represented with schemas.

(F.Exp) $X \Phi$ -d because p (where 'because' can be glossed with 'for the reason that').

(B.Exp) $X \Phi$ -d because they believed that p (where X had a reason to Φ : they believed that p).

The qualifications here ('where . . .') rule out cases where the 'because', of 'because p ' and 'because they believed that p ', is not the 'because' of a reason-explanation.⁹

The obvious way to make a connection between these two would be to say that someone does something because p , if *both* they do it because they believe that p *and* p is actually a reason for them to do it. Thus:

(P) $X \Phi$ -d because (for the reason that) p if

(i) $X \Phi$ -d because they believed that p , and

(ii) A reason for X to Φ was that p .

⁸ 'At least from X 's perspective' has plenty of work to do here. An agent's having some belief may lead them to act in some way, which for all that the belief is true, is actually not a reason—that is, now not a remotely *good* reason—for so acting. The agent then treats a consideration as favouring Φ -ing, although it is not genuinely a reason, but can be seen as favouring Φ -ing only by adopting their perspective on what favours what. A complete account of reasons for acting (which I don't attempt here) would evidently need to speak to these cases.

⁹ At this point I have to assume that we know what a reason-explanation is. Notice that even with the qualifications written in, cases in which a subject's believing something is itself a reason

To see how (P) would work, think about the skater Edna who kept to the edge of the pond because the ice in the middle was thin. According to (P), Edna kept to the edge of the pond because she believed that the ice in the middle was thin, and the fact that the ice there was thin was a reason for her to keep to the edge.

But (P) cannot be right. What conjunct (i) tells us about the agent—that she believed that p —is independent of whether p is true. (Edna could have believed that the ice in the middle of the pond was thin, even if it wasn't.) Conjunct (ii) tells us something about a reason for the agent to Φ which might obtain without registering with the agent. (There might be a reason for Edna to keep to the edge of the pond, even if Edna was not aware of it.) But if p was actually someone's reason for Φ -ing, then the fact that p was something that registered with them. (If it was because of the thin ice in the middle that Edna stayed at the edge of the pond, then the fact that the ice in the middle was thin played a part in keeping her there.) So we can agree that Edna's believing that the ice was thin ensured that she *had* a reason to keep to the edge (a (B)-type one). And we can agree that the thinness of the ice ensures that there *was* a reason for Edna to keep to the edge (an (F)-type one). Still, Edna's believing what she did cannot as such provide her with the actual reason that there was for her to keep to the edge.

Here now is a counterexample to (P), which will serve not only to bring out what is wrong with it, but also to suggest a remedy. The example concerns Edmund who believes that the ice in the middle of the pond is dangerously thin, having been told so by a normally reliable friend, and who accordingly keeps to the edge. But Edmund's friend didn't want Edmund to skate in the middle of the pond (never mind why), so that he had told Edmund that the ice there was thin despite having no view about whether or not it actually was thin. Edmund, then, did not keep to the edge *because* the ice in the middle was thin. Suppose now that, as it happened, the ice in the middle of the pond was thin. This makes no difference. Edmund still didn't keep to the edge *because* the ice was thin. The fact that the ice was thin does not explain Edmund's acting, even though Edmund did believe that it was thin, and even though the fact that it was thin actually was a reason for him to stay at the edge.

Edmund is a familiar sort of character in epistemology. Such characters are usually used to show that justified true belief is not sufficient for knowledge. Edmund here is used to show that someone's having a true belief (even a justified one) which explains their acting is not sufficient for them to have acted for an (F)-type reason. We saw that an (F)-type reason has to have registered with an agent if they are to have acted on it. We see now what this amounts to. A condition of Φ -ing for the reason that p , when one believes that p , is that one *knows* that p . When people act for (F)-type reasons, they do so in virtue of knowing the relevant facts.¹⁰

(see above, where footnote 6 was flagged) are not excluded: in those cases ' X believes that—' substitutes for ' p ' in (F) and in (B).

¹⁰ For plenty of examples showing that one needs to know (relevant) things to do (related) things intentionally, see Gibbons (2001).

1.4 A disjunctive conception of acting for reasons

We wanted to understand how the ethicists' ((F)-type) reasons and the philosophers of mind's ((B)-type) reasons are related. A connection between them is easy to make now that knowledge is on the scene. It can be recorded by giving an account of acting for reasons in which acting from knowledge belongs under the head of acting from belief—an account which might be formulated as follows:

- (DisA) If $X \Phi$ -d because X believed that p , then
 EITHER $X \Phi$ -d because X knew that p , so that $X \Phi$ -d because p .
 OR $X \Phi$ -d because X merely believed that p .

Note that all the instances of 'because' in (DisA) are to be so understood that (DisA) is concerned exclusively with reason-explanations. (Compare the qualifications in (F.Exp) and (B.Exp) above. For the sake of simplicity, I don't write the qualifications in explicitly here.)

'Merely believed that p ' here has to mean 'believed that p without knowing that p '. And this could make it seem that (DisA) relies on little more than that one either knows or doesn't know the things that one believes. Certainly (DisA), which has 'if' and not 'iff', is no sort of analysis.¹¹ What makes (DisA) less than trivial, despite an evident circularity, is the conjunct that belongs in the first disjunct—that $X \Phi$ -d because p . Of course it's a fact that p if X knows that p ; but this conjunct assures us of more than this: it puts it on record that where there is a reason-explanation from X 's knowledge, the fact that p was a reason X had for Φ -ing.

Thus (DisA) connects reasons of the two types. And it can serve to remind us that belief-desire theories lack the resources needed to say what it is to act for an (F)-type reason. If we are to say what it is to act for an (F)-type reason, we need the idea that knowledge is a spring of action.

2

I want to encourage an understanding of (DisA) on which it can complement John McDowell's disjunctive conception of appearances. I am going to assume that a disjunctivism of McDowell's sort is correct and can play the part McDowell gives it. My idea is to build upon it, and perhaps, indirectly, to lend it support. If (DisA) does good clarificatory work when understood in parallel with McDowell's disjunctive conception, and if an attractive picture results when both disjunctive conceptions are in play, then a case of sorts may be made for accepting both.

¹¹ To justify replacing the 'if' with 'iff' I should need at least to justify an assumption (which I take it most belief-desire theorists work with) that propositional knowledge requires belief. There are reasons to be doubtful about this. And it suits me that there is no need to settle whether the assumption is correct: one effect of my not going for 'iff' is that it can be very clear that there is no pretence at analysis. Something that I hope will emerge from section 2 is that a disjunctivist claim can play a dialectical role without having analytic pretensions.

I shall present McDowell's disjunctive conception as a counterpart to (DisA), so that we shall have (DisP), with 'P' for perception, where 'A' was for action. That puts me in a position to show that (DisA) can make a point analogous to McDowell's (section 2.2). It can then be seen that the two disjunctive conceptions have work to do together (section 2.3), and an underlying structure can be uncovered (section 2.4).

2.1 The disjunctive conception of appearances as an analogue of (DisA)

McDowell wants us to appreciate that someone to whom it appears that things are thus and so may actually be someone who perceives how things are. We appreciate this when we say that cases in which it appears to a subject that things are thus and so come in two kinds—cases of genuine perceptual experience, and cases of mere appearance. Cases of the two kinds might be indistinguishable from the point of view of the experiencing subject; but this should not interfere with recognizing their difference—a difference which becomes evident when it is allowed that it may be manifest to a subject in experience that things are thus and so.

Thus McDowell's disjunctive conception simultaneously brings together and distinguishes the experience of a perceiver and the experience of someone who has a mere experience. What the disjunctive conception of acting for reasons does is simultaneously to bring together and to distinguish acting from an actual reason and acting from mere belief. Cases in which an agent acts because they believe that *p* come in two kinds. Cases of the two kinds might be indistinguishable from the point of view of the agent or of an onlooker who is ignorant of the relevant facts; but this should not interfere with recognizing their difference—a difference which becomes evident when it is allowed that an agent may act from knowledge.

If McDowell's idea is formulated in parallel with (DisA), then the notion of knowledge will enter explicitly into its statement. So we have:¹²

- (DisP) It appears to *X* as if things are thus and so if
 EITHER it is manifest to *X* that things are thus and so, so that *X* is well placed
 to know how things objectively are.¹³
 OR it merely appears to *X* as if things are thus and so, so that *X* is not in such
 a position.

This formulation highlights the fact that experiences in which things seem a certain way to a subject are not all on a par epistemologically speaking. Explicitness about this helps to make it clear that a subject's inability to distinguish between experiences of two sorts cannot be adduced in an argument purporting to show that their epistemic situation is the same whichever sort their experience is of. When an

¹² As with (DisA), we have 'if' and not 'iff'. Some might say that this means that Disjunctivism proper is then not really at issue. But (DisP) may be worth considering if it can assist in conveying the point of adopting a disjunctive conception of appearances.

¹³ The qualified '*well placed* to know' allows for cases in which a subject (perhaps because they wrongly suppose that they are under some illusion) fails to take appearances at face value—a subject who, arguably, then fails to perceive how things are.

experience is recorded by saying that things are manifestly thus and so to its subject, it can be clear that, whatever doubts a sceptic might induce in them, the subject, if they judge that things are actually thus and so, does not take a step beyond what their experience entitles them to judge. This at least can be clear if *X*'s being well placed to know how things objectively are is constituted by *X*'s having an experience of the sort they do in satisfying the first disjunct.

2.2 (DisA) as parallel to the disjunctive conception of appearances

The point so far, of (DisA), has been to show that, unless it is to be ruled out that people act for (F)-type reasons, it has to be allowed that knowledge sometimes explains people's acting. Some belief-desire theorists may think this of little moment. 'It can be allowed that a person's knowing something might come into an explanation of their doing what they did', these theorists may say. 'And when questions about an agent's (F)-type reasons are introduced, we need to ask whether the agent knows the things they believe. But this is a separate matter—separate from the phenomenon which interests us belief-desire theorists.'

Knowledge cannot be so easily sidelined, however. Suppose that actually the reason for which a person Φ s is that *p*. (Suppose, for instance, that in fact Sam is headed for Beech Street because [for the reason that] Ann is there.) The belief-desire theorists say that the person Φ s because they believe that *p*; and their explanation can be well and good. (Indeed Sam is headed for Beech Street because he believes that Ann is there.) But the person's acting as they do now depends upon their knowledge. (Something that Sam is doing is *heading for Beech Street because [for the reason that] Ann is there*, and unless he knew that she was there, he would not be doing this.) The person's knowing that *p* is hardly a separate matter, then.

It need not be at all puzzling now that (F)-type and (B)-type reasons should play their roles in concert. I suggested that there could seem to be a puzzle posed by asking: 'Are our actions somehow overdetermined when we act for an (F)-type reason and a (B)-type reason?' Well, the person who acts for a (B)-type reason is moved by believing things to be a certain way. If they act also for an (F)-type reason, then they are satisfactorily related to how things actually are, and the fact that things actually are the way they believe them to be weighs with them. Someone who Φ s because they believe that *p* may actually be someone who Φ s because *p*. Compare the claim which I used to introduce (DisP): someone to whom it appears that *p* may actually be someone who perceives that *p*.

2.3 The two disjunctive conceptions brought together

I turn now to an example in which perceptual experience leads to action. This is with a view to showing how the two disjunctive conceptions—of experience, and of acting for reasons—work together. The example consists of a pair of cases.¹⁴

¹⁴ The cases illustrate what Scott Sturgeon calls 'delusive Rational efficacy' (2000: ch. 1). I treat this phenomenon differently from Sturgeon, and speak to our disagreement in footnote 20 below.

Maja has seen a notice about a child who has lost his pet rabbit in the vicinity and wants it returned. She sees a rabbit and acquires the belief that there is a rabbit at a certain place, so that she has a reason to run in a certain direction. In the other case, Maja again is motivated to catch a rabbit if she sees one, and again she arrives at the belief that there is a rabbit at a certain place. But in this case, Maja is hallucinating. Again she runs towards the place, but now it merely appears to Maja as if there were a rabbit there. There is an explanation of Maja's running as she does which is common to the two cases. Thus, in both cases, Maja comes to believe that there is a rabbit at a certain place because it appears to her as if there were a rabbit there; and she runs in the direction she does because she believes that there is a rabbit there.

No one need quarrel with this common explanation, which makes use, as explanantia, of instances of what occurs on the left-hand-sides of (DisP) and of (DisA). But the disjunctive conception of acting for reasons directs us towards explanations which introduce an agent's (F)-type reasons. In the case in which Maja sees a rabbit, there is another explanation of her running in the direction she does: she runs because *a rabbit is present*. And when we take perceiving-Maja to be acting for this reason, we recognize that there is also another explanation of her acquiring her belief. This now is an explanation towards which a disjunctive conception of appearances might direct us: Maja believes that a rabbit is present because *she sees the rabbit there*.

(DisP) and (DisA) do work in tandem here. Given (DisP), beliefs which are got through experience are sometimes shaped by the facts. Given (DisA), actions which spring from beliefs are sometimes shaped by the facts. When Maja sees a rabbit and chases it, the fact that the rabbit is present plays a role throughout the process: Maja runs because there is a rabbit in her sights.

2.4 The common structure

Let me introduce some terminology, in order to set out the parallels. Say that a case shows up as 'good' when (an instance of) the *first* disjunct of (DisA) or (DisP) is used in its description.¹⁵ And say that a case is described 'neutrally' when (an instance of) the left-hand side of (DisA) or (DisP) is used. So '*X* Φ -d because *X* believed that *p*' and '*It appears to X* as if things are thus and so' give neutral descriptions of cases.¹⁶

We have just seen it to be characteristic of good cases that a distinctive sort of reason-explanation is available, one which shows the subject's or agent's connection with the facts. In the good case, Maja was aware of the fact that a rabbit was present,

¹⁵ Just as the goodness of a case of *X*'s Φ -ing-because-*X*-believes-that-*p* does not consist merely in its being true that *p*, so the goodness of a case of experience does not consist simply in its veridicality. One could introduce some tricks, and construct an example in which Maja has a veridical hallucination as of a rabbit. Evidently the mere presence of a rabbit does not make for a good case.

¹⁶ (DisP) is formulated so as to allow for hallucination. To treat illusions, we should need to introduce different sorts and grades of neutrality, as it were. For instance: someone who says 'The book appeared to *X* to be green' gives a neutral description in one respect, even while showing *X*'s case to be good so far as his perceiving the book is concerned. Obviously a great deal more would need to be said in a worked-out account.

and her acting was then a matter of her responding to this fact. When a case is described neutrally, its goodness, supposing it actually to be a good one, is not apparent. Perceiving-Maja was suitably related to the facts, but is not seen to have been so related when her case is described neutrally. One might say that neutral descriptions of good cases suppress some of the truth.

We saw the limitations of explanations of action which are cast in neutral terms. I said that Edna's believing what she did could not *as such* provide her with the reason that there actually was for her to keep to the edge of the pond. The point is that when Edna is said to have acted because she believed something, it is left open whether she knew it: there is something else one has to know about her in order to know whether she acted for an (F)-type reason. Still, it may be true that Edna has acted for an (F)-type reason, even when this is left unsaid. If an agent has *p* as a reason for acting, then (in such circumstances, and whatever is said) nothing more than what *de facto* suffices for their believing that *p* may be needed for them to act for the reason that *p*. The parallel holds for the case of experience. When it is said only that things appeared some way to Maja, it is left open whether she perceived things to be that way: there is something else one needs to know about her in order to know whether her experience is such as to provide her with a reason for acting. Still, it may be true that Maja has seen a rabbit, even when this is left unsaid. If it appears to a subject that things are a certain way *and* they perceive them to be that way, then (in such circumstances, and whatever is said) nothing more than what *de facto* suffices for their having the experience they do may be needed for them to perceive things to be that way.¹⁷

In philosophy of perception, experience is sometimes thought about in its role of accounting for beliefs. When it comes to philosophy of action, experience must be allowed a role in providing reasons for acting. The factivity of experiences can then come into view, and there is a new point to conceiving of the good cases as a disjunctivist does. From the perspective of an account of agency, veridical perceptions are seen as fitted to explanations which have no use for neutral terms—to explanations of people's acting for reasons that there actually are for them to act. Meanwhile neutral descriptions of experience are seen as relatively lacking in explanatory potential. (The fact that one is prevented from explaining very much when only neutral descriptions are available showed up in our needing to confine ourselves, when we stuck to the neutral, to descriptions of perceiving-Maja's action that can strike us as artificial. For example, we had Maja running 'in a certain direction', where it would be natural to say that she ran *towards the rabbit*. And of course in a good case, someone in a position to demonstrate the rabbit can give an explanation of why she ran towards *that* rabbit [a demonstrated one].)¹⁸

¹⁷ I say '*may be needed*' because it is arguable that someone doesn't perceive that *p* unless she takes it at face value that it appears as if *p*: cf. footnote 13 above.

¹⁸ Martin points out that, by making appeal to (to put the matter in my terms) features proper to good cases which perceptions alone can explain, 'one can . . . rebut the challenge that the disjunctivist's conception of sensory experience is guaranteed to be explanatorily redundant' (2004: 63–4). The features by reference to which Martin makes the point are the particulars in the subject's environment and their properties. I take the point to have added force when the features in question are the facts. My hope then is that (DisA) and (DisP) combine to ensure that the

Someone might wonder how explanations in which experiences are neutrally described can be supposed to work. The explanation of hallucinating-Maja's running as she does can also be given in the case of perceiving-Maja, in whose case there are other, more illuminating, explanations. But if the explanations that are proprietary to good cases tell us more than explanations cast in neutral terms, then how can statements cast in neutral terms play the explanatory role that they do across the board?

Well, if we think that it is manifest to perceiving-Maja that a rabbit is present, then we may think of hallucinating-Maja as under the illusion that it is manifest that a rabbit is present. (Not that a mere appearance always puts someone under an illusion. 'I'm under no illusions' might be the words of someone who takes herself to be hallucinating, and who wants it to be known that she discounts her experience. But the hallucinating-Maja of our example, who comes to have a reason for acting, does not take herself to be hallucinating.) Now, to someone who is under an illusion of its being manifest that *p*, there seems to be a reason for believing that *p*.¹⁹ So hallucinating-Maja is understood when it is known both that it seems to her that there is a reason for believing a rabbit to be present, and that in acquiring the belief she gains a (B)-type reason upon which she acts. Perceiving-Maja, in so far as she is described in neutral terms, can then be exactly like hallucinating-Maja and understood in the same way. By virtue of her experience, there seems to her too to be a reason for a belief which she acquires and acts upon. There is no problem, then, about allowing that the explanatory work done when neutral descriptions of experience are given can be done both in cases which are not good ones and in cases which are.²⁰ This now is comparable to allowing that '*X* Φ -d because *X* believed that *p*' can be true both when '*X* Φ -d because *p*' isn't true and when it is.

inlectably relational character of perceptual experience is writ large: a person's involvement with the facts about how things outside them are may no longer seem to be a matter simply of how the person might be described.

There is much more that can be said about explanation in this area, as (Martin 2004) and subsequent literature show. Relevant are examples, which Williamson has given, which show that an agent's *knowing* something may belong in an account of their acting which is in one good sense the most explanatory powerful account (2000: 62). (I say 'in one good sense' because I think that we can have different purposes in explaining someone's acting, so that explanatoriness might not be measured on a single dimension.)

¹⁹ The 'seem' here is not the 'appear' I have used in descriptions of sensory states. Its appearing to someone that *p* might not result in their seeming to have a reason to believe that *p*; whether or not it does so result depends upon (very roughly) their mindset—upon whether, for instance, they take themselves to be under an illusion.

²⁰ Compare Martin on inherited, or dependent, explanatory potential (2004: 70), and his use of this idea in answering the question why properties which make for goodness aren't screened off.

In Sturgeon's book (2000), Disjunctivism fails as a theory of perception because it fails to provide an explanation of what Sturgeon calls Scene immediacy. Sturgeon, then, will have an objection at this point. For here the Scene immediacy of hallucinations is adduced in accounting for their impact on belief; the delusive Rational efficacy of hallucinations (to continue with Sturgeon's terms) is explained only indirectly; and there is no account of Scene immediacy itself. Well, I suggest that a disjunctive conception of experience can do without the explanatory ambitions of the theory Sturgeon seeks. (Perhaps we understand Scene immediacy well enough as soon as we know what (DisP) is meant to tell us.) But of course there is more to be said, some of which is said in Sturgeon's contribution to the present volume.

I want it to be clear that there is no fault to be found with either of the two conceptions of reasons for acting from which we began. So I return here to the agenda of section 1. I said there that a story about acting for reasons ought to have room for reasons both of the (F)-type and of the (B)-type. And I think that we are now in a position to see that such a story really must accommodate them both. I suggested that the normativity of (B)-type reasons is inherited from the normativity of (F)-type reasons (section 1.2). So too, I shall now suggest, is the *cogency* of (B)-type reasons. It is not merely that reasons for acting could not all be of the (B)-type (cf. section 1.1), but also that our idea of acting for reasons could not be got simply from instances of (B.Exp).

3.1 Explanations from belief

I just gave an account of why there should be an explanation common to the cases of hallucinating- and perceiving-Maja. The key was to find a way to understand Maja's acquiring a belief which is indifferent to the question of whether things are indeed as they appear to her to be. We thought of Maja as seeming to have a reason for the relevant belief. In thinking of her thus whether she hallucinates or perceives, we rely upon our grasp of how things appear to be when they appear to someone to be the way they are. So the understanding we have of a hallucinating subject is rested in our understanding of a perceiving subject. Now when it comes to acting for reasons, we start out equipped with a way of understanding an agent which is indifferent to questions about how things actually are. In finding out what someone believes, we already have a way of rendering their acting explicable which is indifferent to such questions. (Ignorant onlookers, such as the one we encountered in section 1.2, make this clear.) But here again we should allow a sort of derivativeness of understanding.

Often our aim in discovering why a person acted as they did is to reveal the favourable light in which they saw what they did. By adopting the perspective of the agent, so as to come to know what they treated as a reason, we can answer questions about their motivation. We may then arrive at an explanation of their acting which says what (B)-type reason they had. And then the question what (F)-type reasons were in play can lapse. (Compare: if we want only to know how it appeared to a subject as if things were, we have no need to be troubled with how things actually are in their environment.) Still, when one explains a person's acting in some way by saying what they believe, one relies upon a grasp of what is actually a reason for so acting, or of what might be treated as such.²¹ (Compare: when one is informed of how it appears to a subject as if things are, one relies upon a grasp of how things appear when they appear the way they are.) Thus a distinctive sort of explanatory interest is in play when (B)-type reasons are seen to be at work; but the understanding achieved by knowing

²¹ Footnote 8 above touches on the idea of 'treating' something as a reason.

someone's (B)-type reasons is rested in an understanding of the operation of (F)-type reasons.

The point here might be put by saying that (B)-type reasons inherit the explanatory character of (F)-type reasons. The point of the disjunctive conception of acting for reasons, meanwhile, is that acting for (F)-type reasons can be brought under the head of acting for (B)-type reasons. It then seems that reason-explanations on the patterns of (F.Exp) [*'because p'*] and of (B.Exp) [*'because x believed that p'*] must stand or fall together.

3.2 Acting for reasons

Many philosophers, however, have held that we should endorse (B)-type reasons at the expense of the (F)-type, or that we should endorse (F)-type reasons at the expense of the (B)-type. I shall conclude by addressing in turn those who are dismissive of (F)-type reasons, and those who are dismissive of (B)-type reasons.

Belief-desire theorists who take *belief* to exhaust the cognitive states of mind needed in an account of acting for reasons have no place for instances of (F.Exp) in their account. (They may say that they are concerned with 'motivating' reasons, not with 'normative' ones (see section 1.2). And they may think that knowledge can be sidelined (see section 2.2).) Well, when acting for reasons is conceived disjunctively, (F)-type reasons are given a place in an account. And the fact that (B)-type reasons inherit their cogency from (F)-type reasons shows that (F)-type reasons can hardly be left out of account. To leave them out would be to forget that the understanding achieved by knowing someone's (B)-type reasons is rested in an understanding of the operation of (F)-type reasons.

Most of these belief-desire theorists, who think that an account of acting for reasons can make do with *belief* alone on the cognitive side, treat instances of (B.Exp) as claims about what causes what which belong in a naturalistic account of the world's causal workings. But if instances of (B.Exp) are found illuminating only in so far as instances of (F.Exp) can be appreciated, then it is very far from obvious that instances of (B.Exp) are so much as concerned with the causal workings of the world at large.²² It may be then that the dismissal of (F.Exp) on the part of these theorists leads them to an erroneously naturalistic view of human agency. Or again, if these theorists adopt a naturalistic outlook from the outset, then they are predisposed to keep out of sight the idea which (DisA) brings into view—the idea of *reasons* being at work. Their adherence to a kind of naturalism encumbers them with a conception of people's reasons for acting which loses any connection with the normative contexts in which the concept of a reason belongs.

So much for those who would hope to make do without (F)-type reasons in an account of human agency. Recently there has been much suspicion of (B)-type reasons. Some say that the whole idea of a (B)-type reason is misguided.²³ Some say that

²² 'Very far from obvious' puts it very mildly in my own view. For arguments against thinking that reason-explanations draw on naturalistic causal truths, see Hornsby (1993).

²³ For example, Stoutland (2007).

a proper account of agency should make no appeal to the causal role of psychological states.²⁴ And some say that although it might appear that we explain people's acting when we say what they believed, if we then give reason-explanations, these cannot be genuine explanations of such a sort as causal explanations are.²⁵ All of these philosophers are hostile to the conception of reasons for acting that holds sway in philosophy of mind. It could be that some of their hostility is misplaced hostility to the kind of naturalism I have just mentioned. But whatever the source of their wishing to make do with a conception of acting for reasons which leaves out instances of (B.Exp),²⁶ these philosophers face a problem. They fully appreciate the need to accommodate (F)-type reasons in an account of agency. And they think that in the face of (F)-type reasons, (B)-type reasons lose out. But given a conception of acting for reasons in which an understanding of (B.Exp) is rested in an understanding of (F.Exp), and in which (B.Exp) subsumes (F.Exp), (B)-type reasons could hardly lose out to (F)-type ones.

If the argument here is right, then both conceptions of reasons must be admitted. And the naturalistic view of human agency which belief-desire theorists are prone to introduce must not be allowed to stand in the way of endorsing them both. With *knowledge* in the picture, we see how there can be a story of the kind we sought in attempting to combine the ethicists' and philosophers of mind's conceptions—a story in which the normative and the psychological come together.

REFERENCES

- Dancy, J. (2000) *Practical Reality*, Oxford: Oxford University Press.
 Davidson, D. (1963) 'Actions, Reasons and Causes', *Journal of Philosophy*, 60: 685–700. [Reprinted as Essay 1 in his *Essays on Actions and Events*, Oxford: Oxford University Press, 1980.]
 Gibbons, J. (2001) 'Knowledge in Action', *Philosophy and Phenomenological Research*, 62: 579–600.
 Hornsby, J. (1993) 'Agency and Causal Explanation', in J. Heil and A. Mele (eds.) *Mental Causation*, Oxford: Oxford University Press. [Reprinted as Essay 8 in her *Simple Mindedness: Essays in Defence of Naive Naturalism in the Philosophy of Mind*, Cambridge, MA: Harvard University Press, 1997.]
 Hyman, J. (1999) 'How Knowledge Works', *The Philosophical Quarterly*, 49: 433–51.

²⁴ For example, Stout (2004).

²⁵ For example, Dancy (2000).

²⁶ Some of the argument of section 1 above was aimed against denying, what Stoutland and Stout deny, that one can say what reason someone had in saying what they believe.

I surmise that Dancy thinks that anyone who claims that reason-explanation is causal means to make appeal to a feature of reasons' cogency that is not already in evidence when it is known that someone's having a reason to do something has led them to do it. Thus Dancy seems to me to buy into an assumption of (though not the view of) those belief-desire theorists who take the idea that reason-explanation is causal to be part and parcel of treating instances of (B.Exp) as belonging within a naturalistic account of the world's causal workings.

- McDowell, J. (1982) 'Criteria, Defeasibility and Knowledge', *Proceedings of the British Academy*, 68: 455–79. [Reprinted as Essay 17 in his *Meaning, Knowledge and Reality*, Cambridge, MA: Harvard University Press, 1998.]
- Martin, M. G. F. (2004) 'The Limits of Self-Awareness', *Philosophical Studies*, 120: 37–89.
- Parfit, D. (1997) 'Reasons and Motivation', *Proceedings of the Aristotelian Society*, Supplementary Volume, 71: 99–129.
- Stout, R. (2004) 'Internalizing Practical Reason', *Proceedings of the Aristotelian Society*, 104: 229–43.
- Stoutland, F. (2007) 'Reasons for Action and Psychological States', in A. Leist (ed.) *Action in Context*, Berlin and New York: De Gruyter/Mouton.
- Sturgeon, S. (2000) *Matters of Mind*, London and New York: Routledge.
- Williamson, T. (2000) *Knowledge and its Limits*, Oxford: Oxford University Press.

11

On How to Act—Disjunctively

Jonathan Dancy

1

In this essay I revisit some issues which I addressed briefly in my *Practical Reality* (2000). I argued there against disjunctive accounts of acting for a reason. I still think that that is the position I want to take, but it is worth looking again at the issue, in the light of increasingly sophisticated understandings of the nature of disjunctive accounts elsewhere.

The issue is only interesting if we intend any disjunctive account that we come up with in the theory of action to be relevantly similar to disjunctive accounts of experience. It is not very hard to say *something* disjunctive about action. Of course not everyone agrees on the form that a disjunctive account of experience should take; other essays in this volume attest to that. But there seems to be something of a consensus on the following points, which for that reason I will here lay out very snappily.

First, perceptual and non-perceptual experiences are not to share any common element.

Second, perceptual and non-perceptual experiences have no common features or qualities. This point is worth distinguishing from the preceding point for reasons that will emerge below.

Third, non-perceptual experiences are to be characterized entirely as episodes that the experiencer cannot distinguish from perceptual experiences. This negative characterization is all that can be said in general about the non-perceptual, in explication of its nature as *experience*. Non-perceptual experiences not only have no qualities or features in common with perceptions; they have no common features or qualities among themselves either.

Fourth, for this reason, and in this sense, the perceptual case enjoys a certain primacy in our understanding of experience. There are, in the broadest terms, two sorts of experience, and the second is to be understood in terms of the first (that is, as indistinguishable from it).

Fifth, perception reaches out to and essentially involves the perceived object. Perceptual experience is understood relationally. Non-perceptual experience is not understood relationally (because there may be nothing there) though it is understood in terms of something that is relational.

Sixth, and for this reason, a disjunctive account of experience would be empirically false if there were some non-perceptual experiences that are distinguishable (by their owner, and in the relevant way) from perceptions.

There are some friendly rivals to the disjunctive account, which I will consider in a moment. Before that, I want to make a couple of comments on the points above. First, it seemed to be worth distinguishing potential common elements from potential common features or qualities. Talk of elements leads one to think in terms of building blocks. If two items share a common element at least one of them must be metaphysically complex, built out of parts. Two items might have common qualities without being metaphysically complex in that kind of way. If we conceive of a non-perceptual experience in such a way that *the same sort of thing* is present in a perceptual experience, attended of course by the presence of the relevant object (probably acting as a cause), we are conceiving of the perceptual experience as metaphysically complex, and, what is more, as a complex of just the type that the disjunctivist is in the business of trying to get away from. The disjunctivist allows that a perceptual experience is a complex, since it consists partly, but not wholly, of the relevant object, the object perceived. For disjunctivists, then, there is a right way of being complex and a wrong way. The right way is an indissoluble way; whatever is involved in the perceptual experience, in addition to the presence of the object perceived, *could* not be present alone. For if one takes the object away, there is not enough left to stand independently.

Talk of common qualities, by contrast, does not lead one to think in terms of building blocks. Two experiences could both be representational, for instance, without this tending to show that they have some common element. (Just as a picture on the wall and a dream can both be representational without having some common element—though they may for all that have common qualities.) I mention this because it seems relevant to a suggestion that I made some while ago (Dancy 1995), one that has not found favour with leading disjunctivists. This was that we may perhaps start from a workable disjunctive conception of experience of the favoured sort, like this:

A subject S is enjoying an oasis-experience iff either

- (1) an oasis is manifesting itself to S, or
- (2) it is for S just as if (1).

In this version we have, as yet, only characterized the second disjunct as like what it is not. But, I surmised, “there may be available a more explicit characterization of the second disjunct, and in a totally explicit version . . . it would indeed be characterized in that better way” (1995: 436). Tim Crane (2005: section 3.4) responds to this as follows:

... part of the point of disjunctivism is to characterize what genuine perception is, and to show how this can be done without assuming a substantial conception of the ‘common

factor' between perception and hallucination: that is, a conception of a common state which can play a significant psychological or explanatory role. But if Dancy's challenge were answered, by giving a positive account of the hallucinatory experience (for example, in terms of qualia), then this would amount to giving a substantial account of the common factor between perception and hallucination. And this is what the disjunctivist says cannot be done...

It seems to me, however, that Crane here supposes, unnecessarily, that any positive characterization of hallucination, any account of its positive qualities, as it were, would be bound to specify a 'common state'—a sort of building block common to perception and hallucination. That thought is surely a mistake, for two reasons. First, we don't need to think of a quality as a part or state. Second, even if we did identify, in the case of hallucination, something capable of functioning as a part, or state, we would not be thereby forced to allow that the perceptual state itself also had that state as part of it. We could explain the indistinguishability in terms of a common quality that is not a common part.

If this were all that there is to be said, I would be suggesting that the need to deny any common element cannot be presented as a reason not to look for a positive characterization of non-perceptual experience. But Mike Martin has more to say. In his 'The Limits of Self-Awareness' (2004: 46) he does, admittedly, seem to echo Crane's point:

If what marks [the non-perceptual] cases out in the first place is just that they involve the absence of perception, then one may worry that whatever fixes what they have in common with each other will apply equally to any case of perception. That is to say, the further specification of hallucination will be something which is present not only in all cases of illusion or hallucination but also in the case of perception. The disjunctivist will then be left in the unhappy position of conceding that there is a common element to all of the cases. Now if the common element is sufficient to explain all the relevant phenomena in the various cases of illusion and hallucination, one may also worry that it must be sufficient in the case of perception as well. In that case, disjunctivism is threatened with viewing its favoured conception of perception as explanatorily redundant.

Here again, I would say that common qualities need not generate a common element. But Martin had a different, and, I think, a stronger point. He contrasts two accounts of the non-perceptual experience, modest and immodest. The immodest account offers a list of characteristics common to all non-perceptual experiences. The modest account characterizes such experiences as any episode relevantly indistinguishable from a perception. Already we see an advantage in the modest account. For the immodest account can hardly dissent from the modest one. We all have to start from that—from the notion of whatever is, or is indiscriminable from, a perception. The immodest account goes beyond that, and in going beyond it, courts trouble. For even if we suppose that the immodest account succeeds in specifying features sufficient to make an episode an experience, it is vulnerable to the ever-present possibility of there being some further type of episode that is *also* indistinguishable from a perception. The immodest account may, that is, succeed in offering sufficient conditions for an episode to be an experience, but is unlikely to be able to offer necessary ones. The

modest account, by contrast, offers only one condition, and it is both sufficient and necessary.

It will still be the case, on this showing, that if we could conceive of an experience that was not relevantly indistinguishable from a perception, the disjunctive conception would turn out to be inadequate—empirically inadequate, as it were. Are we convinced, then, that this is impossible? A counter-example would have to be an episode that was an experience but which we can still distinguish from a perception. One is tempted to write ‘can distinguish on internal grounds’. Certain ‘external’ ways of distinguishing are ruled out. Asking one’s friends, for example, will not count. Nor will appealing to empirical evidence that experiences of this sort are often, or generally, or even always hallucinatory. More to the point would be Austin’s claim that there are lots of differences between dream-experiences and waking experiences. Presumably what prevents one from appealing to these is that one cannot do so when asleep. It seems to me, however, that there is room for a sort of experience that one knows on ‘internal’ grounds alone to be not a perception. I once suffered briefly from otitis, an inflammation of the inner ear, which has various consequences. One of them was that as I lay in the bath, gazing out at my toes and the taps at the other end, it suddenly began to appear that everything was moving to the left around my point of view as a fixed point, but also that my point of view was moving with it. It now seems to me that I was dealing then not with two experiences, each, as it were, incompatible with the other, but with one experience with contradictory content. Suppose that one allows this as a description of the relevant phenomenon. The question then is whether this is a potential counter-example to disjunctivism. Is an experience with contradictory content one that is distinguishable on ‘internal’ grounds from a perception? It does not seem to me that we can say ‘no’ in answer to the question simply on the grounds that it is, somehow, independently established that the external world cannot contain contradictions.

I certainly don’t think of this as anything approaching a decisive objection to the disjunctivist account of experience. I mention it here merely to reinforce the point, which will be important later, that accounts of this sort are liable to empirical refutation if we are not careful. This is the first point that I want to take forward from the brief survey in this section. The second is that we will be considering later what form a disjunctive conception of action might take, and at that point we will need to ask ourselves whether the reasons given by Crane and Martin for sticking to the pure, unexpanded form of disjunctive account in the theory of experience apply with equal force in the theory of action. In fact my feeling is that it will not be so hard there to produce a set of individually sufficient and jointly necessary conditions.

2

There are other ways of denying the existence of a common element. Timothy Williamson (2000: 44 ff.) has shown us that, though disjunctivism says very little, being largely engaged in denial, there might be what he calls a ‘non-conjunctivist’ position that says even less. What non-conjunctivism does, and all that it does, is to

deny the existence of a common element. It says, of some concept, that there is more than one way in which that concept can get to be instantiated, and that the various different instantiations have no common element. So no instantiation consists in a combination of some other instantiation *plus* some added extra, and there is no part of, or element in, any instantiation that is also a part of all the others.

There is a simple model for non-conjunctivism, the case of colour. Something is coloured if it is red, or blue, or green or . . . : there need be no end to this list. (That is another feature of non-conjunctivism, as compared with disjunctivism proper.) No colour consists in some other colour *plus* some further element or elements. And there is no common element to all the colours. We cannot say that to be red is to be coloured *and* red (though it is, of course, to be coloured in the red way), and no other candidate for the common element seems forthcoming.

It is important that one might still say that all colours have something in common, a common feature; colours, for instance, are a visual phenomenon (whatever that vexed phrase means). It is just that what they have in common is not a common element. And it is also important that none of the conjuncts, at least in the colour case, enjoys the sort of priority that disjunctivism attributes to genuine perception. But that is, no doubt, partly because we are dealing with a very different sort of case.

Not only is there a model for non-conjunctivism, but Williamson offers us an account of knowledge that is neither conjunctivist nor disjunctivist, even though a part of that account is a list of disjuncts, different ways of knowing that *p*. For Williamson, knowledge is the most general factive stative attitude. There are more specific stative attitudes, seeing, remembering, etc., and if I know that *p* I have one of those attitudes to *p*. But seeing that *p* is not knowing that *p plus* some other element, and Williamson's account of knowledge does not consist in the (very open-ended) list of disjuncts, ways of knowing, in the way that according to disjunctivists the account of experience consists in their two-line list.

Williamson does appeal to the example of colour, and it is tempting to generalize that example, and say that the basic model here is that of determinate and determinable. But though I used to think in that way, I now think it is better avoided. This is largely the result of reading David Stanford's impressive essay on the determinate/determinable distinction in the *Stanford Encyclopedia*. What I learned from this is that the distinction, first introduced by W. E. Johnson in his *Logic* (1921), is one of which it is far from easy to give a proper account. It seems, then, better to make such use of the colour model as one wishes and not to generalize it in terms of determinate and determinable.

Of course, once one is in the business of writing lists, all sorts of models are available. Any case of what one might call 'plural constitution' will suit our bill. Something is a good night out if it is one spent at a good play with a close friend, or one spent in a pub with other like-minded citizens, and so on. Or we could look at moral philosophy. An action is wrong iff it is murder, or rape, or torture, and so on. How about this: something is dangerous if it is a loaded gun or if it is indistinguishable from a loaded gun (the police will still shoot you for carrying it). But enough on this score.

3

With these preliminaries in hand, I now turn to the philosophy of action. The disjunctive conception that I considered in my (2000) was not a disjunctive conception of acting. The most general account of acting that I know is that to act is to cause a change; I accept this account and see no need to elaborate it disjunctively. What I did consider was a disjunctive conception of acting for a reason. There were two reasons for taking the possibility of such an account seriously. The first was that there is in the case of action the same pattern of successful and unsuccessful cases as we find in the case of perception, with the same array of possible responses. In both cases, that is, some philosophers are tempted to understand the successful case as what is present in the unsuccessful case *plus* some appropriate extra element, perhaps a suitable cause. Others, feeling that this approach distorts our understanding of the successful case, insist on a very different reading of the relation between the two, and on taking the successful case as the starting point, the primary point of reference. I will call these the Cartesian and the non-Cartesian approaches.

More explicitly, the Cartesian approach in the philosophy of action starts from a case where the agent's apparent reason for acting is that *p*, but where it is not the case that *p*. (This is only one form of unsuccessful case: there are others.) It announces that in such cases we are forced to say that the reason for which the agent acts is not that *p*, but that she believes that *p*. It then announces that this forces us to say the same of the successful case, the case where the agent is right in believing that *p*. In such cases, it remains the case that the reason for which the agent is acting is not that *p*, but that she believes that *p*. The non-Cartesian approach (mine) takes the opposite way in. It starts from a case where the agent is right about whether *p*, and takes it that in such a case the agent's reason for acting is that *p*; it then insists that this is still the agent's reason even in cases where it is not the case that *p*.

Both approaches are awkward, and both are awkward in the sort of way that is usually due to bad theory. As Aristotle said, they leave one saying things that nobody would say unless defending a theory. It is awkward to say that one's reason for action is never such a thing as that the sun has come out; the Cartesian view seems to say this, but it flies in the face of common experience. (In saying this I am *not* meaning to appeal to ordinary language.) But it is also awkward to say that one's reason is that the sun has come out when in fact it hasn't; the non-Cartesian view is in trouble too. It seemed, then, as though an explicitly disjunctive account might resolve the situation, at the cost, of course, of abandoning both the Cartesian and the non-Cartesian views. For what a disjunctive account does is to admit such differences as there may be between successful and unsuccessful cases, but somehow finesse the point by establishing that both are nonetheless perfectly good instances of the definiendum—or better, the intelligendum, since what we are looking for here is not a definition but an explanatory characterization.

The other motivation for thinking in disjunctive terms was that one was thereby enabled to accommodate all the necessary differences between successful and unsuccessful cases without being in danger of breaching an influential maxim in the

philosophy of action: the distinction between true and false should not affect the form of the relevant explanation.¹ This may not be apparent. The idea was that both Cartesian and non-Cartesian approaches respect this maxim by giving the same account of the one sort of case as they give of the other. It is only if we give different accounts of successful and unsuccessful cases that we are in danger of breaching the maxim. But a disjunctive conception can get away with giving different accounts without breaching the maxim. For our intelligendum is just 'A acts for the reason that *p*'. We can say that when A acts, the relevant explanation of his action will just be that he was doing it for the reason that *p*, whether A is right or wrong about that. And so our explanations will have the same form, whether A is right or wrong about whether *p*. In either case we will account for his action in terms of the reasons for which he acted. That we give a disjunctive conception of *this*, that is, that things begin to diverge from here on, involves no rejection of the explanatory maxim.

There seemed, then, to be good reason to look carefully into the prospects for a disjunctive conception of 'A acts for the reason that *p*'. However, in doing this, I was influenced by my earlier views that a disjunctive conception can offer a substantial conception of the second disjunct, and that it can offer more than two disjuncts. If we can say something substantial about a case which is not a successful one, why should we not suppose that there can be more than one sort of unsuccessful case? And, once we do suppose this, surely we should enshrine that possibility in our account. This is exactly what we find in the case of acting for a reason. There is indeed more than one way in which things can go wrong. The agent can be wrong about whether *p*, or wrong about whether if it were the case that *p*, that would be a reason for acting. These thoughts led to my first attempt at a disjunctive (better: trisjunctive²) conception of acting for a reason (Dancy 2000: 140):

A Φ s for the reason that *p* iff

- either* *p* and that *p* is a reason for Φ -ing and A Φ s in the light of the fact that *p*;
- or* it is not the case that *p*, but A takes it that that *p* is a reason for Φ -ing, and A Φ s in the light of his belief that *p*;
- or* *p*, but that *p* is not a reason for Φ -ing but A takes it that that *p* is a reason for Φ -ing, and A Φ s in the light of the fact that *p*.

This was my first account, but there is a mistake in the second trisjunct; this should have read:

- or* it is not the case that *p*, but A takes it that *p*, and A Φ s in the light of his belief that *p*.

One might suspect that some sleight of hand is involved in the phrase 'in the light of', a phrase which I admit is important to my approach to these matters. It is partly

¹ I associate this maxim with Bernard Williams, but in fact it was important to what used to be called the Strong Programme in the sociology of knowledge, notably in Bloor (1991).

² Better in one way, worse in another. The 'dis' in 'disjunctive' does not mean 'two', but 'apart', so that 'disjoint' does not mean 'split into two' but 'parted' or 'separate'. So my coining of 'trisjunctive' is an etymological sport. But I keep with it because it draws a useful contrast.

important because it enables me, even in the first trisjunct, to allow (indeed insist) that A believes that p ; for one cannot act in the light of a fact that one does not recognize to obtain. The point is only that A's belief that p is not, in the first trisjunct, specified as A's reason for Φ -ing. Still, one can, if one wishes, reformulate the whole thing thus:

A Φ s for the reason that p iff

- either* p and that p is a reason for Φ -ing and A's reason for Φ -ing is that p ;
- or* it is not the case that p , but A takes it that p , and A's reason for Φ -ing is that he believes that p ;
- or* p , but that p is not a reason for Φ -ing but A takes it that that p is a reason for Φ -ing, and A's reason for Φ -ing is that p .

This should be contrasted with what we would write if we were trying to copy a disjunctivist conception of experience:

A Φ s for the reason that p iff

- either* (1) p and that p is a reason for Φ -ing and A Φ s in the light of the fact that p (=A's reason for Φ -ing is that p);
- or* (2) it is for A as if (1) were the case.

It is worth listing the respects in which my version is, and is not, a version of disjunctivism proper. First, it offers no common element.³ Second, it mentions no common property or feature. Third, it is at least compatible with the claim that the three trisjuncts are indistinguishable from each other. Fourth, there is a sense in which the successful case is understood as reaching out to the world. Success as represented in the first trisjunct, after all, is not an extra that is added to what is present in the second trisjunct.

But there are further respects in which my version departs from disjunctivism's overall aims, beyond the mere fact that we have trisjuncts rather than disjuncts and so cannot run the pure disjunctive form. First, there is no explicit representation of the supposed primacy of the successful case. Second, it is not clear that the relationality involved in the first trisjunct is suitably like that which disjunctivists find in perception. There is a relation between agent and world required for the successful case, but what remains, if one takes that relation away, is capable of standing alone. Suppose it were not the case that p . Then the second part of the first trisjunct would have to be understood subjunctively, and we would be left with this:

That p , if it were the case, would be a reason for Φ -ing and A's reason for Φ -ing is that p .

³ As it stands, at least. But of course the conjunctive form of the various trisjuncts does open the door, as it were, to thoughts about common elements, and indeed I go on to allow that common to all three is that A believes that p . But even if this does stand as a common element, it is not that one trisjunct (say, the first) consists in the presence of another trisjunct *plus* some further feature, which is what the disjunctive account of perception is primarily concerned to avoid.

One might think that an element of relationality remains, in the expression ‘A’s reason for Φ -ing is that p ’. Now as I understand that phrase, it is not factive; A can act for the reason that p even where A is mistaken about whether p . And if it is not factive, it can remain in place untouched even when it is not the case that p . Others, however, disagree with me about this and hold that ‘A’s reason for -ing is that p ’ is factive, supposing it to entail that p ; and at the moment, in considering the possibility of a disjunctive account, we are trying to accommodate their doubts. And one might, for that reason, say that there is an appropriate relationality in the first disjunct. But, as I say, the matter is debatable.

So my trisjunctive account does not really enjoy some of the more interesting features of true disjunctivism. And I have a further reason of my own for rejecting it, which we have already seen. It seems to me that the second trisjunct, if we take it *au pied de la lettre*, is false. This is, of course, a substantial matter. But the arguments I gave for this conclusion in my (2000) still seem to me to be good ones. They amounted to pointing out that there is a substantial difference between acting for the reason that p and acting for the reason that one believes that p . If my reason for action is that I believe that there are pink rats in my boots, or that I believe that everyone is out to get me, I will probably be calling a psychiatrist. If my reason is that there are pink rats in my boots, I will probably be calling a pest control expert, and if it is that everyone is out to get me, I probably won’t be calling any people in white coats at all. A rather different example is my crumbly cliff. Suppose that I believe that the cliff is crumbly, and so believing, I am more likely to get nervous and lose my grip. This very fact—that I believe it to be crumbly—might be my reason for not climbing it. And this is a different reason from the reason that it *is* crumbly. If my reason were that the cliff is crumbly, that reason would vanish if the cliff ceased to be crumbly; if my reason is that I believe the cliff to be crumbly, that reason persists so long as I have that belief, whether the cliff ceases to be crumbly or not.

For these reasons (all elaborated in Dancy 2000: ch. 6), I reject the suggestion, encapsulated in the second trisjunct, that, where things are not as I suppose them to be, my reason is not that p , but that I believe that p . So I cannot accept the trisjunctive account; it seems to me to involve a substantial philosophical error. My own view now is that one could combine the first two trisjuncts into one, returning us to a disjunctive theory—of a sort:

A Φ s for the reason that p iff

either that p is or would be a reason for Φ -ing and A’s reason for Φ -ing is that p ;

or whether it is true that p or not, that p would not be a reason to Φ , but A takes it that that p is a reason to Φ , and A’s reason for Φ -ing is that p .

Another way of writing this would be:

A Φ s for the reason that p iff

either that p is or would be a reason for Φ -ing, A takes it that p and A Φ s in that light;

or whether it is true that p or not, that p would not be a reason to Φ , but A takes it that p and that that p is a reason to Φ , and A Φ s in that light.

On this version the original trisjunction had a mistaken focus, since one of the driving forces behind it was the possibility that the agent's belief that p be false. As things now stand, however, the issue is not the truth or falsity of that belief, but the light cast on the action by what one believes, whether what one believes is true or not. Would the supposed fact that p cast a favourable light on Φ -ing if it turned out to be the case? This gives us:

A Φ s for the reason that p iff

either A Φ s in the light actually cast by what he believes, namely that p ;
or A Φ s in the light apparently cast by what he believes, namely that p .

But the real point, for me, is that each of these disjuncts effectively ends with the very thing we were trying to give an account of. We could have shortened the whole thing and ended up with just this: A Φ s for the reason that p iff A takes it that p and A Φ s in that light. Here all pretence of disjunctivism has been abandoned.

What about the pure disjunctive account given above? Here it is again:

A Φ s for the reason that p iff

either (1) p and that p is a reason for Φ -ing and A Φ s in the light of the fact that p
 (=A's reason for Φ -ing is that p);
or (2) it is for A as if (1) were the case.

The main problem for this account is that it is empirically false. Its seeming to me as if I am acting in the light of a particular fact is no proof at all that this fact is really the reason for which I am acting. We just don't have this sort of access to our reasons. It might well be that I seem to myself to be acting from pure altruism when in fact my motivation is not at all something to congratulate myself about.

Things would be no better if we had followed the line of reasoning just before, and offered this account:

A Φ s for the reason that p iff

either (1) A Φ s in the light actually cast by what he believes, namely that p ;
or (2) it is for A as if (1) were the case.

This is still empirically false, and for exactly the same reason. So my present conclusion is that no disjunctive account of acting for a reason is defensible.

In this final section I consider suggestions made by Jennifer Hornsby at the conference in Glasgow of which this volume is largely, but not entirely, a record, and adapted in her essay in this collection. I discuss four issues.

4.1

Hornsby's original proposal was for a form of disjunctivism in the theory of action that is very different from the ones I have considered so far, thus:

If X Φ -ed because X believed that p , then

either X Φ -ed because X knew that p ;

or X Φ -ed because X treated [the proposition that] p as an agent who knew that p would treat it, although X didn't actually know that p .

This is in fact a non-conjunctive account which simply borrows Williamson's suggested account of belief (Williamson 2000: 46–8), under which to believe that p is either to know that p or to treat the proposition that p as someone who knows that p would treat it. In her published paper it now becomes (in shortened form):

(D) If X Φ -ed because X believed that p , then

either X Φ -ed because X knew that p ;

or X Φ -ed because X merely believed that p .

My rejection of the disjunctive account of acting for a reason is no reason for me to reject all disjunctive accounts of anything whatever in the philosophy of action. So this new proposal, with its different intelligendum, needs to be considered on its merits. In doing so, I bear in mind that I had two sorts of complaint about the accounts which I considered in the previous section. I maintained in some cases that some part of the account was philosophically misconceived, and in others that the account as a whole was empirically false. Before looking to see whether this new account (in its two different versions) is vulnerable in similar ways, however, it is worth quickly running through the marks of disjunctivism in order to show that what we are dealing with is genuinely disjunctive. I take the first version first, because the two versions differ significantly in this respect. First, the two disjuncts have no visible common element. Second, they have no visible common quality or feature. Third, there is a sort of indistinguishability involved; if X treated the proposition that p as someone who knew that p would treat it, it must be for X as if she knew that p . Otherwise X would presumably treat it differently. Fourth, the first disjunct is world-involving, because of the way in which it treats the concept of knowledge. Fifth, it enshrines the desired primacy of the successful case. And finally, there is the possibility of empirical falsehood, since it might turn out, on further examination, that it is possible to act because one believes that p without satisfying either disjunct.

My own view, at the moment, is that this is just what happens, and so that as things stand we need a further disjunct. (And I will argue below that the reasons for restricting one's account of experience to two disjuncts rather than some other number, say three, do not apply to one's account of intentional action.) It seems to me that I am perfectly well aware of many things that I do not know to be true, but on the basis of which I am willing to act. I am not treating them as someone who knew them would treat them, because, at least to some extent, I retain an open mind. The fact that I am

willing to act on them is, of course, not itself enough to establish that I am treating them as someone who knew them would treat them. If this is right, however, the fault lies really with Williamson's account of belief.

I now turn to the new version. It has the first and second marks of disjunctivism, and the third if, where X merely believes that *p*, it is for X as if she knew that *p*. It has the fourth, but seems to lack the fifth. Finally, it is not obvious that the account is empirically falsifiable. A way to reintroduce the possibility of empirical falsehood would be to provide a substantial characterization of 'merely believing'; for the account as it stands seems to say that X acts because X believes that *p* if either X acts because X knows that *p* or X doesn't know that *p* but acts because X believes that *p* anyway; and this is not terribly informative until we get the desired account of 'merely believing'. But once we get that, as in the earlier version, the possibility of empirical falsehood returns. Hornsby says, however, that her (D) 'makes no general claim about the relation between knowledge and belief'. My own guess, on this point, is that the introduction of any substantial account of 'merely believing' in the guise of a substantial relation between believing and knowing will create a need for a third clause if we are to achieve a set of disjuncts that are jointly necessary and individually sufficient, and thus avoid empirical falsehood. But I note that this is an aim that Hornsby views as optional.

So we already have two ways in which the new version differs from the paradigm case of disjunctivism. It does not enshrine the primacy of the successful case, because it does not attempt to understand merely believing in terms of some relation to knowing (which the first version did). And of course both versions, as they stand, offer only necessary conditions rather than necessary and sufficient conditions, and as such can only claim to offer a partial understanding of the relevant intelligendum.

Hornsby suggests that her account does do something that one might ask of a disjunctive account, namely that it simultaneously brings together and distinguishes acting from knowledge and acting from mere belief, in the sort of way that McDowell's perceptual disjunctivism simultaneously brings together and distinguishes the experience of a perceiver and the experience of someone who has a mere experience. But this seems to be an absolutely minimal conception of disjunctivism; McDowell surely goes further than this in his claim that non-perceptual experience is a case in which it is for the experienter as if he were perceiving.

4.2

I turn now to the question whether there are any philosophical complaints to be made about the way in which the account has been crafted.

Hornsby thinks of her disjunctive thesis as concerned to relate cases in which the agent acts for an 'objective reason' to those in which her reason is a 'subjective reason'. At least, those are the terms in which I would characterize the concern. Hornsby however uses the terms '(F)-type reason' and '(B)-type reason', and I will follow her in this. The agent's reason is an (F)-type reason if it is of the form 'that *p*'; such reasons are

canonically expressed in the form 'A reason for X to Φ was that p '. Subjective reasons, for Hornsby, are expressed in this form: 'X had a reason to Φ : X believed that p '. This distinction between (F)-type and (B)-type reasons involves us in difficult issues. Hornsby starts from two points on which I think everyone should agree. The first point is that someone's acting intentionally is always explained partly by her believing something, or that her believing as she does is always at least relevant to the explanation. This, I would say, is true, though it does not show that the agent's so believing was any part of the reason she had for acting as she did. The second point is that intentional explanations of action can be given in two forms, psychologized or non-psychologized. We can say 'the reason why he is running is that he thinks the train is leaving' or we can say 'the reason why he is running is that the train is leaving'. My own view about these two forms is that they are effectively equivalent, both being expressible as 'the reason why he is running is that, as he supposes, the train is leaving'. Hornsby's view is rather different. She wants to say that the sentence 'X had a reason to Φ : X believed that p ' plays a distinct explanatory role. Someone who thinks he is late when he is not late has a reason to run nonetheless. What is that reason? One might think that it cannot be that he is late, since he is not late. What can it be, then? The only thing left is that he believes that he is late. If this were right, Hornsby's phrase: 'X had a reason to Φ : X believed that p ' would have to be read, more explicitly, as 'X had a reason to Φ , namely that X believed that p ', and *not* as 'X had a reason to Φ , namely what X believed, that is, that p '. But this is not the way that Hornsby wants to take it. We are dealing with a (B)-type reason when something of this form is true: 'X had a reason to Φ : X believed that p '. But the (B)-type reason specified here is not that X believed that p . So the phrase 'X had a reason to Φ : X believed that p ' is not to be understood as 'X had a reason to Φ , namely that X believed that p '. This is the way I was originally tempted to take it, but Hornsby stresses that the reason X had in these cases is the content of her belief, that is, that p , not that X believed that p .

It seems, then, that whether one's reason is (F)-type or (B)-type, it is still that p . An (F)-type reason must be something that is the case, but a (B)-type one need not be. One has a (B)-type reason as soon as one comes to believe that p , whether one is right or wrong on the matter. (Though it seems to be required for this that that p , if it were the case, would be an (F)-type reason, and one could add further conditions as well, such as that one is not transgressing any relevant norms in believing that p .) So apparently something that is not the case can be a reason for one to act, and a reason for which one acts, so long as one believes it (with any further conditions, of course). All we should say is that it is a (B)-type reason rather than an (F)-type one, and canonically expressible in the form 'X had a reason to Φ : X believed that p '.

Consider now the standard worry that if we distinguish between (F)-type and (B)-type reasons in this sort of way we will end up with two sets of reasons to act. Is that true within Hornsby's picture? (F)-type reasons are matters of fact, and they become (B)-type ones once one comes to believe them to be the case. The remaining (B)-type ones are not (F)-type reasons; they are those among the things we believe which would be (F)-type reasons if they were the case. (Note that we are saying here that they would

be (F)-type reasons if they were the case, not that they would be reasons if they were the case; for on Hornsby's account they are reasons already.) What then are we to say to someone who is proposing to do something which there is (F)-type reason not to do, but which she has good (B)-type reason to do? How should we advise her? We cannot, it seems, deny the reality of her (B)-type reasons; they are no less real or genuine than her (F)-type ones. But this seems to lead to two difficulties.

First, it might be that the strength of the (B)-type ones is such as to defeat the (F)-type ones on the other side, and that does not seem to me to be a good result. But it looks as if the only way to avoid it is somehow to downgrade the (B)-type ones. There are two main ways of doing this. We can say that a (B)-type reason is just a consideration that would be a reason if it were the case—which in this case isn't, and so it isn't in fact a reason at all. Or we can say that a (B)-type reason affects the rationality of our action; rationality requires one to act according to one's beliefs, true or false. But it does not go to determine what in fact we have reason to do, which, I think, really means that it is not a reason at all.

The second difficulty is that we outsiders don't really know how to relate a third person's (F)-type reasons to her (B)-type reasons. Once we have worked out what her (F)-type ones are and what her (B)-types ones are, how do we assess them against each other? (The previous worry allowed that reasons of the two sorts are somehow commensurable; the present one questions that.) It is not at all clear how in a case of conflict (and this would be the normal situation) we should advise her. We can't just say 'go with the (F)-type reasons' because we are trying to allow that her (B)-type reasons are equally *reasons* in some way.

Of course all this is a topic for a paper in its own right, and one that I have treated elsewhere (Dancy 2000: ch. 3). Let me just say that where we have both:

X Φ -ed for a reason that X had: X believed that p —and
 X Φ -ed for a reason: that p

Hornsby's approach is close to saying that there are two reasons around, the one that X had and another, real one. This seems to me very uncomfortable, especially because as far as I can see the second claim entails the first.

4.3

I now turn to a different issue. Hornsby's second version, in full dress, runs as follows:

(D) If X Φ -ed because X believed that p , then
 either X Φ -ed because X knew that p (and Φ -ed *because* p);
 or X Φ -ed because X merely believed that p .

The added parenthesis introduces Hornsby's view that when someone Φ s because they merely believe that p , they do not Φ because p , even when they are right about whether it is the case that p or not. So truly believing that p does not itself put you in a position to Φ because p ; it only does that if your true belief is knowledge.

I confess that I am very suspicious of Hornsby's intelligendum 'X Φ -ed because X believed that p '. It needs to be restricted in a certain way if it is to be even relevant

to the topic of intentional explanation, since there many cases where one acts because one believes but in which one's so believing has little to do with any reason for which one acted (though it may be the reason, or part of the reason, *why* one acted). I much prefer the intelligendum with which I started: $X \Phi$ -ed for the reason that p (or something like that). But Hornsby has a quite different attitude to that proposition. She glosses her first disjunct ' $X \Phi$ -ed because X knew that p ' as equivalent both to ' $X \Phi$ -ed because p ' and to ' $X \Phi$ -ed for the reason that p '.⁴ This effectively announces that when one merely believes that p , one cannot act for the reason that p , even if one is right in so believing, and so takes a stance on the very interesting question whether one can act for a reason which one does not know to be the case, but merely believes to be so. Here is the example that Hornsby relies on in suggesting that one cannot do this:

Edmund . . . believes that the ice in the middle of the pond is dangerously thin, having been told so by a normally reliable friend, and . . . accordingly [Edmund] keeps to the edge. But Edmund's friend didn't want Edmund to skate in the middle of the pond (never mind why), so that he had told Edmund that the ice there was thin despite having no view about whether or not it actually was thin. Edmund, then, did not keep to the edge *because* the ice in the middle was thin. Suppose now that, as it happened, the ice in the middle of the pond was thin. This makes no difference. Edmund still didn't keep to the edge *because* the ice was thin. The fact that the ice was thin doesn't explain Edmund's acting, even though Edmund did believe that it was thin, and even though the fact that it was thin actually was a reason for him to stay at the edge. (this volume: Section 1.3)

Hornsby's point (echoed by John Hyman (1999) and Timothy Williamson (2000: 60–4)) is that action on a true belief in a reason-giving fact is not enough for it to be the case that one is acting for that fact as a reason. It is only if one knows that fact that one can be said to be acting for it as a reason. But this seems to me to involve an invalid inference. One cannot argue that Edmund is not keeping to the edge for the reason that the ice is thin from the premise that he is not skating there *because* the ice is thin. In fact, I would say of both scenarios above (the one where the friend is right and the one where he is wrong) that in them Edmund is skating on the edge for the reason that the ice is thin, or that Edmund's reason for skating there is that the ice is thin, or that the reason for which he is skating there is that the ice is thin. I don't see that Hornsby's two scenarios do anything to upset that entirely natural position. And I therefore think it would be a mistake, so far as that goes, to gloss her first disjunct ' $X \Phi$ -ed because X knew that p ' as equivalent to ' $X \Phi$ -ed because p ' and to ' $X \Phi$ -ed for the reason that p '.

There are, however, other reasons to think that one can only act on a reason which one knows to be the case. Suppose that I am meeting my daughter off the train. I don't know that she is on the train, and I know that I don't know this, but I do believe that she is on the train. So I head off to meet her. What would I say if asked what my reason is for meeting this train? I might say 'My daughter is on it'. But I would be more likely to say 'I think my daughter is on it', or 'My daughter should be on it', or 'My daughter is probably on it'. There is some pressure to avoid offering the bald

⁴ The notion of a reason for which one acts is connected to that of an (F)-type reason at the end of section 1.3 of Hornsby's essay.

claim that my daughter is on the train as my reason if I do not know this to be the case. But what if my response to this is to specify my reason thus: 'I think my daughter is on the train'? Are we to say that I have now shifted my ground to something I know? In line with a certain approach to Moore's Paradox, I think that this remark is not to be understood as autobiographical, but as a more guarded version of the claim that she is on the train. It is not, then, that pressure can be put on me to restrict myself, in giving my reasons for doing what I do, to things that I take myself to know. The pressure is more to express, in the way that I give my reasons, the degree of confidence with which I hold them.

So my own view is that there is not sufficient reason to restrict the phrase 'X Φ -ed for the reason that p ' to cases where X knew that p . This releases that phrase for use as the overall intelligendum.

4.4

I end, as promised, with the question whether Mike Martin's reason for insisting that the only permissible characterization of the non-perceptual experience, at least from the disjunctivist's point of view, is that it is any episode relevantly indistinguishable from a perception, applies with equal force in the theory of action. Martin's main reason was that only thus are we able to provide a disjunctive account of experience under which the two disjuncts are individually sufficient and jointly necessary for an episode to be an experience. A positive characterization of the second disjunct would create a danger that, even though all episodes with that character will be experiences, there might yet be episodes of a different sort that are experiences as well. If so, the disjunctive account as a whole will offer conditions that are individually sufficient, but not jointly necessary. One might ask why that is so worrying. Is it merely a fetish with the disjunctive form? If we need to add a third clause, and turn the whole thing trisjunctive, what is wrong with that? One answer to this question is that the process might go on indefinitely, so that prospects of giving a genuine *account* of experience begin to fade. Here we might remind ourselves of Williamson's list of forms of knowledge. There are many different forms of knowledge (perceiving, remembering, proving, intuiting) but the list of these does not, for Williamson, constitute any *account* of knowledge; that account consists in the claim that knowledge is the most general factive stative attitude. Perhaps there is the threat that the process might go on indefinitely. But that is only a threat. There might be ways of averting the threat. In the philosophy of action, it seems not impossible to achieve that. Take one of the versions of my original trisjunctive account:

A Φ s for the reason that p iff

- either* p and that p is a reason for Φ -ing and A's reason for Φ -ing is that p ;
- or* it is not the case that p , but A takes it that p , and A's reason for Φ -ing is that he believes that p ;
- or* p , but that p is not a reason for Φ -ing but A takes it that that p is a reason for Φ -ing, and A's reason for Φ -ing is that p .

It seems to me possible that we could have an *a priori* certainty that no more clauses will be required. What other sort of case could there be?

A rather similar response seems to me appropriate to a concern of Martin's that is voiced in the passage that I quoted earlier, but which I did not pick up at that point. The relevant parts are these:

If what marks [the non-perceptual] cases out in the first place is just that they involve the absence of perception, then one may worry that whatever fixes what they have in common with each other will apply equally to any case of perception [This gives us a common element.] Now if the common element is sufficient to explain all the relevant phenomena in the various cases of illusion and hallucination, one may also worry that it must be sufficient in the case of perception as well. In that case, disjunctivism is threatened with viewing its favoured conception of perception as explanatorily redundant. (Martin 2004: 46)

There is a term of art here that is not sufficiently remarked: 'one may worry that *p*'. One might respond to such remarks 'one may but one need not'. Consider the first worry. Martin's passage here might have started 'If the only thing that marks out the non-perceptual cases is that they involve the absence of perception, whatever else they have in common will apply equally to perception'. If he has said this, his claim would have been undeniable. But he doesn't quite say that. What he does say is rather that if we start by distinguishing a group of cases simply as those that do not involve perception, we (might worry that we) will be unable to find any further feature or element that they, and only they, have in common. Well: one might worry about this, but it is far from certain that it will turn out to be so.

Now turn to the second worry. It is interesting that Martin allows that this worry, which is about explanatory redundancy, arises anyway, independent of the present issue; indeed, he sees it as posing a serious challenge to disjunctivism, and gives an answer to it. If that worry arises anyway, and there is an answer to it, this particular source of worry does not add any difficulty to what we have already. Further, what are 'the relevant phenomena'? The disjunctivist maintains that genuine perception is a world-involving state. Is it that this claim is useless in the explanation of any 'relevant phenomena'? If so, this would be bad news for naïve realism. Luckily, there seems to be room to say that there are certain phenomena that can be explained in either of two ways—though not in both at once.

I conclude that the reasons given for rejecting any positive characterization of the non-perceptual should not deter us in the philosophy of action, and may not even be compelling in the philosophy of perception.

REFERENCES

- Bloor, D. (1991) *Knowledge and Social Imagery*, Chicago: University of Chicago Press.
- Crane, T. (2005) 'The Problem of Perception', in Edward N. Zalta (ed.) *Stanford Encyclopedia of Philosophy* (Spring 2005 Edition) <<http://plato.stanford.edu/archives/spr2005/entries/perception-problem/>>.

- Dancy, J. (1995) 'Arguments from Illusion', *Philosophical Quarterly*, 45: 421–39.
—— (2000) *Practical Reality*, Oxford: Clarendon Press.
Hornsby, J. (this volume) 'A Disjunctive Conception of Reasons for Acting'.
Hyman, J. (1999) 'How Knowledge Works', *Philosophical Quarterly*, 49: 433–51.
Johnson, E. (1921) *Logic*, Cambridge: Cambridge University Press.
Martin, M. G. F. (2004) 'The Limits of Self-Awareness', *Philosophical Studies*, 103: 37–89.
Williamson, T. (2000) *Knowledge and its Limits*, Oxford: Clarendon Press.

PART III
KNOWLEDGE

McDowellian Neo-Mooreanism

Duncan Pritchard

1 MOOREANISM

Consider the following three-part response to the standard formulation of the radical sceptical problem, a response which mirrors in key respects the ‘common-sense’ proposal often ascribed to G. E. Moore, and which is regarded with almost wholesale derision.¹ The first part of the anti-sceptical response is to claim that one knows a paradigm ‘everyday’ proposition—that is, a proposition that all of us would think we know, in those circumstances—such as, to use a famous example, that one has two hands. A claim of this sort is surely intuitive. The second part of the response is to note that since having two hands is inconsistent with the relevant sceptical hypothesis, such as the hypothesis that one is a (handless) brain-in-a-vat (BIV), it follows that if one knows that one has two hands then one also knows the denial of the relevant sceptical hypothesis, in this case that one is not a BIV. Again, this claim is also intuitive. Finally, the third part of the response is the extraction of the anti-sceptical

Earlier versions of the material that make up this essay have been presented at the following conferences: *Contextualism*, Bled, Slovenia (June 2004); *Contextualism*, Amsterdam Free University (October 2004); *Externalism and Internalism in Semantics and Epistemology*, University of Kentucky (April 2005); and *Disjunctivism*, University of Glasgow (June 2005). Earlier versions of this material have also been presented at Faculties of Philosophy at the Universities of York, Birmingham, Copenhagen, Aarhus (Denmark), Rijeka (Croatia), Hull, and Nottingham between 2004 and 2006. I am grateful to the audiences on all these occasions. Special thanks go to Jessica Brown, Tony Brueckner, Earl Conee, Gary Ebbs, Sandy Goldberg, Marie McGinn, Alan Millar, Ram Neta, Jim Pryor, Finn Spicer, two anonymous referees from Oxford University Press, and, especially, Adrian Haddock (it is hard to imagine a more conscientious editor).

¹ For my own part I actually think there is a lot more (if you’ll excuse the pun) going on in Moore’s writings on scepticism than Mooreanism suggests, but I will not be getting into this exegetical issue since it is Mooreanism, rather than Moore’s own nuanced anti-sceptical view, which is my focus here. For Moore’s two key writings on scepticism, see Moore (1925, 1939). For an overview of some of the exegetical issues in this regard, see Baldwin (1993).

conclusion that one knows the denial of the relevant sceptical hypothesis, in this case that one is not a BIV. Here, then, is a rundown of the argument, sticking to the BIV hypothesis throughout:

Mooreanism

M1. I know that I have two hands.

M2. If I know that I have two hands, then I know that I'm not a BIV.

MC. I know that I'm not a BIV.

It is important to the Moorean view that there is nothing more to the stance than the presentation of an argument of this sort. The sceptic has called our knowledge into question, via the presentation of the sceptical hypothesis, and the Moorean, via his opposing argument, has rebutted the sceptic's claims. Thus, there is no case to answer, and hence nothing more that needs to be said. In this sense, then, the Moorean stance is a pre-theoretical proposal, in that it attempts to deal with the sceptical challenge in an entirely common-sense way which avoids the need for a theoretical response to the problem.

There are a number of problems with the Moorean strategy, but I will not be attempting to elucidate them all here.² Instead, I will just mention three key problems with the view.

Perhaps the most common complaint levelled at the Moorean argument is that there is something question-begging about responding to the sceptical problem in this way, in that it simply takes as an unquestioned premise in its argument the denial of the very claim that the sceptic will want to motivate as a conclusion of her argument. Call this the dialectal impropriety objection.³

We can get a grip on what the problem is here by considering the sceptical argument to which the Moorean stance is supposed to be a response. Keeping to the first person, as Moore does, and sticking with the examples used above of having two hands and not being a BIV, we can formulate the opposing sceptical argument as follows:

Scepticism

S1. I don't know that I'm not a BIV.

S2. If I know that I have two hands, then I know that I'm not a BIV.

SC. I don't know that I have two hands.

Both of these premises are intuitive. The first is intuitive because it seems that whether or not we are BIVs is just something that we could never know because sceptical scenarios are defined such that there is nothing in our experiences that could offer us any definitive indication one way or another as to whether we are the victim of such a scenario. The second premise is exactly the same as the second premise in the Moorean argument, which we saw was intuitive above. This premise can be further motivated in terms of the closure principle:

² See Pritchard (2007a) for a fairly comprehensive list.

³ See, for example, Wright (2002) for a complaint against Mooreanism of this sort.

Closure

For all S , φ , ψ , if S knows φ , and knows that φ entails ψ , then S also knows ψ .⁴

Closure certainly seems plausible, since it is hard to see how this principle could fail. How could one know one proposition, know that it entailed a second proposition, and yet fail to know the second proposition? Crucially, however, with closure in play—and given that one knows the relevant entailment, as presumably one does—it follows that if one knows that one has two hands then one also knows that one is not a BIV. We thus get the second premise, motivated in terms of the highly intuitive closure principle.

With these two premises in hand, however, the sceptical conclusion immediately follows. And since this argument can be repeated with any number of everyday propositions (one would just have to vary the sceptical hypothesis to suit), so the full intellectually devastating radical sceptical conclusion—that we are unable to know most of the empirical propositions which we typically think we know—is in the offing.

With the sceptical argument and the Moorean argument set side-by-side, one can see that the debate here encapsulates that old philosophical chestnut that one philosopher's modus ponens is another philosopher's modus tollens. Whereas the Moorean takes his everyday knowledge as secure and argues on this basis that he also has the required anti-sceptical knowledge, the sceptic begins by highlighting the implausibility of anti-sceptical knowledge and argues on this basis that we also lack everyday knowledge. With the debate so construed, however, one can see why the Moorean strategy can seem so dialectically inappropriate. The sceptic has given us an apparently compelling argument for thinking that we lack everyday knowledge. In response, the Moorean simply helps himself to the denial of the contested conclusion and reasons on this basis to the negation of the premise of the sceptical argument. Given that the Moorean argument begins and ends with this strategy, it is little wonder that few find it persuasive. At the very least, some sort of diagnostic story needs to be offered by the Moorean to explain away the intuitive appeal of scepticism, since without such a story we seem to be passing the problem by. Part of this diagnostic story will inevitably involve an epistemological theory to back up the view—a theory which explains, for example, why we can know the denials of sceptical hypotheses after all, contrary to our first intuitions on this score.

A second, and related, difficulty with the Moorean response is that it seems to offer us, at most, a draw with the sceptic, rather than a resolution of the sceptical problem. After all, given that the sceptical argument is just the modus tollens to Moore's modus ponens, and since both arguments have intuitive premises, it appears that the dialectical situation is that we are faced with two opposing arguments of equal force. If this is right, then even despite the Moorean argument we still have just as much reason to be sceptics as to be Mooreans. Put another way, it is still the case on the Moorean view that we have no good reason not to be sceptics. This is a kind of second-order scepticism which, while not obviously reducible to its first-order cousin (which would hold

⁴ One might want to modify this principle in a number of ways in order to deal with potential counterexamples of a trivial sort (such as possible cases where the agent doesn't even believe the entailed proposition), but this unembellished version of closure should suffice for our purposes here.

that we have reason to be sceptics), is still enough to make Moorean anti-scepticism not nearly as intellectually satisfying as it might at first appear. Call this the *impasse* objection.⁵

Finally, a third key problem with Mooreanism is that the Moorean argument seems to consist of a series of assertions which strike one as conversationally inappropriate, if not just plain absurd or contentless. As a number of commentators have noted—most trenchantly Wittgenstein (1969) in his final notebooks—the assertions in question in the Moorean argument seem to offend against our usual usage of the term ‘know’. This phrase plays a very special role in our practices of knowledge self-ascription, but, crucially, not one that seems applicable to the kind of anti-sceptical assertions that the Moorean makes. We never normally say that we know the denials of sceptical hypotheses, and neither do we usually say that we know everyday propositions which are just plain obvious to everyone in that conversational context. Call this the conversational impropriety objection.

Although they by no means form an exhaustive list, these three objections do capture the heart of the difficulties facing Mooreanism, and thus highlight the work that a neo-Moorean view, one which can evade the problems facing Mooreanism, has to do.

2 CONTEMPORARY NEO-MOOREANISM

While no one is persuaded by the Moorean anti-sceptical view, there are positions available in the literature which broadly mirror such a stance, what I have elsewhere termed “neo-Moorean” theories (Pritchard 2002b: section 8; cf. Pritchard 2002c). The chief distinctive aspect of this sort of anti-sceptical theory is that it confronts the sceptical problem head-on by allowing that we can know the denials of sceptical hypotheses. Accordingly, the view avoids scepticism while retaining the closure principle by denying the first premise in the sceptic’s argument, (S1). Furthermore, neo-Mooreanism achieves this while avoiding, as much as possible, epistemological revisionism.⁶

One of the key issues facing neo-Mooreanism is how to explain how we can know the denials of sceptical hypotheses. The standard line in this regard usually adverts to some form of the safety principle for knowledge, as defended, for example, by Ernest Sosa (1999). This principle holds, roughly, that what is essential to knowledge is that

⁵ Wright (1991) discusses this ‘*impasse*’ objection to Mooreanism, and offers an argument which shows that the second-order scepticism which results collapses into first-order scepticism.

⁶ This aversion to revisionism is why standard contextualist anti-sceptical theories—as defended, for example, by DeRose (1995), Lewis (1996), and Cohen (e.g. 2000)—would not count as neo-Moorean proposals, even though they also allow that we can know the denials of sceptical hypotheses. This is because the contextualist achieves this by offering a complex and revisionistic philosophical theory about how our use of the term ‘knows’ is highly context-sensitive. For the neo-Moorean, however, such revisionism is unnecessary, and merely distracts one from offering the straightforward neo-Moorean response to the sceptic that is both required and available. For more on the relative merits of neo-Moorean and contextualist anti-sceptical theories, see Pritchard (2005a: pt. 1; 2007a).

one has a belief that could not have easily been false. The basic idea is that provided sceptical error-possibilities are indeed far-fetched, then it follows that one's true belief that one is not a victim of such an error-possibility will be such that it couldn't have easily been false, and so can count as knowledge.

The underlying motivation for safety comes from the intuition that what is essential to knowledge is that it is non-lucky true belief, where the safety principle captures the heart of this anti-luck intuition. Thus, the claim is that provided that one's environment is epistemically friendly, then one's anti-sceptical beliefs will not be lucky and hence can count as instances of knowledge.⁷

Clearly, however, the neo-Moorean cannot leave the story there, since we still need to be told how the neo-Moorean stance can avoid the problems which we saw afflicting Mooreanism. Moreover, we are also owed an explanation of the evidential basis of this anti-sceptical knowledge. In order to begin filling in some of this detail it is essential that we first factor the epistemic externalism/internalism distinction into this discussion, since the nature of the further detail will be dependent upon which side of this contrast one stands.

By epistemic internalism, I mean access internalism such that what makes an epistemic condition (that is, a condition which, perhaps in conjunction with other epistemic conditions, can turn true belief into knowledge) an internal epistemic condition is that the agent concerned is able to know by reflection alone those facts which determine that this condition has been met. Meeting the justification condition, for example, at least as it is standardly conceived, involves the possession of grounds in support of the target belief, where these grounds—and the fact that they are supporting grounds—is reflectively accessible to the subject. Epistemic externalism denies this, and so holds that there are epistemic conditions which do not demand reflective access on the part of the subject of this sort. I will understand internalism about knowledge as being the view that meeting a substantive internal epistemic condition is necessary for knowledge possession, with externalism about knowledge as the denial of this thesis.⁸

In order to see this distinction in action, consider how it applies to some of the core cases over which internalists and externalists diverge, such as the chicken-sexer case. Here we have an agent who is exhibiting a highly reliable cognitive ability (to distinguish between the sexes of chicks), and yet who has false beliefs about how she is doing what she is doing (she thinks she is using her senses of sight and touch, when actually it is her sense of smell) and who typically lacks good grounds for thinking

⁷ For more on the issue of how to formulate safety, and its connections to anti-luck epistemology, see Pritchard (2005a, *passim*), (2005b, 2007b).

⁸ There are other ways of drawing the classical internalism/externalism distinction of course. One could put the point in terms of supervenience rather than access—see, for example, Conee and Feldman (2000)—or one might weaken the internalist requirement by saying that one only needs reflective access to the supporting grounds for one's belief and not also to the fact that they are supporting grounds, as Alston (1988) suggests (though note that he doesn't regard this view as an internalist thesis as such). The account of the distinction offered here is fairly standard, however, and, I believe, it also gets to the heart of what is at issue in the key debates between epistemic externalists and internalists. For more on the epistemic externalism/internalism distinction in general, see Kornblith (2001).

that she is reliable in this respect. Clearly such an agent does not meet a substantive internal epistemic condition as we have just defined that notion. Nevertheless, that, for the externalist at least, there is still an issue about whether or not the agent has knowledge indicates that what is in question here is the necessity to knowledge of meeting such an internal epistemic condition, which is just as it should be given how I have just characterized the externalism/internalism contrast.⁹

There is also a further assumption which is common to the externalism/internalism debate, albeit not one that (as we will see in a moment) is shared by everyone. This is that what an agent has reflective access to can never be such as to entail any specific empirical fact. One's reason for believing that there is a chair in front of one cannot be, for example, that one sees that there is a chair in front of one (which entails that there is a chair there), but only, say, that it seems to one as if there is a chair in front of one (or some other non-factive description of one's reason), which clearly doesn't entail any specific empirical fact. With this further claim in mind, we get what I will call the classical externalist/internalist thesis. We will return to consider this assumption in more detail below.¹⁰

It ought to be clear straight away that motivating a classical internalist version of neo-Mooreanism is not going to be at all easy. This is primarily because it is hard to see how we could account for our knowledge of the denials of sceptical hypotheses on this view. Loosely speaking, I take it that the worry here is that since there is nothing in our present experiences that we can, as it were, introspectively 'point' to in order to indicate that we are not the victim of a radical sceptical hypothesis, it follows that we are unable to have adequate reflectively accessible grounds to support our beliefs in this respect. Accordingly, on an internalist account, lack of knowledge of the denials of sceptical hypotheses seems to follow fairly quickly.

One can put this point into sharper relief by considering what the evidential basis of our putative anti-sceptical knowledge would be on the classical internalist view. Consider the following 'underdetermination' principle:

Underdetermination

For all S , φ , ψ , if S knows φ , and S knows that φ entails ψ , then S 's evidence for believing φ favours φ over $\neg\psi$.¹¹

In essence, this principle demands that one's knowledge be evidentially supported, where evidential support here means support which favours what is believed over known to be incompatible alternatives (that is, which provides more support for what is believed than it does for the known to be incompatible alternatives). So construed,

⁹ For further discussion of the chicken-sexer example, see Foley (1987: 168–9), Lewis (1996), Zagzebski (1996: sects. 2.1 and 4.1), and Brandom (1998).

¹⁰ This assumption is closely related to the 'new evil genius' intuition that is often put forward in the epistemological literature. Very roughly, this is the claim that the reflective accessible grounds possessed by an agent and her envatted counterpart are identical (or, at least, that the reflectively accessible grounds possessed by the agent can be no better than those possessed by her envatted counterpart). For the classical discussion of this intuition, see Lehrer and Cohen (1983).

¹¹ For a thorough discussion of the underdetermination principle, including some of the different ways in which it can be formulated, see Pritchard (2005c).

the principle seems entirely uncontentious, since it is hard to see how one's evidence could be genuinely supporting evidence if it did not perform this 'favouring' function.

The trouble is, however, that once one feeds sceptical hypotheses into this principle in the context of a classical internalist epistemology, then one immediately generates the sceptical problem. Think, for example, of your belief that you are not a BIV. Clearly, one knows that being a BIV is inconsistent with not being a BIV. With underdetermination in mind, then, it follows that knowing that one is not a BIV requires that one possess reflectively accessible evidence which favours your not being a BIV over being a BIV. But how could there be such evidence on the classical internalist account? After all, it seems that your experiences could be just the same regardless of which of these hypotheses were true. Thus, it follows that you cannot know that you are not a BIV.

By the same token, one cannot know that one has, say, two hands, either. As before, it seems that one's experiences could be just the same regardless of whether one is a BIV who merely seems to have hands or one is in normal circumstances, and so on the classical internalist picture it seems one cannot have any reflectively accessible grounds which would favour the hypothesis that one has hands over the alternative hypothesis that one is a BIV. Accordingly, given that one knows full well that having hands is inconsistent with being a BIV, it follows that one cannot know that one has two hands, and much else besides.¹²

It is unsurprising, then, that there are very few neo-Moorean stances in the literature that are conceived along classically internalist lines.¹³ Classical externalism, on the other hand, is clearly on stronger ground when it comes to motivating a neo-Moorean position. After all, that one lacks adequate reflectively accessible grounds by classical internalist lights for believing that one is not a BIV will not on this view decide the issue of whether or not one can have knowledge of this proposition, since other factors, such as the reliability of the process through which one formed one's belief, can also be relevant. Moreover, on this view one can regard one's evidence as consisting of a lot more than simply that which is reflectively accessible to one. Thus, the path is cleared to allowing one's anti-sceptical knowledge to be evidentially grounded and thus to blocking the underdetermination-based sceptical challenge.¹⁴

¹² One might want to qualify the underdetermination principle in order to accommodate the intuition that not *every* belief needs to be evidentially grounded in order to be an instance of knowledge. Notice, though, that such an amendment to the principle wouldn't necessarily make any real difference to its ability to generate sceptical conclusions. After all, it is surely the case when it comes to *most* of our beliefs in empirical propositions that they are only evidentially grounded to a degree that would support knowledge provided that such evidence is able to play this 'favouring' role. And since the sceptic only needs to call the epistemic status of most of our beliefs in empirical matters into question in order to motivate her sceptical doubt, this weaker construal of the underdetermination principle would, it seems, serve the sceptic's purposes equally well.

¹³ Though see Pryor (2000) and Wright (2004). I critically discuss these views at length elsewhere. See Pritchard (2005d, forthcoming a).

¹⁴ Elsewhere I have developed a neo-Moorean response to scepticism that runs along classical externalist lines in this way. See, for example, Pritchard (2005a: pt. 1) and (2007a). One could also regard Williamson's response to scepticism as falling under this general category. See especially, Williamson (2000).

Taking the externalist route out of the sceptical problem is not without its own difficulties, however, since there clearly is a strong intuitive pull towards epistemic internalism.¹⁵ Accordingly, the worry one might have is that by opting for an externalist version of neo-Mooreanism one fails to stay true to the common-sense credentials of Mooreanism. If such revisionism is unavoidable, then this price might be worth paying to get a way out of the sceptical predicament. My interest here, however, is in whether there is a version of neo-Mooreanism available which is true to our internalist intuitions but which avoids the problems facing classical internalist renderings of the neo-Moorean thesis.

3 MCDOWELLIAN NEO-MOOREANISM

I think that we can find such a position in the work of John McDowell (for example, 1982; 1995), who offers an anti-sceptical thesis which has all the main features of neo-Mooreanism but which does not fit into the classical externalism/internalism contrast and which, as a result, is often either simply misunderstood or just plainly ignored.¹⁶ What is distinctive about it (amongst other things) is that while it is a version of epistemic internalism, in that it demands of a knower that she be in a position to know by reflection alone what the reasons which support her knowledge are (this view would not count the chicken-sexer as a knower, for example), it is not a classical form of internalism since it does not share the further assumption of classical internalism noted above. In particular, McDowell argues that one's reasons can be both empirical and factive—that is, can be reasons for believing an empirical proposition and entail what it is that they are a reason for.

As we noted above, it is tacitly accepted throughout epistemology that what one has reflective access to cannot extend beyond the 'inner' to take in factive empirical reasons in this way. According to McDowell, however, this conventional wisdom is false and leads to a philosophical picture which invites the sceptical challenge. Moreover, the source of the problem with this conventional wisdom, according to McDowell, is the failure to endorse the kind of content externalism that he has in mind.

McDowell's view is best thought of as a form of disjunctivism which is directly cast along epistemic lines. Disjunctivists hold that, contrary to the prevailing consensus on this matter, a veridical perceptual experience and its counterpart phenomenologically indistinguishable non-veridical experience each have a different content. That is, that each experience in such a pairing does not have the same content, even though the two experiences are, *ex hypothesi*, indistinguishable as far as the agent is concerned.

¹⁵ It is a concern of this sort that lies behind 'metaepistemological' challenges to externalist anti-sceptical theories. See, for example, Fumerton (1990) and Stroud (1994).

¹⁶ For some of the (relatively very few) recent discussions of McDowell's response to scepticism, see Brandom (1995), Macarthur (2003), Greco (2004), Wright (2004), and Comesăna (2005). As we will see below, many of these commentators simply misunderstand what McDowell is trying to say.

A pairing of cases that would be relevant for our purposes would be the ordinary (or 'good') case—for example, where one is simply sitting at one's desk and having the associated perceptual experiences—and the counterpart sceptical (or 'bad') case—such as in which one is a BIV being 'fed' one's experiences as if one is sitting at one's desk. On the standard picture, we should treat the experiences in the two cases as having the same content, and thus as being the same experience. The disjunctivist, however, disagrees.

One would naturally expect disjunctivism to have ramifications for one's epistemology, since on most standard views of perceptual knowledge the nature of one's perceptual experience forms part of one's perceptual evidence, and thus if one's experiences are different in the two cases just considered, then so too must be one's evidence. Rather than consider the general issue of whether disjunctivism about perceptual experience has this consequence for perceptual evidence, I want to focus instead on the specific McDowellian version of disjunctivism which directly expresses the disjunctivist point in epistemic terms. McDowell's claim is that the subject's reasons are different in the two cases, even though the two cases are phenomenologically indistinguishable to the subject. I will refer to this formulation of the thesis as epistemological disjunctivism to highlight the fact that the view is explicitly and directly advanced as an epistemological thesis, rather than being a thesis primarily about perceptual experience from which epistemological conclusions of the sort that McDowell has in mind may or may not be drawn.

In order to see McDowell's epistemological disjunctivism in action, consider the following case—my believing, on the basis of (what I take to be) my perceptual experience, that John is in my office. In the good case, in which I am not deceived and so form my belief that John is in my office on the basis of my seeing that this is the case, for McDowell my reason for believing that John is in my office will ordinarily simply be that I see that this is the case, where seeing-that is factive inasmuch as one can only see that something is the case if it really is the case. I may have other reasons in favour of my belief, of course, but the point is that this is the 'master' reason in this regard since as a factive reason it is able to suffice, in conjunction with the relevant belief formed on this basis, for perceptual knowledge.¹⁷

In contrast, consider the bad case in which I have what I would take to be the same experiences as in the good case, but in which I am in fact a BIV. In this case the factive reason would clearly be unavailable to me since I cannot see that John is in my office if I am a BIV. Instead, the supporting reason would be of a different form, such as that it seems to me as if John is in my office. Thus, even though the two cases

¹⁷ Presumably, it is not guaranteed on the McDowellian picture that I possess knowledge in the good case, since while there is in fact no deception taking place in the good case, it may be that there are (misleading) defeaters present which I should take into account and which undermine the possibility of possessing knowledge (perhaps I have been falsely told, for example, that I have been drugged, and I have no reason to doubt this). The good cases that we will consider, however, will be those paradigm cases in which there are no defeaters present, and thus where the agent is guaranteed to have knowledge, at least provided she forms her belief in the target proposition on the basis of the relevant reason(s). What is important for the McDowellian position, I take it, is only that there is a wide class of good cases of this sort, not that *every* good case is of this sort.

are phenomenologically indistinguishable, on this view the nature of the supporting reasons for my beliefs in the two cases are different. McDowell's radical thesis is that even despite this, it remains that one's reasons are reflectively accessible to one, such that in the good case one's reflectively accessible reason in favour of one's belief is the *factive empirical reason* that one sees that John is in one's office.¹⁸

This view clearly has fairly radical ramifications for the sceptical problem. On this picture, since, in the good case, it is possible to have perceptual knowledge on the basis of one's *factive perceptual reasons*, it follows that one's reasons in favour of one's target everyday perceptual belief can favour the proposition believed over sceptical alternatives. That is, one's reflectively accessible grounds in favour of one's everyday perceptual belief in the good case can be such that they entail that one is not in the sceptical scenario, and if this is right then they surely thereby favour one's belief in the everyday proposition over the sceptical alternative, thereby blocking one of the key moves in an underdetermination-based scepticism. Furthermore, in cases in which one has knowledge of the everyday perceptual proposition on this basis, and one knows that the proposition believed entails the falsity of a specific sceptical hypothesis, there ought to be no problem with also allowing that one has knowledge of the denial of this sceptical hypothesis, thereby enabling the McDowellian view to deal with the closure-based sceptical argument as well.

We saw above that there is, in principle at least, a classical epistemic externalist account of knowledge available which allows that the epistemic support one has for one's perceptual beliefs could favour the proposition believed over sceptical

¹⁸ I take it that it is essential to the McDowellian view that one lacks perceptual knowledge in the bad case, so described, even if, as it happens, the belief in question is true—perhaps one is a freshly envatted BIV and John is presently in one's office. If this is the right way to read McDowell (he isn't altogether clear on this score), then he has a straightforward response to perceptual Gettier-style cases. Consider, for example, Chisholm's (1977: 105) 'sheep' example in which an agent is, or so he thinks, looking at a sheep in the next field and on this basis forms the true belief that there is a sheep in the next field. Unbeknownst to him, however, the 'sheep' he is looking at is in fact a large furry dog that looks like a sheep, but this doesn't affect the truth of his belief since there is a sheep in the field which is hidden from view behind a hedge. The agent thus has a true belief which is supported by good reflectively accessible grounds but which, intuitively, is not knowledge. Such cases present the challenge of explaining why perceptual knowledge is lacking in these cases even though—again, intuitively—it would be regarded as present in counterpart cases, such as a case in which the agent in the example just described is actually looking at a sheep in the field.

The McDowellian view, as I am here describing it, has an easy answer to this question, which is that, contrary to first appearances, the agent's perceptual reasons, and thus his reflectively accessible grounds, are different in the two cases, which is why the agent can have knowledge in the one case and not the other. In the good case the agent has perceptual knowledge because he is forming his true belief on the basis of a reflectively accessible *factive empirical reason* (in the example just described, this is his seeing that there is a sheep in the field). In contrast, in the corresponding Gettier-style case, the true belief is formed on the basis merely of it seeming to the agent as if there is a sheep in the field, and reasons of this sort do not by themselves suffice for perceptual knowledge. There is thus a straightforward explanation of why knowledge is lacking in perceptual Gettier-style cases, even though it is present in the non-Gettiered counterpart case. Moreover, note that the McDowellian account—in virtue of allowing *factive perception-based reasons*—offers this response to the perceptual Gettier problem without in the process eschewing the tripartite conception of perceptual knowledge as, essentially, justified true belief, even though Gettier-style cases are usually taken to show that such a view is unsustainable.

alternatives, thereby offering a response to both underdetermination and closure-based sceptical arguments. This was achieved, however, by allowing that such epistemic support was not fully reflectively available to the subject. In contrast, what is distinctive about the McDowellian stance is that the epistemic support in question—the agent’s perception-based reasons—is of its nature reflectively accessible to the agent. Moreover, since McDowell allows that one’s perception-based reflectively accessible reasons can be factive, he doesn’t face the underdetermination problem we posed for classical internalist neo-Moorean stances above. Accordingly, if the position can be adequately motivated, then it offers a way of dealing with the sceptical problem while staying broadly within the internalist epistemological model.

4 MOTIVATING MCDOWELLIAN NEO-MOOREANISM

McDowell offers three main considerations in favour of his view. The first is that it enables us to deal with certain perennial epistemological problems, such as the problem of scepticism. We will consider this claim in more detail below. The second is the contention that epistemological disjunctivism accords with a common-sense way of thinking, and talking, about perceptual reasons. That is, in response to a challenge to a claim to know I might well respond by citing a factive reason in defence of my claim, which suggests that we do, ordinarily at least, allow factive reasons to be our rational support for our knowledge.

For example, suppose I tell my boss over the phone that my colleague is in work today (thereby representing myself as knowing this to be the case), and she expresses scepticism about this (perhaps because she falsely believes that my colleague always skips work when she isn’t there). In response I might naturally say that I know that she’s in work today because I can see that she’s in work—she’s standing right in front of me.

The naturalness of this conversational exchange implies that it is McDowell who is working with the common-sense position, and his detractor who is offering the revisionistic view. Of course, this fact in itself may mean very little when set against the grander theoretical scheme, since sometimes our common-sense convictions are plainly wrong, and it is in any case notoriously hard to read our common-sense convictions off everyday language use in this way. Perhaps, for example, we only cite factive reasons in the way just described because we are implicitly bracketing the kind of error-possibilities which could count against the reason being factive. If that’s right, then the citing of a factive reason has a kind of elliptical form—that is, given that certain shared presuppositions of the conversational context hold, then one’s reason for believing *p* in these cases is that one sees that *p*. This description of events would, however, provide little support for McDowell’s position.¹⁹

These complications aside, the importance of this appeal to common sense lies, I take it, in the fact that if one is able to dress one’s position up in the garments of

¹⁹ I consider this possibility in detail in Pritchard (2003).

common sense, then this puts one's view into a default position relative to opposing views which revise our common-sense convictions. Such a default position is highly defeasible, of course, in large part because of the considerations just raised, but it is important to a thesis like McDowell's since on the face of it, it is this view which represents the radical revisionary option in the light of the opposing prevailing conventional wisdom. If McDowell is right that we often speak in a way that would support his account, then this goes a long way towards levelling the dialectical playing field in this regard.

The third consideration McDowell offers in favour of his epistemological disjunctivism is a negative thesis to the effect that it is far from clear what prevents us from advancing a position of this sort. That is, proponents of the standard view about perception-based reasons seem to take it as given—and hence as not requiring further argument—that one's reasons must be the same in the good case and the counterpart sceptical case. The McDowellian claim, however, is that such a thesis is not obvious at all, and that once the argument is presented it is clear that it is problematic.

We can express the argument for the standard picture of perception-based reasons in terms of the following 'highest common factor' argument:

The Highest Common Factor Argument

- P1. In the 'bad' case, the supporting reasons for one's perceptual beliefs can only consist of the way the world appears to one. (Premise)
- P2. The 'good' and 'bad' cases are phenomenologically indistinguishable. (Premise)
- C1. So, the supporting reasons for one's perceptual beliefs in the 'good' case can be no better than in the 'bad' case. (From (P2))
- C2. So, the supporting reasons for one's perceptual beliefs can only consist of the way the world appears to one. (From (P1), (C1))

The idea behind the argument is that one's perceptual reasons are what is common to the two indistinguishable cases—that is, the highest common factor. The first two premises are granted by McDowell, and it seems fairly clear that the ultimate conclusion, (C2), follows from the first premise and the penultimate conclusion, (C1). The joker in the pack is thus the move from the second premise to the penultimate conclusion.

Classical epistemic externalism certainly has grounds to resist a move of this sort, at least where one's 'reasons' are understood broadly so that they do not presuppose an internalist construal—where they are simply understood as one's perceptual evidence, say. After all, there seems no obvious principled reason why an epistemic externalist should regard an agent's perceptual evidence as being limited in terms of what is phenomenologically distinguishable to the agent. The externalist could hold, for example, that the nature of one's perceptual evidence is in part determined by facts concerning the pedigree of the perceptual mechanism through which that evidence is gained—its reliability in the relevant environment, for example—facts which, on the classical externalist account, are not reflectively available to the subject. On a classical externalist picture of evidence of this sort, the nature of the perceptual evidence

in each case would be different because of the different pedigree of the perceptual mechanisms concerned, even though the two cases are phenomenologically indistinguishable to the agent.

Indeed, a form of disjunctivism about perceptual epistemic support of the general type envisaged by McDowell is available to the classical epistemic externalist, since she could claim that whether that epistemic support is of a factive sort can be dependent upon further environmental facts obtaining which are not reflectively accessible to the subject. On this view, then, either one is in the good case and the epistemic support for one's perceptual belief is factive, or one is in the bad case and the epistemic support for one's perceptual belief is non-factive, merely phenomenal, say (for example, that one seems to see that such-and-such is the case). With such a disjunctivism in play, a way of blocking underdetermination-based scepticism becomes available, since one can claim that in the good case one has adequate evidential support for one's perceptual beliefs which, qua factive evidence, is able to favour the believed hypothesis over sceptical alternative hypotheses that are known to be incompatible, even though the two scenarios are phenomenologically indistinguishable.²⁰

McDowell's idea, however, is that one can help oneself to this sort of picture without thereby conceding that one's epistemic support, now construed again as one's reasons, needs to be understood such that it is reflectively unavailable to the subject. In effect, McDowell wants to claim that we can endorse both the (externalist) claim that one's perception-based reasons in the good case can be factive and the (internalist) claim that one's reasons must be reflectively accessible. McDowell thinks that even someone who endorses the thesis that one's reasons must be reflectively accessible can object to the move from (P2) to (C1). After all, on the picture of reflectively accessible reasons that he outlines, there is a difference in the agent's reasons in the two cases, but the agent's reasons are held to be no less reflectively accessible on account of this. The proponent of the highest common factor argument thus owes us an explanation of what is wrong with such a view before she can help herself to the key move from (P2) to (C1).

Primarily, then, McDowell's contention is that the argument for the highest-common-factor account of perceptual reasons is underdeveloped, thus leaving open the possibility of an epistemological disjunctivism which he exploits. Still, I think many detractors of the McDowellian account are dissatisfied by a negative argument of this sort for the simple reason that they think there is a straightforward problem with epistemological disjunctivism, and thus that the highest-common-factor conception of perceptual reasons is the only game in town.

We can draw this problem out via the following putative *reductio* of the McDowellian view:

²⁰ Speaking loosely, one could describe Williamson's (2000) response to scepticism as being a classical externalist view of this sort. According to Williamson, one's evidence is identified with one's knowledge, and thus evidence is factive. Furthermore, Williamson also claims that the perceptual evidence one's experiences give rise to is not the same in the good and the bad cases. In the good case, one's experiences will give rise to factive perceptual evidence in favour of one's belief, even though the subjectively indistinguishable experiences in the bad case will not generate the same factive perceptual evidence.

Putative Reductio of Epistemological Disjunctivism

- P1. In the 'bad' case, the supporting reasons for one's perceptual beliefs can only consist of the way the world appears to one. (Premise)
- P2. The 'good' and 'bad' cases are phenomenologically indistinguishable. (Premise)
- P3. One's reasons are reflectively available to one. (Premise)
- A1. The supporting reasons for one's perceptual beliefs in the good case can be factive. (Assumption for *reductio*)
- C1*. So, one is able to know by reflection alone that one is in possession of a factive perceptual reason in the good case. (From (P3), (A1))
- C2*. So, for all one can know by reflection alone, one could just as well be in the bad case than in the good case. (From (P2))
- C3*. So, one is unable to know by reflection alone that one is in possession of a factive perceptual reason in the good case. (From (C2*))

If this argument is allowed to stand then, since (C3*) contradicts (C1*), it will follow that the assumption, and with it epistemological disjunctivism, will have to be denied. We have already seen that the first two premises are accepted by McDowell, and he will also accept the third premise too. In addition, the move from (P3) and (A1) to (C1*) ought to be uncontroversial by McDowell's lights. The move to (C2*), however, is certainly problematic for a McDowellian. After all, according to McDowell, there is a difference between what is reflectively available to the agent in the bad case and what is reflectively available in the good case, since it is only in the good case that the agent has the relevant factive perception-based reason available to her. By McDowell's lights, therefore, the inference to (C2*) simply begs the question, meaning that (C2*) and, with it, (C3*), will not be allowed to stand.

This point highlights where the source of the opposition lies when it comes to epistemological disjunctivism and the highest-common-factor conception of perceptual experience. Proponents of the latter view claim that since there is no phenomenological difference between the good case and the bad case, so there can be nothing reflectively available to the subject that could indicate that she is in the good case. According to the epistemological disjunctivist, however, this is not so, since if one is in the good case then it is reflectively accessible to one that one is in possession of a factive perceptual reason, and one could only be in possession of such a reason provided that one is in the good case. Thus, there is something reflectively available to one that entails that one is in the good case even though there is no phenomenological difference between the good case and the bad case.

The proponent of the highest-common-factor conception of perceptual reasons will no doubt find this claim puzzling, and part of the reason for this puzzlement is, I think, that on the face of it this claim seems to entail that one can know, by reflection alone, that one is in the good case (and thus not in the bad case), and this does not seem at all plausible. I think we can show, however, that this is not a consequence of the view, and in so doing provide a crucial plank of support for McDowell's position.

5 MCDOWELL'S 'MCKINSEY' PROBLEM

In order to do this, I first want to consider a second objection that one might level against McDowell's view which, as we will see, is closely related. The objection that I have in mind is one that mirrors the famous 'McKinsey' puzzle that concerns the supposed compatibility of first-person authority and content externalism.²¹ For it seems that on the McDowellian view one can use one's reflective access to one's factive empirical reasons, along with one's reflective knowledge of the relevant entailment (that is, from the factive empirical reason to the empirical fact), to acquire reflective, and thus non-empirical, knowledge of the empirical proposition which the reason is a reason for (for example, that John is in one's office). We can express this argument roughly as follows:

McDowell's Putative 'McKinsey' Problem

- P1'. In the good case, one is able to know by reflection alone that one is in possession of a factive empirical reason. (Premise)
- P2'. One is able to know by reflection alone that if one is in possession of a factive empirical reason for believing the empirical proposition *p*, then *p*. (Premise)
- C1'. So, in the good case, one is able to know by reflection alone the empirical proposition *p*. (From (P1'), (P2'))

Intuitively, this is just a *reductio* of the view, since one ought not to be able to come to know specific empirical truths by reflection alone.²² Indeed, this *reductio* in part explains, I think, why very few commentators have taken McDowell's position seriously and have stuck instead to the conventional wisdom on reflective access to reasons that is implicit in the classical way of drawing the externalism/internalism contrast in epistemology.²³

It is not obvious, however, that McDowell's position is subject to a problem of this sort. The reason for this is that the difficulty only emerges if one acquires non-empirical knowledge by running through an inference of this type, and it is far from clear why, on the McDowellian view, this would be so. In particular, the McDowellian conception of perceptual knowledge in these cases is such that one could only

²¹ For more on this problem, see Nuccetelli (2003).

²² I say 'specific' empirical truths since it is not nearly so obvious that it would be counterintuitive to suppose that one could know *general* empirical truths by reflection alone. Indeed, given the historical popularity of transcendental arguments, it would be odd if it were counterintuitive to suppose that one could know general empirical truths by non-empirical means. To take a contemporary example, consider Davidson's (1986: 314) claim that "belief is in its nature veridical". If this is true a priori, and one can know non-empirically that one has beliefs, then it follows that one can know non-empirically that most of one's beliefs are true, a proposition which is clearly empirical. Although some may find such a conclusion objectionable, it would hardly stand by itself as a *reductio* of the view, and most who do object to this conclusion would surely grant that it is not obviously counterintuitive.

²³ For three recent commentators who have not taken McDowell's non-classical internalist view at face-value and have instead either read him as a classical externalist or classical internalist, see Brandom (1995), Wright (2002), and Greco (2004).

be in a position to run through such an inference provided one already has the empirical knowledge in question, and thus the problem concerning the non-empirical acquisition of empirical knowledge does not arise.

In order to see this, one only needs to recall that for McDowell a factive reason for believing p , in conjunction with a belief that p which is based on that factive reason, will suffice for knowledge that p . Note, however, that reflective accessibility is itself factive, in that if one is in a position to access reflectively that one has a factive reason for believing p , then it must be the case that one has that factive reason for believing p . Furthermore, even though the possession of the factive reason is consistent with non-belief in the target proposition on the McDowellian view,²⁴ it is obviously not going to be possible for someone to run through the reasoning described above without in the process acquiring the reason-based belief in the target proposition, thereby meeting all the conditions required on this view for empirical knowledge in this case. Thus, given the further trivial claim that if one has empirical knowledge of p then one cannot also have reflective (that is, non-empirical) knowledge of p , it follows that there is no prospect of acquiring non-empirical knowledge in this case, and thus McDowell's McKinsey-style difficulty disappears. That is, the two premises of the argument outlined above do not entail the conclusion, since in the circumstances envisaged by the premises one has empirical knowledge of the target proposition which undermines the possibility of acquiring the non-empirical knowledge at issue in the conclusion.²⁵

In effect, the proponent of the *reductio* has failed to notice that while knowledge is, plausibly, closed under known entailment, it does not follow that specifically non-empirical knowledge is closed under (non-empirically) known entailment, such that where one knows the antecedent proposition non-empirically one thereby knows the conclusion non-empirically. As we have just seen, there are counter-examples to this more demanding principle which concern cases in which the nature of the inference in question ensures that the agent's knowledge of the conclusion is empirical rather than non-empirical.²⁶

This point is important for the problem raised in the last section, since the worry expressed there was that if what one is able to know by reflection alone is different in the good case as opposed to the bad case—and one also knows (non-empirically) the import of these differences—then why doesn't it follow that one can know by reflection alone the empirical fact that one is in the good case, and thus the further

²⁴ See, for example, the exchange between Stroud (2002) and McDowell (2002).

²⁵ For more on this point, see Neta and Pritchard (forthcoming).

²⁶ I discuss the differences between these two principles in more length in Pritchard (2002a). I think this stronger closure principle is implausible whatever one's wider epistemological stance, but it is worth noting that from a McDowellian point of view it is particularly contentious because of the way in which it in effect insulates non-empirical knowledge from empirical knowledge, when a key part of the McDowellian programme—see especially McDowell (1986)—is to undermine the sort of (Cartesian) picture that (McDowell thinks) leads to such an insulation. So while I've tried to put my argument in such a form that it does not already presuppose commitment to the general McDowellian picture, were one to set it within that frame then one could read it as highlighting an important *consequence* of a certain theme that runs through McDowell's thought. I am grateful to Adrian Haddock for pressing me on this issue.

empirical fact that one is not in the bad case? As in the 'McKinsey' argument just noted, one is, it seems, gaining non-empirical knowledge of a very specific empirical truth.

Closer inspection of the position reveals that this problem doesn't arise, however, and the forgoing remarks on McDowell's apparent 'McKinsey' problem indicate why. I take it that the argument that is implicit here proceeds in three stages as follows. First, one is able to know by reflection alone that one is in possession of a factive empirical reason for believing *p*. Second, given that one also knows non-empirically that one could only have the factive empirical reason for belief in *p* if one is in the good case, then one can conclude, on entirely non-empirical grounds, that one is in the good case. Finally, given that one knows non-empirically that the good case is inconsistent with the bad case, one can also conclude, again on entirely non-empirical grounds, that one is not in the bad case. Thus, even though there is, *ex hypothesi*, no phenomenological difference between the good case and the bad case, it seems nonetheless possible on the McDowellian view that one can know by reflection alone that one is in the good case as opposed to the bad case.

We can express this argument as follows:

McDowell's Other Putative 'McKinsey' Problem

- P1'. In the good case, one is able to know by reflection alone that one is in possession of a factive empirical reason. (Premise)
- P2''. One is able to know by reflection alone that if one is in possession of a factive empirical reason for believing the empirical proposition *p*, then one is in the good case. (Premise)
- P3''. One is able to know by reflection alone that if one is in the good case, then one is not in the bad case. (Premise)
- C1''. So, in the good case, one is able to know by reflection alone that one is in the good case. (From (P1'), (P2''))
- C2''. So, in the good case, one is able to know by reflection alone that one is not in the bad case. (From (C1''), (P3''))

Given our earlier remarks on McDowell's apparent 'McKinsey' problem, however, the mistake in this reasoning should be obvious. In drawing the penultimate conclusion, (C1''), it is essential that one has the required belief that one is in the good case (if one did not already have it), formed on the basis of the rational support one possesses in favour of this claim. Furthermore, the recognition that one's factive empirical reason for believing the target proposition entails that one is in the good case—(P2'')—highlights that one is in possession of a factive empirical reason for believing that one is in the good case. Accordingly, it follows that one can no more come to know non-empirically via this chain of reasoning that one is in the good case than one can come to know the target proposition non-empirically via the parallel McKinsey-style reasoning, since in running through the reasoning one thereby acquires (if one did not already possess it) empirical knowledge that one is in the good case, thereby excluding the possibility of one possessing non-empirical knowledge for this claim.

That is, the first two premises do not entail (C1'') since the circumstances described in these premises ensure that the agent has empirical knowledge of the fact that she is in the good case, thereby excluding the possibility of non-empirical knowledge at issue in (C1''). Moreover, as before, the error here is to confuse the plausible principle that knowledge is closed under knowledge entailment with the much more demanding—and, as we have seen, false—principle that non-empirical knowledge is closed under (non-empirically) known entailment, a principle which fails in these cases. Given that the penultimate conclusion does not follow, however, then neither does the ultimate one, (C2''). One cannot, on the McDowellian picture, acquire non-empirical knowledge that one is in the good case as opposed to the bad case, and so the worry stated earlier does not emerge. This claim, however, is consistent with the McDowellian thesis that what is reflectively available to one in the good case is not the same as what is reflectively available to one in the bad case. One of the key supposed problems for epistemological disjunctivism—and thus one of the core motivations for the highest-common-factor theory—is thus shown to be illusory.

I think this result should not surprise us, since on the McDowellian view gaining reflective knowledge of one's *factive* empirical reasons is not a purely introspective affair, in contrast to the classical internalist view. Instead, recognizing that one is in possession of a *factive* empirical reason, such as that one sees that such-and-such is the case, is itself a partly world-directed activity since these are, as it were, world-involving states as opposed to being the kind of purely mental states envisaged by the classical internalist. It is little wonder then that there are limits to the kind of non-empirical knowledge that one can gain by reflecting on one's *factive* empirical reasons, since it is inevitable that the drawing of some consequences from one's reflective knowledge in this regard will result (if it is not possessed already) in empirical knowledge of the target proposition being possessed.

As we will see, this general point about the scope of one's non-empirical knowledge by McDowellian internalist lights is important to understanding how McDowell might respond to a certain sort of problem that could be raised for his neo-Moorean response to scepticism. Before we can get into that issue, however, we need to look a little closer at what the McDowellian response to scepticism is.

6 MCDOWELL ON SCEPTICISM

I mentioned above that one of the considerations that McDowell offers in favour of his epistemological disjunctivism is that it enables us to resolve certain perennial epistemological problems, such as the problem of scepticism. Interestingly, though, McDowell himself does not seem to be at all inclined to advance his anti-scepticism in a direct way at all. That is, he does not conceive of the position as issuing in an anti-sceptical argument, such as a neo-Moorean argument, which opposes, and thereby neutralizes, a specific sceptical argument, such as the argument we outlined in section 1. Indeed, he is uncomfortable with the very idea of calling his treatment of scepticism an "answer" to the problem at all, remarking, for example, that his view is "not well cast as an answer to skeptical challenges; it is more like a justification of

a refusal to bother with them” (McDowell 1995: 888). This coyness is odd, since if McDowell’s view is correct, then one can directly answer the sceptical challenge, and it is puzzling why McDowell doesn’t recognize this.

Part of the reason why McDowell expresses his anti-scepticism in this cautious fashion is, I think, because he regards such caution as being mandated by his avowedly ‘diagnostic’ treatment of the sceptical problem, where the diagnosis is that scepticism of the sort that he is interested in “expresses an inability to make sense of the idea of direct perceptual access to objective facts about the environment” (McDowell, this volume: section 2). Indeed, he makes this contrast between answering the sceptical problem and merely diagnosing it explicit in the following passage: “the thing to do is not to answer the skeptic’s challenges, but to diagnose their seeming urgency as deriving from a misguided interiorization of reason” (McDowell 1995: 890). Still, it is far from obvious why a diagnosis of scepticism could not also be (and ought to be) an answer to the sceptical problem.

In order to get a better grip on why McDowell makes this distinction, I think it is important to recall that part of the motivation that he offers for his view is that his disjunctive account of perceptual reasons is the common-sense conception of reasons to which we have been blinded by epistemological theory. The importance of this is that with the debate so construed McDowell can claim that the sceptical problem is the result of abstract philosophizing, rather than the natural intuitive puzzle that it appears to be. Since the visceral pull of scepticism lies in its apparent appeal only to intuition, this exposure of the sceptic as, in effect, a covert philosopher ensures that we can ignore the puzzle with impunity. Indeed, this would be a very Moorean claim to make. If intuition, properly understood, licenses anti-scepticism over scepticism, then there is no need to engage with the sceptical problem further once one has highlighted this fact.²⁷

This, in any case, may well be the official story as to why McDowell’s expresses his anti-scepticism so cautiously. I want to suggest, however, that this official story disguises a tacit motivation that McDowell has for putting his view forward in this way. This is that I think McDowell conceives of an answer to the sceptic as consisting of a series of Moorean-style assertions, and he recognizes—as we did above in section 1—that such assertions are clearly problematic. That is, an answer to the sceptic will, it seems, consist in claiming, for example, that one knows the denials of sceptical hypotheses, and one knows them because one knows everyday propositions and the relevant entailment. With the awkwardness of such assertions in mind, it is little wonder that McDowell prefers simply to remove the urgency from sceptical arguments and then refuse to engage with them thereafter, since this strategy avoids the need to make Moorean assertions of this sort.

In order to see this point in more detail, it is worthwhile considering how classical externalist and internalist neo-Moorean views would go about accounting for the impropriety of Moorean assertions. To begin with, note that, ordinarily at least, one needs good reflectively accessible grounds to back up an assertion if one is to properly

²⁷ This way of reading McDowell makes his strategy for dealing with the sceptical problem similar in key respects to that proposed by Williams (1991).

make that assertion, especially if the assertion involves a claim to know. This means that the conditions under which it is appropriate to claim knowledge and the conditions under which knowledge is possessed will be closely connected on the classical internalist account, but not so closely connected on the externalist account. This is another reason why classical internalists will have great difficulty motivating a form of neo-Mooreanism, since they will find it hard to explain why the Moorean assertions are inappropriate if they are true since, if they are true, they will also be assertions which are appropriately supported by adequate reflectively accessible grounds and so, presumably, ordinarily properly assertible.

On the classical externalist account which allows that knowledge might be possessed even in the absence of good reflectively accessible grounds, in contrast, there will be a class of cases in which agents have knowledge but cannot properly claim that knowledge because they lack adequate reflectively accessible supporting grounds for their claim. Think, for example, of the chicken-sexer case described above. Here we have an agent who has knowledge by classical externalist lights, but who is clearly not in a position properly to assert what she knows—much less assert that she knows it—since she lacks any good reflectively accessible grounds to back up that assertion. If the classical externalist is right in thinking that we can know the denials of sceptical hypotheses, even though we lack good reflectively accessible grounds in support of our beliefs in this regard, then one's anti-sceptical knowledge will be akin to the brute knowledge possessed by chicken-sexers. It is little wonder, then, that Moore cannot properly make the assertions he does, even when what he asserts is true.

The problem facing McDowell is akin to that facing the classical internalist, in that one would naturally think that if one really does possess the knowledge in question, and one can moreover offer conclusive reflectively accessible grounds in favour of that knowledge, then there ought not to be any bar to claiming that one has that knowledge. On the face of it, then, there ought to be nothing at all amiss about Moore's assertions by McDowellian lights, even though they clearly are problematic. It is thus not surprising that McDowell wishes to distance his anti-scepticism from views which try to answer the sceptical challenge, and which thereby find themselves making Moorean assertions.

Still, I think that McDowell can be bolder in his treatment of scepticism, since, as we will see, there is an intermediate form of anti-scepticism available here, one which answers the sceptical argument but which does not do so by making Moorean assertions. What McDowell's view requires, I suggest, is an account of why such Moorean assertions are inappropriate even though true, an account which is in the spirit of McDowell's general epistemological disjunctivism. In the next section I aim to offer just such a thesis.

7 KNOWING AND SAYING THAT ONE KNOWS

To begin with, it is important to note that we rarely convey our knowledge by making assertions which are prefixed with the phrase, 'I know'. Instead, one typically conveys one's knowledge of a proposition simply by asserting the proposition in question.

Adding the further 'I know' phrase is rare, and standardly reflects not just emphasis but also an ability to resolve a particular challenge that has been raised.

For example, one might initially convey one's knowledge of what the time is by simply asserting, say, 'It's 10.22am', but then one might be prompted into the further explicit claim to know this proposition by a challenge to one's original assertion. The usual way to challenge such an assertion is to raise an error-possibility which is held to be salient, such as when someone doubts that this is the time points out that one's watch looks very old and so could well be unreliable. In responding to a challenge of this sort with an explicit knowledge claim one is representing oneself as being in possession of stronger reflectively accessible grounds in support of one's assertion than would be implied simply by making the assertion itself.

Notice, however, that one is also usually representing oneself as being in possession of reflective grounds which speak specifically to the error-possibilities raised. In particular—and this is a point which, I think, has often been overlooked in this regard, despite its epistemic importance—the grounds one needs available to one in making a claim to know in response to this sort of challenge must be such as to discriminate between what is asserted and the relevant error-possibility. This claim is important because, intuitively, the kind of evidential support one needs in order to have knowledge is weaker than this. All that is required here is the sort of 'favouring' evidence that we saw above in our discussion of the underdetermination principle (which is still a strong requirement on knowledge, as we also saw).²⁸

In order to make this point clear, consider the famous 'zebra' case offered by Fred Dretske (1970). Here we are asked to imagine someone who is at the zoo in normal circumstances and sees what looks like a zebra in the zebra enclosure. On this basis, she forms the belief that there is a zebra before her. Clearly, such an agent would normally be attributed knowledge in this case, and we would be perfectly happy with any assertion she might make to the effect that there is a zebra before her—which would represent her as having knowledge of this proposition—since she has adequate reflectively accessible grounds to back up that assertion, such as that she sees that there is a zebra before her.

Similarly, an explicit claim to know that she sees a zebra would also be deemed appropriate in this context, if the circumstances were right. For example, if the original claim that the creature before her is a zebra is challenged in some mundane fashion—perhaps by someone short-sighted who wonders out loud why, since she was expecting to be near the gorilla enclosure, there should be zebras here—then it would be unproblematic for our agent to respond to this challenge by saying that she knows that this creature is a zebra.

It is important to recognize why such an explicit claim to know would be entirely appropriate, given how we have described the situation. The reason for this is not just that the agent is in a position to offer very good (indeed, conclusive) reflectively accessible grounds in favour of the proposition known. Nor is it that she has good reflectively accessible grounds in favour of the proposition known which prefer this proposition over the target error-possibility (that it is a gorilla rather than a zebra),

²⁸ For more on the relationship between knowledge and discrimination, see Pritchard (2006).

which is what her conclusive factive empirical reason secures. This is because the mere fact that one's reflectively accessible supporting reason is that one sees that there is a zebra before one does not speak specifically to the issue of discrimination in the light of the target error-possibility.

It is possible for example, though unlikely in this case, that our agent has the factive empirical reason in question but lacks further good reflectively accessible grounds which would enable her to discriminate between the target proposition and the target error-possibility—that is, between creatures that are zebras and (non-zebra) creatures that are gorillas. Such grounds might be, for example, that zebras and gorillas have very different shapes and gaits. If the agent lacked discriminating grounds of this sort, then the assertion would be deemed inappropriate even though the agent has the supporting factive empirical reason. This is because explicitly claiming knowledge in the light of this challenge will normally generate the conversational implicature that one is able to offer grounds in support of the proposition claimed as known which would suffice to distinguish the scenario described by this proposition from the specified error-possibility. If grounds of this very specific sort are lacking, however, then this conversational implicature is false and the assertion is thereby inappropriate, even though true, and even though it is supported by conclusive and thereby favouring grounds. In this case, however, the agent does have the required grounds.

Here, then, is the crux of the matter: whereas reflectively accessible factive reasons in support of one's belief suffice for knowledge possession, such reasons are not sufficient by themselves to ensure that a corresponding explicit claim to know is appropriate. Recognizing this fact is essential, I claim, to understanding how the McDowellian can regard perceptual knowledge in certain cases as both conclusively supported by reflectively accessible grounds on the one hand, and yet not properly assertible on the other.

As I noted a moment ago, it would be very odd in the case just described to suppose that the agent concerned has the factive supporting reason without in addition possessing the required discriminating reasons. Given the straightforward nature of the discrimination in question, most people in the circumstances described would possess the required discriminating reasons. Nevertheless, there are cases where one is in possession of the factive supporting reason but where one lacks the relevant discriminative grounds, and some of these cases concern scenarios in which one is considering sceptical hypotheses. Accordingly, if one fails to pay due attention to this point then one will be led into denying knowledge to the agent even though on the McDowellian neo-Moorean view such knowledge is possessed.

Imagine, for example, that the error-possibility that the zebra may in fact be a cleverly disguised mule is raised and taken seriously in that conversational context. Since the original assertion that the creature is a zebra has been challenged, it would ordinarily be appropriate for the agent, if she knows this proposition, explicitly to say so, as she did in the case just described where the objector wonders why she isn't presently looking at a gorilla. If the agent now claims to know that there is a zebra before her in the light of this error-possibility being raised, however, then this will generate the implicature that the agent has reflectively accessible grounds which could serve to discriminate between the proposition claimed as known and the target error-possibility.

That is, the agent is representing herself as having grounds which would suffice to enable her to tell the difference between these two creatures (that is, a zebra and a cleverly disguised mule). Such grounds might be, for example, that she has examined the creature at close range and been able to determine that it is not painted. Typically, of course, the agent will not have grounds of this sort available to her, and so her claim to know will be inappropriate because it generates a false conversational implicature.

Notice, however, that this fact alone does not suffice to indicate that the agent lacks knowledge of the target proposition. After all, in the standard case at least, the agent will have grounds which favour the hypothesis that the animal before her is a zebra over the alternative hypothesis that it is a cleverly disguised mule. After all, part of the agent's grounds on the McDowellian view will consist in her seeing that there is a zebra before her, and this ground excludes the possibility that there is a cleverly disguised mule on the scene instead. Moreover, the agent also has further relevant reflectively accessible grounds that she can offer, such as the grounds she has regarding the implausibility of a zookeeper going to such lengths to deceive patrons, and the penalties that would be imposed were such a deception to come to light, as presumably it would eventually. There is no reason to deny knowledge here, or even to deny that such knowledge is appropriately supported by reflectively accessible grounds. Nevertheless, unless the agent concerned has made special checks or has special discriminative abilities, then to claim explicitly to know such a proposition in this context would be inappropriate.

One might at first blush be puzzled by the contrast that is being drawn here. If one can indeed know by reflection alone that one is in possession of a factive empirical ground—such as, in this case, that there is a zebra before one—then how could it be that one is unable to discriminate the zebra case from the target non-zebra case, such as the situation in which one is faced with a cleverly disguised mule? After all, this factive empirical reason entails that one is not faced with a cleverly disguised mule, and one can know that this entailment holds non-empirically. Isn't one then in a position to know, by reflection alone, that one is now looking at a zebra rather than a cleverly disguised mule? And if one can do that, then what more could be required of an agent when it comes to possessing the relevant discriminatory ability?

Indeed, the sort of reflectively accessible grounds required in order to enable one to discriminate between the two scenarios are typically logically weaker than factive empirical grounds. That one has made a special check that the creature before one is a zebra rather than a cleverly disguised mule—one has examined the creature for signs of paint, for example—does not entail that the creature in question is a zebra. Accordingly, one would expect the relative logical strength of the factive empirical reason to be more likely to suffice for the possession of the relevant discriminatory ability rather than the discriminatory reasons actually demanded (I claim) in such cases.

It is at this juncture that the point we made earlier on becomes important. Recall that we noted that McDowell can consistently hold that one can possess reflectively accessible factive empirical grounds which entail that one is in the good case, and know by reflection alone that this entailment holds, and yet fail to know by reflection alone that one is in the good case (and, thereby, that one is in the good case as opposed to the bad case). The reason why this is so is that the principle being

appealed to here—that non-empirical knowledge is closed under (non-empirically) known entailments—is false, and fails in these cases. As we saw above, in making an inference of this sort one inevitably ends up acquiring empirical knowledge of the fact that one is in the good case, thereby excluding the possibility that one acquires non-empirical knowledge of this proposition via this inference. Accordingly, it does not follow from the fact that one can know by reflection alone that one is in possession of a factive empirical reason which entails that one is in the good case that one can know by reflection alone that one is in the good case (or, similarly, that one is in the good case as opposed to the bad case).

The same applies in the situation we have just considered. Suppose one comes to know by reflection alone that one's reason for believing that there is a zebra before one is that one sees that there is a zebra before one, and one also recognizes that this entails that one is not looking at a cleverly disguised mule. Does it follow that one has non-empirical knowledge of the fact that one is not looking at a cleverly disguised mule, thereby indicating that one has reflectively accessible grounds which suffice to discriminate between the zebra hypothesis and the alternative cleverly disguised mule hypothesis? No. One no more has non-empirical knowledge of this proposition in this case than one has non-empirical knowledge that one sees that there is a zebra before one. The supporting reason in question is, remember, an empirical reason, and that supporting reason in conjunction with the relevant belief formed on that basis—a belief required for the inference to go through—will entail that one has empirical knowledge of the fact that one is not looking at a cleverly disguised mule, thereby excluding the possibility of acquiring non-empirical knowledge of this proposition. On the McDowellian account, then, there is the logical space to distinguish between reflectively accessible factive empirical grounds on the one hand, and reflectively accessible discriminating grounds on the other.

This distinction is absolutely vital when it comes to the sceptical case. What is different about the sceptical case is that while there clearly are agents who could properly make the relevant assertions in zebra-style cases because they have the special grounds required (such as a zoologist, for example), there are no such agents in the sceptical case. No one has adequate reflectively accessible grounds—by everyone's lights (including McDowell's)—which would enable them to discriminate between, say, seeing a zebra and being a BIV on a zebra-less planet who merely seems to see a zebra. There are, for example, no special checks that one can make which would be relevant in this regard.

This is why claiming to know the denials of sceptical hypotheses is conversationally inappropriate, since it generates the conversational implicature that one is in possession of reflectively accessible discriminating grounds which are in principle unavailable. Moreover, it also explains why claiming to know even mundane everyday propositions in sceptical contexts would be inappropriate, such as that one has two hands. After all, in such contexts the target error-possibilities at issue will be sceptical hypotheses, and thus in making this assertion one is falsely representing oneself as possessing adequately reflectively accessible grounds which would suffice to enable one to discriminate between having two hands and being a BIV who merely seemed to have hands.

8 CONCLUDING REMARKS

McDowell is thus able to account for the impropriety of Moorean assertions by appealing to independently motivated claims which are in the spirit of his epistemology. With this in mind, it is open to him to argue that, in the good case at least, what Moore says is entirely true while at the same time conceding that what Moore says cannot be properly said in this way. Moreover, he can do this while retaining a key internalist intuition about the importance of reflectively accessible grounds to knowledge possession. McDowell's quietism in the face of scepticism is thus optional, which means that we have available to us a fully-fledged form of McDowellian neo-Mooreanism.

Moreover, note that McDowell's view is also immune to the two other problems facing Mooreanism that we noted above. Since McDowell has diagnosed the appeal of scepticism as ultimately arising out of a theoretical, rather than a common-sense, claim about reasons, the view does not leave us with merely an impasse with the sceptic, since the default position of common sense is with the anti-sceptical thesis rather than its sceptical counterpart. Relatedly, with the putative intuitive support for the opposing sceptical argument undermined, there is also nothing question-begging about the McDowellian neo-Moorean response to that sceptical argument either.

As we have seen, the apparent problems that face McDowell's non-classical internalism are illusory, and thus the choice between an unsatisfactory classical internalist neo-Mooreanism and a revisionary classical externalist neo-Mooreanism is not forced upon us. Furthermore, we have also seen that provided the view is accompanied by a complementary thesis regarding the conditions under which one can properly claim knowledge, then neither is one committed, in signing up to a McDowellian neo-Mooreanism, to subscribing also to McDowell's famous brand of quietism. The path is thus cleared for a robust form of internalist neo-Mooreanism cast along McDowellian lines.²⁹

²⁹ This essay to a certain degree signals a departure from the view I defended in Pritchard (2003); cf. Pritchard (2005a: section 9.2). The apostasy is not complete, however, since I still stand by many of the main charges against McDowell that I levelled in that paper, and this point deserves comment. In particular, I still think that McDowell's quietism in the face of the sceptical challenge is problematic, and also that even by his own lights the issue of epistemic luck that he raises in the context of the sceptical problem is a red herring. Furthermore, I still believe that one can find in Wittgenstein's *On Certainty* an important challenge to McDowell's conception of factive reasons which he never fully engages with himself. That said, I do now hold, as this essay indicates, that there is a plausible anti-sceptical position about knowledge that can be extracted from McDowell's writings on this topic. The key point is that it needs to be shown that McDowell's position does not generate the counterintuitive consequences that it initially seems to, and that it is not essentially wedded to a form of philosophical quietism. This is what I believe I have shown here, and which has prompted the change of heart. (Though note that saying that McDowell is in a position to respond adequately to scepticism about *knowledge* is not thereby to say that he is able to respond to the sceptical problem more generally. I now think that the key sceptical problem is not directed at knowledge possession at all, but rather at a certain kind of cognitive responsibility. For more on this point, see Pritchard (2005a: ch. 9; 2005b).)

REFERENCES

- Alston, W. (1988) 'An Internalist Externalism', *Synthese*, 74: 265–83.
- Baldwin, T. (1993) 'G. E. Moore', in J. Dancy and E. Sosa (eds.) *A Companion to Epistemology*, Oxford: Blackwell.
- Brandom, R. (1995) 'Knowledge and the Social Articulation of the Space of Reasons', *Philosophy and Phenomenological Research*, 55: 895–908.
- (1998) 'Insights and Blindspots of Reliabilism', *The Monist*, 81: 371–92.
- Chisholm, R. M. (1977) *Theory of Knowledge*, 2nd edn, Englewood Cliffs, NJ: Prentice-Hall.
- Cohen, S. (2000) 'Contextualism and Skepticism', *Philosophical Issues*, 10: 94–107.
- Comesaña, J. (2005) 'Justified versus Warranted Perceptual Belief: A Case against Disjunctivism', *Philosophy and Phenomenological Research*, 367–83.
- Conce, E. and Feldman, R. (2000) 'Internalism Defended', in H. Kornblith (ed.) *Epistemology: Internalism and Externalism*, Oxford: Blackwell.
- Davidson, D. (1986) 'A Coherence Theory of Truth and Knowledge', in E. LePore (ed.) *Truth and Interpretation: Perspectives on the Philosophy of Donald Davidson*, Oxford: Blackwell.
- DeRose, K. (1995) 'Solving the Skeptical Problem', *Philosophical Review*, 104: 1–52.
- Dretske, F. (1970) 'Epistemic Operators', *Journal of Philosophy*, 67: 1007–23.
- Foley, R. (1987) *A Theory of Epistemic Rationality*, Cambridge, MA: Harvard University Press.
- Fumerton, R. (1990) 'Metaepistemology and Skepticism', in M. D. Roth and G. Ross (eds.) *Doubting: Contemporary Perspectives on Skepticism*, Dordrecht: Kluwer.
- Greco, J. (2004) 'Externalism and Skepticism', in R. Shantz (ed.) *The Externalist Challenge: New Studies on Cognition and Intentionality*, New York: de Gruyter.
- Kornblith, H. (2001) (ed.) *Epistemology: Internalism and Externalism*, Oxford: Blackwell.
- Lehrer, K. and Cohen, S. (1983) 'Justification, Truth, and Coherence', *Synthese*, 55: 191–207.
- Lewis, D. (1996) 'Elusive Knowledge', *Australasian Journal of Philosophy*, 74: 549–67.
- Macarthur, D. (2003) 'McDowell, Scepticism, and the "Veil of Perception"', *Australasian Journal of Philosophy*, 81: 175–90.
- McDowell, J. (1982) 'Criteria, Defeasibility, and Knowledge', *Proceedings of the British Academy*, 68: 455–79.
- (1986) 'Singular Thought and the Extent of Inner Space', in P. Petit and J. McDowell (eds.) *Subject, Thought, and Context*, Oxford: Clarendon Press.
- (1995) 'Knowledge and the Internal', *Philosophy and Phenomenological Research*, 55: 877–93.
- (2002) 'Responses', in N. H. Smith (ed.) *Reading McDowell: On Mind and World*, London: Routledge.

- (this volume) 'The Disjunctive Conception of Experience as Material for a Transcendental Argument'.
- Moore, G. E. (1925) 'A Defence of Common Sense', in J. H. Muirhead (ed.) *Contemporary British Philosophy*, 2nd series, London: Allen and Unwin.
- (1939). 'Proof of an External World', *Proceedings of the British Academy*, 25: 273–300.
- Neta, R. and Pritchard, D. H. (forthcoming) 'McDowell and the New Evil Genius', *Philosophy and Phenomenological Research*.
- Nuccetelli, S. (2003) (ed.) *New Essays on Semantic Externalism and Self-Knowledge*, Cambridge, MA: MIT Press.
- Pritchard, D. H. (2002a) 'McKinsey Paradoxes, Radical Scepticism, and the Transmission of Knowledge across Known Entailments', *Synthese*, 130: 279–302.
- (2002b) 'Recent Work on Radical Scepticism', *American Philosophical Quarterly*, 39: 215–57.
- (2002c) 'Resurrecting the Moorean Response to Scepticism', *International Journal of Philosophical Studies*, 10: 283–307.
- (2003) 'McDowell on Reasons, Externalism and Scepticism', *European Journal of Philosophy*, 11: 273–94.
- (2005a) *Epistemic Luck*, Oxford: Oxford University Press.
- (2005b) 'Scepticism, Epistemic Luck and Epistemic Angst', *Australasian Journal of Philosophy*, 83: 185–206.
- (2005c) 'The Structure of Skeptical Arguments', *The Philosophical Quarterly*, 55: 37–52.
- (2005d) 'Wittgenstein's *On Certainty* and Contemporary Anti-Scepticism', in D. Moyal-Sharrock and W. H. Brenner (eds.) *Investigating 'On Certainty': Essays on Wittgenstein's Last Work*, London: Palgrave Macmillan.
- (2006) 'Perceptual Knowledge, Relevant Alternatives, and Discrimination', manuscript.
- (2007a) 'How to be a Neo-Moorean', in S. Goldberg (ed.) *Internalism and Externalism in Semantics and Epistemology*, Oxford: Oxford University Press.
- (2007b) 'Knowledge, Luck, and Lotteries', in D. Pritchard and V. Hendricks (eds.) *New Waves in Epistemology*, London: Palgrave Macmillan.
- Pryor, J. (2000) 'The Skeptic and the Dogmatist', *Noûs*, 34: 517–49.
- Sosa, E. (1999) 'How to Defeat Opposition to Moore', *Philosophical Perspectives*, 13: 141–54.
- Stroud, B. (1994) 'Scepticism, "Externalism", and the Goal of Epistemology', *Proceedings of the Aristotelian Society*, Supplementary Volume, 68: 290–307.
- (2002) 'Sense-Experience and the Grounding of Thought', in N. H. Smith (ed.) *Reading McDowell: On Mind and World*, London: Routledge.
- Williams, M. (1991) *Unnatural Doubts: Epistemological Realism and the Basis of Scepticism*, Oxford: Basil Blackwell.
- Williamson, T. (2000) 'Scepticism and Evidence', *Philosophy and Phenomenological Research*, 60: 613–28.
- Wright, C. (1991) 'Scepticism and Dreaming: Imploding the Demon', *Mind*, 397: 87–115.

- Wright, C. (2002) '(Anti-)Skeptics Simple and Subtle: G. E. Moore and John McDowell', *Philosophy and Phenomenological Research*, 65: 330–48.
- (2004) 'Warrant for Nothing (and Foundations for Free)?', *Proceedings of the Aristotelian Society*, Supplementary Volume, 78: 167–212.
- Wittgenstein, L. (1969) *On Certainty*, ed. G. E. M. Anscombe and G. H. von Wright, trans. D. Paul and G. E. M. Anscombe, Oxford: Blackwell.
- Zagzebski, L. (1996) *Virtues of the Mind: An Inquiry into the Nature of Virtue and the Ethical Foundations of Knowledge*, Cambridge: Cambridge University Press.

In Defence of Disjunctivism

Ram Neta

1 DISJUNCTIVISM ABOUT PERCEPTION

Right now, I see a computer in front of me. Now, according to current philosophical orthodoxy, I could have the very same perceptual experience that I'm having right now even if I were not seeing a computer in front of me. Indeed, such orthodoxy tells us, I could have the very same experience that I'm having right now even if I were not seeing anything at all in front of me, but simply suffering from a hallucination. More generally, someone can have the very same perceptual experience no matter whether she is enjoying a veridical perception of some mind-independent object, or merely hallucinating. What differs across these two kinds of case is not the kind of experience that she has, but rather the connections between her experience and the rest of the world. So say most philosophers.

Disjunctivism rejects this widely held picture of perception and hallucination. Here is a typical statement of disjunctivism about perception (Crane 2005: section 3.4):

The disjunctive theory does not deny that there is some true description under which both the perception of a snow-covered churchyard and a subjectively indistinguishable hallucination of a churchyard can fall. It is easy to provide such a true description: both experiences are experiences which are subjectively indistinguishable from a perception of a snow-covered churchyard. Disjunctivists do not deny that such a true description is available. What they deny is that what makes it true that these two experiences are describable in this way is the presence of the same fundamental kind of mental state in the case of perception and hallucination. In the case of the perception, what makes it true that the description applies is that the experience is a perception of the churchyard; in the hallucinatory case, what makes it true that the description applies is that the experience is a hallucination of the churchyard. What the disjunctivist rejects is what J. M. Hinton calls "the doctrine of the 'experience' as the common element in a given perception and an indistinguishable hallucination" (Hinton 1973:

For helpful discussion and comments, I am grateful to Alex Byrne, Adrian Haddock, Fiona Macpherson, Eric Marcus, M. G. F. Martin, Dylan Sabo, and Susanna Siegel.

71). The most fundamental common description of both states, then, is a merely *disjunctive* one: the experience is *either* a genuine perception of a churchyard or a mere hallucination of one. Hence the theory's name.

For the disjunctivist, veridical perception of, say, a tomato is one kind of mental state, whereas hallucination of a tomato is simply another kind of mental state altogether—they are not species of a common genus. Although there are, of course, similarities between the two kinds of mental states, these similarities are not the result of there being a single kind of mental state present in both veridical perception and hallucination. Thus, for the disjunctivist, the similarities between a veridical perception and a subjectively indistinguishable hallucination are not like the similarities between one tiger and another tiger: different creatures of the same species. Nor are they like the similarities between a particular Bengal tiger and a particular Siberian tiger: different creatures of different species, but belonging to the same genus. Rather, the similarities between a veridical perception and a hallucination are like the similarities between a tiger and a hologram of a tiger: the former is an animal of a certain genus and species, whereas the latter is not an animal at all but a fundamentally different kind of thing. There is no genus the species of which include tigers and holograms of tigers. What tigers and holograms of tigers have in common is simply that they are visually indistinguishable, at least from a certain vantage point. Similarly, the disjunctivist says, the only thing that veridical perceptions and hallucinations have in common is that they are 'subjectively indistinguishable'. (Eventually, we'll say more about what 'subjective indistinguishability' amounts to.)

Of course, this analogy, and the statement of disjunctivism quoted above, is open to many different interpretations and elaborations. Disjunctivism, as stated above, is not a very specific doctrine, since there are many different ways of understanding the claim that perception and hallucination involve 'fundamentally different kinds' of experience, or mental state. In this essay, I will not distinguish these various versions, since the points that I want to make apply equally well to all of the versions with which I'm familiar.

Furthermore, different philosophers have very different views about what experience is, and what it can do for us, and so the issue of whether or not disjunctivism is true will look very different to these different philosophers. What one takes to be at stake in the issue of whether or not disjunctivism is true will vary greatly, depending upon one's other views about experience. In this essay, I will try to steer clear of the issue of what's at stake in whether or not disjunctivism is true. The reason I will try to steer clear of this issue is not that I don't regard it as important. On the contrary, I regard it as very important. The reason I try to steer clear of it here is that I would like my discussion of disjunctivism to avoid presupposing one or another of these background views about experience. A necessary cost of such neutrality is that I remain silent on what I take to be a very important issue: the issue of what difference it would make if disjunctivism is true.

Although disjunctivism has won a few adherents in recent years (for example, Hinton 1973; Snowdon 1980–1; McDowell 1982; McDowell 1994; Putnam 1999; Martin 2002, 2004; Travis 2004), it is still overwhelmingly unpopular. Partly, this is

because many philosophers regard the existing arguments for disjunctivism as unconvincing. And partly, it is because many philosophers are persuaded by a particular argument against disjunctivism. In this essay, I'll defend disjunctivism from this popular argument against it. I will also defend one of the current arguments for disjunctivism against a popular objection.

My aim in this essay is not to argue that disjunctivism is true, or that one of the arguments for it is sound or ultimately convincing. Rather, I aim to argue only that the disjunctivism can cogently reply to the most influential argument that's been given against disjunctivism, and can also cogently reply to a particular argument that's been given against one prominent argument for disjunctivism. If we should in the end reject disjunctivism, or reject the particular argument for disjunctivism that I will focus on, it's not for the reasons most commonly given.

2 AN ARGUMENT FOR DISJUNCTIVISM ABOUT PERCEPTION AND A CRITICISM OF THIS ARGUMENT

One prominent argument for disjunctivism—the single pro-disjunctivist argument that I will focus on in this essay—is epistemological. McDowell (1982/1998: 390–1) states the argument this way:

The root idea is that one's epistemic standing on some question cannot intelligibly be constituted, even in part, by matters blankly external to how it is with one subjectively. For how could such matters be other than beyond one's ken? And how could matters beyond one's ken make any difference to one's epistemic standing? . . . When someone has a fact made manifest to him, the obtaining of the fact contributes to his epistemic standing on the question. But the obtaining of the fact is precisely not blankly external to his subjectivity, as it would be if the truth about that were exhausted by the highest common factor.

McDowell's point seems to be this: veridical perception gives us knowledge of independently existing things, and hallucination does not. So, veridical perception puts us in a better epistemic position than hallucination does. But, so the argument goes, the only satisfactory explanation of this fact is an explanation according to which veridical perception and hallucination are *subjectively* different. What is it for veridical perception and hallucination to be *subjectively* different? I take this to amount to the following: a person who is enjoying veridical perception has *reflectively accessible epistemic reasons for belief* that are not possessed by the person who is merely having a hallucination. It is, of course, a substantive and tricky issue what it is for something to be reflectively accessible to a subject. But we can all agree that there is some sense in which, for example, one's own conscious states are reflectively accessible to oneself. Whatever that kind of access is, let's use the phrase 'reflective access' to denote that particular kind of access—whatever its nature may be. So I take McDowell to be saying: the only satisfactory explanation of the fact that veridical perception puts us in a better epistemic position than hallucination does is an explanation according to which veridical perception provides one with different reflectively accessible reasons for belief from those provided one by hallucination. And the only explanation that

involves that claim, so the argument continues, is an explanation according to which veridical perception is a fundamentally different kind of mental state than hallucination. Since the only satisfactory explanation of the fact that veridical perception puts us in a better epistemic position than hallucination does is an explanation according to which veridical perception is a fundamentally different kind of mental state than hallucination, it follows that the only satisfactory explanation of this fact implies disjunctivism about perception.

Many philosophers will resist McDowell's argument by rejecting the premise that the only satisfactory explanation of how veridical perception can put us in a better epistemic position than hallucination does is an explanation according to which veridical perception and hallucination are subjectively different. (Indeed, given my gloss of this subjective difference above—in terms of a difference in reflectively accessible reasons for belief—even such disjunctivists as Martin might resist McDowell's epistemological premise.) These objectors might say, for instance, that the best explanation of this fact is simply this: veridical perception—unlike hallucination—is, or involves, a reliable source of true beliefs. In other words, veridical perception puts us in a better epistemic state than hallucination does, not because of any difference in the experiences involved in each, but because of some difference that's extrinsic to the experiences involved in each. So although there is a real difference between the perception-induced epistemic state and the hallucination-induced epistemic state, this is not a difference in what McDowell would call 'subjectivity'. It is not a difference in what is reflectively accessible to the subject herself. That is one reply that many philosophers would issue to McDowell's argument.

Now, I will not devote any space in this essay to evaluating the relative merits of the various explanations of the fact that perception puts us in a better epistemic position than hallucination does. This is a complicated epistemological issue well beyond the scope of this essay. Instead, I want to focus on a different response that many philosophers have to McDowell's argument for disjunctivism. Recall that, according to McDowell's argument, disjunctivism helps to explain the following fact: perception puts us in a better epistemic position than hallucination does. But many philosophers don't see how disjunctivism could help at all in explaining that fact. Suppose that perception is a fundamentally different kind of mental state than hallucination is. How—these philosophers would ask—could that difference possibly explain the fact that the former puts us in a better epistemic position than the latter? In fact—the objection continues—since the perception and the hallucination at issue are, by hypothesis, subjectively indistinguishable, it's not at all clear how the former could possibly put me in a better epistemic position than the latter, except by appeal to its reliability or veridicality, something extrinsic to the experience itself. The experience is just the same *from my point of view*, the objection has it.

Crispin Wright states this objection in the following form. If veridical perception and hallucination are subjectively indistinguishable, then being in the former state cannot, Wright says, justify me in taking myself to be warranted in believing *p*, if being in the latter state does not equally justify me in taking myself to be warranted in believing that *p* (Wright 2002: 342—3). More generally: if experiences *E*1 and *E*2 are subjectively indistinguishable by their possessor *S*, then there can be no difference

between what being in E1 makes S justified in taking himself to be warranted in believing and what being in E2 makes S justified in taking himself to be warranted in believing. Whatever difference there is in the epistemic state induced in S by his being in E1 and the epistemic state induced in S by his being in E2, it's not a difference in what S is justified in taking himself to be warranted in believing.

Now, the disjunctivist might grant Wright's point about which second-order beliefs one is justified in having, and still say that the difference between the epistemic state induced by perception and the epistemic state induced by hallucination is a difference in what first-order beliefs S is justified in holding—even if it's not a difference in what S is justified in taking himself to be warranted in believing. But, for the sake of argument, I will suppose that the disjunctivist wants to say something stronger than this. I will suppose that the disjunctivist wants to say precisely the thing that Wright claims she's not in a position to say: namely, that the difference between the epistemic state induced by perception and the epistemic state induced by hallucination is actually a difference in what the subject is justified in taking herself to be warranted in believing.

On this interpretation, then, McDowell's argument for disjunctivism proposes an explanation for the fact that perception puts us in a better epistemic position than hallucination does, but this explanation strikes many philosophers in the way that it strikes Wright: as not being a possible explanation at all. It's not merely that McDowell's argument offers an optional explanation, an explanation that's not obviously better—perhaps even worse—than some alternative explanations. Rather, the thought goes, it's that McDowell's argument offers a merely putative explanation that clearly cannot be of any explanatory value whatsoever.

Now, I will not attempt to show that the disjunctivist offers the best, or the only satisfactory, explanation of the fact that perception puts us in a better epistemic position than hallucination does. But I will attempt to show that the disjunctivist's explanation is not open to the popular objection stated by Wright. In short, I will argue that Wright's objection is not a good objection. So, whether or not the disjunctivist's explanation is, in the end, correct, it may well actually have some explanatory value.

Before pursuing that task, I will first examine an argument that's commonly offered for the conclusion that disjunctivism is false: an argument that's sometimes called 'The argument from hallucination'. According to the argument from hallucination, perception and hallucination must be mental states of fundamentally the same kind, and so disjunctivism about perception is false.

3 A COMMON ARGUMENT AGAINST DISJUNCTIVISM ABOUT PERCEPTION

In this section, I will state and examine the argument from hallucination. Perhaps the most prominent recent statement of the argument is in Johnston (2004: 122–3), which states the argument by appeal to the following dramatic example. I quote it in full so that we may then go on to consider various interpretations of Johnston's argument:

You are undergoing an operation for an aneurysm in your occipital lobe. The surgeon wants feedback during the operation as to the effects of the procedure on the functioning of your visual cortex. He reduces all significant discomfort with local anesthetic while he opens your skull. He then darkens the operating theater, takes off your blindfold, and applies electrical stimulation to a well-chosen point on your visual cortex. As a result, you hallucinate dimly illuminated spotlights in a ceiling. (You hallucinate lights on in a ceiling. As yet, you are not at all aware of the lights or the ceiling of the operating theater.) As it happens, there really are spotlights in the ceiling at precisely the places where you hallucinate lights. However, these real lights are turned off, so that the operating theater is too dark to really see anything . . .

While maintaining the level of electrical stimulation required to make you hallucinate lights on in a ceiling, the surgeon goes on to do something a little perverse. He turns on the spotlights in the ceiling, leaving them dim enough so that you notice no difference. You are now having what some call a 'veridical hallucination'. You are still having a hallucination for you are not yet seeing the lights on in the ceiling, the explanation being that they still play no causal role in the generation of your experience. Yet your hallucination is veridical or in a certain way true to the scene before you; there are indeed dim lights on in a ceiling in front of you.

In the third stage of the experiment the surgeon stops stimulating your brain. You now genuinely see the dimly lit spotlights in the ceiling. From your vantage point there on the operating table these dim lights are indistinguishable from the dim lights you were hallucinating. The transition from the first stage of simple hallucination through the second stage of veridical hallucination to the third stage of veridical perception could be experientially seamless. Try as you might, you would not notice any difference, however closely you attend to your visual experience.

Of course, at the level of brain states, there will be some causal explanation for the experiential seamlessness. Whether one's brain is stimulated by the scene before one's eyes or by the direct application of electrical impulses, the effects on one's brain will be very similar in respects relevant to the causation of experience. This explanation in terms of brain states raises another explanatory question, which is our real concern. When we say that either way the effects on the brain are very similar in respects relevant to the causation of experience, we rely upon a picture according to which the differences at the level of brain states make no discernible difference at the level of experience. It is very likely that there are some intrinsic differences between the brain processes in the two cases. The idea of such differences not making a discernible difference at the level of experience begs for a characterization of what is taking place at the level of experience. Accordingly, our question is: What kinds of things can visual experience be a relation to so that in a case of hallucination and a case of seeing there need be no difference which the subject can discern? In itself, appeal to ever so slightly different brain states cannot answer that question.

Johnston goes on to claim that the disjunctivist cannot answer this last question in a satisfactory way. And this, according to Johnston, is a compelling reason to reject disjunctivism. In short, Johnston claims, the disjunctivist cannot provide a satisfactory explanation of the seamlessness of the experiential series described in this example.

Now, it is crucial that we be clear about precisely what it is that the disjunctivist is here alleged to be unable to explain. Johnston describes it as the "experiential seamlessness" present in the example, that the difference in brain states makes "no discernible difference at the level of experience", and that there is "no difference which the subject can discern". Precisely how should we understand these claims?

Now, I take it that Johnston does not mean to be supposing that the surgical patient cannot *by any means* discern a difference between perceiving and hallucinating; the patient could perhaps discern this difference by simply asking the surgeon to announce when the patient is perceiving and when he is hallucinating, or he could discern the difference by looking at some electronic monitor of his brain functions, or in various other ways. I take it that what Johnston means when he says that the surgical patient cannot discern the difference between perceiving and hallucinating is that the patient cannot discern this difference without the help of informants or instruments—in short, *without the assistance of sources of information distinct from his own thinking or experience*. And I suppose that, when Johnston speaks of ‘experiential seamlessness’, what he means is simply (what he takes to be) the lack of anything in the subject’s experience that would, by itself, enable the subject to tell the difference between perceiving and hallucinating.

What is it to ‘tell the difference’ or ‘discern the difference’ between two things? We might reasonably assume that Johnston is using these locutions in such a way that, S’s knowing which of *p* or *q* is true is a sufficient condition for S’s being able to ‘tell the difference’ or ‘discern the difference’ between *p*’s being true and *q*’s being true. For instance, if S knows which of seeing or hallucinating is occurring in her now, then she can discern the difference between seeing and hallucinating. Moreover, I assume that if S knows which of seeing or hallucinating is occurring in her now, and she knows this without the assistance of any source of information distinct from her own thinking or experience, then she can discern the difference between seeing and hallucinating in precisely the way that Johnston is claiming she cannot discern the difference.

If this is what Johnston has in mind, then I take it that his objection to disjunctivism comes to this: the surgical patient in the example above cannot know—at least not without the assistance of sources of information distinct from his own thinking or experience—which of seeing or hallucinating he’s undergoing at any particular moment. It follows that the surgical patient cannot ‘discern the difference’ (in the relevant sense) between seeing and hallucinating, and the disjunctivist cannot satisfactorily explain this inability on the part of the surgical patient. But some anti-disjunctivists can perfectly well explain this inability. They can say, for instance, that the surgical patient is in the very same state of mind both when he’s hallucinating and when he’s perceiving, and this is why the surgical patient cannot know which of the two he’s undergoing, without the assistance of sources of information distinct from his own thinking or experience. So there is a fact—an epistemic fact about the surgical patient and others like him—that some anti-disjunctivists can satisfactorily explain, but that the disjunctivist cannot satisfactorily explain. That, according to Johnston, constitutes a powerful argument against disjunctivism. This is my first interpretation of Johnston’s argument against disjunctivism.

One worry about this first interpretation of Johnston’s argument is that it gives no role to the gradualness of the transition that seems so important to Johnston. What difference could it make to Johnston whether the transition between perception and hallucination is gradual? It might be suggested that what Johnston has in mind is not merely that the surgical patient cannot tell, on the basis of his experience, whether he is perceiving or hallucinating. It might be suggested that what Johnston has in mind

instead is this: the change from hallucinating to perceiving is a gradual change, but this gradualness cannot be modelled by the disjunctivist, who claims that perceiving and hallucinating are fundamentally different kinds of mental events. Now, if this is Johnston's argument, then I don't see how it's supposed to be compelling at all. Consider two things that are fundamentally different in kind: say, a tiger and a hologram of a tiger. If each of those two things is composed of a great many parts, then there could be a gradual transition from one thing to the other thing by a temporal series of replacements of parts of one for parts of the other. We could replace one small part of the tiger with one small part of the hologram, and then continue this series of replacements one-by-one, and thereby make a gradual transition from the tiger to the hologram of the tiger. (We may suppose that the tiger is unconscious or paralysed, so as to enable the replacement to proceed without effect on its behaviour.) Of course, if some epistemicist view of vagueness is correct, then there will in fact be a sharp cut-off point where we no longer have a tiger, and a sharp cut-off point where we finally have a hologram, and these two cut-off points could be distinct. But from the point of view of an observer, this series of transitions would be gradual and the cut-off points would not be noticeable. So I don't see how the gradualness of the transitions from hallucinating to perceiving can play any role in Johnston's argument against the disjunctivist.

It seems to me, therefore, that the second interpretation of Johnston's argument just considered should not be the interpretation we focus on. Later on, I will consider a third interpretation, but I'll do so only after I reply to Johnston's argument, as the first interpretation represents it.

So, to sum up, here's where we've got so far. According to disjunctivism about perception, veridical perception and hallucination involve fundamentally different kinds of mental states. What reason have we to accept this highly unorthodox philosophical view? The epistemological argument that I've chosen to focus on may seem particularly unpromising, as it is open to the popular objection stated by Crispin Wright. Worse yet, Johnston offers an argument which, on at least one interpretation is seemingly compelling, against disjunctivism about perception. It seems, therefore, that the prospects for such disjunctivism are bleak.

In the remainder of this essay, I'll argue that the prospects for disjunctivism about perception are in fact much brighter than they now seem. In the next section, I'll consider how the disjunctivist might easily explain the fact that Johnston (according to the first interpretation of his argument) accuses the disjunctivist of being unable to explain. And I'll then consider a third interpretation of Johnston's argument, and show that the disjunctivist can reply to the argument construed in that third way as well. Finally, I'll return to examine and criticize Wright's objection against McDowell's argument for disjunctivism.

4 FACTORS THAT EXPLAIN IGNORANCE OR TRUTHFUL ATTRIBUTIONS OF IGNORANCE

In this section, I will consider how a disjunctivist might explain the fact that perception and hallucination are subjectively indistinguishable for Johnston's surgical

patient, that is, that the surgical patient does not know which of perception or hallucination he's undergoing at any moment in the series of surgically induced experiences. Let's begin by considering some factors that can prevent us from knowing things that we would otherwise know. As I will argue, the disjunctivist can appeal to some such factors in order to rebut the argument from hallucination (at least as the first interpretation represents it), and also in order to defend his explanation of the fact that perception can put us in a better epistemic position than hallucination does. To see how this works, let's consider what these factors might be. We will do this by appeal to Carl Ginet's classic 'barn façades' example, as described in Goldman (1976).

Consider the following two cases:

Henry 1: Henry is driving through some farmland. He gets a clear, unobstructed look at a barn that's just a few yards from the side of the road. He sees the barn for enough time to recognize it clearly as a barn. His vision is normal, and he is alert and of sound mind. He has no reason whatsoever to suspect that it is not a barn, or to distrust his senses. On the basis of his current experience, Henry is fully confident that that is a barn.

Henry 2: This case is just like Henry 1, except this time the farmland through which Henry is driving is full of barn façades. We can suppose that Henry has not yet seen any of these barn façades, and that he has not the slightest reason to suppose that there are any barn façades around. Again, on the basis of his current experience, Henry is fully confident that that is a barn.

Now, when we think about Henry 1 on its own, we are typically inclined to think that Henry knows that that is a barn. But when we think about Henry 2 on its own, we are typically inclined to think that Henry does not know that that is a barn. And yet the truth-value of Henry's belief that that is a barn is held fixed across the two cases: the belief is true in both cases. So too is Henry's confidence.

What lessons can we learn from this pair of cases? Different philosophers have drawn different lessons. Many philosophers (we may call them 'subject contextualists') say that whether or not Henry knows that *that is a barn* depends on features of Henry's context of which Henry may be unaware (for example, whether or not there are barn façades around). Other philosophers ('attributor contextualists') say that whether or not we can truthfully claim that Henry knows that *that is a barn* depends on features of our context (for example, whether or not we're thinking about Henry 2). Some philosophers accept both of these views. And of course there are various versions of each of these views, corresponding to the different features of the subject's context or the attributor's context that can be identified as relevant, and corresponding to the different ways in which those features can affect (the truth-conditions of our verdicts concerning) Henry's knowledge. I mention these various views not in order to settle on one of them, but only because I want to remind the reader of the variety of factors that might plausibly be thought to lead us truthfully to deny that Henry knows that *that is a barn*.

Notice that I have said that the subject's epistemic status—or what we can truthfully say about the subject's epistemic status—*depends upon* these contextual factors.

I have not said that it is *constituted by* any contextual factors. This distinction between dependence and constitution is important. There are lots of ways in which something can depend upon a factor without being constituted by that factor. For instance, my thinking about epistemology depends upon my being alive. I could not think about epistemology if I were not alive. But my thinking about epistemology is not constituted by my being alive. My being alive is not that in virtue of which I think about epistemology. Again, my eating a meal depends upon my having food available to me, but it is not constituted by my having food available to me. Not all dependence is constitutive dependence. And not all dependence of facts about a subject's epistemic status—or facts about what we attributors can truthfully say about a subject's epistemic status—upon contextual factors is constitutive dependence. Later on, we'll see why this distinction will be important to a McDowellian disjunctivist who wants to reply to Johnston's argument. But we're not in a position to see that just yet.

Now, against the background of our discussion of Henry, let's consider the following argument:

- (1) Henry knows—without the assistance of sources of information distinct from his own thinking or experience—that *that is a barn* (where Henry is using the demonstrative 'that' to refer demonstratively to the barn that he is at that moment seeing).
- (2) If S knows (without the assistance of sources of information distinct from his own thinking or experience) that *p*, and S knows (without the assistance of sources of information distinct from his own thinking or experience) that *p* implies *q*, and S deduces *q* from *p* and so believes *q* while retaining his knowledge that *p*, then S knows (without the assistance of sources of information distinct from his own thinking or experience) that *q*.
- (3) Henry knows—without the assistance of sources of information distinct from his own thinking or experience—that *that is a barn* implies *that is not a barn façade*, and deduces the latter from the former, and so believes the latter while retaining his knowledge of the former.
- (4) Henry knows—without the assistance of sources of information distinct from his own thinking or experience—that *that is not a barn façade*. (1, 2, 3)
- (5) If S knows (without the assistance of sources of information distinct from his own thinking or experience) that *p*, and S knows (without the assistance of sources of information distinct from his own thinking or experience) that not-*q*, then S knows (without the assistance of sources of information distinct from his own thinking or experience) which of *p* or *q* is true.
- (6) Henry knows—without the assistance of sources of information distinct from his own thinking or experience—which of *that is a barn* or *that is a barn façade* is true. (1, 4, 5)

The argument above employs two suppositions about Henry in particular (1 and 3), and two plausible general epistemic principles (2 and 5). The argument concludes that, without relying on any sources of information distinct from his own thinking or experience, Henry can know which of *that is a barn* or *that is a barn façade* is true.

That's just to say, given our gloss above, that Henry can tell the difference between that thing's being a barn and it's being a barn façade, and he can do this without relying on any source of information distinct from his own thinking or experience.

Shall we accept this conclusion? Many philosophers would, I think, reasonably refuse to accept it. But on what grounds could we reasonably refuse to accept it? A few philosophers, such as Dretske (1970), Nozick (1981), and Heller (1999), would reject the epistemic closure principle 2. But they recognize that this rejection of closure requires a good deal of argument, and many philosophers have criticized the arguments provided, and continue to find closure highly plausible. In this essay, I will not entertain the rejection of the closure premise 2.

Nor will I entertain the rejection of the general epistemic principle 5. I take 5 to be obviously true. We don't need an analysis of knowing which, or a comprehensive theory of the 'knows which' construction, in order to recognize that 5 is true.

If we accept 2 and 5, then the only way that we can escape the conclusion 6 is by refusing to grant 1 or 3. But I can see no good grounds for refusing to grant 3, so I'll restrict my attention to grounds that we might have to refuse to grant 1.

On what grounds could we reasonably refuse to grant 1? Perhaps on one or another of the grounds canvassed above, in connection with cases 1 and 2. For instance, we could say this: the cases in which conclusion 6 seems false are precisely those cases in which Henry is prone to mistake a barn façade for a barn (perhaps because of the frequency of barn façades in the relevant landscape). They are cases like Henry 2 above. But if, in those cases, Henry is prone to mistake barn façades for barns, then, even when there happens to be a real barn before him in such a case, he does not know that that is a barn. Had his environment been more cooperative, then he might have known that that is a barn. But if his environment is an epistemically treacherous place—as is Henry 2—then he is prone to mistake barn façades for barns, and consequently does not know that that is a barn. That is one way that we could reject supposition 1—at least for a certain range of cases in which conclusion 6 seems false.

Alternatively, we could resist the conclusion 6 quite generally by claiming that, when we go through the argument 1–6, we move from a context in which we can reasonably suppose that 1 is true to a context in which we cannot reasonably suppose that 1 is true. Henry's context does not change, and neither does Henry's epistemic situations. Rather, what changes is our context, and what words we can use to issue true descriptions of Henry's epistemic situation. These things change by virtue of our going through the argument 1–6. When we entertain supposition 1 on its own, we can truthfully affirm it. But by the time we've reached our conclusion 6, something or other that we've done has moved into a context in which we cannot truthfully affirm 6, or even truthfully affirm 1. Going through the argument has so changed the contents of our words that the sentences that we could truthfully assert when we set out on our argument are sentences that we can no longer truthfully assert. At least, so says the attributor contextualist, in an attempt to explain why we cannot truthfully affirm the conclusion 6. (Of course, different attributor contextualists will locate the context-shifting move, the 'sleight of mind', at different junctures.)

So we've now considered two ways to avoid the conclusion of the argument above. We could reject supposition 1 for certain cases in which we take 6 to be false. We could

reject 1 for those cases on the grounds that, in those cases, Henry's environment makes him prone to mistake a barn façade for a barn, and this proneness to error prevents Henry from knowing that that is a barn. Alternatively, we could reject supposition 1 on the grounds that there's something about going through the argument that puts us in a context in which we can truthfully deny 1. So we can offer a subject contextualist explanation of why 1 is false when Henry is in certain contexts. Or we can offer an attributor contextualist explanation of why we can truthfully deny 1 when we are in certain contexts. Either way, we can give an explanation of why we can truthfully deny 1—why we can truthfully attribute a certain kind of ignorance to Henry.

In the next section, we'll see how this discussion of Henry can help us to reply to Johnston's argument, on our first of the two interpretations of it considered above. We'll also consider a third interpretation of Johnston's argument, and reply to it as well.

5 RESPONSE TO JOHNSTON

Now that we've offered these two options—the subject contextualist option and the attributor contextualist option—for resisting the conclusion of the argument about Henry above, let's see how we might employ one or another of these options to resist the conclusion of the following analogous argument:

(1') The surgical patient knows—without the assistance of sources of information distinct from his own thinking or experience—that *this event is a perception of ceiling lights* (where the surgical patient uses the demonstrative 'that' to refer demonstratively to the event of his now seeing ceiling lights).¹

(2') If S knows (without the assistance of sources of information distinct from his own thinking or experience) that *p*, and S knows (without the assistance of sources of information distinct from his own thinking or experience) that *p* implies *q*, and S deduces *q* from *p* and so believes *q* while retaining her knowledge that *p*, then S knows (without the assistance of sources of information distinct from his own thinking or experience) that *q*.

(3') The surgical patient knows—without the assistance of sources of information distinct from his own thinking or experience—that *this event is a perception of ceiling lights* implies that *this event is not a hallucination of ceiling lights*, and deduces the latter from the former, and so believes the latter while retaining her knowledge of the former.

(4') The surgical patient knows—without the assistance of sources of information distinct from his own thinking or experience—that *this event is not a hallucination of ceiling lights*. (1', 2', 3')

(5') If S knows (without the assistance of sources of information distinct from his own thinking or experience) that *p*, and S knows (without the assistance of sources of

¹ I assume that seeings are events. If they are states instead, then I invite the reader to substitute 'state' wherever I have 'event'. The difference is immaterial to the argument in the text.

information distinct from his own thinking or experience) that not-*q*, then S knows (without the assistance of sources of information distinct from his own thinking or experience) which of *p* or *q* is true.

(6') The surgical patient knows—without the assistance of sources of information distinct from his own thinking or experience—which of *this event is a perception of ceiling lights* or *this event is an hallucination of ceiling lights* is true. (1', 4', 5')

Note that 6', the conclusion of this argument, is incompatible with the fact that Johnston was accusing the disjunctivist of being unable to explain satisfactorily (at least as we understood Johnston's accusation), namely the fact that the surgical patient cannot know whether he's veridically perceiving or hallucinating. The argument from hallucination begins by assuming quite reasonably that 6' is false, and it then claims that the disjunctivist cannot satisfactorily explain the fact that 6' is false. But we're now in a position to see why this accusation is wrong. The disjunctivist is in a perfectly good position to explain the fact that 6' is false.

One way that she could do this is by claiming that the surgical patient's environment makes him prone to mistake a hallucination of ceiling lights for a perception of ceiling lights, and this proneness to error prevents the surgical patient from knowing that this event is a perception of ceiling lights. Just as Henry in Henry 2 is in 'barn façade country', so too is the surgical patient in 'perception façade country'. In both cases, their environments are epistemically treacherous, and render them prone to error. Such proneness to error robs them of knowledge that they would otherwise have. That is one way for the disjunctivist to explain the falsehood of 6'.

Some philosophers might protest that this route is not open to the disjunctivist, for the datum that the disjunctivist must explain is not simply the datum that Johnston adduces—namely that *the surgical patient in the scenario described* cannot know whether he is veridically perceiving or hallucinating. It is rather a different datum—namely that none of us can ever know 'from the inside' whether we are veridically perceiving or hallucinating, for the two events are 'subjectively indistinguishable'. But in order for me to address this protest, I would need to understand it more fully: what is meant by saying that none of us can ever know this 'from the inside'? If it means simply that a subject can never know by reflection alone whether she is enjoying a veridical perception or a hallucination, then this, I take it, is something that McDowell would simply deny. He would say that, under normal circumstances of veridical perception, one can simply tell *by reflection alone* that one is veridically perceiving and not hallucinating. (Of course it does not follow from this that, when one is hallucinating or even seriously prone to hallucination, one can *in those cases* tell by reflection alone whether one is veridically perceiving or hallucinating!) So the claim that one can never tell by reflection alone whether one is veridically perceiving or hallucinating cannot serve as an undefended premise in the anti-disjunctivist's argument against McDowell: that would be to beg the question against McDowell. But then how should we understand these protests? Rather than deal with these issues in the abstract, I am now approaching what may end up being those very same issues more concretely by trying out different interpretations of Johnston's argument against disjunctivism.

Alternatively, the disjunctivist could reject 1' on the grounds that there's something about going through the argument that puts us in a context in which, by the end of the argument, we can truthfully deny 1'. Perhaps entertaining the possibility of hallucinations puts us in a context in which we can truthfully deny 1', no matter how well situated the surgical patient might be. That is another way for the disjunctivist to explain the falsehood of 6', at least relative to the context that we put ourselves in by going through the argument above.

So the disjunctivist can offer a subject contextualist explanation of why 1' and 6' are false in the case of the surgical patient as described by Johnston. Or the disjunctivist can offer an attributor contextualist explanation of why she can truthfully deny 1' and 6' when she is in certain contexts, such as the context that she's in when she completes the argument above. Either way, the disjunctivist can give an explanation of why we can truthfully deny 1' and 6'—why she can truthfully attribute a certain kind of ignorance to the surgical patient. But, given our present understanding of Johnston's objection against disjunctivism, this is precisely what Johnston was saying that the disjunctivist was unable to explain. And so Johnston's objection (at least as we've been interpreting it) is wrong.

Now, it may be objected that the preceding contextualist responses to Johnston's argument are not open to a disjunctivist like McDowell. For McDowell recommends disjunctivism on the grounds that it offers the best—indeed the only satisfactory—explanation of why veridical perception puts us in a better epistemic position than hallucination does. But it is supposed to do so by virtue of being the only response that allows veridical perception to provide one with different reflectively accessible reasons for belief from those which hallucination does. But, it will be said, what is reflectively accessible to a subject is restricted to those things that are constituted solely by factors that are, in some way, *internal to that subject*. And, it may be objected, what reflectively accessible reasons that surgical patient has cannot depend upon contextual factors. So—the objection concludes—the contextualist responses I've just canvassed are not open to McDowell, at least as I read him.

The problem with this objection is that it falsely assumes that, if the truth-values of our epistemic attributions *depend* upon contextual factors, this dependence must be *constitutive dependence*. But this is precisely the confusion that was exposed in the preceding section. Not all dependence is constitutive dependence. The surgical patient's epistemic status—or the truth-values of our attributions of epistemic status—could depend upon contextual factors without being constituted by those contextual factors. It could be that the reasonableness of the surgical patient's belief, in a particular case, that he is perceiving ceiling lights, depends upon the extent to which he is prone, in his actual surroundings, to confuse a hallucination for a perception—but this dependence is not constitutive, so we can maintain the view that veridical perception and hallucination differ in what reflectively accessible reasons for belief they provide, and we can also maintain the view that what is reflectively accessible to a subject is internal to that subject: the veridical perception and the hallucination are both of them—McDowell may allow—internal to that subject. (Of course, McDowell may wish to reject the popular view that only what is internal to a subject is reflectively accessible to a subject. But what I am pointing out here is simply that he *need not*

reject it in order to claim both that there is a difference between the reflectively accessible reasons for belief one gets from veridical perception and the reflectively accessible reasons for belief that one gets from hallucination, and also that some version of subject contextualism or attributor contextualism is true.) So long as such epistemic features are not constituted by those contextual factors, McDowell can accept the claim of dependence. Thus, unless there is an *argument*, and not a mere *assumption*, that the dependence in question is constitutive dependence, there is no obstacle to McDowell's accepting the relevance of such contextual factors to what is reflectively accessible to the subject.

In response to the argument that I've just offered against Johnston, Johnston might reply as follows. The reason that the presence of barn façades in Henry 2 robs Henry of his knowledge that that is a barn is that *barn façades and barns look alike* (at least from a certain angle). The presence of sheep façades, for instance, would not rob Henry of his knowledge that that is a barn, for sheep façades do not look like barns. Analogously, the reason why the occurrence of hallucinations robs the surgical patient of his knowledge that this event is a perception of ceiling lights is that *perceptions of ceiling lights and hallucinations of ceiling lights seem alike* to the one having them. But if perceptions and hallucinations *seem alike* to the one having them, then this can only be because they are, or involve, the very same experience. And this refutes disjunctivism.

I propose this as the third interpretation of Johnston's argument against disjunctivism: perceptions and hallucinations seem alike to the one having them, and the disjunctivist cannot admit this, for if they seem alike to the one having them then this can only be because they are, or involve, the very same experience.

Now, whatever the merits of this argument, it's not an argument that can work against the disjunctivist. From the fact that perceptions of ceiling lights and hallucinations of ceiling lights seem alike (at least to the one having them), it doesn't follow that they are, or involve, the very same experience. If such seeming were all that there was to experience, then seeming alike would suffice for being the same experience. But why should the disjunctivist grant that such seeming is all there is to experience? Recall why it is that we appeal to such seeming in the first place: we appeal to the fact that hallucinations seem like perceptions in order to explain the fact that the frequency of hallucinations can rob the surgical patient of his knowledge that he's perceiving, when he's perceiving. Now, in order for this similarity in seeming to do this bit of explanatory work, does such similarity in seeming require sameness of experience? No. It could do this explanatory work even if there were no sameness of experience. For example, suppose that we take it to be a sufficient condition of hallucination *H*'s seeming like perception of *P* that all of the introspectibly detectable features of *H* are also introspectibly detectable features of *P*, and vice versa. If all of the introspectibly detectable features of *H* were also introspectibly detectable features of *P*, and vice versa, that would help to explain why, when one is *prone* to suffer hallucinations, one cannot know that one is perceiving, even when one happens to be perceiving. But even if all of the introspectibly detectable features of *H* are introspectibly detectable features of *P*, and vice versa, it doesn't follow that *H* and *P* are, or involve, the same experience. So *H* and *P* could seem alike in this sense—all of the

introspectibly detectable features of one are also introspectibly detectable features of the other—without their being, or involving, the same experience. Moreover, their seeming alike could do precisely the explanatory work that the disjunctivist needs it to do without *H* and *P* being, or involving, the same experience.

Johnston might respond to this argument as follows: the fact that *H* and *P* seem alike—their sharing of introspectively detectable features (if that's how we choose to cash out such 'seeming alike')—of course does not imply, but is nonetheless best explained by, the claim that they involve the same experience. But why should we accept this? Tigers and tiger-holograms look alike (that is, share visually detectable features), but this is explained by appeal to their similar impacts on our visual machinery, not by appeal to their being the same kind of thing. Why can't the sharing of introspectibly detectable features by *H* and *P* be equally well explained by appeal to their similar impacts on our introspective machinery? (I'm using the phrase 'introspective machinery' here to refer to whatever causal processes are involved in introspective awareness. We may assume that some causal processes are involved in such awareness. This assumption does not commit me to the claim that our mental states exist, or have their properties, independently of our introspective awareness of them. That claim would follow only from the conjunction of the claim that there is an introspective causal process, and the Humean claim that causes and effects are 'distinct existences'. But I see no reason to accept the latter Humean claim.)

Finally, Johnston might protest that 'seeming the same' cannot amount to, or be explained by, having the same impact on our introspective machinery, for beasts that simply lack introspective machinery might nonetheless have experiences that 'seem the same' in the relevant sense, that is, are subjectively indistinguishable. This objection can be easily met: we can understand the fact that such experiences 'seem the same' to the beast in question in terms of what sort of impact those experiences would make on some standard-issue introspective machinery, whether or not the beast itself has the requisite machinery. Of course, we may treat different sorts of introspective machinery as standard-issue, and so end up counting different sorts of experiences as 'seeming the same' to the beast. Some philosophers will protest that this is to allow an objectionable indeterminacy in our description of the beast's phenomenology. But I don't see what evidence we have for regarding the beast's phenomenology as truthfully describable in terms that are any more determinate than this.

In this section, we've seen how the disjunctivist can respond to Johnston's argument. Johnston alleges that the disjunctivist cannot explain a certain epistemic shortcoming on the part of the surgical patient. But the disjunctivist can explain this epistemic shortcoming in any number of ways, corresponding to the ways in which epistemologists seek to explain Henry's epistemic shortcoming concerning the presence of a barn. So much for Johnston's argument against disjunctivism.

6 RESPONSE TO WRIGHT

Now, let's return to Wright's argument to the effect that the disjunctivist cannot explain how perception can put me in a better epistemic position than hallucination

can. Wright's point, recall, was this: since perceiving one's hand is subjectively indistinguishable from being a brain-in-a-vat that's being electrochemically stimulated to have experiences as of having a hand, it follows that, even when I do perceive a hand, nothing can justify me in taking myself to have a warrant for believing that I see a hand, rather than believing the weaker disjunction, either I see a hand or I am a brain-in-a-vat being electrochemically stimulated to have experiences as of a hand. And this is what McDowell is, on my interpretation, committed to denying. Now, how can the disjunctivist plausibly deny this?

To answer this question, let's consider: what makes Wright's general principle at all plausible? It seems to me that the only possible answer to this question is: reflection upon particular pairs of cases in which a subject is, in one case, perceiving, and in the other case, hallucinating or otherwise having some misleading experiential evidence. For instance, there is Johnston's surgical patient who is sometimes veridically perceiving ceiling lights and is sometimes hallucinating ceiling lights. It seems that, despite the changes in the surgical patient's perceptual state, there is no change in what he is justified in taking himself to be warranted in believing. Such a case may seem to lend support to Wright's general principle, and thus lend support to Wright's argument against the epistemological argument for disjunctivism.

Now, what would the disjunctivist say about such cases? We've already seen that the disjunctivist can explain the subjective indistinguishability of perception and hallucination by appeal to features of the context of the perceiver, or features of the context of the epistemic appraisal, or both. But, no matter which of these factors the disjunctivist appeals to in explaining the subjective indistinguishability of perception and hallucination, it is open to the disjunctivist to appeal to those very same factors to explain our intuitions about what perceivers or hallucinators are justified in taking themselves to be warranted in believing. I'll illustrate this by appeal to the following pair of cases:

Moore 1: G. E. Moore holds up a hand, gets an unobstructed view of it, sees it, is fully confident that he sees it, is alert and of sound mind, and thinks 'I see a hand'.

Moore 2: Just like Moore 1, except that Moore is surrounded by brains-in-vats who are being electrochemically stimulated to have experiences as of being G. E. Moore, and each of whom thinks that it is G. E. Moore seeing his hands.

When we think about Moore 1 on its own, and apart from any consideration of sceptical scenarios, it seems quite plausible to say that Moore is justified in taking himself to be warranted in believing that he has hands. I frequently take myself to be warranted in believing that I have hands. For instance, when I am confident in the success of my plans to hammer some nails into the wall in order to hang a picture, and you challenge my confidence by saying 'how can you be so sure that you will succeed?', I will claim that your challenge is unreasonable, and that I am entirely warranted in being sure that my plans will succeed. But for me to be warranted in being sure that my plans will succeed, I must be warranted in being sure that I have hands—and I know this. So I take myself to be so warranted, and I am typically justified in taking myself to be so warranted.

But when we think about Moore 2, it seems quite implausible to say that Moore (who still veridically perceives his hands) is so justified. What generates the difference? Again, the disjunctivist can appeal to features of Moore's own context in Moore 1 and Moore 2 to explain this difference. Whatever contextual factors operate in Henry 2 to deprive Henry of knowledge that that is a barn, those same factors operate in Moore 2 to prevent Moore from being justified in taking himself to be warranted in believing that he has hands. Thus, by reflection on a case like Moore 2, we may be led to say—truthfully—that Moore's veridically perceiving his hands does not make him justified in taking himself to be warranted in believing that he has hands. But the correctness of this verdict about Moore in Moore 2 need not show that, in Moore 1, Moore's veridically perceiving his hands does not make him justified in taking himself to be warranted in believing that he has hands. Even if Moore's perception in Moore 2 does not render him so justified, it doesn't follow that Moore's perception in Moore 1 does not render him so justified.

In fact, the disjunctivist can grant Wright even more than this. The disjunctivist can even allow that, after thinking about Moore 2, we cannot truthfully admit even that Moore's perception in Moore 1 renders him justified in taking himself to be warranted in believing that he's seeing his hand. But the disjunctivist can allow this without giving up disjunctivism. She can do this by claiming that, when we reflect upon Moore 2, we alter our own context of appraisal, and put ourselves in a context in which we cannot truthfully attribute such justification to Moore in Moore 1. But it doesn't follow from this that we can never truthfully attribute such justification to Moore in Moore 1.

In short, the disjunctivist can appeal either to subject contextualism or to attributor contextualism in order to grant the truth of the intuitive judgements that seem to support Wright's general principle. But the disjunctivist can grant these intuitions without granting Wright's general principle, or Wright's argument against McDowell. The disjunctivist can claim that she is able to explain how perception sometimes—and hallucination never—makes one justified in taking oneself to be warranted in believing that one is perceiving. The disjunctivist can offer such an explanation consistently with allowing that perception sometimes does not make one so justified. And she can offer such an explanation consistently with allowing that, when we attributors entertain sceptical counterpossibilities, then we cannot truthfully ascribe such justification even to someone who does perceive. So once again, disjunctivism can accommodate the intuitions that motivate its critics, and it can do so by appeal to some variety of subject contextualism or attributor contextualism. And such contextualisms are in no way incompatible with the view that epistemic status is not constituted by contextual factors.

This is, of course, not to say, let alone show, that epistemic status is *not* constituted by contextual factors. But again, I am not attempting to show that disjunctivism is true, only that neither Johnston's argument against disjunctivism nor Wright's argument against McDowell's argument for disjunctivism is telling. There may be problems with disjunctivism, or with McDowell's argument for it, but those problems have not yet been exposed by either Johnston or Wright.

7 CONCLUSION

I have argued that Johnston's argument from hallucination does not tell against disjunctivism. I have also argued that the epistemological argument for disjunctivism is not subject to the objection that Wright levels against it. Maybe disjunctivism is true, and maybe it isn't. Maybe it provides the best explanation of the epistemic superiority of perception over hallucination, and maybe it doesn't. In this essay, I have not staked out a position on either of these issues, nor have I staked out a position on why they matter. I hope to have shown only that disjunctivism's dialectical position has been greatly underestimated.

REFERENCES

- Crane, T. (2005) 'The Problem of Perception', in Edward N. Zalta (ed.) *Stanford Encyclopedia of Philosophy (Spring 2005 Edition)*, <<http://plato.stanford.edu/archives/spr2005/entries/perception-problem/>>
- Dretske, F. (1970) 'Epistemic Operators', *Journal of Philosophy*, 67: 1007–23.
- Goldman, A. (1976) 'Discrimination and Perceptual Knowledge', *Journal of Philosophy*, 73: 771–91.
- Heller, M. (1999) 'Relevant Alternatives and Closure', *Australasian Journal of Philosophy*, 77: 196–208.
- Hinton, J. M. (1973) *Experiences*, Oxford: Clarendon Press.
- Johnston, M. (2004) 'That Obscure Object of Hallucination', *Philosophical Studies*, 120: 113–83.
- McDowell, J. (1982) 'Criteria, Defeasibility, and Knowledge', *Proceedings of the British Academy*, 68: 455–79. [Reprinted in his *Meaning, Knowledge, and Reality*, Cambridge, MA: Harvard University Press, 1998].
- (1994) 'The Content of Perceptual Experience', *Philosophical Quarterly*, 44: 190–205.
- Martin, M. G. F. (2002) 'The Transparency of Experience', *Mind and Language*, 17: 376–425.
- (2004) 'The Limits of Self-Awareness', *Philosophical Studies*, 120: 37–89.
- Nozick, R. (1981) *Philosophical Explanations*, Cambridge, MA: Belknap Press.
- Putnam, H. (1999) *The Threefold Cord*, New York: Columbia University Press.
- Snowdon, P. (1980–1) 'Perception, Vision, and Causation', *Proceedings of the Aristotelian Society*, 81: 175–92.
- Travis, C. (2004) 'The Silence of the Senses', *Mind*, 113: 57–94.
- Wright, C. (2002) '(Anti-)Sceptics Simple and Subtle: G. E. Moore and John McDowell', *Philosophy and Phenomenological Research*, 65: 330–48.

Perceptual-Recognitional Abilities and Perceptual Knowledge

Alan Millar

1 THE TOPIC

Disjunctivism is interesting, at least in part, because it provides a way of making sense of a relational conception of sensory experience in the face of an obvious *prima facie* objection. According to the relational conception, when I see something, for instance a book on the table in front of me, I have an experience that is *essentially relational* in that it would not be that very experience unless the book were there on the table.¹ The *prima facie* objection, in its most straightforward form, rests on the idea that there is a perfect hallucinatory counterpart to my experience—an experience such that everything looks to me just as it does in the perceptual case, though there is no book on any table before me. The objection is that the experience in the perceptual case cannot be essentially relational since it is the very same experience as the one in the hallucinatory counterpart case—a case in which no book is present. Under disjunctivism, in the form under consideration here, it is conceded that the way it looks to the subject as if things are in the perceptual case is the same as the way it looks to the subject as if things are in the hallucinatory counterpart case. But it is denied that this entails that the subject would have the same experience in these cases. Consider any statement as to how it looks to the subject that things are. Under disjunctivism, that statement is such that it could be made true by the subject's having an essentially relational experience. Or it could be made true by the subject's having a perfectly hallucinatory experience. And there are other possibilities involving admixtures of partial hallucination or illusion. There need be no psychological state in common

Work for this essay was part of a project on the Value of Knowledge funded by the Arts and Humanities Research Council. I am grateful to the Council for its generous support. For numerous helpful discussions I am also grateful to Duncan Pritchard and Adrian Haddock, who are co-participants in the project. Thanks are also due to those in the audience at the Glasgow Disjunctivism Conference in September 2005, who heard a related talk, and to Fiona Macpherson for comments on an earlier version.

¹ The term 'relational' is used in Campbell (2002: ch. 6) to capture the conception that is in play here.

between these different situations in virtue of which it is true that it looks to the subject as if such-and-such is so.

For the purposes of the present discussion my interest in the relational conception of experience is epistemological. I take it that a major attraction of this conception is that it is thought to provide the only way to do justice to our perceptual knowledge. A guiding idea at this point is that perceptual knowledge involves a certain kind of contact with objects. Paul Snowdon expresses the point in the following way:

perceiving, when it is attached to the right cognitive capacities, crucially enables a certain kind of thought contact to be established with elements in the perceiver's environment. (Snowdon 1990: 150)

The thought-contact of which Snowdon speaks is by means of demonstrative thought, understood as being individuated in part by the very objects that it picks out. Indeed for Snowdon (1990: 143):

for P to see O is for P to stand to it in that experiential relation R which (i) is distinctively visual and (ii) is such that its obtaining means that O can be an object of demonstrative thought by P, given that P is suitably cognitively endowed and attending.

Snowdon is not primarily concerned with epistemology, but with how we should think about the experiences involved in seeing an object. But the way he thinks of *that* naturally suggests a constraint on an adequate account of perceptual knowledge: since perceptual knowledge depends on demonstrative thoughts, the experiences that enable us to have perceptual knowledge must be conceived in such a way that it is intelligible that they enable subjects to have demonstrative thoughts about objects that are *perceptually* picked out. Relationalists, as I shall call them, think that only essentially relational experiences can explain how we can have demonstrative thoughts about objects. They think that if the experience I have as I look at the book on the table could be had by me even if no book were there, then the experience cannot explain how I can so much as think demonstratively about the book, never mind have perceptual knowledge about it.² But, of course, an adequate account of perceptual experience should not only shed light on how such experience enables us to have demonstrative *thoughts*, it should also account for the role of experiences in putting us in *cognitive* contact with objects and facts.³ Part of that story will be about how they enable us to have demonstrative thoughts about objects perceptually picked out. But another and vital part of the story must be about how we are able to recognize those objects for what they are. It will thus concern how experiences, on the part of those suitably endowed with perceptual-recognitional abilities, can put us in cognitive contact with facts about those objects.

² Campbell (2002: 120–4) is explicit that disjunctivism about experience should be understood to accord experiences an explanatory role in relation to the possibility of demonstrative thought.

³ McDowell (1995/1998: 402) speaks of a “cognitive purchase on an objective fact” and Scott Sturgeon (2000: 31) presents the disjunctivist’s account as one relying on the idea that veridical perception is “brute contact between mind and truth”. See also the passage from Child cited in the main text below. I make the metaphor of cognitive contact central to a discussion of disjunctivism in Millar (2007).

The relationalist account of perceptual experience favoured by disjunctivists has been found appealing in part because traditional non-relationalist accounts have seemed to make our knowledge of the world problematic. William Child (1994: 147) sums up the situation as follows.

one cannot have thoughts about Fs unless: either one is (or has been) in direct cognitive contact with Fs; or one can construct a way of thinking of Fs from concepts of kinds of thing with which one is (or has been) in direct cognitive contact. Now on the non-disjunctive conception of experience we are not in direct cognitive contact with the world, since the most basic mental characterization of experience is world-independent.

The present discussion is motivated by the thought that relationalists are right to suppose that our perceptual experiences enable us to have contact with objects in a manner that enables us to have demonstrative thoughts, and know things about, those objects. I present a picture of perceptual knowledge on which a central role is accorded to perceptual-recognitional abilities and abilities to pick out objects perceptually.⁴ I argue that the exercise of these abilities in response to the occurrence of experiences is what puts us in visual and cognitive contact with objects. I then raise the question whether the exercise of these abilities depends on our having essentially relational sensory experiences. It is, of course, agreed on all sides that *seeing* the book on the table is an essentially relational matter. The issue is whether the visual experiences one has when looking at the book have to be conceived as essentially relational. On a traditional (non-relationalist) conception of experience the right answer is, 'No'. That is because seeing the book is a matter, not only of having an appropriate visual experience, but also of having such an experience because of the presence of the book. I am not persuaded that this traditional answer is incorrect so far as it goes. But traditionalists have not in general provided satisfying explanations either of how the judgments we are prompted to make in response to experiences amount to knowledge or of the role of experiences in making thought-contact with objects possible. As I see it an account that invokes perceptual-recognitional abilities, when supplemented by an account of what I shall call perceptual-discriminative abilities, can do the business, without help from a relationalist conception of experience. I do not claim to have refuted the relationalist conception, only to have sketched a plausible picture of perceptual knowledge that undermines one of its motivations.

I shall begin with recognitional abilities and bring in discriminative abilities later.

2 PERCEPTUAL RECOGNITIONAL ABILITIES: A PRELIMINARY SKETCH

Suppose I am a professional gardener and horticulturalist. Working for a client I need to find out whether or not the shrubs growing in a particular plot of a garden are

⁴ Elements of the position outlined are presented in Millar (2007). Here I develop the conception of perceptual-recognitional abilities somewhat further and relate it to issues about justification for belief and about why knowledge matters that are not explored in the latter article.

azaleas. The client particularly wishes to know and is relying on me to tell her. I go to the plot and look. I see, and thereby come to know, that the shrubs are azaleas. In this situation I exercise a rather specific recognitional ability—an ability to tell by looking whether or not something before me is an azalea. This ability depends, of course, on my possessing the concept of an azalea. But it involves a lot more than that since it is an ability that I am not bound to have just in virtue of possessing the concept. It depends on my being able to recognize azaleas by sight. My possession of the concept arguably does not so depend. (The blind can think about azaleas.)

Perceptual-recognitional abilities under this conception are not theoretical posits. We routinely think of people as having the ability to tell by looking, hearing, feeling, and so on, whether or not something is so. Our ascriptions of knowledge are often constrained by consideration of whether the subjects in question have or lack such abilities, and by whether they are in a position to exercise them. Suppose that I am an electrician working with a partner. I need to know whether the bare wires I am looking are live. My partner is near the mains switch box. I ask him to tell me whether or not the power is off, assuming that he can find out by looking at the box. In employing this method of checking whether the power is off I rely on knowing that my partner has the ability to tell such a thing by looking and that he is, or can readily put himself, in a position to exercise this ability in these circumstances.

I take it that the following points, which deserve emphasis, fall within the scope of common-sense thinking or a modest development of such thinking.

- (1) The upshot of the exercise of my ability to tell by looking whether or not the shrubs are azaleas is that I find out whether or not they are. Finding out is simply coming to know. So as a result of looking I come to know either that what I am looking at are azaleas or that they are not.
- (2) If I had judged falsely that the plants in the plot were azaleas I would not have exercised the recognitional ability in question. The general point here is that the notion of the exercise of a recognitional ability is a *success* notion. Success in relation to finding out whether or not *p* is coming to know whether or not *p*. This is not to deny that there are abilities that are abilities to do something a good proportion of the times at which one tries. My ability to hit a darts board within the 25 ring might be like that. In this case success, and thus whether or not the ability has been exercised, is not to be measured by a single shot, but by the frequency of times I hit within the ring when trying to do so. The ability to tell whether or not something is an azalea is not like that. This ability is not *to judge correctly enough of the time*. It is an ability to tell by looking, and thus come to know, whether or not something is an azalea. It is exercised only if one ends up knowing.
- (3) In some circumstances in which I set out to tell whether or not certain shrubs are azaleas I might withhold judgement either way. For instance, the shrubs I am looking at might be shorn of leaves and flowers, and severely pruned. In this case too I would not exercise my ability to tell by looking whether or not

something is an azalea, but I would exercise a closely related ability—an ability to know when to hold back from judging either way. This requires me to be sensitive to whether or not I have a good enough look at the object in question to be able to tell whether or not it is an azalea and sensitive to whether the object in question has an appearance that enables it to be visually recognizable as an azalea.

- (4) Even though I have the ability to tell by looking whether or not something is an azalea I may on occasion judge falsely on this matter in response to my current visual experiences. I may be distracted so that I am neither sufficiently attentive to the features of the plant nor sufficiently attentive to the factors that in more propitious circumstances, or with greater care on my part, would lead me to judge the other way, or suspend judgement. In such cases the aetiology of my judgement will be in important respects like the aetiology of the judgements I make when I judge rightly in response to current visual experiences. But, to repeat, when I do judge falsely, I do not exercise the recognitional ability in question.

3 THE STRUCTURE OF PERCEPTUAL-RECOGNITIONAL ABILITIES

Having a recognitional ability is in part a matter of being prepared to go through a certain judgement-forming procedure in response to suitable prompts and under suitable conditions. For instance, in the situation in which I aim to find out by looking whether or not the shrubs in the plot are azaleas, I look at them, I have certain visual experiences, and in the absence of countervailing factors judge that they are azaleas. The procedure is a highly reliable one in that by and large the judgements that are the outcome of its implementation are true. In the case in hand the implementation is prompted by my trying to determine whether the shrubs in the plot are azaleas. But the same procedure might be implemented as a result of any of a variety of causes. For instance, an especially fine specimen might attract my attention so that I notice that it is an azalea.

In the prevailing conditions I would very likely have judged that the shrubs were not azaleas if they had not been. In that event I would have gone through a judgement-forming procedure that would have involved judging that they are not azaleas in response to experiences of a different sort from those I have in the first scenario. So my recognitional abilities concerning azaleas come close to incorporating a decision procedure under conditions suitable for observation. Not only can I sometimes tell of an azalea that it is an azalea and sometimes tell of something that is not an azalea that it is not an azalea; I count as being able to tell by looking, in suitable conditions for observation, *whether or not* something I am looking at is an azalea. My ability only comes close to incorporating a decision procedure because even when conditions are suitable for observation I may not be able to tell. The shrubs I am looking at may be azaleas but untypical in appearance. I might be stumped if they had

been well burnt due to an out-of-control bonfire. In any case whether the experiences I have lead to a judgement either way, or whether they prompt me to suspend judgement, will depend on a complex set of sensitivities. These will reflect the training and experience I have had, which will have endowed me with relevant knowledge and skill at distinguishing azaleas from shrubs with which they might be confused.⁵

In my usual environment the judgement-forming procedures on which depend my recognitional abilities with regard to azaleas reliably yield true judgements. But we can imagine my being in an environment in which they are not reliable. Suppose that it comes to be the case that a high proportion of what look like azaleas are replicas skilfully made from plastic and silk so that even from fairly close to an expert could not tell the difference between the replicas and the azaleas just by looking.⁶ If I am in such circumstances but unaware of the presence of the replicas then in the face of azaleas or replicas I shall go through the same judgement-forming procedure that I did before the change in my environment, but it will not reliably yield true judgements. The lesson to draw here is that being in command of the judgement-forming procedures that in my actual environment underpin my ability to tell by looking whether or not this or that is an azalea is not sufficient for me to count as having that ability. Indeed, we must think of the ability as being indexed to suitable environments: the ability that I am imagining myself to have, as a horticulturalist in a normal environment, is an ability to tell by looking, of things in that environment and in environments suitably like it, whether or not they are azaleas. In the changed (abnormal) environment, in which I am often led to make false judgements, I lack the ability indexed to that environment, though, of course, I do not lose the ability that is indexed to normal environments.⁷ Even in a normal environment I may judge falsely of something that it is an azalea because unusually it looks just like an azalea but isn't.⁸ So there is an element of luck attached to whether or not the recognitional ability in question is exercised. It is not that it is lucky that going through the procedure and making a judgement as a result leads to a true judgement. Rather, whether or not I count as having exercised the ability depends on the absence of bad luck—the unusual presence of a look-alike.

On the account I have given, an ability to tell by looking whether or not an F is present is relational in this sense: it is constitutive of its exercise that it puts one in cognitive contact with Fs and with the fact that they are Fs, or with things that are not Fs and with the fact that they are not. This contact is such that the experiences I have and the judgements I make are causally dependent on features of the objects in

⁵ This imports a degree of indeterminacy into the notion of an ability to tell by looking whether or not something is an F. People may count as having such an ability relative to an environment in which the only azaleas to be found are in a very limited range, even if they would be unable to identify many sorts of azaleas in other environments in which azaleas outside the limited range are common. Others may have a more finely honed ability.

⁶ I invoke here an example of the fake barn type. See Goldman (1976).

⁷ That abilities are in part individuated by environments should come as no surprise. My ability to leap over a metre-high fence depends on the force of gravity not being much stronger than it is.

⁸ The situation here must be distinguished from that in which replicas abound. Here the environment may be one relative to which I have the ability in question. Where replicas abound, which I cannot discriminate from the real thing by looking, I do not have the ability.

question. The upshot of an exercise of the ability is that the subject knows of something that is present either that it is an F or that it isn't. But note that it is compatible with this that one could have such an ability and never have encountered Fs. Suppose, for instance, that I am undergoing training to identify certain aircraft in flight. I acquire the relevant ability by viewing realistic 3D animations of the aircraft but never see a real one. I pass the training by showing that I can come out with the right answer in response to the simulated presence of aircraft. When I do this I have not been exercising an ability to tell by looking whether or not such-and-such a type of aircraft is present. Since no such objects were present to me I have not exercised any such ability. But I have acquired an ability to tell by looking, in environments in which there will be aircraft, whether or not those at which I am looking are of this or that type.

It might seem that explaining why one gains knowledge that *p* in terms of the exercise of an ability to tell, and thus come to know, by looking, is a bit like explaining why people fall asleep in terms of their having taken a drug with the power to make someone fall asleep. Explaining how I come to know something in terms of the exercise of a capacity to come to know such things might seem no more illuminating than explaining why I fall asleep as a result of the activation of something with the power to make me fall asleep. Note first that even the *virtus dormitiva* explanation is not wholly lacking in explanatory power, since having taken something that makes one fall asleep is different from having naturally fallen asleep or having fallen asleep due to illness. But it sheds no light on what it was about the drug such that taking it makes one sleep. The explanation of particular cases of perceptual knowledge in terms of the exercise of an appropriate perceptual-recognitional ability is not like that. It is not a matter of explaining how knowledge is acquired by some power or other—we know not what—to acquire such knowledge. It is a matter of explaining, for instance, how one acquires knowledge that it's a chaffinch at the bird table through the exercise of an ability to tell that it's a chaffinch. For practical purposes we have a good enough idea of the nature of this ability. It involves being able to tell from the way the bird looks that it is a chaffinch. Thus it involves being responsive to the shape of the bird, its size, how it moves, and so on. All this is at the level of common sense. Reflection shows that the ability exercised on such an occasion implicates a reliable judgement-forming procedure. This further helps to make sense of why the upshot of the exercise of the ability is knowledge, since a case in which the procedure is triggered through an experience the having of which is explained by the presence of a chaffinch will be one in which the judgement made in response to the experience is true and depends on the chaffinch's being present in the way one would expect if the subject had recognized that the bird was a chaffinch. The ability in question, then, is not just an ability to acquire the knowledge in question, though we know not how; it is an ability that has a certain structure, which can be described at a common-sense level and at the more sophisticated level. At this more sophisticated level it is explicit that such abilities depend upon reliable judgement-forming procedures, but also that they are individuated in part in terms of a suitable environment.

4 HOW EXPERIENCES FIT INTO THE PICTURE

The picture I have presented is one on which exercising a recognitional ability is a relational matter. It puts one in cognitive contact with some object such that I know that it is such-and-such. Crucially for present purposes accepting the picture does not commit us to adopting the view that experiences are essentially relational.

To see this we need to consider how it is that in perception objects become available to us as things about which we can think. Part of the story must be that objects are visually picked out—visually discriminated or noticed against their background. On a traditional conception of experience we cannot account for visually picking out an object simply in terms of having a visual experience of an appropriate (non-relational) character since such an experience can be had when no object is picked out. Nor can we account for it just by adding a causal requirement. Adding such a requirement ensures that there is a real connection between the having of the experience and the presence of the object, but it does not to explain why we should think of the subject as picking out the object in virtue of that dependence; nor does it shed light on what kinds of causal dependence count.

On the account I favour visual discrimination is sub-doxastic in that its immediate upshot, in addition to the formation of visual experiences, is not belief or judgement, but a set of potentialities for behaviour, including sub-intentional behaviour. When I pick out people coming towards me in a corridor I am primed so that I shall make appropriate adjustments to avoid bumping into them. These adjustments are sub-intentional, though my intentionally walking along the corridor affects the kind of adjustments that are made. Similarly, when I return a ball in table-tennis, I track the path of the ball and move my hands in the way needed in the setting of the intentional activity in which I am engaged. Though I certainly intend to return the ball, I do not intend to make the precise movement of my arm that enables me to do so. Fortunately my body takes over.

It is crucial for visual discrimination that what is discriminated is not only reflected in the character of the experience but also shapes behaviour by installing a set of potentialities for behaviour directed on the object. Which potentialities are actualized depends on what one is doing. Which objects, and features of objects, are picked out depends on which objects and features of objects are doing the shaping.

The suggested picture might lead one to wonder whether objects necessarily shape behaviour *via the formation of experiences with a particular character*, or whether the shaping of behaviour could proceed in the absence of experiences. I think it is clear that for discrimination to be *visual* discrimination visual experiences must be in play. More importantly, it seems plausible that the having of visual experiences is essential for our ways of thinking about the scenes we encounter. We think about things as looking certain ways and our concepts of how they look are tied to their having a certain kind of experiential impact upon us. My way of thinking about magpies is tied to my conception of them as having a distinctive black-and-white appearance. My

way of thinking of that appearance is inextricable from my ability to call up images relating to it—images I would not have but for experiences gained looking at magpies or visual representations of magpies. A related point is that our ability to keep track of what we have been doing, or of changes in our environment, seems to depend crucially on the character of our perceptual experiences. As I write I have just been to the kitchen to return an empty mug. My complete assurance that that is what I have just done depends crucially on visual recollection—recollection that implicates a kind of edited replay of the visual experiences obtained in the kitchen that contributes to how I think of my having been there.⁹

The view I am presenting does not depend on the assumption that perceptual experiences—experiences obtaining in perceiving something—are essentially relational. So it is a view of which traditionalists about experience can avail themselves. The point of importance is that *visual discrimination* should be essentially relational. It satisfies this condition, according to the account, because visually picking out an object is a matter of the object controlling both the character of the experiences that are obtained and the potentialities for behaviour directed towards the object that are installed. Picking out the object is not, as on a crude causal theory of perception, just a matter of being experientially affected by the presence of the object; it is being in visual contact with the object in a manner that enables one to think about it in a distinctive way and to respond appropriately to its presence, depending on what one is doing. Essential relationality is captured at the visual level without assuming the essential relationality of the experiences implicated.¹⁰

Let us now go back to the side of the story that concerns *cognitive* contact. Let us imagine two scenarios—an epistemically good case and an epistemically bad case. These might be a case in which I am looking at an azalea and tell by looking that it is an azalea, and a case in which I am looking at a skilful artificial replica of an azalea and falsely judge it to be an azalea. On the traditional conception of visual experience, the experiences I have in those cases could be the same. That poses a problem, according to defenders of the relational conception. To their way of thinking if I know in the good case this must be because the experience has disclosed that the thing in question is an azalea. It will not have done that unless it enables me visually to pick out the object in question *and* tell that it is an azalea. But if the experience in the good case is no different from the experience in the bad case then, the argument goes, it cannot have enabled me to tell that it is an azalea. This is an important challenge, and one on which traditionalists have had little to say.

At the visual level there is no asymmetry between the cases. In both cases an object is picked out; in one case it is an azalea; in the other case it is a replica. Despite this there is an *epistemological* asymmetry. In the good case I have exercised the ability to tell by looking that a certain object is an azalea. In the bad case I have not exercised that ability. Of course, in that case I shall have been through a judgement-forming

⁹ The notion of edited replay figures prominently in a discussion of self-knowledge in Lyons (1986).

¹⁰ One can of course acquire perceptual knowledge to the effect that an object is absent from a certain location. In that case the discriminative relation is to the location.

procedure like the one that I go through in the good case. The difference is that in the bad case I have been unlucky; I do not count as having exercised the ability because my judgement is false and thus not a case of telling that the object is an azalea. We should not infer that I have therefore been lucky in the good case. In both cases things that look like azaleas usually are and, given my ability, I could not have easily formed the judgement in the way I did and be wrong. It's just that in the good case I was right and in the bad case I was wrong.

What about the asymmetry between the good case and a corresponding Gettier case in which I form a true judgement from looking at a life-sized photograph? In the good case there is a relation of causal dependence between the experience and the presence of the object and a relation of causal dependence of the judgement upon the experience, and thus indirectly on the presence of the object. In the corresponding Gettier case there is no causal dependence of the experience on the presence of the object, nor therefore of the judgement upon the presence of the object. Invoking causal dependencies here is not *ad hoc* because the account of visual-discriminative and recognitional abilities makes sense of the requirement of causal dependence. The visual-discriminative ability is exercised only if the experience is causally dependent on the presence of the object and the recognitional ability is exercised only if the judgement is causally dependent on the experience, and thus indirectly on the presence of the object. This gives us a constraint on the relevant kinds of causal dependence. They have to be compatible with the exercise of the relevant ability.

On the account I am outlining it is conceded that the experiences gained in perception do not suffice to account for our being in visual or cognitive contact with the objects we see. Rather it is experiences gained in circumstances in which one has exercised the relevant visual-discriminative ability and the relevant recognitional abilities. Exercises of these abilities are essentially relational because they put subjects in touch with objects, events, or scenes. In the case of discrimination the contact, to keep with the metaphor, is visual. In the case of recognition it is cognitive. As I stressed earlier, it is agreed on all sides that perception is relational. The issue has been whether the experiences implicated in perception are essentially relational. I have sought for relationality elsewhere.

5 COMPARISON WITH TRADITIONAL RELIABILISM

Traditional reliabilist accounts of perceptual knowledge have it that perceptual knowledge is true belief formed through the implementation of a reliable belief-forming procedure. In making reliability central such accounts are to that extent correct. But my account differs from traditional reliabilism in a number of respects.

(i) Traditional reliabilism is an attempt to explain what knowledge is in terms of true belief plus something else. I make no such attempt. I explain particular instances of perceptual knowledge in terms of the exercise of perceptual-recognitional abilities,

but the abilities are characterized epistemically—they are abilities to come to know—by perceiving.¹¹

(ii) Under traditional reliabilism the implementation of a reliable procedure is not a success notion. The procedure might be implemented and yet the subject not acquire knowledge; something that is not an F could produce the kind of experiences that, by the procedure, result in a judgement that an F is present. As I have stressed already the exercise of a perceptual-recognitional ability *is* a success matter.

(iii) Traditional reliabilism is problematic if presented as a partial analysis of our ordinary concept of perceptual knowledge. Such an approach suggests a picture of how we ascribe perceptual knowledge on which such ascriptions are ultimately grounded in evidence that the subject has a true belief resulting from the implementation of a reliable procedure.¹² One problem with this is that ordinary thinking about perceptual knowledge does not supply us with a determinate conception of the procedures on which our perceptual recognitional abilities depend. Reflection shows that they must take us from experiences of certain sorts to judgements of a corresponding sort. But we have no idea how to specify with any precision which experiences, against which background of beliefs and sensitivities, are implicated. Despite this, we are often able to tell when people know by looking that an F is present. We do so, for instance, when we see that an F is present and are able to tell from a subject's reactions that he or she has noticed the F and recognized it to be an F. Suppose that I am out with a friend walking his dog. The dog has disappeared among some bushes but emerges. My friend calls out, not in the way he would if trying to locate it, but in the way he would if he had just seen it. Seeing both the dog and him I take in that he is calling to it, not calling out for it. In this situation my judgement that my friend has recognized the dog is not based on any evidence to the effect that, by and large, he forms true beliefs that his dog is present when his dog is in clear view. It is based (a) on the knowledge that my friend knows his dog, and can recognize it by sight, and (b) on my knowledge that on this occasion he has just noticed and recognized it by sight. My knowledge that he knows his dog might be derived simply from seeing the two interact on this occasion. This and my knowledge that he has picked out and recognized his dog by sight are arguably the outcomes of the exercise of complex recognitional abilities, rather than drawn from prior assumptions about the present situation, viewed in the light of my friend's history of acquiring true beliefs as to the presence of his dog. In both cases I tell from visually manifest cues. I tell from the layout of the environment and from the orientation and behaviour of my friend that he has seen and recognized his dog. Telling from these factors should not be conceived as a matter of inferring from assumptions about the layout and about the behaviour and orientation of my friend. I would be hard put to capture by description the indicators

¹¹ I am sympathetic to those who are sceptical about the possibility of analysing knowledge into true belief plus something else. See, especially, Williamson (2000). Compare McDowell (1993/1998; 1995/1998).

¹² I am assuming here that an adequate conceptual analysis would reflect the constraints that in practice govern our dealings with the concept in question. So it would be a problem for such analysis if its import for how we are in practice governed were psychologically unrealistic.

of his having spotted the dog, just as I would be hard put to capture by description the features of a familiar face on the basis of which I recognize whose face it is. The model we should work with is recognition rather than complex reasoning from evidence. This helps to explain how my engagement with my friend's cognitive state can be at the level of his visual-cognitive contact with the dog, which is to say, at the level at which success notions like seeing and recognizing are in play. Reflection on the situation in which the seeing and recognizing have been achieved brings into play the idea that a reliable judgement-forming procedure must have been implemented. But the order of knowledge is from recognition that there has been visual recognition to the conclusion that a reliable judgement-forming procedure has been implemented, rather than vice versa. (I discuss related points below in section 8.)

Recognitional abilities, as I conceive them, bear some resemblance to what are sometimes called *intellectual virtues*. For Ernest Sosa, to have an intellectual virtue relative to a certain environment, with respect to a field of propositions and set of conditions, is to be so constituted that if one is in that environment, in the relevant conditions, and believes or disbelieves some proposition within the relevant field, then one will very likely be right (Sosa 1991: 284). Like recognitional abilities, virtues in this sense are tied to specific subject-matters, and one counts as having them only if one is in a suitable environment. But the exercise of a virtue is not a success notion. One can exercise the virtue in the right environment and in the right conditions and believe or disbelieve wrongly.¹³

There is an obvious kinship between my account of perceptual recognitional capacities and Alvin Goldman's well-known relevant alternatives account of perceptual knowledge.¹⁴ Goldman takes us to have perceptual knowledge that *p* if, roughly, we are caused to believe that *p* non-inferentially through having an appropriate experience in circumstances in which there is no relevant alternative to its being the case that *p*. My own account captures what I take to be a key insight contained within this account. It is the actual environment that determines what relevant alternatives are. If the environment does not throw up things that are not *F*s but which look just like them, or does so only rarely, then a subject can know by looking that an *F* is present without taking steps to determine that what is present is not a look-alike. However, the full definition of perceptual knowledge that Goldman provides is so complex that

¹³ Linda Zagzebski (1999) defines knowledge as belief arising out of acts of intellectual virtue. Acts of intellectual virtue are so understood that a belief does not count as arising from such an act unless it is true. So there is some kinship between my notion of exercising a recognitional ability and her notion of believing from an act of virtue. But, as I understand it, on her view a subject may form a belief virtuously even though the belief is false. In such a case the belief does not count as arising from an act of virtue for the notion of such an act incorporates a success component that the notion of acting virtuously does not. In my scheme exercises of recognitional abilities are by definition successful. A distinct position from Zagzebski's is that in Greco (1999), which builds on work by Sosa. In general I find the account in terms of recognitional abilities more in keeping with common-sense classifications than that in terms of virtues. There is no denying the affinities with virtue theory, however.

¹⁴ See Goldman (1976). Reflection on the fake barn case, which is so central to Goldman's account, undoubtedly contributed to the way I represent recognitional abilities and their relation to reliable judgement-forming procedures.

it is hard to see how it could govern our actual applications of the concept of knowledge. From the standpoint taken here that is a defect. We want an account of knowledge to enable us to understand the role of knowledge and of ascriptions of knowledge in our social interactions.

I do not think there is much doubt that our most natural ways of thinking about perceptual knowledge implicate recognitional abilities as I conceive them. But why should that be? Why should success notions be so crucial to our thinking? At least part of the answer has to do with why knowledge matters to us. I address this below in section 7. Before that I consider how perceptual knowledge as understood here links up with notions like that of justified belief. This is important because without some discussion along those lines the suggested approach to perceptual knowledge might seem to avoid the issues that have preoccupied traditional epistemologists.

6 JUSTIFICATION

When I look at the shrubs and tell that they are azaleas I come to be in a position such that I am justified in being sure that they are azaleas. How should we think of this justification? Should we think of it as explaining how I know or should we think of it as deriving from the fact that I know?¹⁵ The latter, I suggest, is the right way to think of the matter. The explanation for why I know is simply that I have exercised a way of telling—specifically, a perceptual-recognitional ability to tell that the things in question are azaleas. That said, at least in typical situations when we know by looking that something is an F we are aware, that is know, that we have seen the thing to be an F. In those situations the visual experiences that, via the exercise of the relevant perceptual recognitional ability, enable us to know that something is an F, also enable us to tell that we see that thing to be an F. This latter knowledge is made possible by the exercise of a higher-order recognitional ability—an ability to tell whether or not one sees that an F is there. This latter ability is not an introspective ability, for there is no inner scrutiny of anything. Just as I have learned to tell when something is an azalea in response to certain visual experiences, so I have learned to tell when I see an azalea in response to *the very same experiences*. I do not think about or introspect my experiences in the case in which I judge that I see azaleas any more than I do when I tell that the shrubs are azaleas by way of response to the same experiences. I go through a procedure that takes me from experiences to a judgement that I see that the shrubs are azaleas. This procedure is constitutive of an ability to tell, in a suitable environment and under suitable conditions, whether or not I see that certain things are azaleas. The upshot of the exercise of such an ability is that I become apprised of a relational fact about myself and about the azaleas, and not as a result of any inference.¹⁶

¹⁵ For a sharply focused exploration of related issues, see Williamson (2000).

¹⁶ I have reflective access, as one might say, to that in virtue of which I know that the shrub is an azalea, but it is not restricted to modes of access—introspective and a priori—recognized under standard internalist theories of knowledge or justification.

Let us suppose then that I am looking at the azaleas and recognize them for what they are. I am sure that they are azaleas. What justifies this assurance? A natural answer is that what justifies the assurance is my knowing that I see that the shrubs are azaleas. Were I to explain why I know that the shrubs are azaleas it would be natural to cite the fact that I can see that they are. But it is also true that were I to justify thinking that they are azaleas it would be natural to do so in terms of this same fact.

The view I am presenting links being justified with having reasons in an entirely natural way.¹⁷ It is because I know that I see that the shrubs are azaleas that the consideration that I see that they are azaleas can be a reason that I have to believe that they are azaleas. And because I will retain in memory that I have seen that the shrubs are azaleas, the consideration that I have seen that they are can serve as a reason to (continue to) believe that they are.

The preceding remarks naturally lead us to consider whether I am guaranteed to have a good reason to believe that the shrubs are azaleas simply through having exercised the relevant recognitional ability. Is it built into my having exercised the ability that I will know that I have seen that the shrubs are azaleas and on that account will be justified in believing that they are? There is no question that at least typically I would know both that the shrubs are azaleas and that I see that they are, and on that account be justified in believing that they are. But the view I am presenting does not commit us to supposing that perceptual-recognitional knowledge that *p* guarantees that one is justified in believing. Nonetheless, it is important for the use to which we put the concept of knowledge that we can in perceptual cases readily tell that we know. I pursue this further in the next section.

More traditional approaches to the theory of knowledge accommodate the idea that one's reason for believing that *p* might be that one sees that *p*, but the availability of the consideration that one sees that *p* has to be explained in terms of reasons for thinking that one sees that do not entail that one sees. The natural way to do this is in terms of the consideration that it is to one as if one sees that *p*, or some such thing. That raises the question of what entitles one to think that its being to one as if one sees that *p* indicates that one sees that *p*. I suspect that what has made such approaches seem compelling is the thought that justification for the application of success notions like seeing that *p* cannot derive from simply recognizing that one sees that *p*. It is crucial for the view I am presenting that such recognition is commonplace.

7 WHY KNOWING MATTERS

My focus on knowing through the exercise of perceptual-recognitional abilities fits neatly with a plausible conception of why knowing matters in a wide range of social interactions. It can matter to me that I know of the shrubs I am looking at that they are azaleas just because I am called upon to vouch for whether or not they are by my

¹⁷ Compare McDowell (1993/1998) at this point.

client, who wants to know from me. Her confidence in what I tell her is grounded in confidence that I can find out and have done so. My confidence in what I vouch for is grounded in my confidence that I have seen what the shrubs are.

Knowledge enables one to have well-grounded assurance. We can have well-grounded assurance that p in view of the fact that we see that p or have seen that p , and on that account are justified in believing that p . One reason why we value well-grounded assurance is that it enables us legitimately to vouch for something's being so. When we vouch for something's being so we give it to be understood that we have the kind of assurance that is grounded in knowledge. I do not deny that we often say what we think is true when we do not know, and do not take ourselves to know. I do not regard such saying as vouching. We vouch for something's being so in contexts in which we purport to convey knowledge, for instance, in contexts in which someone wants or needs to know something and we wish to oblige them. It is part of the practice of vouching that you should not vouch unless you know.

Another reason why we value the assurance that is grounded in knowledge is that we need our true beliefs not merely to be true but to be secure (not easily vulnerable) as well. We do not want our true beliefs to be like the statues of Daedalus, which slip away when untethered (Plato, *Meno* 97d–e). Suppose that I believe correctly that I have sent off a letter of recommendation for a student, but I do so, not because I remember sending it, but on the grounds that a colleague has told me that I told him that I had sent it just after doing so. But my colleague is confused. *I* did not tell him I had sent off the letter. Someone else had told him that *he* (this someone else) had sent off a letter of recommendation. When my colleague spoke to me he thought it was I who had told him. As things stand my only basis for thinking that I have sent off the letter is his false report. So my belief is liable to be given up if I learn that my colleague was confused, for then I would realize that I don't know. But if I distinctly remember putting the letter in the mail tray it would take a lot more than this to dislodge my belief that I have sent it.¹⁸

It can be practically important that we do not give way to doubt with respect to a true belief. My true belief that I sent the letter was easily abandoned in the face of the realization that my colleague's report was false. I then had to make time-consuming enquiries to settle whether I did or did not send it. Of course, it could be that a true belief that p that is not knowledge that p is not abandoned because the factor that prevents the subject from having knowledge that p never comes to light. That does not tell against the importance of knowledge as opposed to merely true belief. Although merely true belief *might* do as well as knowledge, it is liable not to because of the possibility that the knowledge-preventing factor does emerge.

Knowledge is instrumentally valuable then, not just because the implicated true beliefs are instrumentally valuable in virtue of being true. The instrumental value of knowledge, over and above the instrumental value accruing to the implicated true

¹⁸ The example is similar to Williamson's burglar example (2000: 62). I think it is important that the knowledge plays its characteristic explanatory role in virtue of the subject's recognition that he knows in a certain manner.

beliefs in virtue of their being true, lies in the value of having well-grounded assurance that renders our beliefs secure.

These considerations about why knowledge matters help to explain why success notions like seeing-that are so basic to our thinking about knowledge. We want knowledge when we wish to be assured. We come by assurance as readily as we do because we can so readily tell when we, or others, know something in a wide range of matters that are relevant to our practical affairs. My seeing that the shrubs are azaleas provides me with well-grounded assurance that I know that they are and can thus vouch for this being so. I can *readily* tell that I see and thus can readily tell that I know. As I have already suggested, the very same experiences that figure in my knowledge of the first-order worldly fact that the shrubs are azaleas can also engage a higher-order recognitional ability, resulting in my knowing the relational fact that I see that the shrubs are azaleas. My client's assurance that the shrubs are azaleas derives from her knowledge that I can tell by looking in such matters and have done so on this occasion.

If the concept of knowledge were governed by a more or less complex definition in terms of necessary and sufficient conditions then knowledge would be more like a theoretical posit than it is, at least when perceptual knowledge is at issue. Applying the concept knowledgeably would require that one should know that the conditions for its application have been satisfied. Since it would not be at all clear that we have ready access to whether or not the conditions are satisfied it would be harder to understand how knowledge can play the role it does in interactions such as those I have described.

A related point is that it is important to distinguish between, on the one hand, assurance that *p* deriving from personally seeing that *p*, or knowing that someone else, who was in a position to see that *p*, saw that *p*, and, on the other hand, having a high degree of confidence that *p* deriving from, say, inductive evidence. I may have a high degree of confidence that if I go to a certain café at a certain time I shall see some friends there at that time. My confidence may be grounded in the knowledge that they not only go there regularly at that time but that it is something of a ritual for them to do so. But despite the ritual character of their behaviour their attendance is not perfect; things can crop up that prevent their going. So I do not take myself to *know* that they will be there. I might tell others that it's a good bet that they will be there but I would not vouch for their being there. But if I am at the café with them I could vouch for their being there to any interlocutor with whom I happened to be in telephone conversation. In this latter situation it is not merely that the level of my confidence is sufficiently ratcheted up to make it reasonable to vouch. There is a difference in kind and not merely degree between how things stand with me when I know through seeing that something is so and how things stand when I have good evidence that makes it very likely that something is so. So it would be odd if knowledge could be had in virtue of having a true belief based on evidence that merely made it more or less probable that *p*.¹⁹ Given that available evidence so often does no more than make it more or less probable that something is so, it might be tempting to

¹⁹ This is an important strand in the thinking of McDowell (1982). It is picked up and developed by Travis (2005), who traces it back to Cook Wilson.

think that we would be better not bothering about knowledge and settle for probability. For large tracts of our beliefs this might well be the right approach to take. But our social interactions require us to discriminate between situations in which we, or others, can vouch for something's being so and those in which we, or they, are at best in a position to be reasonably confident.

The foregoing takes us to a further point of contrast with traditional reliabilism about knowledge. Consider a case in which we do not have independent access to whether or not *p*, but wish to tell whether or not somebody knows by looking that *p*. The ascription of knowledge under reliabilism in such a case would be based on the assumptions that a reliable procedure has been implemented and that the chance that the belief is true is, therefore, high. That is not how in practice we think of knowledge in perceptual cases.²⁰ The typical situation in which we ascribe perceptual knowledge that *p* to others, in circumstances in which we do not ourselves know that *p*, is one in which we rely on the following assumptions: (i) that the speaker is in a position to tell whether or not *p* by perceiving, and (ii) that what she tells us on the matter in this situation will be true because it is based on what she has perceived to be so. We reach the conclusion that the speaker knows that *p*, not on the basis of assumptions that if true would raise the chance that what the speaker tells us is true, but on the basis of our having already recognized the speaker as one who knows whether or not *p* and would speak truthfully on the matter. Of course, if this picture is to contrast with one in which a hazardous leap is made to the conclusion that the speaker knows, from something that makes it highly probable that the subject knows, the various assumptions involved in recognizing the speaker as one who knows whether or not *p* should not themselves result from hazardous leaps from evidence. I think that this condition can be met, though the matter demands much fuller treatment than I can give here. Still, it seems fairly clear that we can know that people are in a position to tell whether or not *p* by looking because we know that they are located so as to be in such a position (for instance, beside the mains switch box). And we can know that they would not vouch for its being the case that *p* unless they had perceived that *p* because we know them and can recognize when the situation is one in which what they vouch for is to be relied upon (for instance, whether or not the switch is off).

My claim is that perceptual knowledge can play the role that it does in social interaction because there is a basic level of engagement with the perceptual knowledge of others at which we can tell that they know that something is so as a result of a chain of knowledge-transmission of the sort I have described. At this level we think of subjects as being or having been in perceptual-cognitive contact with the fact. We do not think of their coming to believe that *p* as resulting from the operation of a reliable process that merely raises to a high degree the chance that it's true that *p* and then make a further leap to the conclusion that they know. This is not to deny that

²⁰ The same may be said of cases in which knowledge that *p* is based on some indicator that *p*—a fact that would not obtain were it not the case that *p*. To make this plausible would take us too far from the current topic. It requires that we extend the conception of a recognitional ability to embrace cases in which a fact that *p* is recognized to obtain through seeing that something else is so. A key question would be how to make sense of this without imputing to the subject an inductively justified belief in a covering generalization.

in situations in which subjects tell that p because they see that p a reliable, yet fallible, judgement-forming procedure will have been implemented. On the account of perceptual-recognitional abilities proposed here such a procedure will have been implemented. The point, to repeat, has to do with the level at which we conceptualize what is going on when knowledge is had in perceptual cases.

REFERENCES

- Campbell, J. (2002) *Reference and Consciousness*, Oxford: Clarendon Press.
- Child, W. (1994) *Causality, Interpretation and the Mind*, Oxford: Clarendon Press.
- Goldman, A. (1976) 'Discrimination and Perceptual Knowledge', *Journal of Philosophy*, 73: 771–91.
- Greco, J. (1999) 'Agent Reliabilism', *Philosophical Perspectives*, 13: 273–96.
- Lyons, W. (1986) *The Disappearance of Introspection*, Cambridge, MA: MIT Press.
- McDowell, J. (1982) 'Criteria, Defeasibility and Knowledge', *Proceedings of the British Academy*, 68: 455–79. [Reprinted in McDowell 1998: 369–94. Page references are to this edition.]
- (1993) 'Knowledge by Hearsay', in B. K. Matilal and A. Chakrabarti (eds.) *Knowing from Words: Western and Indian Philosophical Analysis of Understanding and Testimony*, Dordrecht: Kluwer. [Reprinted in McDowell 1998: 195–224. Page references are to this edition.]
- (1995) 'Knowledge and the Internal', *Philosophy and Phenomenological Research*, 55: 877–93. [Reprinted in McDowell 1998: 395–413. Page references are to this edition.]
- (1998) *Meaning, Knowledge and Reality*, Cambridge, MA: Harvard University Press.
- Millar, A. (2007) 'What the Disjunctivist is Right About', *Philosophy and Phenomenological Research*, 74: 176–98.
- Snowdon, P. (1990) 'The Objects of Perceptual Knowledge', *Proceedings of the Aristotelian Society*, Supplementary Volume 63: 121–50.
- Sosa, E. (1991) *Knowledge in Perspective*, Cambridge: Cambridge University Press.
- Sturgeon, S. (2000) *Matters of Mind: Consciousness, Reason and Nature*, London: Routledge.
- Travis, C. (2005) 'A Sense of Occasion', *Philosophical Quarterly*, 55: 286–314.
- Williamson, T. (2000) *Knowledge and its Limits*, Oxford: Oxford University Press.
- Zagzebski, L. (1999) 'What is Knowledge?', in J. Greco and E. Sosa (eds.) *The Blackwell Guide to Epistemology*, Oxford: Blackwell.

Starting Afresh Disjunctively: Perceptual Engagement with the World

Sonia Sedivy

The primary aim of this essay is to elaborate a conceptualist direct realism that strengthens the appeal of taking a disjunctive approach to perception. Conceptualist direct realism holds that perception is a *sui generis* capacity characterized by *sui generis* content: world-involving or fully engaged, genuinely singular, conceptual, determinate. More simply, perceptual experiences are contentful episodes that are distinguished by content that is both conceptual (individuated by conceptual modes of presentation) and determinate, involving its objects and properties. One might say that such contents are both Fregean and Russellian—world-involving yet individuated not only by the worldly facts (or correctness conditions) but also by a perceiver's grasp of the requisite facts. This is an approach to content that follows through on Gareth Evans's (1982) and John McDowell's (1986) groundbreaking work in bringing the best of Frege and the best of Russell together in our explanation of perceptual contents.

It is not new to put some version of direct realism together with a disjunctive approach. Indeed the two typically come together. Both aim to show that perception is a world-involving capacity. Direct realism provides the substantive theoretical framework while the disjunctive insight supplies a key integral component: the insight that the category of how things seem to a subject is an essentially disjunctive category that subsumes both engaged and disengaged states. The disjunctive claim concerns an uncontroversial category of common sense about which we need not make heavy weather—in so far as we can recognize that the category is disjunctive and subsumes two or more categories of mental capacities and their respective kinds of contents. That is, in so far as we recognize that when it comes to how things seem

I would like to thank friends at the University of Toronto—Ronnie de Sousa, Tom Rand, William Seager, Joel Walmsley—and the participants of the 'Disjunctivism: Perception, Action, Knowledge' conference at the University of Glasgow for illuminating discussion of an earlier version of this essay. Special thanks to William Seager, Thomas Baldwin, my editors, Adrian Haddock and Fiona Macpherson, and two anonymous referees for very helpful critical comments. This essay has a companion, or better yet, sister essay with which, like a Siamese twin, it is joined at the hip. The two essays present a package, joined in the account of conceptual, object- and property-involving content that I allude to here but that is developed in more detail in my (2006).

to a subject, she might be perceptually engaged with the world *or* in an illusory state without being able to tell the difference. The intended upshot is that the possibility that a perceiver might be unable to distinguish a perceptual experience from an illusory state does not undermine the fact that in favourable circumstances she is perceptually engaged with her surroundings. The latter is a fact that we— theorists and ordinary perceivers alike— can understand about perception even if, on occasion, any one of us might be unable to distinguish a perceptual experience from a state that mimics engagement.

The disjunctive insight is important because it allows theory of perception to work out a detailed account of perceptual experience as a *sui generis* capacity, and of perceptual content as a *sui generis* kind of content. Theory of perception can proceed with these tasks in so far as it can sidestep the regression to a highest common factor between perceptual experiences and sensory episodes from which perceptions might seem indistinguishable from the first-person perspective. Highest common factor approaches to perception have classically held that because a person might fail to distinguish an illusory episode from a perceptual experience, perceptions yield no greater epistemic warrant than illusions, and this is explained by casting the two capacities as having a common kind of content or common sensory aspect. Contemporary variants diverge along several dimensions, but it is beyond the scope of this essay to address all variants explicitly. Rather, the task here is to present one coherent package—a conceptual direct realist theory of perception that incorporates the disjunctive insight to mutual advantage.

The task is posed by the fact that direct realism and the disjunctive proposal face serious hurdles. Direct realism needs to show that the nature of perceptual experience or “the fabric of [perceptual] consciousness” itself indicates that such experience is relational.¹ Put more in my terms, direct realism needs to show that perception is a *sui generis* relational capacity and that perceptual content is genuinely singular—so that perceptual content could not be just as it is if the world were not involved and if the world were not as it is represented as being. The second and related hurdle is metaphysical. In so far as the metaphysical commitments associated with our scientific picture of the world suggest that the world is not as we perceive it—lacking the technicolour splendour of our perceptions, for example—direct realist theories of perception seem blocked at the start. Direct realist accounts need to scale the metaphysical hurdle and show that their non-objectifying explanations of our selves and the world are compatible with the developing scientific picture. To do so, direct realism needs to explain how it is that the mind-independent individuals we perceive have the properties that we perceive.

This essay will show how *conceptualist* direct realism can answer these principal objections. I will provide a novel argument showing that the “fabric of [perceptual]

¹ See, for example, Michael Martin’s (2004) claim that a “naïve realist” needs to offer reasons to show why we must think of the fabric of consciousness as relational, and therefore not common to perception and hallucination. My argument from perceptual determinacy from the first-person perspective gives those reasons, though the conceptualist direct realism differs in significant respects from the position that Michael Martin labels “naïve realism”.

consciousness" is relational and I will provide an extended two-part response to the metaphysical objection. These arguments aim to show that the disjunctive insight attains its fullest theoretical potential, offering the best explanation of perception, if it is located within a conceptualist form of direct realism.

But I would like to be able to argue in a non-standard way. I would like to take an explanatory approach to perception that is not answerable to sceptical pressure as a precondition at the very outset. Sceptical worries place theory of mind under constraint from epistemology, prioritizing epistemic concerns. In contrast, an explanatory approach may proceed in theory of mind and theory of knowledge simultaneously, developing accounts of the varieties of mental content that work with both third- and first-person facts and show how various kinds of mental contents can play the justificatory or rationalizing roles required of them. One important consequence is that explaining illusory and hallucinatory states does not come up at the outset but rather in the course of developing an explanation of perception to deal with peripheral and limiting cases. Another important feature is that an explanatory approach can be inclusive and non-objectifying. It may include different kinds of explanations without prioritizing facts available to either the third- or the first-person perspectives. And it can strive to show how the non-objectifying considerations that conceptual realists work with are not in tension with scientific explanations.

I will interweave such an explanatory approach within the principal line of argument (which develops a disjunctive conceptual realist account). Though this may seem paradoxical, I believe that taking an explanatory approach addresses rather than evades the third and perhaps most unyielding source of discomfort with direct realist, disjunctive approaches. A disjunctive approach to perception does not show that an individual perceiver can distinguish veridical perceptions from illusory states. It does not follow through on the sceptical trump that gives primacy to the sceptical problem—and that implies that going ahead with theory of mind and metaphysics without a solution to the sceptical challenge is less than adequately warranted. The worry that needs to be addressed is whether a disjunctive approach can legitimately allow us to sidestep the well-worn concessive paths into which we are steered by arguments 'from illusion' given the fact that a disjunctive account does not answer the sceptical challenge as posed.

Suppose we interrupt the usual dialectic briefly to pose a simple-minded question that asks us to suppose something that runs evidently contrary to the historical facts. Reverse the order between the sceptical challenge to perceptual experience and the disjunctive insight. Suppose we had not thought of sceptical challenges to perceptual experience prior to recognizing that the fallibility of perceptual experience suggests a disjunctive account of perceptual experience? What if the disjunctive insight had come first?

Now consider our historical trajectory. Once we have the disjunctive insight at our disposal as a theoretical resource, it does change our sense of the issues. Not in the sense that we think we can answer sceptical challenges directly, but in the sense that we have something like a counterfactual insight about the nature of the problem posed by the fallibility of perception. I am suggesting very simply that: if we already had the disjunctive insight, the sceptical challenge would not have had its

pre-emptive force, its ability to prioritize epistemology over theory of mind and metaphysics. This is because the disjunctive insight allows us to take the fallibility of perception in our stride in the course of developing the best explanation of the nature of perception, rather than at the outset as a precondition for theorizing about perception (or the world) at all. Insights can feedback and supplement one another and the feedback may even be of the strongest, constitutive kind where insights are 'co-constituted' through their feedback effects on one another. The relationship among direct realism, the disjunctive insight, and conceptualism seems to me to be of this mutually co-constitutive kind. That is why I believe that setting out from an explanatory approach—that does not prioritize dealing with issues of perceptual fallibility at the first-person perspective—does not simply disregard the fundamental discomfort but engages with it.

Moreover, undertaking an explanatory approach relies on—without being able to rehearse—much fairly recent argumentation which suggests that: (i) we should not prioritize the first-person perspective at the expense of third-person facts (and vice versa); and (ii) presuppositions concerning the mind have often been at work in epistemic theories even when the epistemic issues seem to force substantive conclusions about the mind and its contents. In other words, taking an explanatory approach in the face of sceptical pressure is not simply historicist insouciance because philosophical arguments, much like artworks, depend on their historical context. Just as Marcel Duchamp could not have offered a 'ready-made' to the Sun King or at a Paris Salon of the nineteenth century, so it is possible to start theory of perception from an explanatory perspective only once twentieth-century philosophical argumentation is in place. Now the task is to examine how illuminating might be the results. And the emphasis needs to be placed on the need to avoid objectifying. That is the heart of conceptualism.²

² See for example, McDowell's view that the problems in theory of mind and knowledge that need to be redressed with a disjunctive approach to perception are posed by "an objectifying conception of the human". This suggestion comes at the end of his initial disjunctivist proposal in 'Criteria, Defeasibility, and Knowledge' (1982) and is elaborated in detail in his later 'Singular Thought and the Extent of Inner Space' (1986): "I want to end by mentioning a source for the attraction of the 'highest common factor' conception that lies, I think, as deep as any. If we adopt the disjunctive conception of appearances, we have to take seriously the idea of an unmediated openness of the experiencing subject to 'external reality' . . . No doubt there are many influences that conspire to give the opposing picture of the 'inner' and 'outer' its hold on us. The one I want to mention is that we are prone to try to extend an objectifying mode of conceiving reality to human beings. In an objectifying view of reality, behaviour considered in itself cannot be expressive or significant; human behaviour no more than, say, the behaviour of the planets. If human behaviour is expressive, the fact resides not in the nature of the behaviour, as it were on the surface, but in its being the outwardly observable effect of mental states and goings-on. So the mind retreats behind the surface, and the idea that the mental is 'internal' acquires a quasi-literal construal, as in Descartes, or even a literal one, as in the idea that mental states are 'in the head' . . . And it is hard to see how the pictured interface can fail to be epistemically problematic in the outward direction too; the inward retreat of the mind undermines the idea of a direct openness to the world, and thereby poses the traditional problems of knowledge about 'external' reality in general. . . . Traditional epistemology is widely felt to be unsatisfying; I think this is a symptom of the error in the 'highest common factor' conception, and more generally, of the misguidedness of an objectifying conception of the human" (1982/1998: 392–4; last italics mine).

1 FANCIFUL OR SENSIBLE?

Let's start by imagining that interstellar travellers might visit earth one day. Considering the way in which they would understand us shows the most straightforward or best explanation of ourselves.

Observing the various life forms, with their various strategies for survival—such as staying rooted in one place or moving about—the intergalactic visitors would find mobile life forms intelligible as follows. The mobile ones survive by acting on their environment and they are able to act on it by means of various sensory and perceptual modalities. The more complex ones—the ones whose interaction with their environment is quite complex (the birds and higher animals, let's say)—are able to behave in complicated ways because they are perceptually and behaviourally engaged with the individuals and with the determinate properties of those individuals that are relevant in their survival strategies. It is clear to our visitors that a complex animal can only behave successfully in its environment if it is engaged with the individuals and determinate properties that its complex behaviour concerns.

Suppose that our intergalactic visitors were fortunate enough to touch down on an African savannah with its wondrous variety of animals, interlocked in the complex daily dance of the food chain. Those who will be eaten, as well as those who will do the eating, walk among the tall grasses. Everyone is careful, moving slowly, fully on the lookout. Stripes and spots are everywhere. The tall grass stripes the already striped bodies of the animals that are on the menu, while the light falling through the leaves and tall grass introduces yet more stripes and spots into the scene. Everything blurs and shimmers slightly with the rising heat. All parties need to note *any* unusual movement of potentially *any* of the myriad stripes or spots, however slight, *as* it occurs—so as to burst into chase or flight. How is this achieved? Could anything short of engagement with their habitat allow the majority of the herd animals to avoid their predators? Could the predators succeed in catching a prey if they were less than engaged? Yet what exactly is it for an animal to be experientially engaged with its habitat?

Not only would intergalactic visitors explain animals as engaged with their habitat, they would also explain engagement as a distinctive world-involving content. To understand why this would be the best explanation it is important to be clear about the alternatives. The straightforward alternative to engagement is that an animal's perceptions have a content that does not involve the represented individuals and *properties* so that their perceptual content could be just as it is while the world is quite otherwise, and that behaviour is initiated by such content—we might call representational content of this kind 'self-standing'.³ There is also another alternative,

³ Intentionalism is a key contemporary theory that holds that perceptual content does not involve properties so that the detailed "phenomenal character" of perceptual experience is abstract rather than determinate, and non-conceptual. Michael Tye's intentionalist theory (1995: 138–40) for example, holds that the "phenomenal character" of perceptual experience is abstract rather than determinate, which means that "no particular concrete objects enter into these contents... What is crucial to phenomenal character is the representation of general features or properties". What it

which holds that animals are engaged with their habitat, but proposes to explain engagement in terms of a content that can be common to illusions or hallucinations such as an existentially quantified content. Neither of these two alternatives can explain continuous update in real time. And this is a primary condition on a theory of perception if one begins, as intergalactic observers would, by striving to explain the role of perception in the life of complex animals. World-involving representational contents update 'for free' along with changes in the world. That is why such contents stand out as the best explanation of perceptual engagement from the outset (so long as one starts from a perspective that focuses on the real-time conditions of animal life). The problem for representational contents that are not world-involving—contents that can be common to hallucinations and illusions—is that they need to be organized in a manner that allows for efficient access for update in real time. (In our history, this problem became vivid in cognitive science research as the frame problem.) World-involving contents trump any attempt to show how straightforwardly self-standing or common-kind representational contents might be organized to yield efficient real-time update—because they update along with the world, thanks to the world. That is why the intergalactic visitors would not resort to positing representational contents that make it very difficult (if not impossible) to explain complex animal activity, when determinate contents that involve the world are exactly what the job description requires. As a cheetah bursts into its 110 kilometres per hour chase and gazelles scatter, our visitors see that only one explanation is adequate.

Of course, the dance of the savannah is not all the intergalactic travellers would visit. It would become clear to them that one species of complex animal is a culture-originating and culture-perpetuating animal whose life activities have been developed into a superabundance of cultural forms. The complexity in this species' activity involves linguistic communication and requires the attribution of complex organismic (whole organism) states to one another, states that can be rendered intelligible in terms of rationalizing propositional contents. Such second-nature animals⁴ act on their environment no less, and of course in order to be able to act and interact they also must be engaged with the individuals and determinate properties with which they interact.

So far, the point of imagining the intergalactic visit is to emphasize that to understand perception is in the first instance to understand the role it plays in the life of organisms. Theory of perception is, in the first instance, integral to the straightforward or best explanation of complex animal life. The best explanation is that complex activity requires a high level of determinate detail, and that such detail

is exactly for objects to enter into content is a nuanced matter. But it is clear that intentionalism denies, whereas I argue, that perceptual contents are determinate, and that the explanation of this fact is that mind-independent objects enter into or are involved in perceptual contents. This is not to suggest that perceptual content has non-representational components or aspects, as Michael Martin (2004) for example, has suggested on behalf of the theory he labels "naïve realism".

⁴ The idea of first and second nature will be explained in the final section, which answers the metaphysical objection. As a first pass, let's say that first nature is the set of inborn potentialities, while second nature is the result of cultural, norm-governed learning—which of course relies on first nature no less than on the cultural norms into which one is trained: that is the point of calling it second.

is available through engagement with it rather than by means of reproducing it. From an explanatory perspective, it is clear that perception and action make up an indissoluble functional unit. What we have historically isolated as ‘perception’ and ‘action’ mutually entail one another as interdependent and co-constituted aspects of an animal’s engagement with the world. Engagement is with individuals. Individuals are determinate (or definite). Their properties are determinate. It is on individuals and their determinate properties that an animal acts and it can only so act in so far as those individuals are perceived to the level of determinacy required for its action—a level of determinacy whose perception is simultaneously made possible by active, self-directed engagement. Co-constitution means that perceptual content would not be possible—and would not be what it is—without complex feedback that is only possible through self-initiated activity (and vice versa of course).⁵ Perceptual engagement might admit limiting cases and peripheral cases, but to explain it one must start from the capacity itself—from its role in the life of the organism—and proceed to understanding the peripheral and limiting cases from there. This is the key point.

Suppose that after a sufficiently prolonged acquaintance we can communicate with our visitors. They tell us how they have come to understand us—they tell us that they understand that we are perceptually and behaviourally engaged with our world. They assure us that it is quite clear that the mental contents that must be attributed to us to render our complex activity intelligible stand in complex justificatory or rationalizing relations. Now suppose we tell them that for hundreds of years we have been troubled by the possibility of imagining that an individual person might have disengaged or non-engaged states that are experientially indistinguishable from her perceptual engagements with the world. It is always possible, we emphasize, that an individual person might not be able to tell for certain when she is in fact in one of those disengaged states. We tell our visitors that it has been a professional occupation for hundreds of years to develop an understanding of perceptual engagement that begins from this putative possibility and that requires explanation of this putative possibility in the first instance en route to understanding perception. Our visitors might be bemused or interested. They might find our approach intriguing, or charming, or quaint. The one thing they don’t do, however, is change their approach. When they get back in their spaceship, the approach characteristic of our epistemic tradition is not the one that would allow them to travel the universe, finding intelligibility where it is to be found.

But they never share their intellectual history with us. They do not tell us whether they also have thought of sceptical challenges to the very possibility of perceptual engagement, and if they have done so, how they deal with them. Does this change anything? All we know is that they accommodate the facts from their explanatory point of view—after all, it is evident from that point of view when an animal is perceptually engaged or when it is in a disengaged state with content that mimics

⁵ For more detailed explanation, see, for example, Hurley (1998: esp. chs. 8 and 9).

engagement. Does this fact—that we do not know how they deal with sceptical problems—change the adequacy of their explanation, given that their explanation is rationalizing and can account for perceptual fallibility?

Like intergalactic travellers who would find us intelligible, what if I propose that perception is perceptual engagement? For an animal whose survival requires complex interaction with its environment, the function of perception is to engage with the determinate particular detail of life: the determinate particular individuals, properties, layouts, events among which it acts. My focus is on us, animals with a second nature that transfigures our first animal nature (and I will return to this way of explaining ourselves in the metaphysically oriented discussion in the fourth section). For us, activity of any kind requires perceptual engagement with the relevant individuals and determinate properties—whether the activity be reading, operating heavy machinery, looking at a painting, playing a computer game, or skiing down an icy mountain. Hiking along a forest path, for example, I need to feel the determinate shapes of the rocks beneath my feet to keep my balance; understanding (of the self-standing, determinable kind) *that there are many differently shaped rocks on the path* would not enable me to walk, it is not what activity requires. Rather, I need to feel “*this* determinate shape” and “*that* curve” and “*that* squishy tilt”. This is what would be clear to the intergalactic travellers. Emphasizing the interrelationship between perception and action from the outset highlights that perceptual experience is distinguished by its determinacy. It shows that theory of perceptual content—in our case no less than in the case of animals—needs to account for the determinacy of perceptual content as one of its primary data.

More precisely, theory of perception needs to explain that perceptions present rather than abstract away from the determinate nature of individuals and properties, and that what is presented must be not only the individuals and properties on which we act, but of course, the individual members of kinds—‘that rock’—and the instances of properties—‘that shape’—on which we act with understanding. If these facts are the starting point for our theorizing, there is no reason to suppose that the *determinacy* of our perceptions entails that those experiences—those engagements—cannot be conceptual like our thoughts. There is nothing about engagement with determinate properties that indicates that such a capability is not or cannot be conceptual. Rather, complex demonstrative concepts—the concepts involved in genuine singular thoughts—provide a model for understanding perceptual content. The model is apt because demonstrative thought or talk involves individuals and properties by means of capacities that pass stringent constraints, so that we call the resulting contents ‘conceptual’. The capacities at issue amount to an *understanding* of the spatial location of the perceived individual or property and this means—though the connection might not be obvious—that there is some *grasp* that individuals are potential bearers of various properties and properties potentially characterize different individuals. That is, the capacities are re-combinatory along both referential and attributive dimensions, a

fact that turns out to be a complex achievement when we recognize all that it entails. But nothing less is involved in thought or experience that is genuinely spatial.⁶

The approach just sketched is a conceptualist form of direct realism. The hallmarks of positions of this type are, most simply, that they claim that (i) perceptual experience is object and property-involving; and that (ii) it is fully conceptual as well. Here is David Wiggins's (1986: 170) more detailed characterization of conceptualist realism:

Conceptualism may be thought of as comprising three claims: first, that the singling out of an object in experience depends upon the possibility of singling it out as a *this such*; second, that there is no surrogate or reductive level (for instance, the level of description of retinal stimulation or whatever), that is, no level distinct from that at which we have objects of this sort and thoughts of them conceived in genuine reciprocity with one another, such that at that reductive level, you could determine what object it was that the subject was impinged upon by; and third, that our cognitive access to reality is always through conceptions that are conceptions of what it is to be this or that sort of thing, these conceptions being a posteriori and at every point corrigible by experience, yet present in advance of the recognition of any particular object as a *this such*.

In part, my point so far is that such realism need not be identified as resulting only from a long philosophical voyage, the sort of voyage that yields contemporary debates about realism. It is also the sensible rather obvious account of ourselves that greets us when we stop to think quite simply about how we are intelligible—which is just to say how we are.

Yet the notion of determinacy and of determinate properties is problematic. To avert misunderstanding, it is important to note at the start that the concept highlights what might seem to be two different points about perception, yet is really one. First, individuals and their properties are determinate. It might be helpful to note that definite and determinate are related concepts. What are 'out there' in the world are determinate or definite individuals and their properties. In so far as complex activity requires engagement with individuals and properties, it is with determinate individuals and determinate properties that complex animals engage. But this is not the only relevant factor when we consider perception. Perceptual experience is not simply individual- and property-involving; it is a mode of engagement. It is not simply the determinate individuals and their properties 'out there, in the world' that our perceptions depend on causally, but those individuals and properties under the relevant perceptual modes of presentation (to put the point using one set of well-known philosophical resources). This is to stress that perceptual engagement is a representational capacity. What we enjoy perceptually is content, albeit a distinctive individual- and property-involving content that presents individuals and properties perceptually. Perceptual content is distinguished by its qualitative richness. This is the second fact that the notion of determinacy captures. Perception's qualitative richness is predominantly cast as a richness of information. But to say that perceptual engagement is determinate—in that it presents and involves determinate individuals and

⁶ For the detailed account, see my (2006).

properties—is to say that the qualitative richness of our experience is in the world rather than something we have ‘received’ along an informational channel. When it comes to perception, the two senses of determinacy belong together. The notion of determinacy relates the ‘richness’ of perceptual content and the world-involving nature of such content.⁷

Consider that when we touch something, running our fingers over it, we gain information about the object’s fully determinate nature—its determinate shape, contours, and texture—by engaging with it. What we think of as our distance senses—perceptual capacities that give us information about objects that are at some distance from us such as vision or hearing or smell—are no different. When we think about running our fingers along a surface, we can readily imagine that enabling sub-personal (central nervous system) processes might be best explained as forming representations of the felt surface. But we are not inclined to think that what we experience is, in the first instance, a representation of the surface, rather than the surface itself—or that the representational content we experience does not involve its object, so that the content could be just as it is even if the surface were quite different or if there were no surface at all. I am suggesting that the so-called distance senses are similar. Distance senses engage us with the world, an engagement that is presumably enabled by complex sub-personal representational processes. But what we are engaged with are the objects and properties out there in the world rather than representations of them. When I touch a table and when I look at a table, or when I hear a coin drop onto a table—these capacities form a kind, they are perceptual engagements with the world. Such capacities are made possible in part by sub-personal processes that may involve and produce representations of the external environment. But it is very important to be clear that what the sub-personal processes make possible is that the animal—the whole animal, the cheetah or the person, let’s say—is engaged with its environment.

This point is in no tension with the evident fact that *what* an animal engages with—which individuals and determinate properties—varies from species to species and also across individuals. As David Wiggins (1986: 171) put a somewhat related point, this is no more problematic or “exciting” than noting that “the size and mesh of a net determine not what fish are in the sea but which ones we shall catch”. If a hippo, alligator, house cat, and human walk down the same path, the determinate properties with which their feet engage will vary along with the thickness of their skin as well as the number of receptors within the skin, etc. (Their engagement will also vary with the nature of their understanding of their surroundings, but I am setting aside the issue of conceptual articulation for now—David Wiggins’s issue about the determinacy of conceptual thought and experience. I will return to it in the last section of the essay but it is not the focus of the point I am making here.) If my feet become thickly calloused, then the determinate properties *out there* with which I engage be different from those I engaged before the calluses grew. The hippo’s feet, needless to say, will engage at least some different determinate properties

⁷ I am grateful to William Seager for indicating the need for this preliminary note about the notion of determinacy. It should also help make clear how the position I am developing differs from others such as intentionalism, for example.

of the surface. Imagine running your fingers along a cloth of raw silk with its slightly stubbly texture. Now imagine running your fingers along that same cloth of raw silk with the thickly calloused fingers that come from decades of hard labour out of doors. The cloth is the same. It is the same 'reservoir' or 'repository' of textures. Which determinate properties I engage with and how I do so depends in part on the relevant mode of perceptual presentation—my skin and its condition in this case—as well as on the properties of the cloth. But regardless of the condition of my skin, it is some determinate property with which this perceptual capacity engages. The same point can be made for our distance senses. Before I put on my glasses, the lettering on a small medicine bottle is thin, grey, and indistinct. As I put on my glasses, the lettering thickens, turning a nice rich black as each shape becomes crisply contoured. But the grey blurs are no less determinate than are the rich black letters that my glasses help bring into view, and that my eyes engaged with on their own a few years ago.

2 THE FIRST HURDLE FOR DIRECT REALISM: IS THE CHARACTER OF PERCEPTUAL EXPERIENCE RELATIONAL?

An argument for engagement from the determinate character of perceptual experience

This argument also begins from the determinacy of perceptual experience to show that perception is a mode of engagement and that perceptual content is world-involving—or in other words, that the “fabric of [perceptual] consciousness is relational” and so “not common to perception and hallucination” (Martin 2004: 19). But the argument switches perspective. It starts from what perceptual experience is like at the first-person perspective, rather than from the intersubjectively evident facts that one could examine, like intergalactic visitors.⁸

Here is the first-person datum. Recently, I spent the month of October in the Canadian 'north', finding myself walking along the same path, seeing the same stretch of forest every afternoon. The deciduous forest all around would take my breath away with its beauty, as the colours changed day by day from green to a jumble of green, gold, bright red, and orange; to all-gold across a stretch of many days; and then finally to a sparse sprinkling of copper when the gold leaves fell. I would stop and try to take in the beauty, to impress it within myself so that I could always remember it. But

⁸ The argument can stand independently of the theoretical orientation just examined, but it fits with and supports it as well. Starting from an intersubjectively available explanatory stance that looks for intelligibility simply counsels using all the relevant considerations and facts. Nothing in such an orientation suggests that facts that are evident from a first-person perspective cannot be used or that they can be trumped by considerations from a third-person perspective; and vice versa, first-person facts cannot trump available facts from a third-person perspective. A truly explanatory perspective needs to be comprehensive—it needs to explain and to use facts from both first- and third-person perspectives without privileging either. This point cuts both ways, so that third-person facts cannot be used to trump those available only from a first-person perspective, as, for example, does Dennett's (1978b, 1991) so-called “heterophenomenology”. It is only so-called phenomenology precisely because first-person facts are simply trumped by considerations that arise from a third-person perspective.

I couldn't. I would stop and look, then tell myself I should keep moving, but stop immediately just to look again. I could never retain the detailed beauty within myself even across the instant that I would look away, only to look back immediately again. This made me realize that it is impossible to remember a complex natural scene even across the span of a second or two, or even just a split second. I couldn't help but think about this fact as I would walk along. And I realized that this is true of all perception, of course: we just don't notice the fact or highlight it when we turn to forming a theoretical understanding of perception.

Similarly, you can examine whether this is true of your perception. You will find that from moment to moment you cannot retain a preceding perceptual experience in its determinacy (let alone recall it with equal determinacy days or months later). Pick up just a single leaf, and as you hold it in your hand, move your eyes away or close them. The determinate nature of the leaf ceases to be available. This is true however 'simple' or 'complex' the content of the perceptual experience might be—a single leaf or blade of grass, or a forest vista.

The datum that needs explaining is that perceptual experience is determinate and one is not able to *retain* a perceptual experience with equal determinacy even momentarily as one's eyes shift, let alone remember or recall the experience with the same determinacy as the original. This is a constant, ever-present fact about perception that needs to enter at the outset into our conceptualization of the nature of perception. It is not something that we suppose might be the case. It is the very character or "fabric" of perceptual experience.⁹

To explain it, one needs to consider possible varieties of mental contents. One uncontroversial category is content that can occur in the absence of the represented individuals and property instances. This category subdivides into varieties depending on whether it satisfies one or both of the following two dimensions: whether the content (i) depends on the represented individuals and properties causally and whether it (ii) involves those objects and properties in the representational content. Descriptive contents and existentially quantified are examples. But there is also the somewhat controversial category of 'genuinely singular' content that cannot occur in the absence of its objects because such content *both* depends on the represented individuals causally *and* involves those objects in the representational content.

The datum I am pointing out indicates that the contents of perceptual experience are genuinely singular. That is, the failure to retain perceptual contents indicates that

⁹ I believe that this datum holds across perceptual modalities though the detailed case cannot be presented here. One might suppose that audition poses a significant counterexample that should make us rethink the visual case as well. It is well known that composers and musically gifted people have rich auditory images of music they are composing or have played. Hence, one might object that auditory images—at least of the musically gifted—are just as determinate as auditory perceptions. By way of a swift response, consider Beethoven. Why was Beethoven made so unhappy by his loss of hearing? If his auditory images had matching determinacy to auditory perceptions, he would have missed nothing of significance by not being able to hear his pieces actually played. The only reason that his hearing loss was as devastating a tragedy as it was to him is that, in fact, an auditory image—even of someone like Beethoven—does not match the determinacy of hearing an orchestra, or chamber group, or single human voice realize a musical composition. Many thanks to Christopher Peacocke for raising this objection in conversation.

such contents depend on their objects and property instances causally, and involve their objects and property instances in the representational content. Otherwise, there would be no 'in principle' reason for failure to retain individuals and properties with the same determinacy as the initial experience. Representational content that does *not* depend on *and* involve its objects and properties can be maintained or recalled in the absence of the represented individuals and properties. Since the initial representation of the relevant individuals and properties does not involve them, such self-standing representational content can be maintained or recalled in their absence, in principle, barring problems of representational cost, storage, and access.

Moreover, the failure to retain a perceptual content in its original determinate detail cannot be explained as a failure of memory—as a failure to recall self-standing contents that is owing to contingent factors such as costliness of storage or complexity of retrieval. There is no time interval at issue. The factors that affect retrieval—such as the 'costliness' of storing complex content and organizational problems of storing complex content in a way that allows for efficient retrieval—do not come into effect instantaneously. Perceptual experience confronts us with two problems, only one of which might be explainable as a deficit of memory. What is at issue is not simply why the experiential content cannot be *recalled* with equal determinacy after a time interval, but why it cannot be *maintained* or *retained* even momentarily in the absence of the presented individuals and properties.¹⁰

To cast this point in the simplest terms, if I (or rather, my sub-personal, central nervous system processes) could make my own pictures or movies at the determinate detail of perceptual experience, and if perceptual experience consists in such movies, then there is no reason I couldn't hold onto and 'watch' or utilize a given 'frame' of such a movie a split second longer. If relevant sub-personal processes make determinate movies, there should not be the differences that we can note at every moment of our waking lives.¹¹

¹⁰ One of the reviewers has objected that if the failure at issue is a failure to retain the content of a perceptual experience, this is an issue about perceptual recollection and why such recollection does not have the determinate detail of a perceptual experience. The suggestion is that "it might be argued that perceptual recollection involves the retention of a representation of a perceptual experience, rather than retention of the content of a perceptual experience". I am grateful for the objection, but I do not see how this explains away my account. If the content of a perceptual experience is determinate, then why would the retention of a representation of a perceptual experience—having determinate content—switch to less determinate content? If the determinacy of the perceptual content were not explained in terms of its object involving nature, then the point about feasibility of representational costs seems to hold. Whatever representational cost is feasible up to the point the eyes turn away should be feasible for a moment longer as one's eyes turn away.

¹¹ Here is a simple illustrative way to cast this point that raises an instructive objection. My claim is that if I (or rather, my sub-personal central nervous system) could make my own pictures or movies at the determinate detail of perceptual experience, and if perceptual experience consists in such movies, then there is no reason I couldn't hold onto and 'watch'—or at least utilize—a given 'frame' of such a movie a split second longer. Someone might object that when we watch movies in movie theatres, what we see are discrete frames (or pictures) with a shuttered interval of 13 milliseconds between each frame. This seems to suggest that we do retain an image for a small time interval—13 milliseconds. But who is the 'we' in question? Movie-watching demonstrates that: (i) human sub-personal processes can maintain an image for 13 milliseconds, whereas (ii) a person

That is why the datum I am pointing out—that we cannot retain or maintain perceptual contents in their determinacy even momentarily—points to facts about perceptual experience rather than memory. There is an important lacuna in the initial descriptions of perceptual experience that theories of perception use, a hard issue that must not be conflated with questions about memory.¹²

What about action? Careful scrutiny of lived experience includes experience during activity; it is not restricted to reflection while we sit, perusing books in our hands or looking through our windows. The unity of our perception and action is very much alive at the first-person viewpoint. It is not something that can be observed only from the third-person viewpoint. From the first-person perspective, no less than from the third-person viewpoint, activity makes it clear that one key issue is change and the fact that complex action must deal with change as it occurs in real time. The activity in the first person datum I offered—that of looking at the leaves of trees while walking along a peaceful forest trail—does not raise in any obvious way the need for dealing with change as it occurs. But any number of different examples could be raised where this need would be vivid—such as hunting a deer in that same forest, or preferably, watching birds, or skiing down its icy hillsides after the leaves have fallen—as well as examples that switch scenarios altogether such as our need to drive on multi-lane highways at 120 kilometres per hour. (One might be tempted to return to the example with which we started: the daily dance of gazelles and their kin with lions and cheetahs on an African savannah. When it comes to thinking about the functional unity—or constitutive interdependence—of perception and action, such examples might seem more vivid, but only if one overlooks what it is like to run along a forest path or ski down an icy mountain when one theorizes.) From the first-person perspective, one experiences changing determinate features, and one can understand that it is necessary to experience changing determinate features as changes occur. How is this best achieved? By continuous update of representational content that does not involve the individuals and properties? Or by content that involves the individuals and properties? Considering the costliness of updating a representation of a complex natural scene—and the organizational difficulty of ensuring immediate access to just the needed portion of the ambient scene (already mentioned as the frame

does not detect intervals of 13 milliseconds. The phenomenological datum stands. As a person turns her eyes away, she can no longer retain what she just saw in anything like the determinate detail of the experience itself. Many thanks to William Seager for suggesting this objection as a way of clarifying my point.

¹² Note that if one were inclined to press objections from well-worn (or worn-out) evil-genius types of scenarios, this datum can be used in rebuttal. If the evil genius is ‘feeding’ us our determinate experiences there would be no reason—stemming from the nature of experience itself—to suppose that he wouldn’t feed them to us as our eyes shift (or close) while trying to retain the preceding experience in its determinacy. There is no reason to suppose this other than that he wanted to mislead us precisely in the direction of the sort of theory I am proposing, one that argues from the premise that we fail to retain or maintain perceptual experiences with equal determinacy to the conclusion that determinacy is a function of engagement. But at this point, the argument from an evil genius is resorting to ad hoc additions to pre-empt a specific theory rather than relying on facts that anyone might putatively discover about the character of her experiences.

problem)—suggests that the best solution to the problem, and hence the best explanation, is that one is perceptually engaged with the world, engagements that involve the individuals and properties that one's activities concern. Culture-originating and culture-perpetuating animals are no different from other animals in this respect. The constitutive interdependence of perception and action indicates that the best explanation for real-time activity is that it is mutually supported by and supports object- and property-involving contents—contents that up-date 'for free' along with the changing surroundings.

This is the abbreviated case showing that the very fabric of perceptual consciousness is relational—or in my terms, that the determinate content of perceptual experience is engagement that involves determinate individuals and properties. It is an opening. The case proceeds from what perceptual experience is like from the first-person perspective and goes on to show that only one explanation is adequate to the facts. The more elaborated case will presumably be multifaceted and multi- or interdisciplinary in nature, bringing together further research and explanations of different kinds that highlight different factors.

Briefly, here are but two of the needed tasks. One is to work out a precise account of perceptual content at the personal level. What needs explaining is how perceptual engagement is both conceptual and object- and property-involving. This possibility has been viewed with scepticism despite ground-breaking work by Gareth Evans (1982) and John McDowell (1986). Both have given detailed arguments for the "possibility of Fregean senses which are object-dependent" (McDowell 1986: 143). That is, both have argued that the leading insight from Frege—that senses or modes of presentation determine the individuation of mental contents—can be combined with the Russellian idea that there are genuinely singular contents that involve the object. Russell's insight was that "in order to entertain a proposition one must know how one's thinking represents things as being" (McDowell 1986: 140), an insight that entails that we can think singular thoughts only in so far as we know which object (or property) is being presented. This stringent requirement on genuine singular thought might seem to drive us inwards to a "highest common factor" (McDowell 1982/1998: 392–4) that we cannot fail to know. But appeal to the disjunctive insight shows that the retreat is not needed. Genuine singular thought can be intrinsically connected to the world (or engaged with it, in my terms) in a way that we can understand, even if we can, on occasion, be wrong in ways that we might not be able to discern. Together, these insights lead us to recognize the possibility that there are "Fregean thought-constituents (singular senses) which are object-dependent, generating an object-dependence in the thoughts in which they figure" (McDowell 1986: 142). These ideas, if fully developed—as, for one example, I have developed them in Sedivy (2006)—can explain the fact that perceptual content is both conceptual and object- and property-involving.

If we turn towards sub-personal explanations, it is clear that a variety of research is already beginning to show the constitutive interdependence of perception and action. Detailing aspects of that interdependence will show what it takes for perception to be a mode of engagement. One interesting issue here is that further research needs to

show whether our central nervous system *cannot* or mostly *does not* represent at the level of detail required for complex activity. That is, the reorientation I am suggesting leaves it an open matter awaiting empirical investigation whether:

- (a) The central nervous system (or CNS) *cannot* represent the determinate detail that is characteristic of perceptual experience and required for complex activity in a natural environment by means of self-standing contents that do not involve their objects. The computational costs and the cost for update in real time for action are too high. Humans mistakenly suppose that they can visualize with equal determinacy; and from the first-person viewpoint it mistakenly seems that dreams and hallucinations match the determinacy of perceptual experiences.
- (b) The central nervous system *does not* represent the determinate detail that is characteristic of perceptual experience and required for complex activity when the animal is perceptually engaged with the environment, because the computational costs associated with representational contents that are not object-involving is too great and update in real time is either not possible or too costly. However, the CNS can represent at the same complexity of determinate detail during relatively brief and intermittent intervals by means of representational contents that are not object-involving.

Bringing these future elaborations into view gives but a sense of the multifaceted framework whose construction we can pursue. What we can already appreciate is how the determinacy of perceptual content propels us towards a new understanding of perception as a mode of engagement.

3 THE ROLE OF THE DISJUNCTIVE INSIGHT: EXPLAINING PERIPHERAL AND LIMITING CASES

Our trajectory shows that in an explanatory approach to perception, issues of fallibility need not come up at the very outset but at the point of refining the explanation to deal with errors or illusions, as well as hallucinations and dream states. The disjunctive insight is required at this juncture in the explanatory sequence. The approach taken here also indicates that determinate object-involving contents are the central condition and perceptual illusions, hallucinations, etc. are peripheral and limiting cases.¹³ The secondary status of the latter sorts of states to fully engaged experiences follows from the fact that determinate content needs to involve its objects and properties.

¹³ This does not deny that the disjunctive insight might figure in a transcendental case for the interpenetrating nature of mind and world, a case for the very possibility of genuinely singular experiences and thoughts. Recall that this essay displays what I am calling a 'straightforward best explanation' of perception, while elaborating a direct realism that is integral to approaches that might begin from a transcendental case. The disjunctive insight may also figure in a transcendental case that starts from some undeniable facts of what perceptual experience is like, and goes on to show that only world-involving perceptual capacities could deliver such experiences. For an example, see John McDowell's essay in the present volume. Typically such arguments have started from the

Understood as peripheral cases, perceptual errors or illusions do not pose special explanatory challenges for this sort of approach (as they do for purely causal theories of mental contents) since perceptual content results from 'Fregean senses' or modes of presentation as well as the involved objects. But this approach has interesting implications for our understanding of hallucinations and dreams. The fact that determinate content needs to be object-involving entails that hallucinations and dreams have a secondary, dependent status, as I indicated.¹⁴ It follows that dream states, for example, might have contents that either *are* indistinguishable or *seem to be* indistinguishable from waking states even though they are in fact not as determinate (see empirical possibilities (a) and (b) above). Other kinds of disengaged states might also have contents that mimic the sorts of contents we enjoy during normal engagement. But such disengaged states can only be intermittent or relatively isolated occurrences in contrast to the normal engaged, object-involving condition. Moreover, note the qualification that a hallucination or dream might only *seem* to be indistinguishable from an ongoing perceptual engagement. If determinate content is object- and property-involving, then disengaged states would not be determinate even though they might seem to be determinate—most would not be determinate or none would be determinate (again, see empirical possibilities (a) and (b) above). There is no denying this challenge to the transparency of the mental (the idea that when it comes to our mental states we cannot be mistaken about their contents or 'what they are like') even if a single essay cannot tilt at all the windmills in this domain.

In short, because an illusion, hallucination, or dream might at the moment of its occurrence be or seem to be as determinate as a perceptual experience, a theoretical specification of the nature of perception needs to take account of this fact with a disjunctive account of perceptual experience from the first-person viewpoint.¹⁵ If we explain that

(i) Visual perception (for example) is a mode of perceptual engagement,

we need to go on to explain that

(ii) human beings are capable of a variety of fully engaged and disengaged states in the visual modality (as in any other perceptual modality);

apparently spatial or objective nature of perceptual experience—since perceptual experience is at the very least as of spatial objects. But the determinate character of perceptual experience might also be a starting point for a transcendental case, as the previous section indicates. In fact, I have argued (2006) that the objective and the determinate character of perceptual experience is one complex phenomenon. But that issue lies beyond the scope of this essay, as does the possibility of using perceptual determinacy for a transcendental argument.

¹⁴ It is also possible to argue from the objective, spatial nature of perceptual experiences to the same conclusion: that engagement is the central condition that makes a variety of peripheral or limiting cases possible. These two facts about perceptual content—its objective, spatial nature and its determinacy—are connected, as I indicated, though I cannot make the case here (2006).

¹⁵ There is also a sense in which our ordinary mature understanding of ourselves has a disjunctive structure. A person understands that she might occasionally be in a disengaged state that mimics engagement without being able to figure this out, and this insight needs to be available to her as she engages with her world.

and that

- (iii) from the first-person perspective, perceptual experience is marked by the disjunctive possibility that one might on isolated occasions be in a disengaged state rather than engaged, without being able to figure it out oneself—a possibility that the experiencing person needs in some sense to grasp.

Locating the disjunctive insight within the approach rather than at the outset—as an explanation of peripheral and limiting phenomena made possible by the central engaged condition—does not imply that theory of knowledge is less important than we have thought. The point is to show how theory of mind and theory of knowledge work together as mutual partners within an explanation of perception. In contrast, epistemically motivated disjunctive theories of perception that begin from sceptical scenarios leave open the thought that if we could solve our epistemic problems in an alternative way, well then, so much the worse for a disjunctive structure—it can fall by the wayside. A disjunctive approach is on a firmer, less optional, footing in so far as it is required by the facts of perception, and other facts of the engaged human condition. Alternative ways of explaining a set of facts is always an issue as well, but I am suggesting that the determinacy of perception doesn't allow for alternative explanation: determinate perceptual experience is the central case that allows for peripheral or limiting cases that might mimic determinate engagement in its absence.

Finally, there is the question about the status of the disjunctive thesis—is it only epistemic or metaphysical as well? I must confess to discomfort on this point. The thesis is epistemic since it denies that perceptions and other states from which perceptions might be indistinguishable have the same epistemic standing. Does the thesis hold that these various kinds of states do not share any important metaphysical properties (in so far as being indistinguishable is not taken as an important metaphysical property)? I have argued that the states at issue are contentful states with contents of different kinds and I am uncomfortable about talking about contents in metaphysical terms (Sedivy 2004). Just as I have avoided any mention of 'disjunctivism', preferring to discuss the disjunctive insight (since that is what I believe is at stake), so I would prefer to leave metaphysics out of the discussion when it comes to experiential contents and stick to detailed analysis of the different varieties and the relationships between them.

4 THE SECOND HURDLE FOR DIRECT REALISM: IS IT THE WORLD WE ENGAGE WITH?

In so far as our scientific picture or our basically empiricist metaphysics holds that what is out there is quite different from our ways of understanding or perceiving it, it could not be mind-independent individuals and properties that our perceptual experience involves. Hence this dominant metaphysical framework blocks direct realism in any form, including conceptualist varieties. After all, conceptualist realism claims that our perceptual, conceptual engagement with the world captures determinate individuals and properties in its net. Here are two answers to the metaphysical

discomfort with the conceptual, direct realist proposal that perception engages us with the world—that perceptual modes of presentation can get the world’s determinate individuals and properties in their net. Both answers turn on legitimating a “non-objectifying” (McDowell 1982/1998: 392) view of the world as well as of our selves.

4.1 First response: from interpretivism

A broadly interpretivist theoretical framework requires multiple explanations not only of ourselves but also of the world with which we interact—though this seems not to have received as much airtime in theory-of-mind discussions. The requisite variety of explanations will include not only the ‘objectifying’ ones suitable for explaining the sub-personal processes and external conditions that enable an animal’s or person’s capacities, but also the ‘non-objectifying’ explanations appropriate to persons and world, animals and habitats.

I am not suggesting the need to subscribe to either Davidson’s (1984) or Dennett’s (1978a, 1987) particular forms of interpretivism, but rather to point out what their brands of interpretivism and indeed any broadly interpretivist approach are committed to. It is useful to point out this commitment, precisely because many people do buy into either a Dennettian or a Davidsonian interpretivist approach, or into a more broadly conceived framework along such lines. The broader framework to which many of us subscribe even if we don’t think that Davidson or Dennett get it quite right can be summed up as the insight that entities—individual members of kinds or determinate instances of properties—figure in explanations. What it is for there to be such entities is to figure in a certain kind of explanation, an explanation that captures patterns that cannot be captured by any other kind of explanation.

Davidson and Dennett make clear that complex systems require more than one kind of explanation. What is it to be a person or an intentional system? It is to be intelligible as a person or intentional system—but not only as a person or intentional system. The complexity is such that a complete understanding would involve other explanations as well—sub-personal or sub-intentional system explanations that give us insight into the workings of the parts of persons that causally enable a person’s complex capacities. To be intelligible as a person (I am going to drop the longer intentional system phrase from here on) requires attributing representational and presentational capacities, states, and utterances. But what are these capacities, states, and utterances about? They are about the shiny red apples that we like to eat, the grey clouds that indicate rain, the eye-stopping scarlet that a maple leaf becomes. In so far as a person’s true beliefs are about the beautiful progression of colours in a particular deciduous forest in the autumn, then it is a beautiful progression of colours in that autumn forest that is in her surroundings. What surrounds the sub-personal states that help make those beliefs possible? The relevant explanatory factors might be light waves and surfaces that reflect different portions of light waves, but this explanation is not quite settled yet and we need not take sides for our purposes here. Just as we recognize that we need to give more than one account about

persons, we also need to give more than one account about the surroundings that the person's beliefs, wishes, fears, and perceptual experiences involve and concern. And there need be no conflict between these explanations. It is wavelengths of light that affect my sub-personal workings. It is colours that affect me. A person's environment, unlike the surroundings of her sub-personal parts, consists of the colours, sounds, tastes, smells, buildings, roads, cars, and other individuals and properties on which she acts—in short, a person's environment is the world. This brings me to the second answer.

Let's quickly summarize first. Interpretivism makes salient two points: (1) that entities figure in explanations; and (2) that complex systems or organisms require more than one kind of explanation—a point that carries over to the surroundings of a complex system, what it engages with, as well as the system itself.

These two points concern us as well as animals and other complex systems. Yet we also need to be able to distinguish between the habitat that first-nature animals engage with, and the world with which second-nature or culture-originating animals engage. Depending on how interpretivism is developed, for example in Davidson's hands or in Dennett's, it might be able to make this further distinction between persons and world on the one hand, and animals and habitat on the other. I would like to highlight a different set of resources for this insight.

4.2 Second response: from Wittgensteinian Realism with a Human Face

Wittgensteinian Realism with a Human Face develops distinctive resources—which focus on human practices and their dimension of rule-following—for understanding ourselves specifically as inseparable from the world.¹⁶ What makes this realism is that it attempts to give us a new understanding of objectivity that replaces our sense of the objective as that which does not involve ineliminable mention of our selves. The objective—be it facts, values, or simply 'the world'—is co-constitutive with ourselves, so the very fabric of the world is not something whose specification can be independent of reference to us and our activities. Realism with a human face gives us the extra bit that we need to address the worry that it can't possibly be the world—the world of brightly coloured, annoyingly revved-up sports cars zooming by—that we are engaged with perceptually and through our bodily actions. If successful, realism with a human face shows how a non-objectifying explanation can contribute to our understanding of our selves and our relationship with the world.

¹⁶ This wonderful phrase is Hilary Putnam's (1990). What is lovely about it, in part, is that it suggests that there is nothing naïve about such realism, even if it doesn't do so well during our time, rather like socialism with a human face, crushed as it is by tanks at one point in its history, at another by the rush of multinational corporations. But what makes Putnam's phrase theoretically important is that it highlights the fundamental issue when we think about the possibility of realism and declares its stance on that issue. Realism is problematic because we cannot but recognize the contingency of our engagement with the world, from the nature of our perceptual modalities to the practices we strive to refine. Realism with a human face argues that it is precisely our contingency that makes realism possible.

Wittgenstein develops the idea of mutuality between persons and world by focusing on human practices or forms of human life activities and on the fact that such practices are largely rule-governed, or informed and constrained by norms and standards. Just as rule-following transfigures human animals into persons (human animals with a second nature), so rule-following opens up the domain of objective facts and values within which human lives take place and with which human beings must deal. Ways of acting come together with what those actions involve and concern, and ways of judging come together with what those judgements involve and concern.

Here is a quick sketch. It falls into three parts. First, I will quickly go over Wittgenstein's leading ideas about rule-following. This will allow two further sketches: of the way in which rule-following yields second nature, and of the way in which rule-following brings "into view" but not "into being" (McDowell 1996: 285) an objective world—of determinate individual members of kinds and determinate instances of properties—that is inseparable from second nature. The core phenomenon of rule-following shows that the two—second natures and objective world—are like two sides of single coin; or, in other words that realism comes with a human face.

If it seems odd that this essay should now turn into a somewhat extended discussion of Wittgenstein's thought—well, that is part of the point, too. Starting out by asking how we are intelligible (perhaps considering intergalactic visitors along the way) is in no sense disjoint from learning from Wittgenstein's later work. Contemporary theory of mind that strives to understand perception and the nature of perceptual content is in no sense disjoint from learning from Wittgenstein's later work. Neither is contemporary epistemology.

(i) Wittgenstein's discussions of rule-following try to transform our understanding of ourselves and of our place in the world by bringing the dimension of communal practices and customs into focus. Practices come into view when we see that we are following rules throughout our various activities—many that we take to be cognitive, solitary achievements—and that following rules involves membership in a practice, a public communal way of going on. This is Wittgenstein's controversial suggestion that "To obey a rule, to make a report, to give an order, to play a game of chess, are *customs* (uses, institutions)" and "'obeying a rule' is a practice" (Wittgenstein 1953: sections 199 and 202). In what has been cast as an argument against the possibility of private language or private rule-following, Wittgenstein shows that rule-following requires a public context that functions—in myriad ways that he explores and details—as the repository of criteria of correctness that are independent of any individual rule-follower but that are brought to bear on each. If criteria of correctness were not independent of each individual, then one might claim any behaviour to follow a certain rule. Consider the opposite: the absence of extant practices in which it is determinate what the rules specify. In such a case, if it seems to me that adding by 1 up to 1000 and by 2 after 1000 follows the rule '+1', for example, then—it does! But if this were all rule-following amounts to, then the idea that we follow determinate

rules is empty—without established practices the very notion of a determinate rule would be empty.

And hence also ‘obeying a rule’ is a practice. And to *think* one is obeying a rule is not to obey a rule. Hence it is not possible to obey a rule ‘privately’: otherwise thinking one was obeying a rule would be the same thing as obeying it. (Wittgenstein 1953: section 202)

Though this is the most controversial, it is only one of the interwoven considerations pointing in the same direction. Criteria of correctness that are independent of any individual are extant in our practices, in the way we do things at the level of communities, and are internalized by individuals through training. This does not suggest that all there is to criteria of correctness is communal agreement in practice, or that communal agreement determines what is correct and what is incorrect. Rather, the struggle is to understand how it is that human practices bring criteria of correctness “into view” though not “into being”, to echo John McDowell’s phrasing.¹⁷ The second and third steps—showing how rule-following opens up both second nature and an objective world—strive to give substance to that distinction.

(ii) The second point in my summary above—about internalization through training—is just as key as the first about the need for criteria of correctness that are independent of individuals. Recognizing that rule-following requires established practices is only part of the point; the other part is to recognize the extent to which all aspects of ourselves involve the internalization of rules through training and hence involve the context of established practices. Many of the interweaving threads in Wittgenstein’s examinations of rule-following strive to elicit this recognition. The picture of ourselves that emerges is of a piece—perhaps surprisingly—with Aristotle’s account of moral virtue as second nature in the *Nicomachean Ethics*.¹⁸ The affinity is clear as soon as we notice that both Aristotle and Wittgenstein stress the role of *training*, the process of instilling immediate non-interpretive rule-following. Training instils ways of acting and reacting across the multitudinous situations that comprise our forms of life. Hence it is the process by means of which we acquire what we might call a ‘second nature’, following Aristotle. The concept of second nature is the notion of acquired ways of acting that are instilled through training to the point of seamless automaticity. It is precisely the idea of unreflective, immediate rule-following—of cultural norm-governed ways of being. “When I obey a rule, I do not choose. I obey the rule *blindly*” (Wittgenstein 1953: section 219). Aristotle develops the idea of second nature with respect to our moral nature. Wittgenstein’s investigations indicate the same kind of account, though not in the same terms, but take it further by exploring how myriad activities involve rule-following.

¹⁷ In the following two sections I am trying to capture with a somewhat different focus the point that John McDowell has articulated vividly with this distinction: “acquiring a second nature brings the demands of reason into view (a view that may be skewed or inaccurate in some respects). It does not bring them into being.” (1996: 285).

¹⁸ Here I am of course springing off from McDowell’s suggestion to this effect in his (1994), especially Lecture IV.

According to Aristotle, we become moral persons by being trained—when we are uncomprehending children—into doing the right thing across various circumstances to the point of automaticity, the point where we just react to our circumstances by doing the right thing. Only at this point—where we *react* to the moral possibilities in our circumstances with the right action—can we reflect on our actions; understand why a certain way of acting—being truthful, for example—is morally virtuous or correct, and freely choose it. It is only at that juncture—where a reaction is also a free choice—that an action is moral. Aristotle's point is that one can't understand why one should be moral unless one already is moral. One can't freely choose moral actions unless one already acts morally.

This is a picture of our selves as animals with a second nature: our first animal nature (which, for example, responds to reward and punishment and enables us, members of one species, to respond similarly in similar circumstances) becomes transfigured into a norm-governed second nature through training into ways of acting and reacting. Wittgenstein's examinations of rule-following suggest the same sort of account. We have seen that these investigations culminate, in part, in the recognition that when I understand another's speech or say something myself I am participating in "*customs* (uses, institutions)" (1953: section 199). Rather than focusing on the fact that we are moral agents, Wittgenstein examines the fact that we are 'thinking beings' by showing that thinking—applying concepts, making descriptions or reports with those concepts, etc.—involves rule-following no less than what we think of as moral action. Aristotle's account presupposes a cultural communal context. This might not seem startling when moral training is at issue. But the point becomes controversial when its generality is pointed out: rule-following of any kind involves training, which requires a larger context where we are held responsible to norms of correct behaviour. By showing that our various cognitive capabilities involve following rules—and involve second-nature reactions and judgments—Wittgenstein's investigations show that practices supply the required context for every aspect of our emergent second-nature selves. It takes a village to raise a child.¹⁹

But how is this realism? What about the world and our place in it? How do the ubiquity of rules and practices yield realism? If criteria of correctness are extant in practices, which are norm-governed ways of doing things at the level of communities, then what vouchsafes the practices? These questions lead to the third step of showing how rule-following not only transfigures animals into second nature, but opens up the world in which they act as well.²⁰

(iii) Wittgenstein stresses that our practices are ours—with the full historical contingency that this entails. Yet Wittgenstein's point is that stress on the contingency of our practices also shows the objectivity of what we thereby do.

¹⁹ It is in his epistemological work that Wittgenstein (1969) investigates the ubiquity of rule-following that is constitutive of our second nature.

²⁰ This last question of course raises the possibility of interpreting Wittgenstein as an anti-realist rather than a realist, as some currently do. But it is not my responsibility in this essay to offer that putative interpretation.

“So you are saying that human agreement decides what is true and what is false?” It is what human beings *say* that is true and false; and they agree in the *language* they use. That is not agreement in opinions but in form of life. (Wittgenstein 1953: section 241)

If language is to be a means of communication there must be agreement not only in definitions but also (queer as this may sound) in judgements. This seems to abolish logic, but does not do so.—It is one thing to describe methods of measurement, and another to obtain and state results of measurement. But what we call “measuring” is partly determined by a certain constancy in results of measurement. (Wittgenstein 1953: section 242)

For there to be measurement of length is for there to be users of measurement with a perceived need for measurement who can agree in what it is to follow the constitutive rules of measurement. And the agreement must be unproblematic, seamless—something that we take for granted once we have been taught how to measure (that is, once we have been trained into the practices in which measurement is an integral part—or practices of measurement, in shorthand). There must even be some constancy in the measurements that are obtained, as Wittgenstein points out. But to make all this clear is also to make clear that the results of measurement are not thereby determined as well. In so far as I have been trained into the practices of measuring, I can use a measuring tape to figure out the length of this room—but this in no way determines the length of the room. That the room is *metres* or *yards* long is a contingent fact determined by our practices of measurement. There is even a sense in which the fact that the room is measurable and has a *length* is contingent on there being practices of measurement of length at all. But *what* length the room is measured to be in metres or yards is not contingent on our practices. While this point tends to be clear and unproblematic when it comes to measurement in metres or yards, it seems to become elusive and controversial when applied to language: to language games and to the words that fulfil the same sort of role across our myriad practices as metres or yards do when we measure. But if we accept the point about measurement, we cannot reject what is the same point concerning language.

The agreement, the harmony, of thought and reality consists in this: if I say falsely that something is *red*, then, for all that, it isn't *red*. And when I want to explain the word “red” to someone in the sentence “That is not red”, I do it by pointing to something red. (Wittgenstein 1953: section 429)²¹

Yes, forms of life that involve explicitly identifying colours are our forms of life. Forms of life that involve practices of measurement are our forms of life. But that does not mean that practices of measurement are merely ours—as if we impose measurement on a world that stands constitutively apart from us as a bare presence. Forms of life that involve practices of measurement do so because the world is such as to

²¹ Section 429 of Wittgenstein (1953) is set up by its companion section 428, which discusses how we lead ourselves astray into apparently metaphysical waters when we stop to think and to theorize about the ability of our meanings or our thoughts to reach out to the facts: “‘This queer thing, thought’—but it does not strike us as queer when we are thinking. Thought does not strike us as mysterious while we are thinking, but only when we say, as it were retrospectively: ‘How was that possible?’ How was it possible for thought to deal with the very object itself? We feel as if by means of it we had caught reality in our net” (Wittgenstein 1953: section 428).

be measurable—given that, of course, there are forms of life in that world in which practices of measurement have a place—though, of course, there are forms of life in which measurement has a place only in so far as the world is measurable.

We can begin to grasp this mutuality by following Wittgenstein's stress on activity. He emphasizes repeatedly across various contexts that 'at bottom' are activities—rule- or norm-governed ways of acting in the world that have been instilled through training. His image of *striking 'bedrock'* evokes this inescapable fact of the human condition: that when we press to justify our ways of going on, we find that our justifications run out in our ways of acting—this is how I act, this is how I have been trained to react. And it is in the world that we act, of course. That action is in the world and on the world—enabled and constrained by the world—is implicit in the notion of action. Hence it is perhaps easy to overlook.²² But it gives substance to the contrast between bringing something—a certain fact, for example—into being and bringing it into view. That wheat is nutritious is not brought into being by our activities. But it is brought into view by our activities and cannot be in view without them, or analogous practices of harvesting seeds, breaking them down, etc. This is the point about measurement and colours, cited from Wittgenstein's text above.

Our activities are constituted in a reciprocally evolving relationship between the possibilities for action in the world and our possibilities. This is the sense in which forms of life—we might more fully say forms of life activities—are co-constitutive for persons and world; the sense in which it is forms of life that open a mutuality of persons and world, languages and facts, actions and values. In so far as forms of life are co-constitutive, facts or values are not one-sidedly of our making or up to us—the persons who are co-constituted with the world. This is realism with a human face. In struggling with such realism we are struggling with changing our understanding of objectivity. The notion of the objective—of objective facts or objective descriptions of such facts, for example, is not of a description that does not involve an essential (or ineliminable) reference to ourselves, the beings to whom such facts can be evident or who articulate such descriptions. We have a fairly recent though potent history of supposing that an objective description is one that must not mention us and that the objectivity of the world is the objectivity of a bare presence that does not involve us in any way. On this sort of view, any aspect that requires that its specification involves ineliminable mention of our selves is not objective, not part of the 'fabric of the world', but a projection of ours onto the bare presence that is really the world. This is the view that Wittgenstein tries to displace by detailing the mutuality between the objective—be it individual members of kinds or determinate instances

²² "How can he know how he is to continue a pattern by himself—whatever instruction you give him?—Well, how do I know?—If that means 'Have I reasons?' the answer is: my reasons will soon give out. And then I shall act, without reasons" (Wittgenstein 1953: section 428); "The end is not an ungrounded presupposition; it is an ungrounded way of acting" (Wittgenstein 1969: section 110); "'How am I able to obey a rule?'—if this is not a question about causes, then it is about the justification for my following the rule in the way I do. If I have exhausted the justification I have reached bedrock, and my spade is turned. Then I am inclined to say: 'This is simply what I do'. (Remember that we sometimes demand definitions for the sake not of their content, but of their form. Our requirement is an architectural one; the definition a kind of ornamental coping that supports nothing.)" (Wittgenstein 1953: section 217).

of properties, facts, descriptions, values, or simply ‘the world’—and ourselves. Detailing the mutuality between persons and the world shows that realism is available to us, so long as it is realism with a human face.

Let’s be clear how realism of this stripe was contained in the notion of rule-following and second nature with which we began. Just as rule-following transfigures human animals into persons (human animals with a second nature), so rule-following *opens up* the domain of objective facts and values within which human lives take place and with which human beings deal (by harvesting and ‘measuring’ and otherwise). Ways of acting come together with what those actions involve and concern, and ways of judging come together with what those judgements involve and concern. Ways of measuring, one might say, come together with what ways of measuring concern—the measurable. Training opens up a mutuality of activity with circumstances and objects of activity just as it makes immediately available a mutuality of judging with circumstances and objects of judging. To be trained into second nature is to react to one’s circumstances appropriately. This is equivalent to saying that in so far as I *immediately react* to objects or layouts as having lengths, length is *immediately there*. Reactions and circumstances of reactions entail one another. This was Aristotle’s point. His account of what it takes to be moral also addresses our meta-ethical anxieties, explaining how morality can be a real dimension of the world. What is it for there to be moral possibilities in our circumstances? It is for us to react to moral possibilities in our circumstances so that we can rationally assess and choose them. The Aristotelian notion of second nature puts in place a mutuality of moral agents and moral circumstances. What we might call its Wittgensteinian extension puts a broader mutuality in place—of activity with circumstances and objects of activity and judging with circumstances and objects of judging. Contingency and objectivity are not strange bedfellows after all. And immediacy is a key ingredient in their kinship.

Wittgensteinian realism with a human face and interpretivism are but two entry-points to the idea that the world of objective facts and values is inseparable from what we are. Individual members of kinds and determinate instances of properties make up this world and we are perceptually engaged with them.

5 CONCLUSION

Considerations of intelligibility, explanation, and interpretation; conceptualist realism and Wittgensteinian realism with a human face; Fregean senses and Russellian singular thoughts; individual- and property-involving modes of presentation and their sub-personal enabling conditions—this essay has touched on these and more besides. They are just some of the ideas and approaches that might cross-fertilize in understanding our perceptual engagement with the world. It is in this complex, I am suggesting, that the disjunctive insight finds its proper realization, in so far as that insight beckons us to take a fresh start.

This essay has focused on the determinacy of perceptual experience—which is evident from both third- and first-person explanatory perspectives—in order to undertake a new start. It has also detailed a conceptualist approach to our selves as

explanatorily inseparable from the world to show the contributions that a non-objectifying approach can make. Here is a summary of the explanatory package—a package that could not be on offer without the disjunctive insight at its core.

- (1) The best explanation of the perceptual capacities of complex mobile organisms is that such capacities engage the animal with the individuals with which it interacts and with the determinate properties that are relevant to its activities.
- (2) The best explanation of the fact that human perceptual experience is determinately detailed to a point that outruns *retention* (let alone recall) is that human perceptual experience is engagement with individuals and determinate properties. This means that the content of perceptual experience is object- and property-involving (as well as conceptual).
- (3) To account for the possibility that, from the first-person perspective, human perceptual engagement might be or seem to be indistinguishable on occasion from various disengaged states, experiences need to be construed disjunctively as either perceptual engagements or disengaged conditions whose content might mimic the individual- and property-involving content distinctive of perceptual engagement.
- (4) The disjunctive approach is required by the facts of perception, by the need to explain perceptual engagement and its peripheral and limiting cases. (It may have other theoretical uses as well, such as playing a pivotal role in transcendental arguments.)
- (5) Since (i) the explanations in which persons figure are also necessarily the explanations in which the world figures, and (ii) persons and world are mutually constitutive and inseparable, it is the world—of individual members of kinds, determinate instances of properties, facts, and values—that we represent, that we understand and with which we are actively and perceptually engaged.

Isn't this the understanding we already share?

REFERENCES

- Davidson, D. (1984) 'Radical Interpretation', in his *Inquiries into Truth and Interpretation*, Oxford: Clarendon Press.
- Dennett, D. (1978a) 'Intentional Systems', in his *Brainstorms*, Cambridge, MA: MIT Press.
- (1978b) 'Two Approaches to Mental Images', in his *Brainstorms*, Cambridge, MA: MIT Press.
- (1987) 'True Believers', in his *The Intentional Stance*, Cambridge, MA: MIT Press.
- (1991) 'A Method for Phenomenology', in his *Consciousness Explained*, Boston: Little, Brown, and Company.
- Evans, G. (1982) *Varieties of Reference*, Oxford: Clarendon Press.
- Hurley, S. L. (1998) *Consciousness in Action*, Cambridge, MA: Harvard University Press.

- McDowell, J. (1982) 'Criteria, Defeasibility, and Knowledge', *Proceedings of the British Academy*, 68: 455–79. [Reprinted in McDowell 1998: 369–94. Page references are to this edition.]
- (1986) 'Singular Thought and the Extent of Inner Space', in P. Petit and J. McDowell (eds.) *Subject, Thought and Context*, Oxford: Clarendon Press.
- (1994) *Mind and World*, Cambridge, MA: Harvard University Press.
- (1996) 'Reply to Gibson, Byrne, and Brandom', in E. Villaneuva (ed.) *Perception, Philosophical Issues*, 7: 283–300.
- (1998) *Meaning, Knowledge, and Reality*, Cambridge, MA: Harvard University Press.
- Martin, M. G. F. (2004) 'The Limits of Self-Awareness', *Philosophical Studies*, 120: 37–89.
- Putnam, H. (1990) 'Realism with a Human Face', in his *Realism with a Human Face*, Cambridge, MA: Harvard University Press.
- Sedivy, S. (2004) 'Minds: Contents without Vehicles', *Philosophical Psychology*, 17, 2: 149–79.
- (2006) 'Nonconceptual Epicycles', *European Review of Philosophy*, 6: 33–66.
- Tye, M. (1995) *Ten Problems of Consciousness, A Representational Theory of the Phenomenal Mind*, Cambridge, MA: MIT Press.
- Wiggins, D. (1986) 'On Singling out an Object Determinately,' in P. Petit and J. McDowell (eds.) *Subject, Thought, and Context*, Oxford: Clarendon Press.
- Wittgenstein, L. (1953) *Philosophical Investigations*, 3rd edn, trans. G. E. M. Anscombe, Upper Saddle River, NJ: Prentice-Hall.
- (1969) *On Certainty*, trans. Denis Paul and G. E. M. Anscombe, New York: Harper Torchbooks.

16

The Disjunctive Conception of Experience as Material for a Transcendental Argument

John McDowell

1

In *Individuals* (1959) and *The Bounds of Sense* (1966), P. F. Strawson envisaged transcendental arguments as responses to certain sorts of scepticism. An argument of the sort Strawson proposed was to establish a general claim about the world, a claim supposedly brought into doubt by sceptical reflections. Such an argument was to work by showing that unless things were as they were said to be in the claim that the argument purported to establish, it would not be possible for our thought or experience to have certain characteristics, not regarded as questionable even by someone who urges sceptical doubts. So the argument's conclusion was to be displayed as the answer to a "How possible?" question. That has a Kantian ring, and the feature of such arguments that the formulation fits is the warrant for calling them "transcendental".

Barry Stroud responded to Strawson on the following lines.¹ Perhaps we can see our way to supposing that if our thought or experience is to have certain characteristics it does have (for instance that experience purports to be of a world of objects independent of us), we must *conceive* the world in certain ways (for instance as containing objects that continue to exist even while we are not perceiving them). But it is quite another matter to suggest that by reflecting about how it is possible that our thought and experience are as they are, we could establish conclusions not just about how we must conceive the world but about how the world must be. Stroud (1994/2000: 158–9) writes:

Even if we allow that we can come to see how our thinking in certain ways necessarily requires that we also think in certain other ways, and so perhaps in certain further ways as well, . . . how can truths about the world which appear to say or imply nothing about human thought or

¹ See Stroud's 1968 essay 'Transcendental Arguments', reprinted in his (2000). Several other essays in that collection are very helpful in clarifying the picture.

experience be shown to be genuinely necessary conditions of such psychological facts as that we think and experience things in certain ways, from which the proofs begin? It would seem that we must find, and cross, a bridge of necessity from the one to the other. That would be a truly remarkable feat, and some convincing explanation would surely be needed of how the whole thing is possible.

According to Stroud, Kant's explanation is transcendental idealism. As Stroud reads it, transcendental idealism explains how that "bridge of necessity" can be crossed by saying that the world of which the transcendently established claims are true is "only the 'phenomenal' world which is somehow 'constituted' by the possibility of our thought and experience of it".

Perhaps this might be better put by saying there is no bridge to cross. But then how satisfying a response to scepticism can be provided by such arguments? On this reading transcendental idealism does not so much respond to sceptical worries as brush them aside. Or perhaps it amounts to a concession that they are well placed. As Stroud (1994/2000: 159) puts it:

[T]here is the challenge of saying in what ways idealism is superior to, or even different from, the sceptical doctrines it was meant to avoid. How it differs, for example, from Hume's view that we simply cannot avoid believing that every event has a cause, and cannot help acting for all the world as if it were true, but that it is not really true of the world as it is independently of us.

And even if Stroud does not succeed in raising our suspicions of transcendental idealism, Strawson is anyway suspicious of it. In *The Bounds of Sense*, Strawson claims to preserve fundamental Kantian insights, but outside the idealist frame in which Kant formulated them. So Strawsonian transcendental arguments are expressly not equipped with what Stroud identifies as the Kantian apparatus for explaining how that "bridge of necessity" can be crossed. Stroud suggests, accordingly, that the Strawsonian arguments can yield only conclusions on the near side of the bridge. They uncover structural connections *within* our thought or experience, enabling us to argue that our thought or experience must be a certain way as a condition for the possibility of their being a certain other way.

That need not deprive the arguments of all force against scepticism. Suppose that whether things are a certain way comes within the scope of sceptical doubts. If we can establish that we must conceive things as being that way for it to be possible that our thought or experience has some characteristic that a sceptic would not or could not deny that it has, then we will have made some headway against that sceptical worry. This falls short of claiming to have shown that things must be that way for our thought and experience to be as they are. But with an argument of this more modest kind, we will have shown that, given the characteristic of our thought or experience that is the unquestioned starting point of the argument, there is no possibility of our being rationally required to discard the conviction that the sceptical argument was supposed to undermine.

Strawson has come to share Stroud's doubts about crossing that "bridge of necessity". It is not that he has given up the Kantian project, an inquiry into how it is possible that our thought and experience are as they are. But he has come to

approach the project in something like the way Stroud recommends, as tracing connections *within* how we conceive and experience things, rather than between how we conceive and experience things and how things must be. The aim of the investigation, as Strawson more recently sees it (Strawson 1985: 21), is to establish “a certain sort of interdependence of conceptual capacities and beliefs; for example, . . . that in order for self-conscious thought and experience to be possible, we must take it, or *believe*, that we have knowledge of external physical objects or other minds”.

2

This territory has been much worked over.² I am not going to work over it any more; I have sketched this picture of the state of play, in a certain region of recent discussion of transcendental arguments, only to bring out a contrast. I am not going to consider transcendental arguments of either of the two kinds that have come into view so far: neither the ambitious kind, in which the aim is to establish the truth of general claims about the world; nor the modest kind, in which the aim is to establish only that we cannot consistently go on taking it that our thought and experience are as they are in the relevant respects while withholding acceptance of the relevant claims about the world.

Instead I want to consider a different approach to one sort of scepticism. I want to suggest that this different approach can be pursued through a kind of transcendental argument that belongs to neither of those two types.

The scepticism in question is scepticism about perceptually acquired knowledge of the external world. And the approach in question is diagnostic. The diagnosis is that this scepticism expresses an inability to make sense of the idea of direct perceptual access to objective facts about the environment. What shapes this scepticism is the thought that even in the best possible case, the most that perceptual experience can yield falls short of a subject’s having an environmental state of affairs directly available to her. Consider situations in which a subject seems to see that, say, there is a red cube in front of her. The idea is that even if we focus on the best possible case, her experience could be just as it is, in all respects, even if there were no red cube in front of her. This seems to reveal that perceptual experience provides at best inconclusive warrants for claims about the environment. And that seems incompatible with supposing we ever, strictly speaking, *know* anything about our objective surroundings.³ The familiar sceptical scenarios—Descartes’s demon, the scientist with our brains in

² For a helpful survey, see Stern (2000).

³ Stroud regularly depicts scepticism about the external world as arising like this. See, for example, his (1996/2000: 131): “[The philosopher] chooses a situation in which any one of us would unproblematically say or think, for example, that we know that there is a fire in the fireplace right before us, and that we know it is there because we see that it is there. But when we ask what this seeing really amounts to, various considerations are introduced to lead us to concede that we would see exactly what we see now even if no fire was there at all, or if we didn’t know that there was one there.” See also Stroud (1984).

his vat, the suggestion that all our apparent experience might be a dream—are only ways to make this supposed predicament vivid.

Suppose scepticism about our knowledge of the external world is recommended on these lines. In that case it constitutes a response if we can find a way to insist that we *can* make sense of the idea of direct perceptual access to objective facts about the environment. That contradicts the claim that what perceptual experience yields, even in the best possible case, must be something less than having an environmental fact directly available to one. And without that thought, this scepticism loses its supposed basis and falls to the ground.

It is important that that is the right description of what this response achieves. We need not pretend to have an argument that would prove that we are not, say, at the mercy of Descartes's demon, using premises we can affirm, and inferential steps we can exploit, without begging questions against someone who urges sceptical doubts. As I said, the point of invoking the demon scenario and its like is only to give vivid expression to the predicament supposedly constituted by its not making sense to think we can have environmental facts directly available to us. But if it does make sense to think we can have environmental facts directly available to us, there is no such predicament. And now someone who proposes those scenarios can no longer seem to be simply emphasizing a discouraging fact about our epistemic possibilities. When we reject the scenarios—if we choose to bother with them at all—we need no longer be hamstrung by a conception of argumentative legitimacy controlled by that understanding of their status. An accusation of question-begging need no longer carry any weight. We can invert the order in which scepticism insists we should proceed, and say—as common sense would, if it undertook to consider the sceptical scenarios at all—that our knowledge that those supposed possibilities do not obtain is sustained by the fact that we know a great deal about our environment, which would not be the case if we were not perceptually in touch with the world in just about the way we ordinarily suppose we are.

Similarly, there is no need to establish, without begging questions against scepticism, that in any particular case of perceptual experience we actually are in the favourable epistemic position that scepticism suggests we could never be in. That would similarly be to accept tendentious ground rules for satisfying ourselves in given cases that we have knowledge of the environment. If we can recapture the idea that it is so much as possible to have environmental states of affairs directly presented to us in perceptual experience, we can recognize that such ground rules reflect a misconception of our cognitive predicament. And then our practice of making and assessing claims to environmental knowledge on particular occasions can proceed as it ordinarily does, without contamination by philosophy. There need no longer seem to be any reason to discount the fact that in real life the assessment is often positive.

Perhaps most people will find it obvious that reinstating the sheer possibility of directly taking in objective reality in perception would undermine a scepticism based

on claiming that perceptual experience can never amount to that. (I shall consider an exception later.)

But what does this have to do with transcendental arguments? Well, it depends on how the undermining move is defended. And it can be defended by an argument that is broadly Kantian, in the sense in which the arguments I was considering at the beginning are broadly Kantian. The argument aims to establish that the idea of environmental facts making themselves available to us in perception must be intelligible, because that is a necessary condition for it to be intelligible that experience has a characteristic that is, for purposes of this argument, not in doubt.

The relevant characteristic is that experience purports to be of objective reality. When one undergoes perceptual experience, it at least appears to one as if things in one's environment are a certain way.

Consider Wilfrid Sellars's discussion of "looks" statements in *Empiricism and the Philosophy of Mind* (Sellars 1997: 32–53). Sellars urges something on the following lines. In order to understand the very idea of the objective purport of visual experience (to single out one sensory modality), we need to appreciate that the concept of experiences in which, say, it looks to one as if there is a red cube in front of one divides into the concept of cases in which one sees that there is a red cube in front of one and the concept of cases in which it merely looks to one as if there is a red cube in front of one (either because there is nothing there at all or because although there is something there it is not a red cube).

At least implicit here is a thought that can be put as follows. In order to find it intelligible that experience has objective purport at all, we must be able to make sense of an epistemically distinguished class of experiences, those in which (staying with the visual case) one sees how things are—those in which how things are makes itself visually available to one. Experiences in which it merely looks to one as if things are thus and so are experiences that misleadingly present themselves as belonging to that epistemically distinguished class. So we need the idea of experiences that belong to the epistemically distinguished class if we are to comprehend the idea that experiences have objective purport. If one acknowledges that experiences have objective purport, one cannot consistently refuse to make sense of the idea of experiences in which objective facts are directly available to perception.

The scepticism I am considering purports to acknowledge that experiences have objective purport, but nevertheless supposes that appearances as such are mere appearances, in the sense that any experience leaves it an open possibility that things are not as they appear. That is to conceive the epistemic significance of experience as a highest common factor of what we have in cases in which, as common sense would put it, we perceive that things are thus and so and what we have in cases in which that merely seems to be so—so never higher than what we have in the second kind of case.⁴ The conception I have found in Sellars can be put, in opposition to that, as a disjunctive conception of perceptual appearance: perceptual appearances are either objective

⁴ On the idea of the highest common factor, see, for example, my (1994: 113).

states of affairs making themselves manifest to subjects, or situations in which it is as if an objective state of affairs is making itself manifest to a subject, although that is not how things are.⁵

Experiences of the first kind have an epistemic significance that experiences of the second kind do not have. They afford opportunities for knowledge of objective states of affairs. According to the highest common factor conception, appearances can never yield more, in the way of warrant for belief, than do those appearances in which it merely seems that one, say, sees that things are thus and so. But according to the Sellarsian transcendental argument, that thought undermines its own entitlement to the very idea of appearances.

The highest common factor conception is supposedly grounded on a claim that seems unquestionable: the claim that from a subject's point of view, a misleading appearance can be indistinguishable from a case in which things are as they appear. That might be taken as a self-standing claim about the phenomenology of misleading appearance, available to be cited in explaining the fact that subjects can be misled by appearances. So taken, the claim is open to dispute (Austin 1962). But the right way to take it is as simply registering the fact that, on that interpretation, it is supposed to explain: the undeniable fact that our capacity to get to know things through perception is fallible.⁶

The claim of indistinguishability is supposed to warrant the thought that even in the best case in which a subject, say, has it visually appear to her that there is a red cube in front of her, her experience could be just as it is even if there were no red cube in front of her. But we need a distinction here. When we say her experience could be just as it is even if there were no red cube in front of her, we might just be registering that there could be a misleading experience that from the standpoint of her experience she could not distinguish from her actually veridical experience. In that case what we say is just a way of acknowledging that our capacity to acquire knowledge through perceptual experience is fallible. It does not follow that even in the best case, the epistemic position constituted by undergoing an experience can be no better than the epistemic position constituted by undergoing a misleading experience, even one that would admittedly be indistinguishable. The acknowledgement of fallibility cannot detract from the excellence of an epistemic position, with regard to the obtaining of an objective state of affairs, that consists in having the state of affairs present itself to one in one's perceptual experience. This is where the disjunctive conception does its epistemological work. It blocks the inference from the subjective indistinguishability of experiences to the highest common factor conception, according to which neither of the admittedly indistinguishable experiences could have higher epistemic worth than that of the inferior case. And the transcendental argument shows

⁵ On the disjunctive conception, see Hinton (1973), Snowdon (1980–1), and my (1982) and (1986).

⁶ I have revised what I first wrote in this connection, partly in response to an objection from Costas Pagondiotis. I have been influenced here by Sebastian Rödl.

that the disjunctive conception is required, on pain of our losing our grip on the very idea that in experience we have it appear to us that things are a certain way.⁷

4

This transcendental argument starts from the fact that perceptual experience at least purports to be of objective reality, and yields the conclusion that we must be able to make sense of the idea of perceptual experience that is actually of objective reality. I have urged that that is enough to undermine a familiar sort of scepticism about knowledge of the external world.

Now there may be a temptation to object that this argument assumes too much. Should it be left unquestioned that perceptual experience purports to be of objective reality?

There is plenty of room to argue that it is proper to start there. The sceptical arguments Descartes considers, for instance, do not question the fact that perceptual experience yields appearances that things are objectively the case. Descartes's arguments question only our entitlement to believe that things are as they appear to be. The highest common factor conception owes its attractiveness to the subjective indistinguishability of experiences all of which can be described in terms of the appearance that things are objectively thus and so. This supposed basis for scepticism does not need a more minimal picture of experience.

But what if we do decide that we ought to confront a more wholehearted scepticism, a scepticism willing to doubt that perceptual experience purports to be of objective reality? Well then, the transcendental argument I have been considering cannot do all the work. But it can still do some of the work. If this is the target, we need a prior transcendental argument, one that reveals the fact that consciousness includes states or episodes that purport to be of objective reality as a necessary condition for some more basic feature of consciousness, perhaps that its states and episodes are potentially self-conscious. Strawson's reading of the Transcendental Deduction in Kant's first Critique might serve, or perhaps the Transcendental Deduction itself. It would take me too far afield to go into this here. The point is just that we cannot dismiss an argument that pivots on the disjunctive conception of perceptual appearance, on the ground that it does not itself establish the characteristic of perceptual experience that it begins from.

⁷ The essential thing is that the two sides of the disjunction differ in epistemic significance, whereas on the highest common factor conception the "good" disjunct can afford no better warrant for perceptual claims than the "bad" disjunct. This difference in epistemic significance is of course consistent with all sorts of commonalities between the disjuncts. For instance, on both sides of the disjunction it appears to one that, say, there is a red cube in front of one. In his (2002: 341, fn. 12, and associated text), Crispin Wright makes needlessly heavy weather of this.

In a recent paper, Crispin Wright (2002: 331) argues that as a response to scepticism, replacing the highest common factor conception of perceptual experience with a disjunctive conception is “dialectically quite ineffectual”.

Wright starts from a helpful account of why G. E. Moore’s “proof of an external world”—at least if taken at face value—is as unimpressive as nearly everyone finds it.⁸ Moore moves from the premise “Here is a hand” to the conclusion, which is indeed entailed by that premise, that there is an external world. Wright takes Moore to suppose that his premise is itself grounded on something yet more basic: something Moore could express by saying “My experience is in all respects as of a hand held up in front of my face”. And Wright’s diagnosis of what goes wrong in Moore’s argument is that the warrant this ground supplies cannot be transmitted across the acknowledged entailment from “Here is a hand” to “There is an external world”. The warrant that “My experience is as of a hand” provides for “Here is a hand” is defeasible, and it is defeated if the sceptic is right and we are, for instance, at the mercy of Descartes’s demon. We can allow it to warrant the premise of Moore’s entailment only if we already take ourselves to be entitled to accept the conclusion of the entailment. So the whole argument is question-begging.

Wright now turns to the disjunctive conception. He sums up his verdict on it as follows (2002: 346–7):

In brief: whether our perceptual faculties engage the material world directly [the thesis that the disjunctive conception is aimed at protecting] is one issue and whether the canonical justification of perceptual claims proceeds through a defeasible inferential base is another. One is, so far, at liberty to take a positive view of both issues. And when we do, the I-II-III pattern [the pattern of Moore’s argument, augmented with a formulation of the ground for the premise of Moore’s entailment] re-emerges along these lines:

I *Either* I am perceiving a hand in front of my face *or* I am in some kind of delusional state
 II Here is a hand
 Therefore
 III There is a material world.

It is clear that this is a mere variation on Moore’s argument as Wright reconstructs it. In this version too, the support I provides for II is defeasible. That we take it not to be defeated depends on our already taking ourselves to be entitled to accept III. So it would be question-begging to suppose the argument provides any support for III.

But what does this have to do with the disjunctive conception? The point of the disjunctive conception is that if one undergoes an experience that belongs on the “good” side of the disjunction, that warrants one in believing—indeed presents one with an opportunity to know—that things are as the experience reveals them as being. When one’s perceptual faculties “engage the material world directly”, as Wright puts it, the result—a case of having an environmental state of affairs directly present to one

⁸ Moore may intend something more subtle. But I shall not consider this possibility.

in experience—constitutes one's being justified in making the associated perceptual claim. It is hard to see how any other kind of justification could have a stronger claim to the title "canonical". And this justification is *not* defeasible. If someone sees that P, it cannot fail to be the case that P. So if one accepts the disjunctive conception, one is *not* at liberty to go on supposing that "the canonical justification of perceptual claims proceeds through a defeasible inferential base".

In urging the contrary, Wright constructs an argument whose starting point is the whole disjunction. Of course he is right that the whole disjunction could provide at best defeasible support for a perceptual claim. But what he has done is in effect to cast the whole disjunction in the role in which the supposed case for scepticism casts the highest common factor. And the point of the disjunctive conception is precisely to reject the highest common factor picture of the justification for perceptual claims.

I do not mean to suggest that a I-II-III argument starting from the "good" disjunct would be any more impressive as an augmentation of Moore's "proof" than the I-II-III argument Wright considers, starting from the whole disjunction. I shall come to that in a moment. The point for now is that Wright is wrong to claim that the disjunctive conception leaves one free to think perceptual claims rest on defeasible inferential support.

What has gone wrong here?

Wright apparently assumes that a dialectically effective response to scepticism would need to be what Moore—again, if we take his performance at face value—tries to produce: that is, an argument that directly responds to the sceptic's questioning whether there is an external world. Such an argument would need to start from a premise available without begging a question against the sceptic, and it would need to transmit warrant legitimately from that premise to the conclusion that there is indeed an external world. And only the whole disjunction is non-question-beggingly available as a premise for such an argument.

But the point of the disjunctive conception is not to improve our resources for such arguments.

At one point (341) Wright acknowledges, in a way, that when I appeal to the disjunctive conception I do not claim to be directly answering sceptical questions. The acknowledgement is backhanded, since Wright describes my disclaimer as "an official refusal to take scepticism seriously". It is worth pausing over this description. The wording would be appropriate if in order to take scepticism seriously one had to attempt direct answers to sceptical questions. But that seems simply wrong. Surely no one takes scepticism more seriously than Stroud. And Stroud (1980: 56) thinks "the worst thing one can do with the traditional question about our knowledge of the world is to try to answer it".

Wright notes my suggestion that the disjunctive conception "has the advantage of removing a prop on which sceptical doubt . . . depends", as he puts it. But he treats this as a mere lapse from the "official refusal", as if removing a prop could only be offering an answer to a sceptical question. Only on that assumption could noting the inefficacy of the re-emergent I-II-III argument, the argument that starts from the

whole disjunction, seem relevant to the anti-sceptical credentials of the disjunctive conception.

The disjunctive conception cannot improve on Moore in the project of proving that there is an external world. Wright is correct about that.

This is not, as Wright has it, because the disjunctive conception allows us to go on holding that “the canonical justification of perceptual claims proceeds through a defeasible inferential base”. As I have insisted, the disjunctive conception is flatly inconsistent with that thesis. The canonical justification for a perceptual claim is that one perceives that things are as it claims they are, and that is not a defeasible inferential base.

The point is, rather, that if one lets the sceptic count as having put in doubt whether there is an external world in which things are pretty much as we take them to be, it becomes question-begging to take oneself, on any particular occasion, to *have* the *indefeasible* warrant, for a claim such as “Here is a hand”, constituted by, for instance, seeing that there is a hand in front of one. In the dialectical context of an attempt to show that the sceptical scenarios do not obtain, the indefeasible warrant for “Here is a hand” constituted by seeing that there is a hand in front of one can no more be transmitted across the entailment to “There is a material world” than can the defeasible warrant Wright considers in his diagnosis of Moore. In Moore’s argument as Wright reconstructs it, the fact that the warrant’s support for “Here is a hand” is not defeated depends on our already taking ourselves to have grounds for the conclusion supposedly reached by entailment from there. In the argument I am considering now, our conviction that we have the warrant at all depends on our already taking ourselves to have grounds for the conclusion. This, incidentally, suggests a different account, which seems no less plausible than Wright’s, of the implicit warrant for the premise Moore actually starts from. In any case, whether or not it is what Moore has in mind, an argument that starts from one’s seeing a hand in front of one would be just as useless for Moore’s purpose—if, again, we identify his purpose by taking his performance at face value.

But all this is irrelevant to the anti-sceptical power of the disjunctive conception. What the disjunctive conception achieves is indeed to remove a prop on which sceptical doubt depends. That is Wright’s wording, but he does not allow it to carry its proper force. The prop is the thought that the warrant for a perceptual claim provided by an experience can never be that the experience reveals how things are. The disjunctive conception dislodges that thought, and a sceptical doubt that depends on it falls to the ground. There is no need to do more than remove the prop. In particular, as I explained before, there is no need to try to *establish* theses like the conclusion of Moore’s argument, with the ground rules for doing so set by scepticism. The idea that such theses are open to doubt now lacks the cachet of simply emphasizing an epistemic predicament constituted by its being impossible for experience to reveal to us how things are. There is no such predicament, and now it is perfectly proper to appeal to cases of ordinary perceptual knowledge in ruling out the sceptical scenarios, or—better—in justifying a common-sense refusal to bother with them.

Wright might be tempted to seize on what I have just said as vindicating his talk of my “official refusal to take scepticism seriously”. But like Stroud, I hold that the way

to take scepticism seriously is not to try to disprove the sceptical scenarios. We take scepticism seriously by removing the prop, thereby entitling ourselves to join common sense in refusing to bother with the sceptical scenarios.⁹

Considering the form “*Either* I am perceiving thus-and-such *or* I am in some kind of delusional state”, Wright (2002: 346) offers this reconstruction of the sceptical reasoning that, according to him, survives the disjunctive conception:

[I]n this case it is our practice to treat one in particular of the disjuncts as justified—the left-hand one—whenever the disjunction as a whole is justified and there is, merely, *no evidence* for the other disjunct! That’s a manifest fallacy unless the case is one where we have a standing reason to regard the lack of any salient justification for a disjunct of the second type as a reason to discount it. And—the sceptical thought will be—it’s hard to see what could count as a standing reason except a prior entitlement to the belief that delusions are rare. But that’s just tantamount to the belief that there is a material world which, at least on the surfaces of things, is pretty much revealed for what it is in what we take to be normal waking experience. So, the Sceptic will contend, that broad conception once again emerges as a rational precondition of our practice, even after the disjunctive adjustment to the concept of perception; and on its warrantedness depends whatever warrant can be given for our proceeding in the way we do. Since it cannot be warranted by appeal to the warrant for specific perceptual claims—Moore’s proof being no better in this setting than before—the Sceptic may now focus on the apparent impossibility of any kind of direct warrant for it, and the dialectic can proceed essentially as before.

It is clearly correct that our practice of assessing the credentials of perceptual claims could not be rational if we were not entitled to the “broad conception” according to which the external world is pretty much the way we take ourselves to experience it as being. But it is tendentious to suppose it follows that the rationality of our practice is in jeopardy unless the “broad conception” can be warranted *in advance of the practice* without begging questions against scepticism. And it is wrong to suppose the disjunctive conception leaves unchallenged the idea Wright here exploits, that the justification for a perceptual claim must go through the whole disjunction, exploiting some supposed standing reason for discounting the “bad” disjunct. The justification for a perceptual claim is an entitlement to the “good” disjunct. What entitles one to that is not that one’s experience warrants the whole disjunction, plus some supposed ground for discounting the “bad” disjunct. That would commit us to trying to reconstruct the epistemic standing constituted by perceiving something to be the case in terms of the highest common factor conception of experience, plus whatever ground we can think of for discounting the “bad” disjunct. I think Wright is correct that that is hopeless; if we see things this way, the sceptic wins. But the disjunctive conception eliminates the apparent need for any such project, because it contradicts the highest common factor conception.

⁹ In writing here of a common-sense refusal to bother with the sceptical scenarios, I am echoing a remark at McDowell (1994: 113), in the passage Wright cites to document the “official refusal”: “The aim here is not to answer sceptical questions, but to begin to see how it might be intellectually respectable to ignore them, to treat them as unreal, in the way that common sense has always wanted to.” Of course it takes work to reach such a position. This attitude can look like a “refusal to take scepticism seriously” only given the picture of what it is to take scepticism seriously that Stroud rejects.

What does entitle one to claim that one is perceiving that things are thus and so, when one is so entitled? The fact that one is perceiving that things are thus and so. That is a kind of fact whose obtaining our self-consciously possessed perceptual capacities enable us to recognize on suitable occasions, just as they enable us to recognize such facts as that there are red cubes in front of us, and all the more complex types of environmental facts that our powers to perceive things put at our disposal.

Of course we are fallible about the obtaining of such facts, just as we are fallible about the facts we perceive to obtain. I can tell a zebra when I see one—to take up an example Wright (2002: 342–4) borrows from Fred Dretske. If what I believe to be a zebra is actually a cunningly painted mule, then of course I do not recognize it as a zebra, as I suppose, and I do not have the warrant I think I have for believing it is a zebra, namely that I see it to be a zebra. My ability to recognize zebras is fallible, and it follows that my ability to know when I am seeing a zebra is fallible. It does not follow—this is the crucial point—that I cannot ever have the warrant for believing that an animal in front of me is a zebra constituted by seeing that it is a zebra. If the animal in front of me is a zebra, and conditions are suitable for exercising my ability to recognize zebras when I see them (for instance, the animal is in full view), then that ability, fallible though it is, enables me to see that it is a zebra, and to know that I do. My warrant is not limited to the disjunction “Either I see that it is a zebra or my visual experience is misleading in some way”. That is the highest common factor conception, and fallibility in our cognitive capacities cannot force it on us.¹⁰

6

Transcendental arguments of Stroud’s ambitious type aim to establish large-scale features the world must have for it to be possible that thought and experience are as they are. Those of his modest type aim to establish large-scale features we must conceive the world to have for it to be possible that thought and experience are as they are.

The argument I have considered belongs to neither of these types. It does not offer to establish anything about how things are, let alone must be, in the world apart from us, so it is not vulnerable to Stroud’s doubts about arguments of the ambitious type. But the way it makes itself immune to those doubts is not by weakening its conclusion to one about structural features we must conceive the world to have. The conclusion is rather one about how we must conceive the epistemic positions that are within our reach, if it is to be possible that our experience is as it is in having objective purport. That frees us to pursue our ordinary ways of finding out how things are in the world apart from us. The specifics of what we go on to find out are not within the scope of what the argument aims to vindicate.

¹⁰ A misconception of the significance of fallibility on these lines is the topic of the passage in my (1994: 112–13) that Wright comments on in his (2002: 341, fn. 13). His remarks there seem to me to miss, or ignore, the dialectical context of the passage on which he is commenting.

That might seem to distance this argument from much in Kant, who is presumably the patron saint of transcendental arguments. In sketching the argument, I have not needed to connect it with the question 'How is synthetic a priori knowledge possible?', or with an investigation of the principles of the pure understanding. But there is still the fact that the argument displays its conclusion as a necessary element in the answer to a 'How possible?' question about experience. Moreover, Sellars's account of how experience has its objective purport, which the argument exploits, is strikingly Kantian, in the way it represents the content of an experience as the content of a claim. Sellars links the fact that experience is of objective reality with the fact that to make a claim is to commit oneself to things being objectively thus and so. This talk of claims is Sellars's counterpart, after the "linguistic turn", to Kant's invocation of judgement. So perhaps the argument I have been considering can be seen as belonging to a minimal Kantianism. In the argument's background is an explanation of the objective purport of experience in terms of the fact that experience exemplifies forms that belong to the understanding. But in the argument as I have considered it so far, we exploit that Kantian thought without needing to concern ourselves either with how the world must be or with how we must conceive the world to be. Of course this is not the place to try to take this any further.

REFERENCES

- Austin, J. L. (1962) *Sense and Sensibilia*, Oxford: Clarendon Press.
- Hinton, J. M. (1973) *Experiences*, Oxford: Clarendon Press.
- McDowell, J. (1982) 'Criteria, Defeasibility, and Knowledge', in his *Meaning, Knowledge, and Reality*, Cambridge, MA: Harvard University Press, 1998.
- (1986) 'Singular Thought and the Extent of Inner Space', in his *Meaning, Knowledge, and Reality*, Cambridge, MA: Harvard University Press, 1998.
- (1994) *Mind and World*, Cambridge, MA: Harvard University Press.
- Sellars, W. (1997) *Empiricism and the Philosophy of Mind*, with introduction by R. Rorty and Study Guide by R. Brandom, Cambridge, MA: Harvard University Press.
- Snowdon, P. (1980–1) 'Perception, Vision, and Causation', *Proceedings of the Aristotelian Society*, 81: 175–92.
- Strawson, P. F. (1959) *Individuals: An Essay in Descriptive Metaphysics*, London: Methuen.
- (1966) *The Bounds of Sense*, London: Methuen.
- (1985) *Skepticism and Naturalism: Some Varieties*, New York: Columbia University Press.
- Stern, R. (2000) *Transcendental Arguments and Scepticism: Answering the Question of Justification*, Oxford: Clarendon Press.
- Stroud, B. (1968) 'Transcendental Arguments', in his *Understanding Human Knowledge*, Oxford: Clarendon Press, 2000.
- (1980) 'Reasonable Claims: Cavell and the Tradition', in his *Understanding Human Knowledge*, Oxford: Clarendon Press, 2000.

- (1984) *The Significance of Philosophical Scepticism*, Oxford: Clarendon Press.
- (1994) 'Kantian Argument, Conceptual Capacities, and Invulnerability', in his *Understanding Human Knowledge*, Oxford: Clarendon Press, 2000.
- (1996) 'Epistemological Reflection on Knowledge of the External World', in his *Understanding Human Knowledge*, Oxford: Clarendon Press, 2000.
- Wright, C. (2002) 'Anti-Sceptics Simple and Subtle: G. E. Moore and John McDowell', *Philosophy and Phenomenological Research*, 65: 330–48.

Comment on John McDowell's 'The Disjunctive Conception of Experience as Material for a Transcendental Argument'

Crispin Wright

1 THE DISJUNCTIVE CONCEPTION OF EXPERIENCE

Descartes was surely right that while normal waking experience, dreams, and hallucinations are characteristically distinguished at a purely phenomenological level—by contrasts of spatial perspective, coherence, clarity of image, etc.—it is not essential that they be so.¹ What it is like for someone who dreams that he is sitting, clothed in his dressing gown, in front of his fire can in principle be subjectively indistinguishable from what it is like to perceive that one is doing so, fully conscious and awake. The same holds for multi-sense hallucination and, it is assumed, would hold of the experience of an envatted brain in the usual postulated scenario.

This thought—that normal perceptual experience allows in principle of perfect phenomenological counterfeit—is, to the best of my knowledge, nowhere seriously challenged in John McDowell's writings.² What he rejects is an idea that builds upon

Thanks to Annalisa Coliva, Lars Gundersen, Adrian Haddock, Aidan McGlynn, Duncan Pritchard, and Elia Zardini for helpful comments and criticisms.

¹ Descartes writes: "Nevertheless, I must remember that I am a man, and that consequently I am accustomed to sleep and in my dreams to imagine the same things that lunatics imagine when awake, or sometimes things which are even less plausible. How many times has it occurred that the quiet of the night made me dream of my usual habits: that I was here, clothed in a dressing gown, and sitting by the fire, although I was in fact lying undressed in bed! It seems apparent to me now, that I am not looking at this paper with my eyes closed, that this head that I shake is not drugged with sleep, that it is with design and deliberate intent that I stretch out this hand and perceive it. What happens in sleep seems not at all as clear and distinct as all this. But I am speaking as though I never recall having been misled, while asleep, by similar illusions. When I consider these matters carefully, I realise so clearly that *there are no conclusive indications by which waking life can be distinguished from sleep* that I am quite astonished, and my bewilderment is such that it is almost able to convince me that I am sleeping" (my italic). The passage is from *First Meditation*; see pp. 145–6 of the Haldane and Ross edition of *The Philosophical Works of Descartes*, vol. 1 (Cambridge: Cambridge University Press, 1968).

² Though see the text associated with (McDowell, this volume: fn. 6).

and would be potentially explanatory of it: the Lockean idea that, as far as the states enjoyed by the experiencing subject are concerned, there is actually no generic distinction between dream, hallucination, and wakeful perception—that one and the same *type* of state of consciousness is involved in all three cases, and that which (if any) of the three a particular occurrence of the type falls under is a matter of how it is caused. On this model, the distinction between Descartes's notional fully lucid dream and the corresponding raft of perceptions of his dressing-gowned, sedentary state is like that between a certain kind of sunburn and nettle rash. For Disjunctivism, however, the Lockean—as McDowell likes to say, 'Highest Common Factor'—conception of perceptual experience and its potential counterfeits is a *conceptual* error. There is no single type of state of consciousness present in each of dreaming, perceiving, and hallucinating, whose instances fall under one or other of those characterizations purely by virtue of their aetiology. In perception one is directly aware of the items perceived—the states of mind involved are essentially external world involving.³ The various relevant kinds of counterfeit states, by contrast, involve no essential relation to the kinds of thing of which, in perception, one is directly aware. Subjectively, they can be such as to allow of being mistaken for states involving the kind of external world relationship essential to perception. But this mistakability, however tempting its explanation by means of the idea of "Highest Common Factor"—a common front presented by both perceptual awareness and its counterfeits—does not require that idea. A waxwork may, in certain circumstances, be mistakable for a living human being, but it is no part of being human to have a waxwork component.⁴

This is all familiar. What is new in McDowell's present contribution is more about how he thinks the replacement of the Lockean conception of perceptual experience with the Disjunctive Conception helps with the treatment of sceptical doubt, and about how he conceives of the authority for the Disjunctive Conception itself. In what follows I'll briefly review McDowell's offerings on both these points, before turning to his critical reactions to the misgivings expressed in my (2002) (henceforth "Simple and Subtle") concerning the bearing of Disjunctivism on the sceptical paradoxes.

2 WHAT ENFORCES THE DISJUNCTIVE CONCEPTION?

McDowell (this volume) offers a "transcendental argument" on behalf of the Disjunctive Conception, albeit one he is at pains to distinguish from two older models.

³ It might be rejoined that so much is also true for Locke: for Locke too, perceptual states are essentially external-world-involving, since the concept of a perceptual state is one which invokes an external worldly cause. That is true. Where Locke and the Disjunctivist differ is over the claim that the *state of mind* involved in perception is one which essentially has an external worldly cause.

⁴ To insist that perception and, say, dreaming have no common factor in the sense gestured at is of course, consistent with recognizing that they can present "all sorts of commonalities" (McDowell, this volume: fn. 7)—including, crucially, the presentation of appearances that P. The vital point for McDowell is that the "epistemic significance" of perceptual appearances differs from that of appearances within a dream. It is a question how exactly this differential epistemic significance is supposed to follow from the disjunctivist conception as just outlined. I will argue later that sceptical doubt on the point has no need of the Lockean conception.

Transcendental arguments of the traditional, Kantian stripe aim to show that it is a necessary condition for experience to have a certain—uncontested—characteristic that the real world should have certain (perhaps sceptically contested) characteristics. A notionally more modest kind of transcendental argument, canvassed by Barry Stroud as an antidote to the failure of nerve that the Kantian model may inspire, aims to show that it is a commitment of our *thinking* of experience as having a certain—normally uncontested—characteristic that we *should conceive* of the real world as having some (perhaps sceptically contested) characteristic.⁵ The idea is that a successful example of this latter kind of argument might still be of some anti-sceptical effect if the characteristic of experience in question is one on which a given sceptical argument itself depends in some way. McDowell distinguishes his new transcendental argument from each of these forms.⁶ Its premise is, once again, that our (perceptual) experience has a certain feature, to wit: *objective purport*. But its advertised conclusion is nothing about the world, nor about any respect in which we are, supposedly, committed to conceiving of the world. Rather, it is the Disjunctive Conception of experience itself—or more specifically, it is the conception of perceptual experience as consisting in a direct apprehension of characteristics and relations exemplified by items in the local material environment.⁷

In McDowell's view, then, the very possession of "objective purport" by episodes of our experience enforces a conception of them as, in the best possible case, episodes of direct awareness of the external world. So the Lockean conception, which holds that they are never that, is refuted unless there is a case to be made that experience lacks objective purport. The crux, of course, is what is to be understood by "objective purport". McDowell seems to regard it as sufficient to settle the matter that our experience involves *representational* states whose contents are articulated in material conceptual vocabulary. To (seem to) perceive the material environment in my office is to receive

⁵ Stroud (1968).

⁶ There are some puzzles about the intended relative modesty of the Stroud model. How is it to be argued that *our taking it that* that experience has a certain feature commits us to conceiving of the world as being a certain way except by showing that experience's *having* the feature entails that the world is *in fact* that way—so we are taken back to the more ambitious Kantian form of argument? Moreover, why would the effect of any successful argument of the Stroud type not be merely to deepen the aporia generated by the sceptical paradox? For on the one hand we shall then have a transcendental argument that conceiving of experience as having a certain feature commits us to P; and on the other hand, a sceptical argument that somehow depends exactly on that same conception yet appears to show that there is no warrant for P. That hardly seems like progress unless the idea is merely to silence an actual sceptical adversary in debate—'See: you are yourself committed to endorsing the conclusion you claim is unwarranted.' But the fact is that the dialectical situation is not that of confrontation with any such adversary: we are dealing with *paradox*. If the plausible-seeming materials that the paradox exploits themselves engender commitments which it seems to undermine, that deepens rather than solves it. But I will not pursue these issues further here. See Wright (1991) on "the Adversarial Stance".

⁷ A picky reader might wonder why McDowell's argument doesn't actually perfectly fit the Kantian template. Premise about experience: experience has objective purport. Conclusion about the external world: the world is such as to present situations of which we can, in the best case, be directly aware. The answer, I suppose, is that McDowell's argument only gives a weakened conclusion: the external world, *if such there be*, is such as to present situations of which we can, in the best case, be directly aware. Still, it is not clear why it doesn't fit the weaker, Stroudian template—McDowell's emphasis on its alleged novelty of form is a puzzle.

a succession of representations that P—for example, that the screensaver on my computer is active, that there are some gaps on the bookshelf opposite, etc.—where the concepts configured in the relevant instances of ‘P’ are concepts of material things and their observable characteristics. McDowell’s idea appears to be that if—as he takes it to be so—perceptual experience does indeed consist in the reception of appearances that P, where P is a content of this kind, then there is no alternative but to conceive of the (in principle possible) successful case—where appearance matches reality—as consisting in a direct awareness of the latter.

It is much less my purpose here to assess this idea than to review the work to which McDowell wants to put his conclusion. Nevertheless, one wonders whether so simple a train of thought can possibly do serious damage to a way of thinking about perception that, though I have characterized it as Lockean, is arguably as old as Plato’s Cave. McDowell (this volume: section 4) does allow for a possible challenge to the claimed objective purport of experience, “a scepticism willing to doubt that perceptual experience purports to be of objective reality” and concedes that in that case his transcendental argument “cannot do all the work”, so that “we need a prior transcendental argument, one that reveals the fact that consciousness includes states or episodes that purport to be of objective reality as a necessary condition for some more basic feature of consciousness, perhaps that its states and episodes are potentially self-conscious”. But now it is crucial what objective purport really amounts to. No Lockean, after all, ever denied that we are accustomed to *use* concepts of material reality to describe the flickerings of coloured light and shadow on the screen of inner experience. McDowell’s idea has presumably to be that its carrying of contents articulated in terms of concepts of external material reality is somehow *essential* to perceptual experience—that such content is not a matter of *interpretation* (as it would be for Plato’s Cave dwellers). Since some of the most ancient of sceptical concerns about our relationship to the external are precisely rooted in the opposed, interpretative conception, it may seem that much of the real action is indeed focused on the unargued premise for McDowell’s transcendental argument, and that the latter’s significance is therefore modest at best.

Actually—on the most straightforward construal of objective purport—it is less than modest. There is, on that construal, a decisive worry about the ability of McDowell’s transcendental argument to reach its intended conclusion. For if it suffices for objective purport simply to involve the occurrence of appearances that P for instances of ‘P’ articulated in terms of concepts of external material reality, then dreams and hallucinations have objective purport too. And in that case its possession of objective purport no longer provides a bridge to the idea that a certain kind of experience is apt to constitute *direct awareness* of the kind of objective material it purports to represent—since such direct awareness is exactly what dreams and hallucinations are essentially *not* apt to constitute.

There is significant reason to doubt that McDowell’s present contribution accomplishes much to motivate the Disjunctive Conception. In what follows, however, I am going to assume that a relevant (if you will, transcendental) argument can somehow be made to run, so that we are impelled by unquestioned features of our experience to conceive of perceptual experience both as direct awareness of the external

and as possible. Our concern will be with the prospects for the continuation of challenging forms of sceptical paradox even in this setting: with whether Disjunctivism really does “dislodge a prop” on which sceptical doubt depends, or whether—to give the reader a peek at the story’s end—it at best presents temporary obstacles to its formulation.

3 DISJUNCTIVISM AND CARTESIAN SCEPTICISM

McDowell (this volume: section 2) thinks that an enforcement of the Disjunctive Conception is enough to undermine a scepticism that

expresses an inability to make sense of the idea of direct perceptual access to objective facts about the environment. What shapes this scepticism is the thought that even in the best possible case, the most that perceptual experience can yield falls short of a subject’s having an environmental state of affairs directly available to her. Consider situations in which a subject seems to see that, say, there is a red cube in front of her. The idea is that even if we focus on the best possible case, her experience could be just as it is, in all respects, even if there were no red cube in front of her. This seems to reveal that perceptual experience provides at best inconclusive warrants for claims about the environment. And that seems incompatible with supposing that we ever, strictly speaking, know anything about our objective surroundings.

One way of putting the point would be to say that the Lockean conception of experience enforces the idea that, even in the best case, perceptual knowledge, if possible at all, involves a defeasible inference from matters in one’s own *cognitive locality*—matters of which one *is* capable of a direct cognitive apprehension—to matters that lie outside it.⁸ So the Lockean conception sets up the framework for a type of sceptical argument—*Humean* scepticism, as expressed in what I have termed the I-II-III argument—whose focus is on this strict insufficiency of our evidence for a certain kind of conclusion and on the epistemic status of the collateral presuppositions whose acceptance seems to be required in order to rationalize the ampliative evidential transitions in question.⁹

The Lockean conception certainly provides a platform for a sceptical challenge of this kind. It is, however, another question whether it provides the only possible platform for a I-II-III version of external world scepticism. The principal contention of “Simple and Subtle” to which McDowell takes exception was that this is not so—that a sceptical paradox of essentially the same structure can be developed even under the aegis of the Disjunctive Conception. I shall discuss the grounds of McDowell’s resistance to this claim below. But before we come to that, it merits emphasis that the strongest effect that a shift to the Disjunctive Conception can possibly have, in point of provision of anti-sceptical weaponry, is to extend the sphere of our *cognitive locality*—the sphere of matters directly accessible to our cognitive powers—beyond the usual inner Cartesian limits to include the local material environment. Such a shift,

⁸ For more on ‘cognitive locality’ see Wright (2004: 172–4 and 201–2).

⁹ All this is set out in detail in “Simple and Subtle” at pp. 337–40 and Wright (2004: 169–75).

accordingly, can accomplish nothing against those forms of sceptical argument that are precisely designed to address what we are normally pleased to regard as our epistemic achievements concerning the cognitively local. It is, in my opinion, a shortcoming of much of McDowell's discussion of these matters that he tends to write of scepticism as if it were a "frame of mind", or a kind of rootless anxiety or preoccupation, when the truth—as manifested in the recent and contemporary discussion—is that we have to deal with a number of specific, sharply formulable paradoxes, differing in detail in significant ways. One such form of paradox, designed specifically for the service of scepticism concerning the cognitively local, is the version of the Cartesian dreaming argument formulated and discussed in my (1991). An outline of its essential structure is offered in the Appendix. A version of this argument will engage the presumed products of *any* form of cognitive activity which allows in principle of subjectively wholly convincing counterfeit. Since that much is not challenged by the Disjunctive Conception, our claims to perceptual knowledge fall within the scope of the Cartesian paradox—as does memory, on direct realist construals of that faculty, and, especially disconcertingly, intellection.

As noted, the form of perceptual scepticism against which McDowell believes that his transcendental argument might be of some purpose is diagnosed as expressing "an inability to make sense of the idea of direct perceptual access to objective facts about the environment" (McDowell, this volume: section 2). If the Disjunctive Conception can be enforced, then that idea *does* make sense. And he dismisses the "familiar sceptical scenarios"—Descartes's Demon, the scientist with our brains in his vat, the suggestion that all our apparent experience might be a dream—as just so many ways of making vivid the Lockean predicament of our enclosure within a cognitive locality that allegedly excludes the local material environment. This is not, however, the role of such scenarios in the most important form of perceptual scepticism—the Cartesian scepticism developed on the model outlined in the Appendix—in whose service they are characteristically invoked. Their role is as reminders not of our alleged Lockean predicament but as notional exemplars of the possibility of subjectively perfect counterfeiture of perceptual activity, even when it is direct-realistically construed. McDowell (this volume: section 3) writes that

Perhaps most people will find it obvious that reinstating the sheer possibility of directly taking in objective reality in perception would undermine a scepticism based on claiming that perceptual experience can never amount to that.

That is indeed obvious, and—contrary to McDowell's parenthetical remark immediately following—I am no exception to the generalization. The point remains, simply, that perhaps the epitome of sceptical argument—elicited from Descartes's text, with minor variations, by each of Barry Stroud (1984: ch. 1), James Pryor (2000), and myself (1991)—involves no reliance whatever on the alleged senselessness, or impossibility, of direct perceptual acquaintance with the material world. The Disjunctive Conception of experience, even if it can indeed be enforced by transcendental argumentation supplementary to that offered by McDowell, provides us with nothing to address mainstream perceptual scepticism in this mould.

4 DISJUNCTIVISM AND HUMEAN SCEPTICISM

To the main business of this essay. Whatever the situation with other sceptical paradoxes, does the I-II-III form of sceptical argument, when directed against perception as a source of warranted belief concerning the local material environment, perforce rely upon the Lockean conception of experience as the best but defeasible ground for local environmental claims, or can it, as I suggested in “Simple and Subtle”, be reconstructed in the context of the Disjunctive Conception?

Here in outline is the case made in “Simple and Subtle” for thinking that it can. Consider first the original trio:

I My experience is in all respects as if P (say: as if there is a hand in front of my face).

II P (There is a hand in front of my face).

III There is an external material world.

Here, proposition I is to be understood as recording the occurrence of a Lockean experience, offering at best a defeasible ground for proposition II. So the question arises, what if anything renders the inference from I to II, or the transition from the occurrence of the experience to acceptance of II, rational? The sceptical contention is that its rationality, or otherwise, depends upon the nature of one’s collateral information. Minimally, that had better include that there is a material world at all, and that sense-experience is a reliable indicator, for the most part, of those of its characteristics that it is able to represent. But if that is correct, then this information needs to be in place as stage-setting for the rationality of any particular inference, or transition from experience to belief, of the type schematized by I and II. The problem is then—according to the sceptical argument—that it is impossible to understand how one might *come by* the information that there is a material world at all, still less one broadly veridically represented in ordinary sense-experience, except on the basis of acquiring knowledge or grounds for the truth of propositions of type II. Since sense-experience itself represents our best shot at such knowledge, we find ourselves trapped in a justificatory circle. Treating any type I proposition, or experience recorded thereby, as sufficient grounds for the corresponding type II proposition presupposes a warranted acceptance of III; but the later acceptance could be warranted only on the basis of prior warrant for propositions of type II.¹⁰

My point in “Simple and Subtle” was that no obstacle to the formulation of essentially this paradox is posed by a shift to the Disjunctive Conception. We merely reformulate I disjunctively as

I: Either I am perceiving a hand in front of my face or I am in some kind of delusional state of that phenomenal character,

retaining II and III as before. So reformulated, the inference from I to II, if it is to be rationally justified, will depend—or so the sceptical argument will contend—on the

¹⁰ For detailed development of the paradox, see the references at footnote 8.

justifiability of discounting the right-hand, uncongenial disjunct. And if the uncongenial disjunct is to be justifiably discountable in general, it appears that we had better have evidence that delusional states are exceptional—that, for the most part, states of the relevant phenomenal character constitute genuine perceptual apprehension of features of the local material environment. Yet that, in effect, is just another way of articulating the presupposition that featured in the original, Lockean formulation of the argument: the presupposition that there is a material world and that our perceptual experience is a mostly reliable guide to those aspects of it which it seems fitted to disclose. True, the nature of the “disclosure” varies between the two conceptions of experience. But that doesn’t affect the dialectical functioning of the paradox.

Very well. So why does McDowell believe that, by adopting the Disjunctive Conception, we acquire the resources to avoid the paradox? In “Simple and Subtle”, after developing the above analogy, I summarized the situation as follows:

The key point is that the disjunctivist makes a tendentious assumption in supposing that to conceive of perceptual experience as a form of direct acquaintance with reality is automatically to eliminate the idea that in the justification of perceptual statements, there is any role for claims weaker than perceptual statements . . . In brief: whether our perceptual faculties engage the material world directly is one issue and whether the canonical justification for perceptual claims proceeds through a defeasible inferential base is another. One is, so far, at liberty to take a positive view of *both* issues . . . When we do, the I-II-III pattern re-emerges along [the indicated] lines.

The thought was, in brief, that to think of our perceptual faculties as providing, in the best case, means of direct cognitive awareness of the material environment is so far a commitment to no particular view about the justificational architecture of perceptual claims, or about how, in the best case, a claim to be perceiving the local environment—and thereby to have found out that P—is itself to be justified. To accept the Disjunctive Conception is to take on board the idea that, in the best case, one is capable of a direct awareness of the “layout of reality”. But it is—so I was suggesting—another question, what constitutes *justification* for a claim about those aspects of reality which, in the best case, one’s direct perceptual awareness is awareness of.

This is the crux of the dispute. McDowell needs to make a case that the I-II-III template does indeed misrepresent the justificational architecture of perceptual claims. That case must be that when disjunctivism is incorporated into the picture, we are no longer grounded, so to speak, at stage I, with mandate only for the disjunction and requiring collateral information to progress beyond it. Rather, we now start, in effect, at stage II. McDowell (this volume: section 5) accordingly writes that

The point of the disjunctive conception is that if one undergoes an experience that belongs on the “good” side of disjunction, that warrants one in believing—indeed presents one with an opportunity to know—that things are as the experience reveals them as being. When one’s perceptual faculties “engage the material world directly” . . . the result—a case of having an environmental state of affairs directly present to one in experience—*constitutes one’s being justified* [my emphasis] in making the associated perceptual claim. It is hard to see how any

other kind of justification could have a stronger claim to the title “canonical”. And this justification is *not* defeasible. If someone sees that P, it cannot fail to be the case that P. So if one accepts the disjunctive conception, one is *not* at liberty to go on supposing that “the canonical justification of perceptual claims proceeds through a defeasible inferential base”.

And later:

The point for now is that Wright is wrong to claim that the disjunctive conception leaves one free to think perceptual claims rest on defeasible inferential support.

McDowell here says exactly what he has to say. He has to say that there is a better kind of warrant for perceptual claims than that provided by the information lodged in proposition I of the reconfigured I-II-III argument, together with collateral grounds for the supplementary claims, including III, argued to be necessary if that information is to add up to a warrant for II. But the obstinate fact is that, for all his assertion to the contrary, nothing in the Disjunctive Conception per se enforces that view of the matter. Indeed it had better not. For there is an evident gap between direct awareness of a situation in virtue which P is true and the acquisition of warrant for the belief that P, even for one sufficiently conceptually *savvy* to ensure that the direct awareness presents as an appearance that P.

One consideration which opens the gap is very familiar in the externalist tradition (with which McDowell’s thinking in these passages is in effect, but unacknowledgedly, coincident). Driving in Barn Façade County, but in all innocence of the locally distinctive “layout of reality”, the one-in-a-thousand real barn confronting me at the turn of the road ahead may draw my perceptual attention. I am directly aware of the barn, its location, the colour of its roof, its approximate dimensions, and so on. Yet scores of writers have scrupled over the suggestion that, in the circumstances, I *know* that there is a barn up the road, of such-and-such approximate dimensions, and with such-and-such a colour of roof. Whether they are right to do so, and if so why, are issues. The circumstances are such as to encourage false judgements about the things of which I am directly aware, they are such that I am merely lucky to be right in this instance, they are such that my claim is defeasible by improved information. McDowell doubtless has proposals to make about the proper handling of such veteran examples. But by writing in his present contribution as if perceptual uptake were *tout court* sufficient for warranted belief, he writes as though the issues they raise do not exist.

Direct awareness of states of affairs that make P true is one thing; warranted belief that P, for one fully apprised of what it takes for P to be true, is something else. One plausible additional necessary condition for the latter, not—as it happens—motivated by barn-façade type examples, is that the belief that P not be held irresponsibly or irrationally. What it takes, however, for a belief to avoid those failings is here exactly the key issue. Suppose I knowingly participate in double-blind trials of a new hallucinogenic drug. Half the participants receive the drug, the other half an identical-looking and identical-tasting vitamin pill. The consciousness of those who ingest the drug will move—so it is predicted—seamlessly into a completely plausible, sustained, multi-sense hallucination, in which all the experiences “make sense” in

the context of the subjects' recollection of their previous waking experiences of the day prior to ingestion of the pill, and indeed of the general fabric of their lives. This will happen—if the drug performs as expected—within a few minutes of ingestion, and the hallucination will then be sustained for several hours.

I take the pill and then am invited to sit down. Within what seems a very few minutes, a nurse invites me to join the other trialists in a recreation area. I am, seemingly, led down a corridor and through a pair of swing doors into what appears to be a converted gymnasium, with sauna, swimming-pool, snooker table, and a pool-side restaurant area with a wide-screen television showing footage of an international cricket match that, when I see it, I seem to remember was due to be played today.

Suppose that all this is, as it happens, a perfectly normal, veridical waking experience. Nevertheless it would, in the circumstances, be unwarranted for me to believe any of the mundane propositions of whose truth-conditions I am perceiving the worldly satisfaction. I do not know that these propositions are true. I am not, in the circumstances, even justified in claiming that they are true (the trialists, remember, are knowingly divided 50–50 but blind between those who take the drug and those who take the vitamins).

Here, then—to spell it out—is the salient sceptical rejoinder to McDowell's notion that the sheer possibility of direct perceptual engagement with the world, and the consequent Disjunctive Conception of experience, suffices to pre-empt the development of the I-II-III paradox for perceptual knowledge of matters external to us. Even for one apprised of all relevant concepts, and disposed to believe the proposition in question, the perceptual apprehension of a state of affairs that makes it the case that P is one thing, and possession of warrant for taking it to be the case that P is another. The former, even when a subject's belief that P is based on that very apprehension, is insufficient for warrant if aspects of the subject's collateral information conspire to make the belief that P somehow irrational or irresponsible. We have just reviewed a case where, as it may seem, that would be so. In the circumstances of the example, my collateral information makes it as likely as not, for all I can tell, that my perceptual faculties are systematically but wildly malfunctioning; in these circumstances, my actual perceptual apprehension, as normal, of a whole range of local environmental states of affairs does *not* provide a warrant for the beliefs which my experiences may dispose me to form—indeed, if I am rational, it will inhibit those dispositions. It is, however, a key feature of the example that my collateral information stops short of justifying me in believing that my perceptual faculties are actually malfunctioning. Rather, it forces me to regard the question as open. And—the crucial point—the mere openness of the issue seems to be enough to neutralize the evidence my faculties provide. The point the example makes is not—what McDowell can take in his stride—that once one has all-things-considered sufficient grounds to doubt that one's perceptual faculties are working properly—that is, sufficient grounds to opine that they are not—one cannot warrantably believe what seems to be disclosed by one's senses. *Mandated agnosticism* seems to be sufficient to do the damage.

It merely remains to harness the last consideration to the I-II-III paradox. The sceptical contention will be that, in the light of the possibility in principle—granted by McDowell—that instances of the 'bad' member of the disjunction may counterfeit

instances of the ‘good’ disjunct to phenomenological perfection, one is in fact no better placed *in general*—at least as far as the evidence of one’s own subjectivity is concerned—than the hypothetical subject of the imaginary double-blind trial. The hypothetical subject has balanced evidence for and against the possibility that he is out of perceptual touch with the real world. But a special case of balanced evidence is the situation of *no evidence* either way. And that, the sceptical claim will be, is exactly our predicament in general. For it is undeniable that if the alternatives are to suppose that my current experiences are elements of veridical perceptual activity and to suppose that they are a marvellously convincing counterfeit, then the subjective quality of the experiences itself—what it is like to undergo them—can indeed provide no rational motive for either view. *If*, therefore, there is reason for a view, it must originate elsewhere—and the suggestion of the final section of “Simple and Subtle” was that it must be found in our grounds, if any, for our ordinary picture of the material world: one broadly accessible to us via our perceptual faculties, and, most of the time, successfully accessed thereby. But the crucial point is that, once it is accepted that mandated agnosticism about the proper functioning of our perceptual faculties defeats the warrant supplied even by veridical experience as conceived by Disjunctivism, we owe an argument that we are in a better position than one of mandated agnosticism if we are to lay claim to the warrants which experience potentially provides. This is where the dialectic goes—some distance past the point where McDowell would like it to stop.

It is for this reason that the I-II-III paradox, originally composed for the Lockean setting, transfers smoothly to the Disjunctivist context. It cannot be finessed by McDowell’s simple stratagem of treating sheer apprehension as the “canonical warrant” for perceptual claims. Of course, if direct realism is true, sheer perceptual apprehension is a *component* of the very best kind of warranting situation for such claims that it is possible to enter into. But there are other necessary conditions that need to be met before apprehension can underwrite warranted belief. The sceptical argument is a case that they are unmet. It must be addressed differently.

5 TAKING SCEPTICISM SERIOUSLY

McDowell, it seems, is provoked by my attribution to him of “an official refusal to take scepticism seriously”. He does, apparently, believe that the Disjunctive Conception provides a rationally principled way of finessing sceptical doubt. What, he explains, he does not take seriously is the project of attempting to prove the propositions which scepticism brings into doubt—the forlorn project, apparently embarked on by Moore,¹¹ of trying to prove the existence of the external world (or, I suppose, that of other minds, or the reality of the long past, or the inductive amenability of Nature) by purely philosophical means. I am pleased to withdraw any imputation to McDowell of a cavalier dismissal of a serious philosophical problem and to welcome him to the

¹¹ It is, of course, debatable whether Moore really had this project. At footnote 8 of his contribution to this volume, McDowell flags a hint of a misgiving about it. See also Coliva (forthcoming).

company of those who would like to do something about it. Let me close, however, by emphasizing why the project of making a case for propositions that scepticism brings into doubt must indeed be part and parcel of any satisfactory response to it.

Lest that seem too shocking, let me immediately stress that I do not hereby mean to allow that Moore had the right project. He did not. But the reason why he did not is because the conclusion of his "Proof"—that there is an external material world—is *not*, in any relevant sense, a proposition that scepticism brings into doubt. One reason why McDowell has been inclined to underestimate scepticism is because he persistently fails to see any daylight between the forlorn projects—the Quixotic pursuit of the resources to establish a priori that there is matter, many minds, a substantial past, etc.—and what he thinks *is* possible, viz. the provision of alternatives to various of the preconceptions which sceptical paradoxes exploit, thereby showing that those preconceptions are not mandatory. Yet there is daylight, and the project we really need is in fact neither of these but falls in the daylight area.

Distinguish *doubt* and rationally enforced *agnosticism*. The thrust of material world scepticism is indeed that we do not know or even have sufficient evidence to believe that there is an external material world at all and that there is, moreover, no escaping from this predicament. The conclusion that—it is alleged—we should draw is that no attitude other than agnosticism on the matter is justified (even if it is an attitude we cannot actually psychologically sustain). But if 'doubting' is taken to involve believing that not, rather than merely not believing, then scepticism does *not* bring the existence of the material world into doubt. What it brings into doubt are *epistemic* claims—for instance, that we know, or have sufficient reason to believe, that there is an external material world. *These* doubts are ones that any satisfactory response to scepticism *must* remove. Somehow or other, considerations have to be marshalled to show that we are at least in a position rationally to claim the kind of knowledge of material reality with which we routinely credit ourselves. Modern epistemological scepticism—that is, the sceptical tradition in epistemology of the last 400 years—is preoccupied not with the reality of the subject matter of our large regions of knowledge—matter, other minds, the past—but with the credentials of our epistemic claims themselves. Cartesian doubt is already a *second-order* doubt—a doubt about the extent of the knowledge we can rationally lay claim to. When that is the orientation of our inquiry, it is obvious from the start that merely canvassing possibilities of direct cognitive acquaintance is unlikely to result in the placing of any very powerful philosophical levers. The nagging voice of intellectual conscience will rejoin: "Well yes, but even if that—a mode of direct acquaintance—is how we *have* to conceive of our experience, in the best case, what reason is there to take it that the best case is common, or even that it ever occurs at all?" We could have run the I-II-III argument, reconfigured for the disjunctive setting, in second-order terms *ab initio*. It would have sufficed merely to prefix each line with the operator: *I am in position rationally to claim knowledge that*. Thus:

- I I am in position rationally to claim knowledge that either I am perceiving a hand in front of my face or I am in some kind of delusional state of that phenomenal character.

- II I am in position rationally to claim knowledge that there is a hand in front of my face.
- III I am in position rationally to claim knowledge that there is an external material world.

No doubt I is true—subjectivity alone puts one in a position rationally to claim knowledge of the original disjunction. The question for McDowell is: what if anything puts me in a position rationally to claim knowledge that here is a hand—what makes II true? His proposal in the first-order case was that the actual perceptual apprehension of the hand would be knowledge-constituting. I explained above how scepticism can question that. But even if one repudiates that explanation, there is surely no plausibility whatever in the thought that the mere perceptual apprehension of the hand in front of my face is all it takes to put me in a position *rationally to claim* knowledge that it is there. For to lay claim to the latter knowledge, I have to lay claim to the relevant episode of awareness as one of genuine sensory experience. The mere fact of perceptual apprehension, if that is what it is, cannot *per se* make it rational to claim that that is what it is; and the rational defensibility of the belief that knowledge is what one has depends, in the circumstances, on the rationality of the latter claim.

It goes with McDowell's tendency to prefer to think of scepticism as a kind of anxious preoccupation,¹² rather than a family of paradoxes, that he tends to see it as something that needs, in the best later-Wittgensteinian tradition, diagnosis and therapy. I think this orientation tends to skew the philosophical discussion. But for what it is worth, if diagnosis is the objective, then it needs to be recognized that the basic, troubling thought in the vicinity is not that dreaming, hallucination, and veridical experience are all states of the same kind, distinguished only by their causes. It is, rather, exactly what Descartes said: that there are no "conclusive indications" by which to tell these states apart, and that this imperils our right to claim what we take to be our normal, commonplace cognitive achievements which depend on it being the 'good', rather than the alternative, 'bad' kinds of state that we normally occupy. It is the supposedly possible phenomenology of subjective indistinguishability, rather than a supposed "highest common factor", that is at the base of the problem. Once the root concern is thus properly identified, it should be obvious that the Disjunctive Conception has no materials to address it.

6 APPENDIX: THE CARTESIAN DREAMING PARADOX

Read 'Dxt' as x is undergoing a fully lucid, coherent dream at t at t, and 'WxtP' as: x is warranted in believing P at t, where such warrant is taken as requiring possession of all-things-considered sufficient evidence for P at t.

The premises for the paradox are then two. For fixed x and t (think in terms of 'you' and 'now'), premise 1 is

¹² Of course, no one is actually *anxious* about any of this!

$\sim W_{xt}[\sim D_{xt}]$

—you do not right now have all-things-considered sufficient evidence for the claim that you are not right now undergoing a fully lucid, coherent dream.

Premise 2 is

$W_{xt}[W_{xt}P \rightarrow \sim D_{xt}]$

—you do right now have warrant (that is, all-things-considered sufficient evidence) for the claim that if you right now have warrant for the claim that P, then you are not right now dreaming.

The classic Cartesian motive for premise 1 is simply the possible subjective indistinguishability of waking life and dreams. A subtler motive is provided by the following train of thought. That you are right now not dreaming is an empirical claim. So all-things-considered sufficient evidence for it has to be empirical evidence. To get empirical evidence you have to carry out some empirical procedure. It is plausible that the evidence produced by such a procedure may not rationally be regarded as any stronger than one's independent evidence that the procedure was carried out competently and with due diligence¹³—and hence that it was carried out at all. Dreaming excludes the genuine execution of empirical procedures. Hence the strength of warrant generated by executing such a procedure is rationally limited by the strength of one's independent warrant that one did not dream its execution! Hence no such procedure can generate a warrant that one is not (fully lucidly and coherently) dreaming.

A motivation for premise 2 may be given as follows. First, restrict attention to any P for which x's situation and state of information at t is such that in order to have sufficient evidence for P, she will have to *perceive* (so, for example, the range of P comprises all propositions about your local visible environment which you have never considered before). Then reflect that perceiving is *conceptually excluded* by dreaming. To be sure, in a dream you may seem to perceive things which are actually true—that it is, for instance, raining outside. But in order to count as perceiving that it is raining outside, you have to meet counterfactual constraints of sensitivity to the rain which in a dream you necessarily fail.¹⁴ It follows that $W_{xt}P \rightarrow \sim D_{xt}$ is true. But you have just run through the foregoing reasoning at t ('now'). So you right now have sufficient evidence for the conditional. And that is just what premise 2 says.

The proof-theory for the operator 'W' need involve no more than *closure* across (known) entailment—there is warrant for things which follow from what is warranted—and (as an optional extra) *iterativity*—if one is warranted in believing something, that one is so is itself something which one is warranted in believing.

¹³ This is the *Proper Execution Principle* of Wright (1991).

¹⁴ If, in a state of apparent sleep, you turn out to satisfy those constraints—however exactly they should be specified—we will say that you were not really asleep—or at any rate, that even if in a sleep-like state, you were not genuinely dreaming but rather engaged in some form of unusual perception compatible with sleep. Dreaming proper necessarily involves *disconnection*.

The paradox is then very immediate:

Assume $W_{xt}W_{xt}P$

Then $W_{xt}\sim D_{xt}$ —from premise 2, logic and Closure

So $\sim W_{xt}W_{xt}P$ —from premise 1 and *reductio*

That is as much as to say that no one ever is warranted in thinking that he has warrant for a proposition which it would, at their time and informational circumstances, take perception to justify. In effect, claims to be perceiving that P are never warranted if, in context, P could only be warranted by perception. That is a bad result, though it takes an appeal to iterativity to convert it into the properly sceptical conclusion,

$\sim W_{xt}P$

—one never has warrant for any proposition which it would, in one's context, take perception to justify, so that, in effect, perception is never a source of warrant.

This form of paradox will run against any cognitive faculty whose operations can in principle be mimicked by a kind of subjective state whose 'outputs' would be evidentially worthless.

REFERENCES

- Coliva, A. (forthcoming) 'The Paradox of Moore's "Proof of an External World"', *The Philosophical Quarterly*; online version available at <<http://www.blackwell-synergy.com/doi/abs/1001111/j.1467-92132007.513x>>.
- Descartes (1968) 'First Meditation', in E. Haldane and G. Ross (ed.) *The Philosophical Works of Descartes*, vol. 1, Cambridge: Cambridge University Press.
- McDowell, J. (this volume) 'The Disjunctive Conception of Experience as Material for a Transcendental Argument'.
- Pryor, J. (2000) 'The Skeptic and the Dogmatist', *Noûs*, 34: 517–49.
- Stroud, B. (1968) 'Transcendental Arguments', *Journal of Philosophy*, 65: 241–56. Reprinted in his *Understanding Human Knowledge*, Oxford: Oxford University Press, 2000.
- (1984) *The Significance of Philosophical Scepticism*, Oxford: Oxford University Press.
- Wright, C (1991) 'Scepticism and Dreaming: Imploding the Demon', *Mind*, 100: 87–116.
- (2002) '(Anti)-Sceptics, Simple and Subtle: G. E. Moore and J. McDowell', *Philosophy and Phenomenological Research*, 65: 330–48.
- (2004) 'On Epistemic Entitlement', *Proceedings of the Aristotelian Society*, Supplementary Volume 78: 167–212.

Index

- abilities 198–200, 211, 213
 - perceptual-discriminative, *see* discrimination
 - recognition 10, 332–7, 339–47, 387
- access, *see* direct access/availability
- accessibility of reasons, *see* reasons
- action 18–22, 85–6, 231–42, 244–61 *n*; *see also* bodily movement; disjunctivism; trying
 - basic 231–42, 234 *fn.* 8
 - rational explanation of 20–1, 244–61, 267–78
 - reasons for, *see* reasons
- animals 50, 96, 107, 154–5, 159, 175, 210–16, 222
- appearances 4, 7, 8 11–12, 18–19, 35, 38–9, 42–3, 45, 64–6, 80 *fn.* 39, 117, 138, 170, 380–2, 391 *fn.* 4, 393, 398, *see* disjunctivism
- Alston, William 64 *fn.* 14, 91, *fn.* 58, 92, 117 *fn.* 3, 143, 287 *fn.* 8, 308
- Alvarez, Maria 231, 237 *fn.* 12, 242
- Aristotle 96, 110, 200, 369–70
- Armstrong, David 15, 23, 63 *fn.* 11, 92, 150, 153–6, 162–3, 166
- Austin, J. L. 37–8, 56, 58, 63–5, 68–9, 91 *fn.* 58, 92, 265

- Bach, Kent 231, 237, 242
- Baker, G. P. 177, 179
- Baldwin, Thomas 283 *fn.* 1, 308
- Bennett, Jonathan 82 *fn.* 45, 92, 233, 242
- Berkeley, George 169, 179, 233
- Bishop, John 231, 242
- Bloor, David 268 *fn.* 1, 278
- bodily movement 19–22, 191, 230
- Brandom, Robert 288 *fn.* 9, 290 *fn.* 16, 297 *fn.* 23, 308
- Brewer, Bill *ix*, 13, 18, 26, 89, 92, 168–80, Broad, C. D. 148, 166
- Byrne, Alex *ix*, 4, 10, 18, 21, 25, 57–94, 115 *fn.* 2, 118 *fn.* 4

- Campbell, John 78 *fn.* 37, 92, 207, 233, 330 *fn.* 1, 331 *fn.* 2, 347
- Candlish, Stewart 242
- Cartesianism 1, 11–13, 21–2, 70–1, 87, 91, 267–8, 390–1, 394–5, 401–3
- causal theory of perception 1, 2–4, 8, 11, 12, 19, 39–40, 53, 54–5, 72, 95–109, 148, 391, 402

- Child, William 65 *fn.* 15, 72 *fn.* 26, 92, 96 *fn.* 4, 110, 181, 203, 227 *fn.* 1, 242, 331 *fn.* 3, 332, 347
- Chisholm, Roderick 117 *fn.* 3, 143, 292 *fn.* 18, 308
- Cohen, Stewart 286 *fn.* 6, 288 *fn.* 10, 308
- Comesãna, Juan 290 *fn.* 16, 308
- common factor 7, 12, 14–16, 17, 20, 44, 48–51, 59–60, 61, 63–4, 66–8, 69, 71, 72–3, 74, 78 *fn.* 35, 79, 81, 84 *fn.* 49, 85–6, 87, 88–90, 99, 101, 104 *fn.* 18, 109, 144–7, 153–6, 166, 170, 181, 183, 202, 205–8, 212, 217, 218, 228, 262–5, 269 *fn.* 3, 294–5, 348, 362, 379–87, 391, 391 *fn.* 4, 401–2
- concepts 44, 48, 54, 96, 106, 155, 159–60, 175–8, 199–200, 393, 399
 - conceptual and non-conceptual content 91 *fn.* 58, 175–6, 348–51, 362–3
- Conee, Earl 287 *fn.* 8, 308
- consciousness 5, 7 *fn.* 8, 10, 14, 64, 66, 68, 86, 97–103, 107–110, 112–42, 148, 183, 189, 191, 196–8, 203, 211–14, 222, 382, 391, 393, 398; *see also* phenomenal character; self-consciousness
- content 2, 9, 13, 14–18, 68, 71, 75, 83 *fn.* 47, 84, 86 *fn.* 52, 87 89–91, 95–6, 101–8, 117–19, 122, 123 *fn.* 8, 138, 149, 154, 158–9, 168–77, 183, 188–9, 197, 352–3 *fn.* 3, 388, 392–3; *see also* intentional theory
- singular content, *see* object-dependent content
- contextualism 286 *fn.* 6, 322–8
- Cook Wilson, John 345 *fn.* 19
- Crane, Tim 65 *fn.* 15, 80 *fn.* 39, 92, 99 *fn.* 8, 110, 118 *fn.* 4, 143, 263, 265, 278, 311, 329

- Dancy, Jonathan *ix*, 20, 21, 23, 28, 56, 69, 92, 242, 245, 249 *fn.* 7, 260, 260 *fn.* 26, 262–79
- Davidson, Donald 231 *fn.* 5, 246 *fn.* 3, 260, 308, 366–7, 374
- Davis, Wayne 231, 242
- defeasibility 4–7, 9–10, 19, 64, 290–302, 383–5, 394, 396, 397, 398
- Dennett, Daniel 159, 166, 358 *fn.* 8, 366–7, 374
- DeRose, Keith 286 *fn.* 6, 308

- Descartes, Rene 11, 105 fn. 19, 185, 378–9, 382, 383, 390, 390 fn. 1, 391, 395, 402, 404; *see also* Cartesianism
 determinate/determinable 266
 direct access/availability 14–15, 37–40, 62–3, 64 fn. 13, 82, 104 fn. 18, 105–8, 168–79, 218, 221, 378–81, 383–4, 391–401; *see also* direct realism
 direct realism 13, 38, 40, 82, 168–79, 348–75, 400; *see also* naïve realism; direct access
 discrimination, *see* indiscriminability/indistinguishability; perceptual-discriminative abilities
 disjunctivism:
 epistemological 4–13, 17–19, 21–2, 35, 36, 57, 65–8, 91, 290–307, 311–29, 330–2, 337–9, 391 fn. 4, about acting for reasons 20–2, 244–61, 267–78
 perception-illusion disjunctions 16, 18, 40–2, 45, 60–1
 about bodily movements 19–20, 22
 disjunctive conception of appearances 4, 7, 8, 11, 12, 18–9, 252–58, 351 fn. 2, 363–5, 380–2, 391 fn. 4
 disjunctive theory of action 1–2, 18–21, 230–42
 experiential 2–4, 8, 13, 17, 19–20, 21, 35, 37, 57, 63, 95–111, 112–43, 144–66, 168–79, 181–203, 205–23, 227, 332–3, 376–89, 390–404
 extreme 182–202, metaphysical 21–2, 35, 57, 67, 68–91, 95–111, 112–43
 phenomenal 13–18, 21, 141–66, 168–79, 205–23
 pure 139–42, 144
 dreams 149–50, 185–92, 196 fn. 19, 203, 379, 390–5, 402–3, 404
 Dretske, Fred 58 fn. 1, 63 fn. 11, 82, 88 fn. 53, 92, 303, 308, 321, 329, 387
 Dummett, Michael 4

 epistemic externalism 287–90, 292–5, 302, 339–42, 398–99
 epistemic internalism 287–300, 313–15
 epistemology 6, 10, 36, 40–1, 66–7, 69, 70 fn. 24, 78, 83, 84, 91, 120–43, 156, 183, 186, 201, 205–23, 208, 250–61, 266, 271–8, 283–404, 381, 401; *see also* disjunctivism, epistemological
 Evans, Gareth 58 fn. 1, 89 fn. 55, 92, 348, 362, 374
 events 70, 79, 83, 84, 85, 183, 202, 230–42
 intrinsic events 230–41
 mere events 230–41
 experience:
 content view of experience 2 fn. 1, 5, 8–9, 12–13, 14–18, 68, 71, 87, 89–91, 95–6, 101–8, 122, 169–79, 183, 388, 392–3
 object view of experience 8–9, 12–13, 14, 63–5, 66, 79–81, 116–17, 171–9
 expression 11, 18, 154, 164
 externalism, *see* epistemic externalism
 extrinsic 8, 13, 103, 133

 facts 5–13, 19–20, 22, 69, 82 fn. 45, 103, 114–19, 123, 129, 131, 141, 169–78, 194, 198–203, 205–6, 214–23, 255–8, 372–3, 378–9, 380, 387, 394–5
 fallibility and infallibility 6–7, 11–12, 22, 76–7, 106, 197 fn. 22, 381, 387
 Farkas, Katalin 100 fn. 9, 101 fn. 12, 110, 210 fn. 9, 223
 Feldman, Richard 287 fn. 8
 Fish, William ix, 18, 27, 144–67, 206–23
 Fodor, Jerry 172 fn. 8, 179
 Foley, Richard 288 fn. 9, 308
 Frege, Gottlob 348, 362; *see also* Fregean thoughts
 Fregean thoughts 11–12; *see also* Frege, Gottlob
 Fumerton, Richard 290 fn. 15, 308
 fundamental:
 disjunctivist commitment 18
 kind 16–17, 61 fn. 9, 79–81, 83, 84, 100 fn. 10, 107 fn. 22, 186, 202, 207–8, 228
 nature 3, 39, 116, 154, 168, 170, 171

 Geach, P. T. 216, 223
 Gettier cases 251, 276–7, 292 fn. 18, 339
 Gibbons, John 251 fn. 10, 260
 Ginet, Carl 319
 Goldman, Alvin 319, 329, 335 fn. 6, 341, 347
 Goodman, Nelson 229, 242
 Graff, Delia 71 fn. 25, 92, 133 fn. 16, 143
 Greco, John 290 fn. 16, 297 fn. 23, 341 fn. 13, 308, 347
 Grice, H. P. 2, 23, 53, 90, 177–8, 179

 Hacker, P. M. S. 177, 179
 Haddock, Adrian ix, 1–24, 202 fn. 25, hallucination 13–18, 37, 39, 45, 58–61, 61–3, 64 fn. 14, 67 fn. 20, 68–72, 73, 74–5 fn. 28, 80, 83–6, 88–91, 95, 98–100, 109 fn. 23, 113–14, 118, 125, 131, 137–8, 144–66, 168, 170, 173–4,

- 178–9, 181–203, 205–23, 227–30, 244, 252–8, 262–5, 277–8, 291–7, 311–22, 323–8, 390–1, 393, 398–9, 402
- Hamlyn, David 230, 242
- Hawthorne, John 58 fn. 2, 75 fn. 31, 89 fn. 55, 90 fn. 57, 92, 93, 118 fn. 4, 127 fn. 12, 143
- Heller, Mark 321, 329
- Hellie, Benj 71 fn. 25, 92, 131 fn. 15, 143
- highest common factor, *see* common factor
- Hinton, J. M. 11, 3, 4, 23, 35–56, 59–61, 62, 63–5, 66–8, 68–9, 72–3, 81, 83 fn. 46, 88, 91, 92, 96 fn. 4, 99 fn. 8, 101 fn. 11, 110, 116, 181, 203, 227 fn. 1, 242, 311–12, 329, 381, 388
- Horgan, Terence 233, 242
- Hornsby, Jennifer ix, 20–1, 23, 28, 230, 232, 242, 244–61, 271–7
- Hurley, Susan 354 fn. 5, 374
- Hyman, John 231, 237, 242, 247 fn. 6, 260, 276, 279
- illusion 6, 7 fn. 8, 14, 18, 40–8, 54, 58–63, 64–5 68 fn. 21, 69–70 fn. 23, 73, 90–1, 98, 113–15, 168–79, 181, 205, 216 fn. 16, 253 fn. 13, 390 fn. 1
- indefeasible, *see* defeasibility
- indiscriminability/indistinguishability 6, 15, 16, 18, 37, 39, 58–61, 62, 64–6, 70–1, 73–7, 84, 85 fn. 50, 99 fn. 8, 100, 102, 118–42, 144–66, 170, 178, 181–203, 206–23, 228, 264–5, 272, 277–8, 290–6, 302–7, 311–15, 318–28, 381, 382, 390, 402, 403
- infallibility, *see* fallibility
- infants 26, 50, 98, 101, 127–9, 159, 175–6,
- inferential justification, *see* justification
- inner and outer 5, 10, 11–12, 18–19, 21–2, 37, 44, 87, 151, 173, 351 fn. 2, 393, 394
- intentional theory 13–18, 68 fn. 21, 71–2, 75, 85, 89–90, 117–18, 352–3 fn. 3, 357 fn. 7; *see also* content
- strong intentionalism 15–16, 18
- weak intentionalism 15–16
- intentionality, *see* content
- internalism, *see* epistemic internalism
- intrinsic 3, 8, 12–13, 15, 16, 17, 19–20, 21, 97–104, 108–9, 133, 145–7, 156, 177, 184, 200–1, 214; *see also* events
- introspection, 13, 15, 16, 18, 58, 70–1 74–5, 82, 83 fn. 47, 85, 100 fn. 9, 121, 123, 125, 127, 140–2, 150–1, 153, 154, 170, 173–4, 178, 182, 205–13, 218–23; *see also* knowledge,
- introspective; reflection;
- self-consciousness
- intuitive criterion of difference 11
- irrationality, *see* rationality
- Jackson, Frank 63 fn. 11, 88 fn. 53, 92
- James, William 240–2
- Jastrow, Joseph 175, 179
- Johnson, Mark 58 fn. 2, 64 fn. 14, 89 fn. 55, 91 fn. 58, 92, 315–8, 322–6, 328–9
- Johnson, W. E. 266, 279
- justification 6, 7 fn. 9, 16, 19, 128, 283–307, 313–15, 342–3, 383–6, 396–8, 399, 401, 404; *see also* warrant
- Kant, Immanuel 23, 376, 377, 380, 382, 388, 389, 391–2
- Kornblith, Hilary 287 fn. 8, 308
- Kovakovich, Karson 58 fn. 2, 75 fn. 31, 89 fn. 55, 90 fn. 57, 92, 118 fn. 4, 127 fn. 12, 143
- knowledge 1, 2, 4–14, 18–20, 21–2, 35, 41, 44, 64, 70, 74 fn. 28, 77, 78 fn. 35, 81 fn. 41, 88, 100 fn. 9, 104 fn. 18, 120–31, 134–40, 146, 150, 163, 181–2, 210–11, 213–14, 218–23, 250–1, 252 fn. 11, 253–60, 266, 271–8, 283–9, 302–7, 343–7, 378–89, 396, 401–2
- introspective knowledge 10, 127–8, 131 134–8, 150, 193, 198; *see also* introspection; reflection
- perceptual knowledge 4–5, 7, 11, 19, 64, 252–4, 290–302, 311–29, 331–43, 385, 394–5, 399
- Langsam, Harrold 65 fn. 15, 81 fn. 43, 92, 117 fn. 3, 143
- Lehrer, Keith 288 fn. 10, 308
- Lewis, David 169 fn. 4, 171 fn. 6, 179, 199, 204, 286 fn. 6, 288 fn. 9, 308
- Locke, Don 60, 64 fn. 14, 93
- Locke, John 150, 167, 169, 174, 179
- Logue, Heather ix, 4, 10, 18, 21, 25, 57–94, 115 fn. 2, 118 fn. 4,
- Lowe, E. J. 4, 25–6, 95–111, 148, 167, 242
- Lycan, William G. 150, 167
- Lyons, William 338 fn. 9, 347
- Macarthur, David 308
- McCann, Hugh 234–5, 243
- McDowell, John ix, 1, 2–13, 16, 18–19, 21–2, 23, 24, 30, 36, 56, 60–2, 64–8, 72, 80 fn. 39, 89 fn. 55, 90 fn. 56, 91, 93, 96 fn. 4, 98 fn. 7, 99 fn. 8, 110, 116–21, 143, 170, 180, 181, 183, 204,

- McDowell, John ix, (*cont.*)
 207 fn. 3, 224, 227 fn. 1, 243, 244,
 252–4, 261, 273, 290–302, 305–9,
 312–5, 318, 323–9, 331 fn. 3, 340
 fn. 11, 343 fn. 17, 345 fn. 19, 347, 348,
 351 fn. 2, 362–3, 366, 368–9, 374,
 376–89, 390–404
- McKinsey problem 297–300
- Macpherson, Fiona x, 1–24, 239 fn. 13
- Martin, C. B. 199, 204
- Martin, M. G. F. 1, 13–18, 21, 23, 58 fn. 2,
 59 fn. 3, 4, 60–1, 64 fn. 14, 65
 fn. 15, 66 fn. 17, 69–90, 93, 97 fn. 6,
 99 fn. 8, 100 fn. 9, 102 fn. 14, 107
 fn. 22, 111, 116, 121 fn. 6, 123 fn. 8,
 125, 126, 127 fn. 12, 135 fn. 17, 135
 fn. 18, 136 fn. 20, 143, 156, 167, 173
 fn. 10, 180, 181–3, 186–7, 193, 195
 fn. 18, 198, 201–3, 205–24, 227 fn. 1,
 228–9, 243, 256 fn. 18, 257 fn. 20,
 261, 264–5, 277–9, 312, 329, 349 fn.
 1, 353 fn. 3, 358, 375
- Marx, Karl 235, 243
- Millar, Alan x, 4, 10, 11, 15, 24, 29, 70
 fn. 24, 93, 330–47
- Moore, G. E. 24, 283, 309, 383–6, 389,
 400–1, 404
- naïve realism 13, 16, 17, 36–40, 79–83,
 85–6, 91, 97 fn. 6, 99, 107–8, 161,
 222 fn. 18; *see also* direct realism
- narrow sense 29, 46, 54–5, 184 fn. 9, 185,
 236–42; *see also* wide sense
- negative characterisation or epistemic fact 43,
 46, 48, 69, 79, 83–4, 96, 109, 118,
 124, 146–66, 178, 208, 211, 218–23,
 182–203, 262
- neo-Mooreanism 268–307
- Neta, Ram x, 10, 11, 29, 298 fn. 25, 309,
 311–29
- Noë, Alva 148, 167
- non-conceptual content, *see* concepts,
 conceptual and non-conceptual content
- non-conjunctivism 265–6, 272
- non-inferential justification, *see* justification
- Nozick, Robert 321, 329
- object-dependent content 11, 12, 71, 89, 90
 fn. 57, 331–2, 348, 352–74
- other minds 4–5, 19, 64, 65–6, 378, 400–1
- outer, *see* inner and outer
- Parfit, Derek 247 fn. 6, 261
- Peacocke, Christopher 59, 93, 360 fn. 9
- Penfield, Wilder 148–9, 167
- perception 1–4, 13–18, 19, 21, 35–55,
 59–61, 65–6, 69, 72–5, 79–88, 91
 fn. 58, 95–108, 113–36, 144–66,
 168–79, 181–203, 205–17, 227–30,
 244, 252–8, 262–5, 277–8, 291–7,
 311–8, 322–8, 330–47, 348–65,
 379–81, 386, 391, 393, 395–6, 403
 fn. 14, 404
- perceptual access, *see* direct access/availability
- perceptual content, *see* content
- perceptual-discriminative abilities 337–9
- Philipps, Ian 175 fn. 14, 180
- Pietroski, Paul 232, 243
- Pitcher, G. 38–40, 56
- phenomenal character 1, 2, 5, 8, 13–18,
 21–2, 71, 76, 86–7, 97 fn. 6, 98, 100
 fn. 9, 102 fn. 14, 104, 107, 108, 120–2,
 133–42, 145–9, 149–62, 175–78,
 182, 194–5, 205–22, 229–30, 352–3
 fn. 3, 383 fn. 8, 381, 390, 396–7, 400,
 401, 402; *see also* consciousness
- phenomenology, *see* phenomenal character
- Plato 344
- practical reason, *see* reasons
- Price, H. H. 64 fn. 14, 81 fn. 43, 93, 243
- Pritchard, Duncan x, 10, 11, 29, 283–310
- propositional attitudes 14, 117
- Pryor, James 289 fn. 309
- Putnam, Hilary 68, 93, 98 fn. 7, 99 fn. 8,
 111, 121 fn. 7, 143, 207 fn. 3, 224, 312,
 329, 367 fn. 17, 375
- Pylyshyn, Zenon 89, 93
- qualia, *see* phenomenal character
- rationality 11, 20, 21–2, 127, 186, 217, 223,
 369 fn. 18, 377, 386, 396, 398–400,
 400–402, 403
- reasons 4–10, 20, 21–2, 66, 157, 163–4
 accessibility of 5, 10, 11, 287–300, 302–7,
 313–5
 for action 1, 20, 22, 244–61, 267–78
 for belief 4–5, 21–2, 290–302, 313–15
 disjunctivism about, *see* disjunctivism
 motivating 249–50
 normative 249–50
- recognition 10, 145, 162, 175, 176 fn. 16,
 189; *see also* abilities, recognitional
- reflection 73–82, 100 fn. 9, 106, 108,
 120–38, 142 fn. 25, 175, 182–203,
 211–12, 222 fn. 18, 287, 290, 296–9,
 305–6, 323, 327–8, 336, 340–1, 344,
 361, 376
- relational 16–17, 85–6, 97, 105, 110,
 116–42, 146, 156–7, 162–5, 170 fn. 5,
 171, 183, 194–6, 201, 205–7, 209,

- 223, 256–7 fn. 18, 269–70, 330–2, 337–9, 349, 391
 representation, *see* content; intentional theory
 representational theory, *see* content; intentional theory
 Robinson, Howard 59 fn. 4, 64 fn. 14, 65 fn. 15, 81 fn. 43, 83 fn. 48, 93, 148, 161, 167,
 Rosenthal, David 154, 167
 Ruben, David-Hillel x, 20, 27, 227–43
 Russell, Bertrand 11, 64, 81, 91, 93, 244, 348
 scepticism 1, 4, 50, 51, 64, 66–7, 83, 98, 104–6, 120, 149, 154, 229, 283–90, 295 fn. 20, 300–7, 350, 376–88, 391–404
 scientific explanation/knowledge/picture/theory 9, 52, 55, 96, 110 fn. 25, 259–60, 365–6
 Sedivy, Sonia x, 13, 30, 348–75
 Segal, Gabriel 89 fn. 55, 93,
 self-consciousness 10, 210, 378, 382, 387, 393; *see also* introspection; knowledge, introspective; reflection
 Sellars, Wilfred 23, 380–1,
 sense-datum theory 13–17, 37–8, 51, 59, 63–73, 77, 82, 83 fn. 47, 84, 85, 105, 138,
 Shoemaker, Sydney 151, 167
 Sider, Ted 171 fn. 6, 180
 Siegel, Susanna x, 18, 27, 74 fn. 28, 75 fn. 31, 32, 77 fn. 34, 89 fn. 55, 93, 123 fn. 8, 123 fn. 9, 127 fn. 12, 135 fn. 17, 143, 150, 167, 205–24
 singular content, *see* object-dependent content; content
 Smith, A. D. x, 18, 27, 45 fn. 9, 62 fn. 10, 63, 65 fn. 15, 81 fn. 43, 89 fn. 55, 93, 147–8, 167, 181–204
 Snowdon, Paul x, 1, 2–4, 7–8, 11–12, 13, 21, 22, 24, 26, 35–56, 60–1, 66, 72, 80, 81 fn. 40, 88, 94, 95 fn. 2, 96 fn. 4, 99 fn. 8, 111, 116, 140 fn. 24, 202 fn. 24, 227 fn. 1, 243, 312, 329, 331, 347, 381 fn. 5, 388
 Sosa, E. 213, 224, 309, 341, 347
 Soteriou, Matthew 80 fn. 39, 89 fn. 54, 94, 146 fn. 26, 176 fn. 17
 Spener, Maja 142 fn. 25, 143,
 Stern, Robert 378 fn. 2, 388
 Stoneham, T. 169 fn. 3, 180
 Stout, Roland 260 fn. 24, 261
 Stoutland, Fred 259 fn. 23, 261
 Strawson, P. F. 2, 24, 30, 53, 89, 94, 376–8, 382, 388
 Stroud, Barry 290 fn. 15, 298 fn. 24, 309, 376–8, 384, 385, 386 fn. 9, 387, 388, 392, 395, 404
 Sturgeon, Scott x, 18, 26, 37 fn. 4, 65 fn. 15, 74 fn. 28, 75 fn. 31, 94, 112–43, 144, 151, 167, 254 fn. 14, 257 fn. 20, 261, 331 fn. 3, 347
 Thau, Michael 63, 64 fn. 14, 65, 94
 Thornton, Tim 118 fn. 4, 143,
 transcendental 4, 8–9, 11, 89, 363–4 fn. 13, 376–88, 390–5, 404
 transparency of experience 11, 13–17, 82–3, 99, 108,
 Travis, Charles 89, 91 fn. 58, 94, 312, 329, 345 fn. 19, 347
 trying 232, 235; *see also* action
 Tye, Michael 71, 94, 159, 167, 352 fn. 3, 376
 Vendler, Z. 82, 94
 veridical and non-veridical experience 13–18, 37, 44, 58, 61, 65, 68, 69, 73–5, 77–8, 79–80, 83, 84–6, 88, 89–90, 99–100, 104, 106, 113–14, 119, 122, 123 fn. 8, 128–30, 133, 136, 144–66, 170, 181 fn. 1, 182–3, 201–2, 205–22, 262–5, 381, 396, 399–400, 402
 Von Wright, G. H. 230, 243
 warrant 1, 64, 66, 376, 378, 381, 382 fn. 7, 383–7, 392–404; *see also* justification
 wide sense 12, 184 fn. 9, 236–241; *see also* narrow sense
 Wiggins, David 356–7, 375
 Williams, Bernard 268 fn. 1
 Williams, Michael 301 fn. 27, 309
 Williamson, Timothy 57, 66, 67, 68, 70, 71 fn. 25, 74 fn. 28, 78, 81 fn. 41, 85–6, 94, 126–7, 129 fn. 14, 131–2, 143, 146–7, 150, 167, 172 fn. 9, 209 fn. 8, 210 fn. 9, 224, 257 fn. 19, 261, 265–6, 276, 279, 289 fn. 14, 295 fn. 20, 309, 340 fn. 11, 342 fn. 15, 344 fn. 18, 347
 Wittgenstein, Ludwig 4, 64, 65, 175, 180, 234–5, 286, 310, 367–73, 375, 402
 Wright, Crispin x, 7, 10, 24, 30–1, 284 fn. 2, 286 fn. 5, 289 fn. 13, 290 fn. 16, 297 fn. 23, 309–10, 314, 318, 326–9, 382 fn. 7, 383–7, 389
 Yablo, Stephen 86 fn. 51, 94
 Zagzebski, Linda 288 fn. 9, 310, 341 fn. 13, 347

